

Computer Vision 1

<https://tinyurl.com/dhcp7h3w>

Neighborhood Processing Image Filtering Fourier Transform

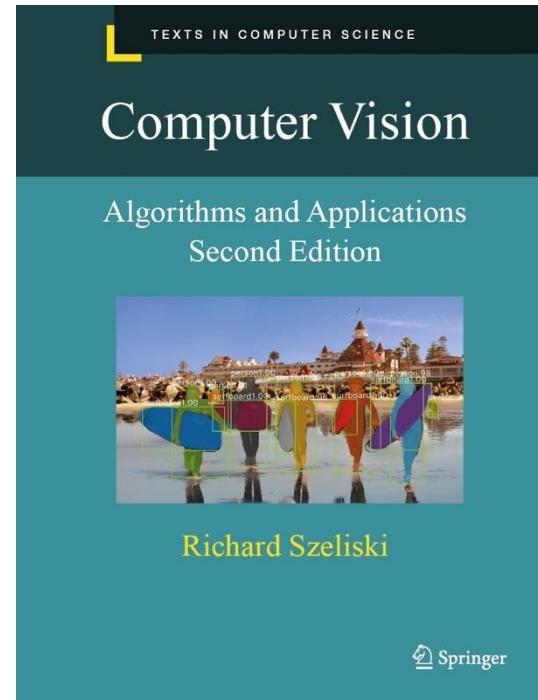
Dr. Martin Oswald, Dr. Dimitris Tzionas, Dr. Arun Mukundan,
[m.r.oswald, d.tzionas, a.mukundan]@uva.nl

Outline

- **Neighborhood image processing**
 - Connected Component Analysis
 - Mathematical Morphology (Dilation, Erosion, Opening, Closing)
- **Linear Image Filtering**
 - Correlation and Convolution
 - Frequently used Filters
 - Gaussian Filter
- **Nonlinear Image Filtering**
 - Median Filter
 - Bilateral Filter
- **Frequency Domain Analysis**
 - Fourier Transform
 - Sampling and Aliasing

Textbook

- 3.2 Linear filtering
- 3.3 More neighborhoos operators
- 3.4 Fourier transforms



[<https://szeliski.org/Book/>]

Outline

- **Neighborhood image processing**

- Connected Component Analysis
- Mathematical Morphology (Dilation, Erosion, Opening, Closing)

- **Linear Image Filtering**

- Correlation and Convolution
- Frequently used Filters
- Gaussian Filter

- **Nonlinear Image Filtering**

- Median Filter
- Bilateral Filter

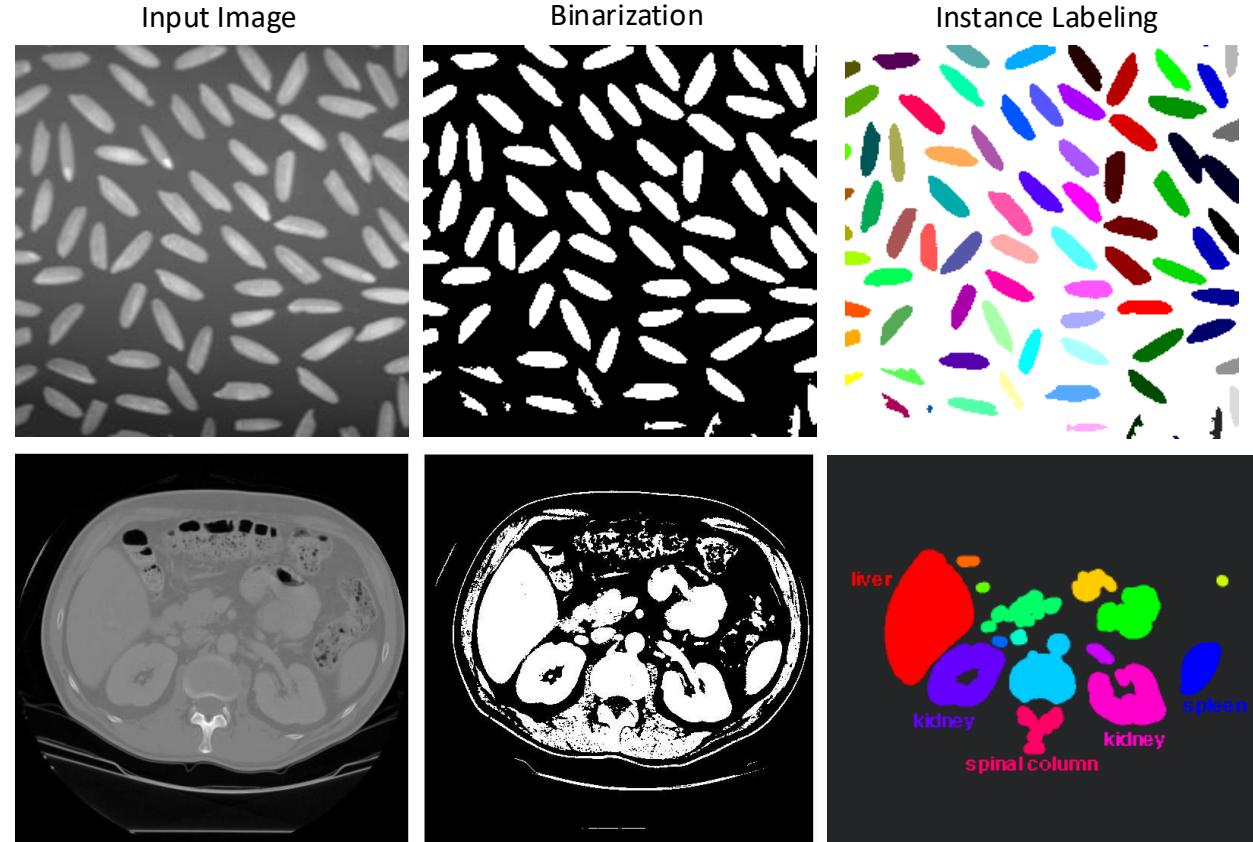
- **Frequency Domain Analysis**

- Fourier Transform
- Sampling and Aliasing

Simple Image Processing Examples

Instance Segmentation and Counting

- Rice grains
- Organs



Thresholding

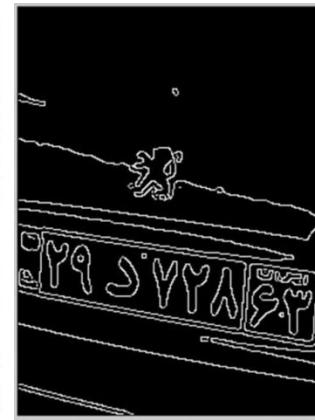
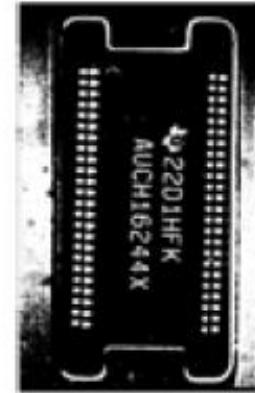
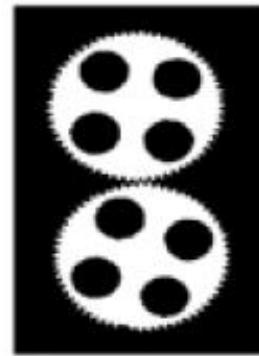
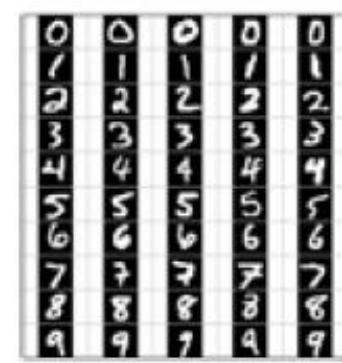
Binarize input signal:



Input Image

Binarized Image
$$B(i,j) = \begin{cases} 255 & \text{if } I(i,j) > \tau \\ 0 & \text{otherwise} \end{cases}$$
 Threshold

Why binary ?

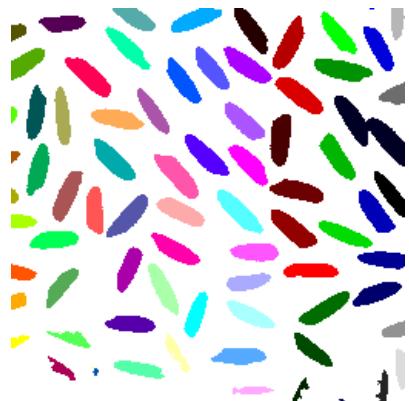
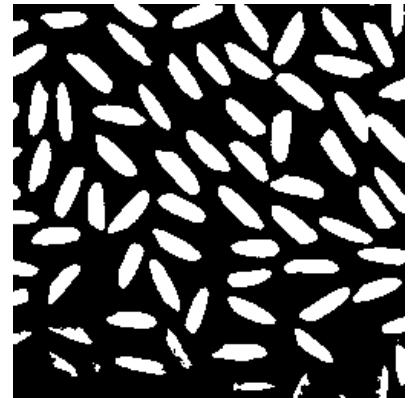


Outline

- Neighborhood image processing
 - Connected Component Analysis
 - Mathematical Morphology (Dilation, Erosion, Opening, Closing)
- Linear Image Filtering
 - Correlation and Convolution
 - Frequently used Filters
 - Gaussian Filter
- Nonlinear Image Filtering
 - Median Filter
 - Bilateral Filter
- Frequency Domain Analysis
 - Fourier Transform
 - Sampling and Aliasing

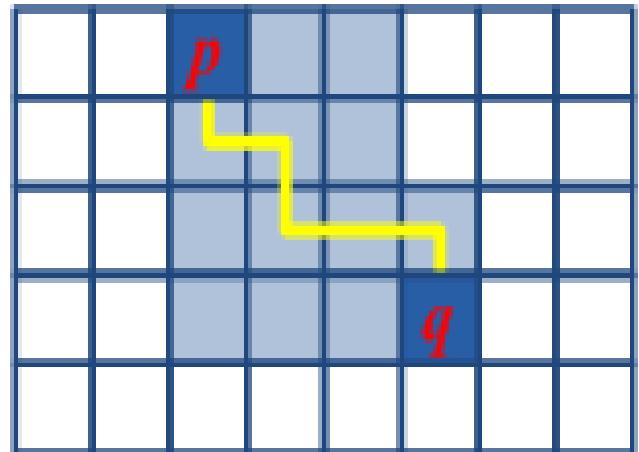
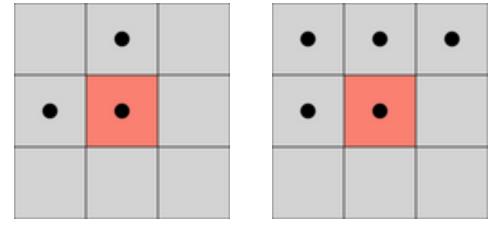
Connected Components Analysis

- Connected Components Analysis is a multiple-class labeling operation.
- It performs the change from pixels to regions.
- All pixels in the same region are given the same label.



Connected Component Operators

- Definition of Connected Component
 - Two pixels p and q belong to the same connected component C if there is a sequence of 1-pixels (p_0, p_1, \dots, p_n) , where
 - $p_0 = p$
 - $p_n = q$
 - $p_{i-1}, p_i : i = 1, \dots, n$ are neighbors



Region property computation

For each connected region, we can compute many of its properties, e.g.

- Areas
- Perimeter
- Centre of gravity
- circularity,
- Major axis
- Minor axis
- mean and standard deviation of radial distance
- bounding box
- extremal axis length from bounding box
- second order moments (row, column, mixed)
- ...

Region property computation

First Pass:

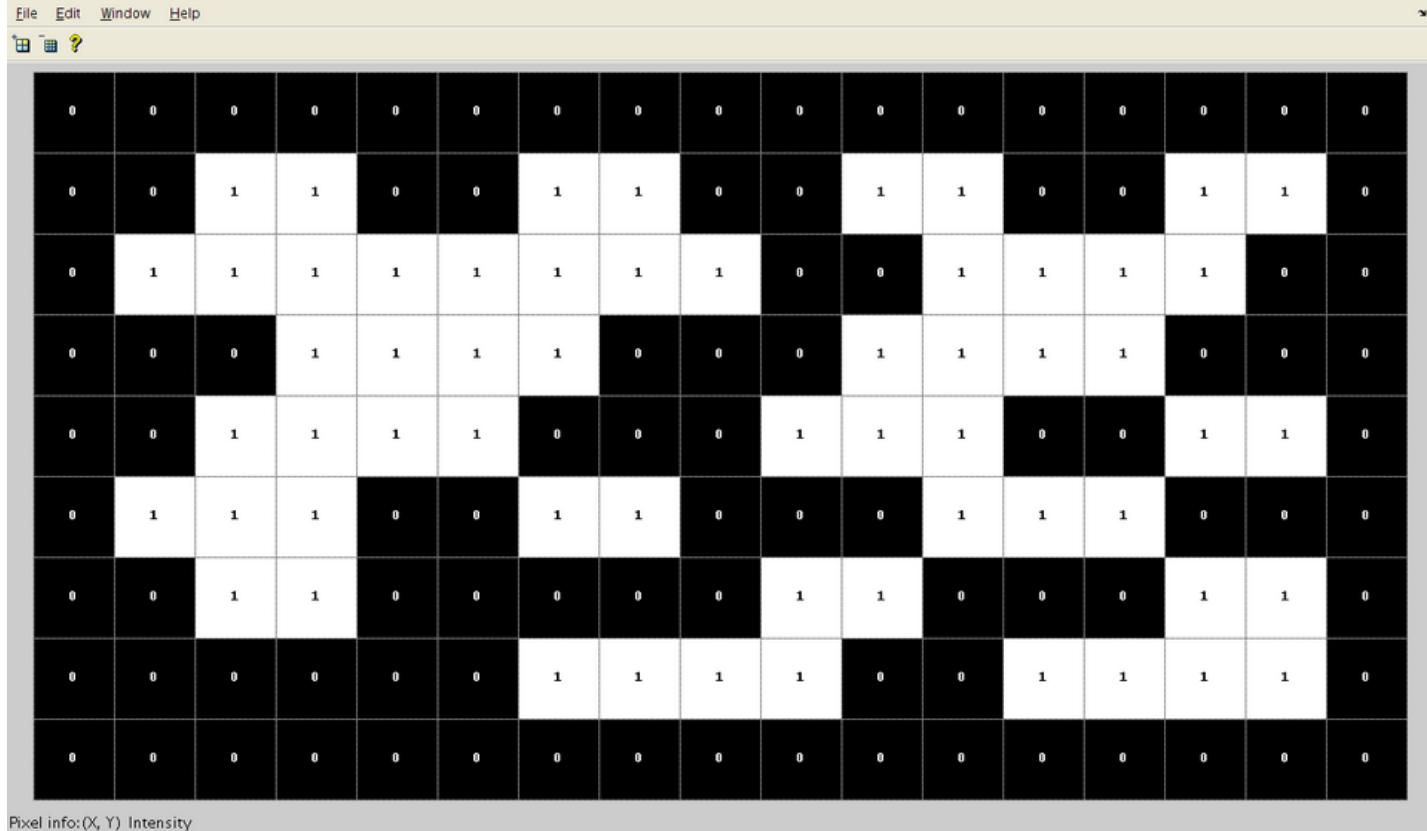
1. Iterate through each element of the data by column, then by row (Raster Scanning)
2. If the element is not the background
 1. Get the neighboring elements of the current element
 2. If there are no neighbors, uniquely label the current element and continue
 3. Otherwise, find the neighbor with the smallest label and assign it to the current element
 4. Store the equivalence between neighboring labels

Second Pass:

1. Iterate through each element of the data by column, then by row
2. If the element is not the background
 1. Relabel the element with the lowest equivalent label

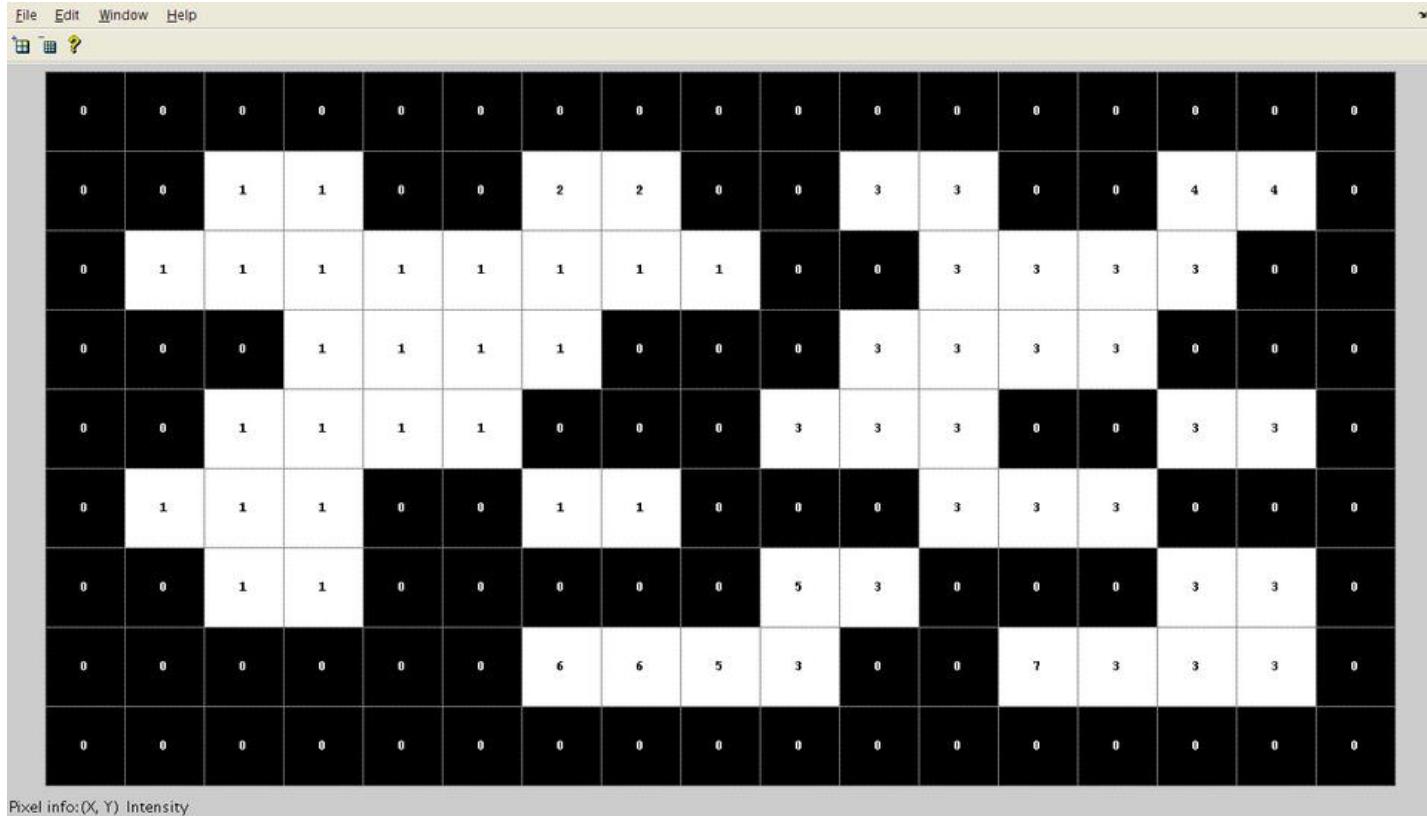
Region property computation

Input



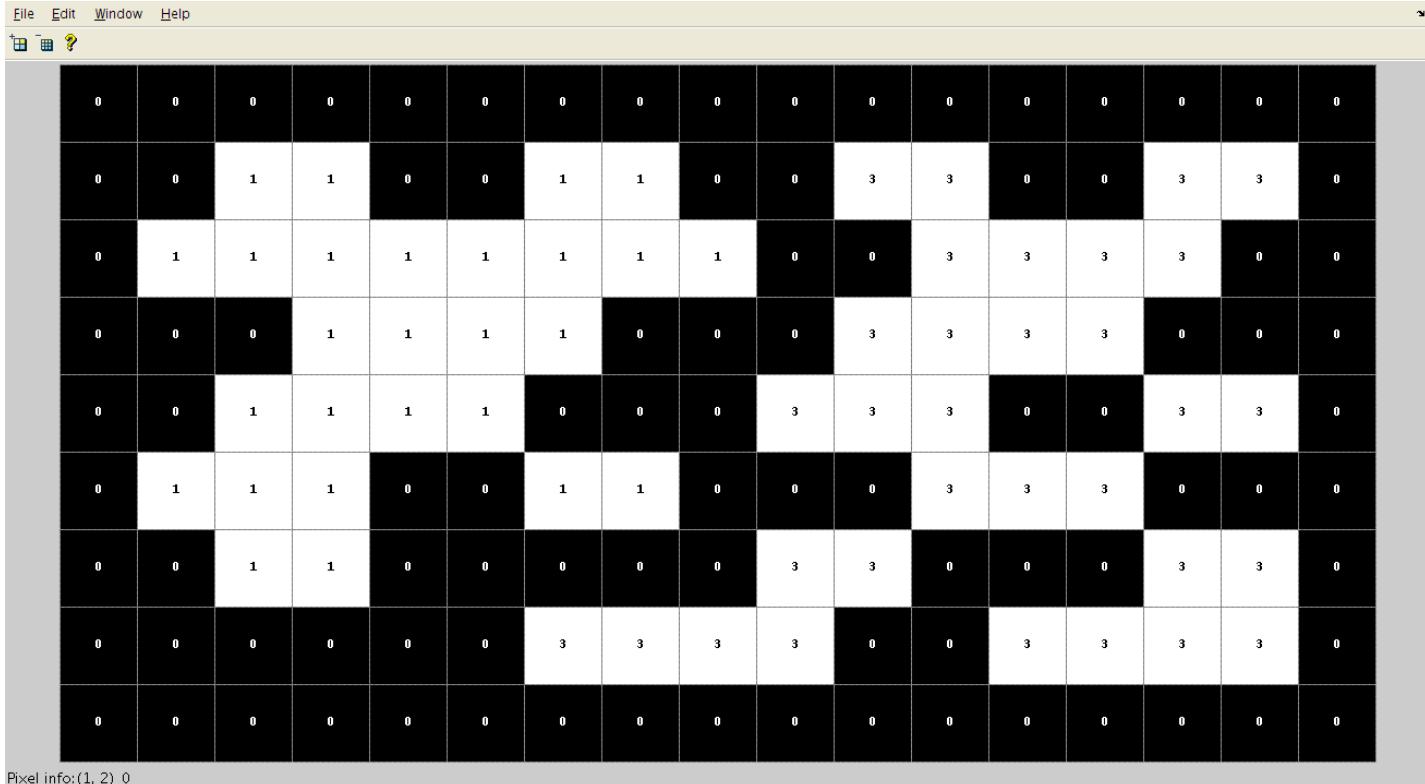
Region property computation

After First Pass



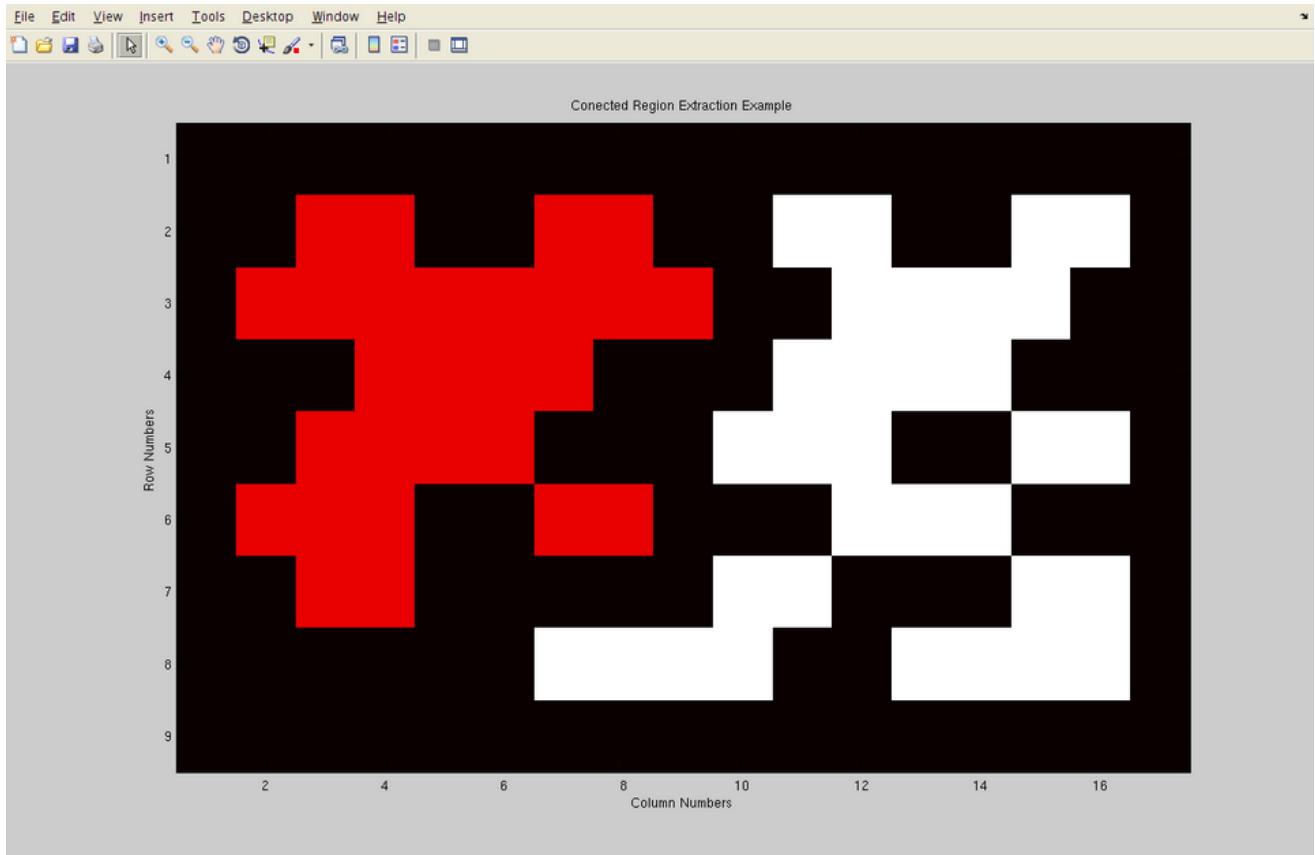
Region property computation

After Label Merging



Region property computation

Final Labeling



Outline

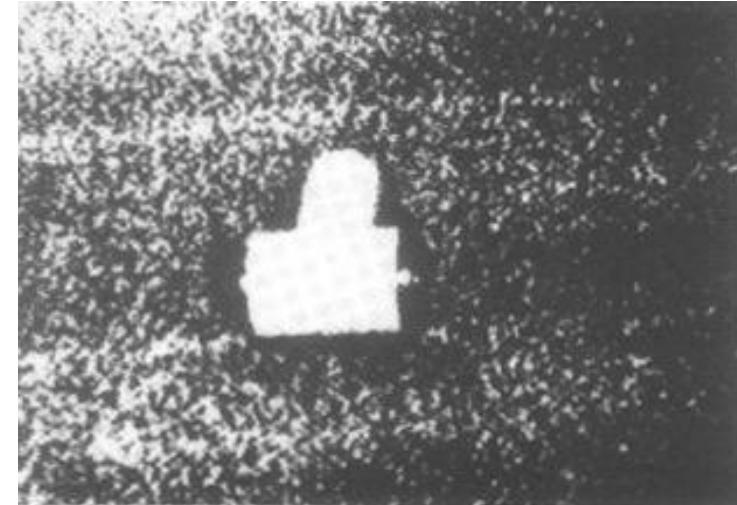
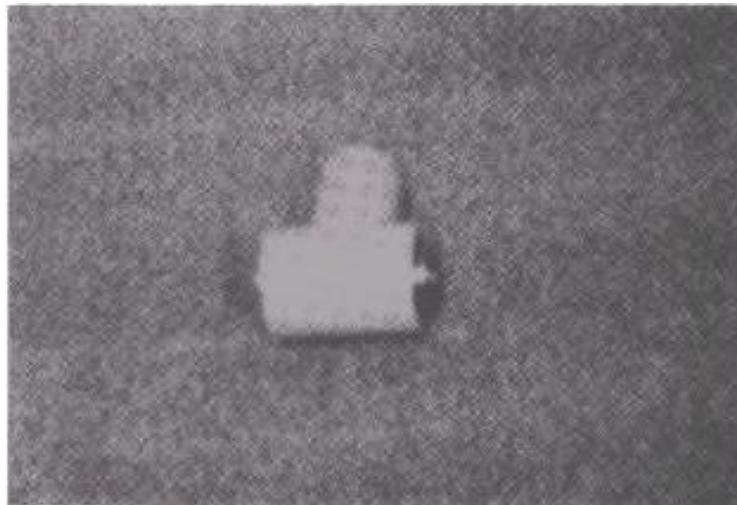
- Neighborhood image processing
 - Connected Component Analysis
 - Mathematical Morphology (Dilation, Erosion, Opening, Closing)
- Linear Image Filtering
 - Correlation and Convolution
 - Frequently used Filters
 - Gaussian Filter
- Nonlinear Image Filtering
 - Median Filter
 - Bilateral Filter
- Frequency Domain Analysis
 - Fourier Transform
 - Sampling and Aliasing

Mathematical Morphology

- **Morphology:** originally, a branch of biology science that studies the form and structure of animals and plants
- Mathematical Morphology in image processing is used to extract image components for representation and description of region shape, such as boundaries, skeletons, and the convex hull.

Motivation: remove small noise regions

- Example result by a binarization algorithm



Mathematical Morphology

In binary image processing, mathematical morphology consists of two basic operations,

dilation , erosion

and several composite operations

closing , opening, . . .

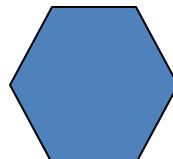
Structuring Elements

A **structuring element** is a shape mask used in the basic morphological operations.

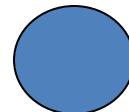
They can be any shape and size that is digitally representable, and each has an **origin**.



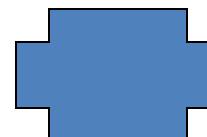
box



hexagon



disk



something

box(length, width)

disk(diameter)

Dilation

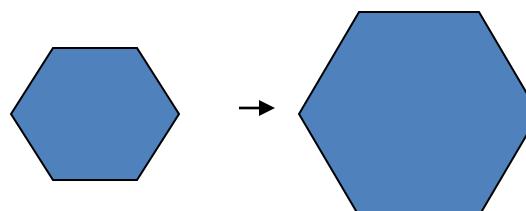
`dilate(B,S)` takes binary image B, places the origin of structuring element S over each 1-pixel, and **ORs** the structuring element S into the output image at the corresponding position.

Dilation

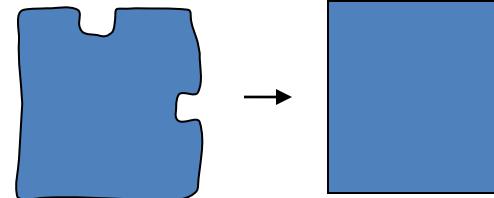
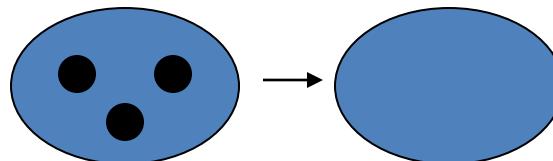
Dilation **expands** the connected sets of 1s of a binary image.

It can be used for

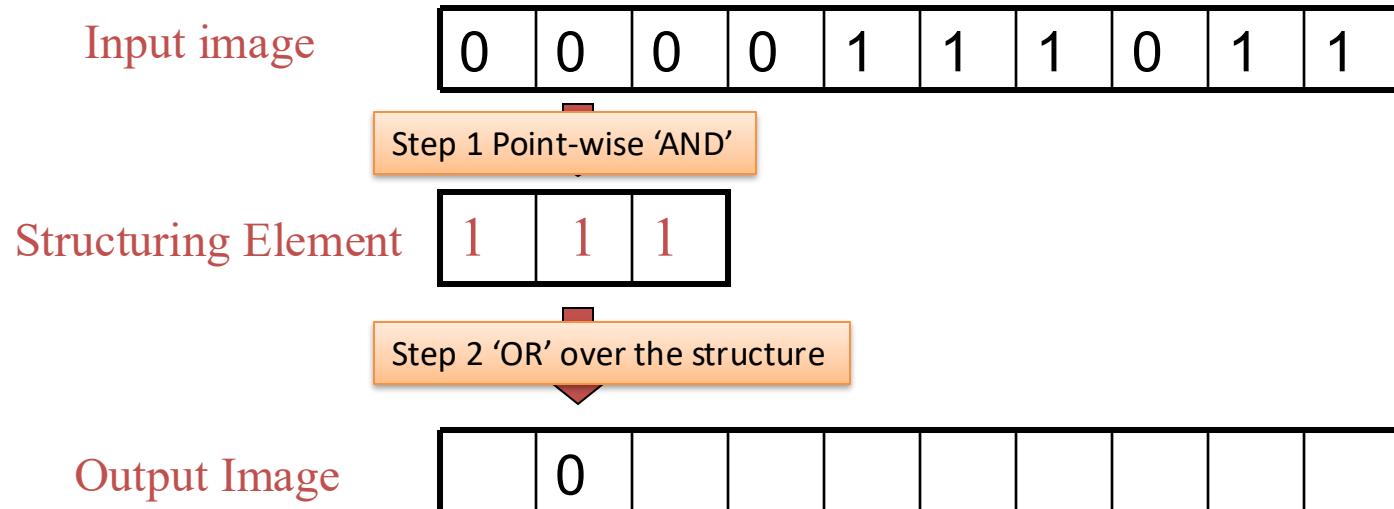
1. growing features



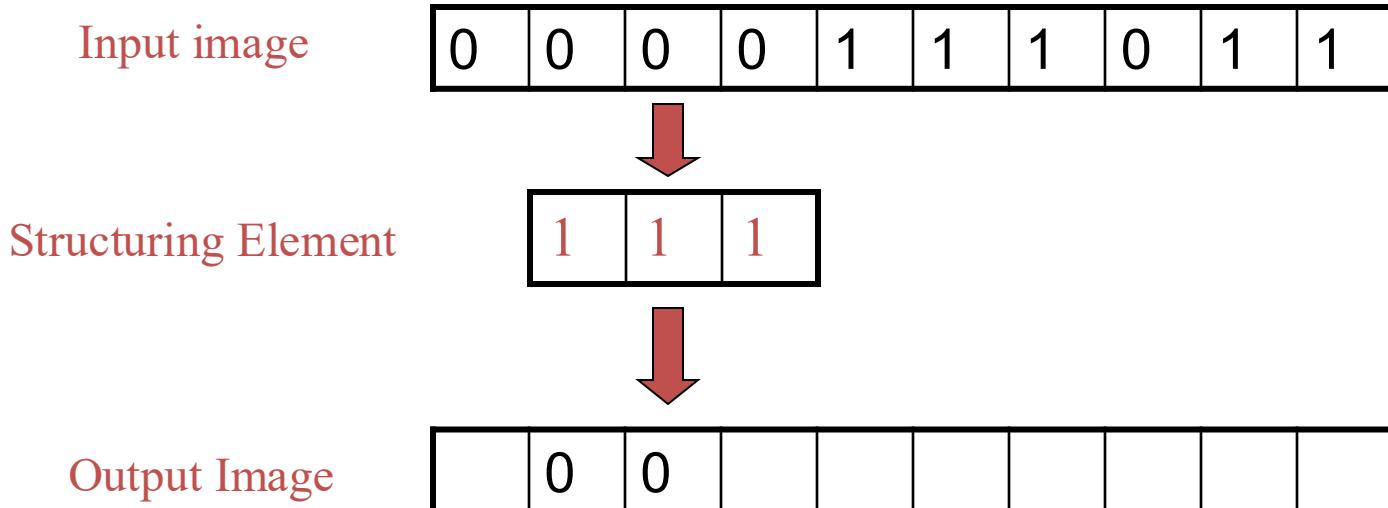
2. filling holes and gaps



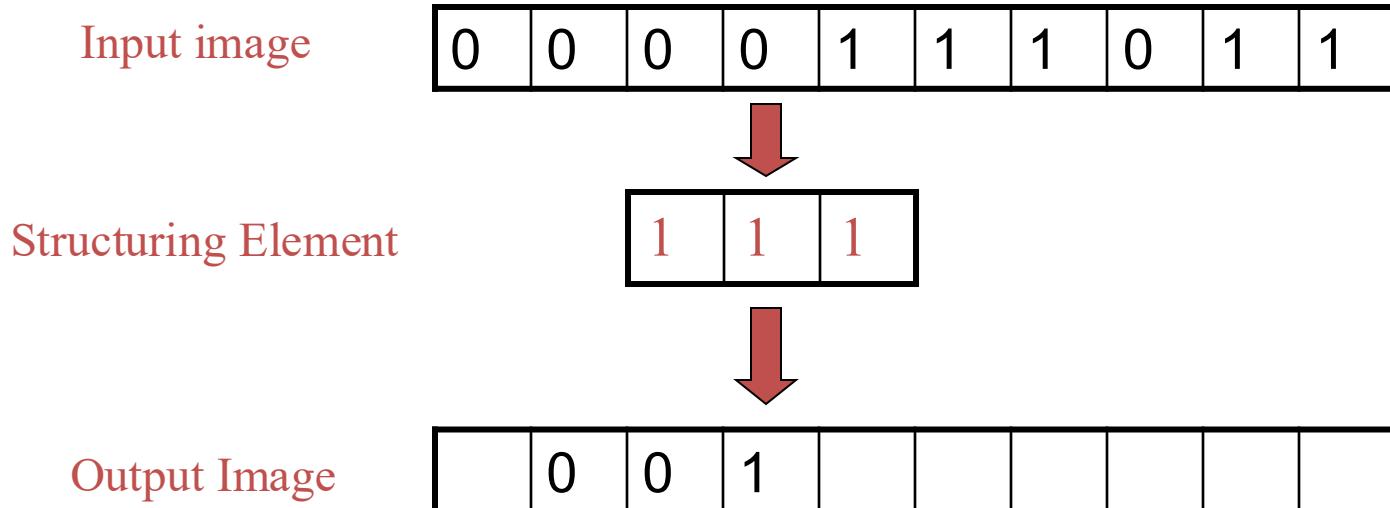
Example for Dilation



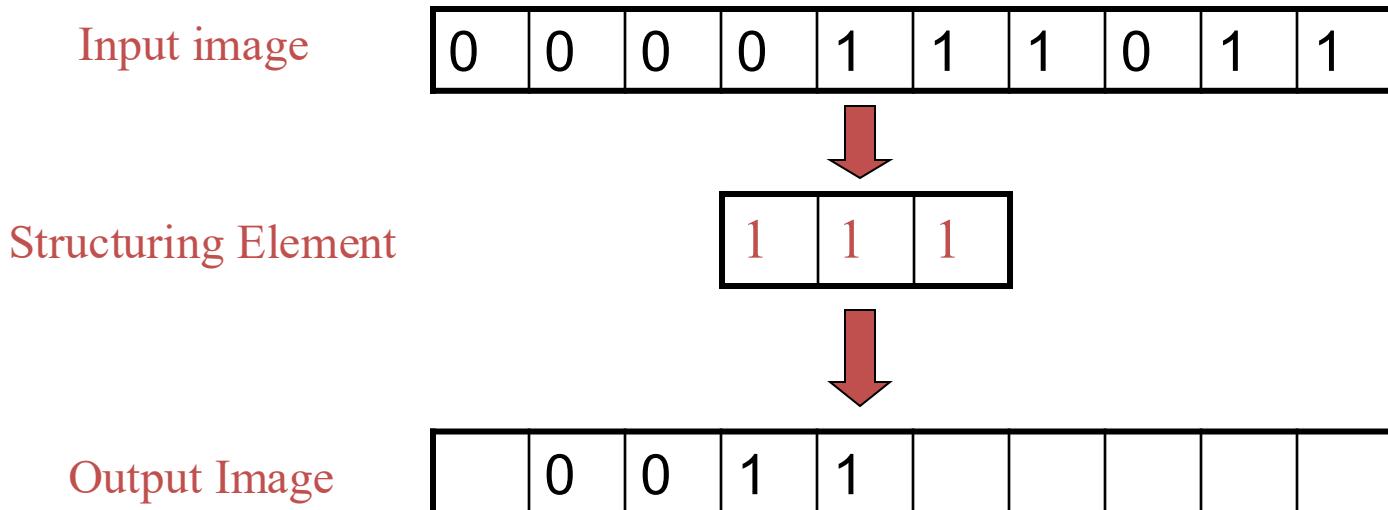
Example for Dilation



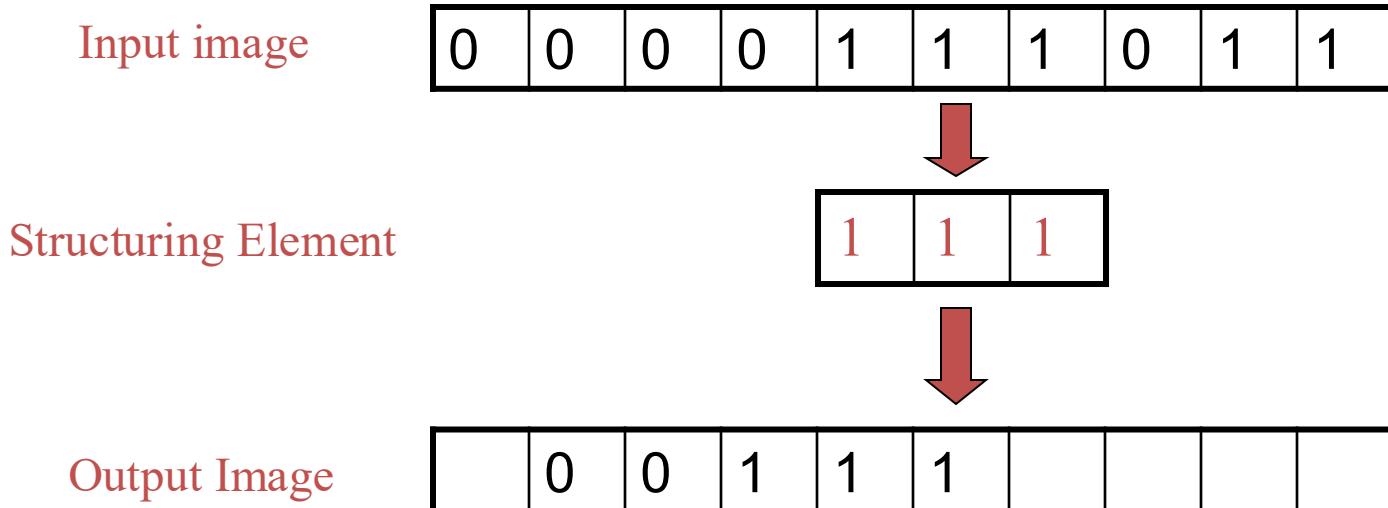
Example for Dilation



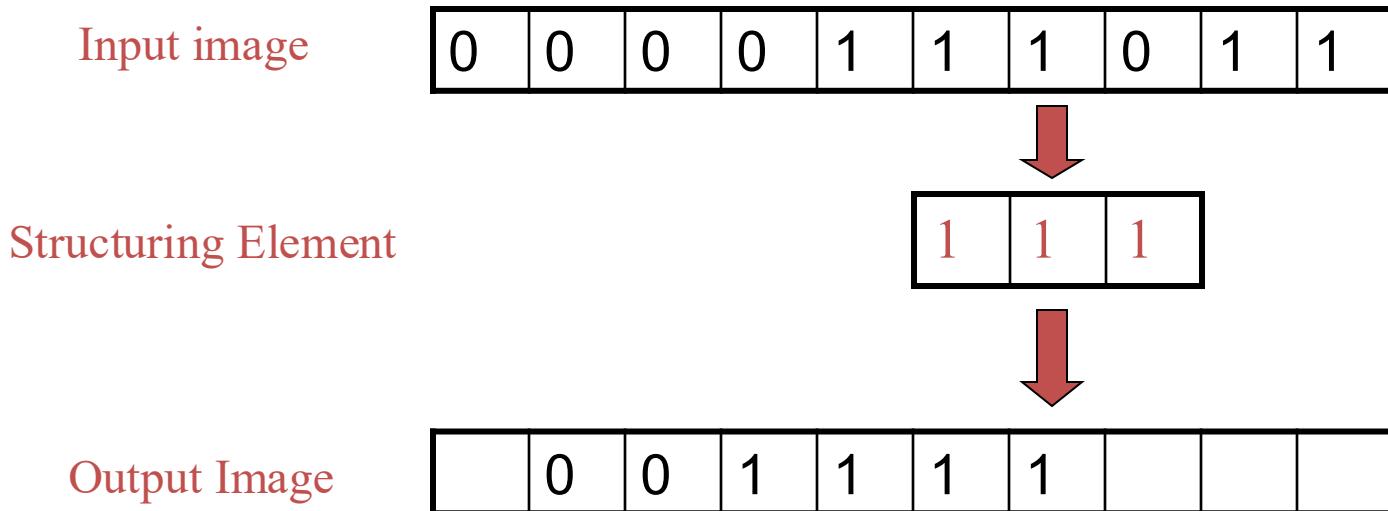
Example for Dilation



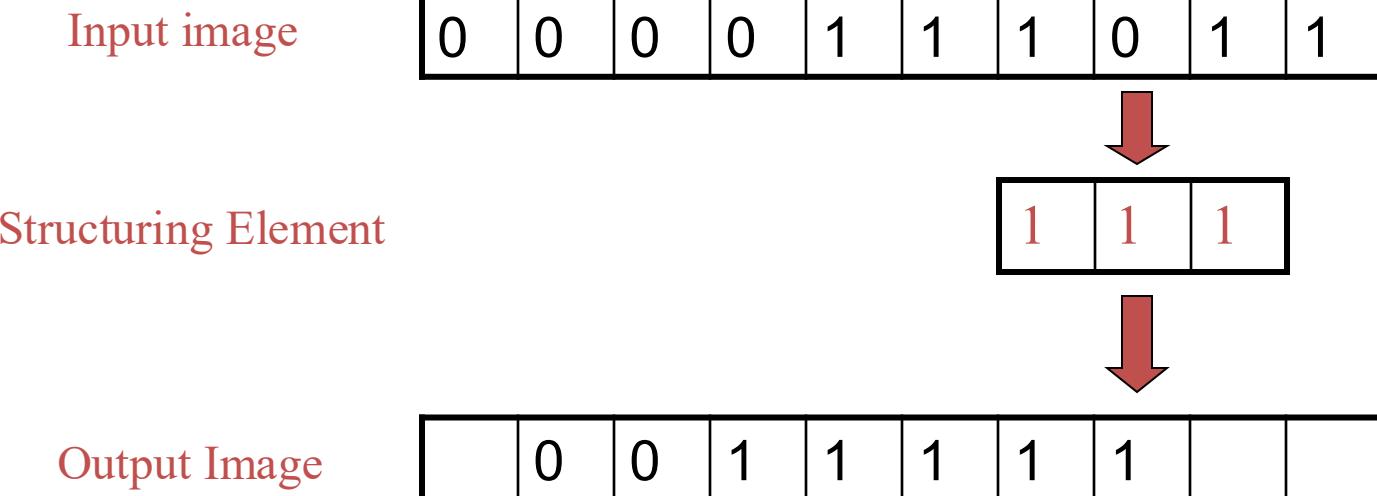
Example for Dilation



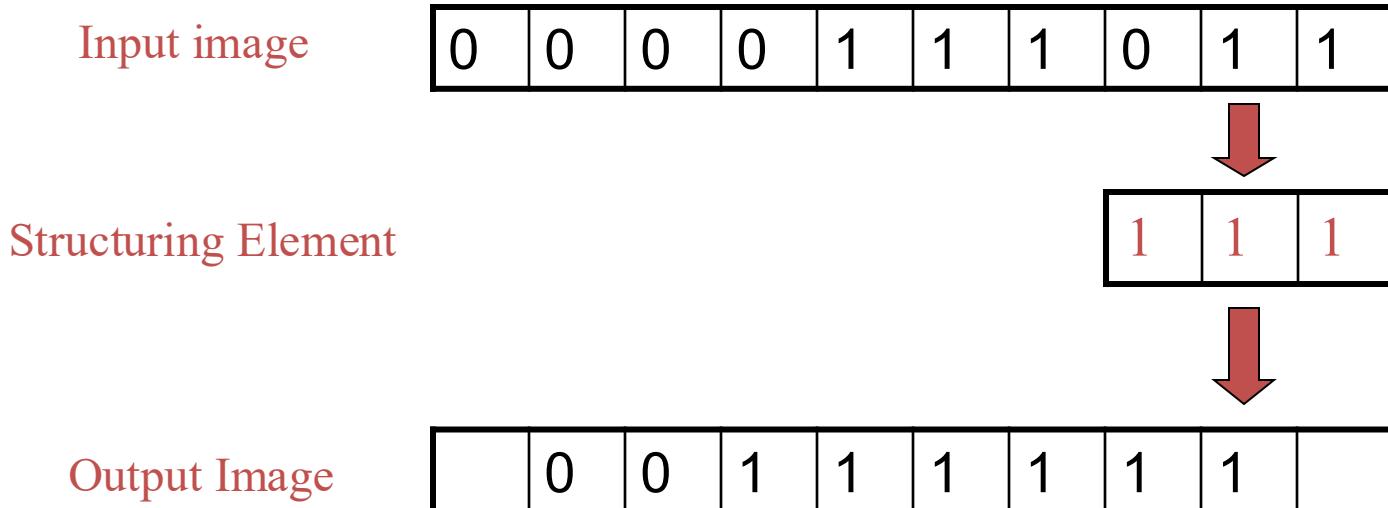
Example for Dilation



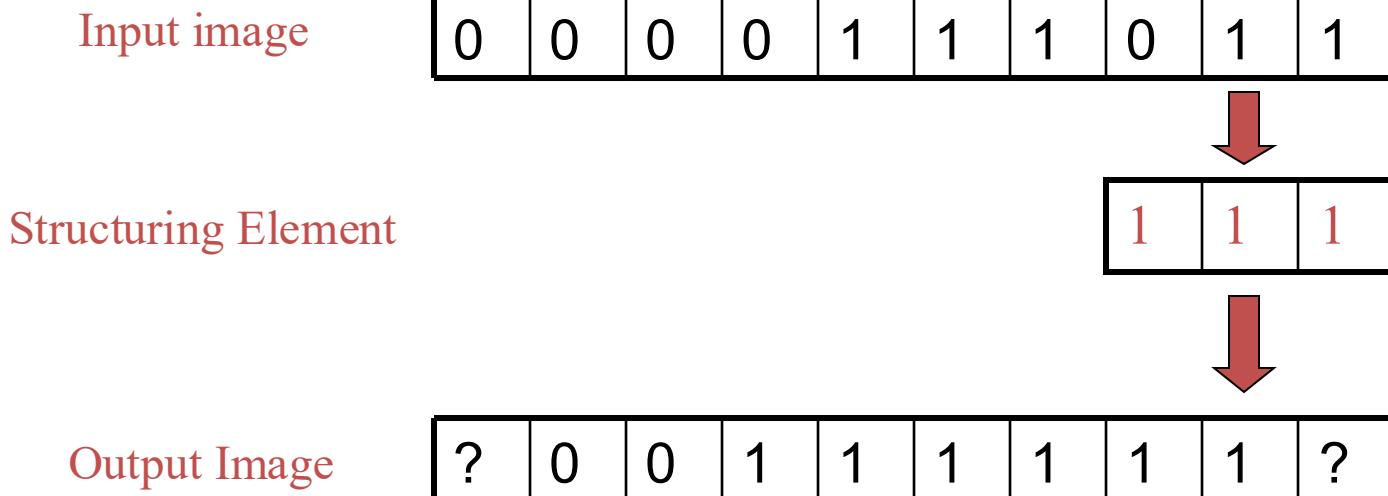
Example for Dilation



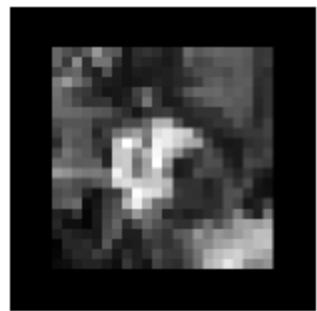
Example for Dilation



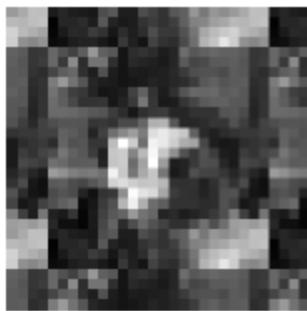
Example for Dilation



How to handle borders?



zero



wrap

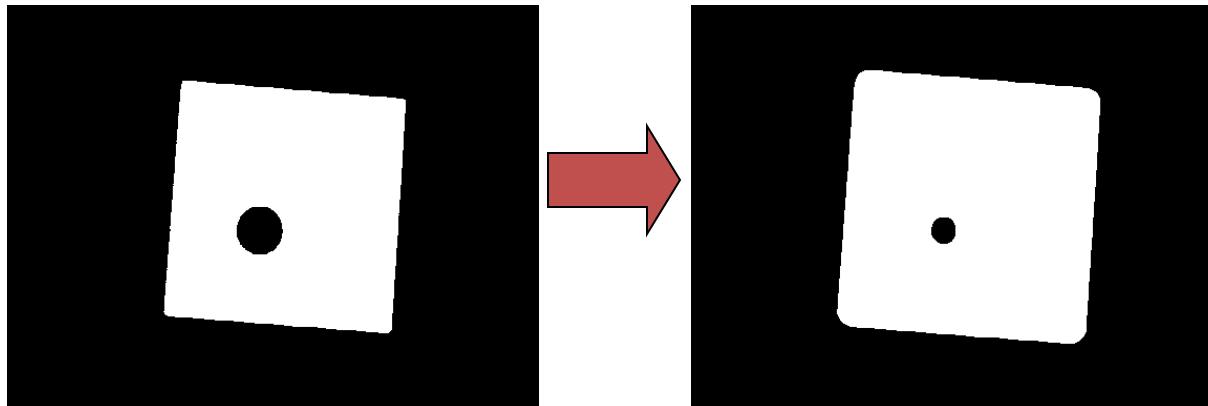


clamp



mirror

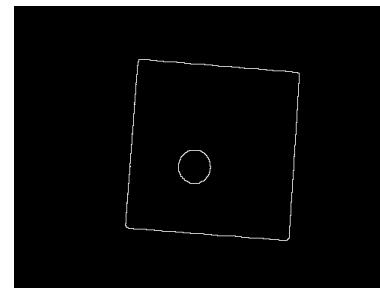
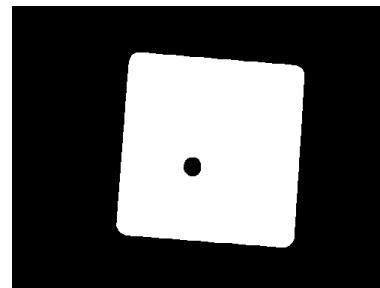
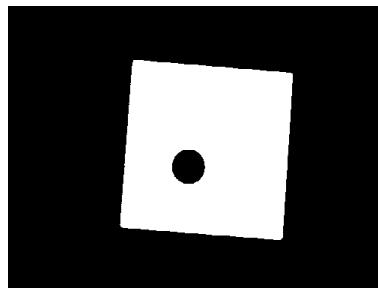
Another Dilation Example



- Image get lighter, more uniform intensity

Edge detection example

- Edge Detection
1. Dilate input image
 2. Subtract input image from dilated image
 3. Edges remain!



Erosion

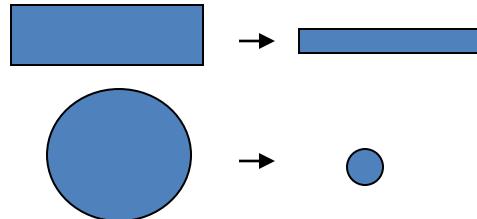
`erode(B,S)` takes a binary image B, places the origin of structuring element S over every pixel position, and **AND** a binary 1 into that position of the output image, if and only if every position of S (with a 1) covers a 1 in B.

Erosion

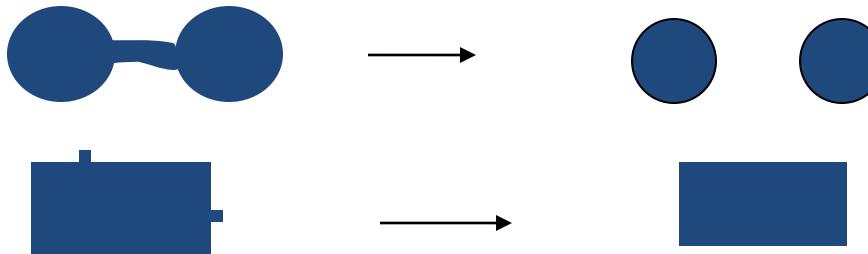
Erosion **shrinks** the connected sets of 1s of a binary image.

It can be used for

1. shrinking features

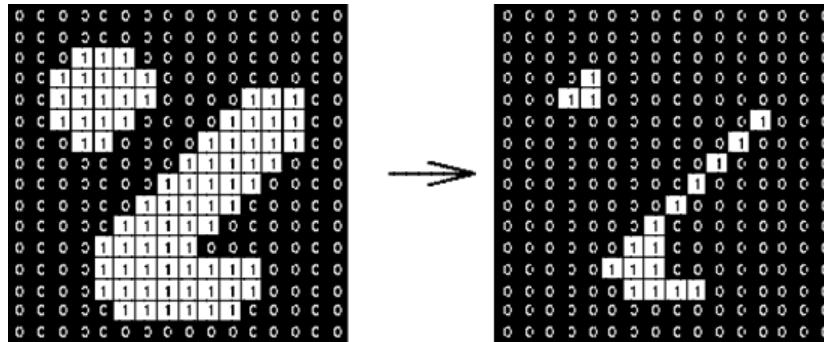


2. Removing bridges, branches and small protrusions



Erosion

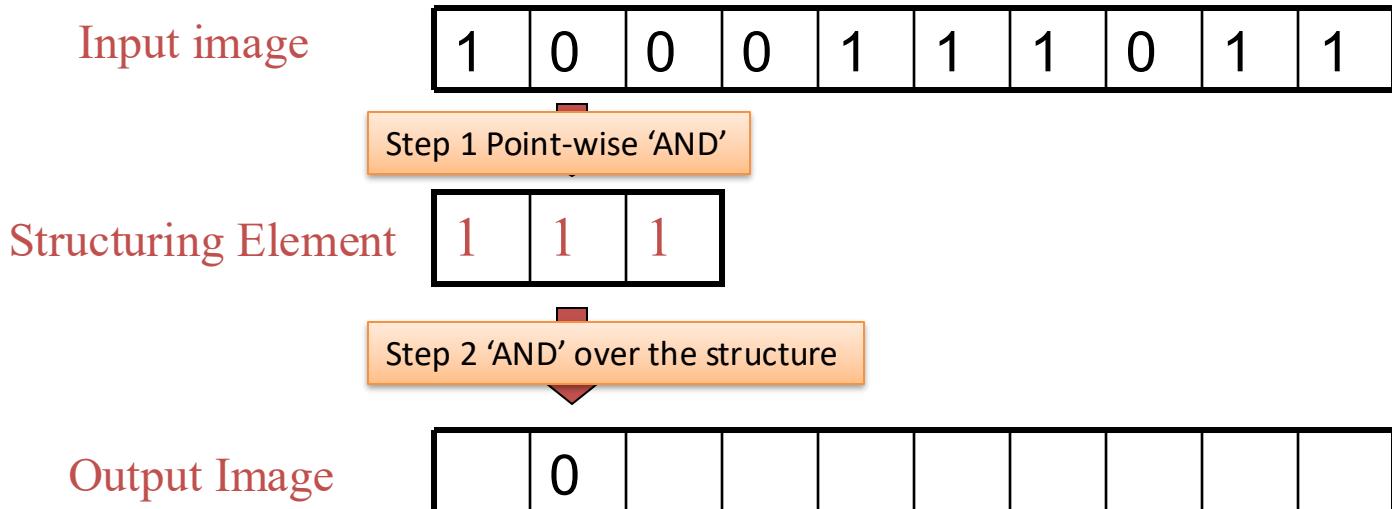
- **Erosion** is a basic morphological operation



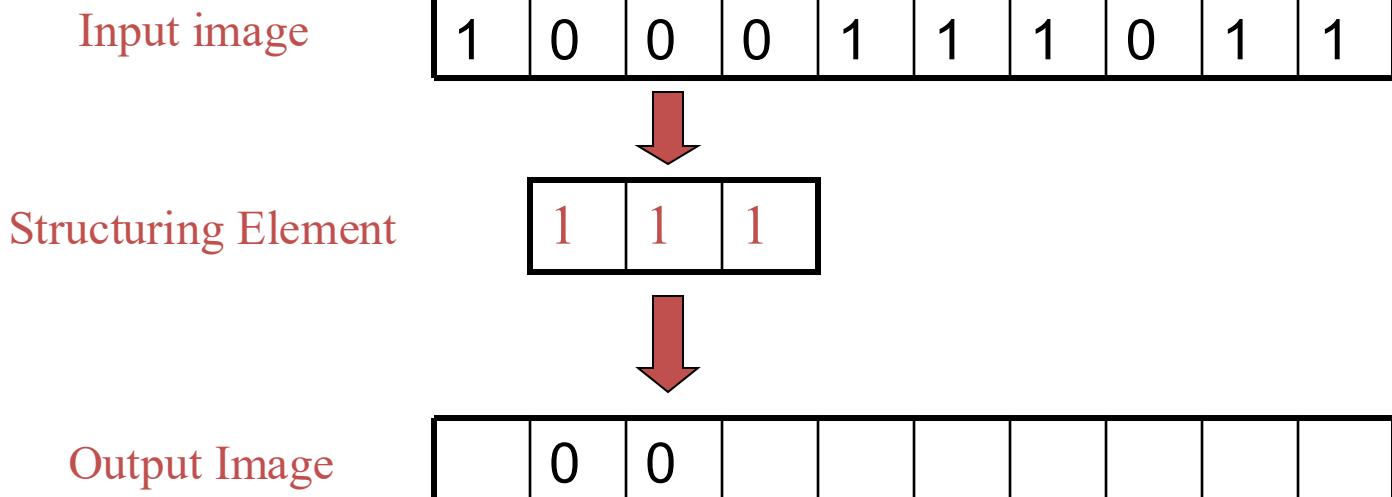
- Applied Structuring Element:

1	1	1
1	1	1
1	1	1

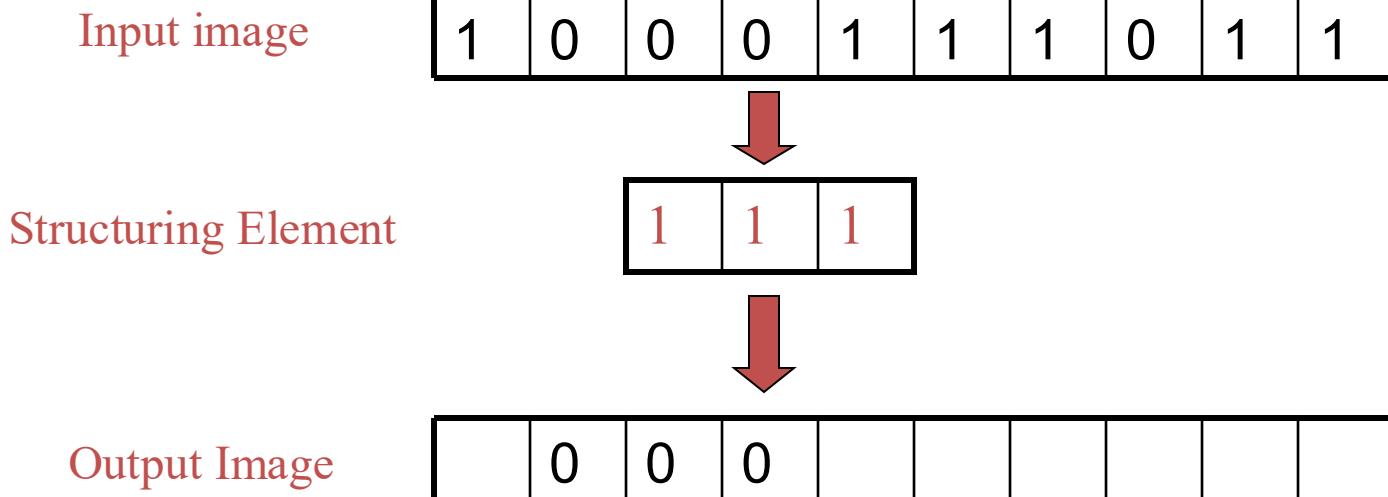
Example for Erosion



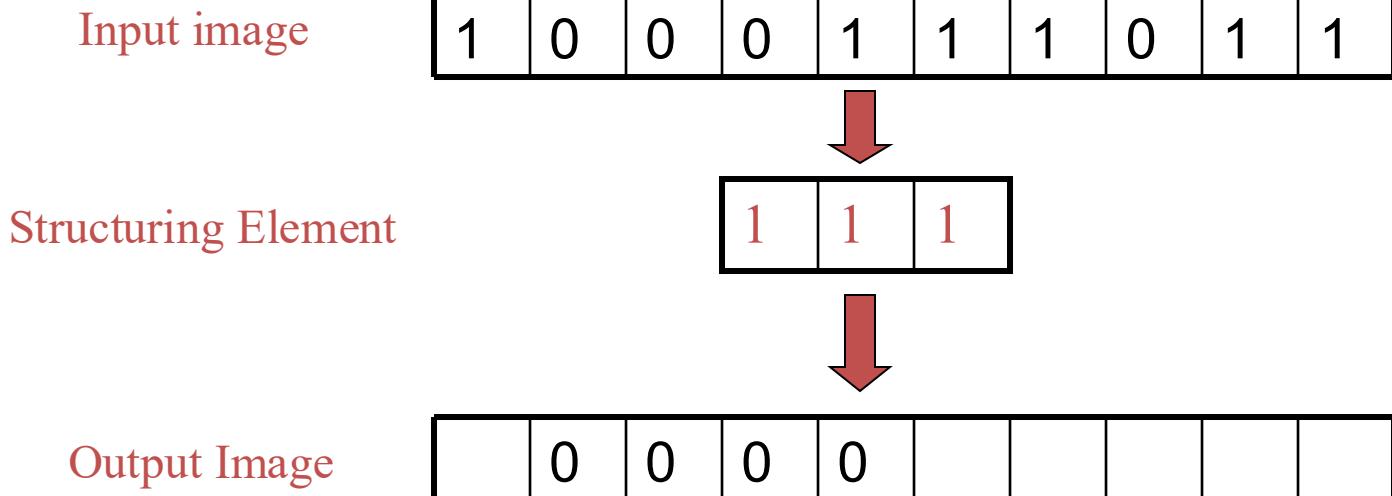
Example for Erosion



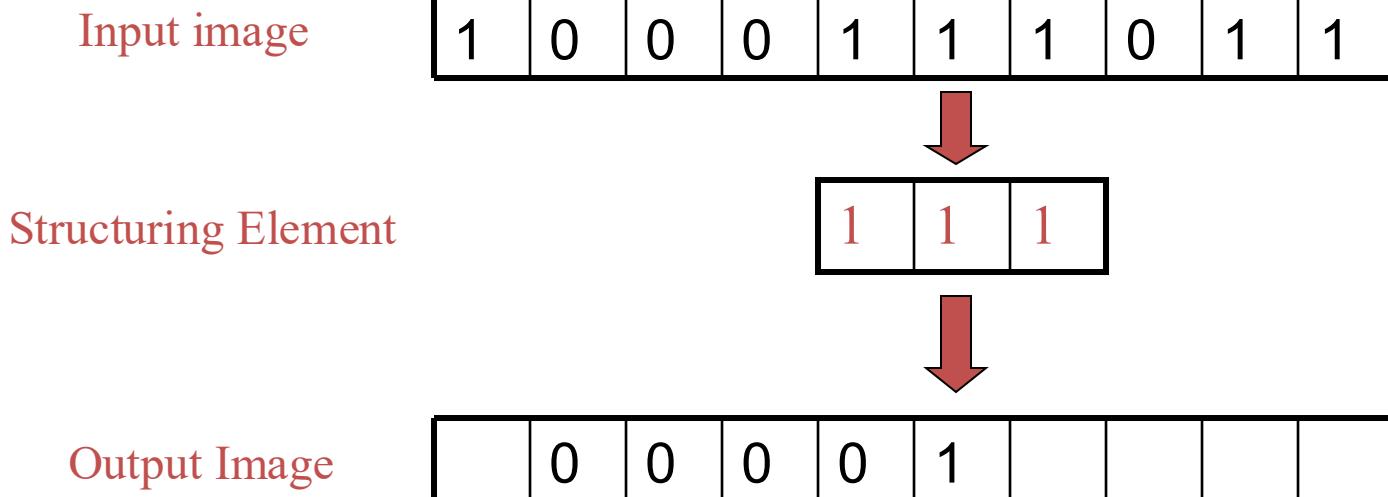
Example for Erosion



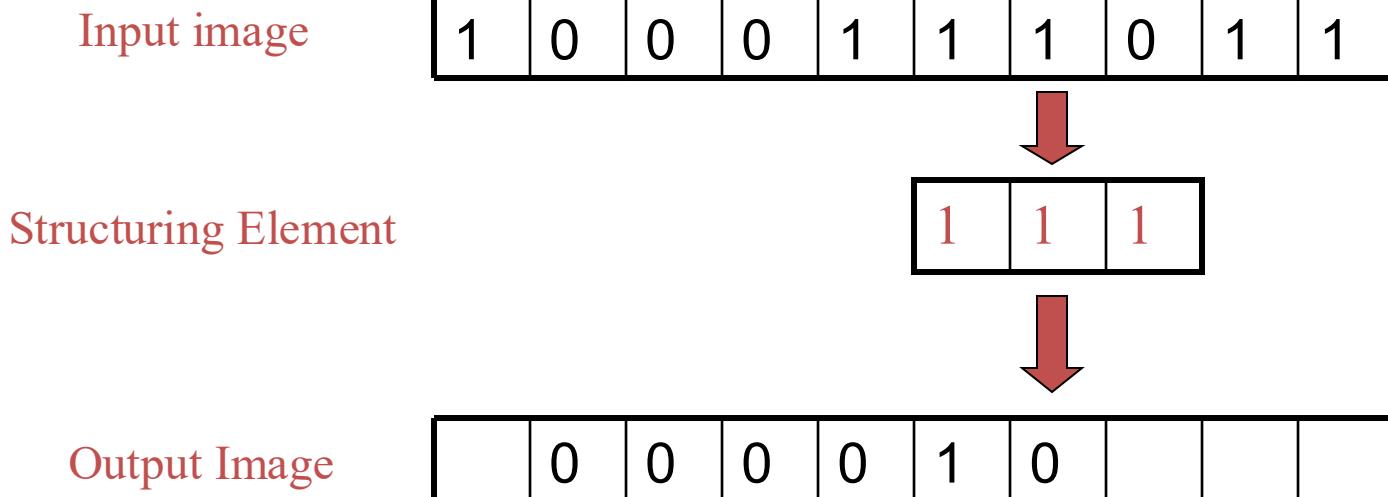
Example for Erosion



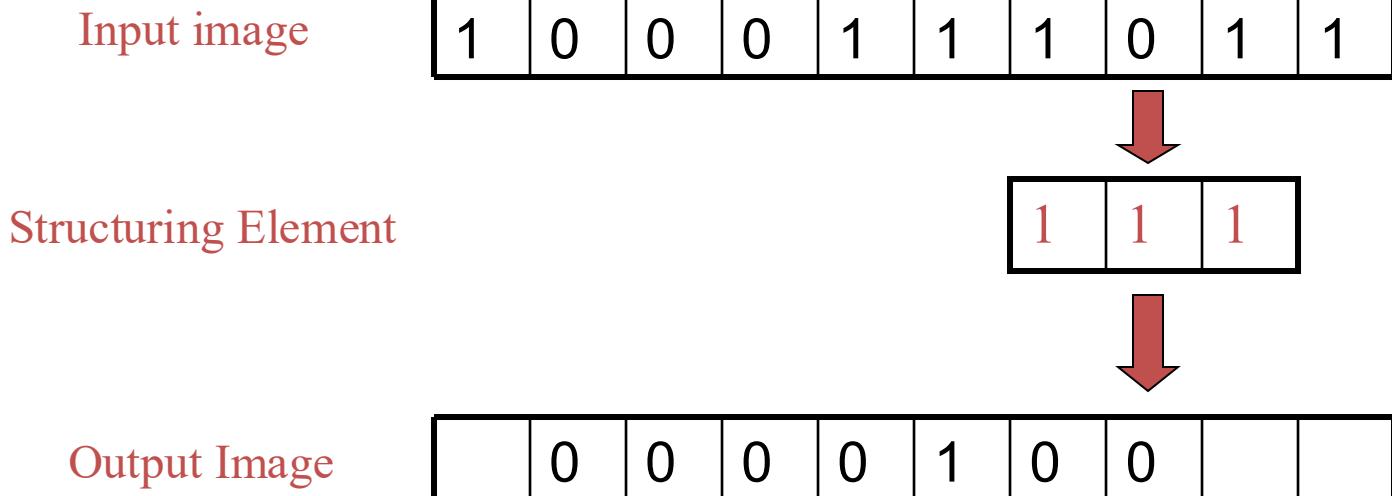
Example for Erosion



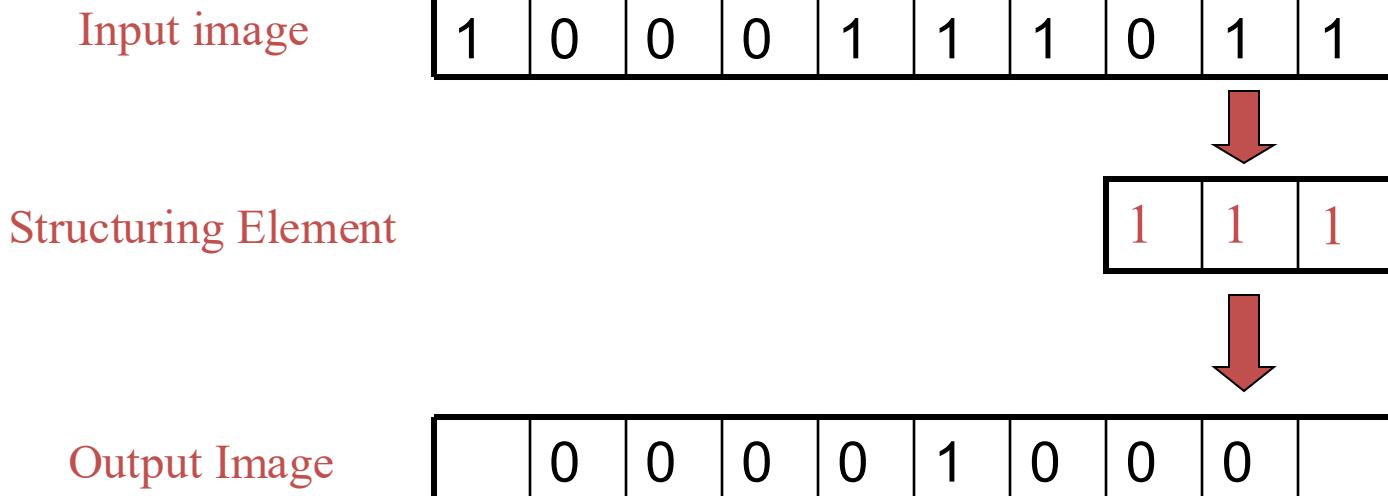
Example for Erosion



Example for Erosion

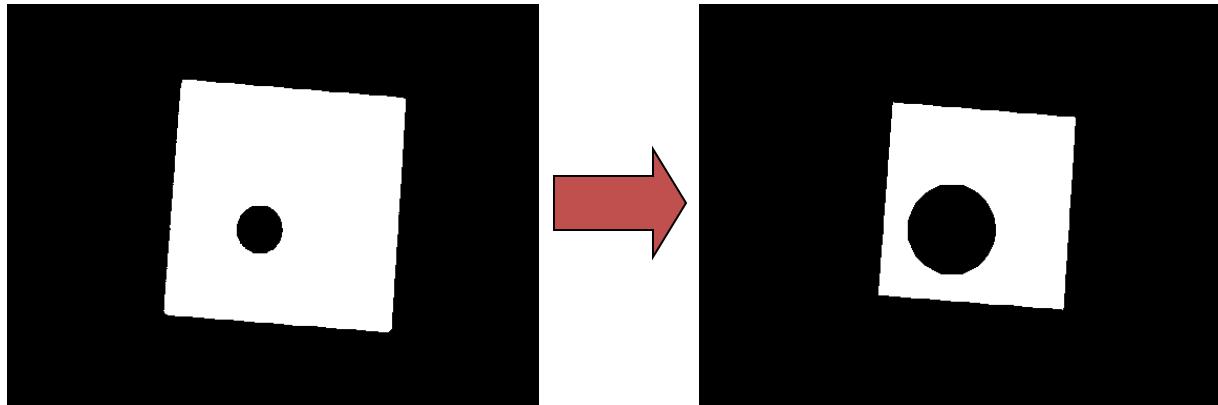


Example for Erosion



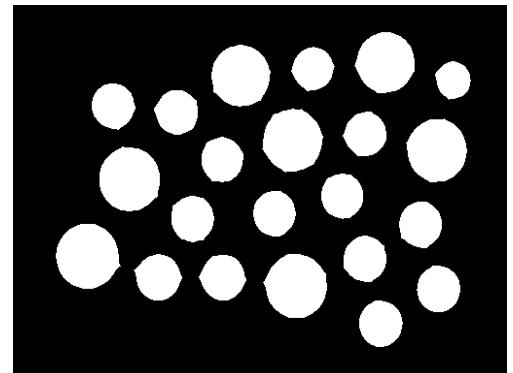
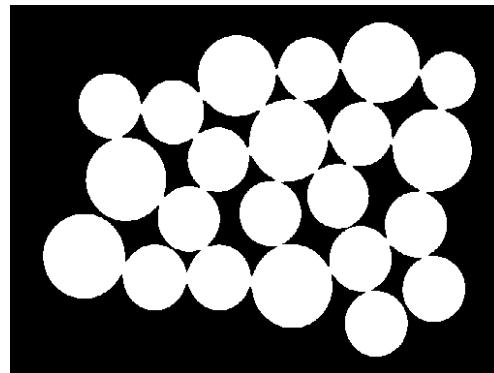
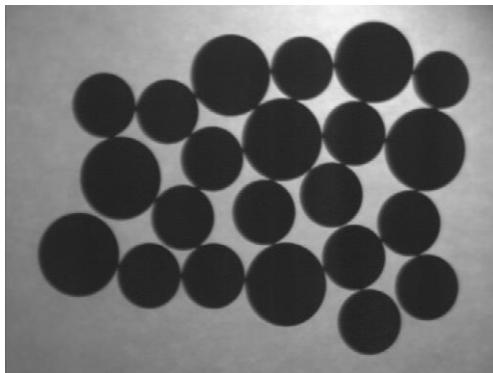
An example of erosion

- White = 0, black = 1, dual property, image as a result of erosion gets darker



Counting Coins

- Counting coins is difficult if they touch each other!
- Solution: Binarization and Erosion separates them!

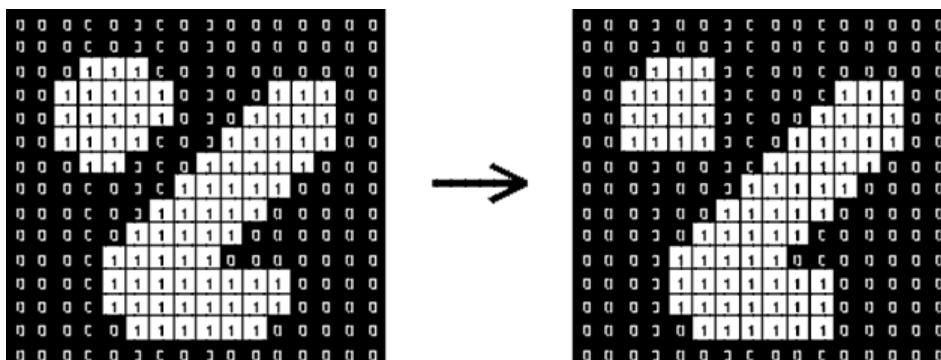


Opening

- Similar to Erosion
 - Spot and noise removal
 - Less destructive
- **Erosion next dilation**
- *the same structuring element for both operations.*

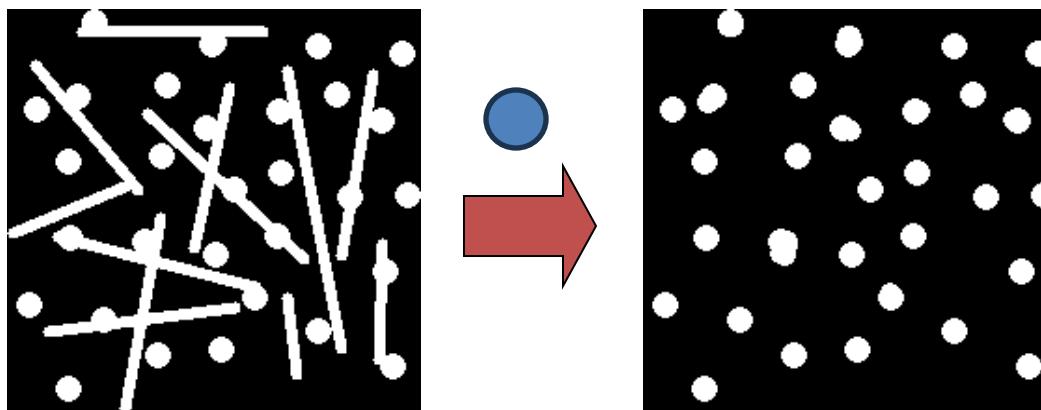
Opening

- Structuring element: 3x3 square

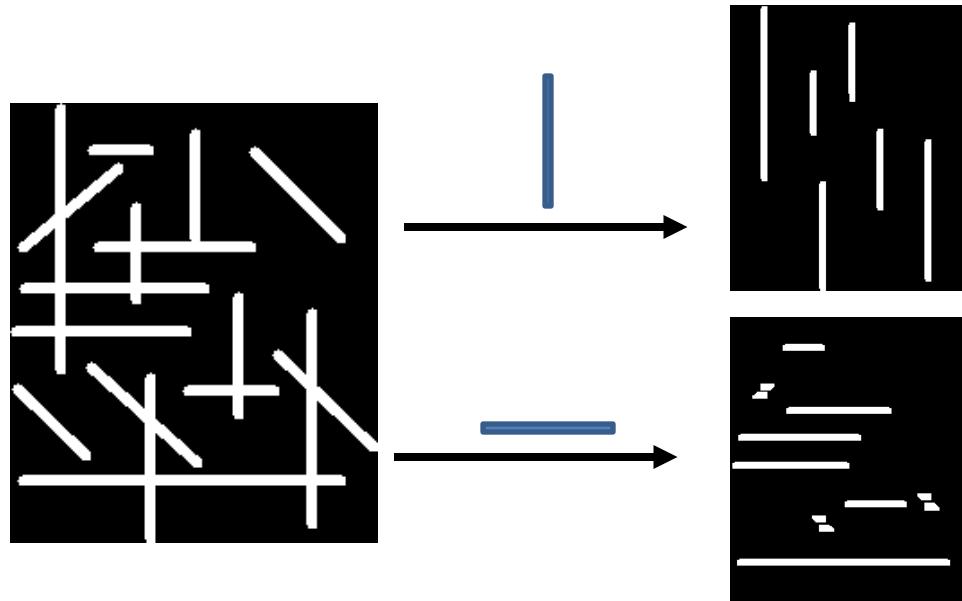


Opening Example

- Opening with a 11 pixel diameter disc
- Erosion plus dilation

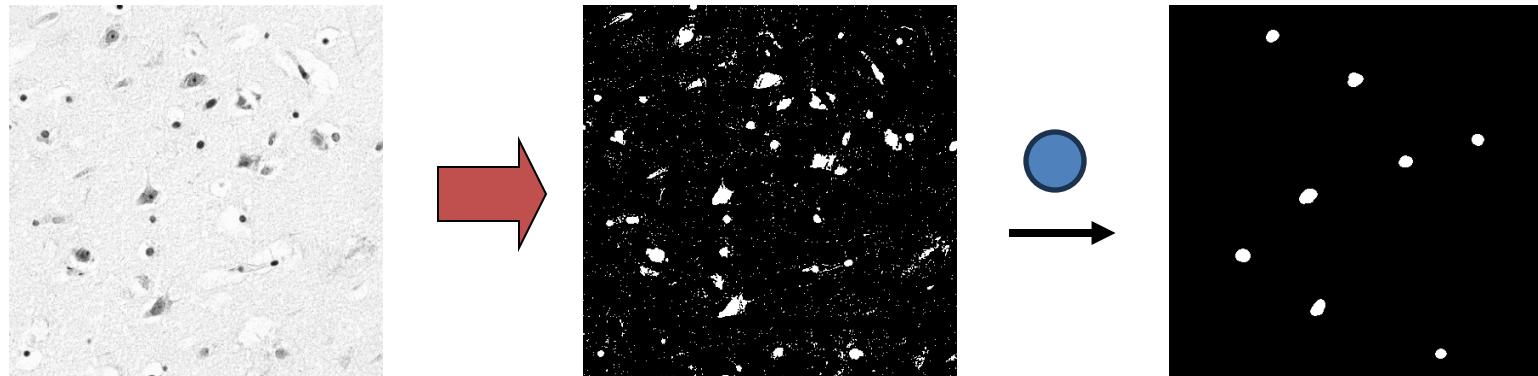


Opening Example



Use Opening for Separating Blobs

- Use large structuring element that fits into the big blobs
- Structuring Element: disc

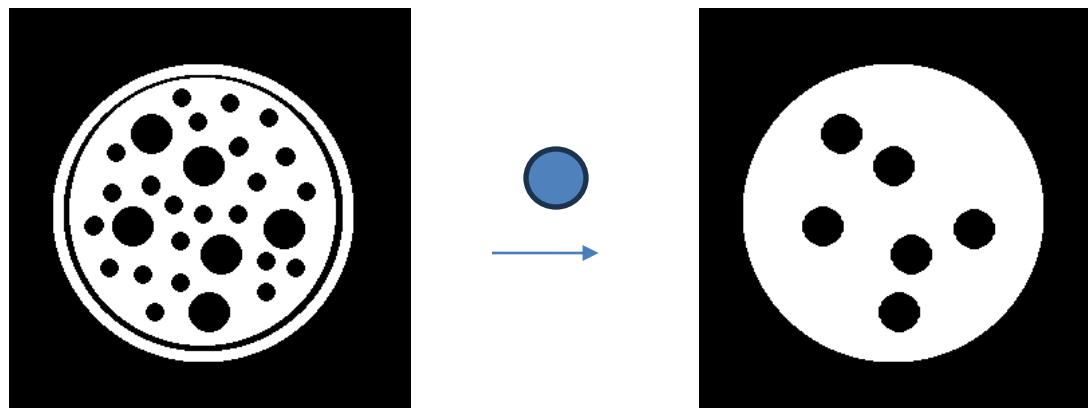


Closing

- Similar to Dilation
 - Removal of holes
 - Tends to enlarge regions, shrink background
- **Dilation next erosion.**

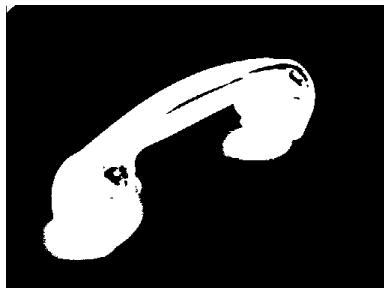
Closing Example

- Closing operation with a 22 pixel disc
- Closes small holes in the foreground



Closing Example

1. Threshold
2. Closing with disc of size 20



Opening & Closing

- Derived from the two basic operations
 - Dilation
 - Erosion
- Opening and closing are dual operations
 - i.e. opening the foreground pixels with a particular structuring element
 - is equivalent to closing the background pixels with the same element.

Outline

- Neighborhood image processing
 - Connected Component Analysis
 - Mathematical Morphology (Dilation, Erosion, Opening, Closing)
- Linear Image Filtering
 - Correlation and Convolution
 - Frequently used Filters
 - Gaussian Filter
- Nonlinear Image Filtering
 - Median Filter
 - Bilateral Filter
- Frequency Domain Analysis
 - Fourier Transform
 - Sampling and Aliasing

Filtering

Goal: Remove unwanted parts of a signal, and keep only task-relevant information.



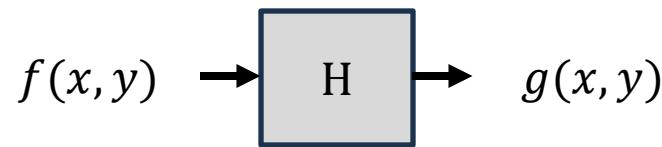
For images, new pixel values are typically computed based on some function of the local neighborhood of the pixels.



Linear Filtering

A linear filter in can be written as

$$g(i, j) = \sum_{k,l} f(i + k, j + l) \cdot h(k, l)$$



Linear, if and only if: $H(a(i, j) + b(i, j)) = H(a(i, j)) + H(b(i, j))$

$$H(C \cdot a(i, j)) = C \cdot H(a(i, j))$$

Motivation: Noise reduction

- Given a camera and a still scene, how can you reduce noise?



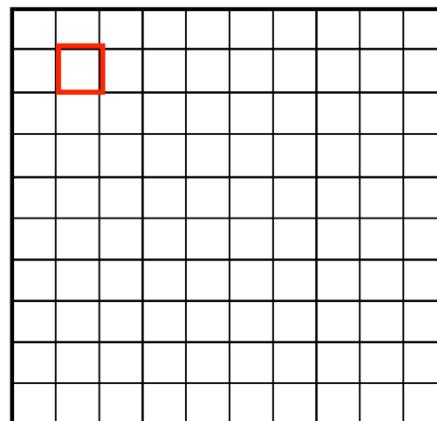
Mean filtering

$$\begin{array}{c} \begin{matrix} & & \\ & & \\ & & \end{matrix} \\ H \end{array} * \begin{array}{c} \begin{matrix} 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 90 & 90 & 90 & 90 & 90 & 0 & 0 \\ 0 & 0 & 0 & 90 & 90 & 90 & 90 & 90 & 0 & 0 \\ 0 & 0 & 0 & 90 & 90 & 90 & 90 & 90 & 0 & 0 \\ 0 & 0 & 0 & 90 & 90 & 90 & 90 & 90 & 0 & 0 \\ 0 & 0 & 0 & 90 & 0 & 90 & 90 & 90 & 0 & 0 \\ 0 & 0 & 0 & 90 & 90 & 90 & 90 & 90 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 90 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{matrix} \\ F \end{array} = \begin{array}{c} \begin{matrix} & & & & & & & & \\ 0 & 10 & 20 & 30 & 30 & 30 & 20 & 10 & \\ 0 & 20 & 40 & 60 & 60 & 60 & 40 & 20 & \\ 0 & 30 & 60 & 90 & 90 & 90 & 60 & 30 & \\ 0 & 30 & 50 & 80 & 80 & 90 & 60 & 30 & \\ 0 & 30 & 50 & 80 & 80 & 90 & 60 & 30 & \\ 0 & 20 & 30 & 50 & 50 & 60 & 40 & 20 & \\ 10 & 20 & 30 & 30 & 30 & 30 & 20 & 10 & \\ 10 & 10 & 10 & 0 & 0 & 0 & 0 & 0 & \end{matrix} \\ G \end{array}$$

Mean filtering / Moving average

 $F[x, y]$

0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	0	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0

 $G[x, y]$ 

Mean filtering / Moving average

 $F[x, y]$

0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	0	90	90	90	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0

 $G[x, y]$

0	10										

Mean filtering / Moving average

 $F[x, y]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

 $G[x, y]$

0	10	20								

Mean filtering / Moving average

 $F[x, y]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

 $G[x, y]$

0	10	20	30							

Mean filtering / Moving average

 $F[x, y]$

0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0

 $G[x, y]$

0	10	20	30	30

Mean filtering / Moving average

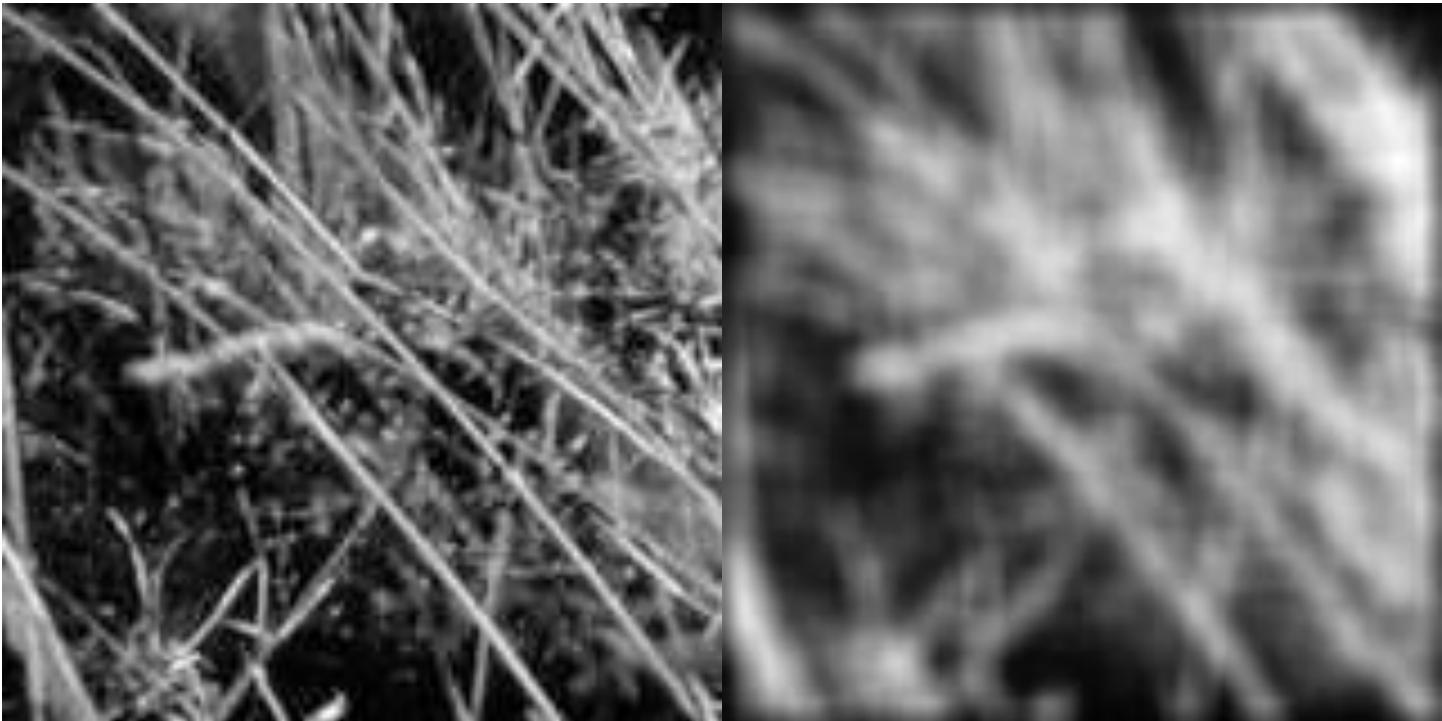
$F[x, y]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

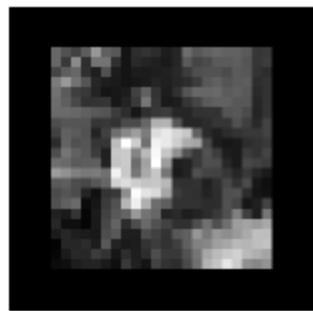
$G[x, y]$

	0	10	20	30	30	30	20	10		
	0	20	40	60	60	60	40	20		
	0	30	60	90	90	90	60	30		
	0	30	50	80	80	90	60	30		
	0	30	50	80	80	90	60	30		
	0	20	30	50	50	60	40	20		
	10	20	30	30	30	30	20	10		
	10	10	10	0	0	0	0	0		

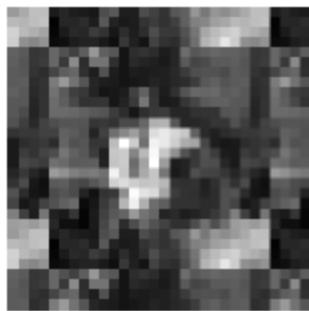
Mean filtering / Moving average



How to handle borders?



zero



wrap



clamp



mirror

Linear filters: examples



$$\text{Original} \quad * \quad \begin{array}{|c|c|c|} \hline 0 & 0 & 0 \\ \hline 0 & 1 & 0 \\ \hline 0 & 0 & 0 \\ \hline \end{array} \quad = \quad ?$$

Linear filters: examples

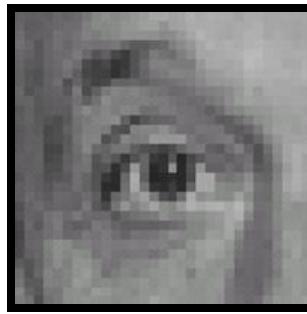


Original

*

0	0	0
0	1	0
0	0	0

=



Identical image
(no change)

Linear filters: examples

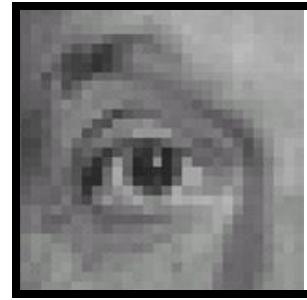


$$\text{Original} \quad * \quad \begin{array}{|c|c|c|} \hline 0 & 0 & 0 \\ \hline 1 & 0 & 0 \\ \hline 0 & 0 & 0 \\ \hline \end{array} \quad = \quad ?$$

Linear filters: examples

 $*$

0	0	0
1	0	0
0	0	0

 $=$ 

Original

Shifted left by 1 pixel

Linear filters: examples

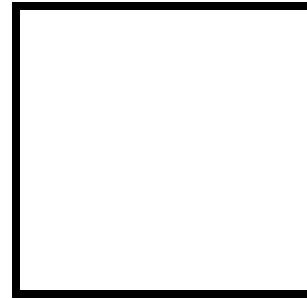


$$\text{Original} \quad * \quad \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array} \quad = \quad ?$$

Linear filters: examples

 $*$

1	1	1
1	1	1
1	1	1

 $=$ 

Original

Linear filters: examples

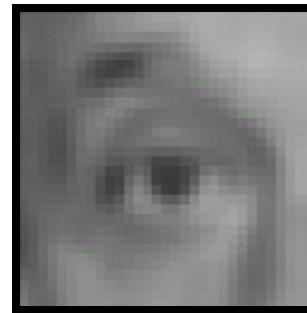


$$\text{Original} * \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} = ?$$

Linear filters: examples



$$\ast \frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array} =$$



Linear filters: examples

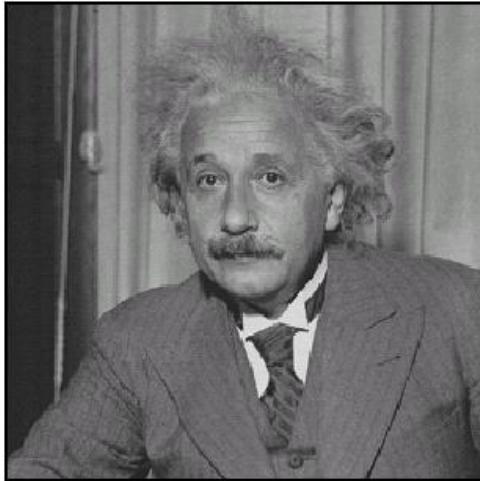


$$\text{Original} * \left(\begin{array}{ccc} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{array} - \frac{1}{9} \begin{array}{ccc} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{array} \right) = \text{Sharpened Image}$$

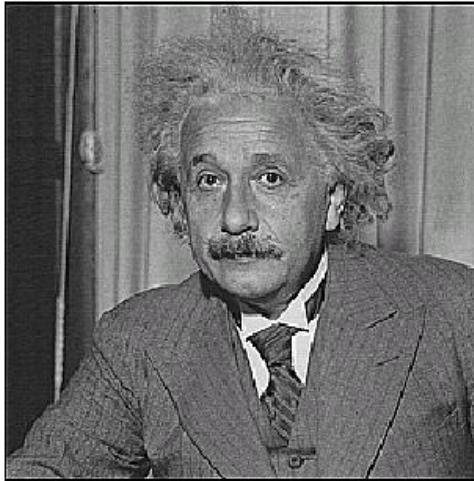


Sharpening filter
(accentuates edges)

Sharpening



before



after

Edge filter

$$\begin{bmatrix} -1 & 1 \end{bmatrix}$$

$$\frac{\partial \mathbf{I}}{\partial x} \simeq \mathbf{I}(x, y) - \mathbf{I}(x - 1, y)$$

 \otimes

$$\begin{bmatrix} -1 & 1 \end{bmatrix}$$

 $=$

$$\mathbf{h}[m,n]$$

$$\mathbf{g}[m,n]$$



$$\mathbf{f}[m,n]$$

Edge filter

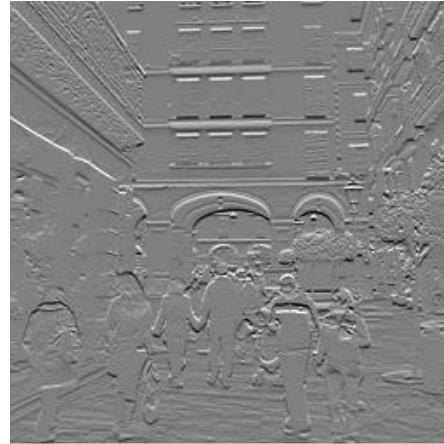
$$[-1 \ 1]^T$$

 \otimes

$$[-1, 1]^T =$$

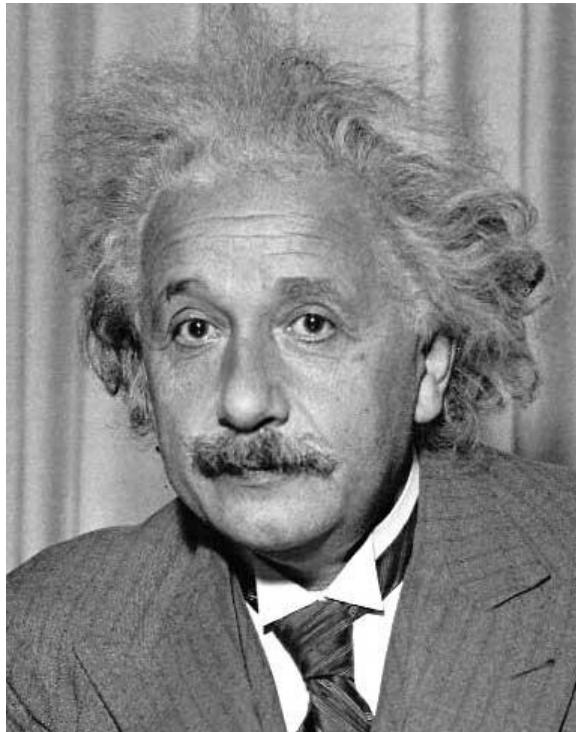
$$h[m,n]$$

$$g[m,n]$$



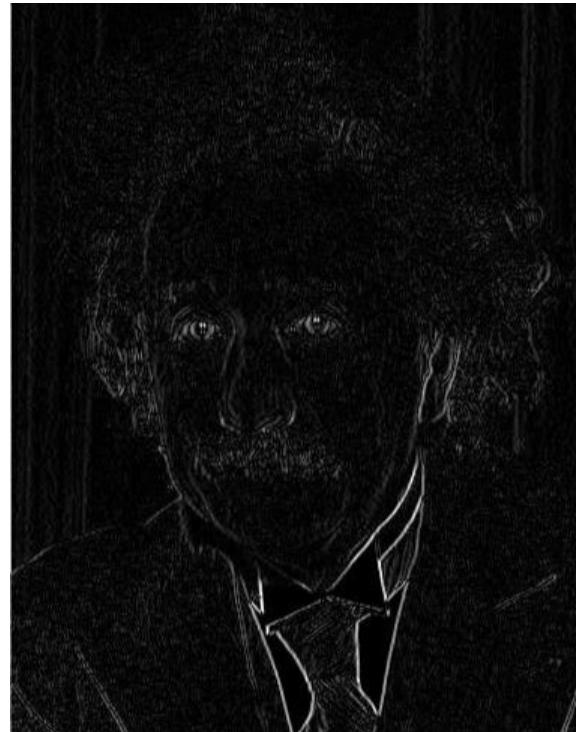
$$f[m,n]$$

Derivative filters



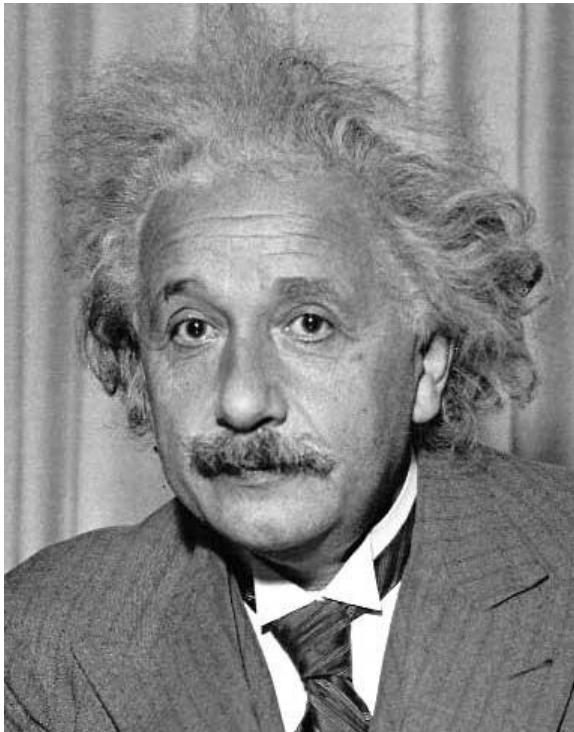
1	0	-1
2	0	-2
1	0	-1

Sobel



Vertical Edge
(absolute value)

Derivative filters



1	2	1
0	0	0
-1	-2	-1

Sobel



Horizontal Edge
(absolute value)

Motion blur filter

 $g[m,n]$

$$\square \quad \text{---} =$$

 $h[m,n]$  $f[m,n]$

Motion blur filter

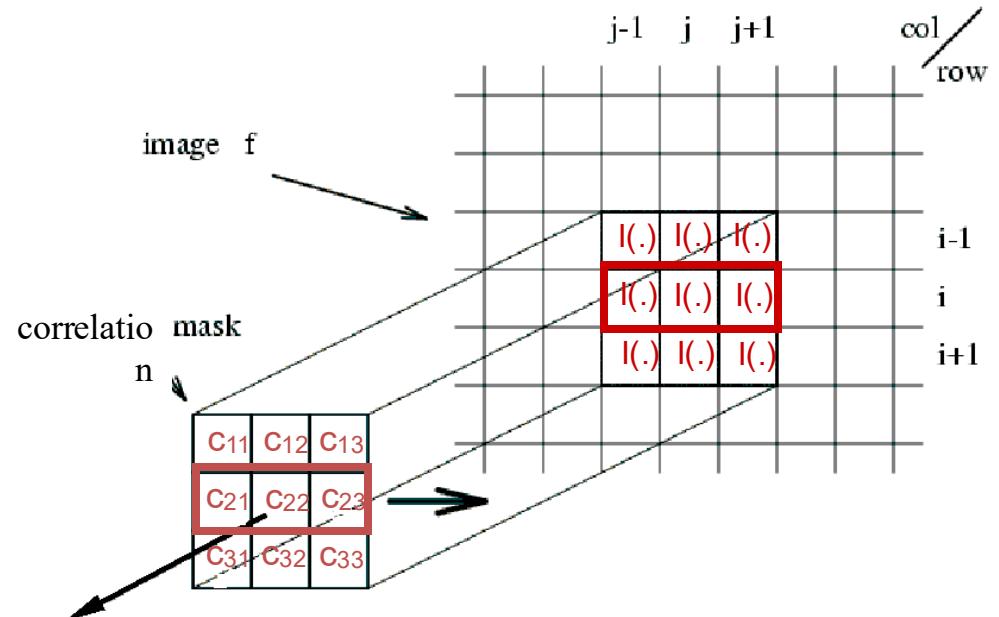
 $g[m,n]$  $h[m,n]$  $=$  $f[m,n]$

Outline

- Neighborhood image processing
 - Connected Component Analysis
 - Mathematical Morphology (Dilation, Erosion, Opening, Closing)
- Linear Image Filtering
 - Correlation and Convolution
 - Frequently used Filters
 - Gaussian Filter
- Nonlinear Image Filtering
 - Median Filter
 - Bilateral Filter
- Frequency Domain Analysis
 - Fourier Transform
 - Sampling and Aliasing

Correlation

Motivation: Template Matching



$$\begin{aligned}
 o(i,j) = & C_{11} l(i-1, j-1) + C_{12} l(i-1, j) + C_{13} \\
 & C_{21} l(i, j-1) + C_{22} l(i, j) + C_{23} \\
 & C_{31} l(i+1, j-1) + C_{32} l(i+1, j) + C_{33}
 \end{aligned}$$

Correlation

- Linear operation of *correlation*:

$$I' = K \circ I$$

$$I'(x, y) = \sum_{(i, j) \in N(x, y)} K(i, j) I(x + i, y + j)$$

- Represent the linear weights as an image, K

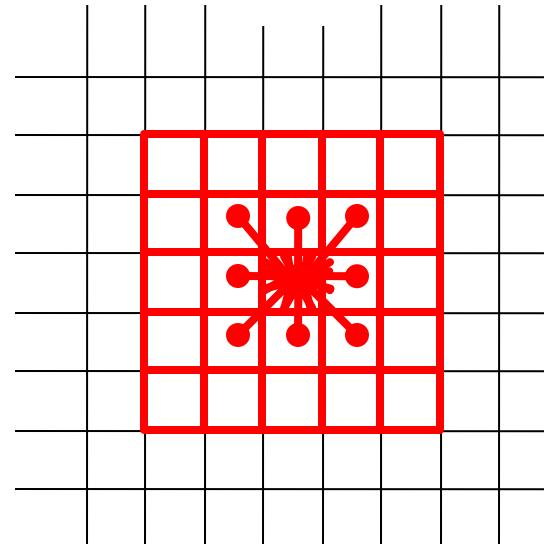
Convolution

Motivation: point spread function



Kernel

K(-1,-1)	K(0,-1)	K(1,-1)
K(-1,0)	K(0,0)	K(1,0)
K(-1,1)	K(0,1)	K(1,1)



$$\begin{aligned} I'(x,y) = & K(1,1)I(x-1,y-1) + K(0,1)I(x,y-1) + K(-1,1)I(x+1,y-1) \\ & + K(1,0)I(x-1,y) + K(0,0)I(x,y) + K(-1,0)I(x+1,y) \\ & + K(1,-1)I(x-1,y+1) + K(0,-1)I(x-1,y) + K(-1,-1)I(x+1,y+1) \end{aligned}$$

Convolution

- Linear operation of *convolution*:

$$I' = K * I$$

$$I'(x, y) = \sum_{(i, j) \in N(x, y)} K(i, j)I(x - i, y - j)$$

- Represent the linear weights as an image, K
- Same as correlation, but with kernel reversed

Correlation vs. Convolution

Correlation

$$I'(x, y) = \sum_{j=-k}^k \sum_{i=-k}^k K(i, j) I(x + i, y + j)$$

Convolution

$$\begin{aligned} I'(x, y) &= \sum_{j=-k}^k \sum_{i=-k}^k K(i, j) I(x - i, y - j) \\ &= \sum_{j=-k}^k \sum_{i=-k}^k K(-i, -j) I(x + i, y + j) \end{aligned}$$

So if $K(i, j) = K(-i, -j)$, then Correlation == Convolution

Mathematical definition of image convolution

- Image $f(x, y)$ and kernel $k(x, y)$
- Continuous form, $(x, y) \in \mathbb{R}^2$

$$k(x, y) * f(x, y) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} k(u, v) f(x - u, y - v) du dv$$

- Discrete form, kernel size (W_k, H_k) , Image size (W_f, H_f) ,

$$k(x, y) * f(x, y) = \sum_{i=0}^{W_k-1} \sum_{j=0}^{H_k-1} k(i, j) f(x - i, y - j)$$

- Commutative:

$$k(x, y) * f(x, y) = f(x, y) * k(x, y)$$

Mathematical definition of cross correlation



- Image $f(x, y)$ and kernel $k(x, y)$
- Continuous form, $(x, y) \in \mathbb{R}^2$

$$k(x, y) \otimes f(x, y) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} k(u, v) f(x + u, y + v) du dv$$

- Discrete form, kernel size (W_k, H_k) , Image size (W_f, H_f) ,

$$k(x, y) \otimes f(x, y) = \sum_{i=0}^{W_k-1} \sum_{j=0}^{H_k-1} k(i, j) f(x + i, y + j)$$

- If $k = f$, the operation is called auto-correlation
- Relation with convolution

$$k(x, y) \otimes f(x, y) = k(-x, -y) * f(x, y)$$

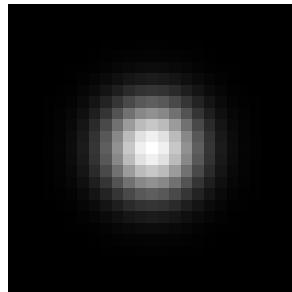
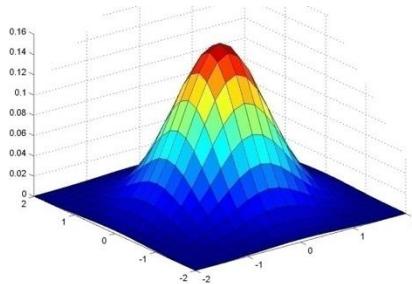
Image filtering

Useful for...

- Enhance images
 - De-noise, resize, increase contrast, etc.
- Extract information from images
 - Texture, edges, distinctive points, etc.
- Detect patterns
 - Template matching

Important linear filter: Gaussian

- Weight contributions of neighboring pixels by nearness

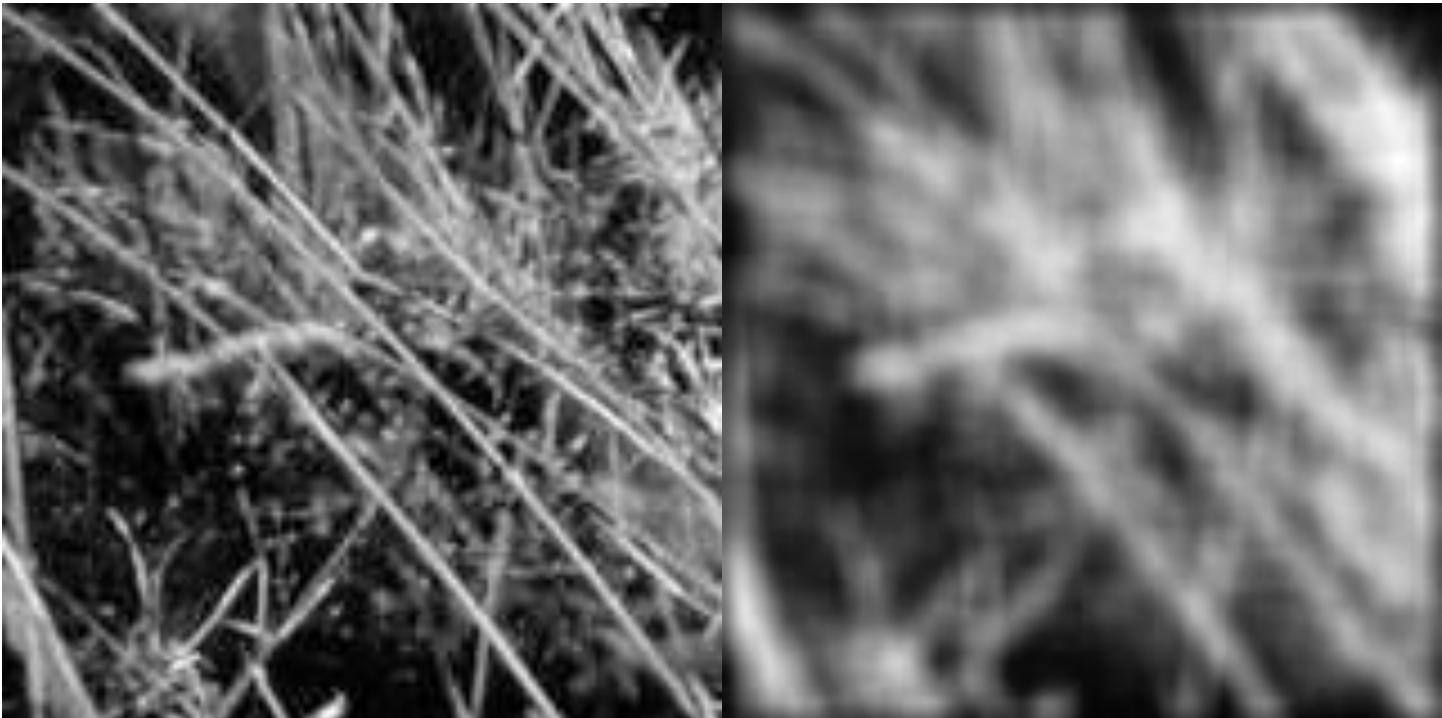


0.003	0.013	0.022	0.013	0.003
0.013	0.059	0.097	0.059	0.013
0.022	0.097	0.159	0.097	0.022
0.013	0.059	0.097	0.059	0.013
0.003	0.013	0.022	0.013	0.003

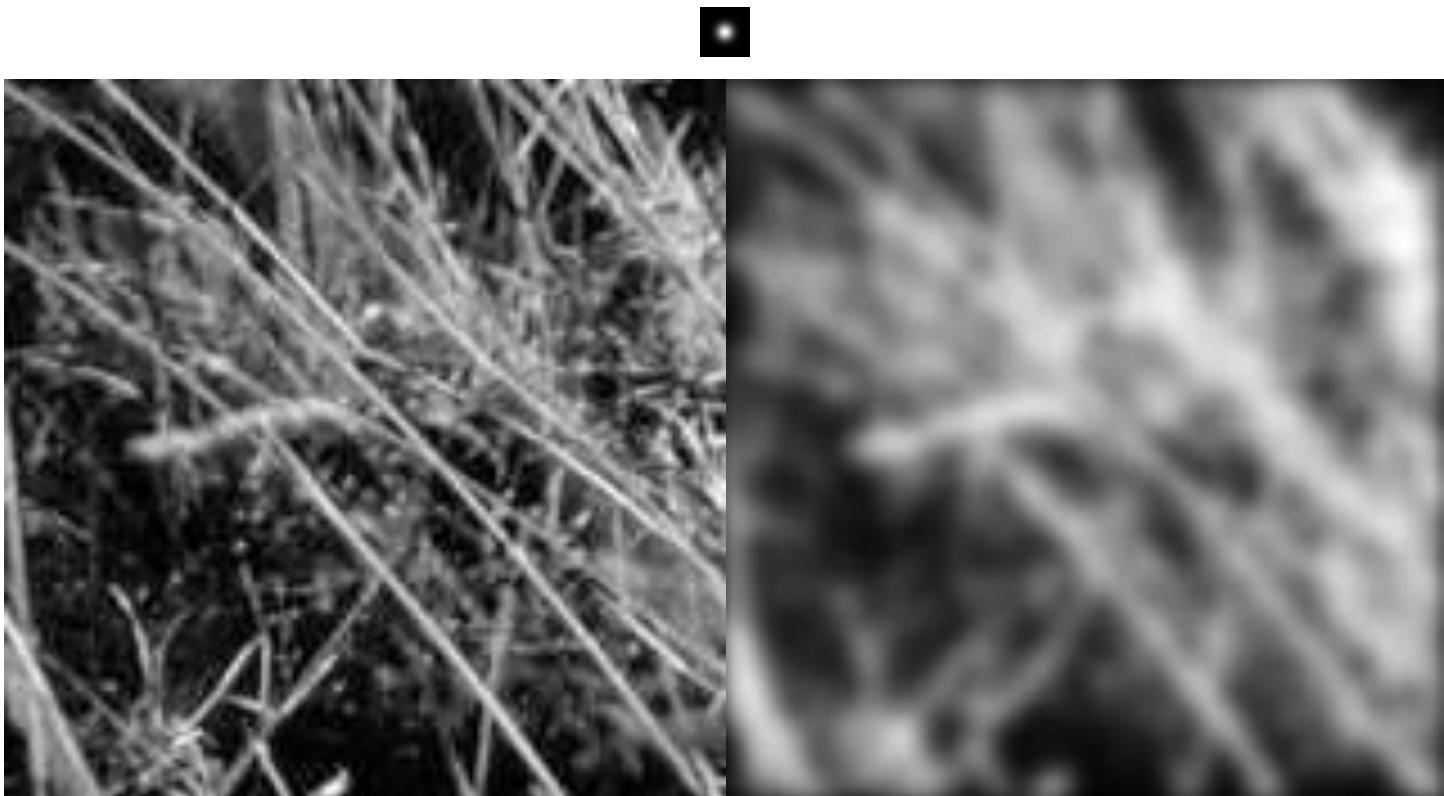
$5 \times 5, \sigma = 1$

$$G_\sigma = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

Smoothing with box filter

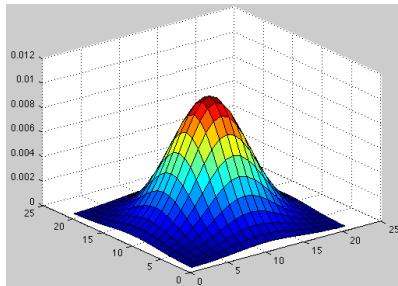


Smoothing with Gaussian filter

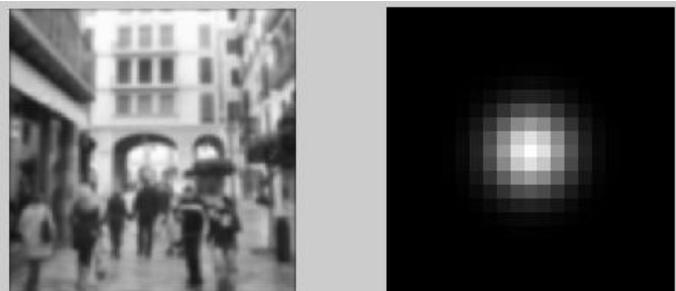


Gaussian filter: change scales

$$G(x, y; \sigma) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}}$$



$\sigma=2$



$\sigma=4$

Properties of Gaussian filter

- Rotational symmetry treats features of all orientations equally (isotropy).
- Efficient: Rule of thumb is kernel width $\geq 5\sigma$
 - Separable
 - Cascadable: Approach to large σ comes from identity

$$G_{\sigma_1} * G_{\sigma_2} = G_{\sqrt{\sigma_1^2 + \sigma_2^2}}$$

Separability of the Gaussian filter

$$\begin{aligned} G_\sigma(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

The 2D Gaussian can be expressed as the product of two functions, one a function of x and the other a function of y

In this case, the two functions are the (identical) 1D Gaussian

Separability example

2D convolution
(center location only)

$$\begin{matrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{matrix} * \begin{matrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{matrix}$$

The filter factors
into a product of 1D
filters:

$$\begin{matrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{matrix} = \begin{matrix} 1 \\ 2 \\ 1 \end{matrix} \times \begin{matrix} 1 & 2 & 1 \end{matrix}$$

Perform convolution
along rows:

$$\begin{matrix} 1 & 2 & 1 \end{matrix} * \begin{matrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{matrix} = \begin{matrix} & 11 & \\ & 18 & \\ & 18 & \end{matrix}$$

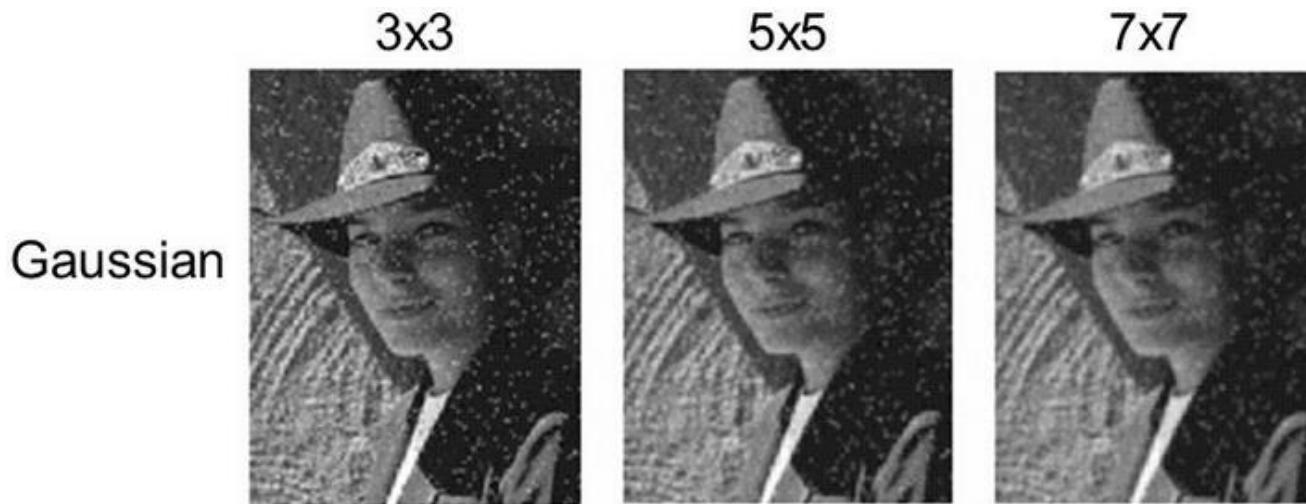
Followed by convolution
along the remaining column:

Outline

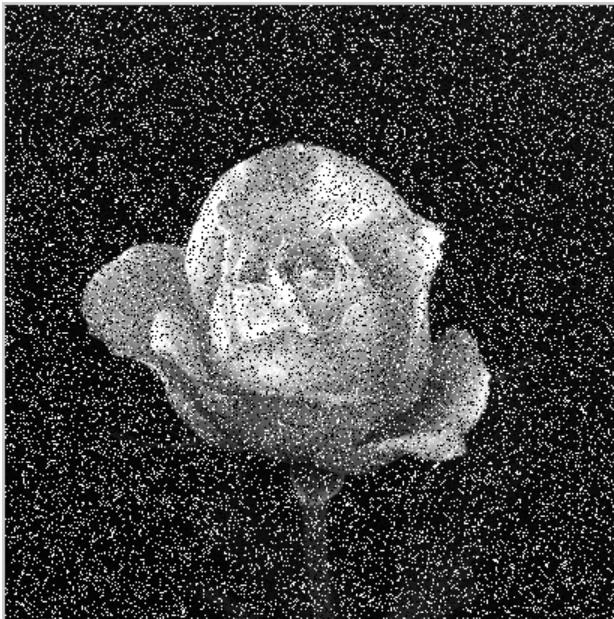
- Neighborhood image processing
 - Connected Component Analysis
 - Mathematical Morphology (Dilation, Erosion, Opening, Closing)
- Linear Image Filtering
 - Correlation and Convolution
 - Frequently used Filters
 - Gaussian Filter
- Nonlinear Image Filtering
 - Median Filter
 - Bilateral Filter
- Frequency Domain Analysis
 - Fourier Transform
 - Sampling and Aliasing

Salt-and-Pepper Noise

- Noise with random occurrence of white and black pixels (signal independent).
Can be due to transmission errors, corrupted pixels, faulted memory locations, etc.



How to better preserve image edges?



Use Median Instead of Average

- Toy Example (1D):

1, 5, 2, 80, 7

1, 2, **5**, 7, 80

Mean = 19

Median = 5

Median Filter

What is the value of the yellow box after 3x3 median filtering?

1	2	2	2	3
2	5	2	4	3
2	2	2	2	3
4	3	7	5	3
3	3	3	1	1

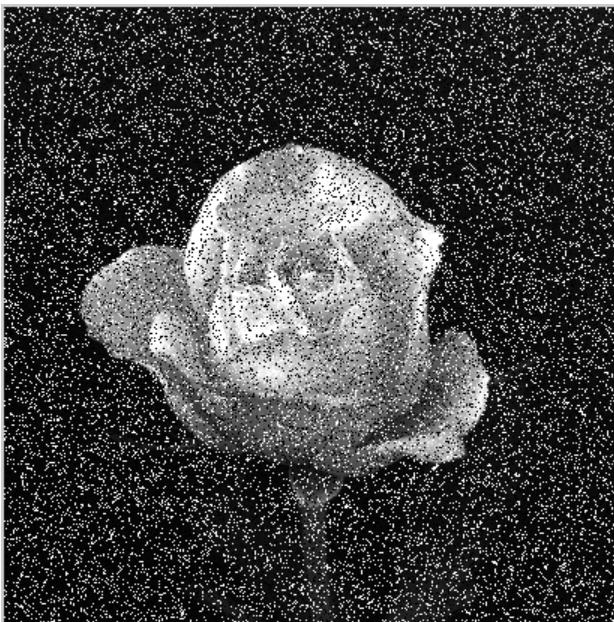
Sorted:

2, 2, 2,

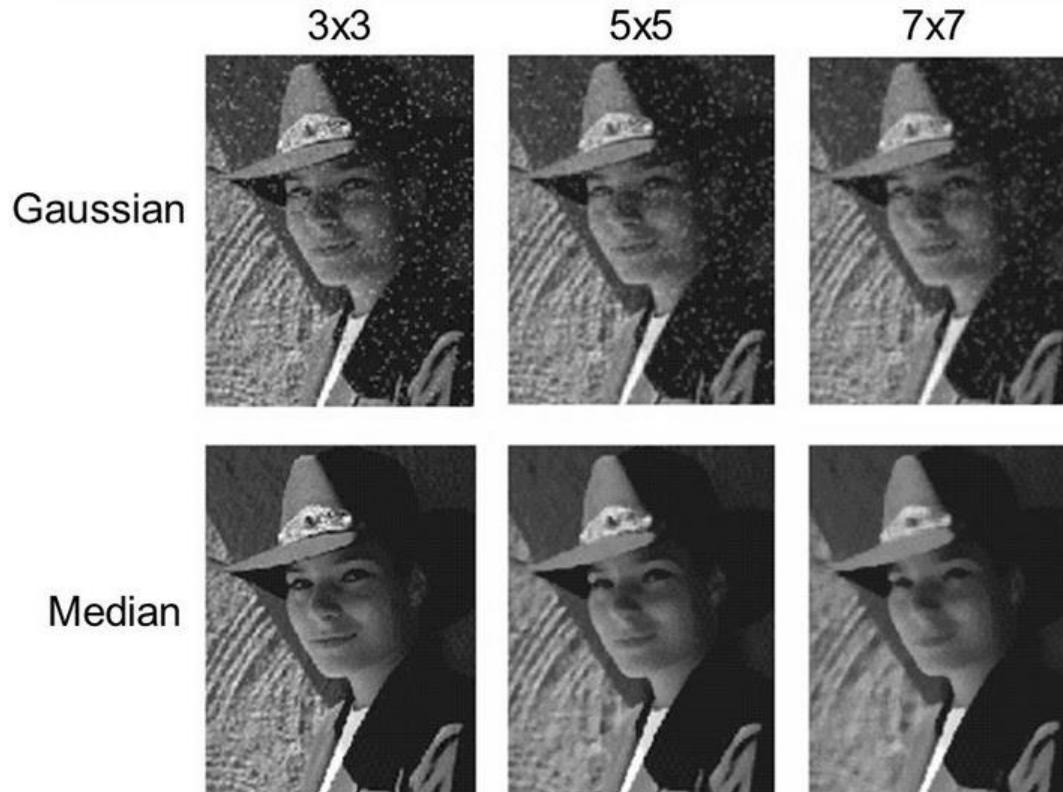
2, 3, 4,

5, 5, 7

Median Filter Result



Gaussian vs. Median Filtering



Bilateral Filter

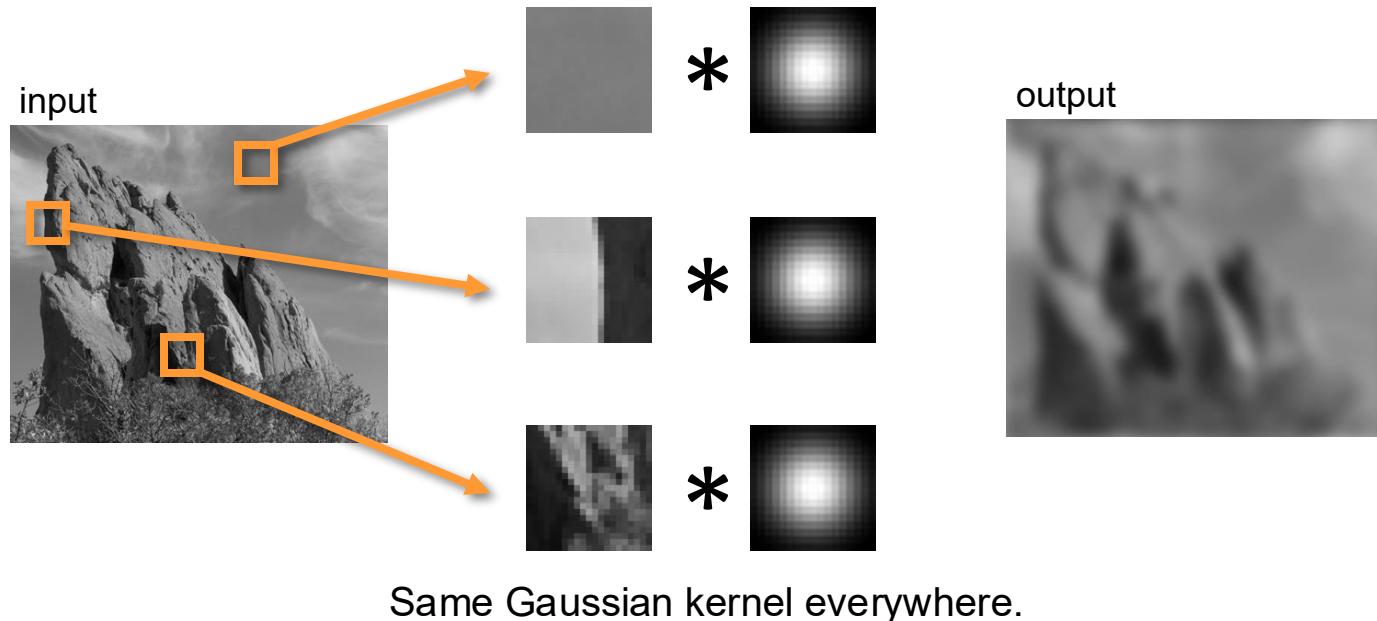
Bilateral Filter (Tomasi et al.1998)

- Goal: Smooths regions while preserving edges



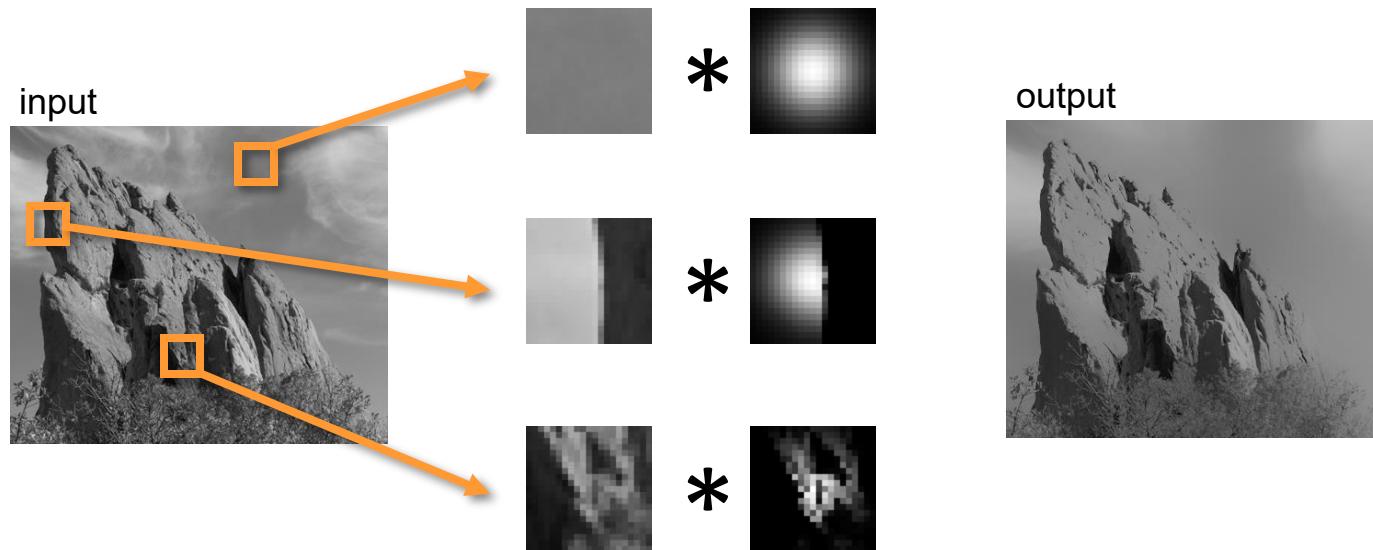
Bilateral Filter

Motivation: In Gaussian filtering the kernel should blur less across image edges



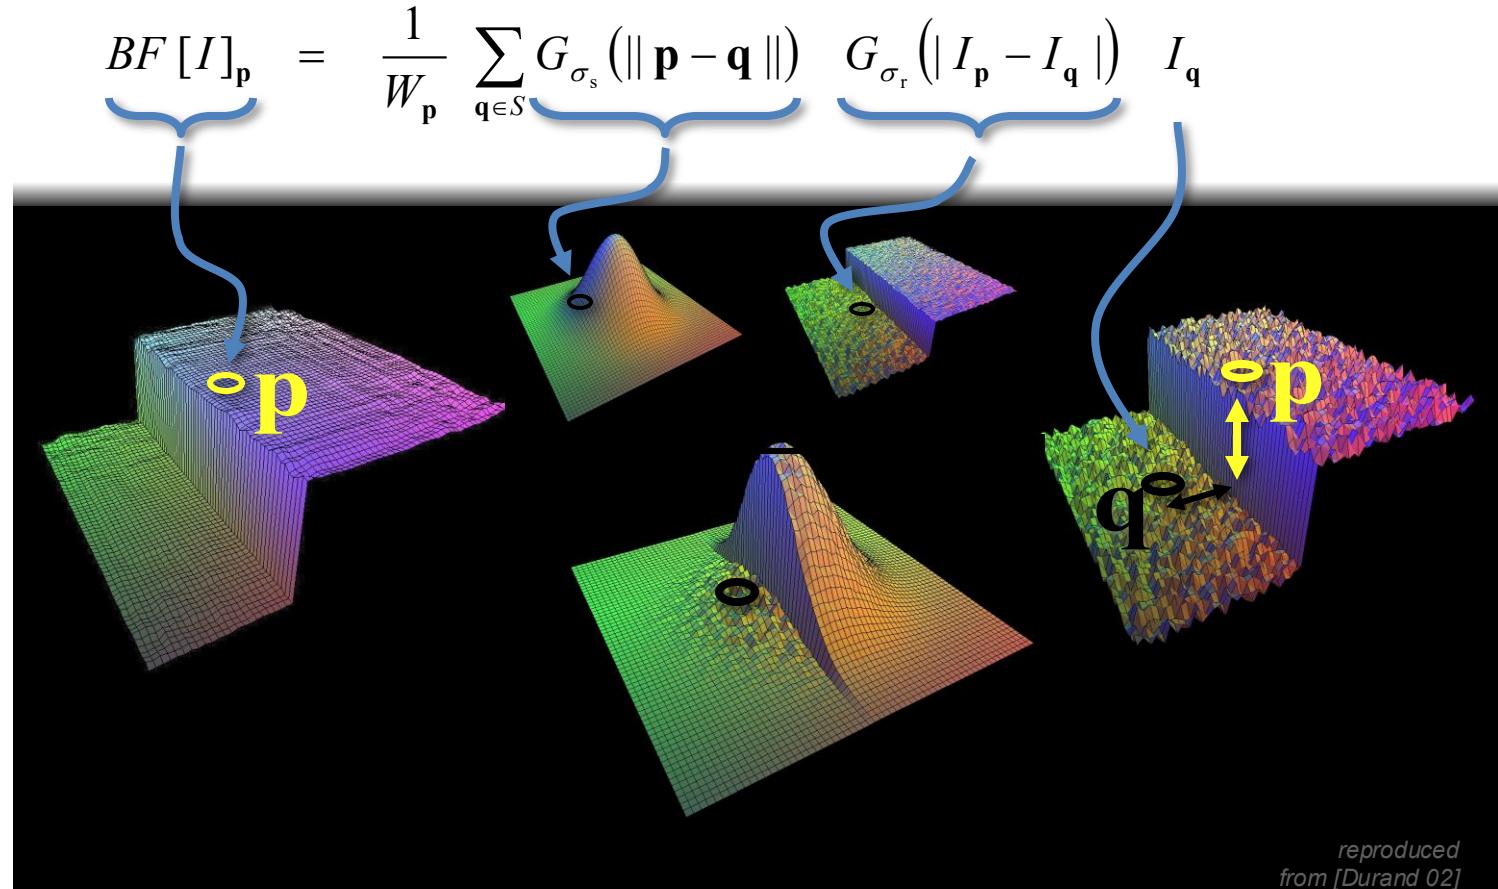
Bilateral Filter

Motivation: No averaging across edges.



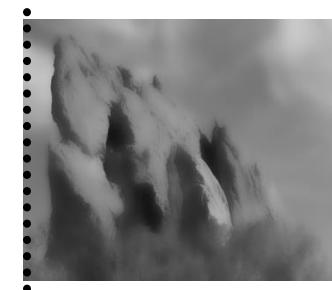
The kernel shape depends on the image content.

Bilateral Filter on a Height Field



Varying the Space Parameter

input

 $\sigma_s = 2$  $\sigma_r = 0.1$ $\sigma_s = 6$  $\sigma_s = 18$  $\sigma_r = 0.25$  $\sigma_r = \infty$ (Gaussian blur)

Outline

- Neighborhood image processing
 - Connected Component Analysis
 - Mathematical Morphology (Dilation, Erosion, Opening, Closing)
- Linear Image Filtering
 - Correlation and Convolution
 - Frequently used Filters
 - Gaussian Filter
- Nonlinear Image Filtering
 - Median Filter
 - Bilateral Filter
- Frequency Domain Analysis
 - Fourier Transform
 - Sampling and Aliasing

Let's Enhance that Image!



Image Scaling

This image is too big, can we generate a half-sized version?

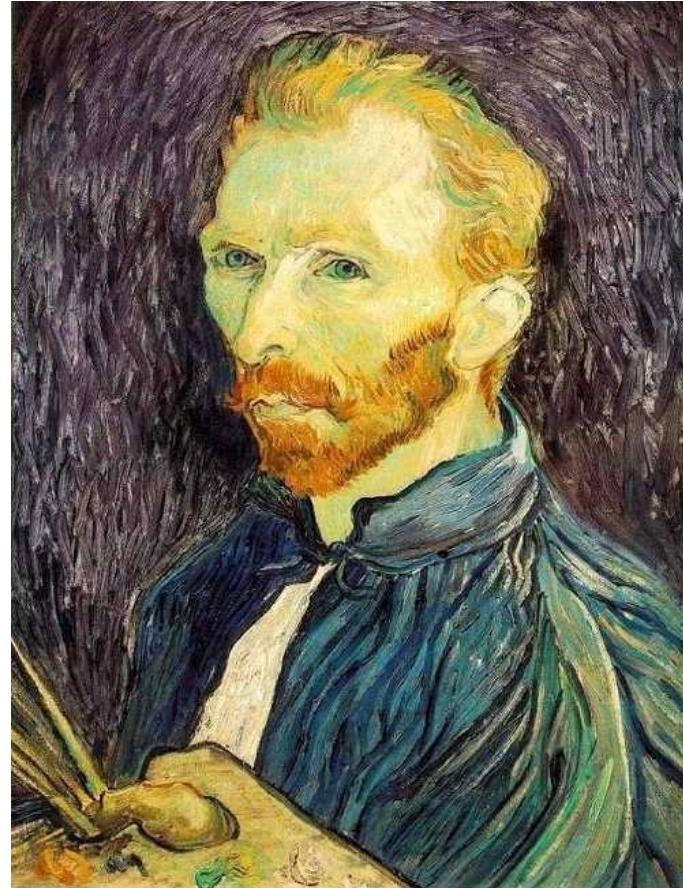
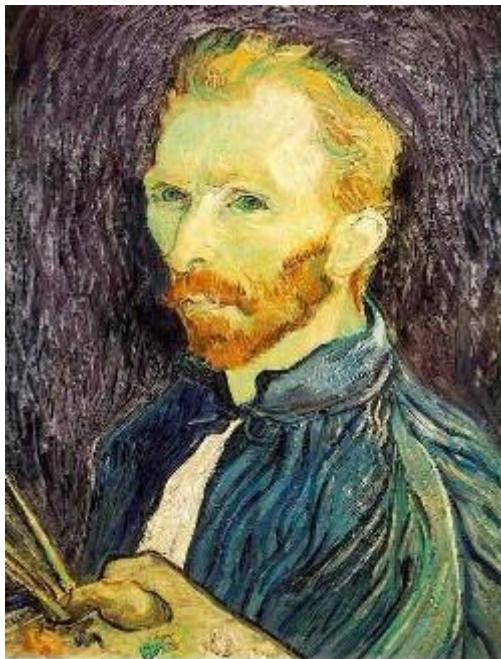


Image sub-sampling



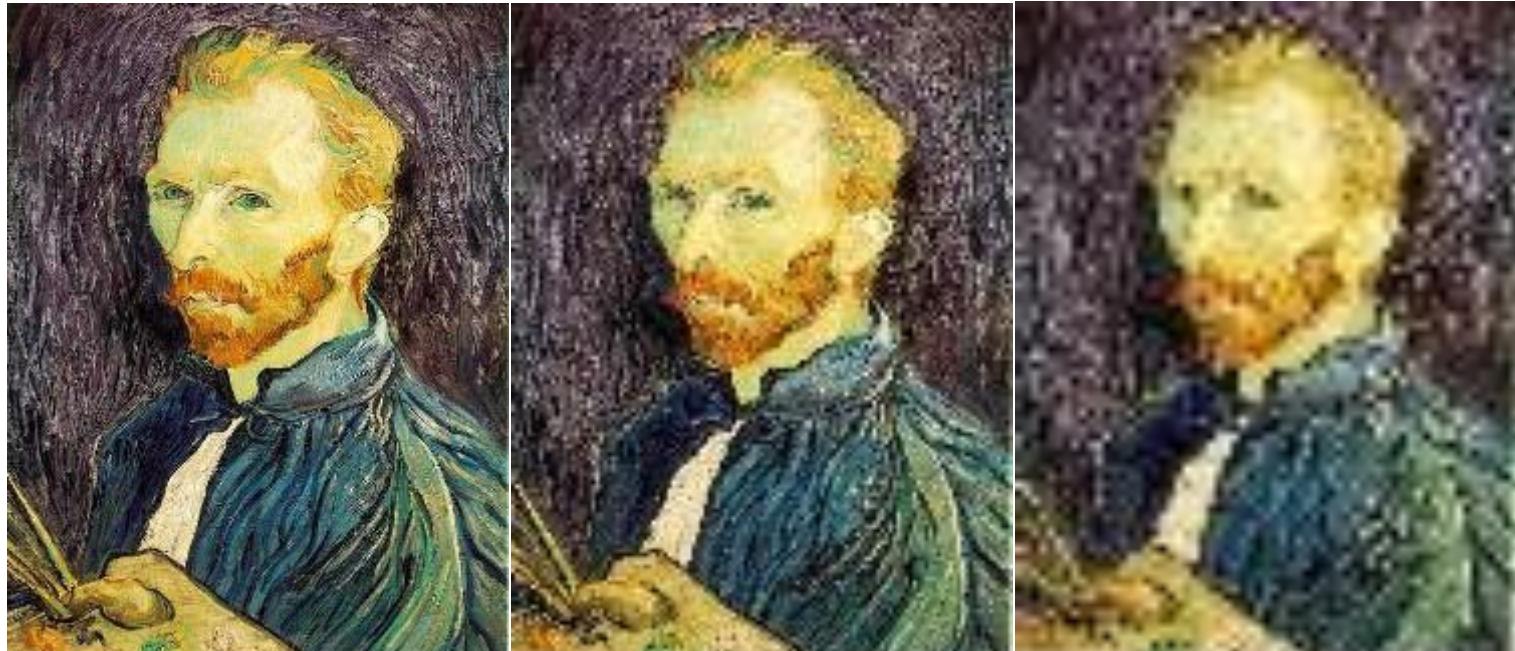
1/4



1/8

Throw away every other row and
column to create a 1/2 size image
- called *image sub-sampling*

Image sub-sampling



1/2

1/4 (2x zoom)

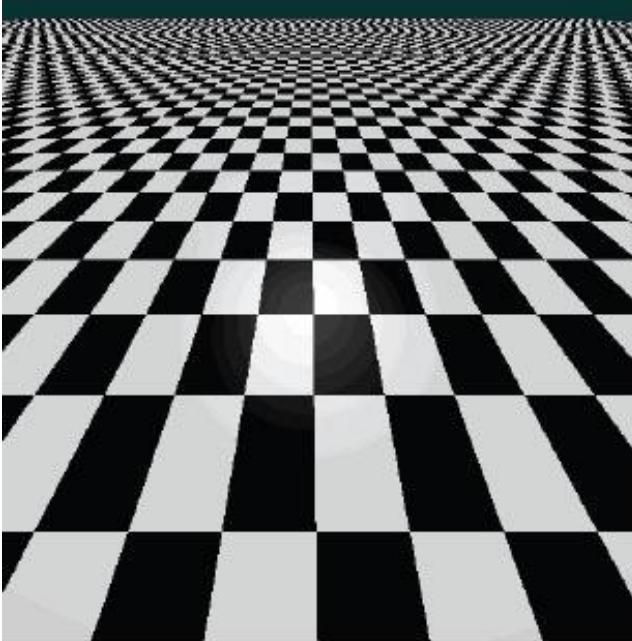
1/8 (4x zoom)

Why does this look so crulty?

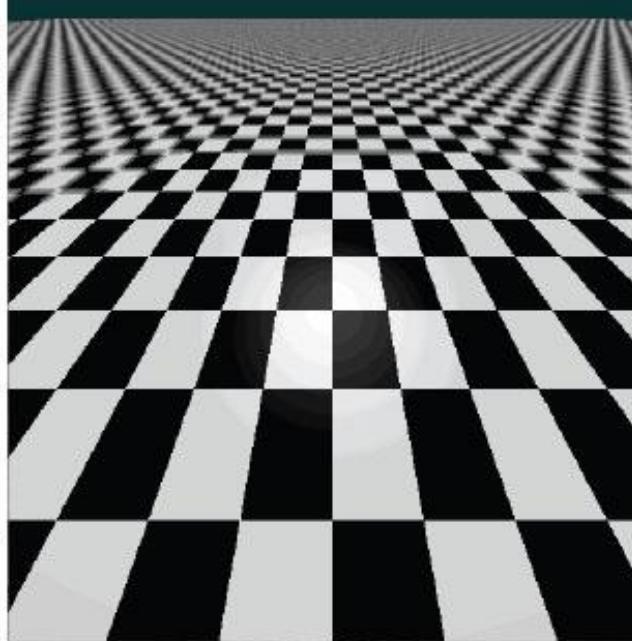
Image sub-sampling – another example



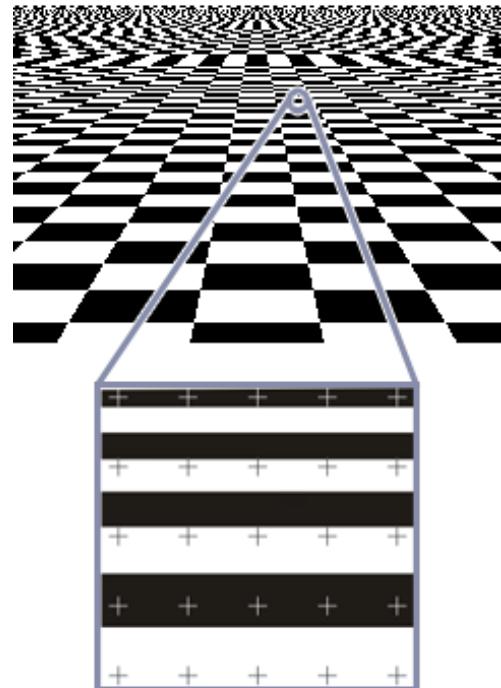
Image sub-sampling



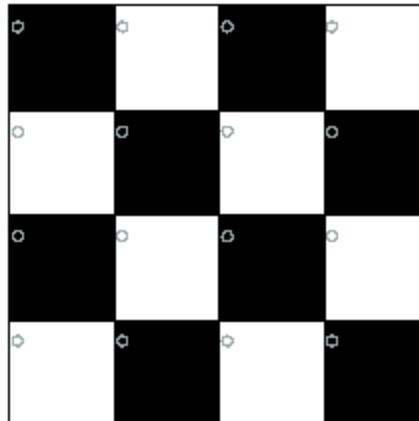
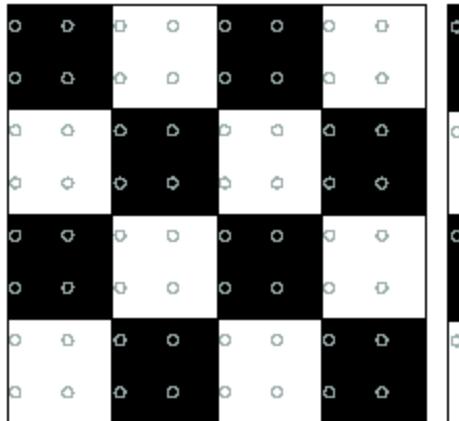
aliasing effects



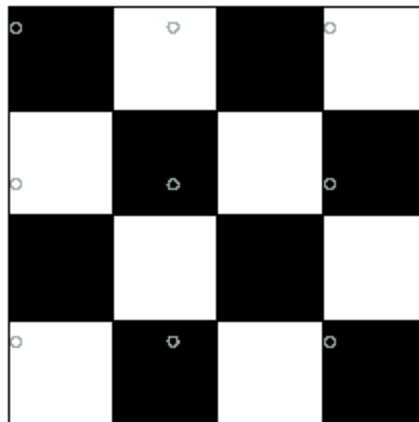
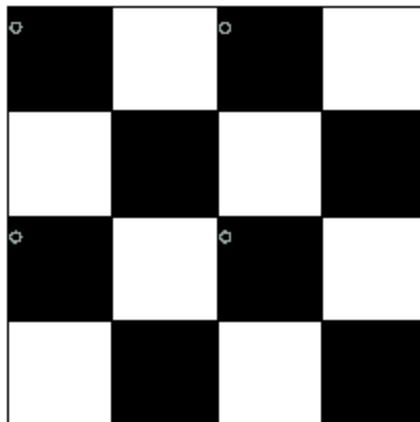
anti-aliasing by over-sampling



Nyquist limit – 2D example

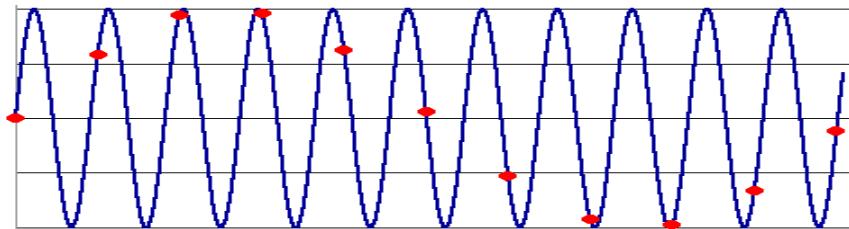


Good sampling



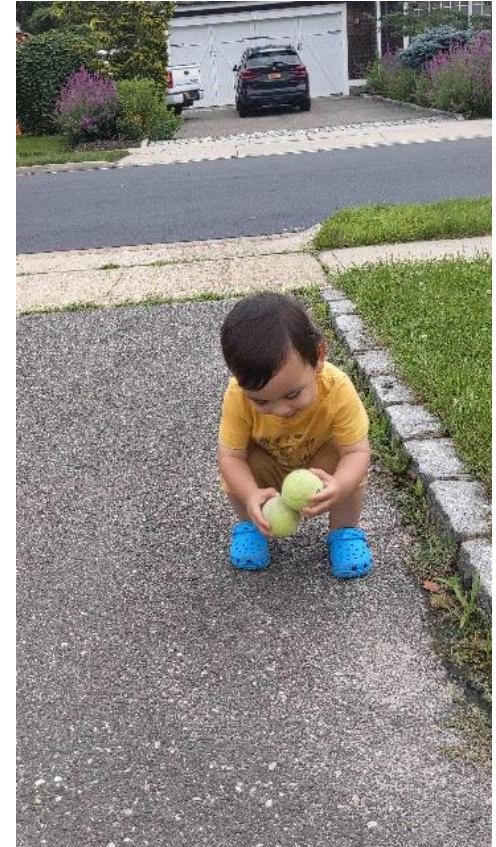
Bad sampling

Aliasing

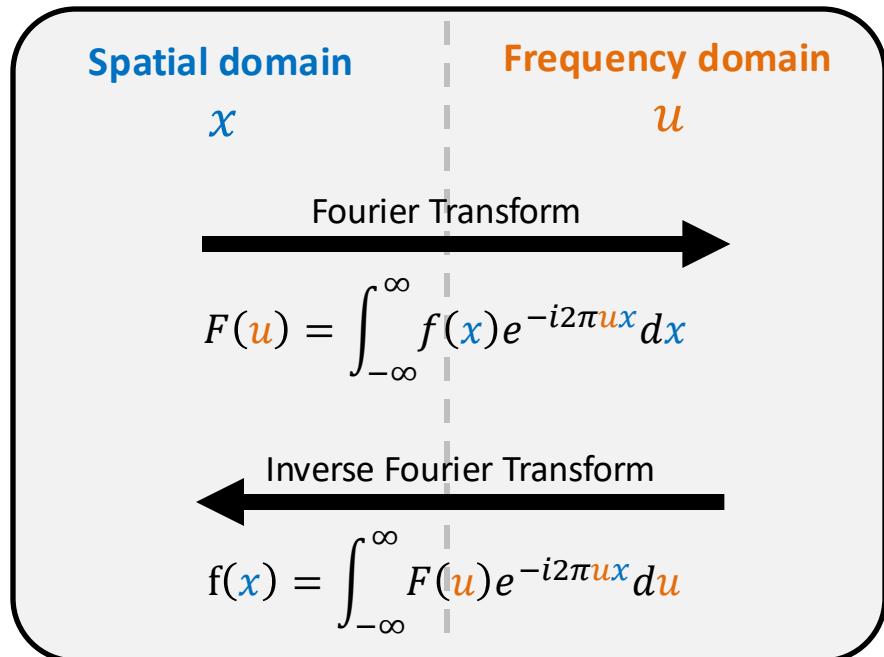


- Occurs when your sampling rate is not high enough to capture the amount of detail in your image
- Can give you the wrong signal/image—an *alias*
- To do sampling right, need to understand the structure of your signal/image
- To avoid aliasing:
 - sampling rate $\geq 2 * \text{max frequency in the image}$
 - said another way: $\geq \text{two samples per cycle}$
 - This minimum sampling rate is called the **Nyquist rate**

Aliasing in Practice



Fourier Transform



$$e^{i\theta} = \cos\theta + i \sin\theta$$

$$i = \sqrt{-1}$$

Note: This is the continuous case for 1D.

Fourier
Analysis

Fourier
Synthesis



Jean-Baptiste Joseph Fourier
(1768-1830)

Fourier Transform

Fourier Transform is expressed with **complex** numbers.

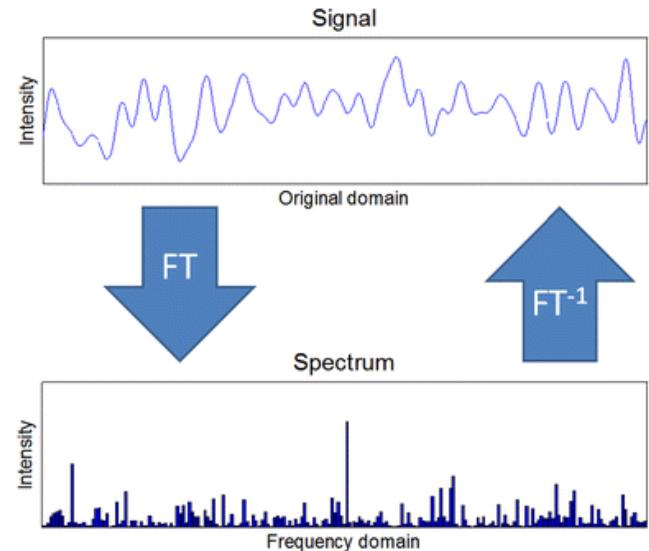
Here, $F(u)$ holds both the **Amplitude** and **Phase** of the sinusoid of frequency u .

$$F(u) = \int_{-\infty}^{\infty} f(x)e^{-i2\pi ux} dx$$

$$F(u) = \Re\{F(u)\} + i \cdot \Im\{F(u)\}$$

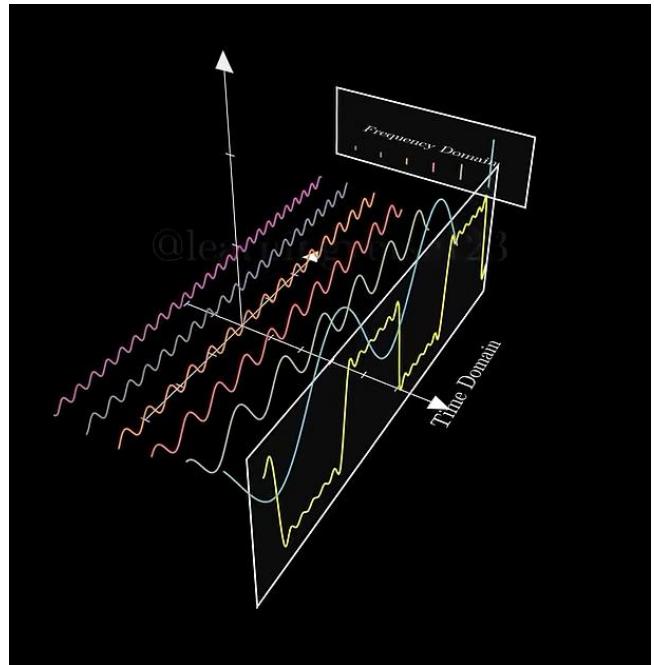
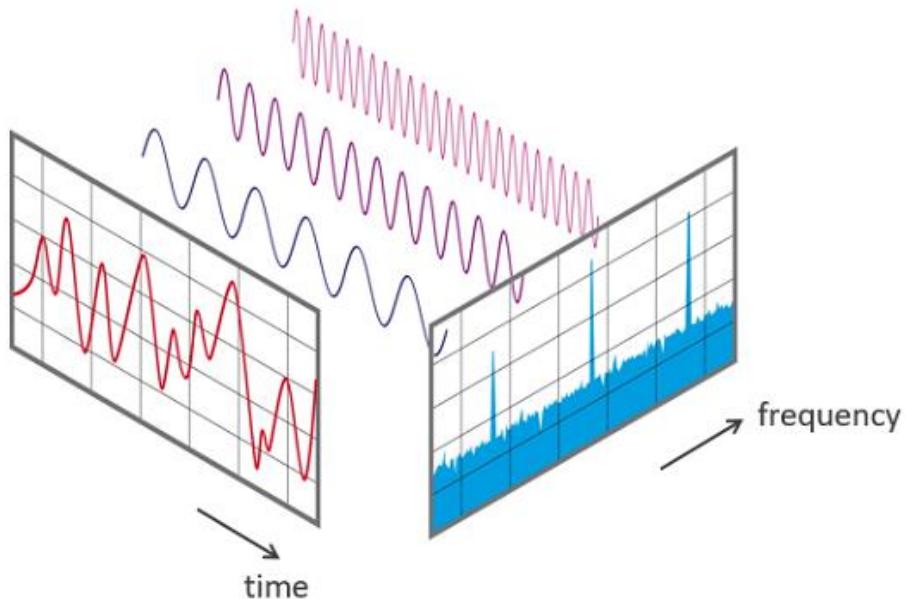
Amplitude: $A(u) = \sqrt{\Re\{F(u)\}^2 + \Im\{F(u)\}^2}$

Phase: $\phi(u) = \text{atan2}(\Re\{F(u)\}, \Im\{F(u)\})$



Fourier Analysis

Decompose a given function into its frequencies and phases

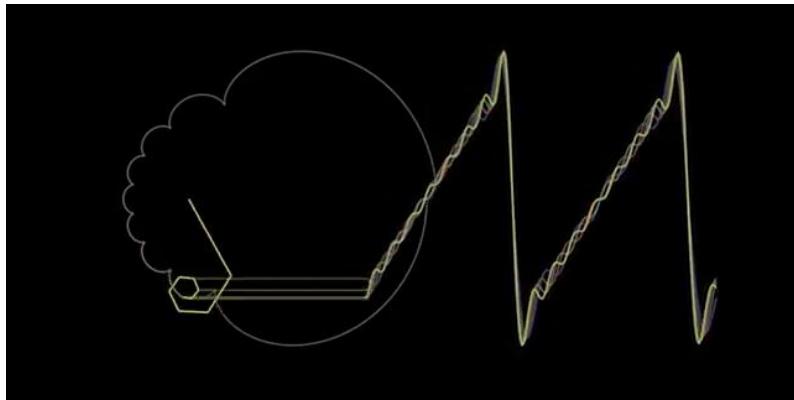


[learningverse123 - [YouTube](#)]

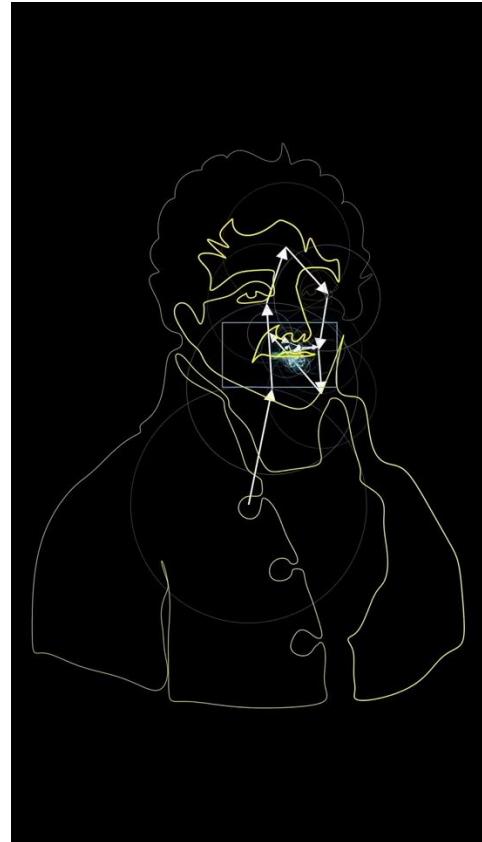
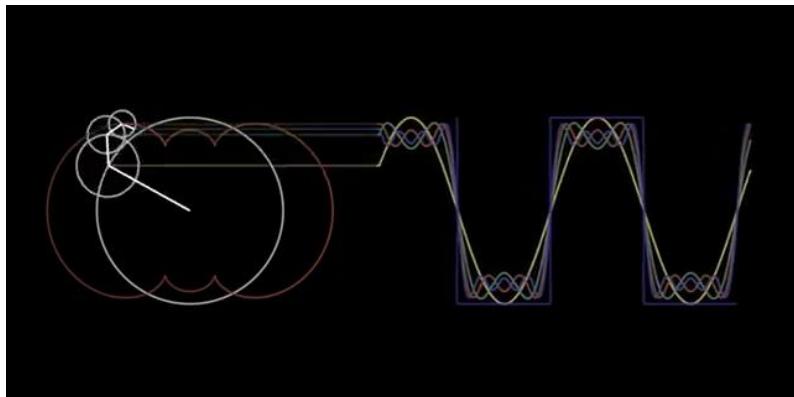
Fourier Synthesis

Compose a function from chosen frequencies and phases.

- Sawtooth:



- Square wave:



Fourier Transform

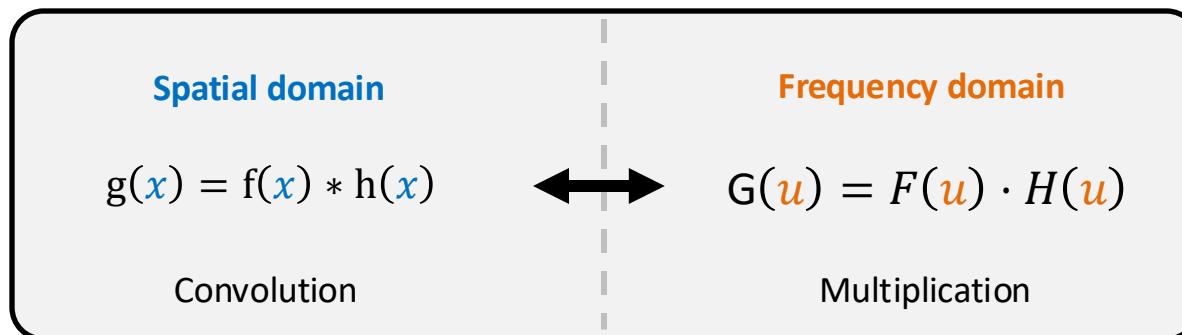
Important Facts

- The Fourier transform (FT) is linear
- There is an inverse FT
- The FT of a Gaussian is a Gaussian

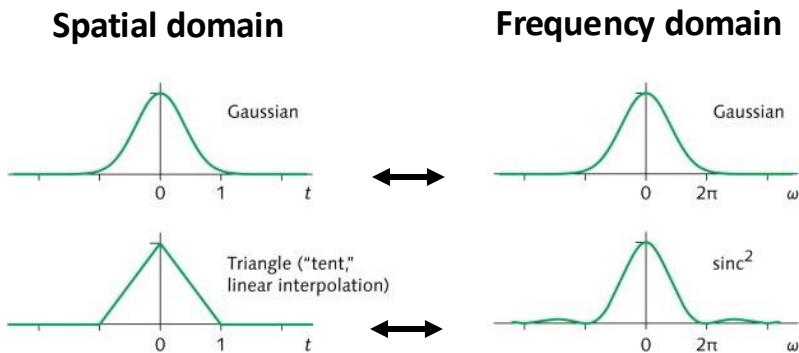
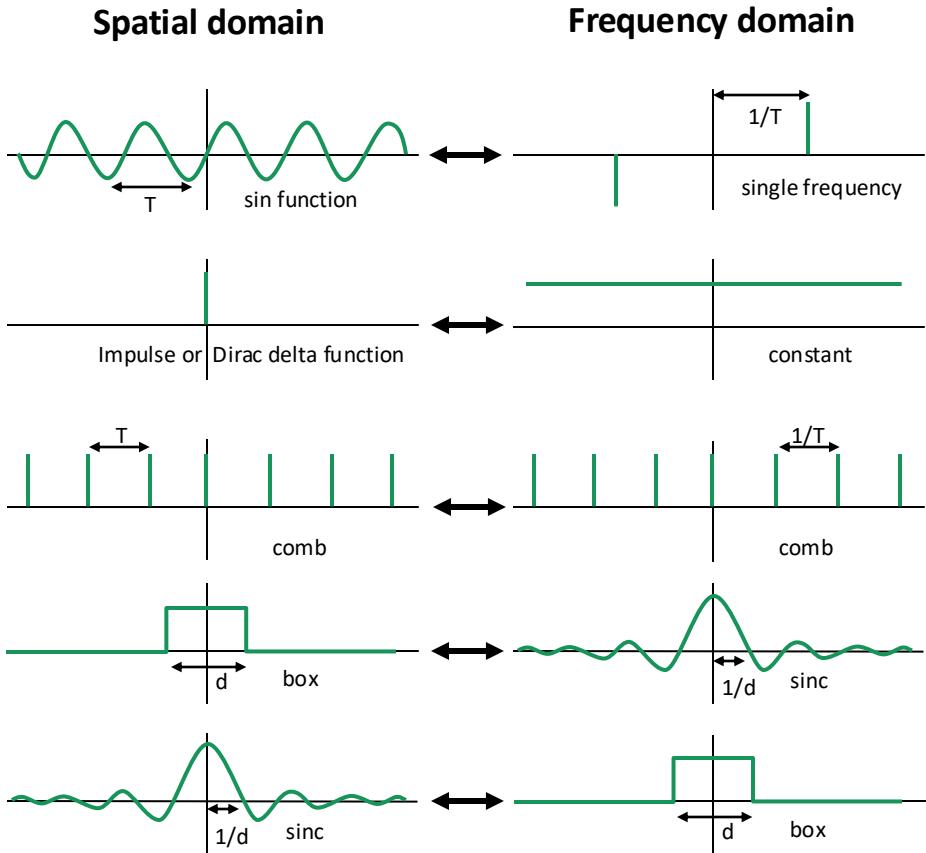
Convolution Theorem

The convolution theorem

1. The Fourier transform of the convolution of two functions is the product of their Fourier transforms
2. The Fourier transform of the product of two functions is the convolution of the Fourier transforms



Fourier Transform of Important Functions



Dirac delta function:

$$\delta(x) = \begin{cases} 0 & \text{if } x \neq 0 \\ \infty & \text{otherwise} \end{cases}$$

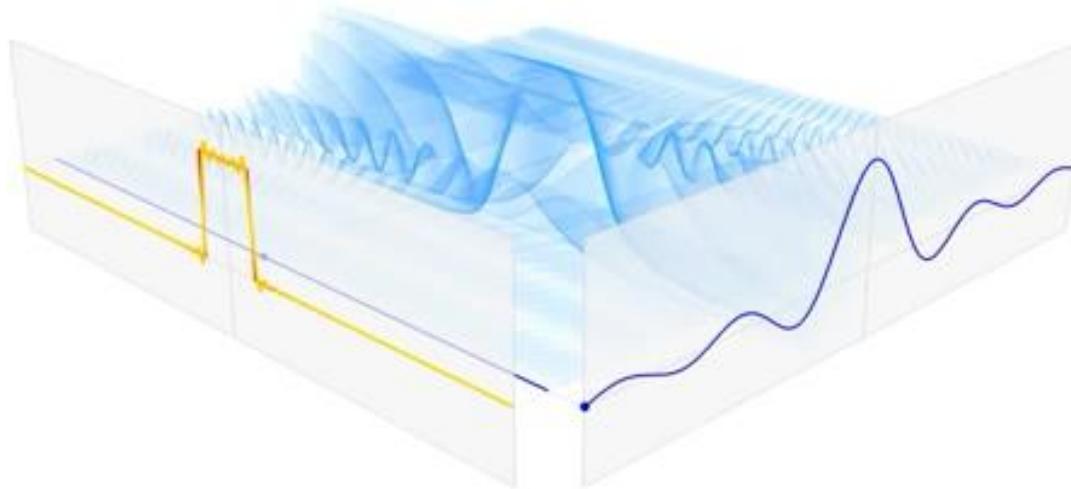
with $\int_{-\infty}^{\infty} \delta(x) dx = 1$

Sinc function:

$$\text{sinc}(x) = \frac{\sin(x)}{x}$$

Fourier Transforms of sinc and rect Functions

$$\int_{-\infty}^{\infty} f(x) e^{-2\pi i x \xi} dx$$



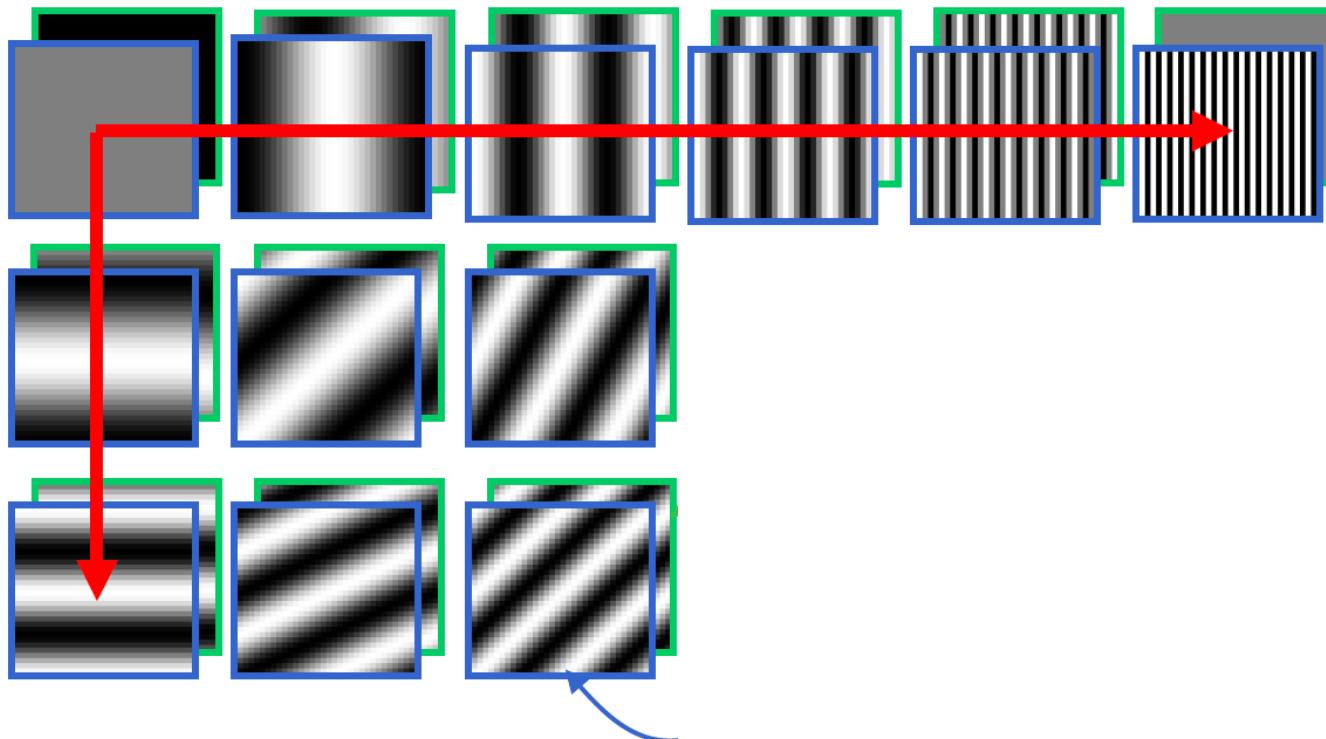
The Discrete Fourier Transform (2D)

The *Discrete Fourier Transform* of $f(x, y)$, for $x = 0, 1, 2 \dots M-1$ and $y = 0, 1, 2 \dots N-1$, denoted by $F(u, v)$, is given by the equation:

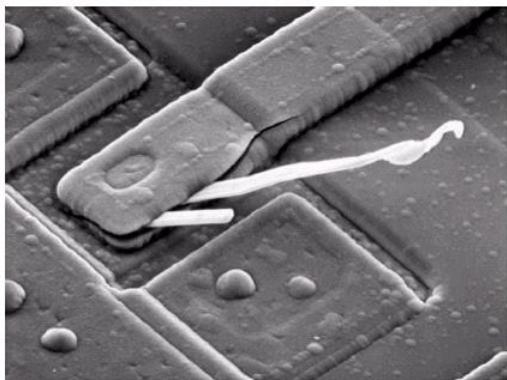
$$F(u, v) = \sum_{x=0}^{M-1} \sum_{y=0}^{N-1} f(x, y) e^{-j2\pi(ux/M + vy/N)}$$

for $u = 0, 1, 2 \dots M-1$ and $v = 0, 1, 2 \dots N-1$.

Represent Images by a set of Fourier basis functions

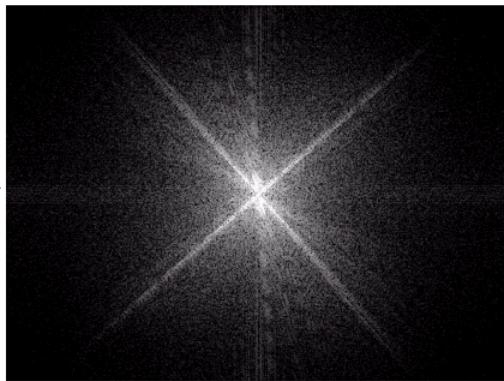


DFT & Images (cont...)



Scanning electron microscope
image of an integrated circuit
magnified ~2500 times

DFT



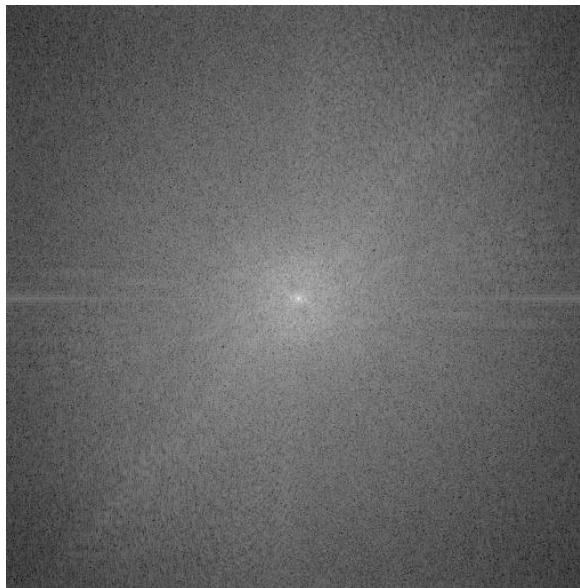
Fourier spectrum of the image

Fourier Space Visualization

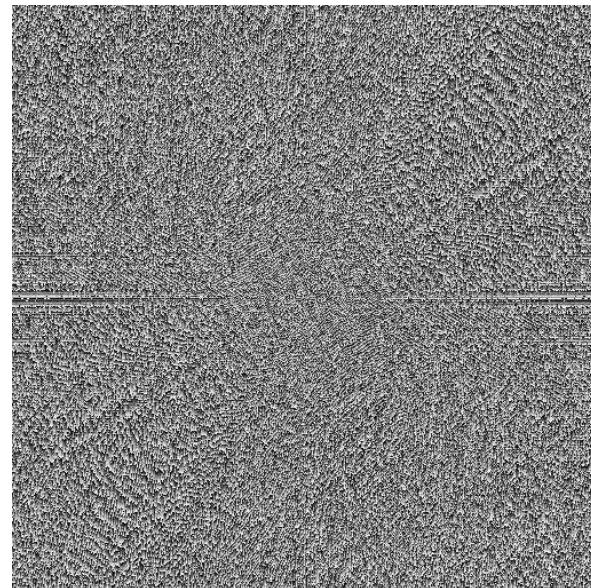
Input



Magnitude



Phase

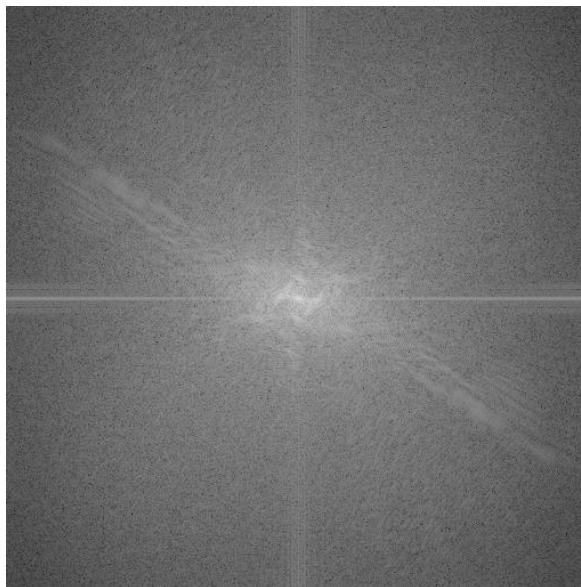


Fourier Space Visualization

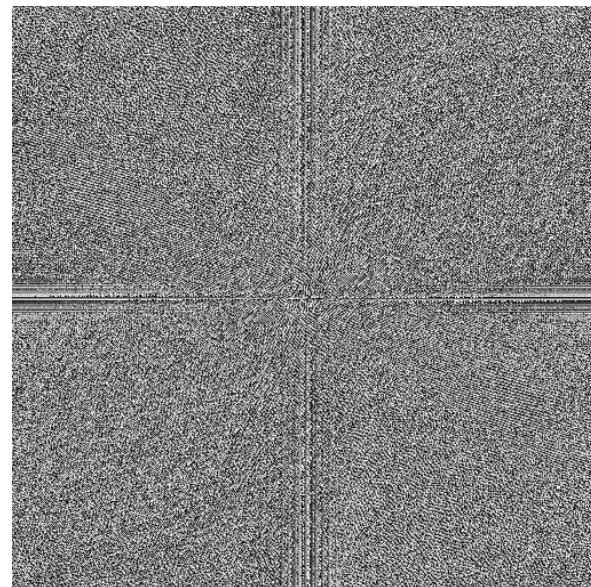
Input



Magnitude



Phase

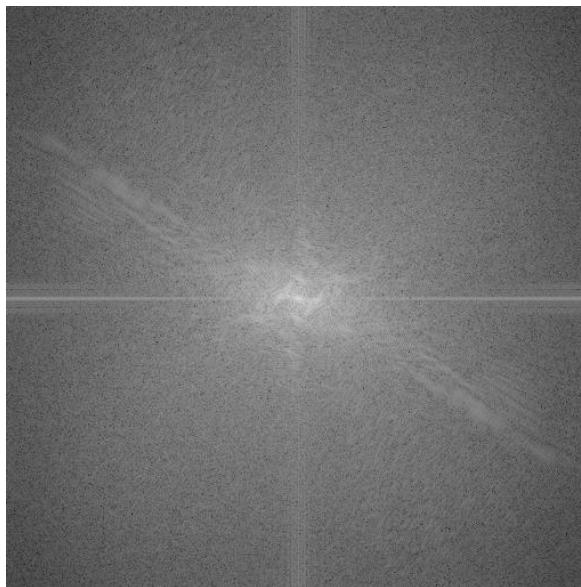


Fourier Space Visualization

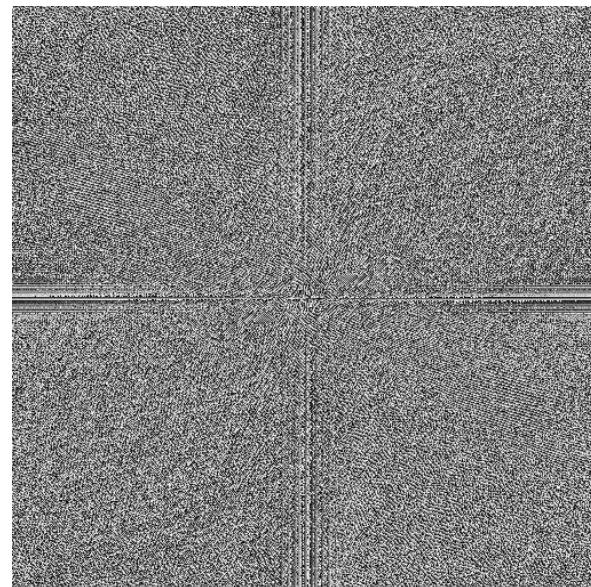
Input



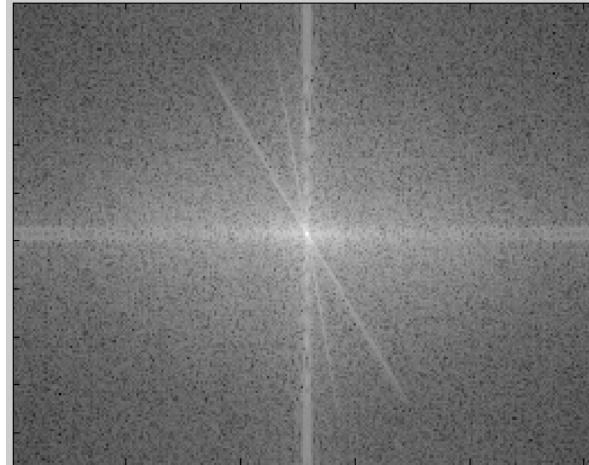
Magnitude



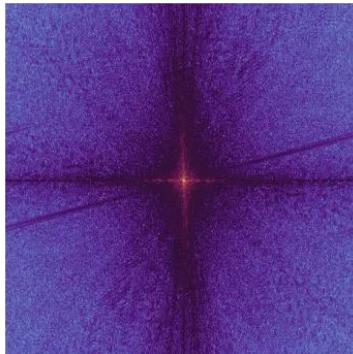
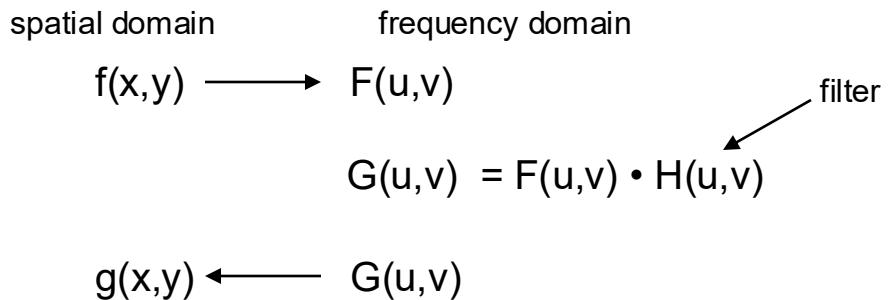
Phase



Example



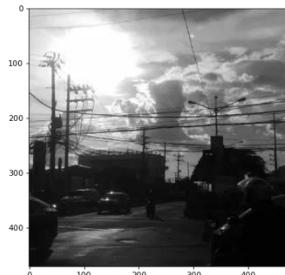
Low Pass Filter



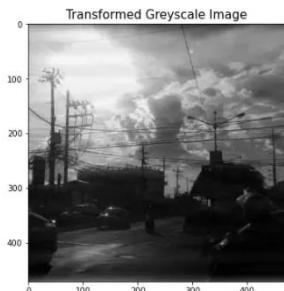
Low Pass Filter

spatial domain

$$f(x,y)$$



Dark horizontal lines

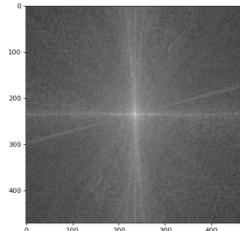


Thinner horizontal lines

$$g(x,y)$$

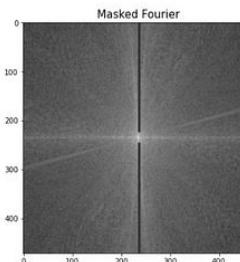
Frequency domain

$$F(u,v)$$



filter

$$G(u,v) = F(u,v) \cdot H(u,v)$$

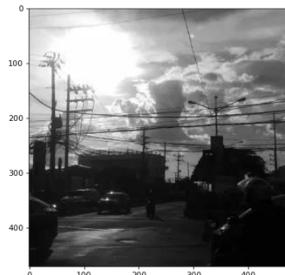


$$G(u,v)$$

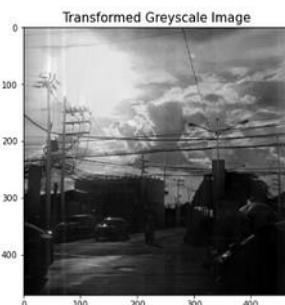
Low Pass Filter

spatial domain

$$f(x,y)$$



Dark vertical poles

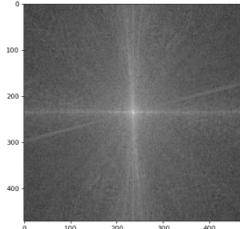


Lighter vertical poles

$$g(x,y)$$

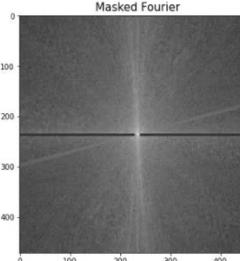
Frequency domain

$$F(u,v)$$



$$G(u,v) = F(u,v) \cdot H(u,v)$$

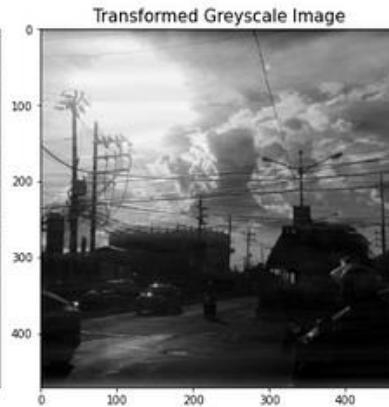
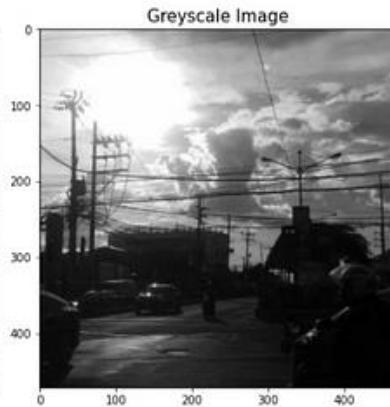
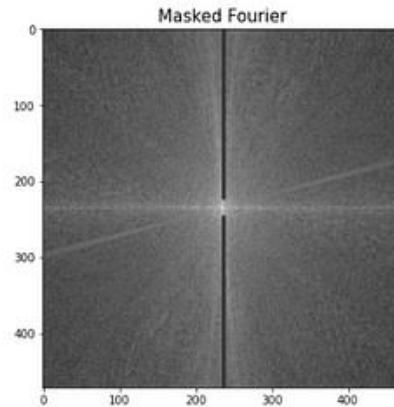
filter



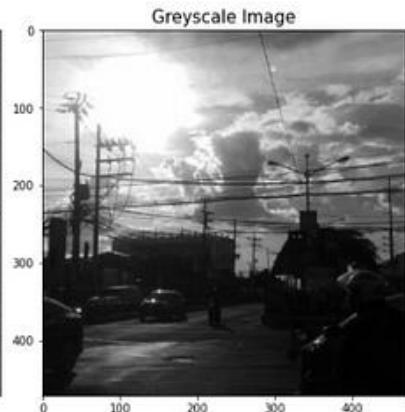
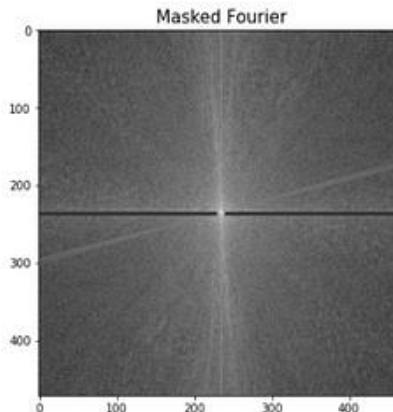
$$G(u,v)$$



Low Pass Filter

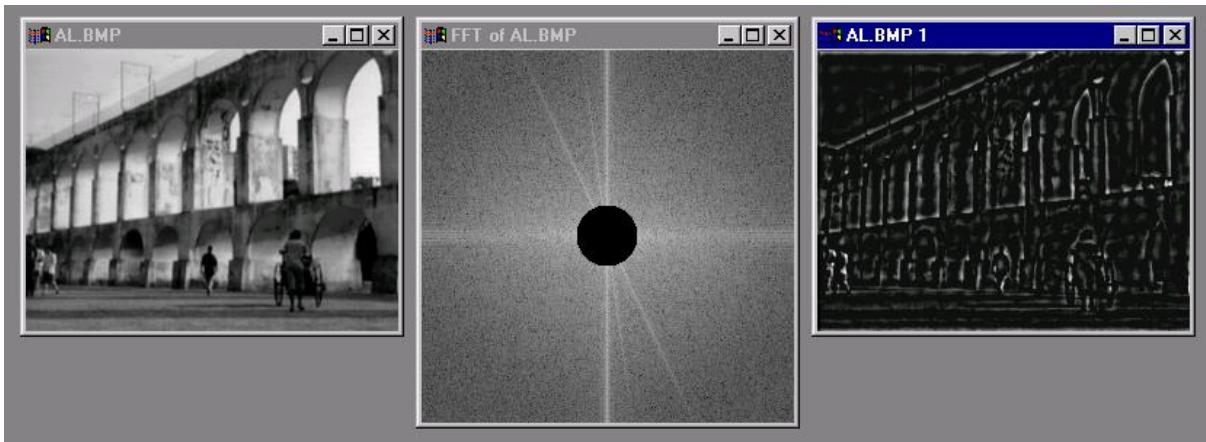


Thinner horizontal lines

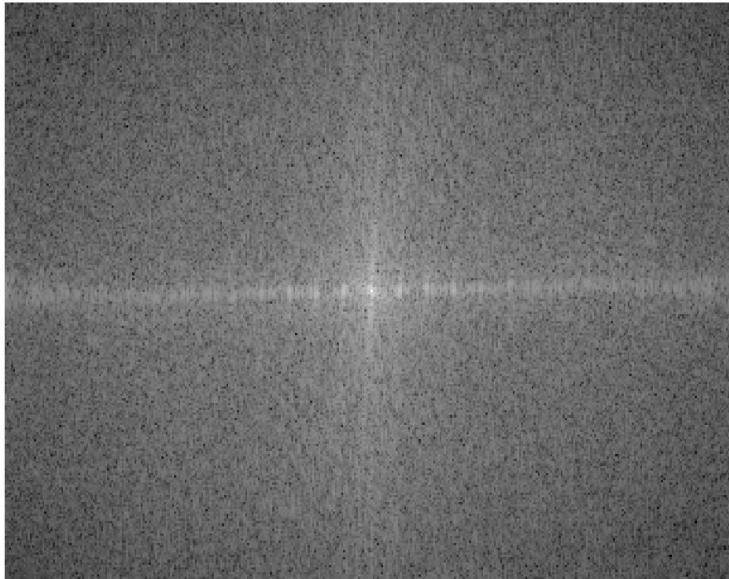


Lighter vertical poles

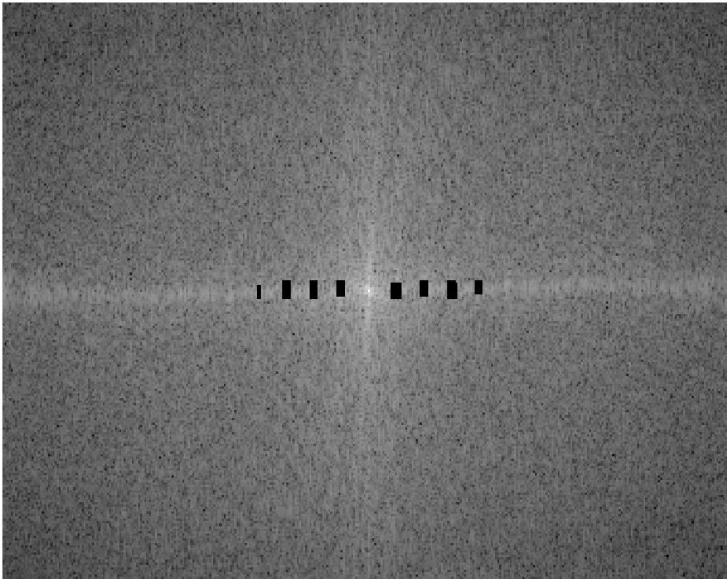
Low and High Pass filtering



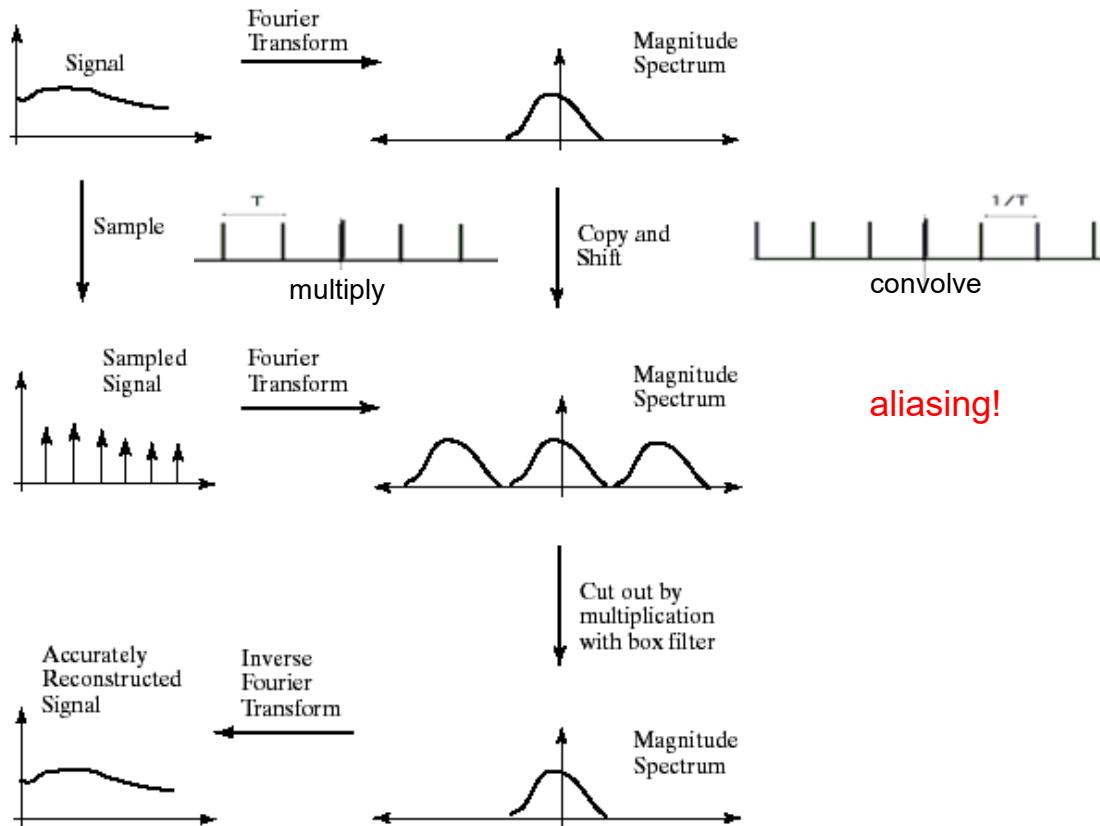
Frequency Filtering



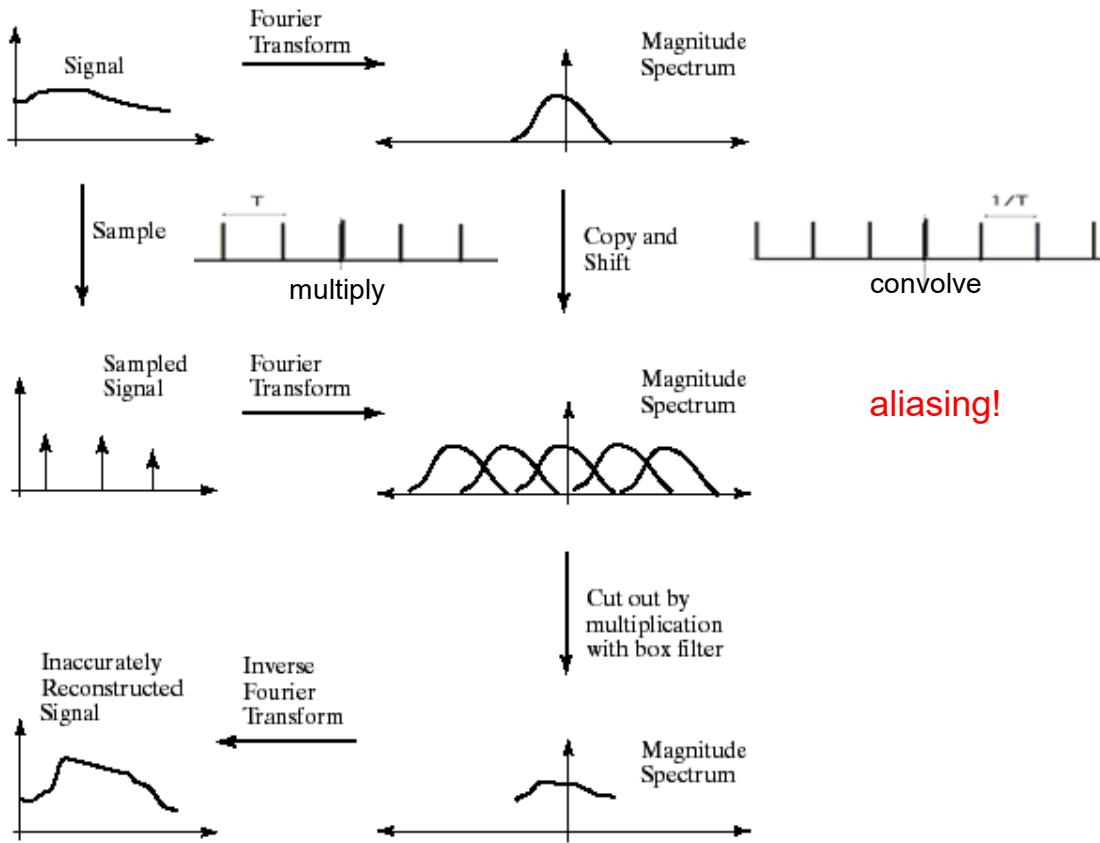
Frequency Filtering



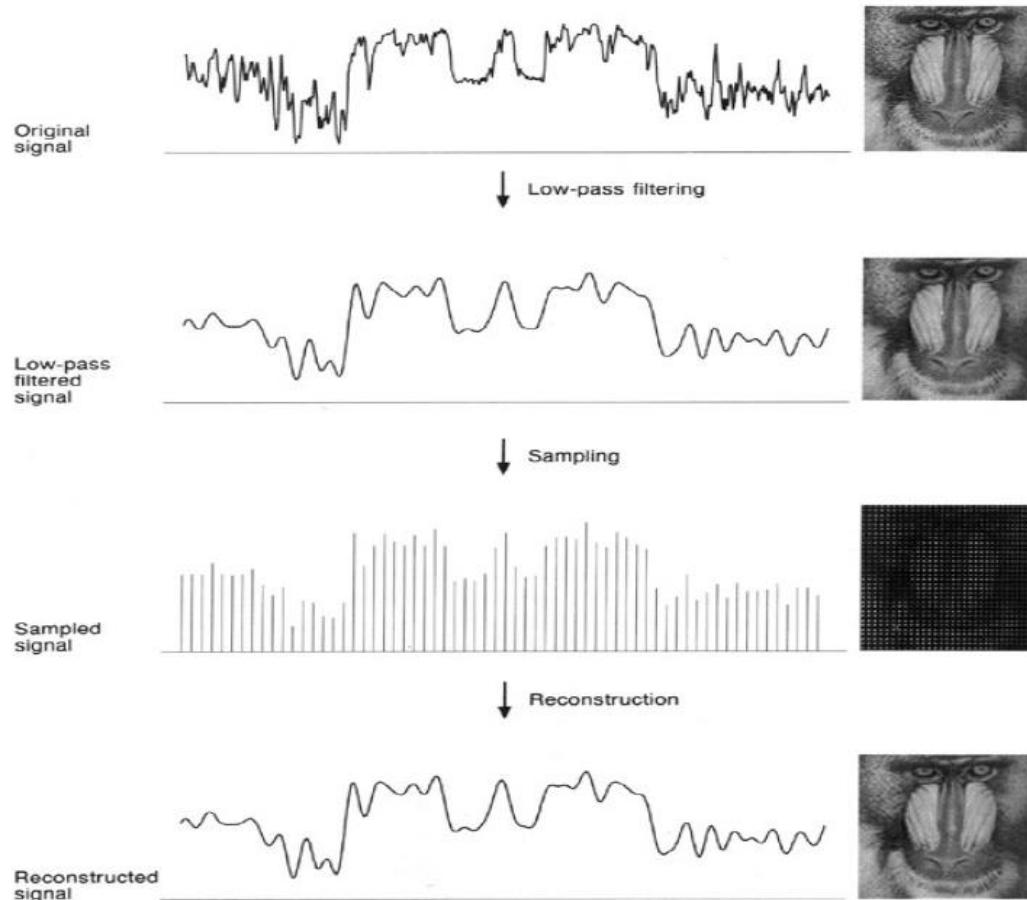
Aliasing



Aliasing

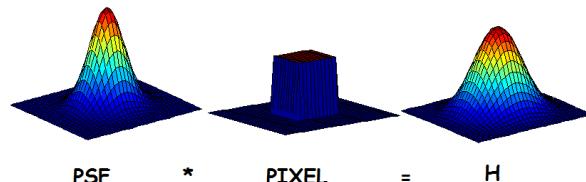


Proper sampling



Application: Spatial super-resolution

- lens+pixel=low-pass filter (desired to avoid aliasing)



- Low-res images = $D \circledast H \circledast G$ (desired high-res image)
 - D: decimate, H:lens+pixel, G: Geometric warp
- Simplified case for translation: $LR = (D \circledast G) \circledast (H \circledast HR)$
 - G is shift-invariant and commutes with H
 - First compute $H \circledast HR$, then deconvolve HR with H
- Super-resolution needs to restore attenuated frequencies
 - Many images improve S/N ratio ($\sim \sqrt{n}$), which helps
 - Eventually Gaussian's double exponential always dominates



Additional Resources

- Convolution
 - 3Blue1Brown: [YouTube](#)
- Fourier Transform
 - 3Blue1Brown: [YouTube](#)
 - Shree Nayar: [YouTube](#), [YouTube](#)

Disclaimer

Many of the slides used here are obtained from online resources (including many open lecture materials) without appropriate acknowledgement. They are used here for the sole purpose of classroom teaching. All the credit and all the copyrights belong to the original authors. You should not copy it, redistribute it, put it online, or use it for any other purposes than for this course.