

Computer Vision 1

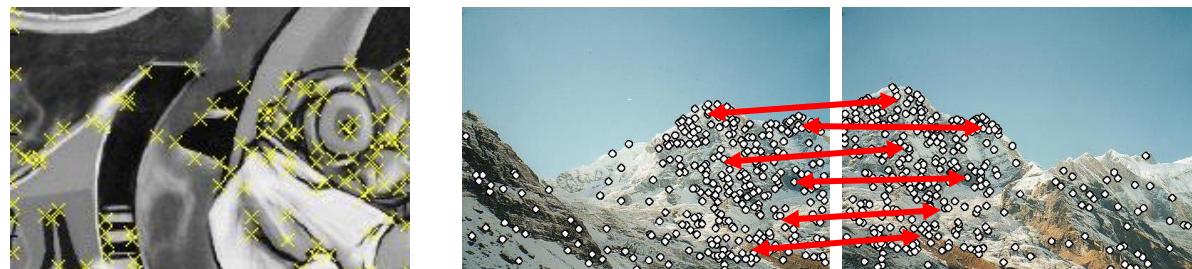
HC3a

Local Features Edges, Lines, Corners

Dr. Martin Oswald, Dr. Dimitris Tzionas, Dr. Arun Mukundan,
[m.r.oswald, d.tzionas, a.mukundan]@uva.nl

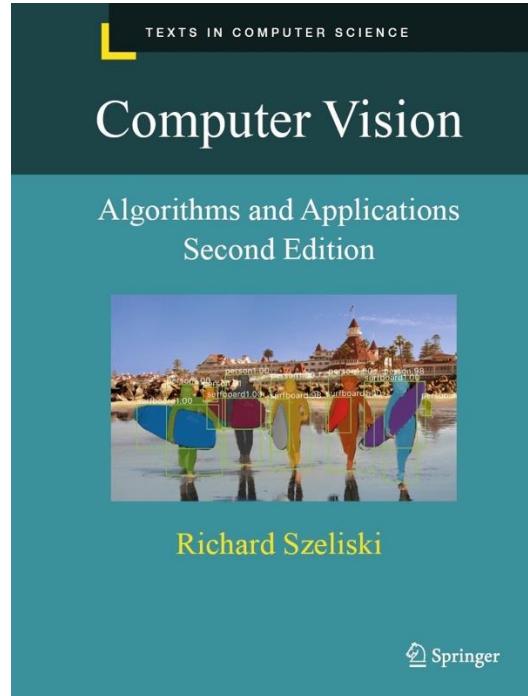
Outline

- Edge Detection
 - Derivatives of Image
 - Canny Edge Detector
- Line Fitting
 - Least Squares
 - RANSAC
 - Hough Transform
- Corners
 - Harris Corners
 - SIFT
 - Applications



Textbook

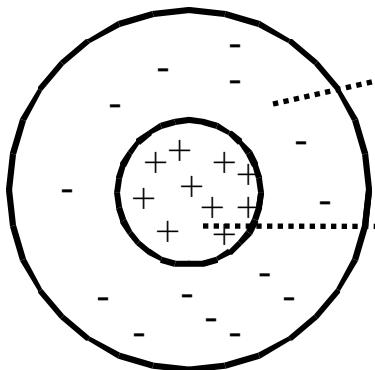
- Sec. 7.2
- Sec. 7.4
- Sec. 7.1



Outline

- Edge Detection
 - Derivatives of Image, Derivatives of Gaussian
 - Canny Edge Detector
- Line Fitting
 - Least Squares
 - RANSAC
 - Hough Transform
- Corners
 - Harris Corners
 - SIFT
 - Applications

Ganglion Cell – Week 1 Refresher



On-Center
Ganglion Cell

Inhibitory
(-) region

Excitatory
(+) region

When illuminated ...

- ... **reduces** cell's output
(think: Contribute **negatively**)
- ... **increases** cell's output
(think: Contribute **positively**)

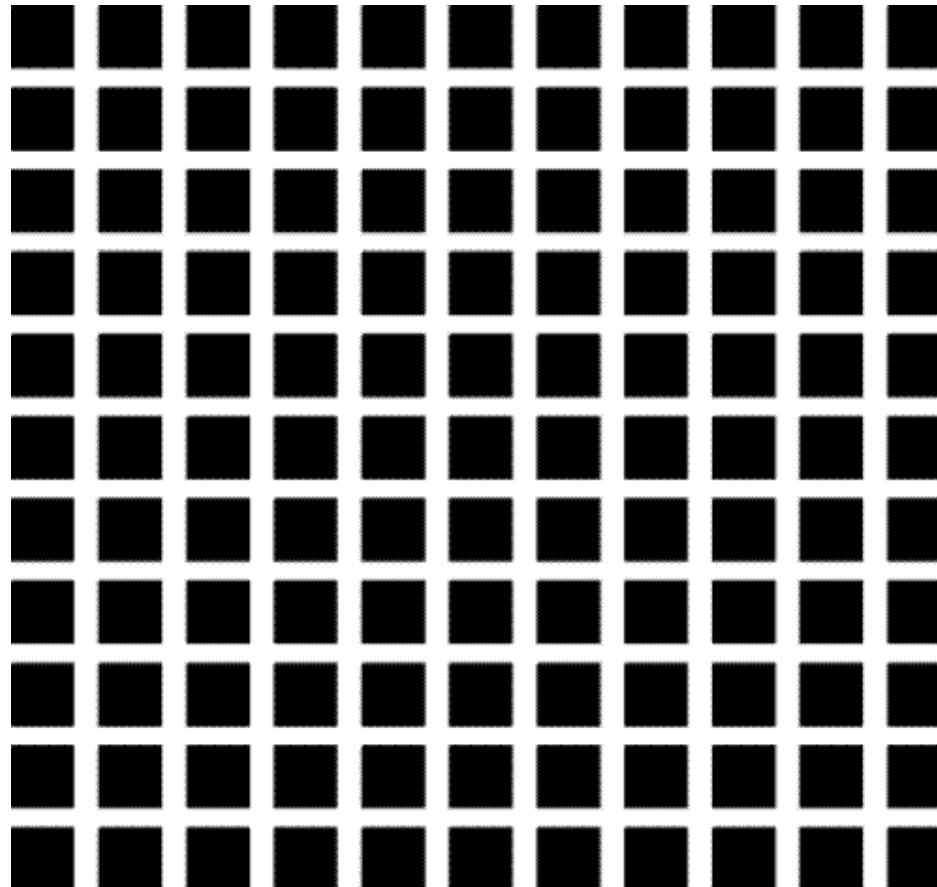
Total output:
Compute **Net Signal**

Ganglion Cells – Illusion – Hermann Grid

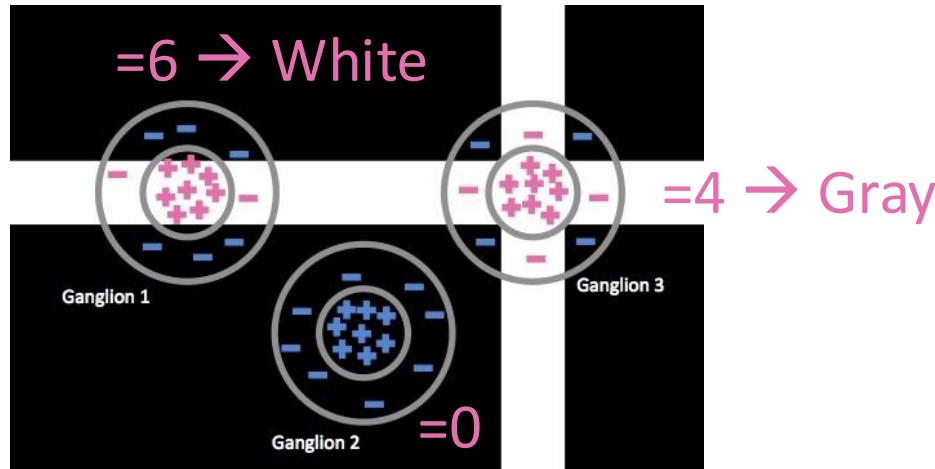
White grid on
black background

Illusion

Faint **grey** spots @
intersections
in **peripheral** areas!



Ganglion Cells – Illusion – Hermann Grid

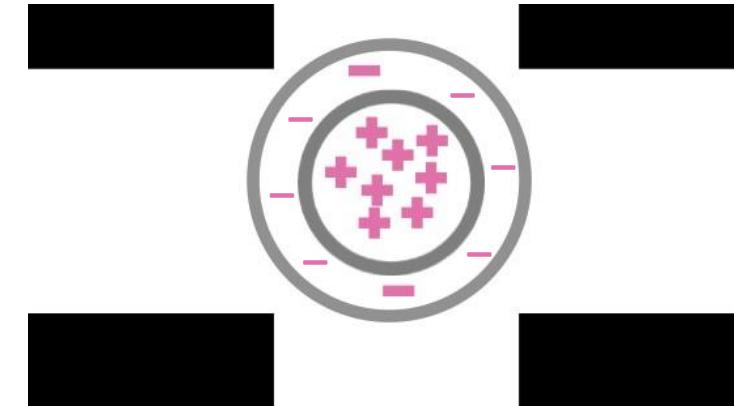


Illusion – Grey spots @ peripheral intersections

Pink → Inputs **stimulated** by light

Blue → Inputs that **not stimulated**

Compute **net signal** between activated + and -

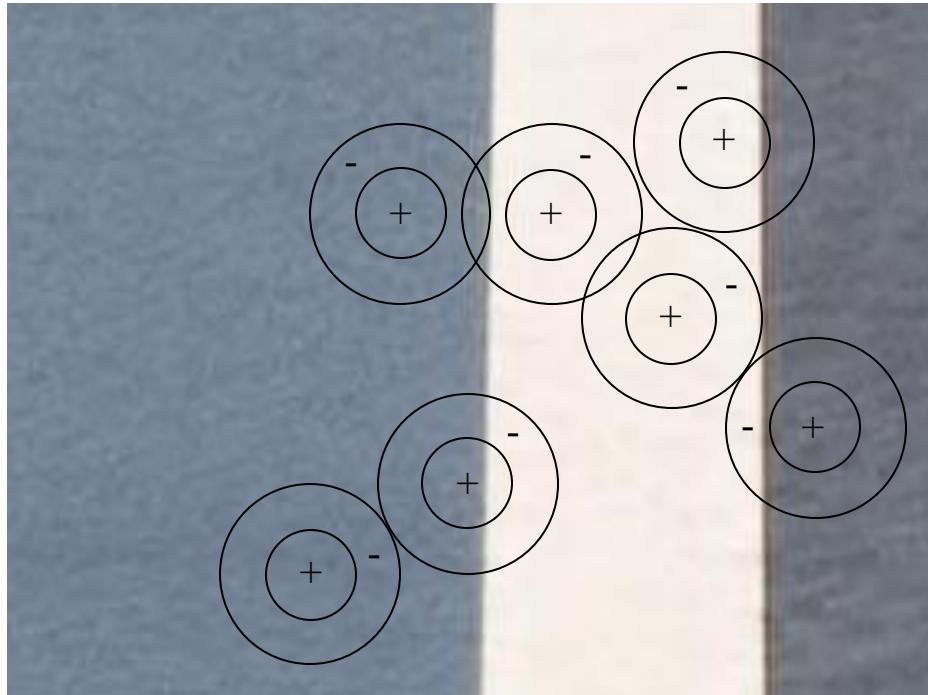


Illusion **disappears** for intersection we directly look at!

Fovea → Receptive fields smaller
→ Cells @ intersection no longer surrounded by inhibitory components of grid

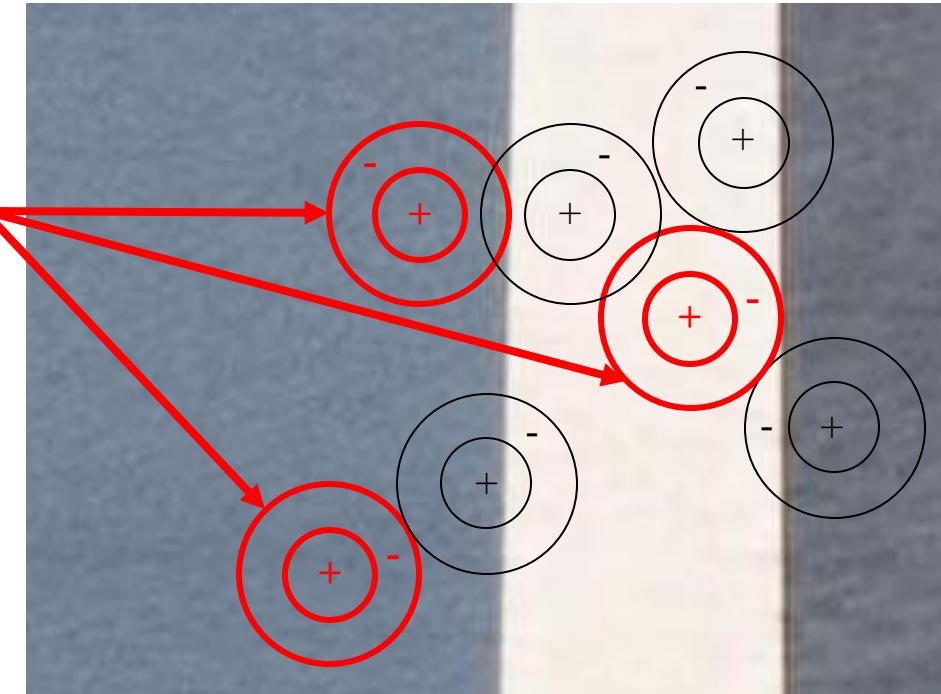


Ganglion Cells – Responses

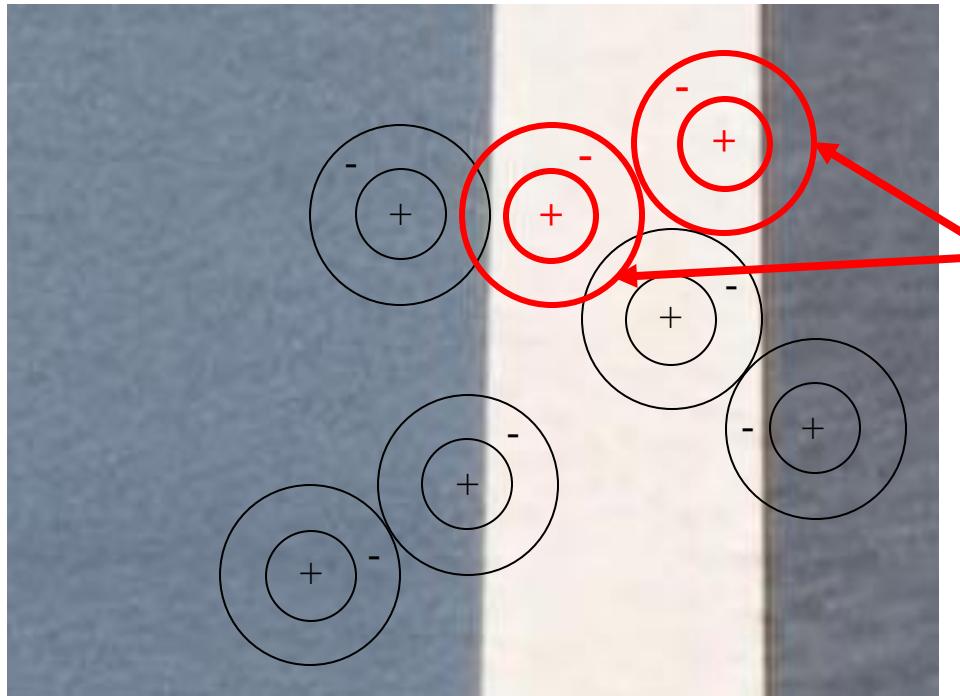


Ganglion Cells – Responses

Illuminated: Uniform
Response: Low

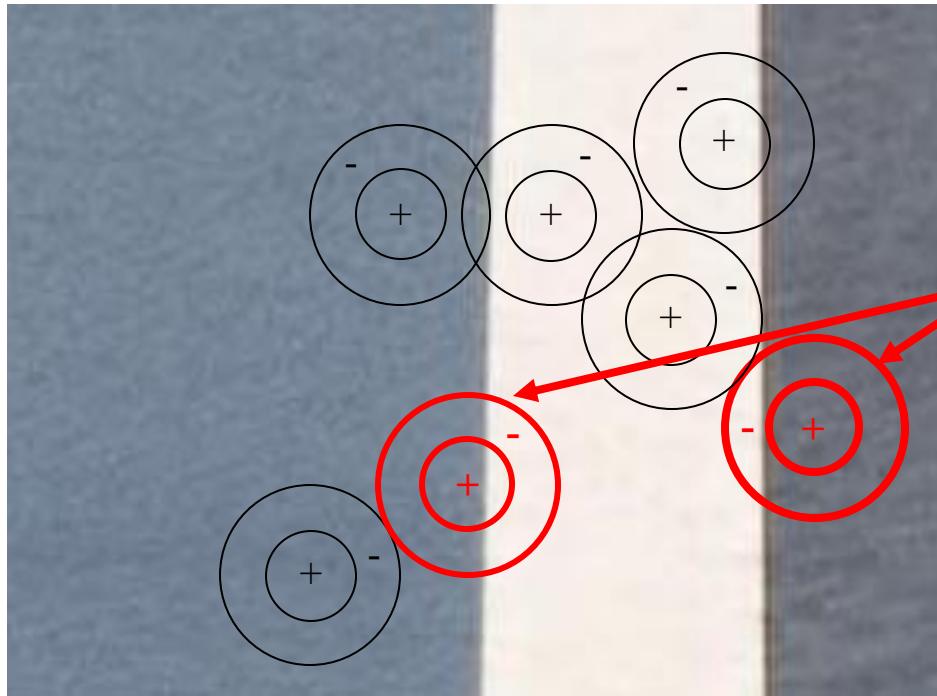


Ganglion Cells – Responses



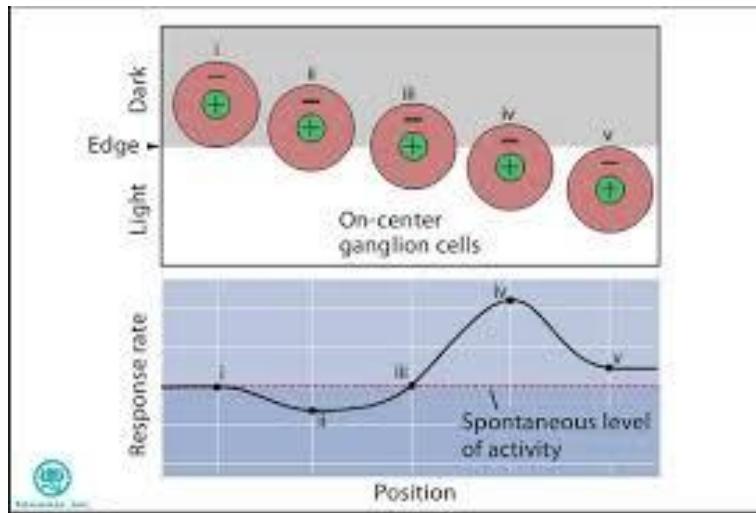
+ Illuminated: Full
- Illuminated: Partial
Response: High

Ganglion Cells – Responses



+ Illuminated: ~None
- Illuminated: Partial
Response: Low
(maybe negative)

Ganglion Cell – Responses



Ganglion Cells – Key Take-Away

Input image
(cornea)

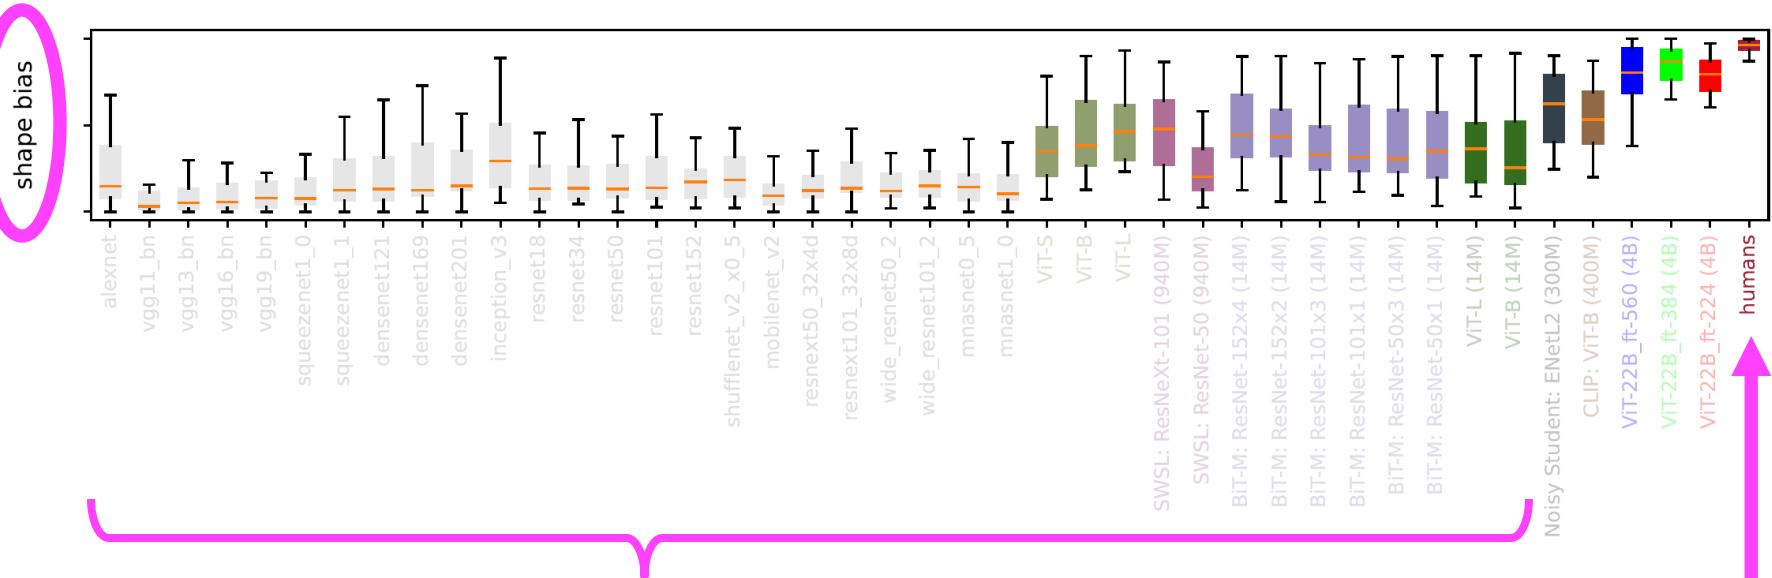


“Neural image”
(retinal ganglion cells)



Ganglion cells respond to **edges**

Human Perception – Bias towards Shape



Most **Computer Vision** models:

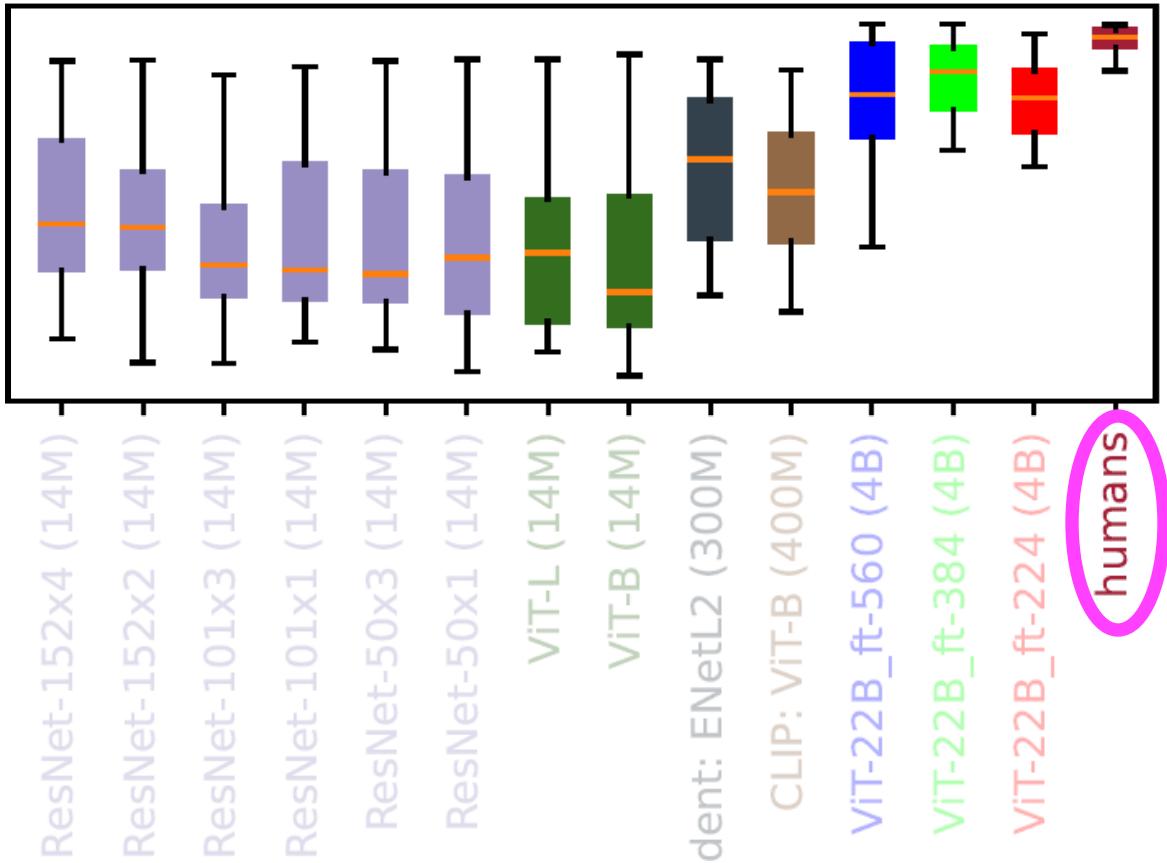
- High bias towards **Texture**
- Low bias towards **Shape**

Human perception
is biased towards
Shape

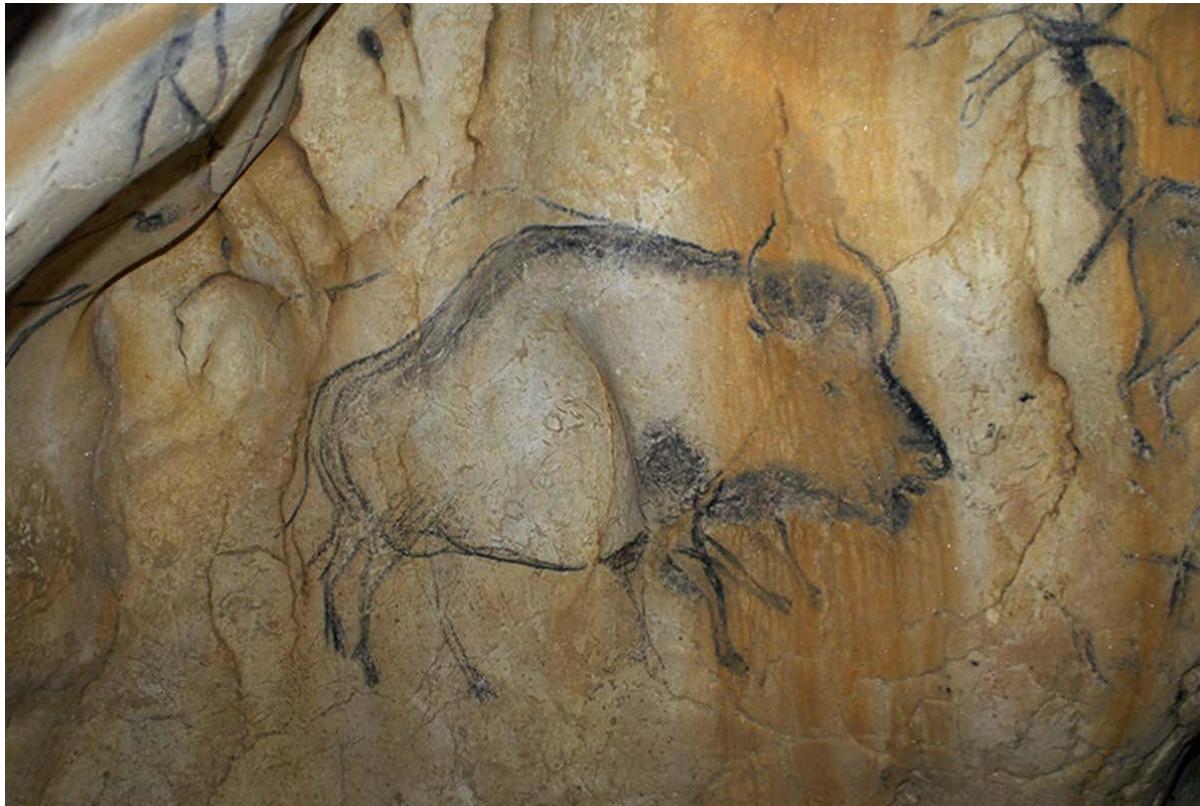
Human Perception – Bias towards Shape

CVN

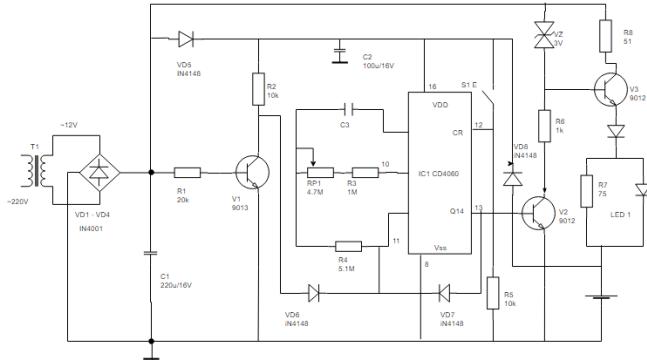
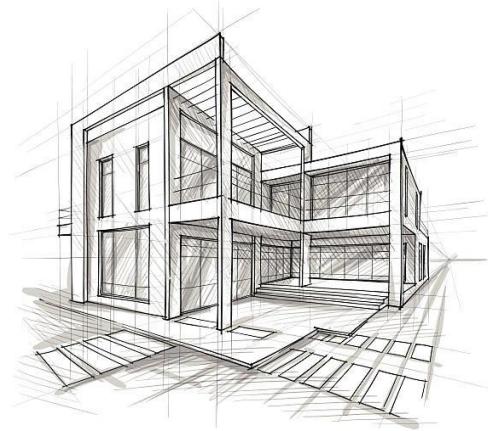
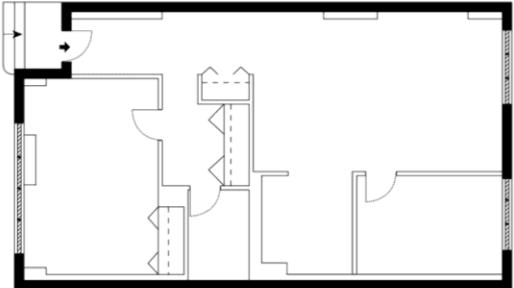
Big 'gap' between
• Humans
• Computer Vision



Early Edge Drawings

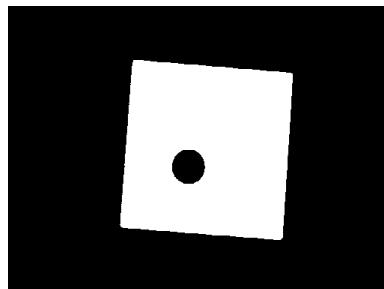


Typical Edge Drawings

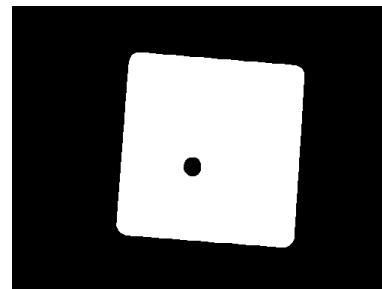


Edge Detection – Naive Way

1. Dilate input image
2. Subtract the input image from the dilated one
3. Edges arise!

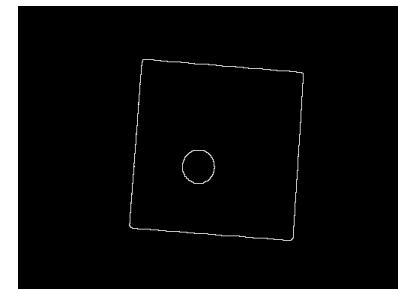


Input Image



Dilated Image

(exaggerated for
visualization purposes)

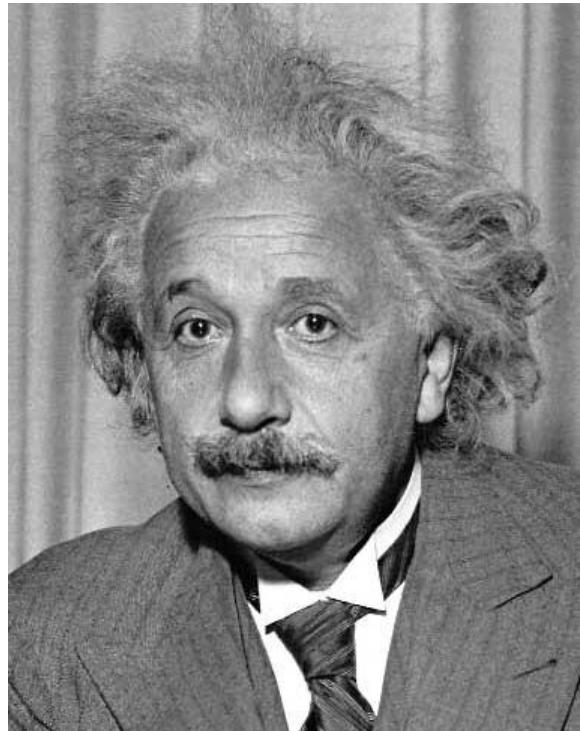


Edges

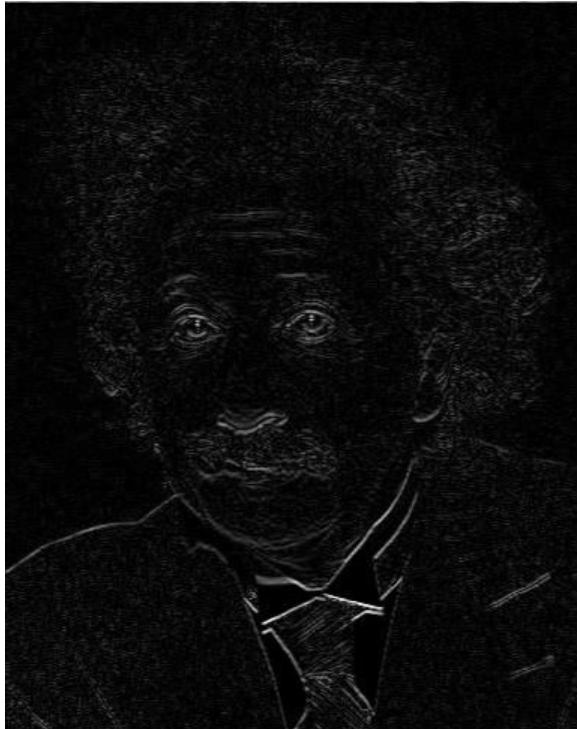
image ‘arithmetics’:

`Edges = dilated - input_img`

Derivative Filters



Sobel

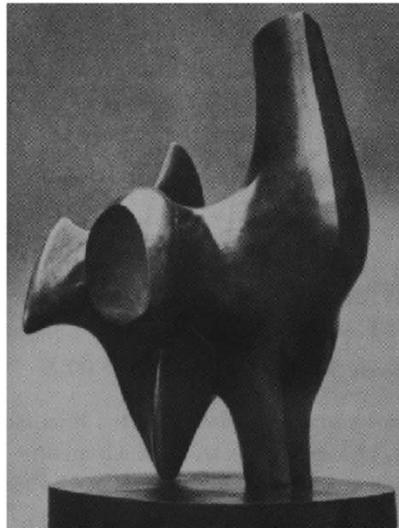


Horizontal Edges
(absolute value)

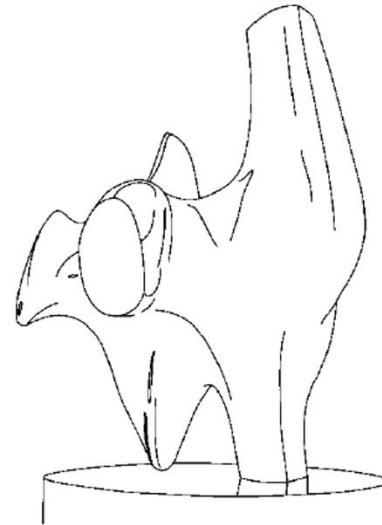
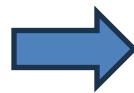
Outline

- Edge Detection
 - Derivatives of Image, Derivatives of Gaussian
 - Canny Edge Detector
- Line Fitting
 - Least Squares
 - RANSAC
 - Hough Transform
- Corners
 - Harris Corners
 - SIFT
 - Applications

Edge Detection



Input



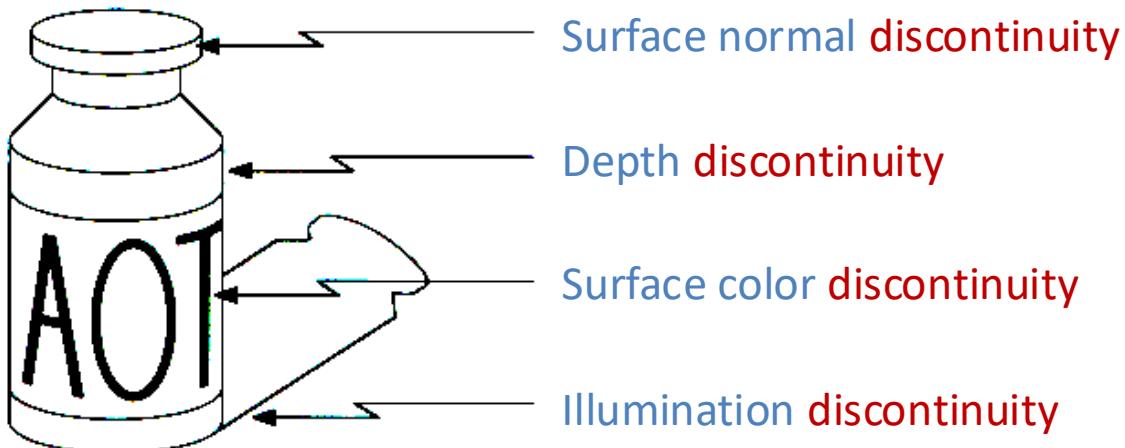
Output

Convert **2D image** → **Set of curves**

- Extracts **salient features** of the scene
- More **compact** than pixels

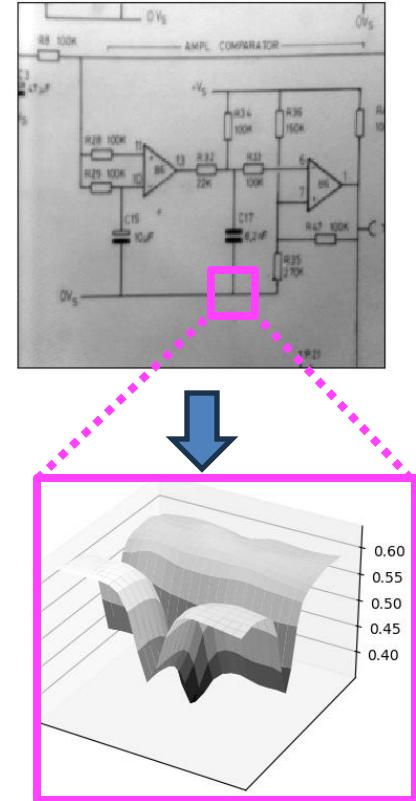
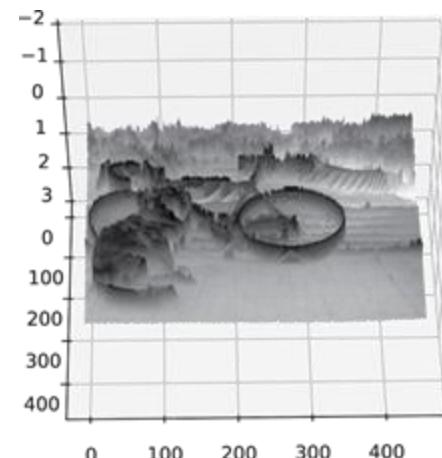
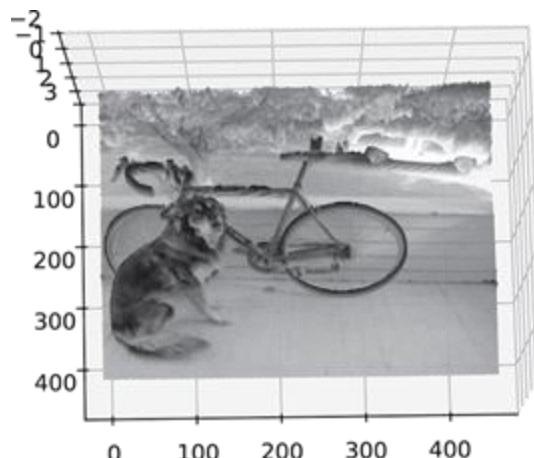
Edges – Origin

Caused by various factors:



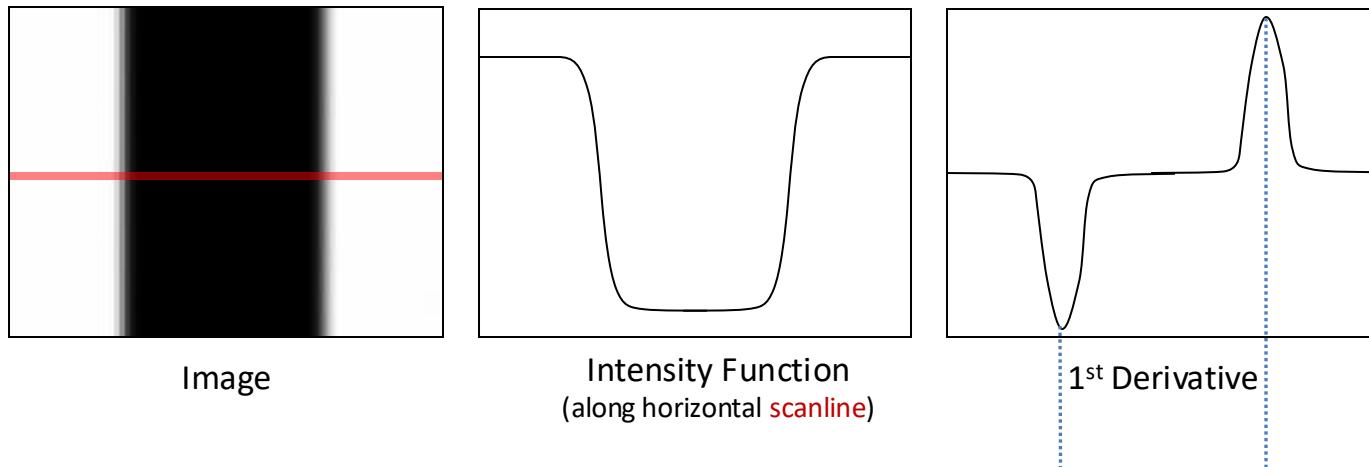
What is an Edge?

- Image is a **function**
- Think of the gray tones as **HEIGHTS**
- **Edges** are **rapid changes** – look like **steep cliffs**



Edge Detection – 1D Function

Edge → A place of **rapid change** in image intensity function



Edges correspond to Extrema of Derivative



Image Derivatives

How can we differentiate a *digital* image $F[x,y]$?

- **Option 1** ☹ :

- First, reconstruct a *continuous* image, f
- Then, compute the *derivative*

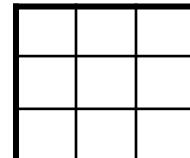
- **Option 2** ☺ :

- Take **discrete derivative** (finite difference):

$$\frac{\partial f}{\partial x}[x, y] \approx F[x + 1, y] - F[x, y]$$

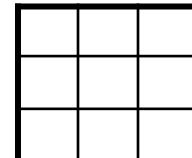
How would you implement this as a *linear filter*?

$$\frac{\partial f}{\partial x}:$$



$$H_x$$

$$\frac{\partial f}{\partial y}:$$



$$H_y$$

Image Gradient

Image Gradient: $\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$

Gradient → Points in **direction** of
most rapid increase
in image intensity



$$\nabla f = \left[\frac{\partial f}{\partial x}, 0 \right]$$

$$\nabla f = \left[0, \frac{\partial f}{\partial y} \right]$$

$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

Edge **strength** → **Gradient magnitude**:

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

Gradient direction:

$$\theta = \tan^{-1} \left(\frac{\partial f / \partial y}{\partial f / \partial x} \right)$$

Image Gradients

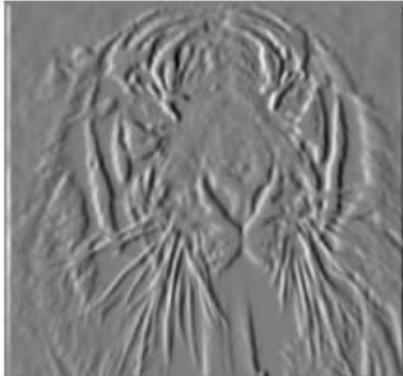
Image



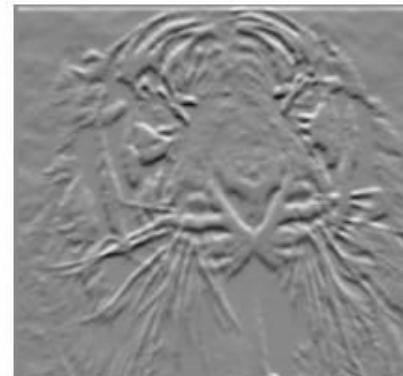
Gradient
Magnitude



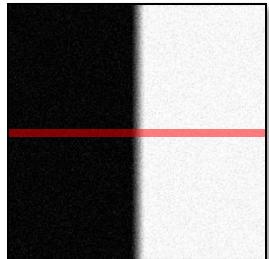
X-direction
Gradients



Y-direction
Gradients



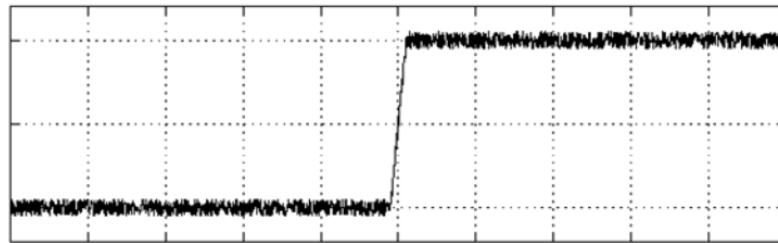
Effect of Noise



Noisy image

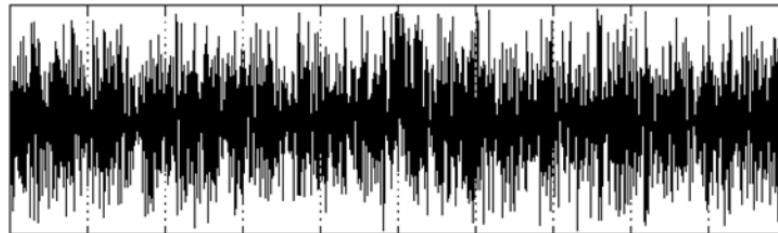


$f(x)$



Where is
the edge?

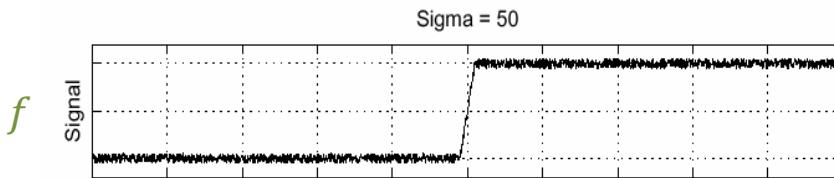
$$\frac{d}{dx} f(x)$$



Derivation
amplifies noise
(like a high-pass filter)

Solution – Pre-smooth Image

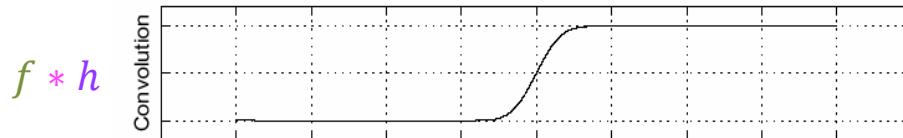
Noisy image



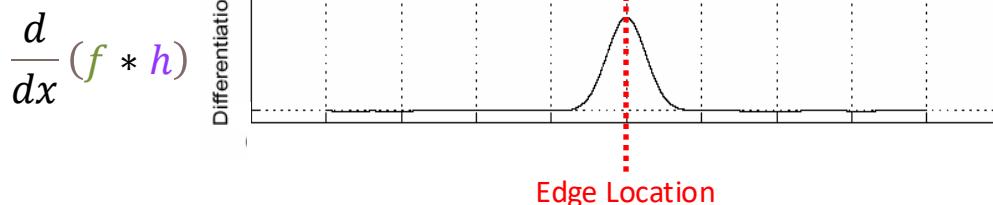
Smoothing w. Gaussian kernel



Convolve to smoothen



Differentiate & find extrema



To find edges, find peaks in $\frac{d}{dx} (f * h)$

But these are
3 steps



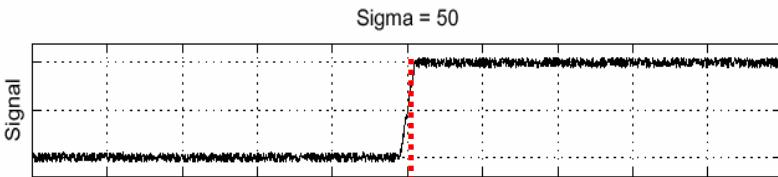
And we must
differentiate
every time



Pre-smooth & Pre-Differentiate

Noisy image

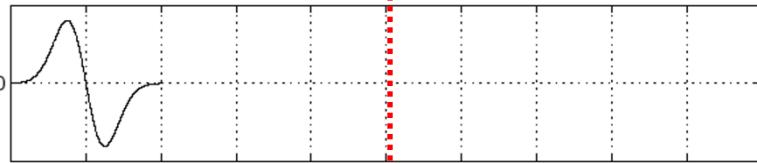
f



Differentiate kernel offline
(only once)

$\frac{dh}{dx}$

Kernel



Online convolution h
pre-diff kernel

$f * \frac{dh}{dx}$

Convolution

Edge Location

Convolution is
Associative!

$$\frac{d}{dx} (f * h) = f * \frac{dh}{dx}$$



To find edges,
find peaks in $f * \frac{dh}{dx}$



1 step less online!
Remaining step is cheaper!

Differentiating
Gaussian kernel h :

- Is exact – Compute first analytically, and later discretize
- Do it offline, only once

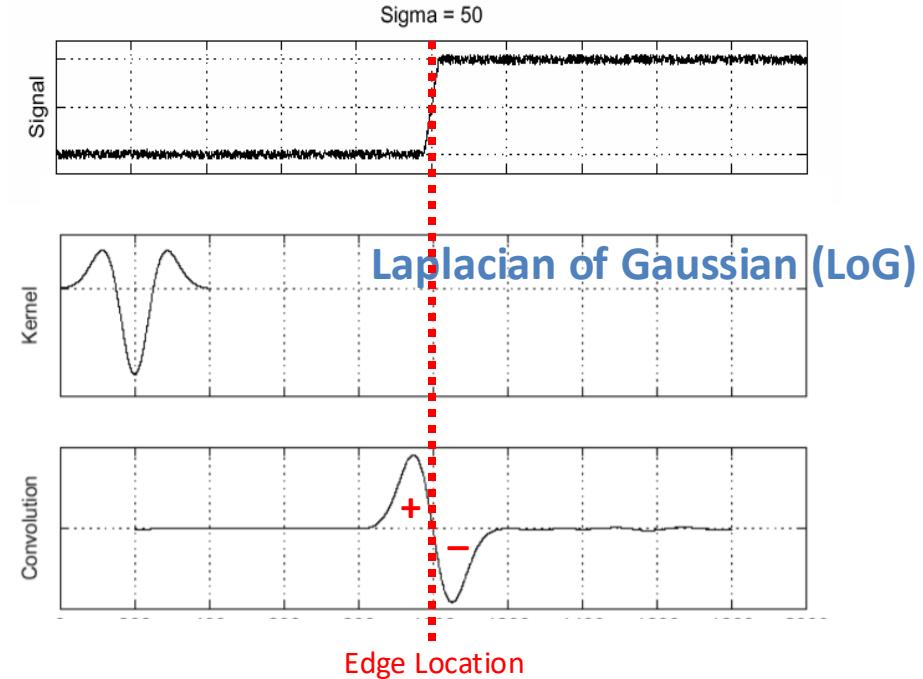


Extra Robustness

Edge as **zero-crossing**
of 2nd order derivative using
Laplacian of Gaussian (LoG)



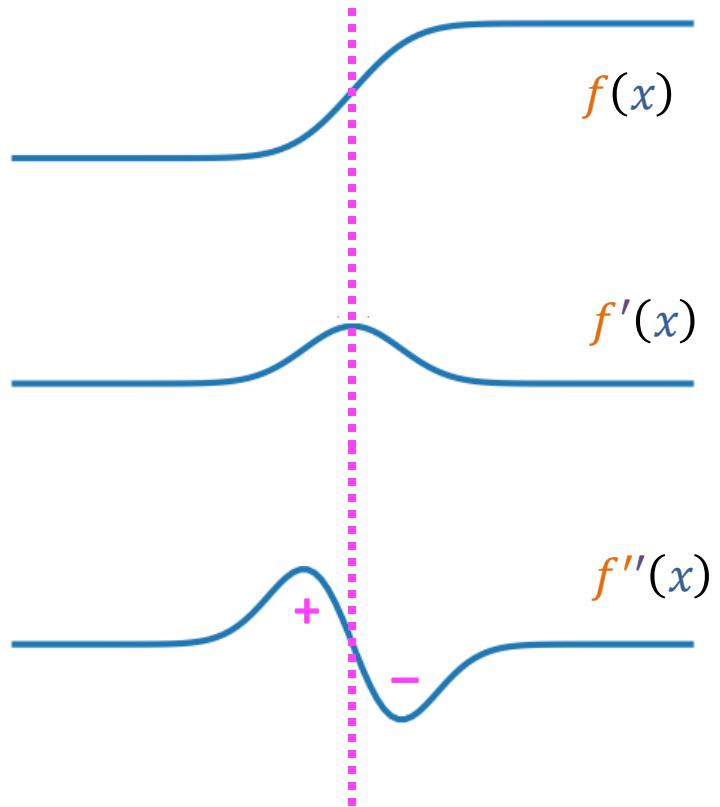
$$f * \left(\frac{d^2}{dx^2} h \right)$$



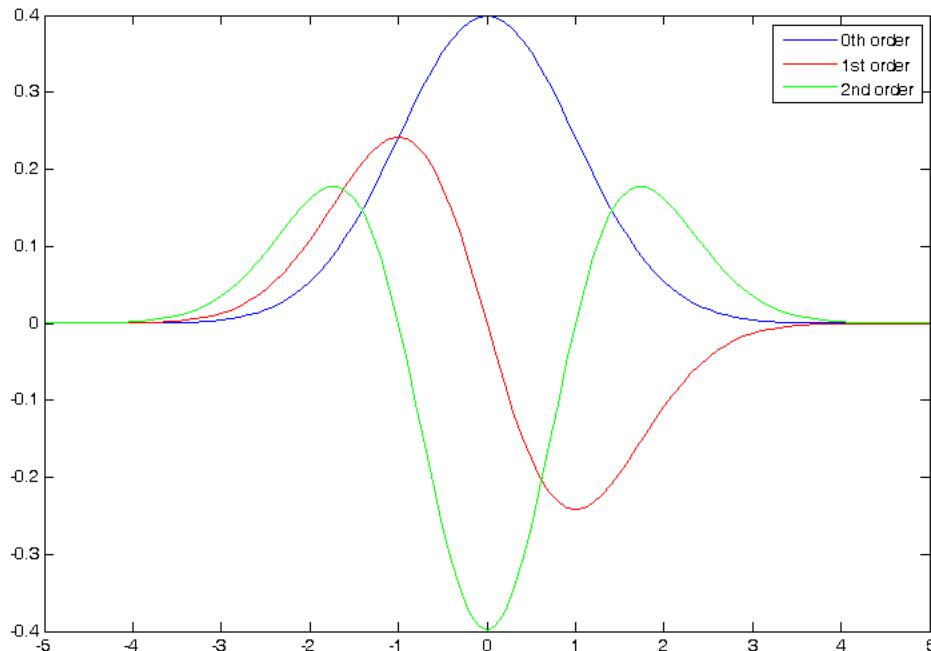
Extra Robustness

Extra robustness → Add a second condition!

- Ideal case
 - **1st order** derivative: local extremum $|f'(x)| \gg 0$
 - **2nd order** derivative is 0 $f''(x) \approx 0$
- Discrete signal in practice:
 - **1st order** derivative: local extremum
 - **2nd order** crosses 0 → consider left/right samples!



1D Gaussian & Derivatives



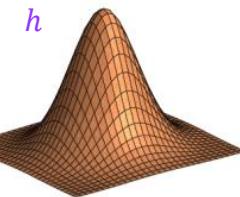
$$h_\sigma(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{x^2}{2\sigma^2}}$$

$$h'_\sigma(x) = \frac{d}{dx} h_\sigma(x) = -\frac{x}{\sigma^2} h_\sigma(x)$$

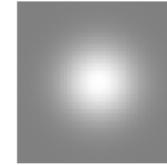
$$h''_\sigma(x) = \frac{d^2}{dx^2} h_\sigma(x) = \left(\frac{x^2}{\sigma^4} - \frac{1}{\sigma^2} \right) h_\sigma(x)$$

2D Gaussian & Derivatives

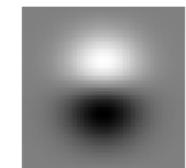
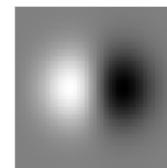
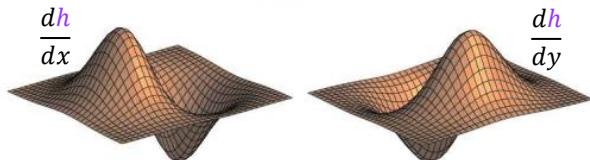
Gaussian kernel



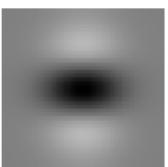
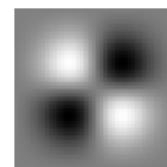
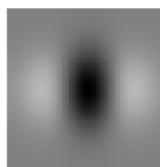
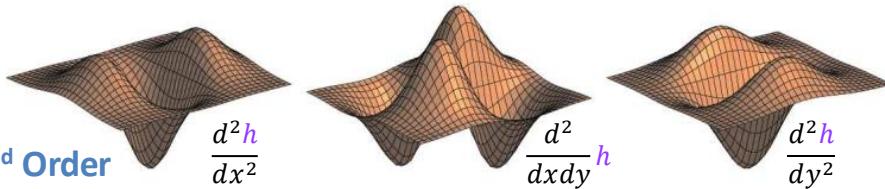
Height along 'z' re-encoded
as bright/dark value



1st Order Derivatives



2nd Order Derivatives

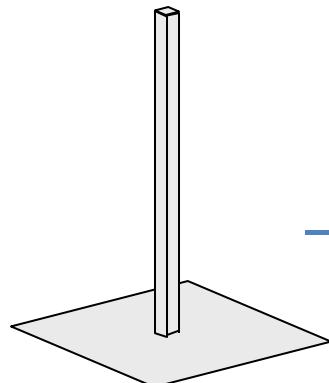


Laplacian

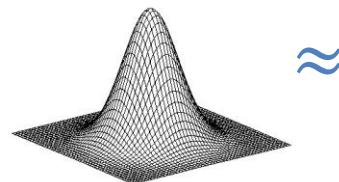
$$\nabla^2 h = \frac{d^2 h}{dx^2} + \frac{d^2 h}{dy^2}$$

Difference of Gaussians (DoG)

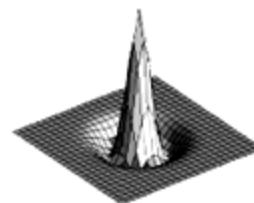
- Laplacian of Gaussian (LoG) can be approximated as Difference of Gaussians (DoG)
- Spoiler: Used later today for SIFT keypoints



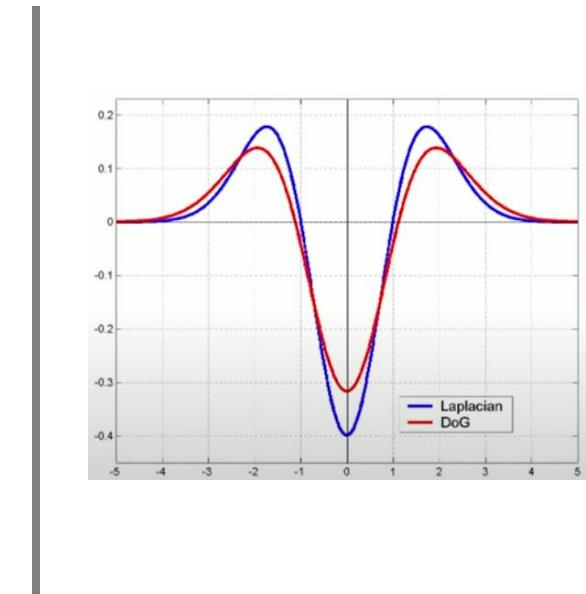
A narrow Gaussian
(or an impulse function)



A wide Gaussian



Laplacian of
Gaussian (LoG)

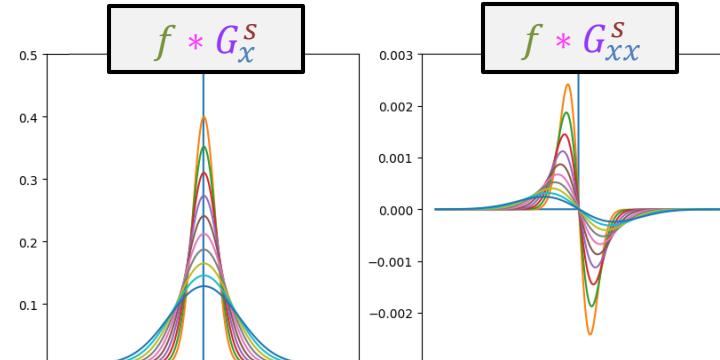


Gaussian Derivatives & Scale

Example for Step Function f

f
$s = 10.00$
$s = 11.36$
$s = 12.92$
$s = 14.68$
$s = 16.68$
$s = 18.96$
$s = 21.54$
$s = 24.48$
$s = 27.83$
$s = 31.62$

1st order

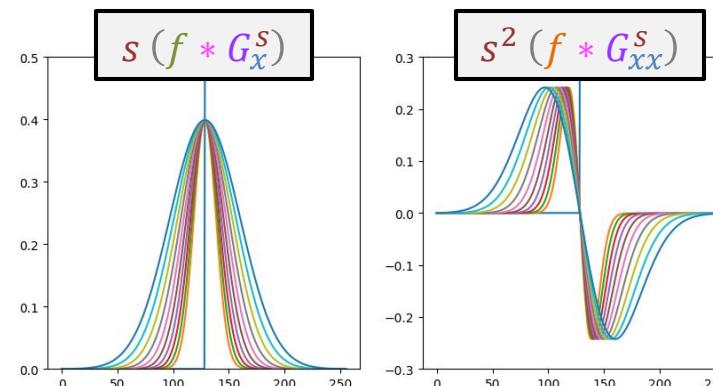


2nd order

Amplitude decreases
as scale increases



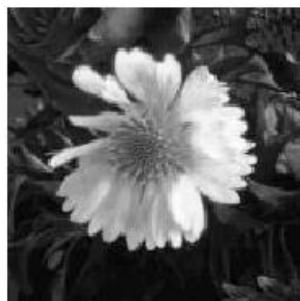
Comparing derivatives
across scales → problem!



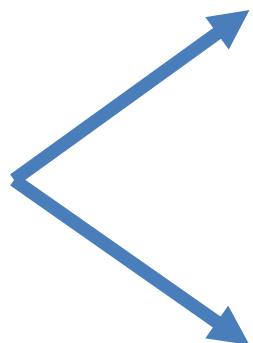
← Scale-normalized
Derivatives
 k -th order

$$s^k(f^0 * \partial^k \dots G^s)$$

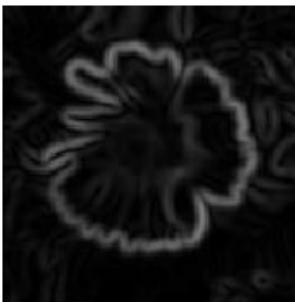
Gaussian Derivatives & Scale



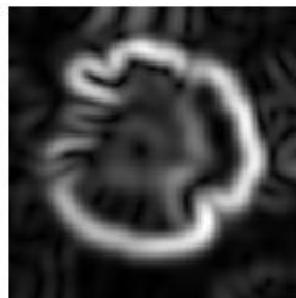
Gradient
norm →



$s = 1.00$ $s = 2.15$ $s = 4.64$ $s = 10.00$



Scale-normalized
Gradient norm →

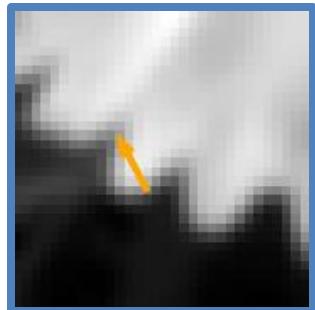


$$s^k (f^0 * \partial^k \dots G^s)$$

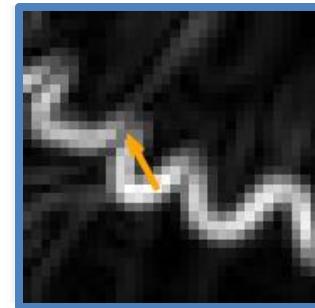
Gaussian Derivatives & Scale



Image in
scale space →



$s = 1.00$



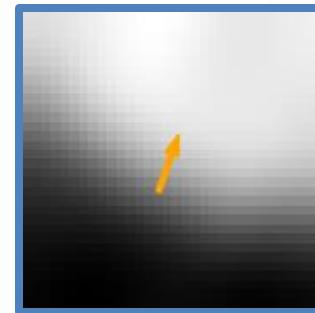
Scale-normalized
Gradient norm →

$$s^k(f^0 * \partial^k \dots G^s)$$

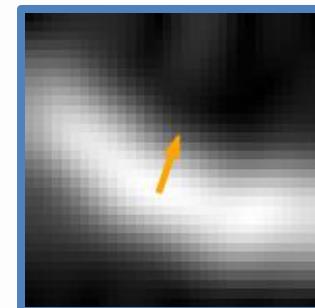


⚠ Gradient at each point → depends on **scale** of Gaussian conv

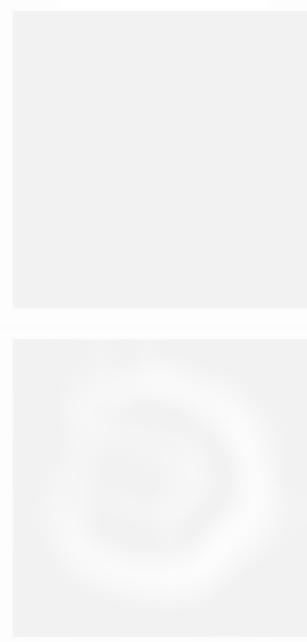
$s = 2.15$



$s = 4.64$

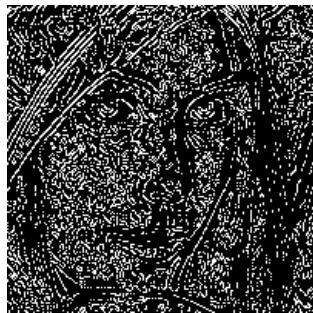


$s = 10.00$

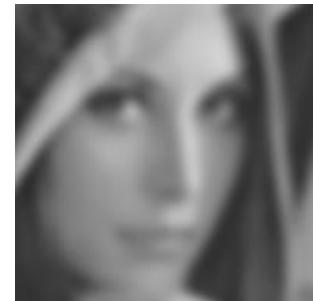
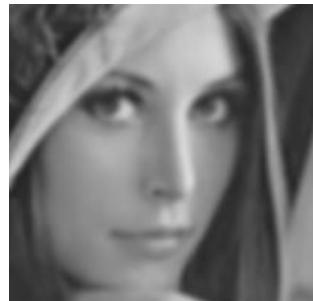


Change of Sigma

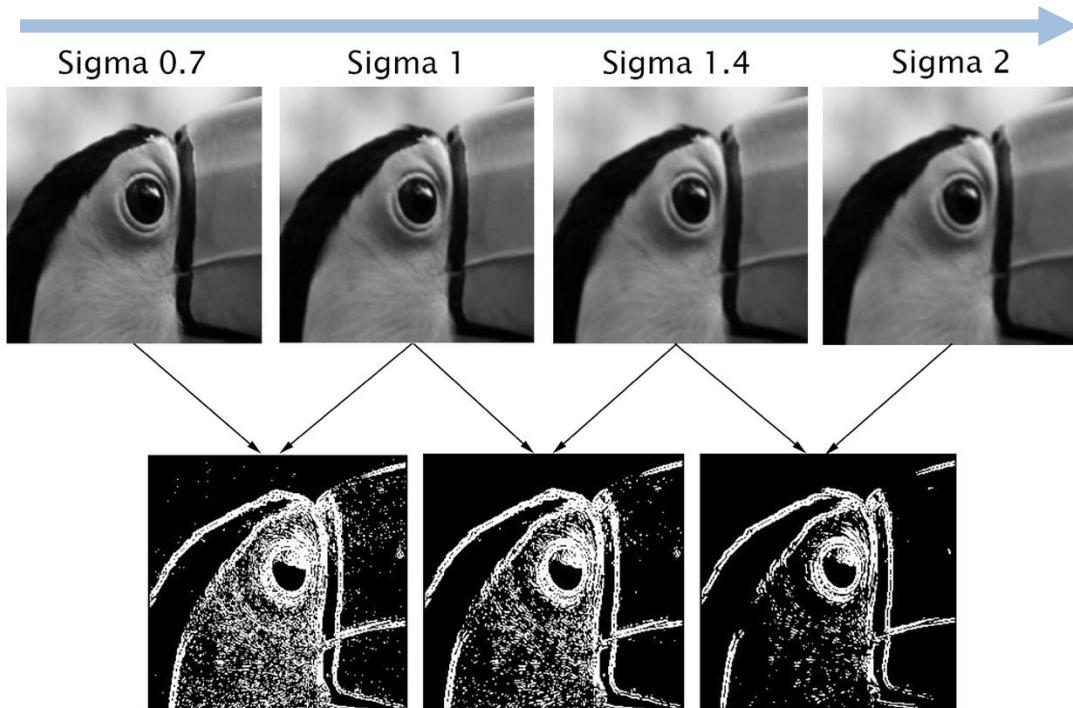
Small
Sigma



Large
Sigma



Change of Sigma – DoG



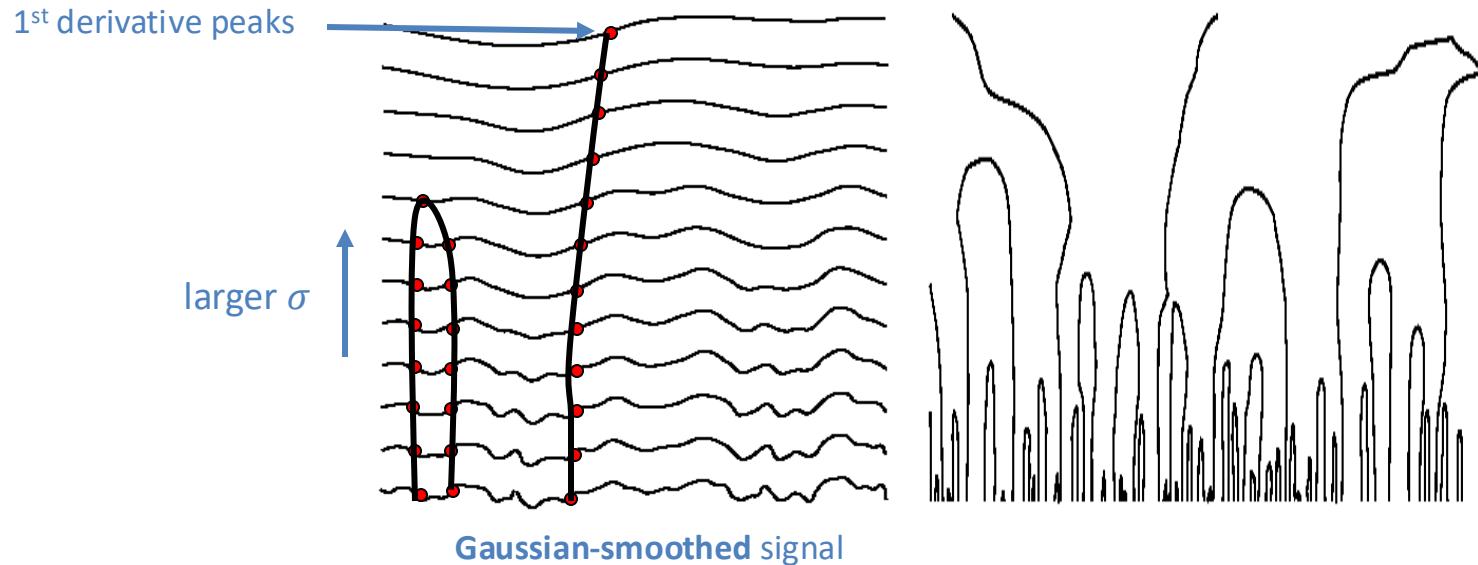
Difference of Gaussians (DoG)

Blurring with bigger-sigma
Gaussian suppresses
high-frequency info

Subtracting preserves info
in the range of frequencies
preserved in blurred imgs

DoG acts as a spatial
band-pass filter

Edges & Scale [Witkin 83]

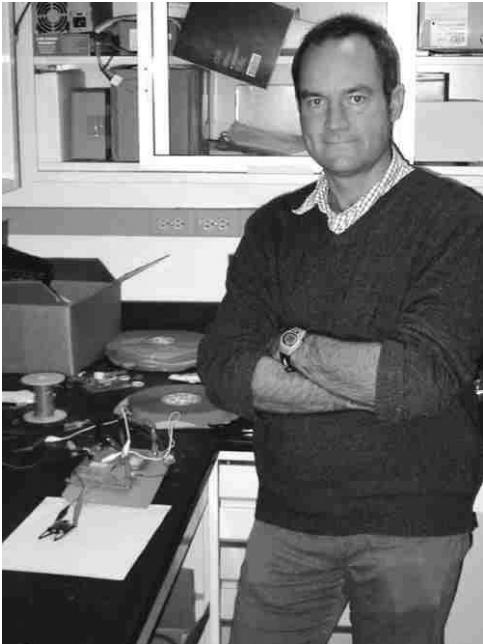


- When the scale σ increases:
 - Edge position may shift
 - Two edges may merge
 - An edge may *not* split into two

Outline

- Edge Detection
 - Derivatives of Image, Derivatives of Gaussian
 - Canny Edge Detector
- Line Fitting
 - Least Squares
 - RANSAC
 - Hough Transform
- Corners
 - Harris Corners
 - SIFT
 - Applications

Canny Edge Detector



[John Canny](#)



One of the most successful edge detection methods

Canny Edge Detector – Input Image



Image credit:
Joseph Redmon

Canny Edge Detector – Gradient Magnitude

CV

After first
pre-smoothing for
cleaner gradients
(Conv with DoG)

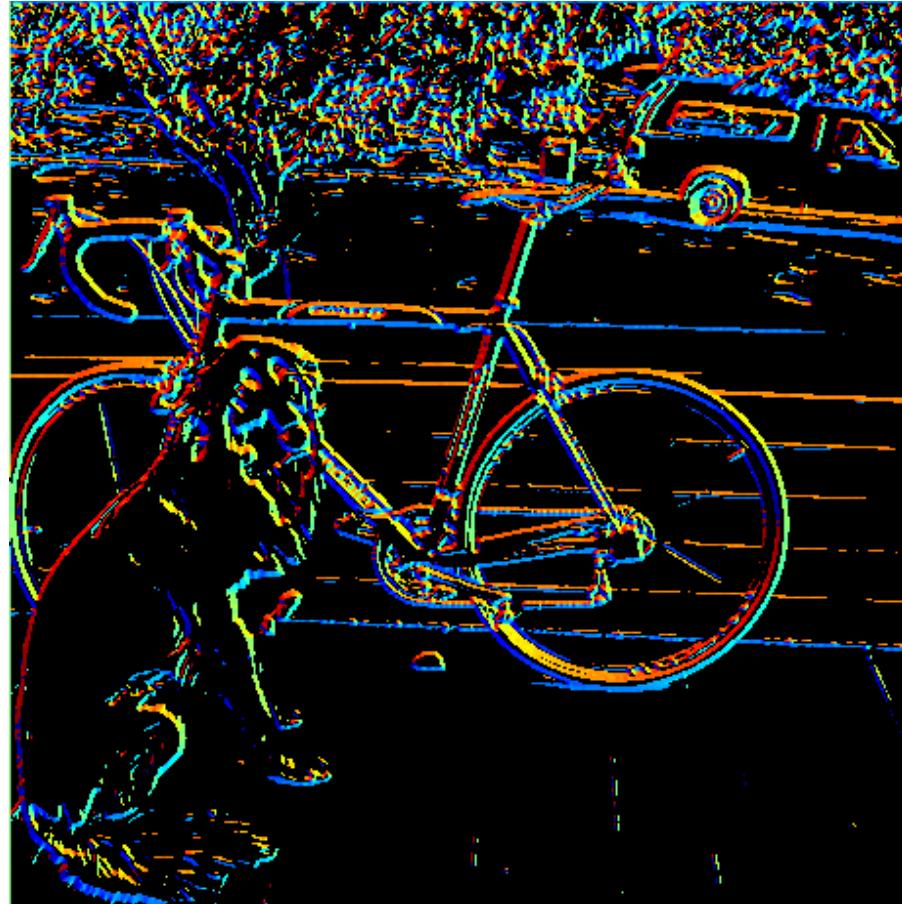


Where is the edge?

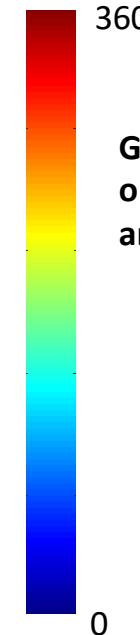
Canny Edge Detector – Gradient Orientation

CV

Get orientation
(threshold at minimum
gradient magnitude)



$$\theta = \text{atan2}(gy, gx)$$

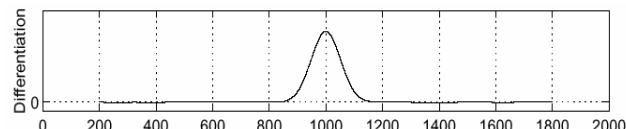
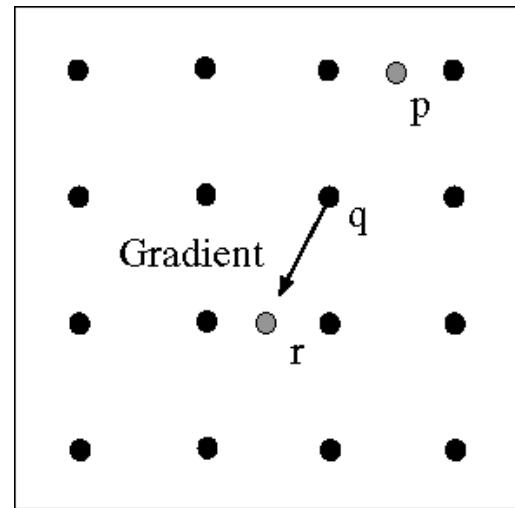
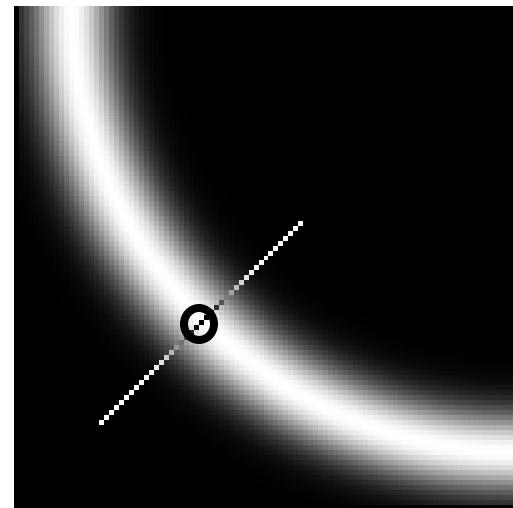


Gradient
orientation
angle

Canny Edge Detector – Non-Max Suppression

Check if pixel is **local maximum** along **gradient direction**

Requires **interpolating** pixels **p** and **r**



Canny Edge Detector – Non-Max Suppression



Before
Non-max Supr.



After
Non-max Supr.

Canny Edge Detector – Non-Max Suppression

CV

- Still, some **noise**
- Only want **strong** ‘edgels’
(edge pixels with value **R**)
- 2 thresholds **T** and **t**, 3 cases:
 - $R > T \rightarrow$ strong edge
 - $t < R < T \rightarrow$ weak edge
 - $R < t \rightarrow$ no edge
- Why two thresholds?



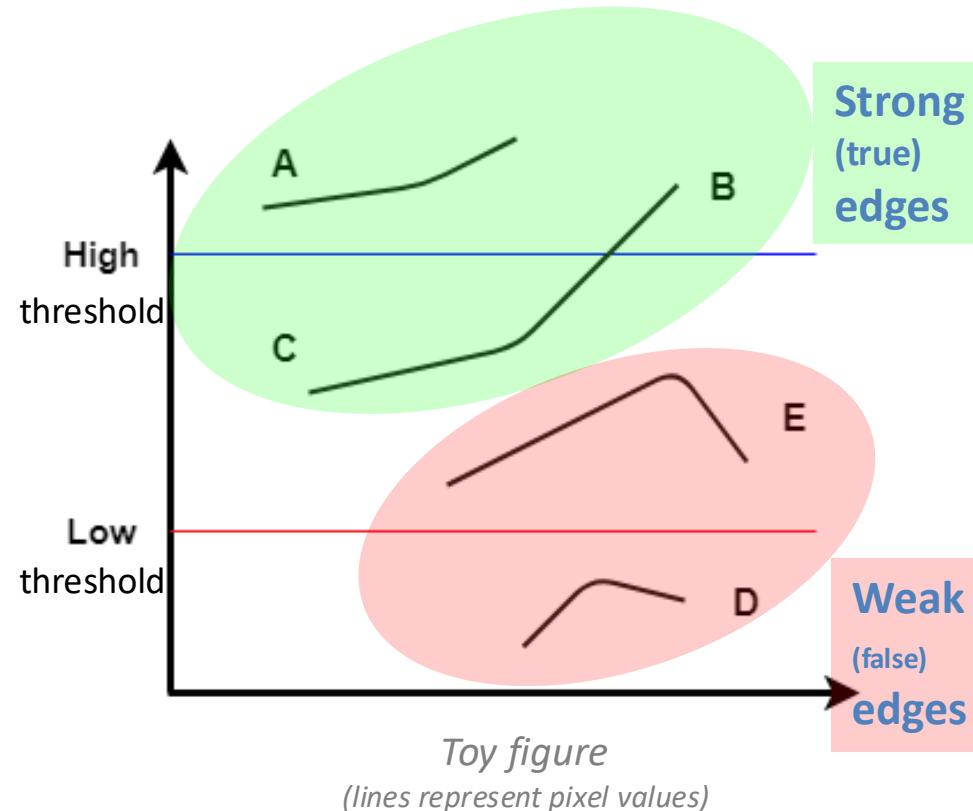
After

Non-max Supr.

Canny Edge Detector – Thresh. & Link Edges

CV

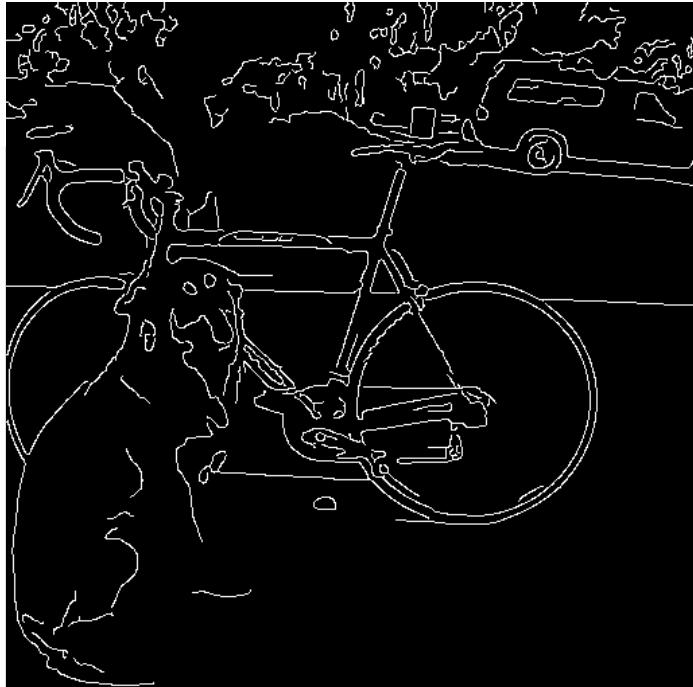
- Need to threshold (hysteresis) resulting edges and link them
- Strong edgels → edges (**A, B**)
- Weak edgels → edges only if connected to strong ones (**C**)
- Look in neighborhood (usually 8 closest)



Canny Edge Detector



Non-max Suppression
through edge following



Final Output
after thresholding edges

Canny Edge Detector – Full Example



**Input
Image**



**Gradient
Magnitude**

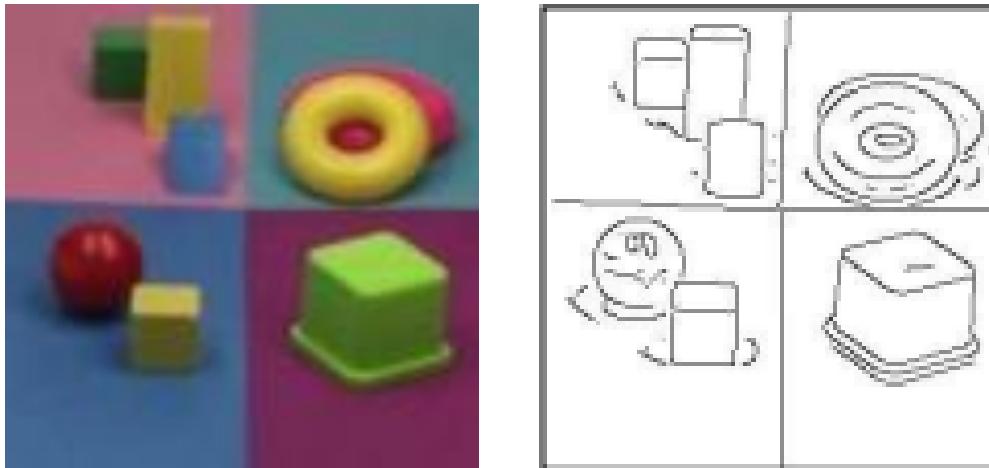


Non-max Suppression



**Final Output
after hysteresis
(thresholding)
and linking edges**

Canny Edge Detector



Also applicable on **color** images

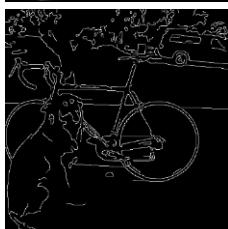
Canny Edge Detector – Summary



MATLAB: `edge(image, 'canny')`

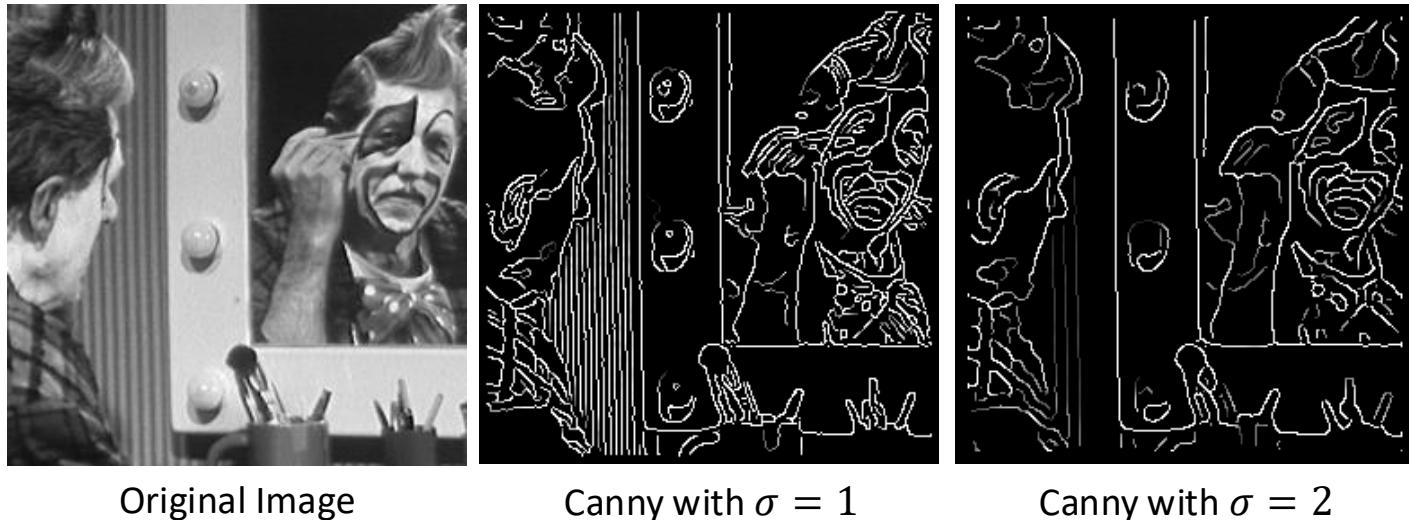


OpenCV: `cv.Canny(I, low_thr, high_thr)`



1. Filter image with Derivative of Gaussian
2. Compute gradient magnitude & orientation
3. Non-maximum suppression
4. Linking and thresholding (hysteresis):
 - Define two thresholds:
 - High threshold → Use to start edge curves
 - Low threshold → Use to continue them

Canny Edge Detector



- Choice of σ depends on desired behavior:
 - Large σ detects 'large-scale' edges
 - Small σ detects fine edges

Canny Edge Detector – Demo

from disk from url

barbara car_pad cells cervin einstein lena panda valve-wi...

Sigma 3

Low threshold 5%

High threshold 50%

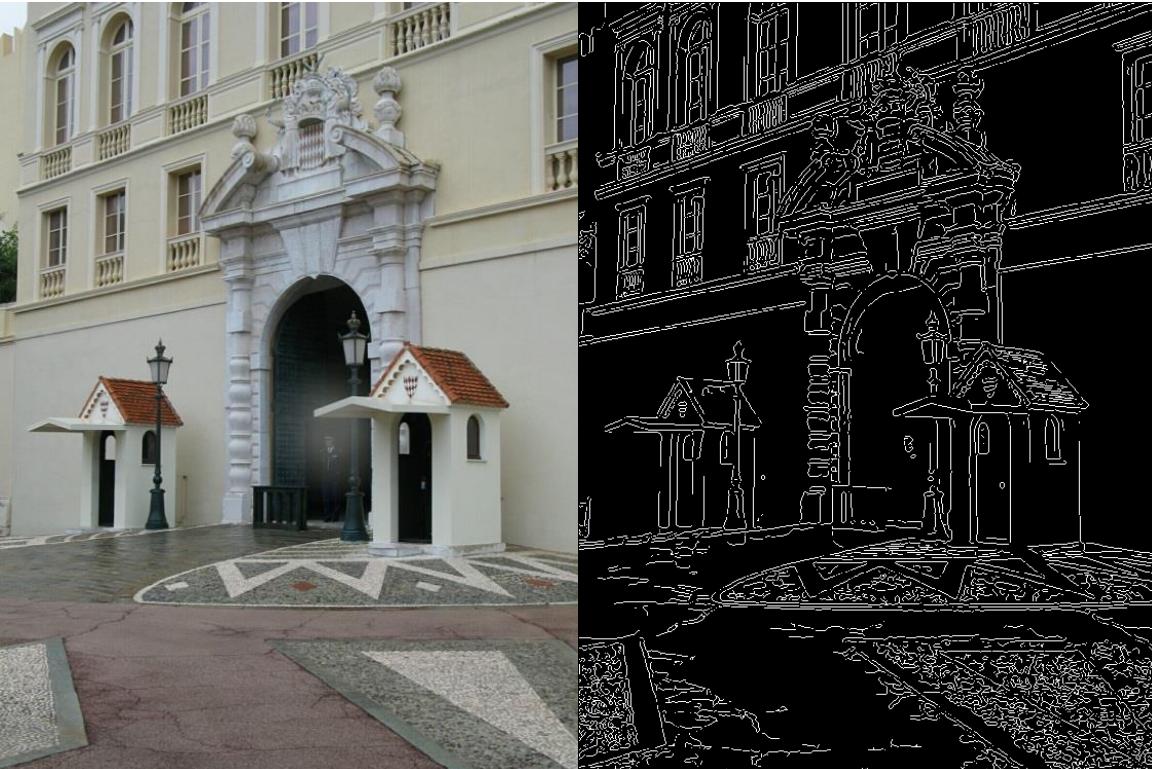
Steps 0/5



Outline

- Edge Detection
 - Derivatives of Image, Derivatives of Gaussian
 - Canny Edge Detector
- Line Fitting
 - Least Squares
 - RANSAC
 - Hough Transform
- Corners
 - Harris Corners
 - SIFT
 - Applications

Image Structure



Binary Edges Not Enough

- Fragmented edges
- False positives
- Noisy edge orientations
- Undefined Membership
which points belong to
which line ?

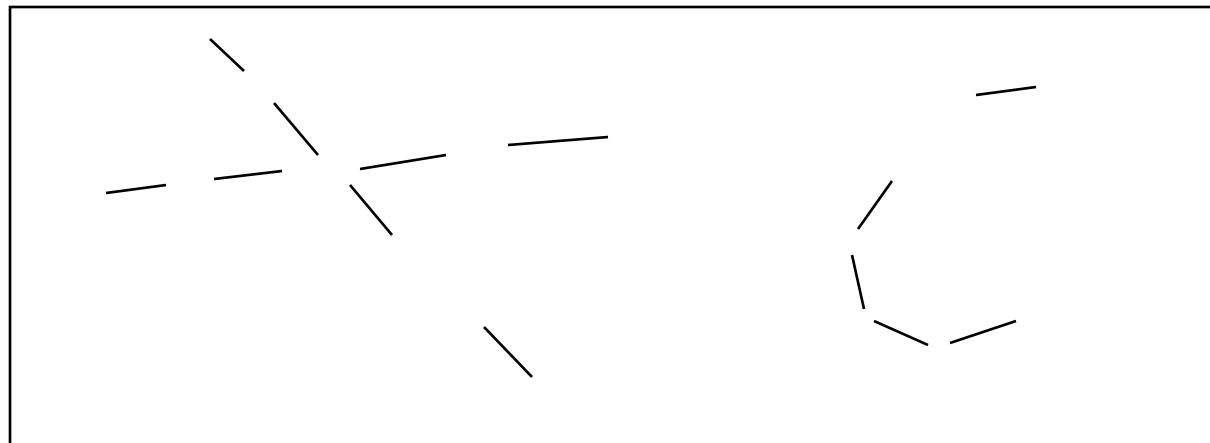
Image Structure

Example of
Vanishing
Lines



Line Fitting – Goal

- From **fragmented edges**
- To **straight lines** or **curves**

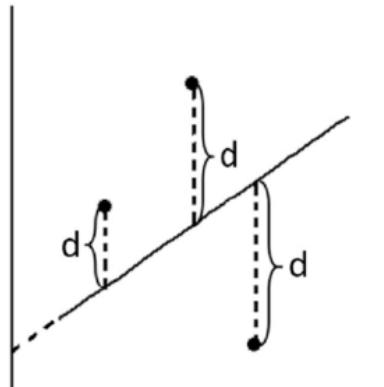


Outline

- Edge Detection
 - Derivatives of Image, Derivatives of Gaussian
 - Canny Edge Detector
- Line Fitting
 - Least Squares
 - RANSAC
 - Hough Transform
- Corners
 - Harris Corners
 - SIFT
 - Applications

Line Fitting - Ordinary Least Squares

- **Input:** Set of points belonging to a straight line
- **Output:** Find the line equation
- **Nugget:** Find the line parameters that minimize error in a least-square sense

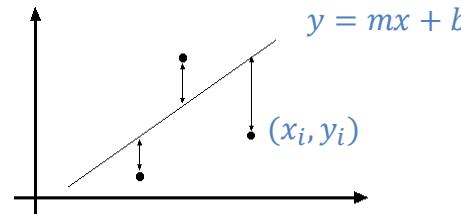


Least Squares

Line Fitting - Ordinary Least Squares

Input: Data points $(x_1, y_1), \dots, (x_n, y_n)$

Line model $y_i = mx_i + b$



Output: Find $p = (m, b)$ that **minimizes**

$$e = \sum_{i=1}^n (y_i - mx_i - b)^2$$

Minimize Error:

$$\begin{aligned} e &= \sum_{i=1}^n \left(y_i - [x_i \quad 1] \begin{bmatrix} m \\ b \end{bmatrix} \right)^2 \\ &= \left\| \begin{bmatrix} y_1 \\ \vdots \\ y_n \end{bmatrix} - \begin{bmatrix} x_1 & 1 \\ \vdots & \vdots \\ x_n & 1 \end{bmatrix} \begin{bmatrix} m \\ b \end{bmatrix} \right\|^2 \\ &= \|\mathbf{y} - \mathbf{Ap}\|^2 \\ &= \mathbf{y}^T \mathbf{y} - 2(\mathbf{Ap})^T \mathbf{y} + (\mathbf{Ap})^T (\mathbf{Ap}) \end{aligned}$$

$$\frac{de}{dp} = 2\mathbf{A}^T \mathbf{Ap} - 2\mathbf{A}^T \mathbf{y} = 0$$

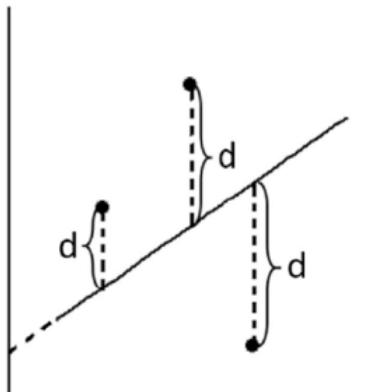
Pseudoinverse of \mathbf{A}

$$\therefore p = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{y}$$

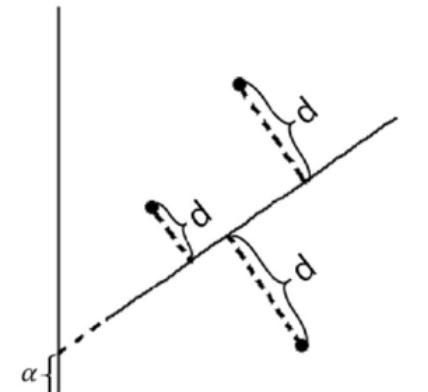
Matlab: $\mathbf{p} = \mathbf{A} \setminus \mathbf{y}$

Ordinary Least Squares – Problems

- **Problem:** Fitting performance affected by **slope**
Fails completely for **vertical lines**
- **Solution:** ‘Total least squares’ using a **perpendicular distance**

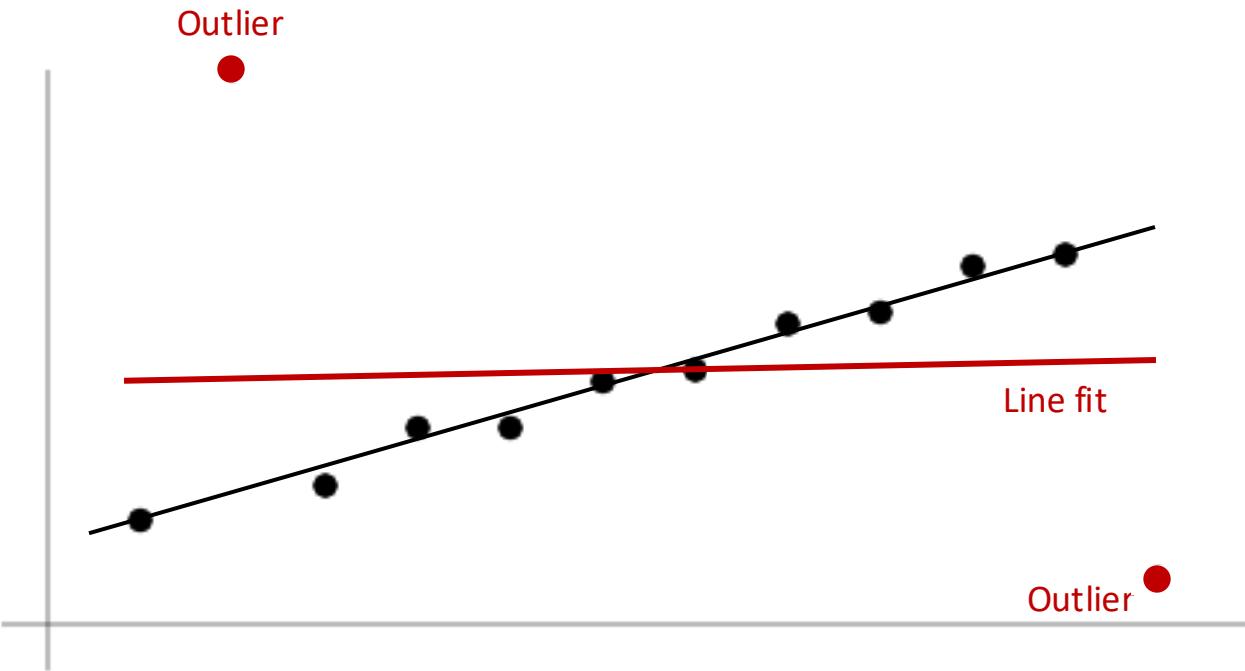


Least Squares



Total Least Squares

Least Squares Fitting



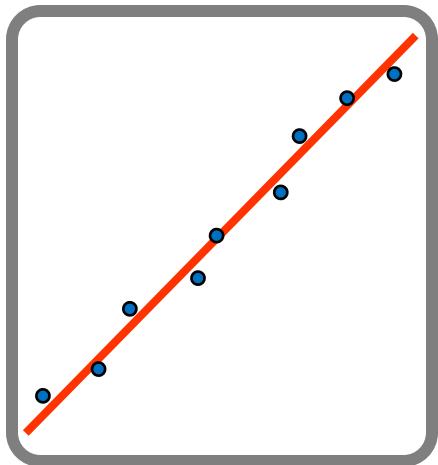
Heavily sensitive
to outliers



Outline

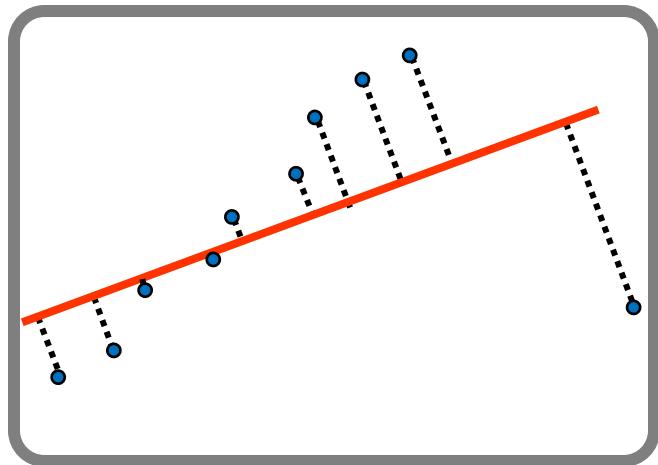
- Edge Detection
 - Derivatives of Image, Derivatives of Gaussian
 - Canny Edge Detector
- Line Fitting
 - Least Squares
 - RANSAC
 - Hough Transform
- Corners
 - Harris Corners
 - SIFT
 - Applications

RANSAC – Motivation



*2D point → n-dim data point
2D line → parametric model*

*How do we
fit a model
to data points?*



*Detecting outliers
is crucial!*



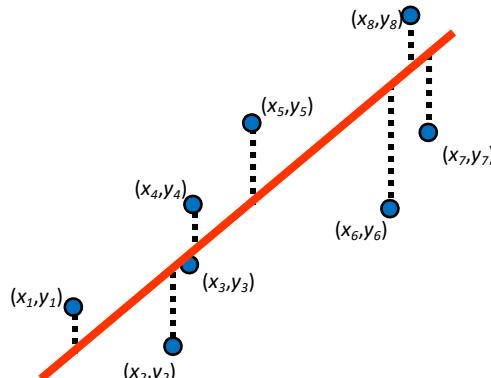
Model Fitting – 2D Line

Line fitting

- Input:** Set of n 2D points (x_i, y_i)
- Output:** Line $\Theta = (a, b)$ minimizing **fitting error**

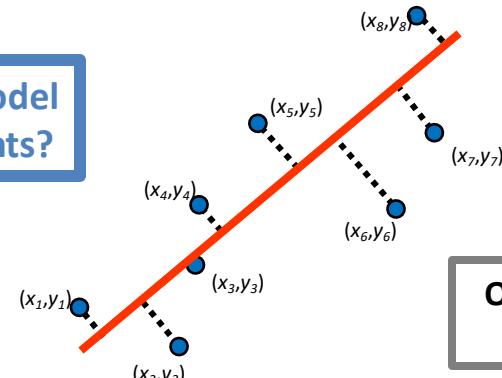
Line equation: $y = ax + b$

Fitting cost: YOU choose it 😊



Vertical Distance

How 'far' is my model from my data points?



Orthogonal Distance

Fitting Cost - Norm

Let:

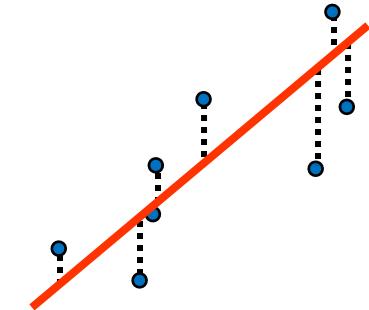
$$\mathbf{x} \in \mathbb{R}^n$$

$$\mathbf{x} = (x_1, \dots, x_n)$$



L_p norm:

$$\|\mathbf{x}\|_p = \left(\sum_{i=1}^n |x_i|^p \right)^{1/p}$$



$$\begin{aligned} \mathbf{x} &= (x_1, x_2, x_3) \\ &= (-3, 0, 4) \end{aligned}$$

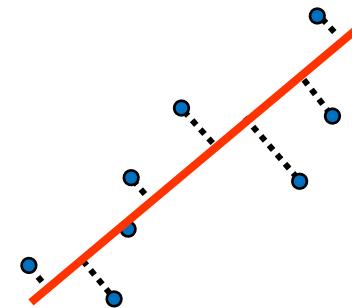


$$\|\mathbf{x}\|_1 = ?$$

$$\|\mathbf{x}\|_1 = |-3| + |0| + |4| = 7$$

$$\|\mathbf{x}\|_2 = ?$$

$$\|\mathbf{x}\|_2 = \sqrt{(-3)^2 + 0^2 + 4^2} = \sqrt{25} = 5$$



Fitting Cost - Norm

Let:

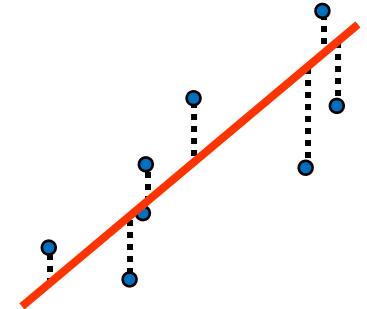
$$\mathbf{x} \in \mathbb{R}^n$$

$$\mathbf{x} = (x_1, \dots, x_n)$$



L_p norm:

$$\|\mathbf{x}\|_p = \left(\sum_{i=1}^n |x_i|^p \right)^{1/p}$$



For $n = 1$:

(1D case)

Without proof here

$$(L_2): \arg \min_u \sum_{i=1}^n |u - x_i|^2$$

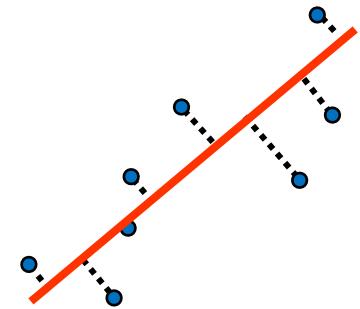
Solution \rightarrow Average

$$u = \frac{1}{n} \sum_{i=1}^n x_i$$

$$(L_1): \arg \min_u \sum_{i=1}^n |u - x_i|$$

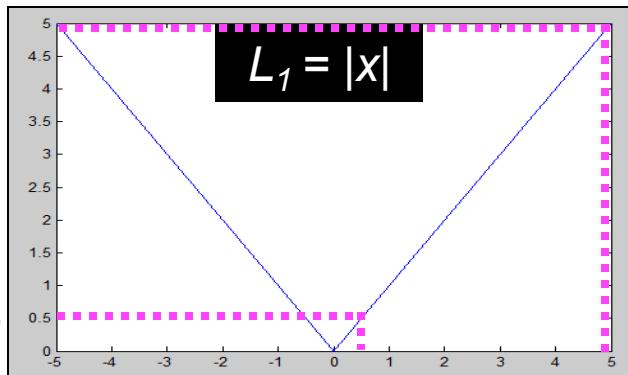
Solution \rightarrow Median

if $a < b < c < d$,
the median of the list $\{a, b, c, d\}$
is the mean of b and c ; i.e., it is $(b + c)/2$.



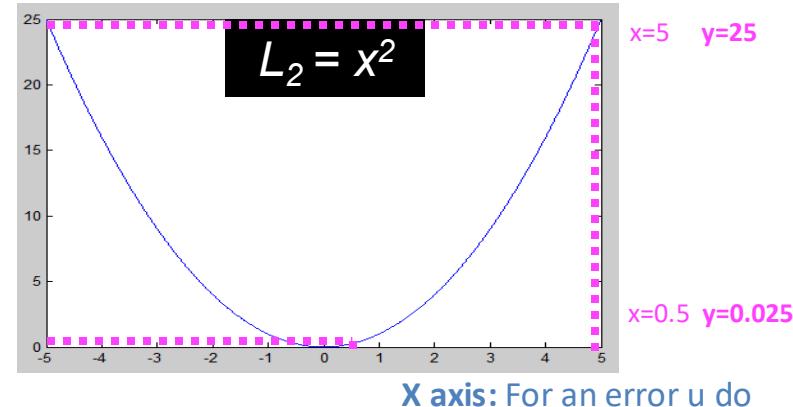
Fitting Cost - Norm

x=5 y=5
x=0.5 y=0.5



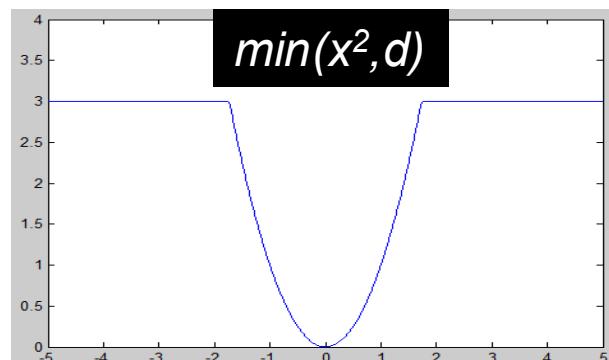
Do you need ...
... heavy penalty
on big errors?
Choose $L_2 \rightarrow \rightarrow$
... noticeable penalty
on tiny errors?
← ← Choose L_1

Y axis: How much 'penalty' u pay



X axis: For an error u do

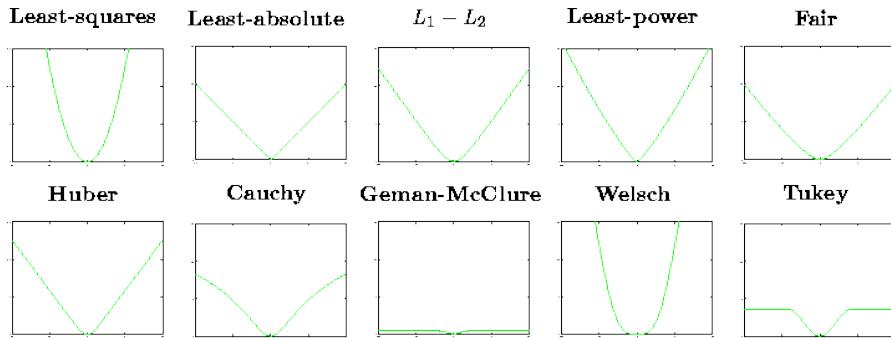
$d \rightarrow$
'Robustifier' $\rightarrow \rightarrow$
Influence of
outliers 'saturates'



Which one to
choose?
Design/play with
your cost function

Cost Functions

Go beyond just norms → aka '*cost functions*'



Rules of thumb (but look things up on your own):

- **L2, L1, Huber**: These are the most popular, depending on the applications
- **Huber**: Best of both worlds. Treats small errors like L2 (tolerates them) and large errors like L1 (robustness).
- **Dimitris**: For my problems – I use a lot the Geman-McClure one. Chosen after experimentation.

YOU *choose/define* it 😊

Play with parameters & judge
for your own task/data !

type	$\rho(x)$
L_2	$x^2/2$
L_1	$ x $
$L_1 - L_2$	$2(\sqrt{1+x^2/2} - 1)$
L_p	$\frac{ x ^p}{p}$
“Fair”	$c^2[\frac{ x }{c} - \log(1 + \frac{ x }{c})]$
Huber	$\begin{cases} x^2/2 & \text{if } x \leq k \\ k(x - k/2) & \text{if } x \geq k \end{cases}$
Cauchy	$\frac{c^2}{2} \log(1 + (x/c)^2)$
Geman-McClure	$\frac{x^2/2}{1+x^2}$
Welsch	$\frac{c^2}{2} [1 - \exp(-(x/c)^2)]$
Tukey	$\begin{cases} \frac{c^2}{6} (1 - [1 - (x/c)^2]^3) & \text{if } x \leq c \\ (c^2/6) & \text{if } x > c \end{cases}$

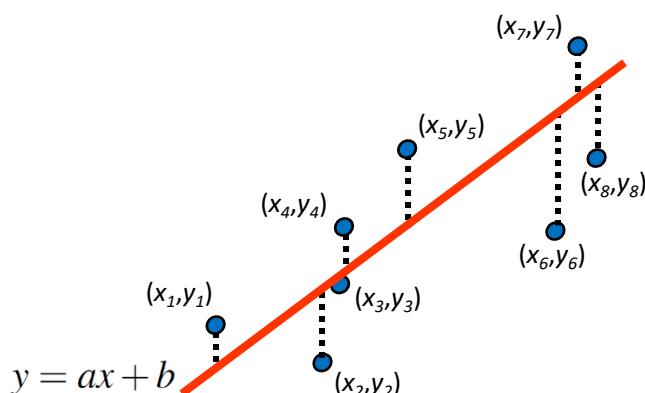
Model Fitting – Under Noise

Data points: (x_i, y_i)

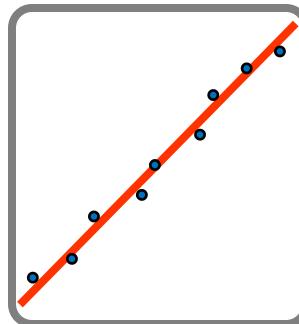
Line model: $y = ax + b$

Parameters: $\Theta = (a, b)$

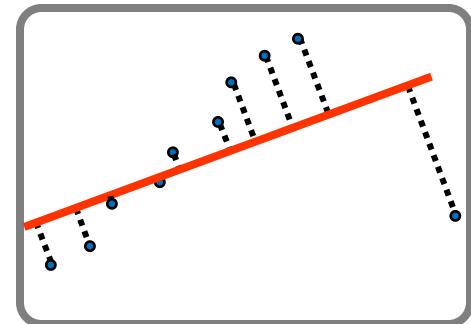
L_1 cost:
$$\min_{\Theta=(a,b)} \sum_{i=1}^n |y_i - (ax_i + b)|$$



But remember...



Only inliers



Effects due to single outlier

Detecting outliers is very important!

Solution: RANSAC

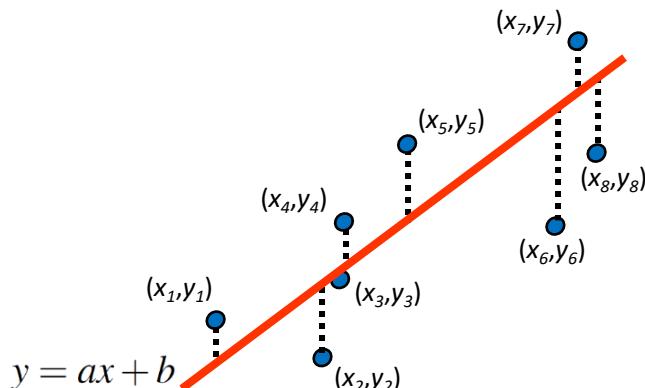
Model Fitting – Under Noise

Data points: (x_i, y_i)

Line model: $y = ax + b$

Parameters: $\Theta = (a, b)$

L_1 cost:
$$\min_{\Theta=(a,b)} \sum_{i=1}^n |y_i - (ax_i + b)|$$



Motivation



Detecting outliers
is very important!



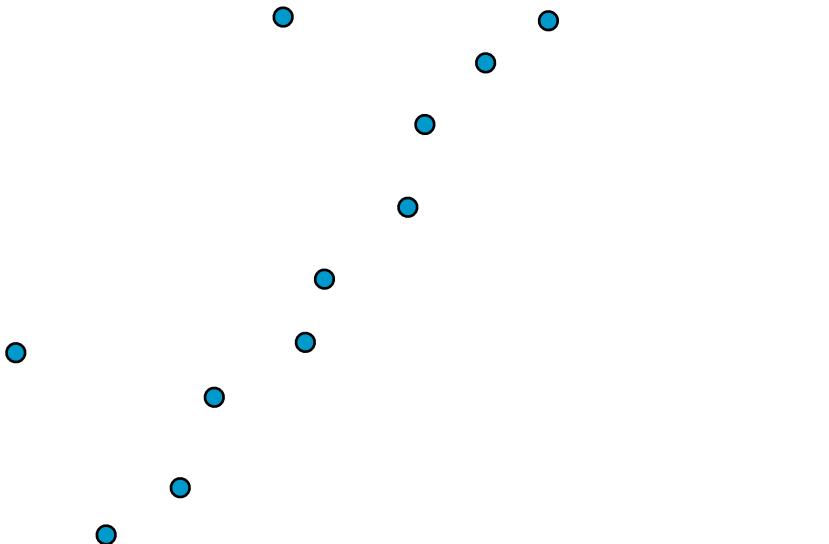
Solution:
RANSAC



*Also for Vision
(coming soon)*



RANSAC – Example – Lines



Input:

Set of points

Output: Estimate line equation & determine **inliers/outliers**

RANSAC – Example – Lines

Which one do you like the most?

Automate this with RANSAC

RANSAC – Example – Lines

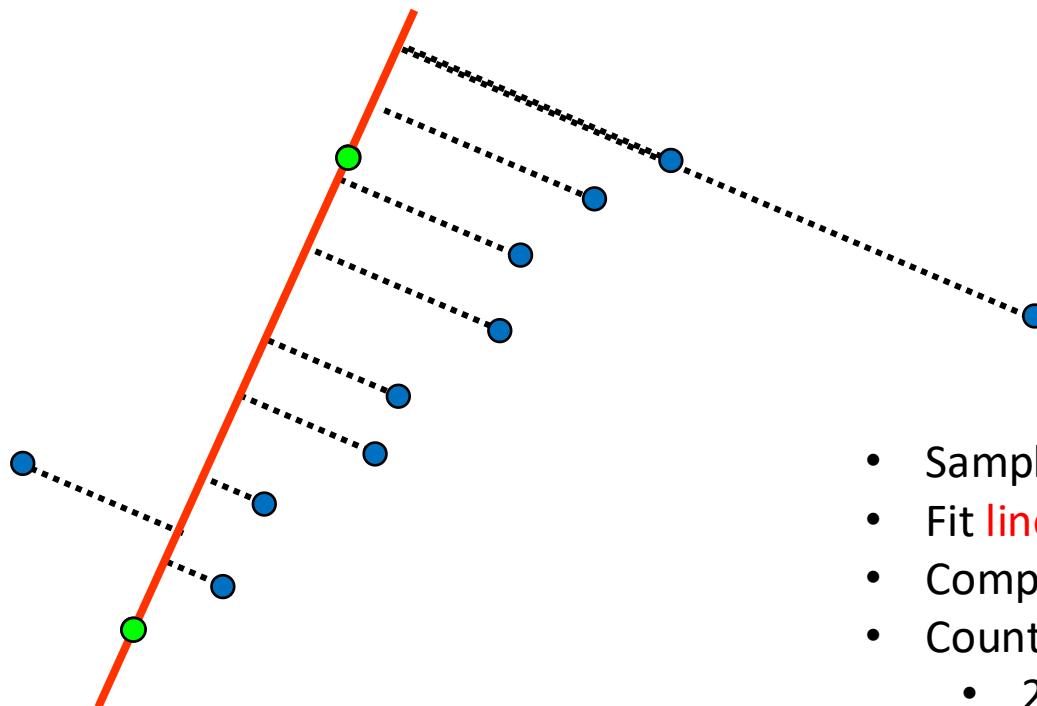


RANSAC - RANdom SAmple Consensus
by Fischler and Bolles, 1981

Very *easy* method and *good* results

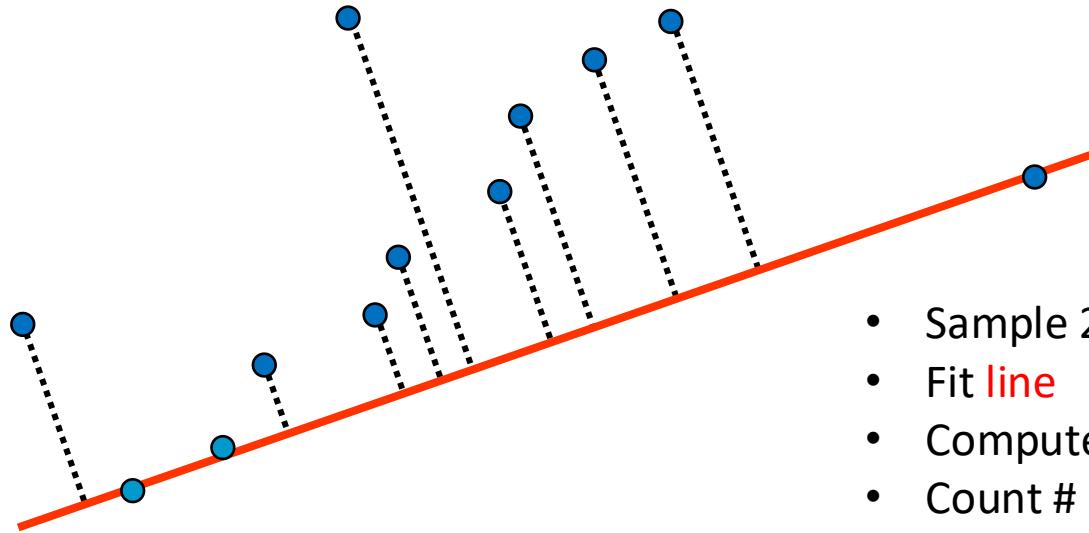
Aims to **maximize** the **number #** of
inliers ('consensus set maximization')
through **random sampling**

RANSAC – Example – Lines



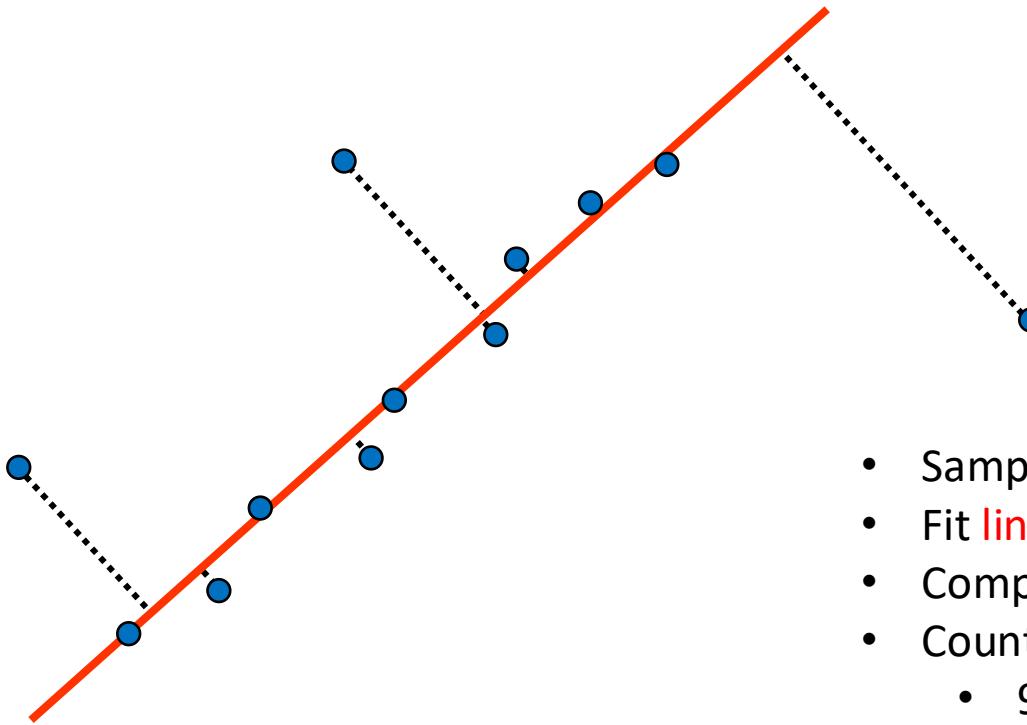
- Sample 2 points randomly
- Fit **line**
- Compute distances
- Count # of **inliers**
 - 2 out of 13

RANSAC – Example – Lines



- Sample 2 points randomly
- Fit **line**
- Compute distances
- Count # of **inliers**
 - 3 out of 13

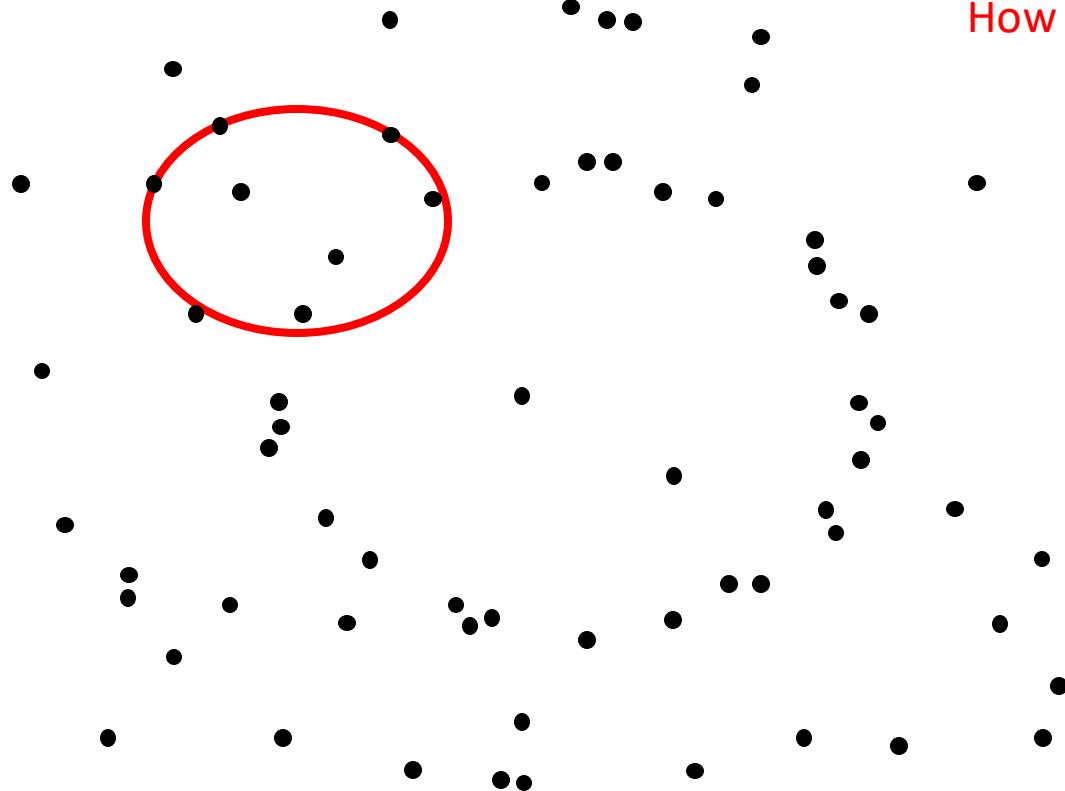
RANSAC – Example – Lines



- Sample 2 points randomly
- Fit **line**
- Compute distances
- Count # of **inliers**
 - 9 out of 13



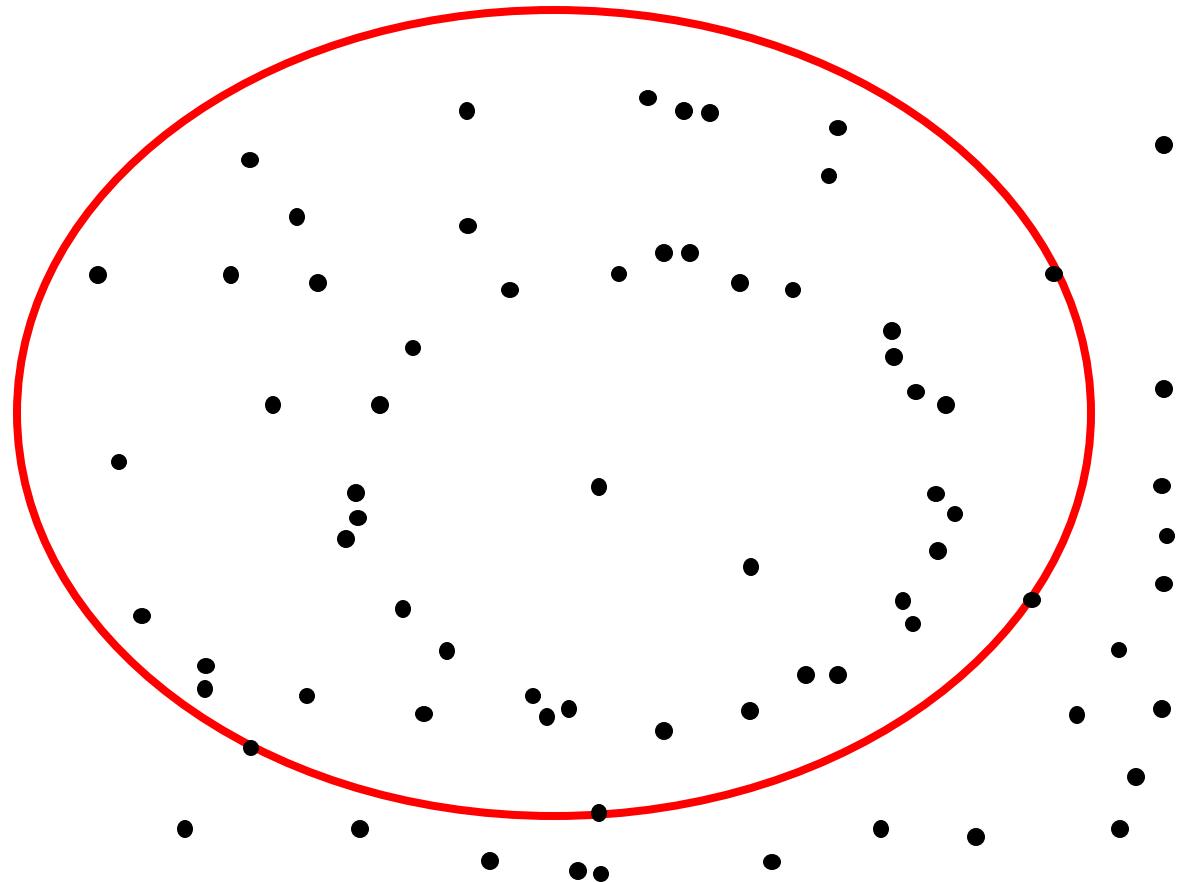
RANSAC – Example – Circles



How many points for a circle?

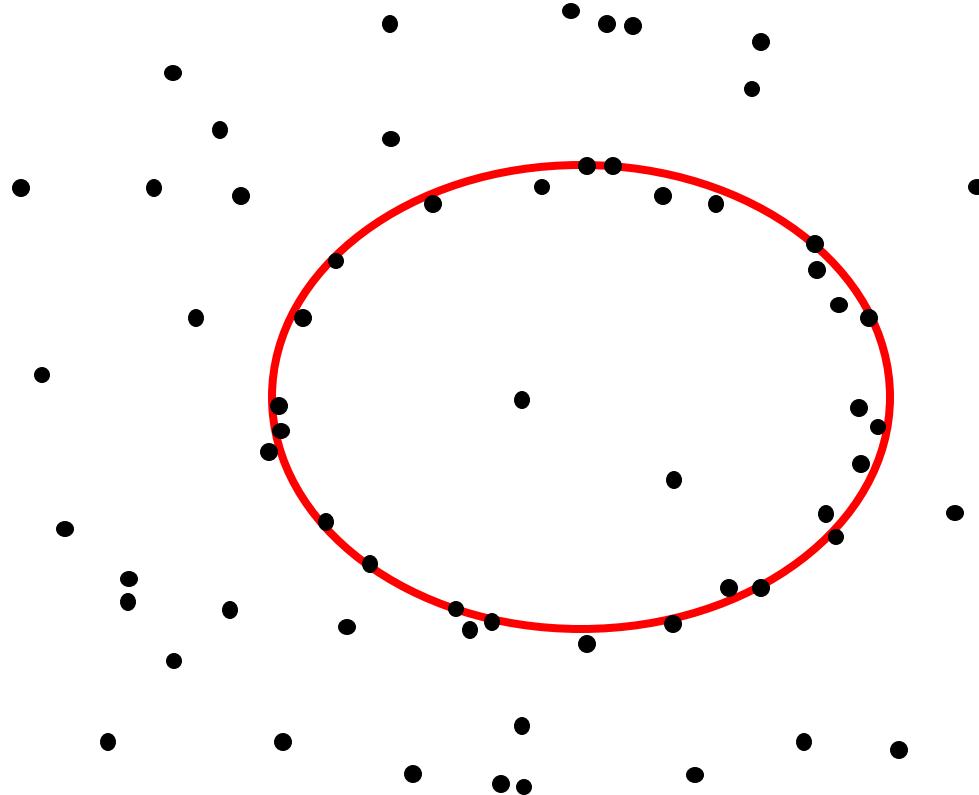
- Sample 3 points randomly
 - Fit **circle**
 - Compute distances
 - Count # of **inliers**
- 5

RANSAC – Example – Circles



- Sample 3 points randomly
 - Fit **circle**
 - Compute distances
 - Count # of **inliers**
- 4

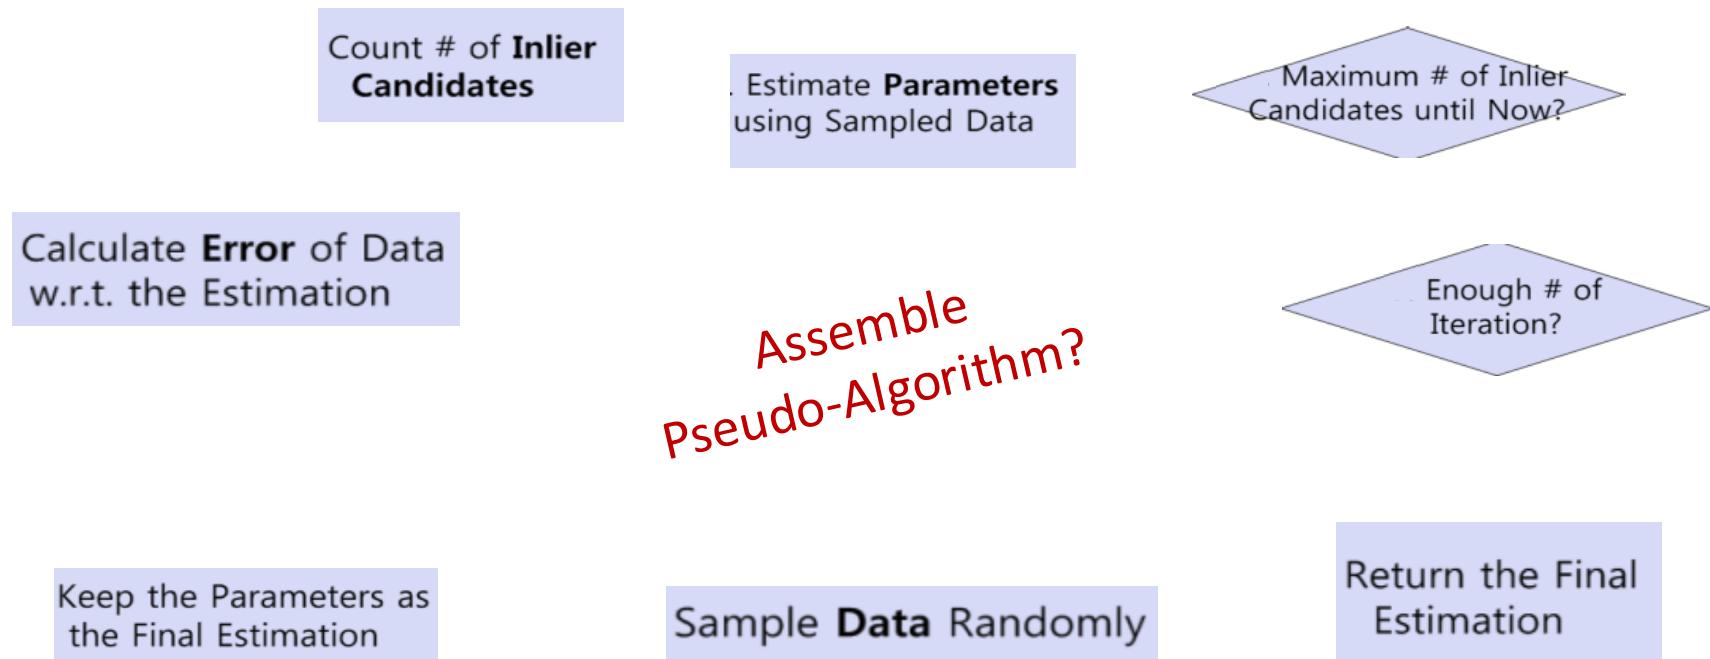
RANSAC – Example – Circles



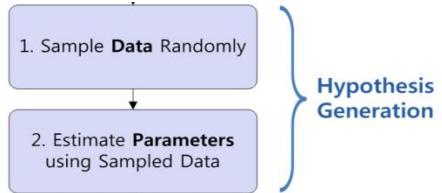
- Sample 3 points randomly
- Fit **circle**
- Compute distances
- Count # of **inliers**

-23 ✓

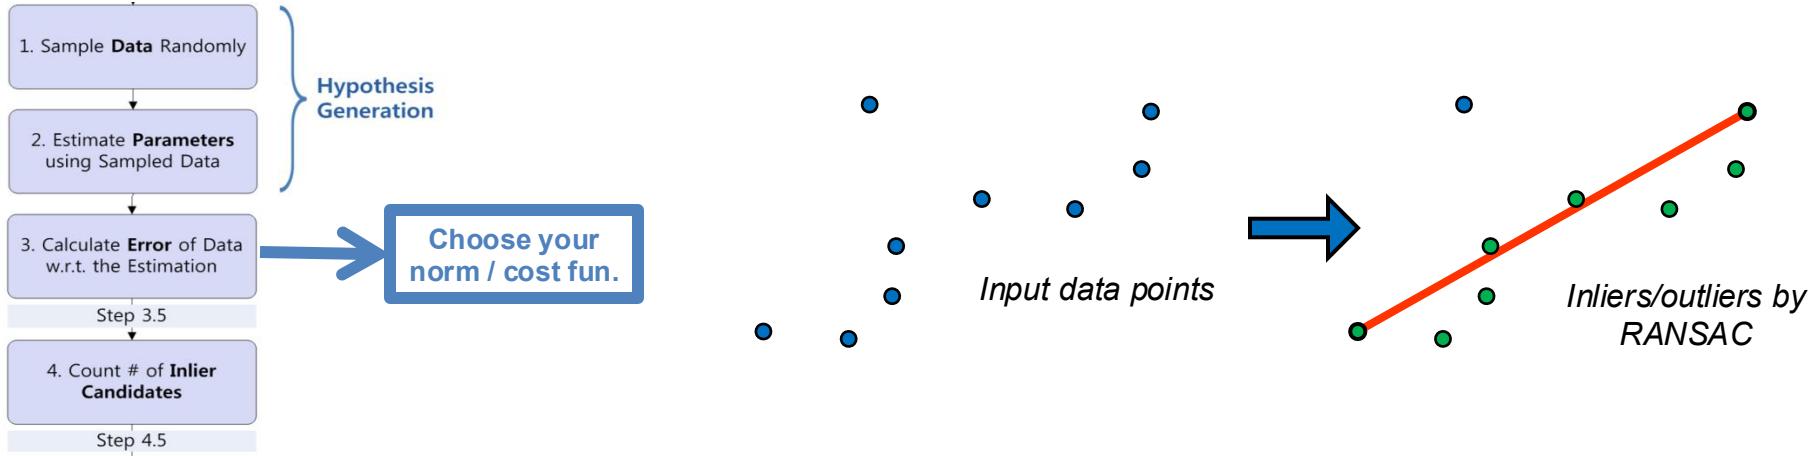
RANSAC – Quiz



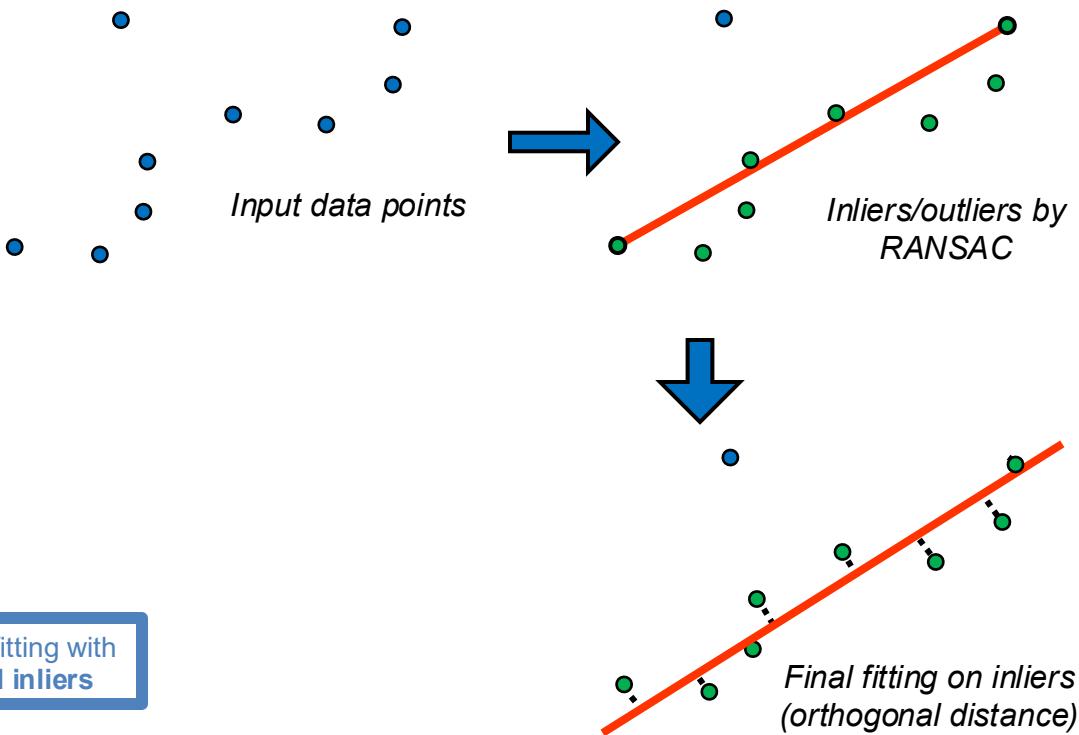
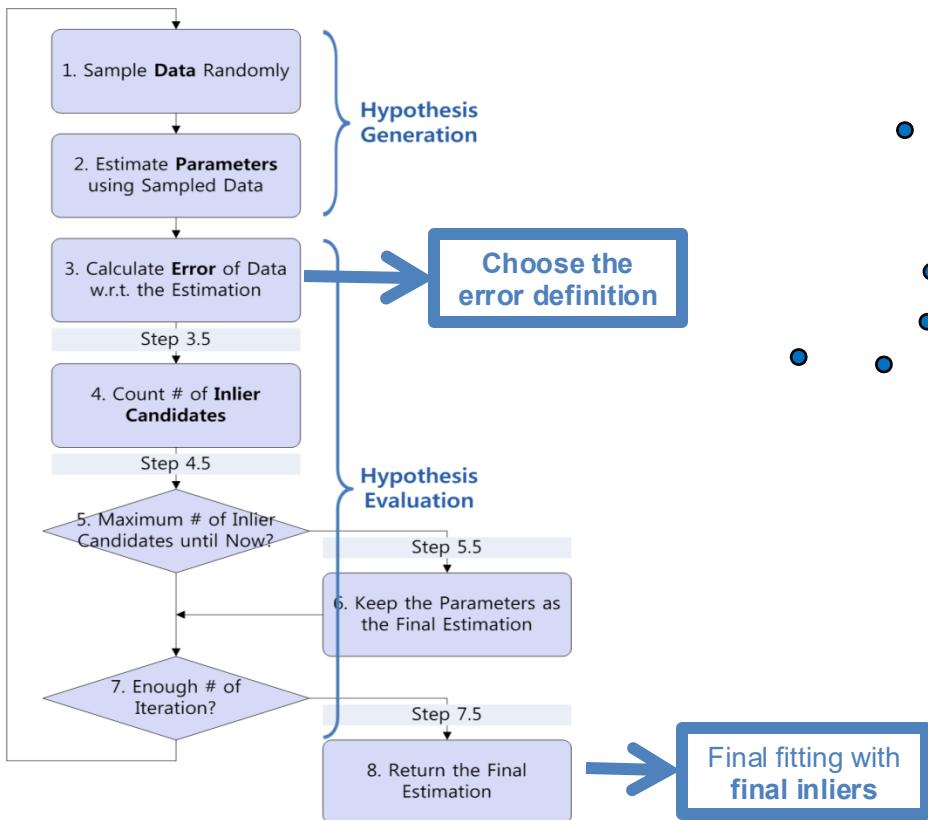
RANSAC – Overview



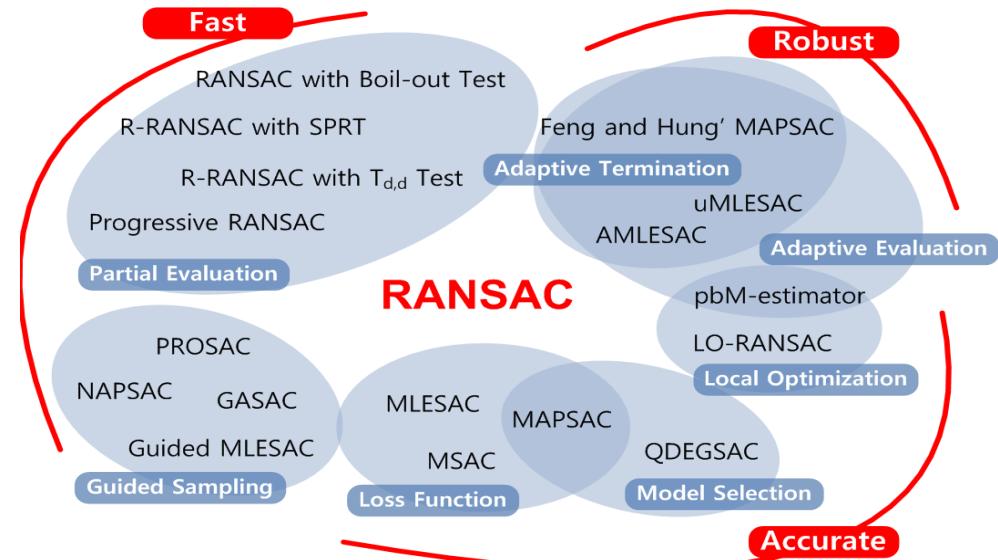
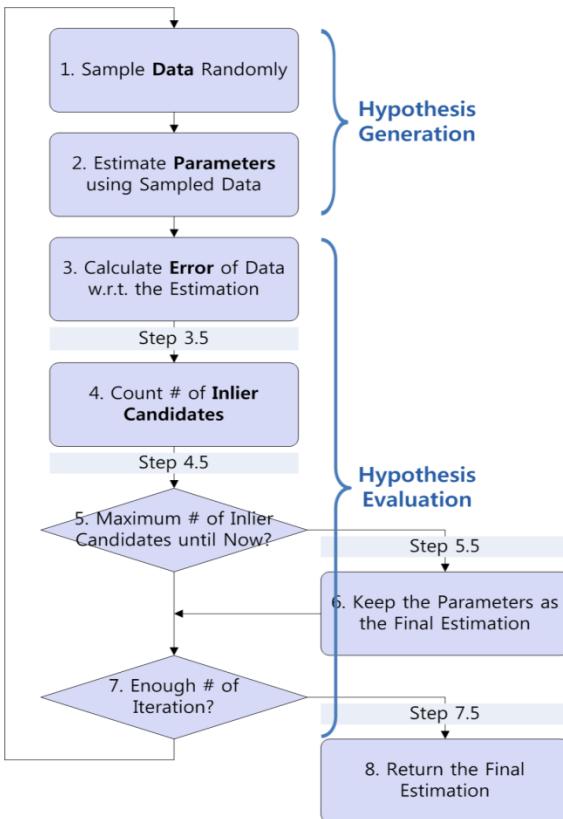
RANSAC – Overview



RANSAC – Overview



RANSAC – Family



Outside the course scope

RANSAC – Family – References

- Martin A. Fischler und Robert C. Bolles, "Random Sample Consensus: A Paradigm for Model Fitting with Applications to Image Analysis and Au
Cartography", Comm of the ACM, 1981
- D. Capel. "An effective bail-out test for RANSAC consensus scoring", BMVC, 2005
- S. Choi and J.-H. Kim. "Robust regression to varying data distribution and its application to landmark-based localization", IEEE SMC, 2008 [uMLESAC]
- O. Chum and J Matas. "Randomized RANSAC with $T_{d,d}$ test", BMVC, 2002
- O. Chum and J. Matas. "Matching with PROSAC - Progressive Sample Consensus", CVPR, 2005
- O. Chum and J. Matas. "Optimal randomized RANSAC", TPAMI, 2008.
- O. Chum, J. Matas, and S. Obdrzalek. "Enhancing RANSAC by generalized model optimization", ACCV, 2004.
- C.L. Feng and Y.S. Hung. "A robust method for estimating the fundamental matrix". In Digital Image Computing: Techniques and Applications, 2003
- J.-M. Frahm and M. Pollefeys. "RANSAC for (quasi-)degenerate data (QDEGSAC)", CVPR, 2006.
- A. Konouchine, V. Gaganov, and V. Veznevets. "AMLESAC: A new maximum likelihood robust estimator", Int. Conf. on Computer Graphics and Vision (GrapiCon), 2005.
- J. Matas and O. Chum. "Randomized RANSAC with sequential probability ratio test", ICCV, 2005
- D.R. Myatt, P.H.S Torr, S.J. Nasuto, J.M. Bishop, R. Craddock. "NAPSAC: High noise, high dimensional robust estimation - it's in the bag", BMVC, 2002
- D. Nister. "Preemptive RANSAC for live structure and motion estimation", ICCV, 2003
- R. Raguram, J.-M. Frahm, and M. Pollefeys. "A comparative analysis of RANSAC techniques leading to adaptive real-time random sample consensus", ECCV, 2008
- V. Rodehorst and O. Hellwich. "Genetic Algorithm SAmples Consensus (GASAC)- a parallel strategy for robust parameter estimation", CVPR Workshop, 2006
- R. Subbarao and P. Meer. "Beyond RANSAC: User independent robust regression", CVPR Worksop, 2006
- B. J. Tordoff and D. W. Murray. "Guided-MLESAC: Faster image transform estimation by using matching priors", TPAMI, 2005
- P.H.S. Torr and A. Zisserman. "MLESAC: A new robust estimator with application to estimating image geometry", CVIU, 2000
- And many others



Outside the
course
scope

RANSAC – Parameters

$$p = 1 - (1 - w^n)^k$$

$$k = \frac{\log(1 - p)}{\log(1 - w^n)}$$

p → Reliability of RANSAC result

k → Number of iterations

w → Probability that a point is an inlier

n → Minimal number of samples for model (line, circle, projective transf., etc)

We stitch 2 images (spoiler: later today)

Find the number of RANSAC iterations for

- 95% reliability, and
- 20% inlier probability
- for SIFT corresp. (spoiler: n=4 corr.)

$$\begin{aligned} k &= \frac{\log(1-p)}{\log(1-w^n)} & p &= 0,95 \text{ --- RANSAC reliability} \\ &= \frac{\log(1-0,95)}{\log(1-0,20^n)} & w &= 0,20 \text{ --- inlier probability} \\ &= 1870,8344 & n &= 4 \text{ --- we need 4 corresp.} \\ &\approx 1871 && \text{as minimum to stitch images with the projective transform.} \end{aligned}$$



RANSAC – Parameters

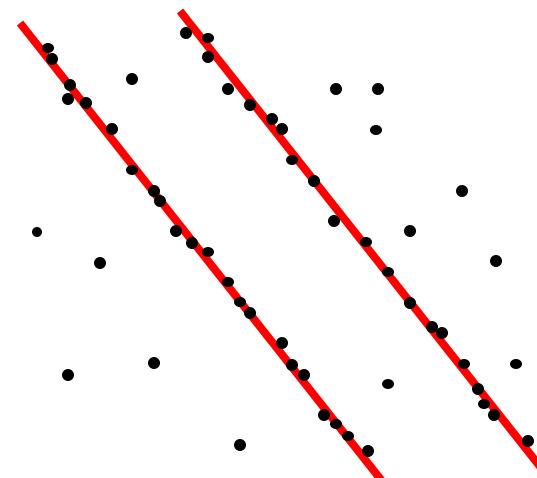
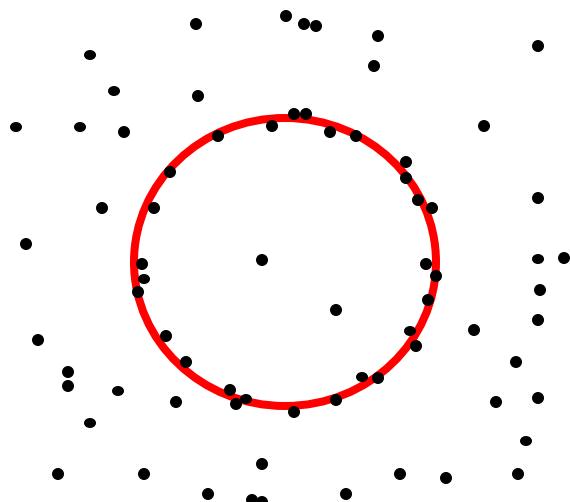
Sample size	proportion of outliers						
	5%	10%	20%	25%	30%	40%	50%
2	2	3	5	6	7	11	17
4	3	5	9	13	17	34	72
5	4	6	12	17	26	57	146
7	4	8	20	33	54	163	588
8	5	9	26	44	78	272	1177

The **number of RANSAC iterations** required to ensure ...

[input] ... with a **probability 0.99**, and a given number of **minimal points** and a **proportion of outliers**,

[output] that **at least one sample has no outliers**.

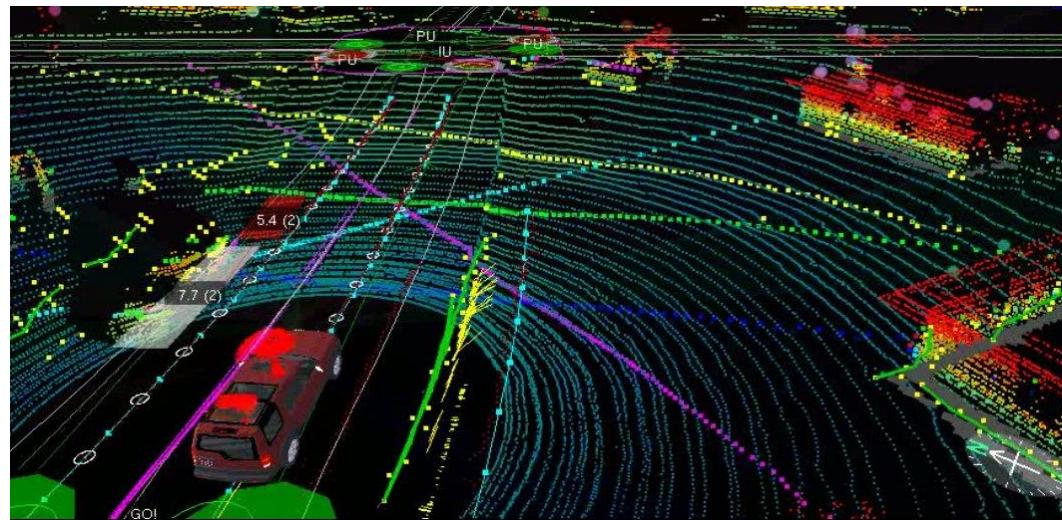
- Think about
 - Advantages & disadvantages of RANSAC
 - How to apply for detecting ellipses, planes, parallel lines, etc?



RANSAC – Applications

3D Ground plane extraction on LIDAR 3D point cloud

How many points?



RANSAC – Applications



RANSAC – Pros & Cons

Advantages

- Robust estimation of model parameters
- Applicable for large number of parameters
- High degree of accuracy
- Even with outliers
- Easy to implement



Disadvantages

- No upper bound on runtime (so cap w. threshold)
- Accuracy may be suboptimal for low iter. Threshold
- Suboptimal when # of inliers <50% (so, check variants)
- Requires problem-specific thresholds
- Runtime grows quickly with
 - % of outliers and
 - number of parameters
- Not good for getting multiple line fits ←

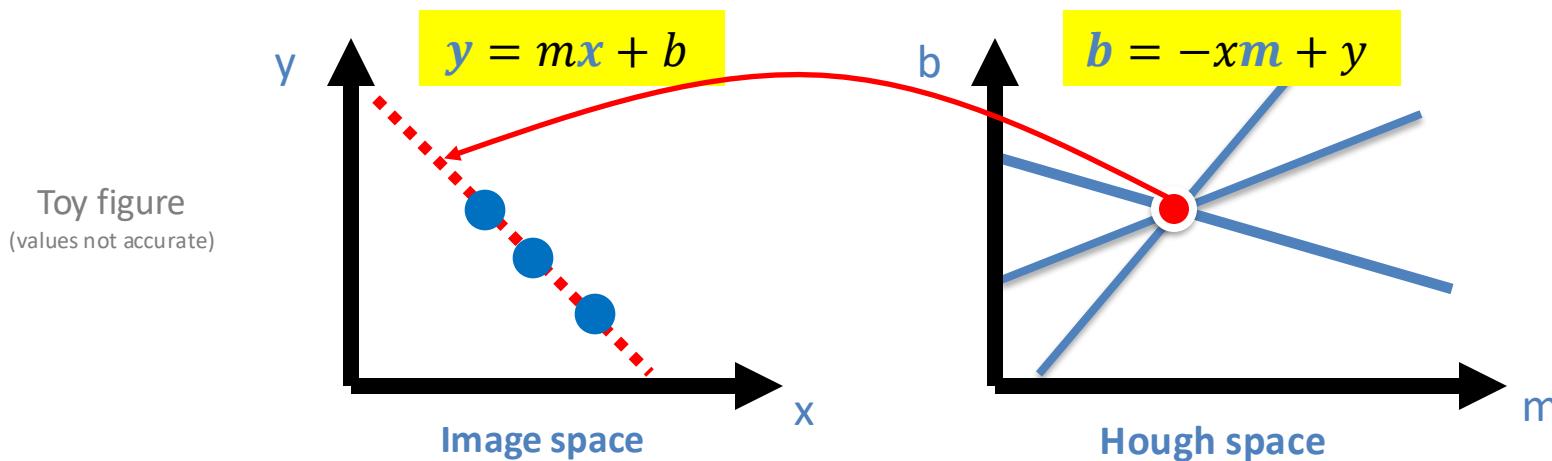
Solution?

Outline

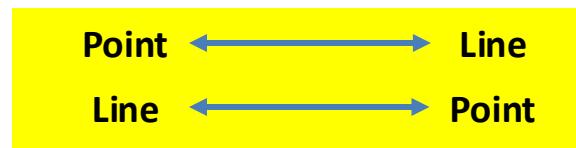
- Edge Detection
 - Derivatives of Image, Derivatives of Gaussian
 - Canny Edge Detector
- Line Fitting
 - Least Squares
 - RANSAC
 - Hough Transform
- Corners
 - Harris Corners
 - SIFT
 - Applications

Hough Transform – Fitting Multiple Lines

- Given a binary **edge** image
- Find the **lines** (or curves) that best explain the data points in **parameter space**
- Parameter space → Called '**Hough space**'



P.V.C. Hough,
'Machine Analysis of
Bubble Chamber
Pictures', 1959

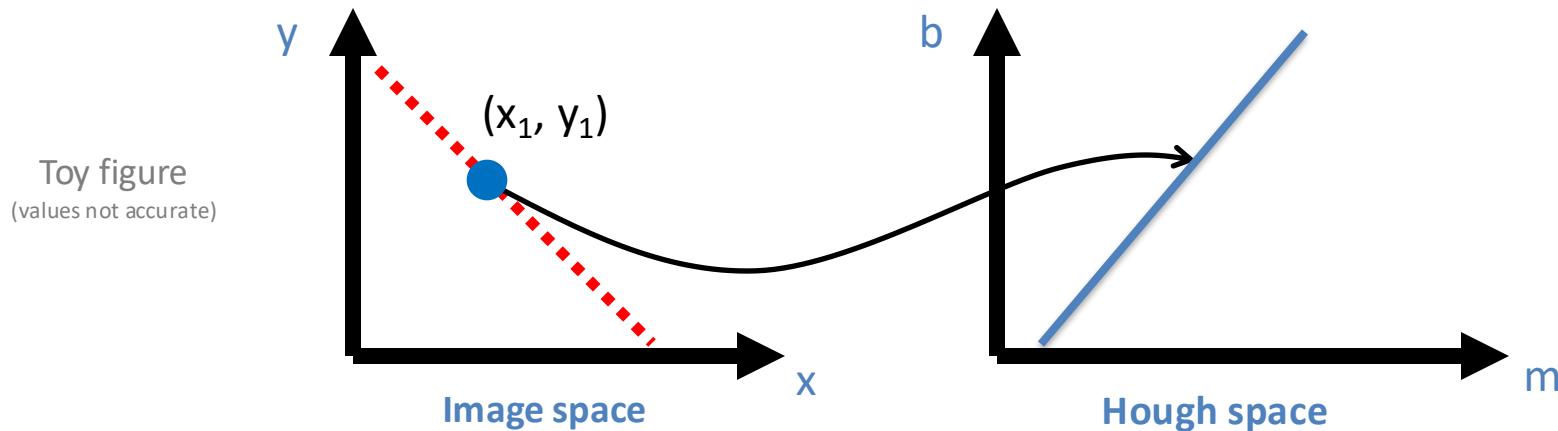


Map Point (x_1, y_1) → Line in Hough Space

$$y_1 = mx_1 + b$$

$$b = -x_1m + y_1$$

Rewrite
equation in
terms of
 m and b

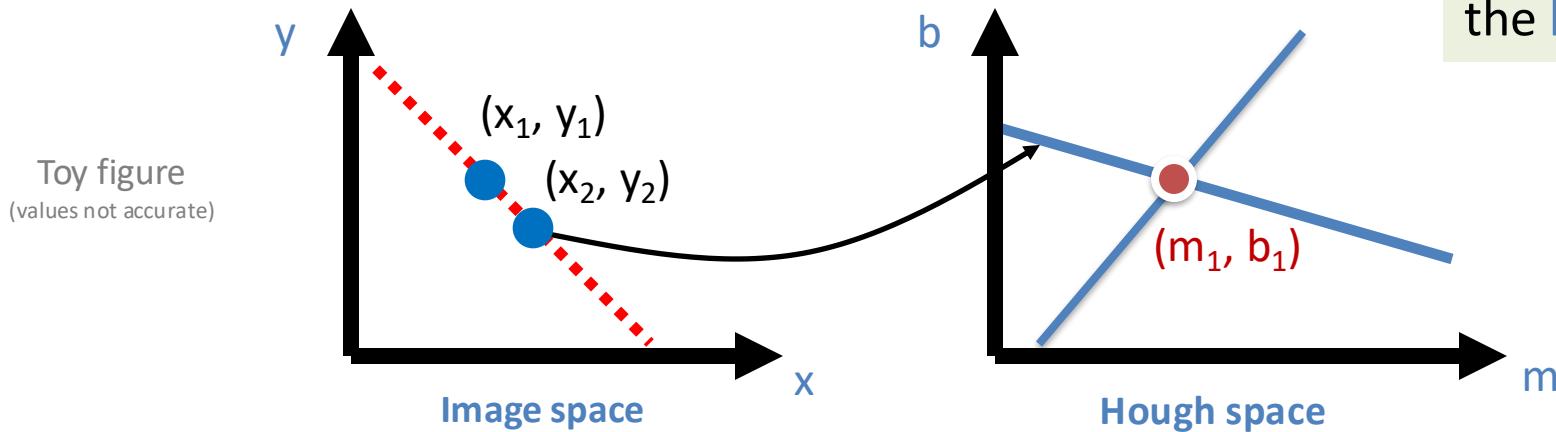


Map Point (x_2, y_2) → Line in Hough Space

$$y_2 = mx_2 + b$$

$$b = -x_2m + y_2$$

Intersection point satisfies (x_1, y_1) & (x_2, y_2) & by extension the line

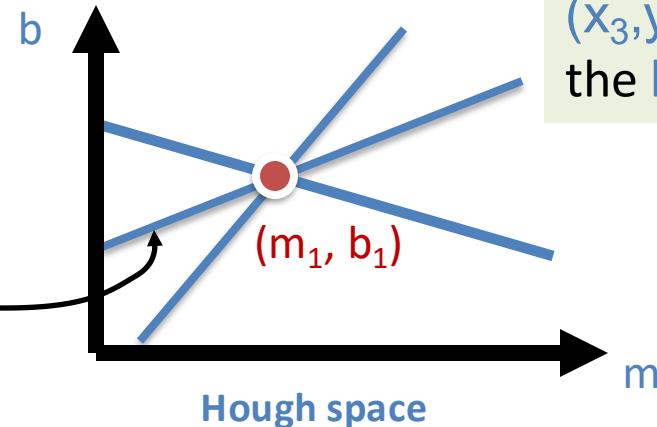
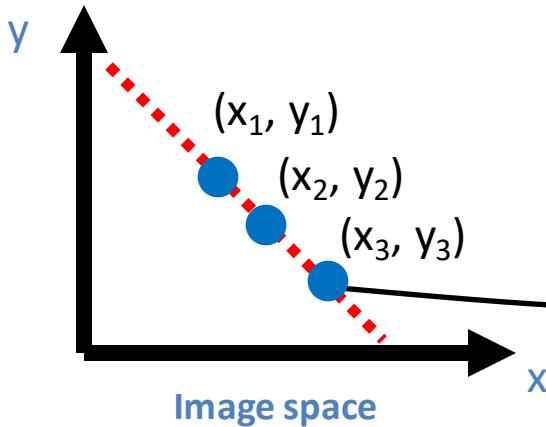


Map Point (x_3, y_3) → Line in Hough Space

$$y_3 = mx_3 + b$$

$$b = -x_3m + y_3$$

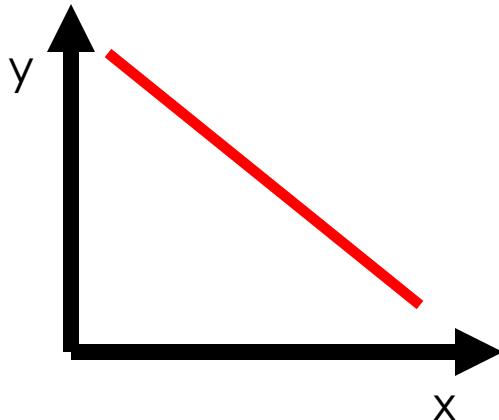
Toy figure
(values not accurate)



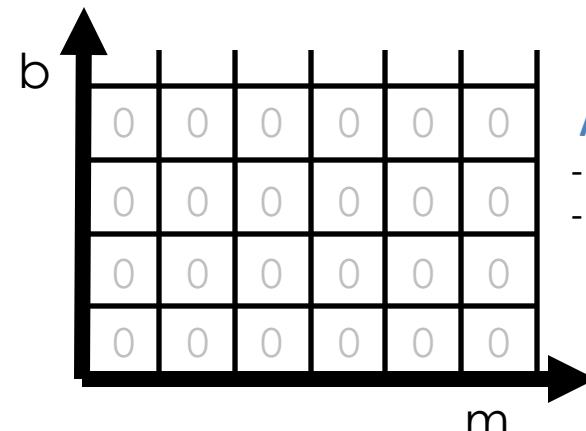
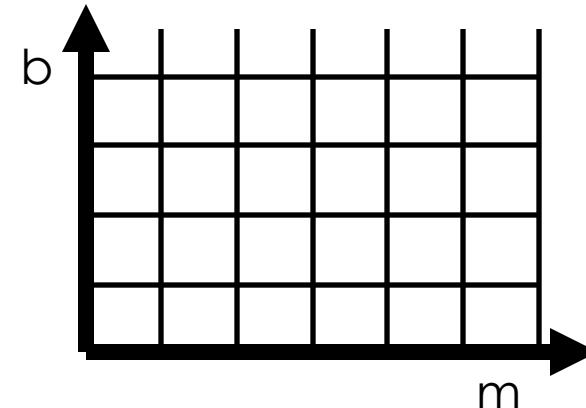
Intersection point (m_1, b_1) satisfies points (x_1, y_1) , (x_2, y_2) , (x_3, y_3) & thus the line

Hough Transform – Accumulator

Toy figure
(values not accurate)



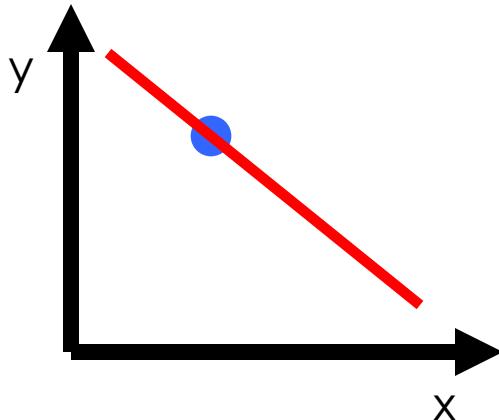
- For each edge^{el} → - Draw a line in Hough space
- Add +1 in Accumulator cells



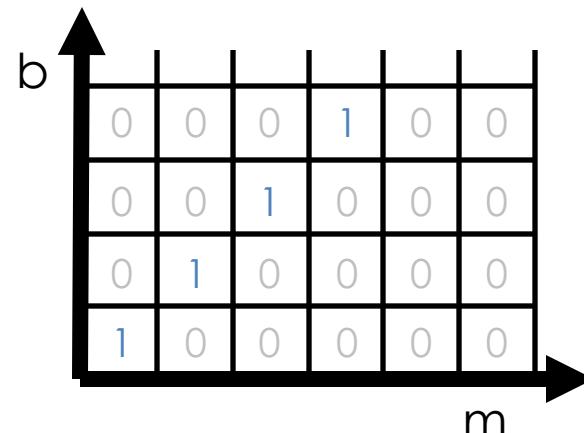
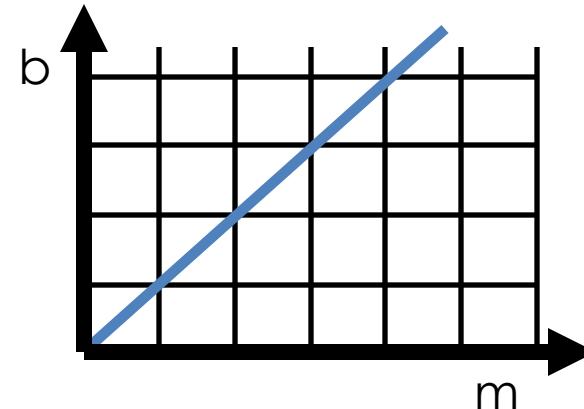
Accumulator
- Discrete grid
- Initialize w 0s

Hough Transform – Accumulator

Toy figure
(values not accurate)

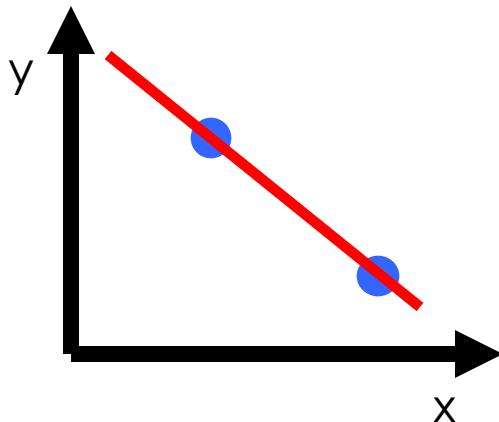


- For each edgel → - Draw a line in Hough space
- Add +1 in Accumulator cells

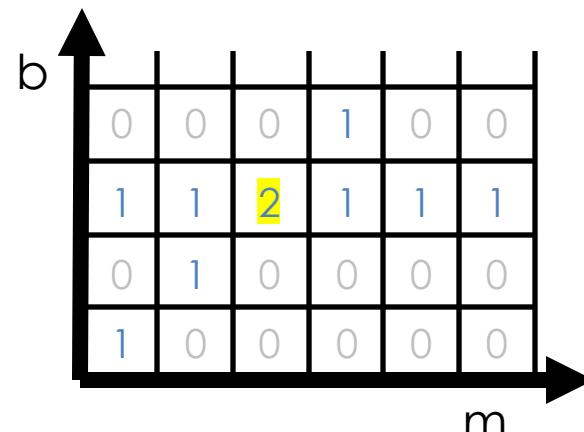
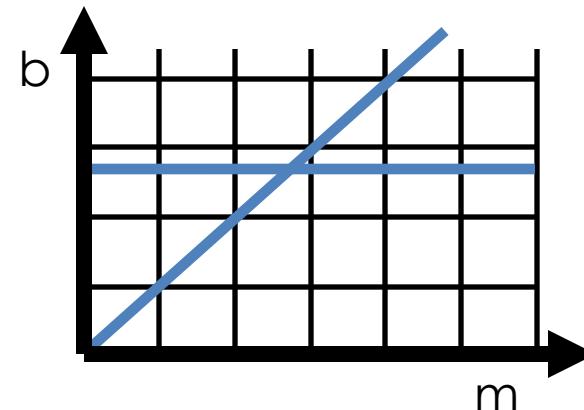


Hough Transform – Accumulator

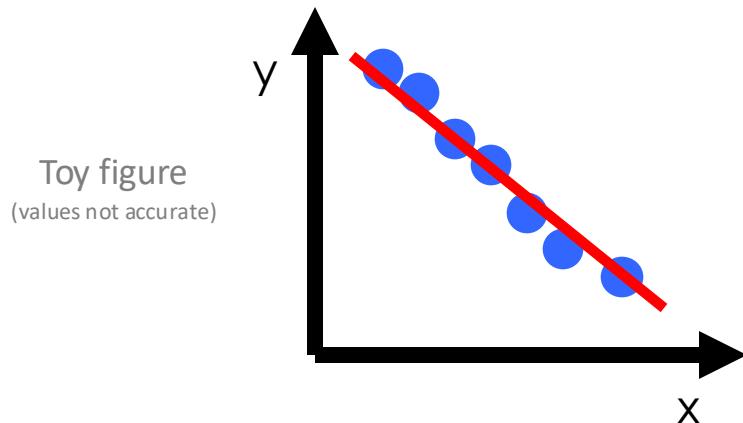
Toy figure
(values not accurate)



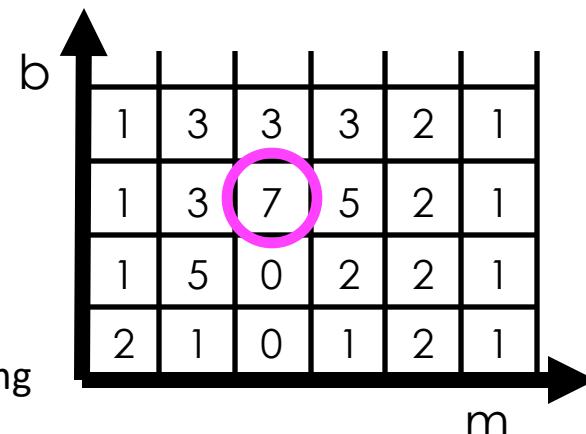
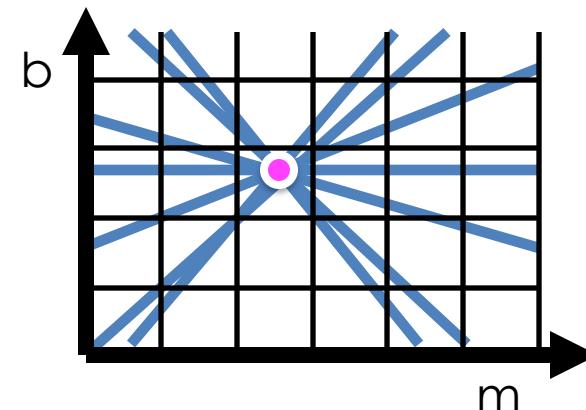
- For each edgel → - Draw a line in Hough space
- Add +1 in Accumulator cells



Hough Transform – Accumulator



- For each edgel → - Draw a line in Hough space
- Add +1 in Accumulator cells
- Process all pixels → Accumulate values ('voting')
- Find local maxima → Can be many points
due to many lines in image
- Map each maximum Point in Hough space → Line in Img



Hough Transform

- Detect multiple lines in image (from multiple local maxima)
- Easily extended for circles & ellipses
- Computationally efficient
- Problem: Parameters (m, b) are unbounded (slope m can be infinite)
Need an infinitely large Accumulator (memory hungry)
- Solution? See next

Hough Transform – Polar Form

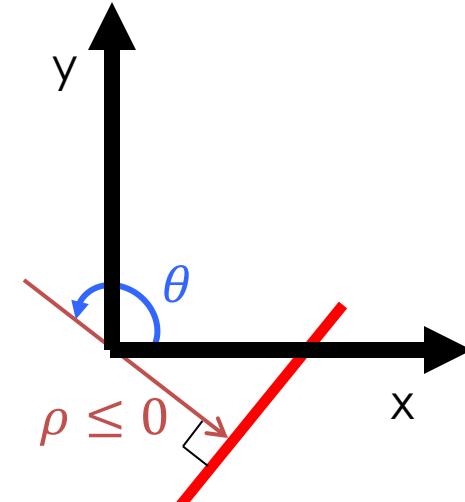
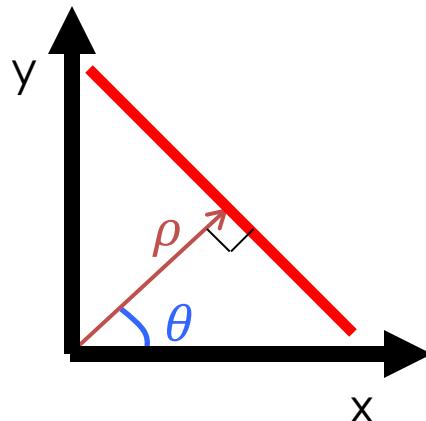
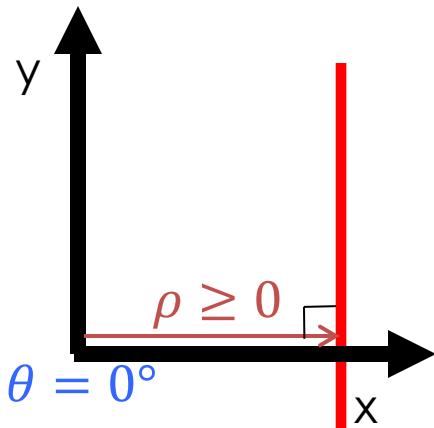
Parameter space → **Polar Representation**

$$x \cos \theta + y \sin \theta = \rho$$

All parameters finite

$0 \leq \theta < 180^\circ$

ρ can be negative/positive

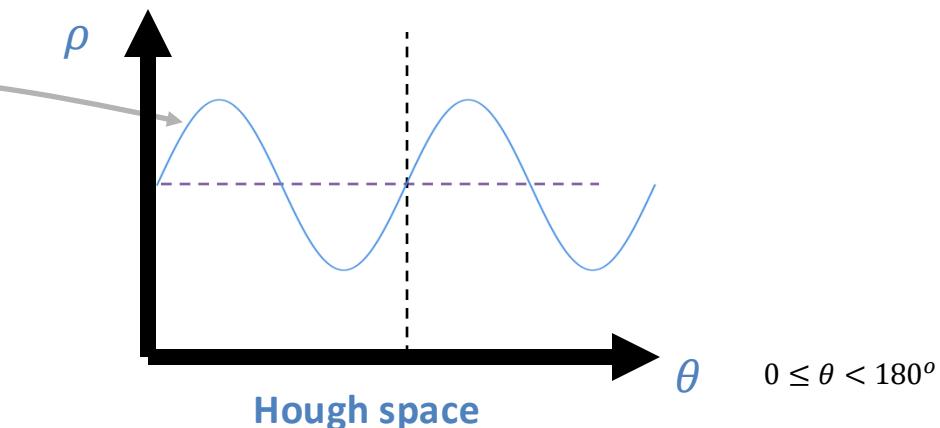
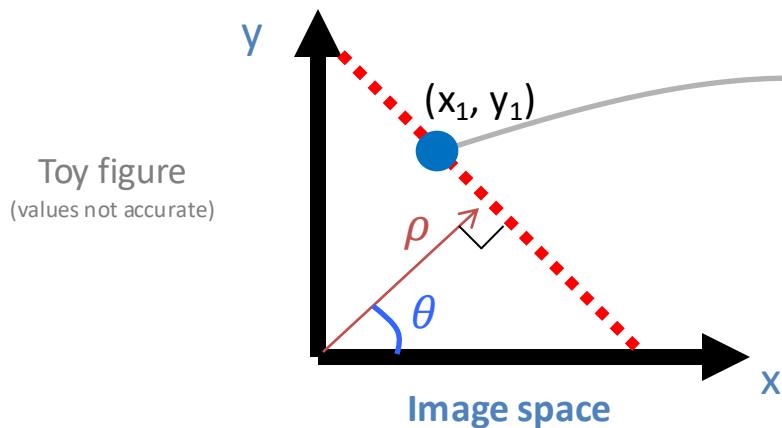


Map Point (x_1, y_1) → Sinusoid in Hough Space

CV

$$x_1 \cos \theta + y_1 \sin \theta = \rho$$

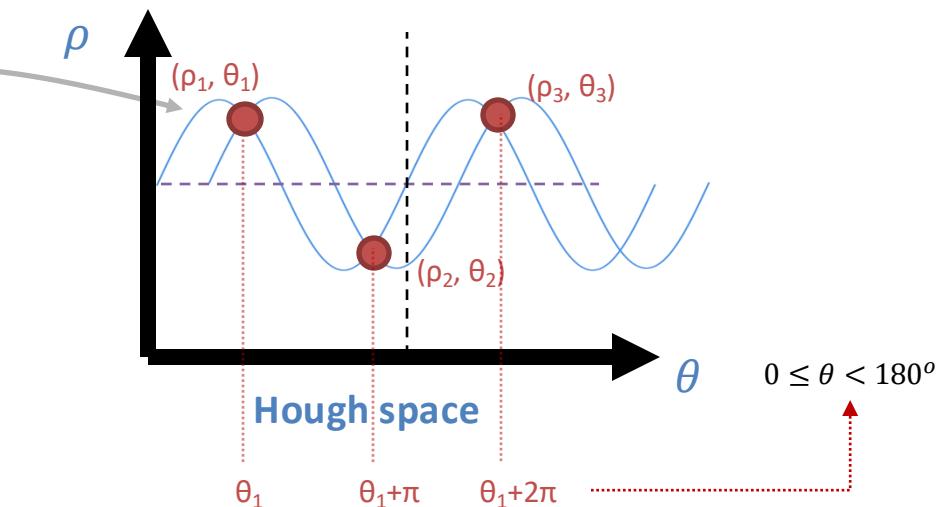
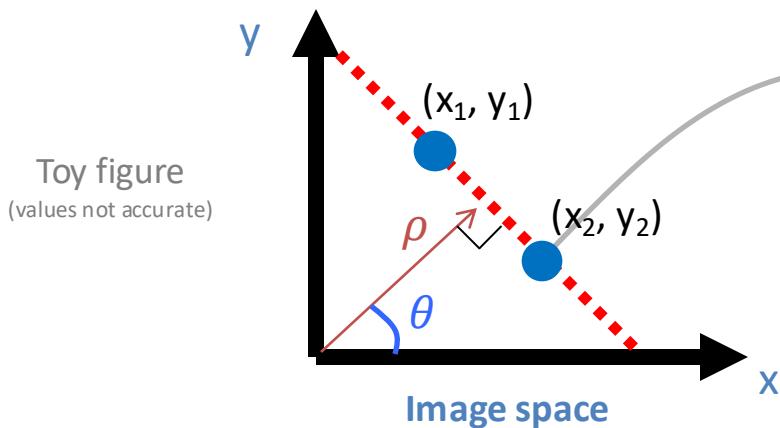
$$\begin{aligned}\rho &= x_1 \cos \theta + y_1 \sin \theta \\ &= \alpha_1 \sin(\theta + \beta_1)\end{aligned}$$



Map Point (x_2, y_2) → Sinusoid in Hough Space

$$x_2 \cos \theta + y_2 \sin \theta = \rho$$

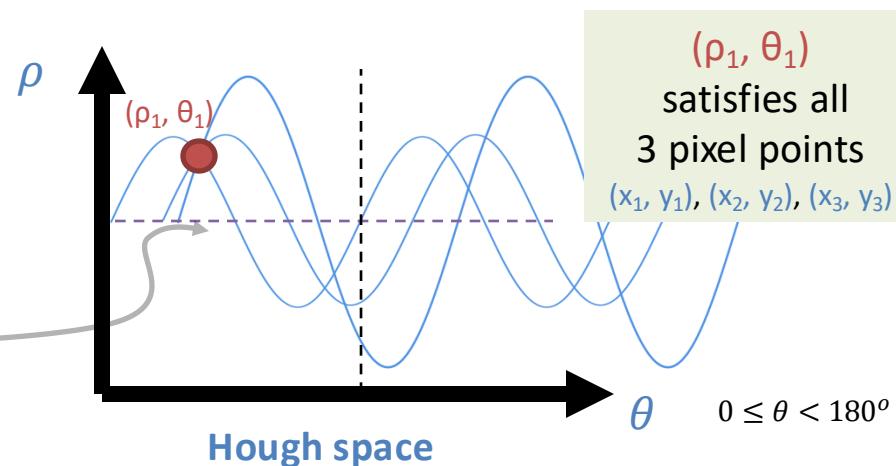
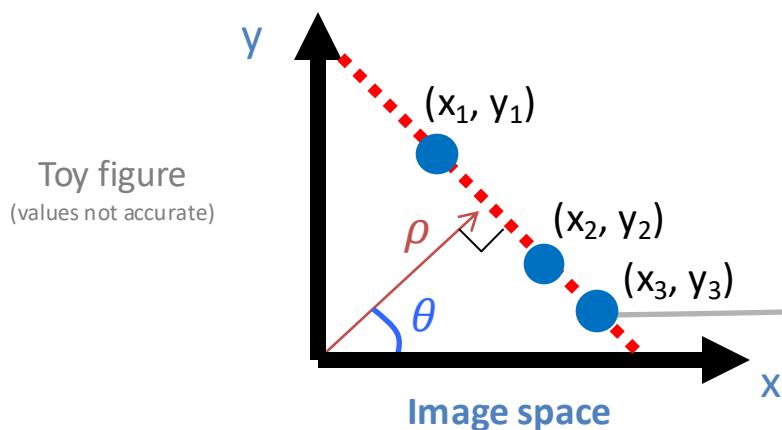
$$\begin{aligned}\rho &= x_2 \cos \theta + y_2 \sin \theta \\ &= \alpha_2 \sin(\theta + \beta_2)\end{aligned}$$



Map Point (x_3, y_3) → Sinusoid in Hough Space

$$x_3 \cos \theta + y_3 \sin \theta = \rho$$

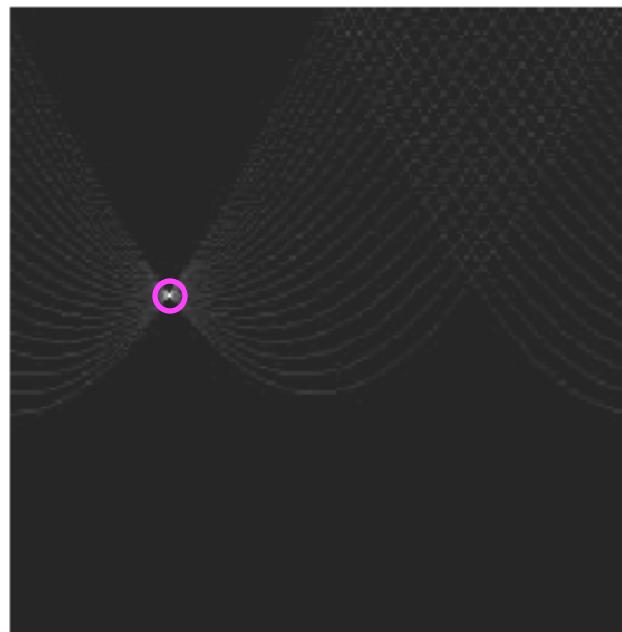
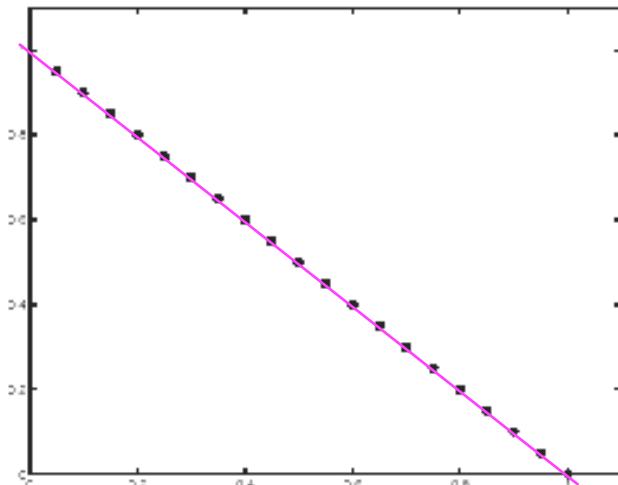
$$\begin{aligned}\rho &= x_3 \cos \theta + y_3 \sin \theta \\ &= \alpha_3 \sin(\theta + \beta_3)\end{aligned}$$



Hough Transform – Example

Clean data points → ‘Clear’ peak in Hough space

Data Points

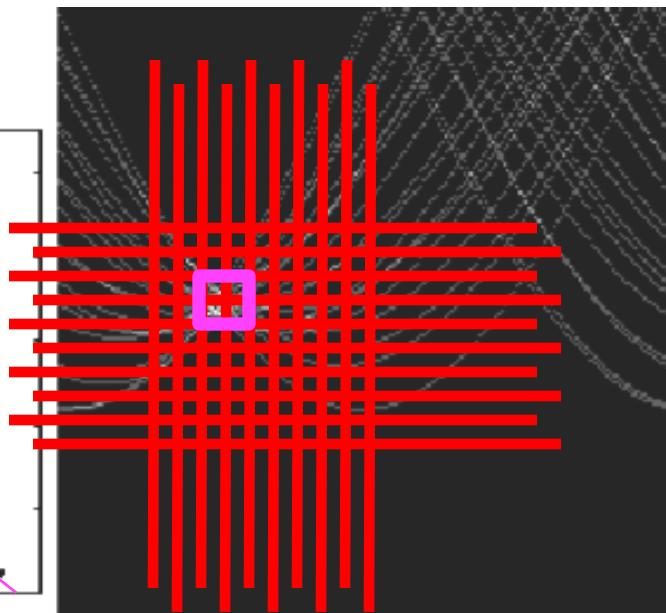
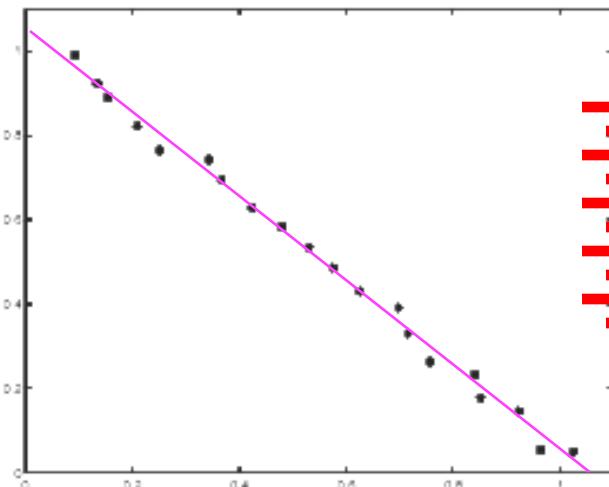


Votes in Hough Space
(regular grid)

Hough Transform – Example

Noisy data points → ‘Fuzzy’ peak in Hough space → Adjust grid size (sparser)

Data Points

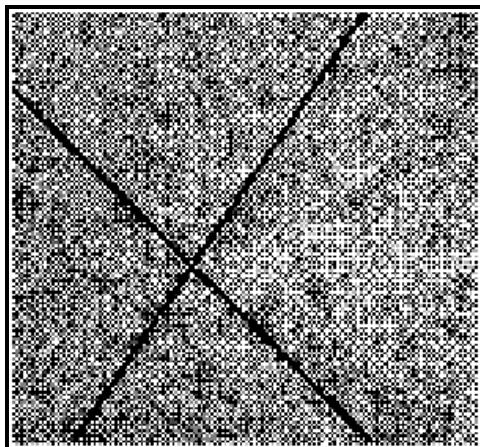


Votes in Hough Space
(regular grid)

Cell size – Trade Offs:

- **Too big:** Diff. lines merged
- **Too small:** Lines missed due to noise

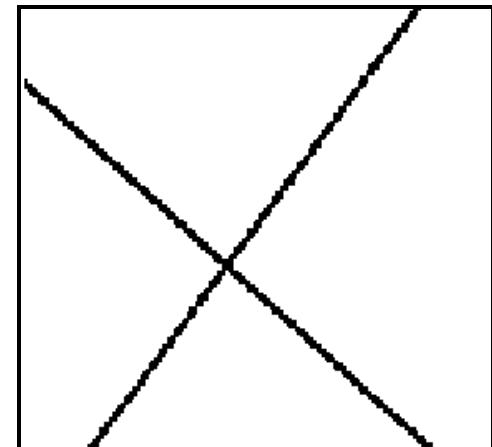
Hough Transform – Example



Image



Canny Edge
Detection



Line Detection
w Hough Transform

(Step 1/3) Image → Canny Edges



Image

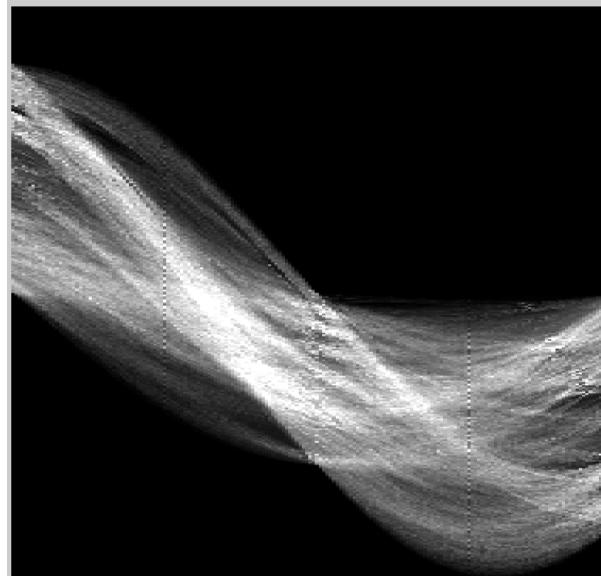


Canny Edge
Detection

(Step 2/3) Canny Edges → Hough Votes



Canny Edge
Detection

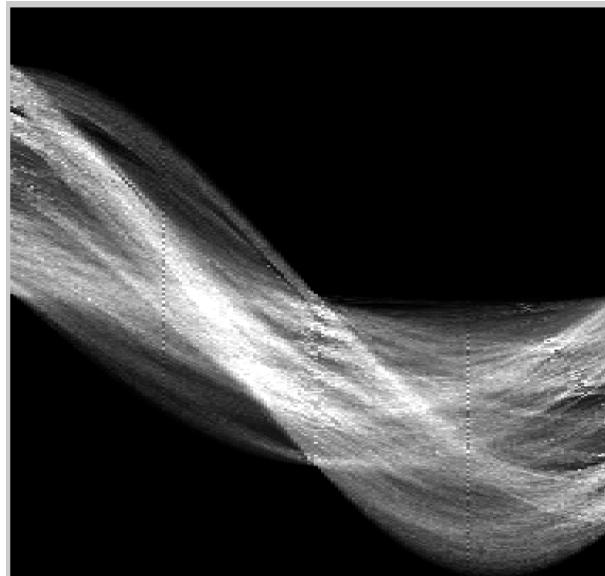


Accumulator
for Hough space

(Step 3/3) Hough Votes → Lines



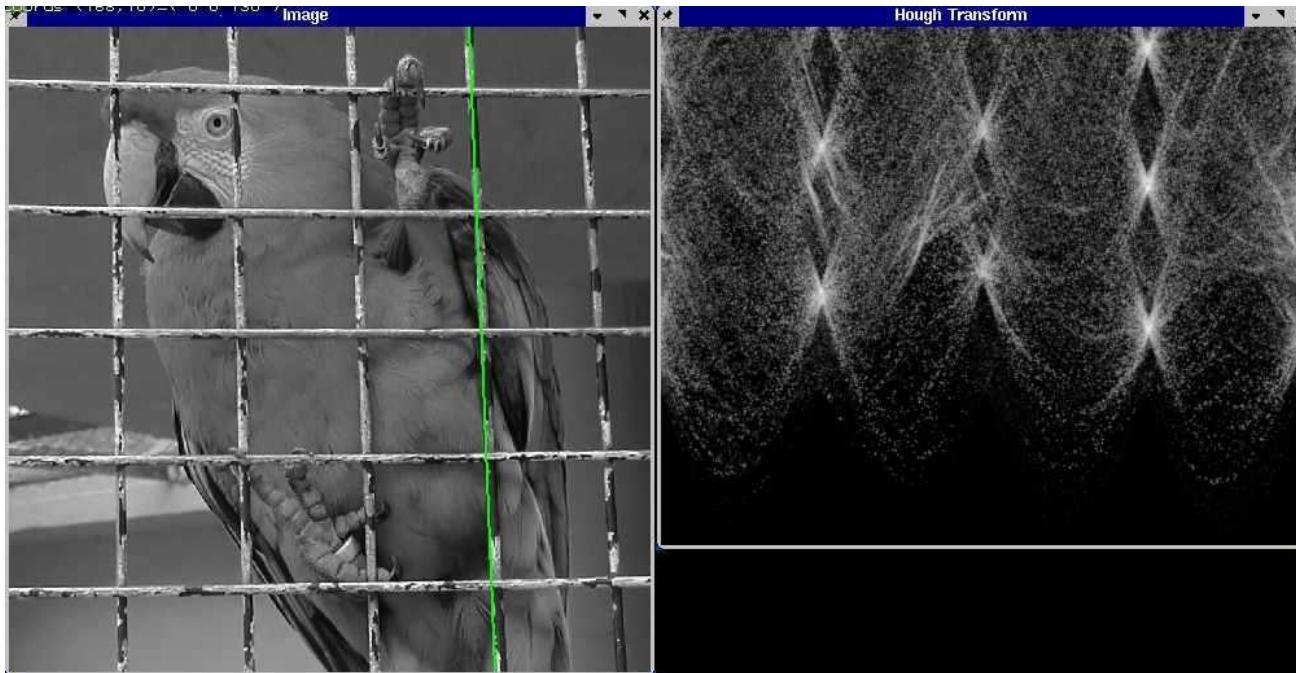
Line Detections



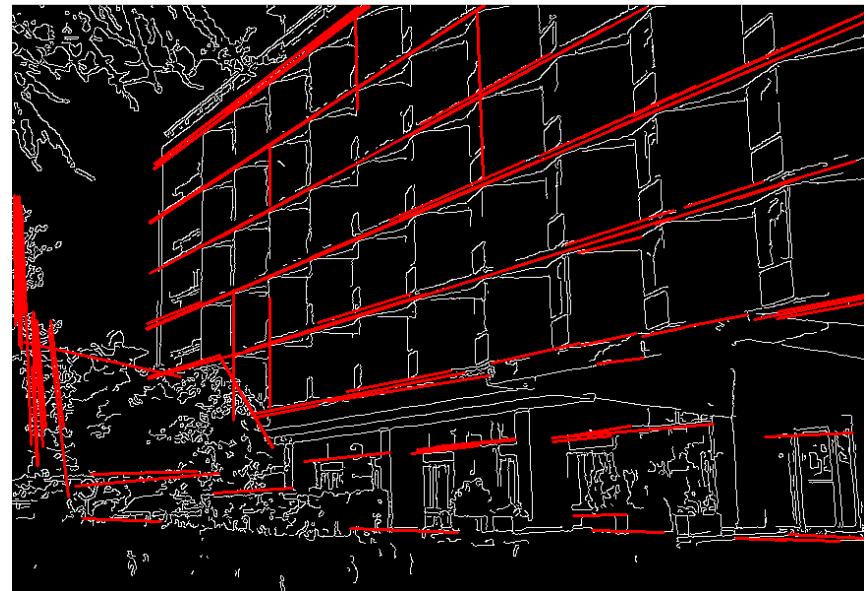
Accumulator
for Hough space

- Find [peaks](#) in Hough space
- Convert these to [lines](#) in pixel space

Hough Transform – Example



Hough Transform – Example



Probabilistic Hough Transform
(OpenCV)

Hough Transform – Pros & Cons

Pros 😊

- All points processed independently → Can cope with occlusion
- Some noise robustness (outliers unlikely to contribute consistently to single bin)
- Can detect multiple instances of a model in a single pass

Cons 😞

- Search-time complexity increases exponentially with # of model parameters
- Non-target shapes can produce spurious peaks in parameter space
- Quantization – Difficult to pick a ‘good’ grid size

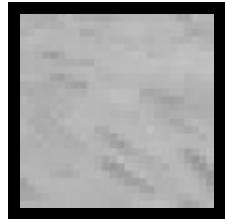
Hough Transform – Extension for Circles



Outline

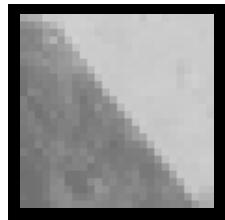
- Edge Detection
 - Derivatives of Image, Derivatives of Gaussian
 - Canny Edge Detector
- Line Fitting
 - Least Squares
 - RANSAC
 - Hough Transform
- Corners
 - Harris Corners
 - SIFT
 - Applications

Interest Points



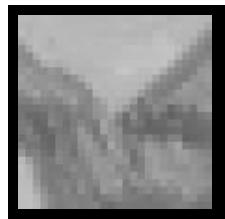
0D structure: ‘Flat’ areas

→ not useful for matching ☹



1D structure: Lines / Edges

→ edge, can be localized in 1D
subject to the aperture problem ☹



2D structure: Corners / Interest Points

→ can be localized in 2D
good for matching ☺

Interest Points have **2D** structure

Why Extract Features?

- Motivation → Panorama Stitching
 - We have 2 images
 - How to combine them?

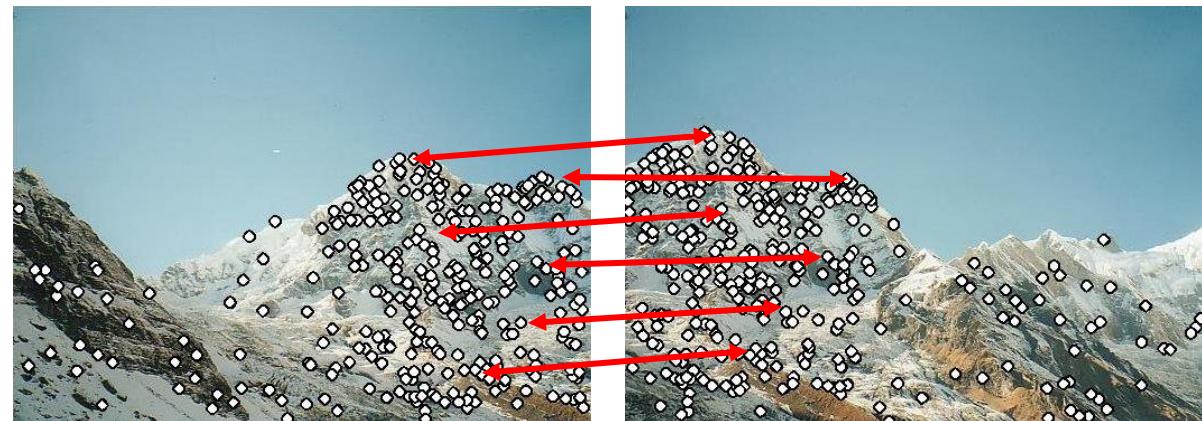


Why Extract Features?

- Motivation → Panorama Stitching
 - We have 2 images
 - How to combine them?

Step 1 → Extract Features

Step 2 → Match Features



Why Extract Features?

- Motivation → Panorama Stitching

- We have 2 images
- How to combine them?

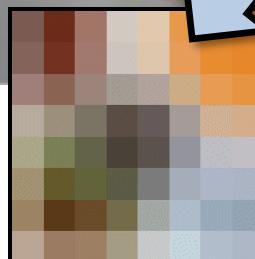
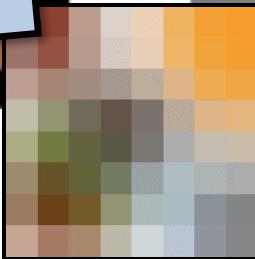
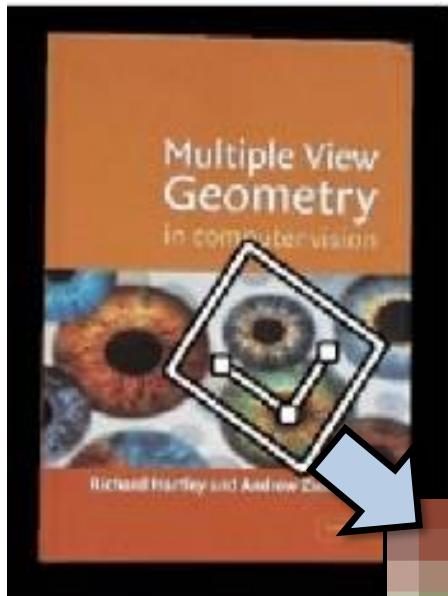
Step 1 → Extract Features

Step 2 → Match Features

Step 3 → Align Images



Feature Matching



Outline

- Edge Detection
 - Derivatives of Image, Derivatives of Gaussian
 - Canny Edge Detector
- Line Fitting
 - Least Squares
 - RANSAC
 - Hough Transform
- Corners
 - Harris Corners
 - SIFT
 - Applications

Harris Corner Detector – Mathematics

Change of intensity
for shift $[u, v]$

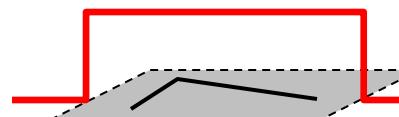
$$E(u, v) = \sum_{x,y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

Diagram illustrating the components of the Harris corner detector equation:

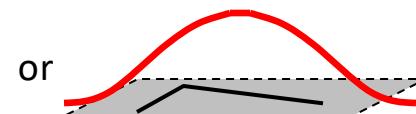
- Window function**: Represented by a blue rounded rectangle.
- Shifted intensity**: Represented by an orange rounded rectangle.
- Intensity**: Represented by a purple rounded rectangle.

Arrows point from each component label to its corresponding term in the equation.

Window function $w(x, y) =$

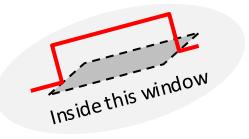


1 → in window
0 → outside



Gaussian

Harris Corner Detector – Mathematics



$$\begin{aligned}
 E(u, v) &= \sum_{x,y} [I(x-u, y-v) - I(x, y)]^2 \\
 &= \sum_{x,y} [I(x, y) - I_x u - I_y v + \varepsilon - I(x, y)]^2 \\
 &\approx \sum_{x,y} [-I_x u - I_y v]^2 \\
 &= \sum_{x,y} I_x^2 u^2 + I_x u I_y v + I_y v I_x u + I_y^2 v^2 \\
 &= \sum_{x,y} [u \quad v] \begin{bmatrix} I_x^2 & I_x I_y \\ I_y I_x & I_y^2 \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}
 \end{aligned}$$

$$E(u, v) \cong [u, v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

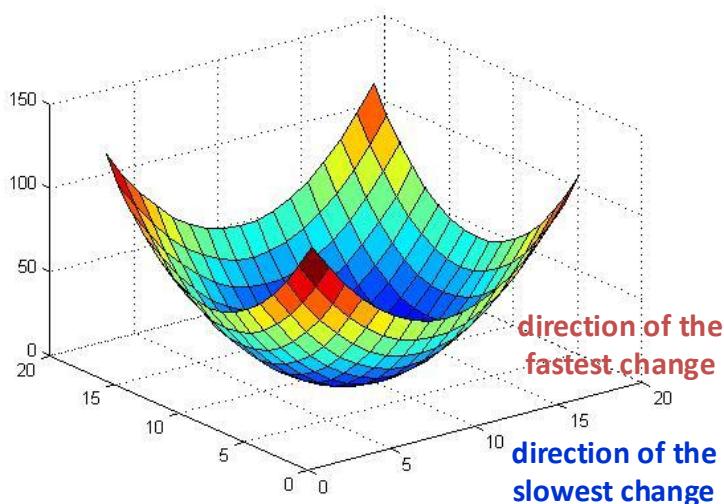
$$M = \sum_{x,y} \begin{bmatrix} I_x^2 & I_x I_y \\ I_y I_x & I_y^2 \end{bmatrix}$$

Approximate Hessian
Encodes image ‘curvature’
Product of 1st order derivatives
instead of 2nd order derivatives

M depends on
img properties

Quadratic Form & Eigenvalues

$$E(u, v) \cong [u, v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

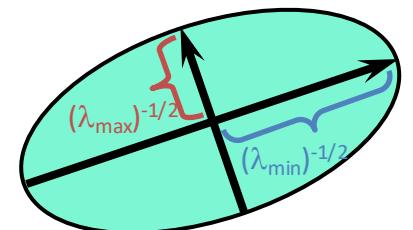


Measure of corner response:

$$R = \det M - k(\text{trace } M)^2$$

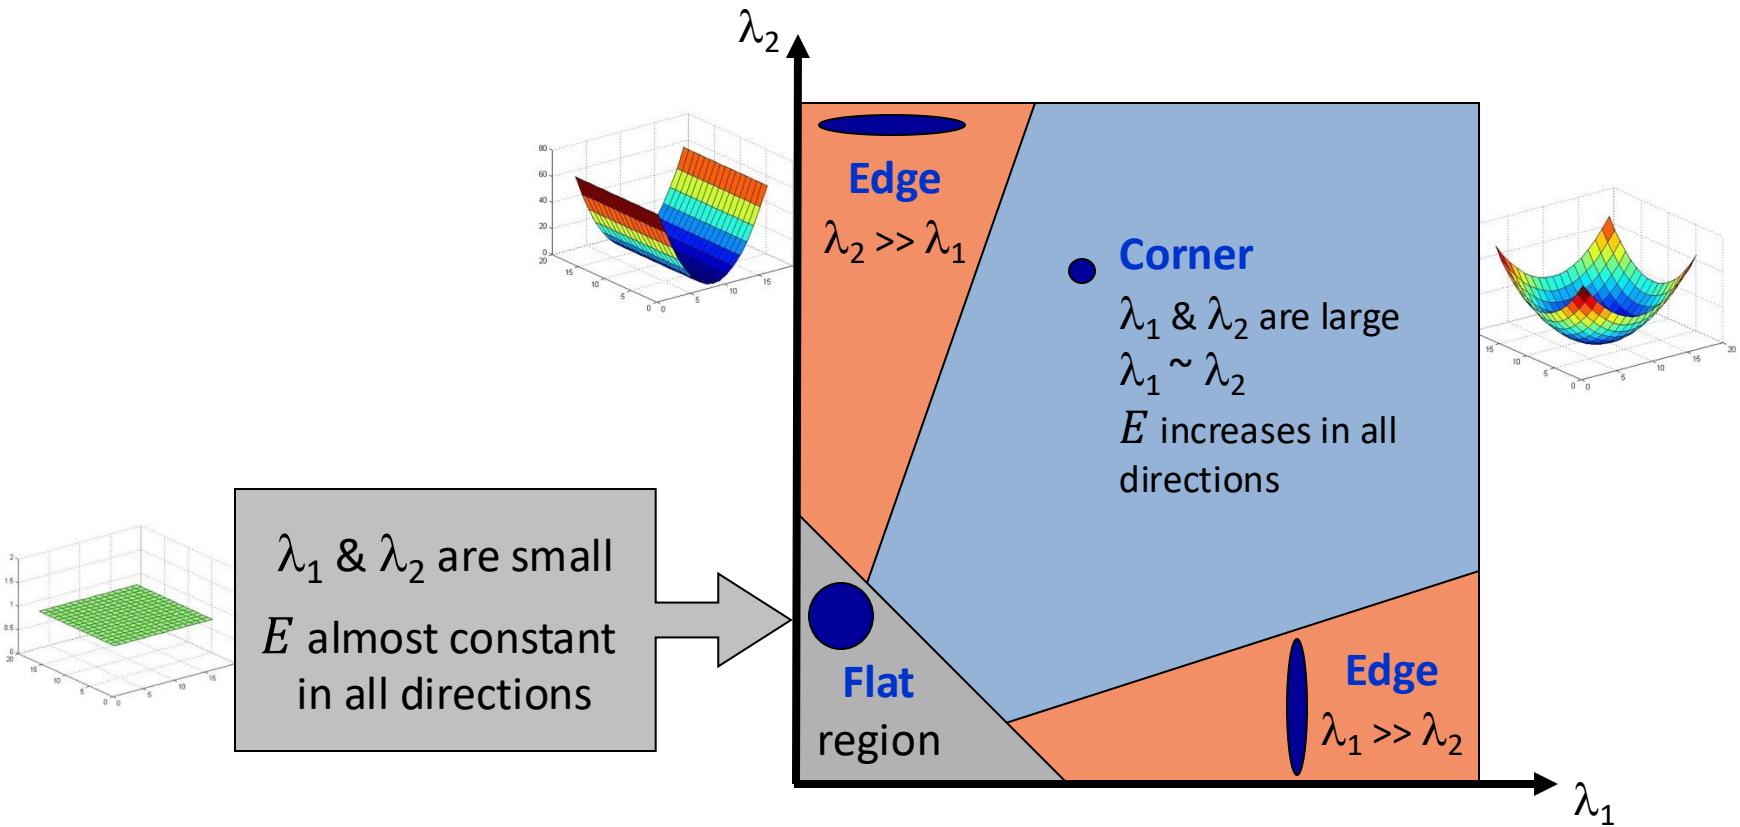
$$\det M = \lambda_1 \lambda_2$$

$$\text{trace } M = \lambda_1 + \lambda_2$$



($k \rightarrow$ empirical constant)
 $(k = 0.04 - 0.06)$

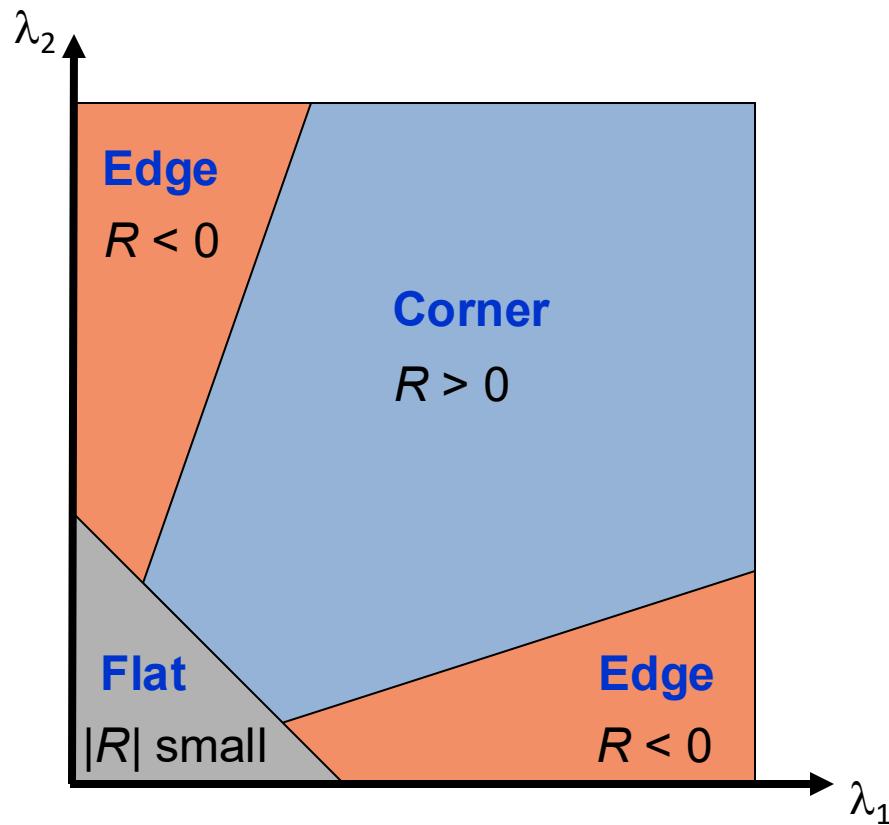
Eigenvalues of M – Classify Image Points



Harris Detector

$$R = \det M - k(\text{trace} M)^2$$

- R depends only on eigenvalues of M
- R is large for a corner
- R is negative with large magnitude for an edge
- $|R|$ is small for a flat region



Harris Detector – Summary

1. Compute x and y derivatives of image

$$I_x = G_\sigma^x * I \quad I_y = G_\sigma^y * I$$

2. Compute products of derivatives at every pixel

$$I_{x2} = I_x \cdot I_x \quad I_{y2} = I_y \cdot I_y \quad I_{xy} = I_x \cdot I_y$$

3. Compute the sums of the products of derivatives at each pixel

$$S_{x2} = G_{\sigma'} * I_{x2} \quad S_{y2} = G_{\sigma'} * I_{y2} \quad S_{xy} = G_{\sigma'} * I_{xy}$$

4. Define at each pixel (x, y) the matrix

$$H(x, y) = \begin{bmatrix} S_{x2}(x, y) & S_{xy}(x, y) \\ S_{xy}(x, y) & S_{y2}(x, y) \end{bmatrix}$$

5. Compute the response of the detector at each pixel

$$R = \text{Det}(H) - k(\text{Trace}(H))^2$$

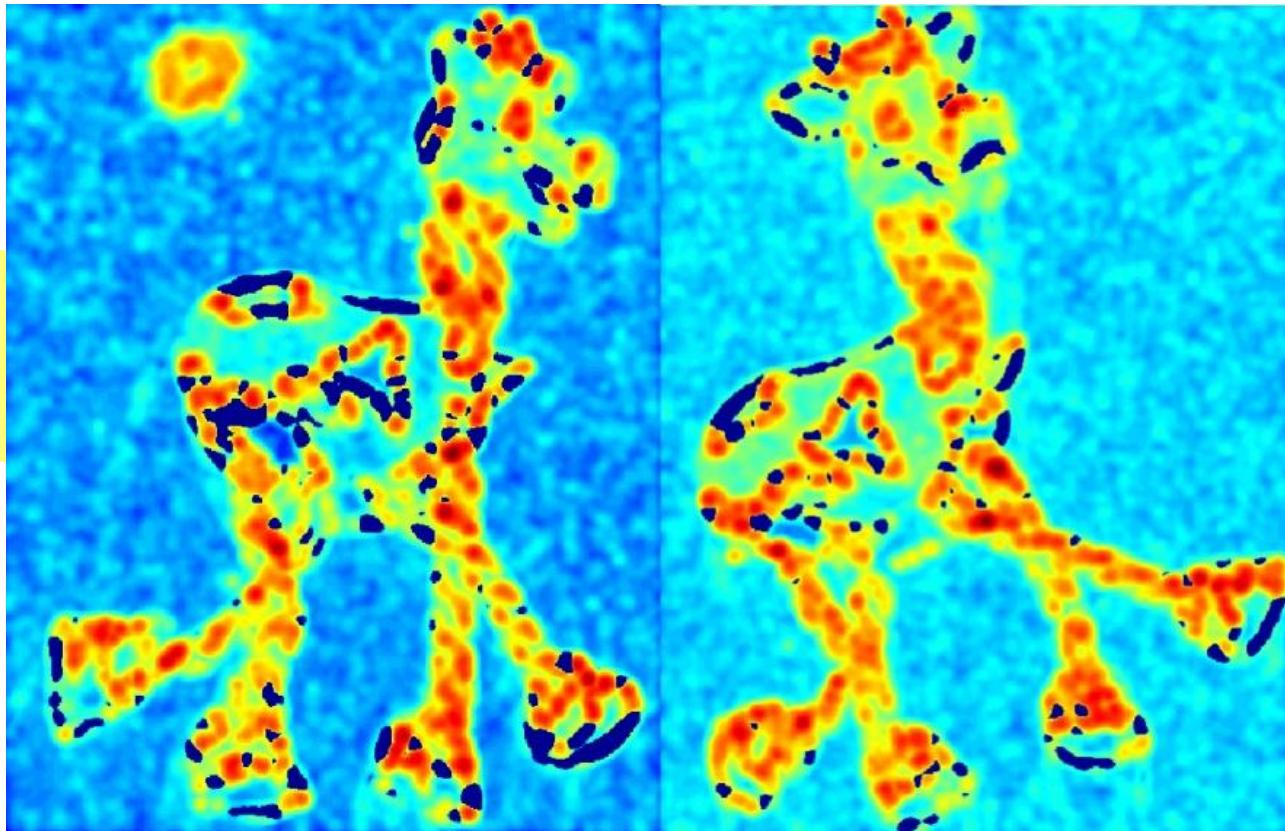
6. Threshold on value of R. Compute nonmax suppression

Harris Detector – Workflow



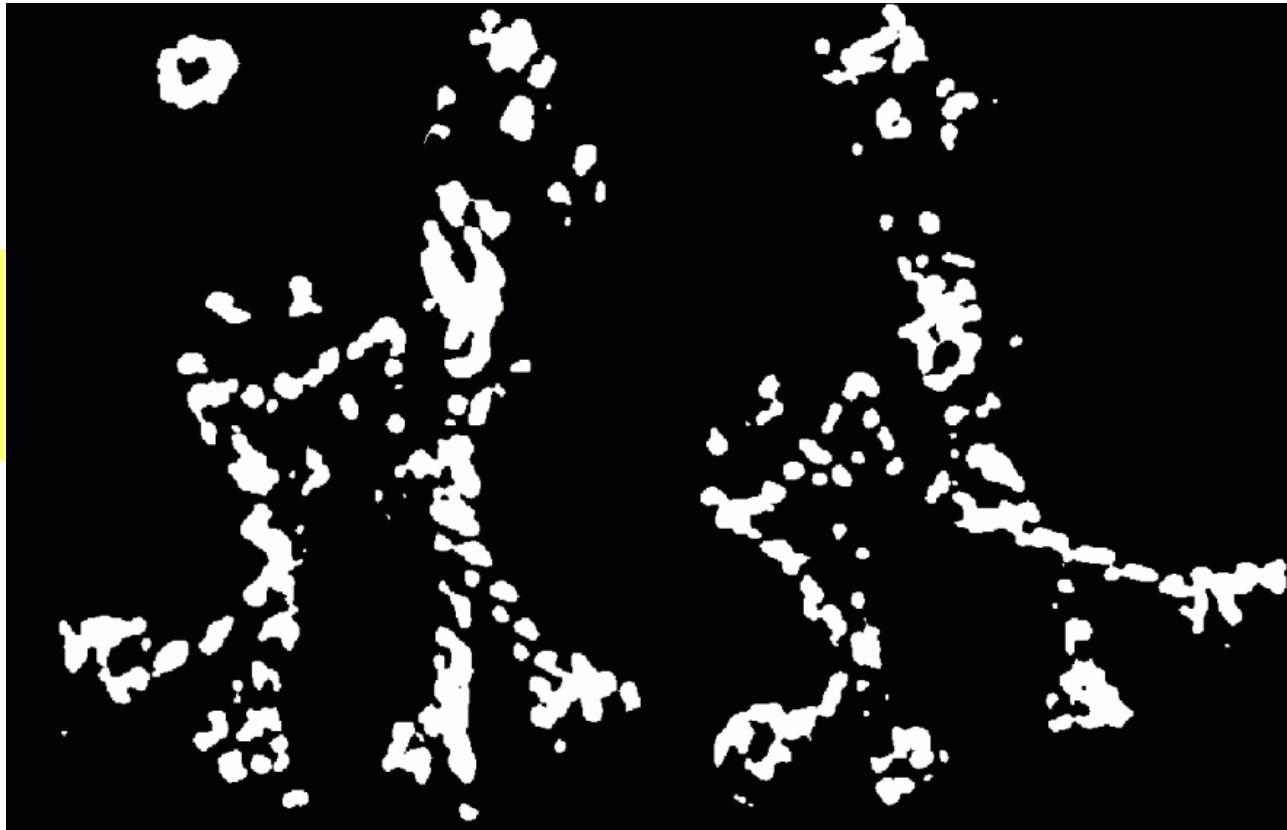
Harris Detector – Workflow

Compute
corner response
 R



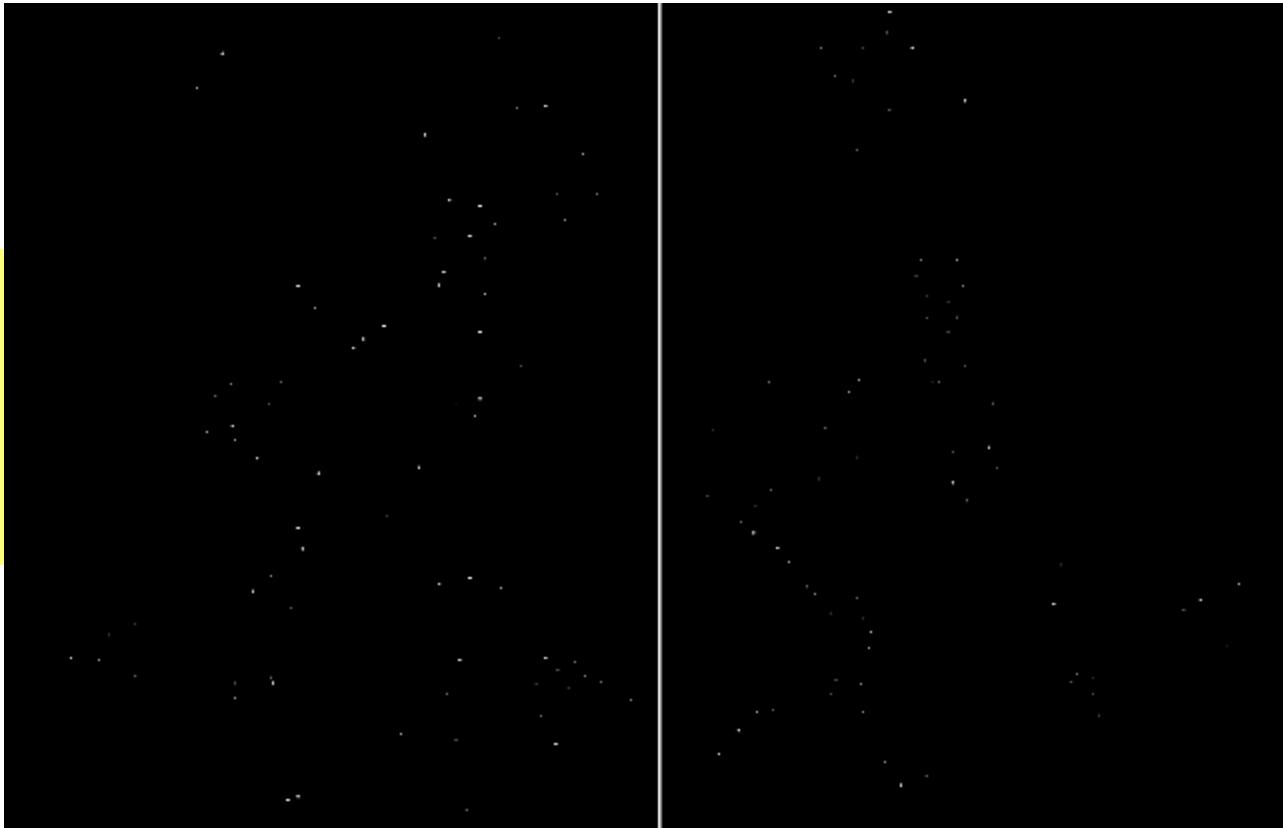
Harris Detector – Workflow

Points w. large
corner response
 $R >$ threshold



Harris Detector – Workflow

Take only points
at local maxima
of R
(non-maxima
suppression)



Harris Detector – Workflow



Outline

- Edge Detection
 - Derivatives of Image, Derivatives of Gaussian
 - Canny Edge Detector
- Line Fitting
 - Least Squares
 - RANSAC
 - Hough Transform
- Corners
 - Harris Corners
 - SIFT
 - Applications

Interest Points & Image Scale

Can we detect the
same interest points
despite changes in
image scale?



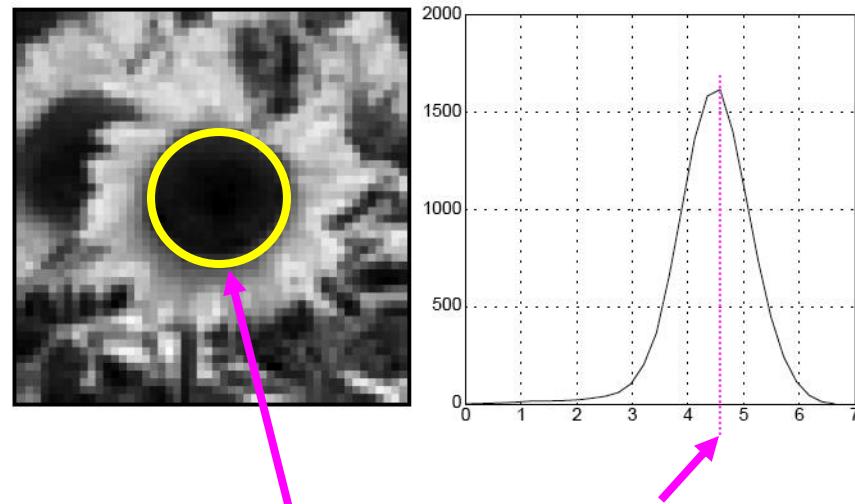
Blob Detection in 2D

Characteristic Scale → The scale that produces peak of Laplacian response

Nugget:

- Convolve image with Gaussians of varying scale
- Detect max responses across scale space

A narrower or wider Gaussian → would have a weaker response for this flower



Example

$\frac{3}{4}$ the
Original Size



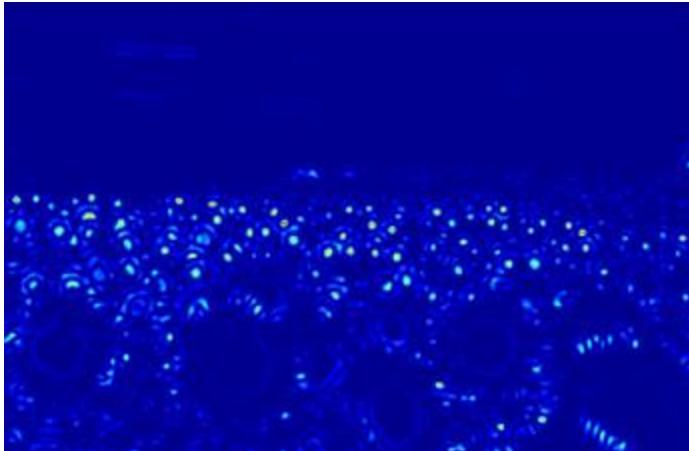
Original
Size



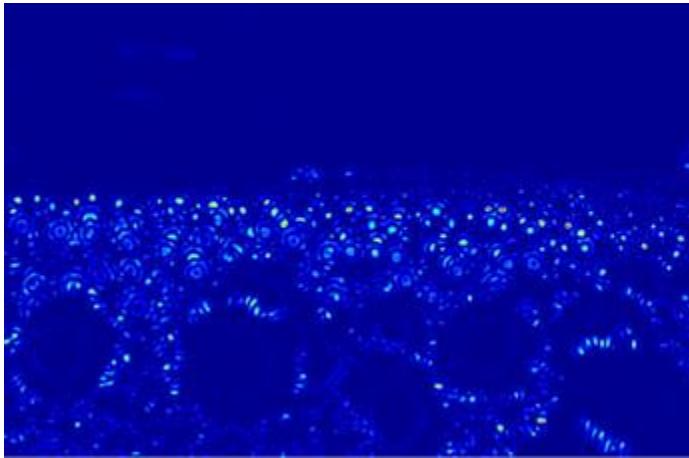
Example

Image at $\frac{3}{4}$ the
original size →

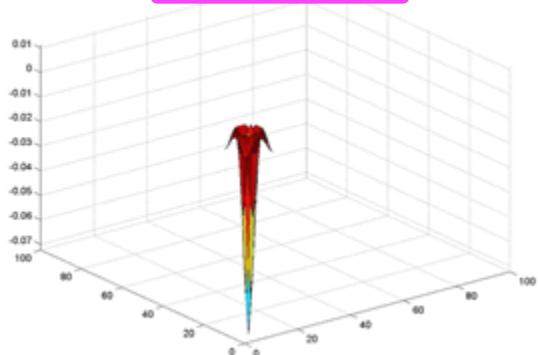
$\frac{3}{4}$ the
Original Size



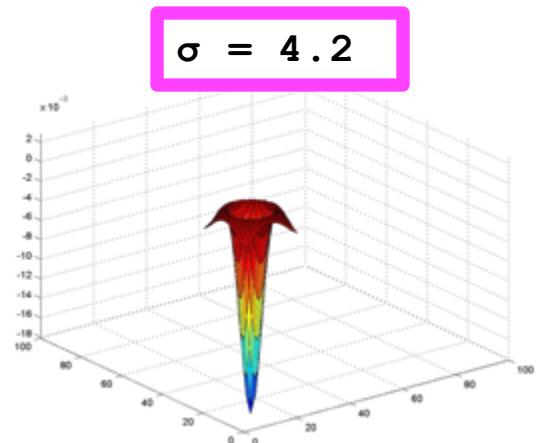
Original
Size



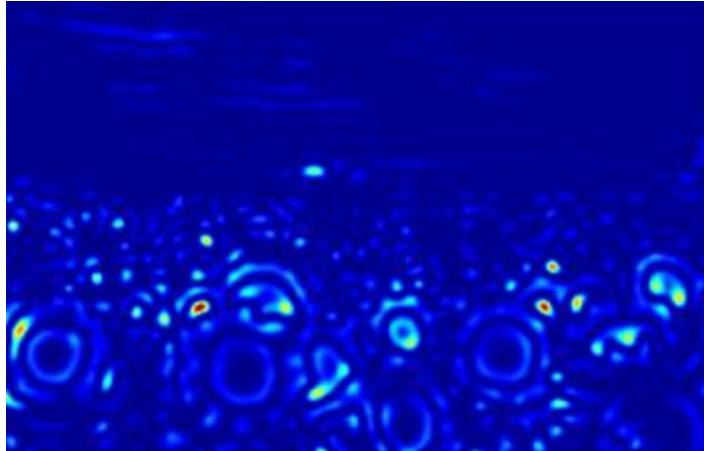
$$\sigma = 2.1$$



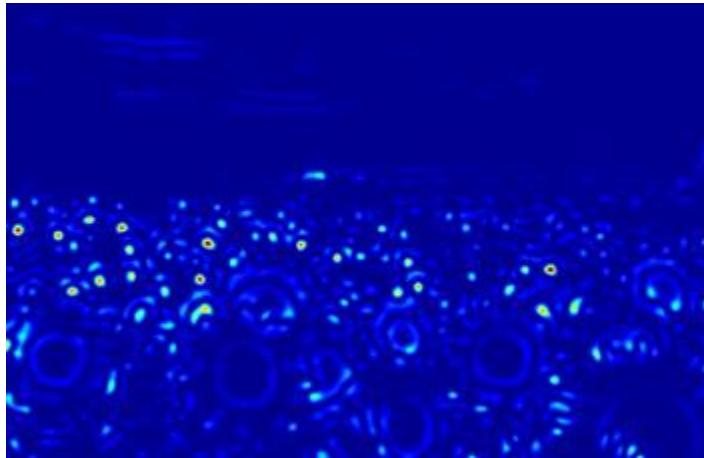
Example



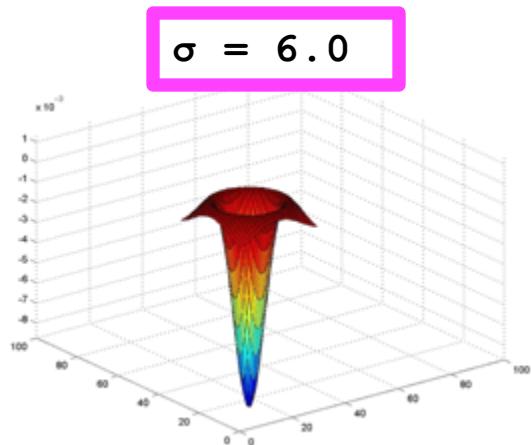
$\frac{3}{4}$ the
Original Size



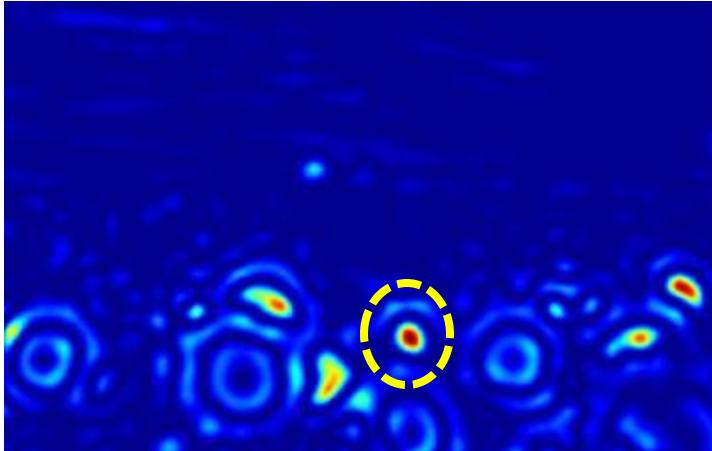
Original
Size



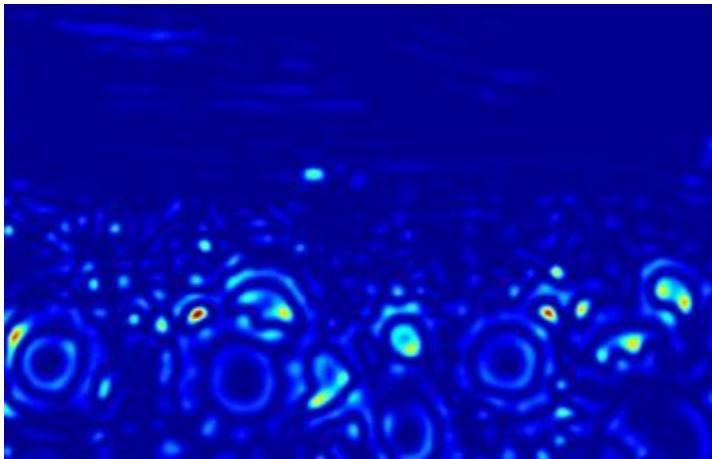
Example



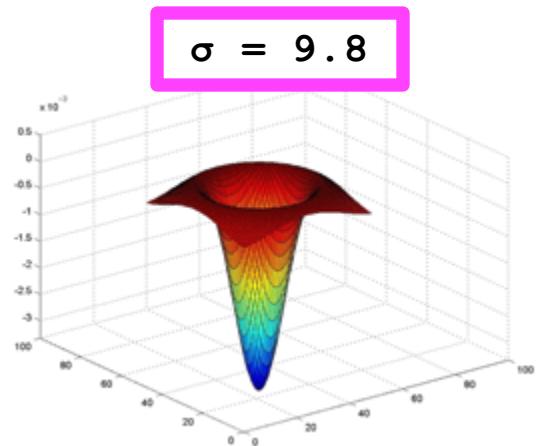
$\frac{3}{4}$ the
Original Size



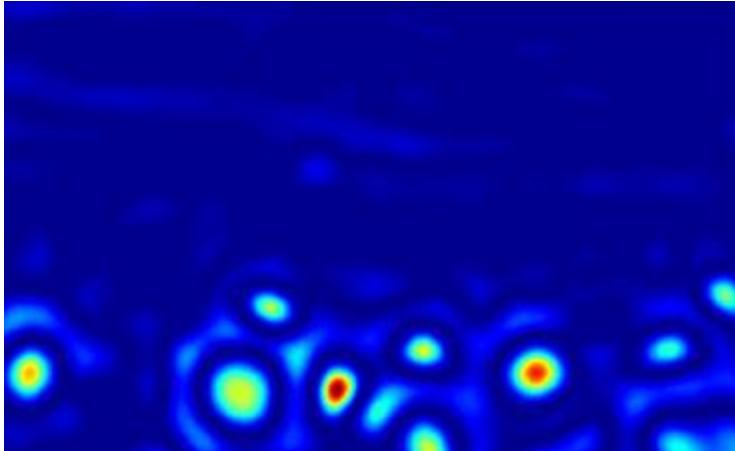
Original
Size



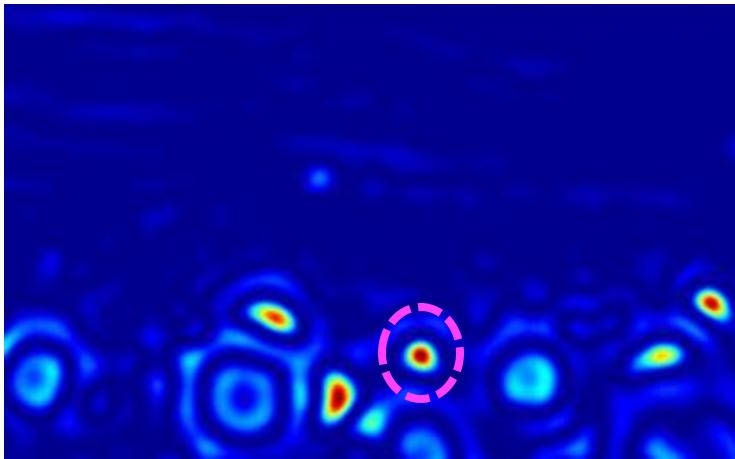
Example



$\frac{3}{4}$ the
Original Size

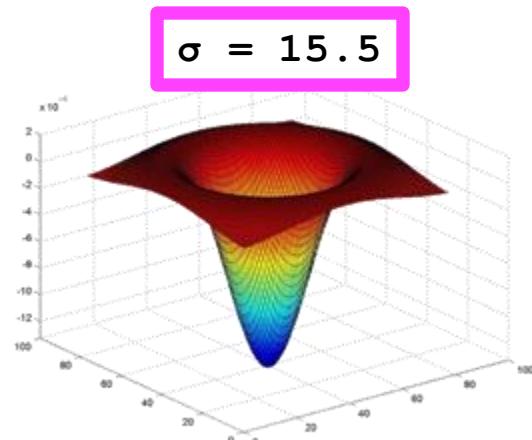


Original
Size

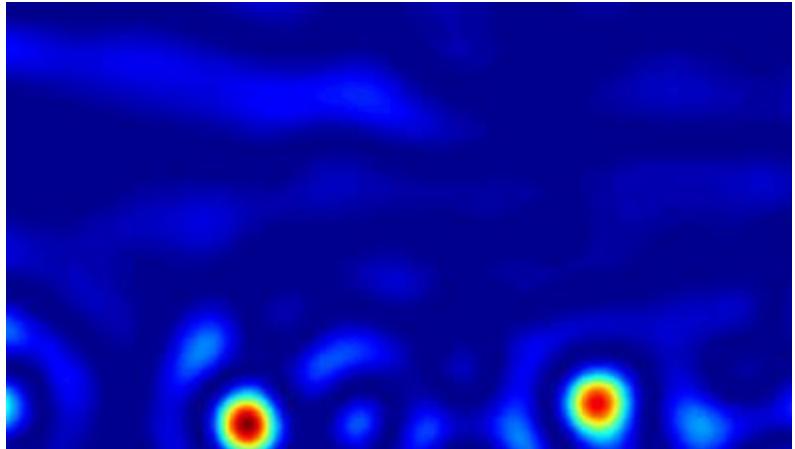


Example

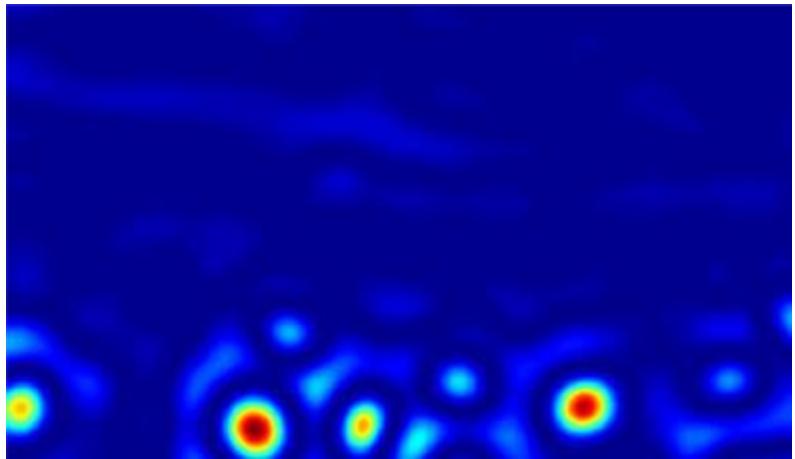
A LoG of **increasing sigma** produces stronger activations for **bigger-looking sunflowers** (lying **closer to the camera**)



% the
Original Size



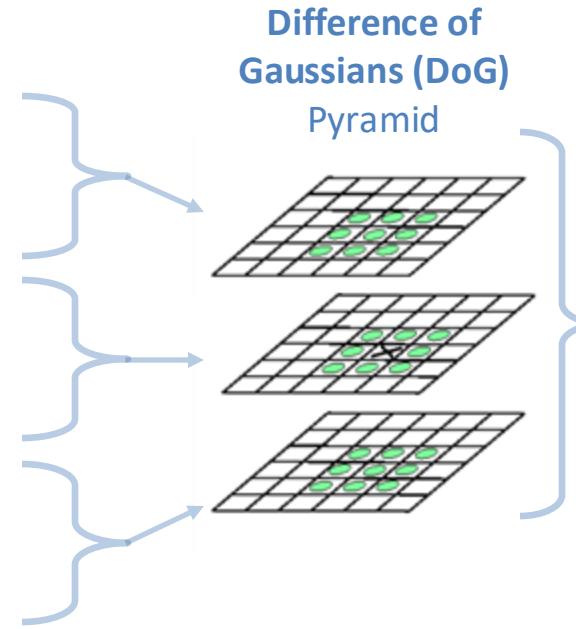
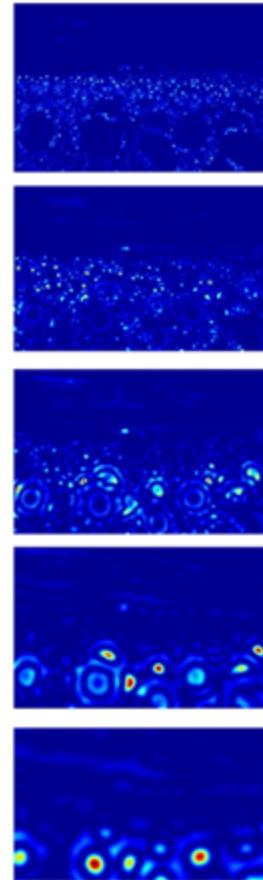
Original
Size



SIFT Keypoints



Gaussian Scale Space s

$$L_{xx}(\sigma) + L_{yy}(\sigma)$$


Convolution with Gaussian is rotation invariant

Keypoints:
local **extrema** in
 $3 \times 3 \times 3$ grid i.e. in both:

- Spatial space (x, y)
- Scale space (s)

Keypoints are !
Scale Invariant

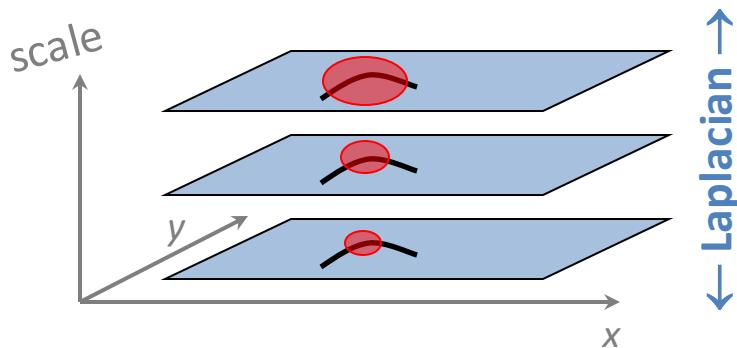
Keypoints are !
Rot. Invariant

Two Popular Scale Invariant Detectors

Harris-Laplacian^[1]

Find local maximum of:

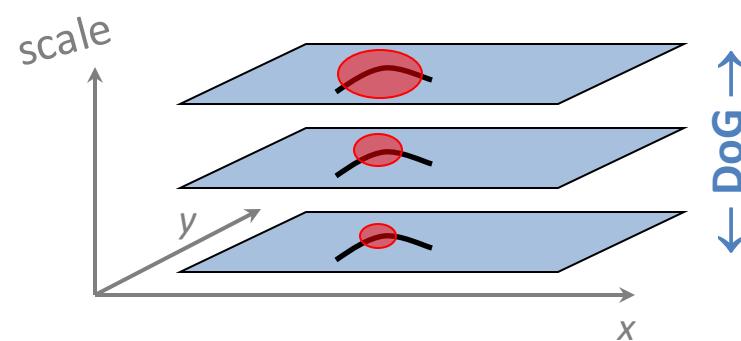
- Harris corner detector @ Image space
- Laplacian @ scale space



SIFT (Lowe)^[2]

Find local maximum of:

- Difference of Gaussians (DoG) @ both
 - Image space and
 - Scale space



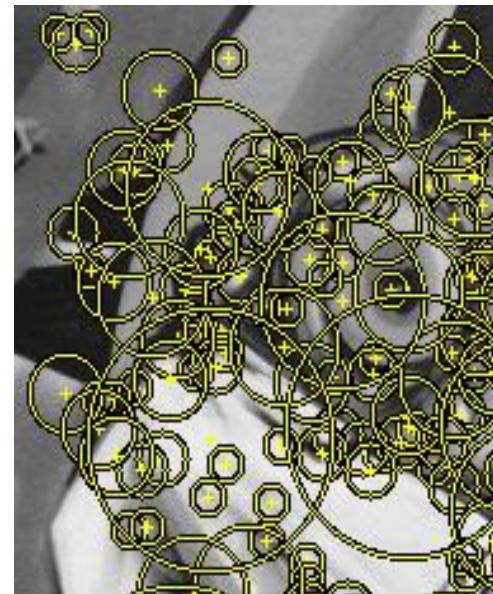
[1] K.Mikolajczyk, C.Schmid. 'Indexing Based on Scale Invariant Interest Points'. ICCV 2001

[2] D.Lowe. 'Distinctive Image Features from Scale-Invariant Keypoints'. IJCV 2004

Comparison



Harris

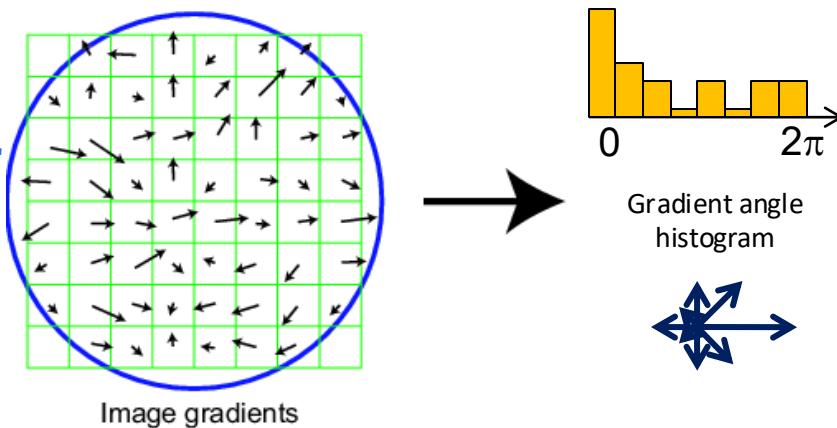


LoG

SIFT – Descriptor

Histogram computed in
Gaussian-weighted
neighborhood

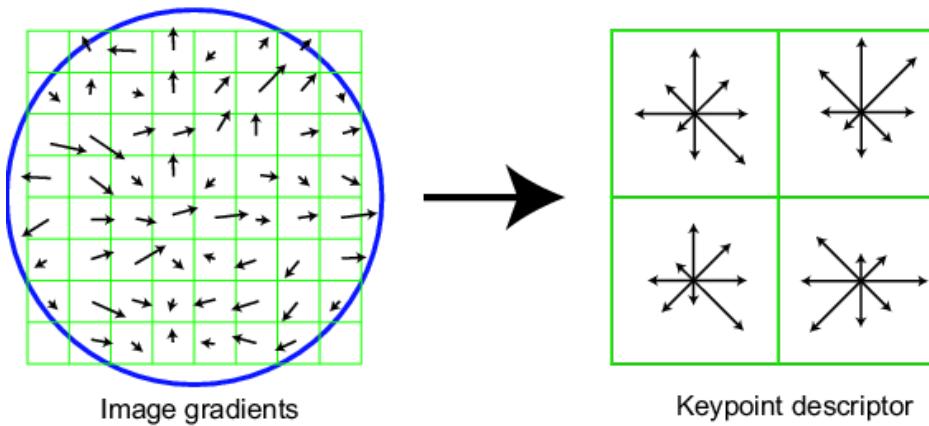
Contributes to Rot.
Invariant for keypoints



Basic idea:

- Take **16x16 window** around detected keypoint (8x8 example shown here)
- For each window **pixel** – Compute **edge orientation** (angle of the gradient - 90°)
- Throw out weak edges (threshold gradient magnitude)
- Create **histogram** of surviving edge orientations

SIFT – Descriptor



Full version of SIFT:

- Divide **16x16** window → **4x4** cell grid (**2x2** example shown here)
- For each cell, compute **orientation histogram** (8 orientation bins)
- 16 cells (4x4 grid) * 8 orientations (per cell) → **128 dimensional descriptor**

SIFT – Rotation Invariant (Robust)

For each patch:

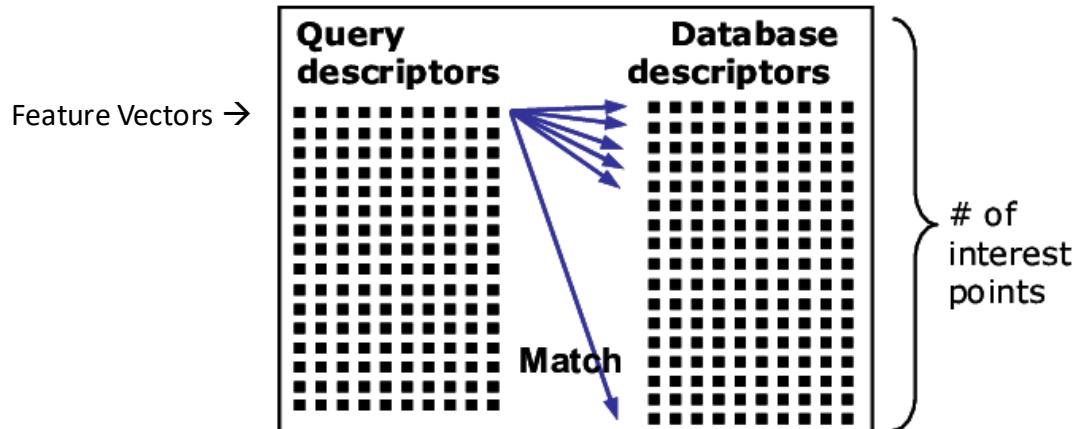
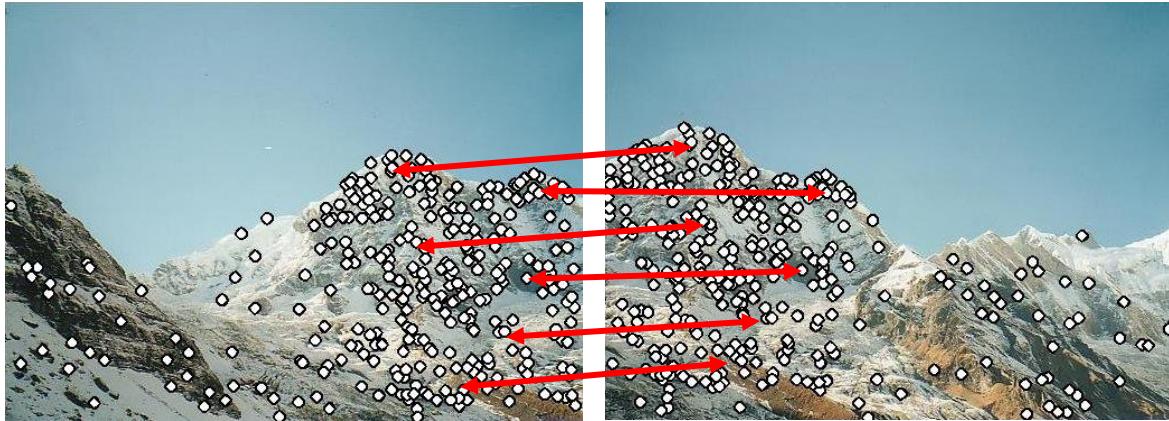
- Compute **histogram** of its gradient directions/orientations
- Find **dominant orientation**
- **Descriptor** represented relative to this **orientation**



**Descriptor is
Rotation Invariant!**

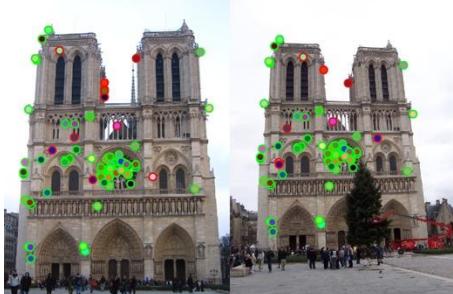


SIFT – Features Matching – Brute-Force



Keypoints –VS– Features

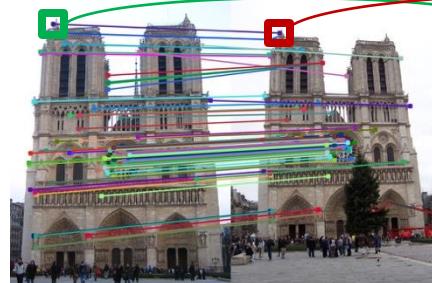
Keypoint Detector (Finding ‘interesting’ points)



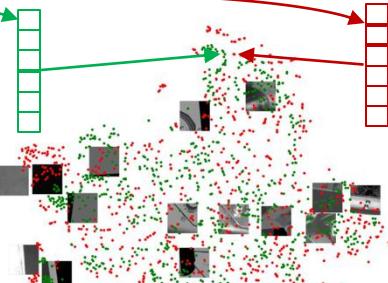
[S. Shah, [Local Feature Matching](#)]

Examples (namedrop): Harris corners, Shi-Tomasi corner detector, SUSAN corner detector, FAST feature detector, Laplacian of Gaussian, Difference of Gaussian, Superpoint, ...

Feature Descriptor (Describing area around keypoint)



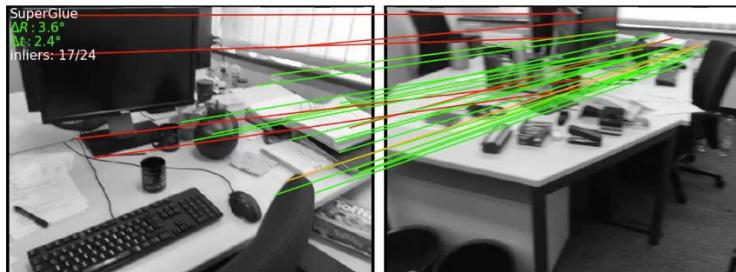
[S. Shah, [Local Feature Matching](#)]



[Ghimire et al., Applied Sciences, 2021]

Examples (namedrop): SIFT – Scale Invariant Feature Transform, SURF – Speeded-up Robust Features, BRIEF – Binary Robust Independent Elementary Features, ORB – Oriented FAST Rotated BRIEF, Superpoint, ...

Goal: Feature Matching



[Sarlin et al., SuperGlue, CVPR 2020]

SIFT properties

- **Invariant** to:
 - Scale → Detect scale using DoG
 - Rotation → Dominant direction using histogram
 - **Partially invariant** to
 - Illumination changes: Gradient magnitude is thrown away, only use direction
 - Camera viewpoint (affine transform)
 - Occlusion
 - Clutter
- ← (check all icons) 

Bringing Everything Together

SIFT & RANSAC
for Panoramas
(aka Image Stitching)

Image Stitching – Overview

1. Detect keypoints – OpenCV tuple:

$$(x, y, s, \theta, \bar{d})_i$$

- *position* (x, y) in image
- *scale* s of detection
- *orientation* θ
- *feature descriptor* \bar{d}



Embeds local image area
into a *128D latent space*



Features nearby in this
space → *image areas*
that *look similar*

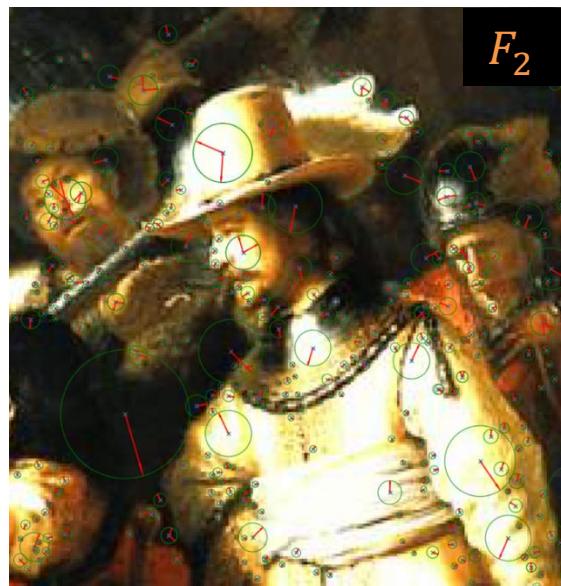


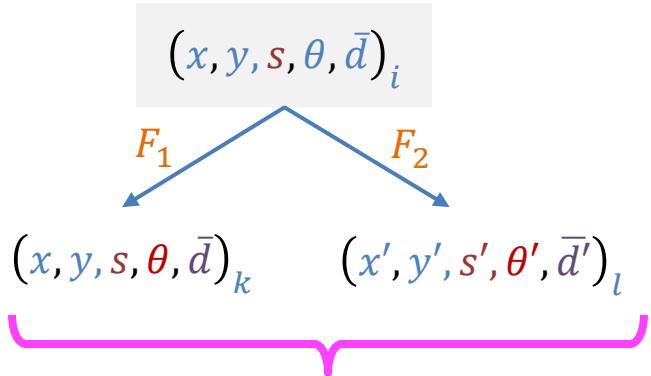
Image Stitching – Overview

1. Detect keypoints – OpenCV tuple

$$(x, y, s, \theta, \bar{d})_i$$

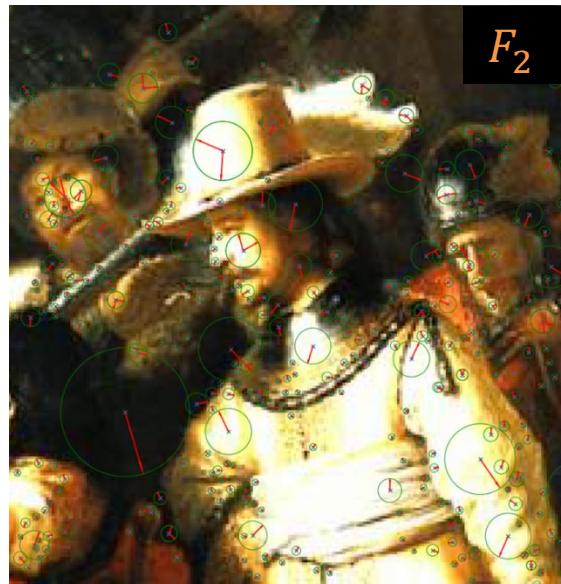
2. Compute a feature per keypt

3. Establish keypt *correspondences*



Corresp.
In case:

$$\|\bar{d}_k - \bar{d}'_l\| \leq \varepsilon$$



F_1

F_2

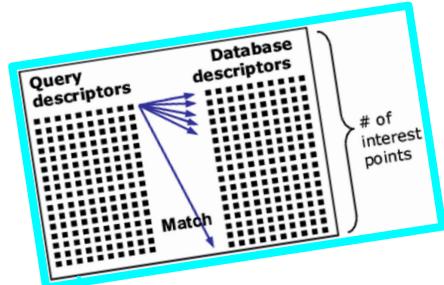
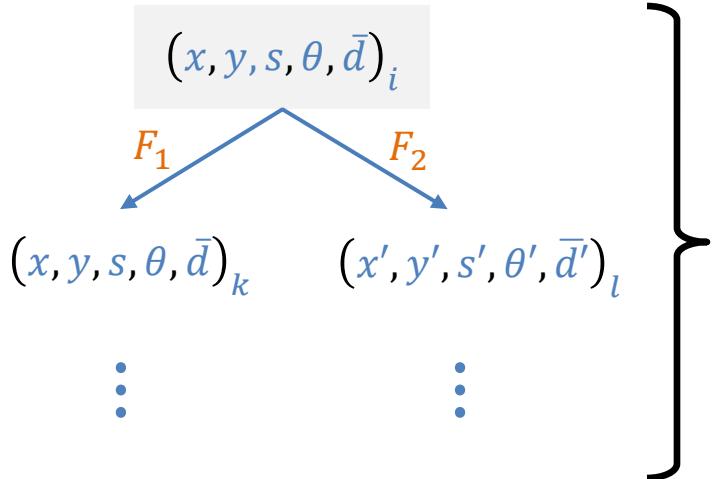
Image Stitching – Overview

1. Detect keypoints – OpenCV tuple

$$(x, y, s, \theta, \bar{d})_i$$

2. Compute a feature per keypt

3. Establish keypt *correspondences*



Compute for all pairs $\{k, l\}$:

$$D_{kl} = \|d_k - d'_l\|$$

Brute-force 'Search':

$$(x, y, x', y') = \arg \min_{k, l} D_{kl}$$

Iterate using as query:

1st point, 2nd, 3rd, ...

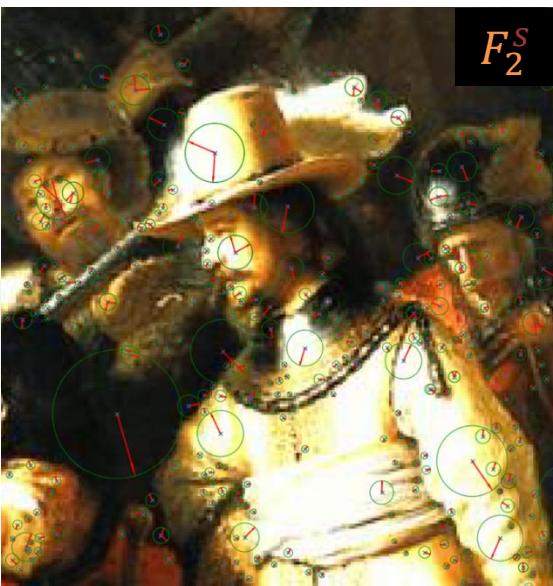


Image Stitching – Overview

1. Detect keypoints – OpenCV tuple

$$(x, y, s, \theta, d)_i$$

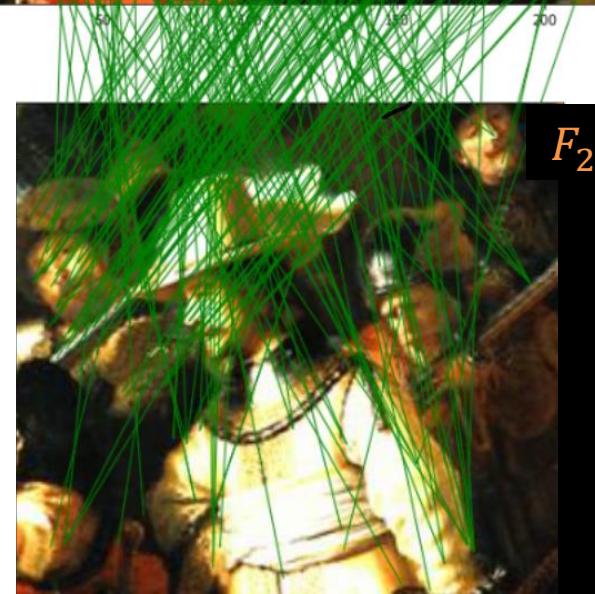
2. Compute a feature per keypt

3. Establish keypt *correspondences* → Can be *noisy*

4. Filter corresp. & find
transf. with *RANSAC*

Solution

Random subsets produce
diff. Projective Trans. T



F_1

F_2

Image Stitching – Overview

RANSAC

RANdom SAmple Consensus

Goal: Find **subset** of corrsp.
that produce a **good**
Proj. Transform T

A *candidate* corrsp. $\{(x, y), (x', y')\}$
is **true positive** if it complies with T :

$$\begin{bmatrix} x \\ y \\ 1 \end{bmatrix} \approx T \begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} \quad \text{within a small error margin}$$



Steps

- a. Randomly select 4 corrsp: $\{(x, y), (x', y')\}_{i \in \{1,2,3,4\}}$ → Compute (hypothesis) Proj. Transf. T
- b. Count number # of *inliers* for current hypothesis/model → For other corrsp. evaluate whether:
- c. Store model with max inliers (corr., # inliers, fit quality, trnsf. T)
- d. Iterate max N times
- e. For best model – use *all inliers*

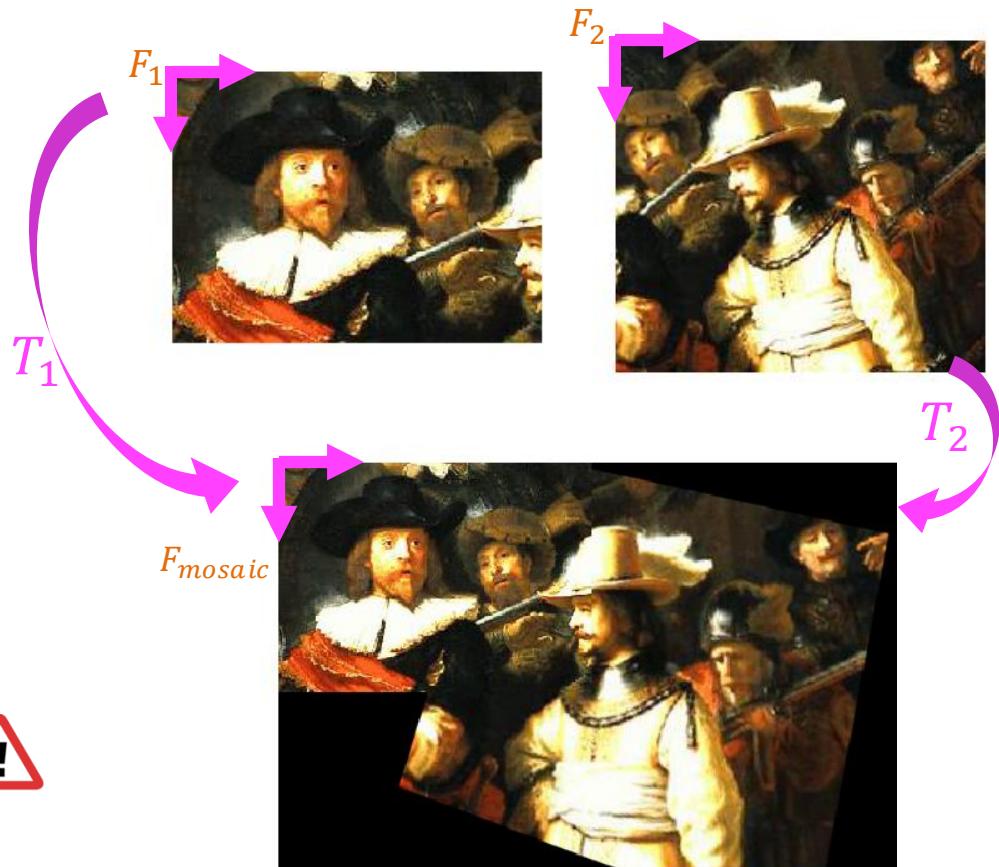


Image Stitching – Overview

1. Detect keypoints – OpenCV tuple

$$(x, y, s, \theta, d)_i$$

2. Compute a feature per keypt
3. Establish keypt *correspondences*
4. Filter corresp. & find transf. with *RANSAC*
5. Warp F_1 and F_2 into frame F_{mosaic} with their respective *final transf.* T



Note: here frames of F_{mosaic} & F_1 overlap for simplicity $\rightarrow T_{F1} = I$



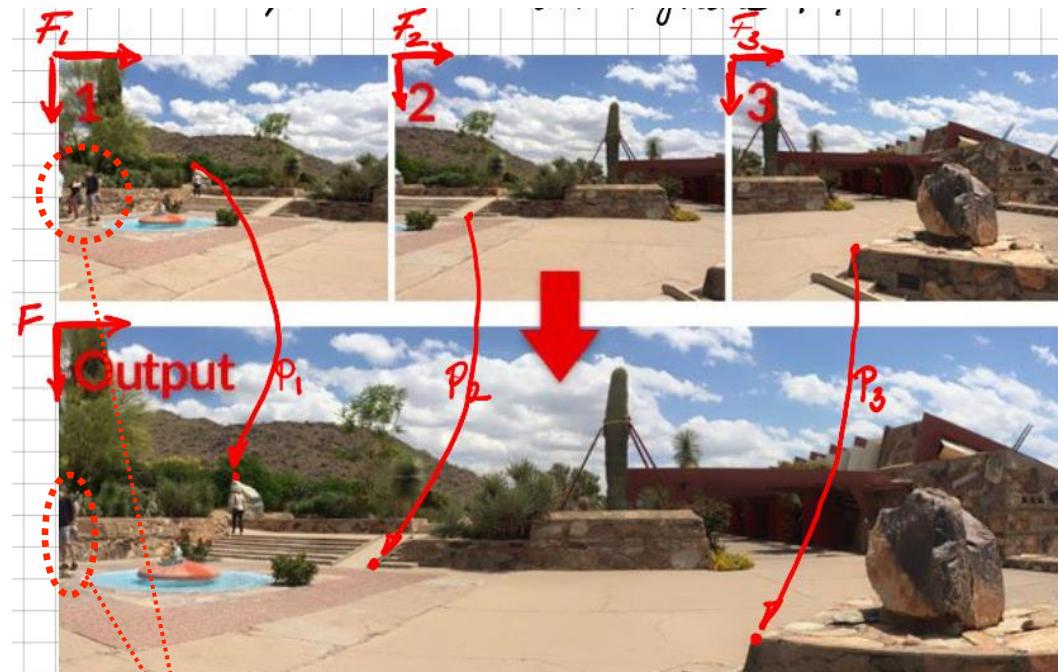
Image Stitching – Overview

1. Detect keypoints – OpenCV tuple

$$(x, y, s, \theta, \bar{d})_i$$

2. Compute a feature per keypt
3. Establish keypt *correspondences*
4. Filter corresp. & find transf. with *RANSAC*
5. Warp F_1 and F_2 into frame F_{mosaic} with their respective *final transf.* T

 Generalizes for > 2 images



Note: frames of F_{mosaic} & F_1 do not necessarily overlap

Outline

- Edge Detection
 - Derivatives of Image, Derivatives of Gaussian
 - Canny Edge Detector
- Line Fitting
 - Least Squares
 - RANSAC
 - Hough Transform
- Corners
 - Harris Corners
 - SIFT
 - Applications

Applications

- Feature points are used for:
 - Image alignment
 - 3D reconstruction
 - Motion tracking
 - Robot navigation
 - Indexing and database retrieval
 - Object recognition



Multi-band Blending



P. J. Burt and E. H. Adelson
[A multiresolution spline with application to image mosaics](#)

ACM Transaction on Graphics (ToG), 1983

Multi-band Blending



P. J. Burt and E. H. Adelson
[A multiresolution spline with application to image mosaics](#)

ACM Transaction on Graphics (ToG), 1983

Slide: Derek Hoiem 167

Multi-band Blending



P. J. Burt and E. H. Adelson
[A multiresolution spline with application to image mosaics](#)

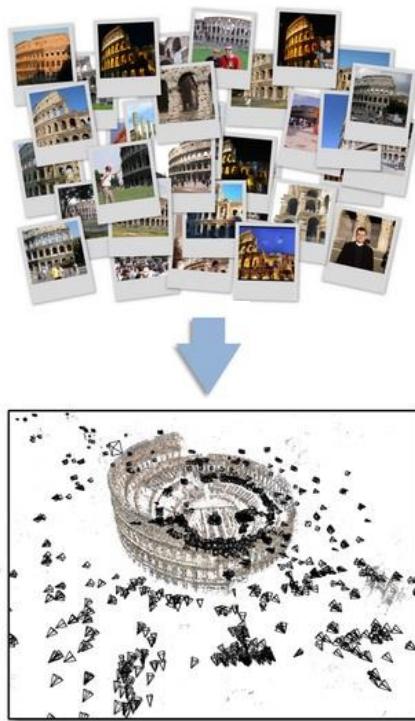
ACM Transaction on Graphics (ToG), 1983

Automatic Mosaicing



<http://www.cs.ubc.ca/~mbrown/autostitch/autostitch.html>

Example: Structure from Motion



Jan-Michael Frahm et al.
[Building rome on a cloudless day](#)
ECCV 2010



<https://youtu.be/4cEQZreQ2zQ>

Closing Remarks

Action Points

- This week:
 - Thursday Lecture: Optical Flow & Motion
 - Finish Quiz-2
 - Finish Lab-2 Assignment
 - Start Lab-3 Assignment

Disclaimer

Many of the slides used here are obtained from online resources (including many open lecture materials) without appropriate acknowledgement. They are used here for the sole purpose of classroom teaching. All the credit and all the copyrights belong to the original authors. You should not copy it, redistribute it, put it online, or use it for any other purposes than for this course.