



# Computer Vision 1

HC5a

# Retrieval, Detection, Segmentation

Dr. Martin Oswald, Dr. Dimitris Tzionas, Dr. Arun Mukundan,  
[m.r.oswald, d.tzionas, a.mukundan]@uva.nl

# Last Lecture

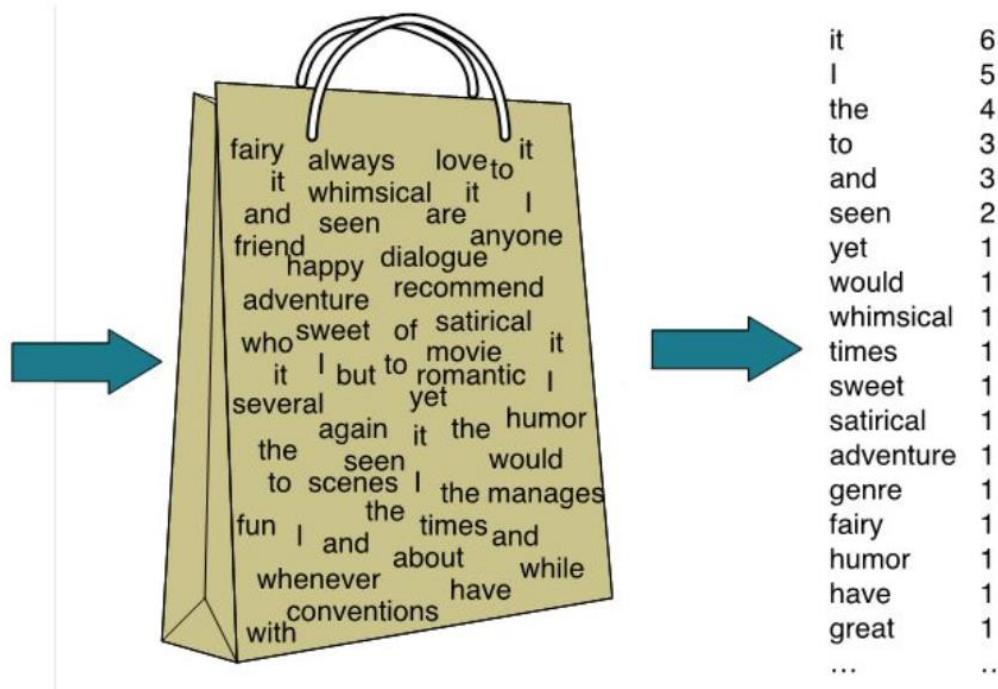
---

## Recap

# How many features are needed?

## The Bag of Words Representation

I love this movie! It's sweet, but with satirical humor. The dialogue is great and the adventure scenes are fun... It manages to be whimsical and romantic while laughing at the conventions of the fairy tale genre. I would recommend it to just about anyone. I've seen it several times, and I'm always happy to see it again whenever I have a friend who hasn't seen it yet!



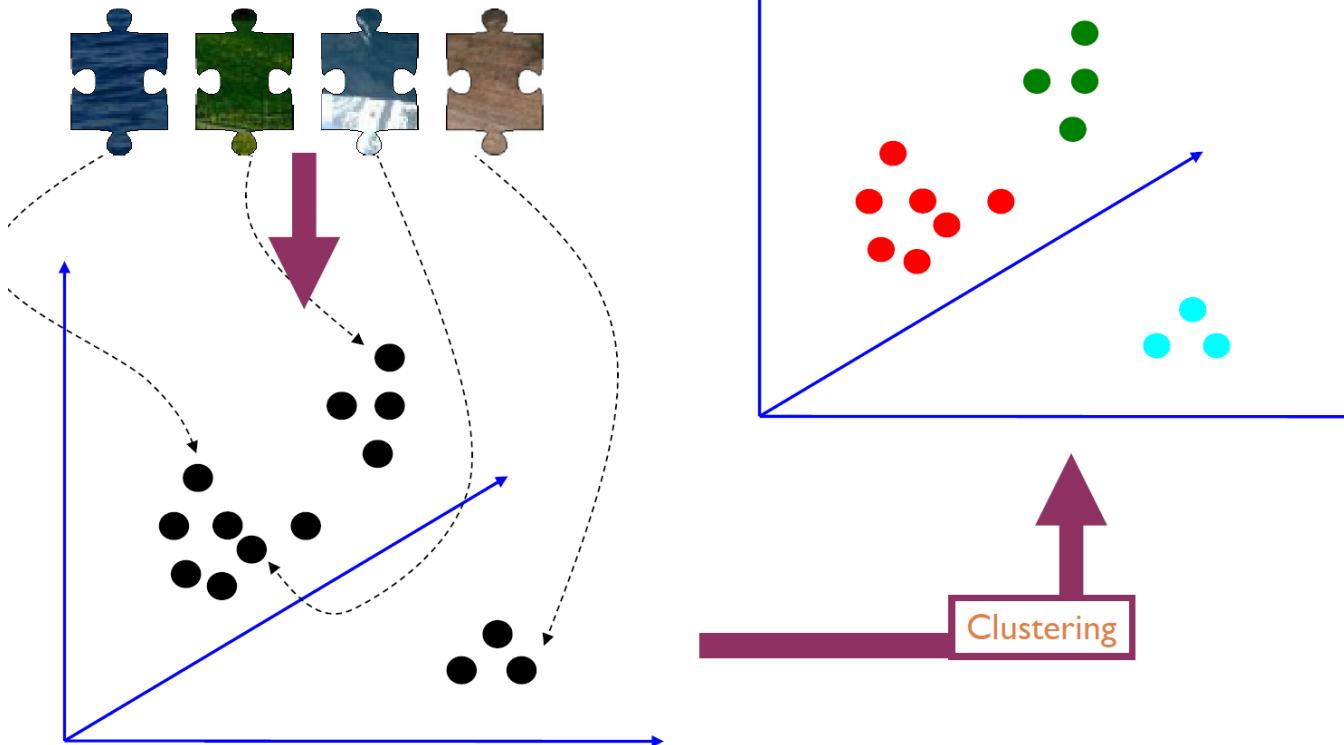
# Bag of Visual Words



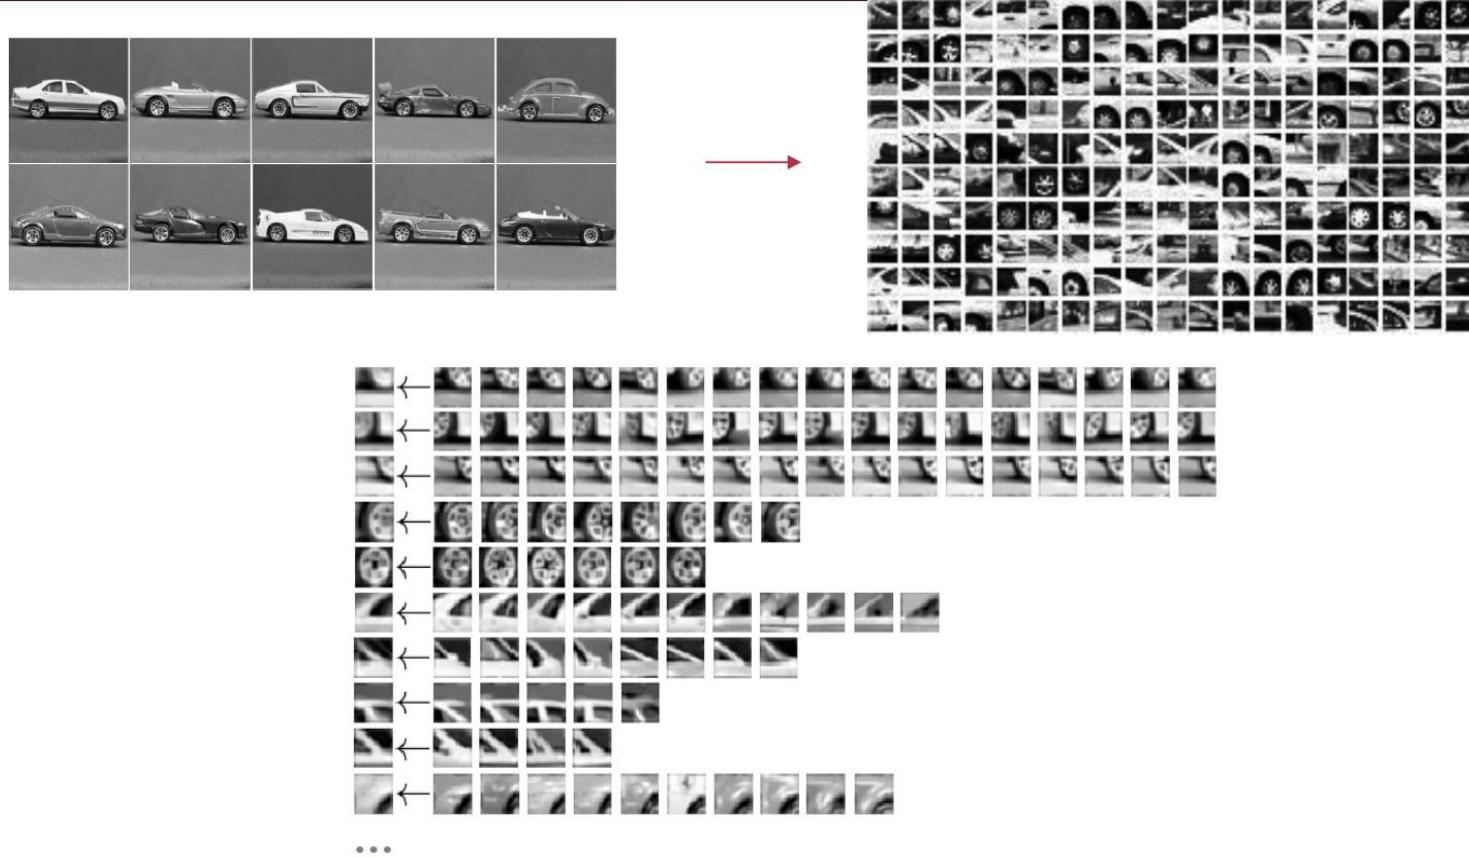
Slide credits: Jan van Gemert

# Feature Clustering

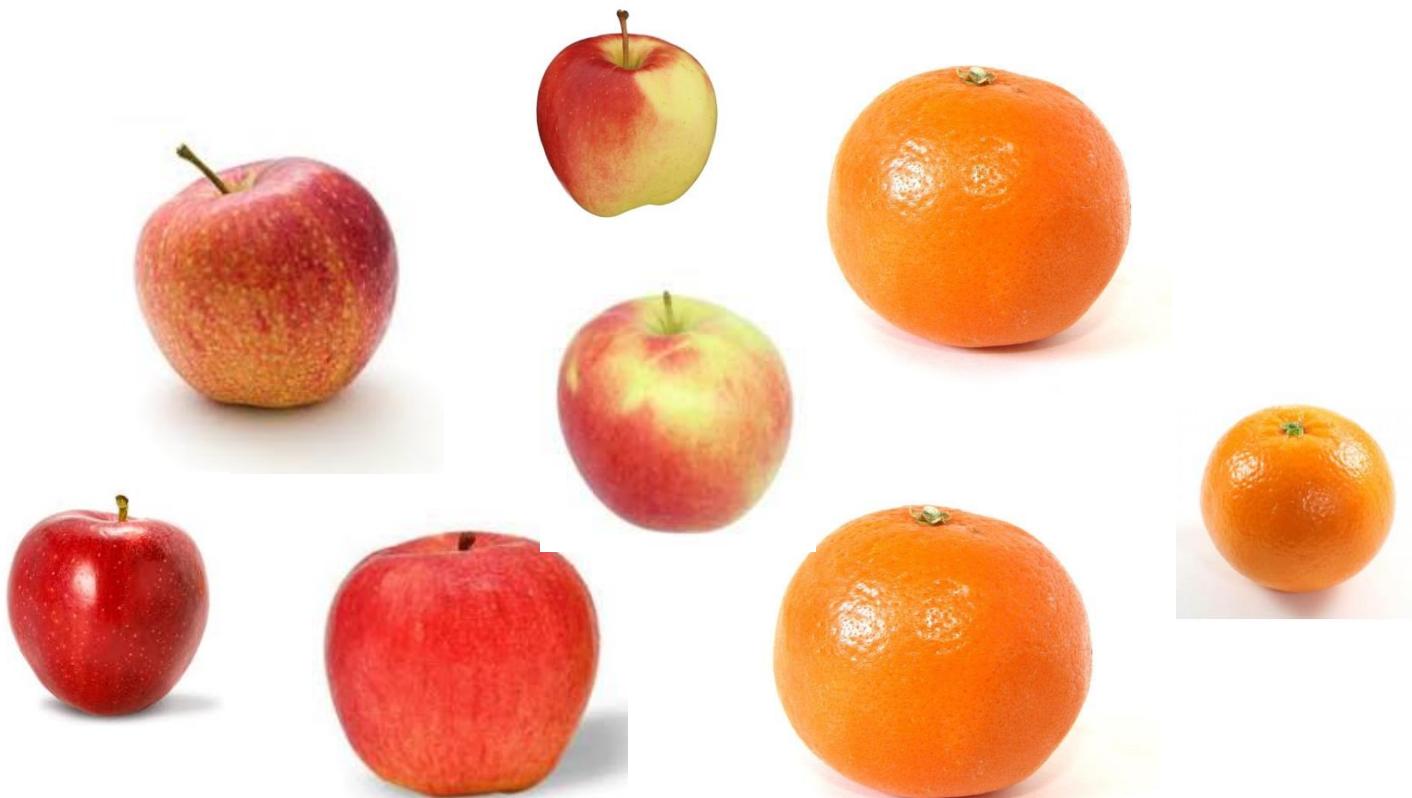
## CLUSTERING FOR BOW



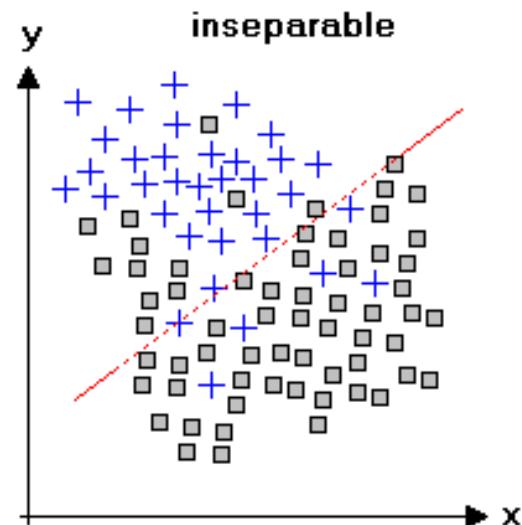
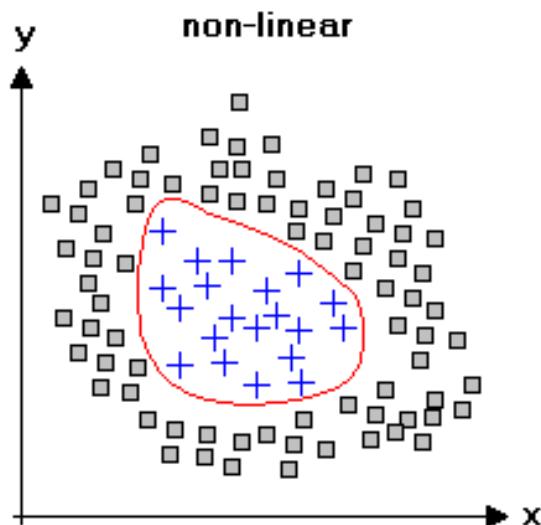
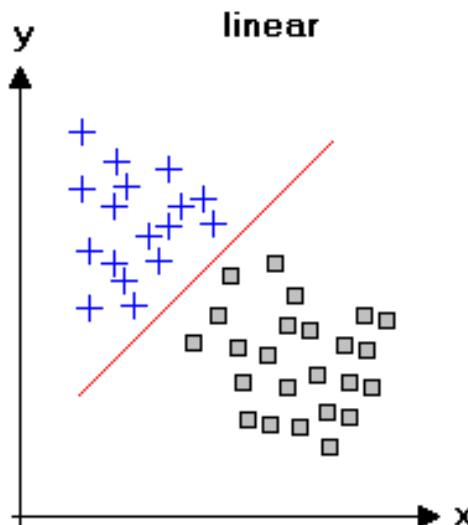
# Codebook Example



# Apple – Orange Classification



# Feature Spaces - Separability



# Classification - Performance Evaluation

- Accuracy:

$$\text{Accuracy} = \frac{\text{TN} + \text{TP}}{\text{TN} + \text{FP} + \text{TP} + \text{FN}}$$

- Precision and Recall:

$$\text{Precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$

$$\text{Recall} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

- F1 score: geometric mean of precision and recall

$$\text{F1} = 2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$$

# Precision and Recall

How many selected items are relevant?

$$\text{Precision} = \frac{\text{true positives}}{\text{selected elements}}$$

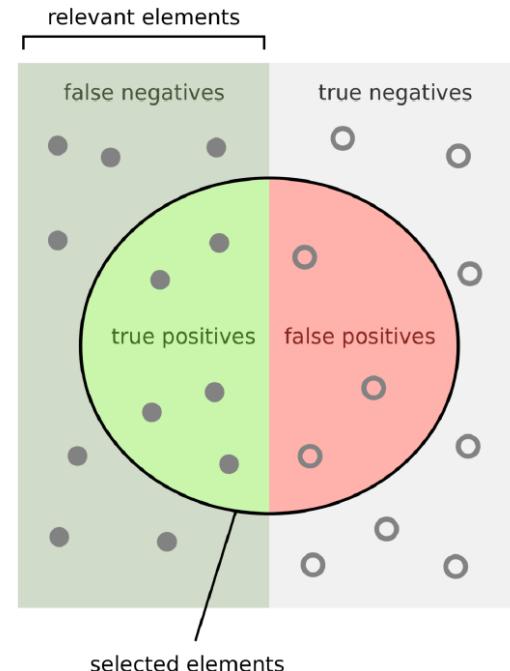


How many relevant items are selected?

$$\text{Recall} = \frac{\text{true positives}}{\text{relevant elements}}$$

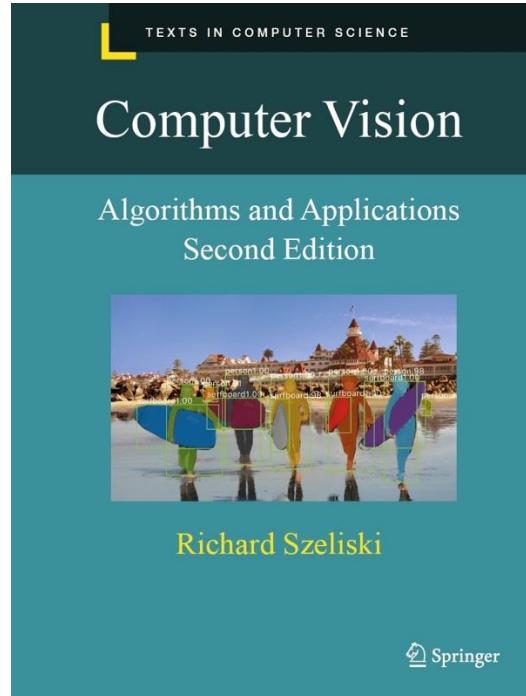


Combine P+R into: F1-Measure



# Textbook

- Chapter 6.3
- Chapter 6.4 (introduction)
- Chapter 6.4.2



Richard Szeliski

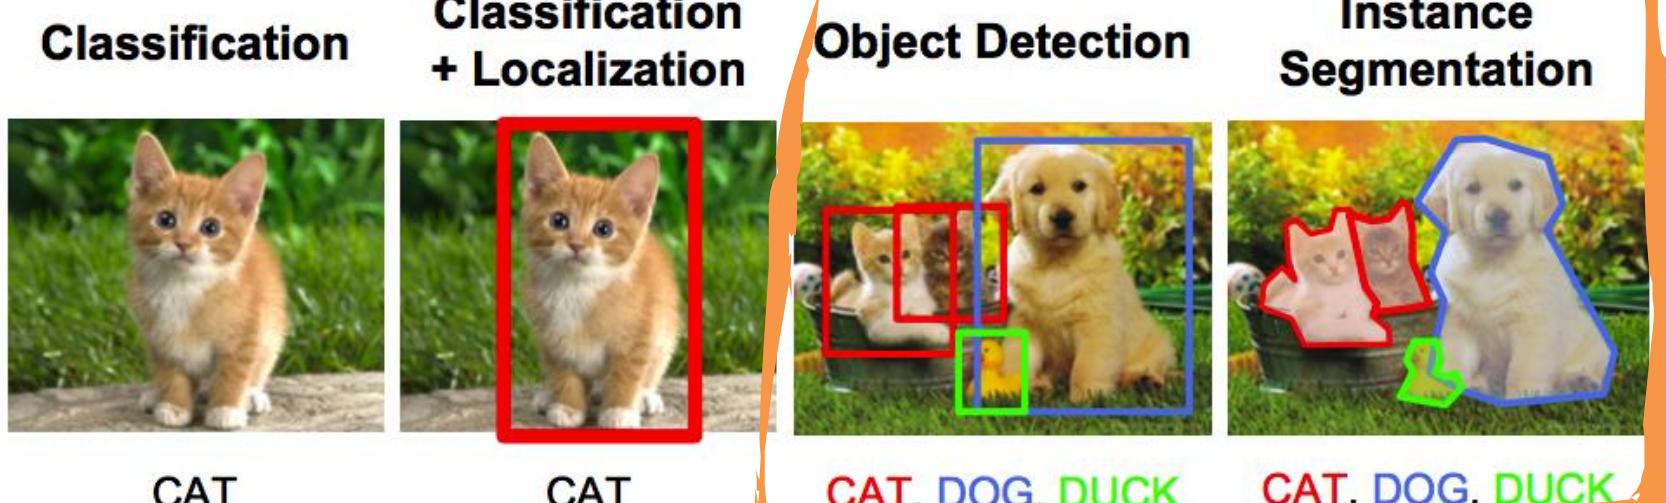
 Springer

# Image Understanding

AN IMAGE IS WORTH A THOUSAND WORDS



# Image Understanding



- Object Classification: classify the main object in an image if the image has one main object.
- Visual Image Retrieval: retrieve images from the library through a query image.
- Object Detection: identify whether there is an object in an image and recognize the object.
- Image Segmentation: segment an image into regions according the pixels the object belong to.

## Today: Object Detection + Segmentation

# Today's Agenda

---

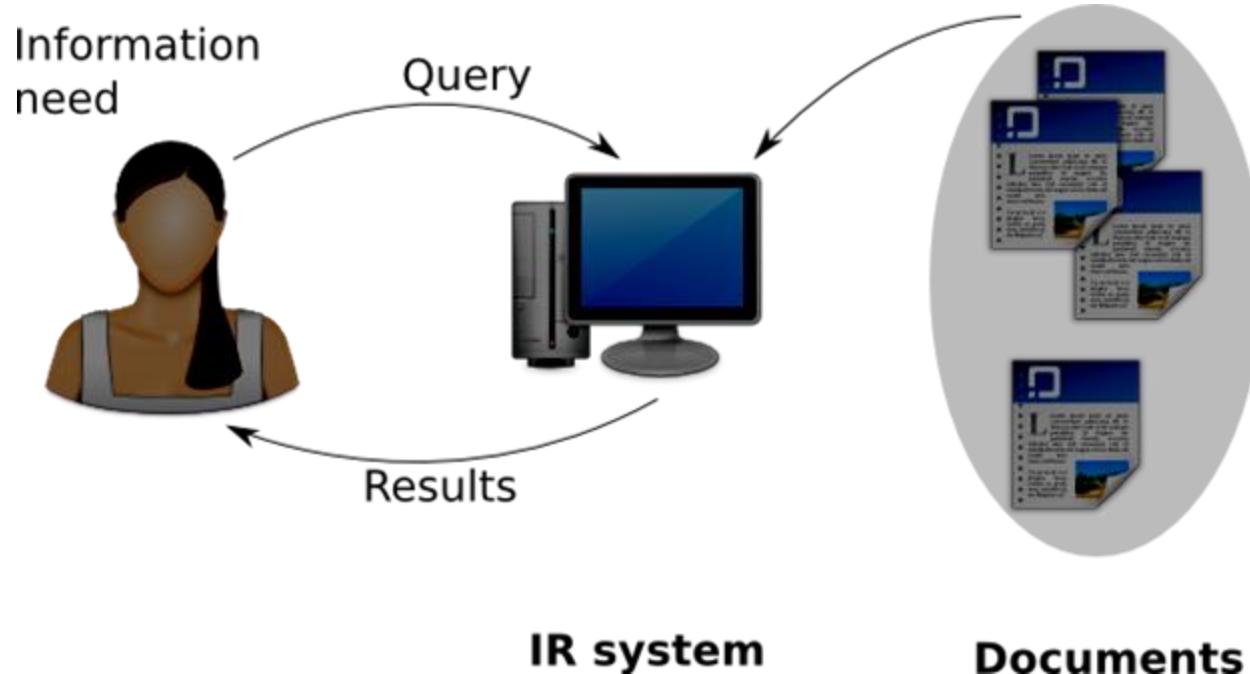
- **Image Retrieval**
- **Object Detection**
- **Image Segmentation**

# Today's Agenda

---

- **Image Retrieval**
- **Object Detection**
- **Image Segmentation**

# What is Image Retrieval?



# What is Image Retrieval?



# Image Retrieval



## Google Search

Search term: “apple”

Bing Search



# Image Retrieval



Top1

[Home - Apple](#) - ...  
apple.com

Top2

[Apple](#)  
apple.com

Top3

[Apple Trade In - Apple](#)  
apple.com

Top4

[Apple](#) - [Buy Premium 100% Natural Honeycrisp Apples](#) - Walmart  
walmart.ca[History of Apple Inc.](#) - Wikipedia  
en.wikipedia.org[Honeycrisp Apple, One Month Old](#) - Amazon  
amazon.com[Apple Watch, iPhone, iPad, iPod Touch, and more](#) - Apple  
apple.com[Buy Apple Pencil \(2nd Generation\) - Apple](#)  
apple.com[Apple](#) - YouTube  
youtube.com[Apple](#) - Home | Facebook  
facebook.com[Apple](#) | LinkedIn  
linkedin.com[11-inch iPad Pro Wi-Fi 64GB](#) - Apple  
apple.com[Apple](#) - Home | Facebook  
facebook.com[AppleCare+](#)[Buy iPhone Accessories](#) - Apple  
apple.com[Buy AirPods with Wireless Charging Case](#) - Apple  
apple.com[Popular Varieties](#) - Usapple.org  
usapple.org[iPhone 11](#) - Apple  
apple.com[iPhone XR](#) - iPhone BOGO Offer from AT&T  
att.com[History of Apple Inc.](#) - Wikipedia  
en.wikipedia.org[Apple credit card: read the fine print](#) - CNBC  
cnbc.com

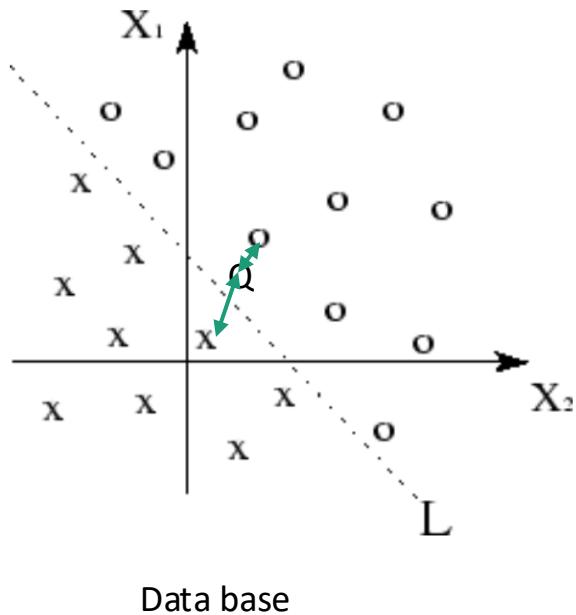
# Bag of Words for Image Retrieval

- Offline: construct BoW Vocabulary
- Offline: All images in the dataset -> BoWs
- Create description from Query image
- Compare to all images in the dataset
- Rank according to similarity

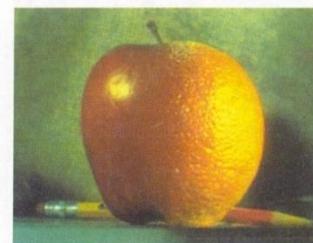


Slide credits: Jan van Gemert

# Retrieval Example



1. A Query: Q



2. Extract feature

3. Get similarity

Top 1: o

Top 2: x

...

# Retrieval Evaluation

## Precision@K

- Set a rank threshold K
- Compute % relevant in top K
- Ignores images ranked lower than K
- Example:
  - Prec@3 of 2/3
  - Prec@4 of 2/4
  - Prec@5 of 3/5

Precision

$$\text{P@K} = \frac{\overbrace{TP}{\text{Precision}}}{{\underbrace{TP + FP}}_{\text{K}}} = \frac{TP}{K}$$



# Retrieval Evaluation

## Average Precision@K

- Consider rank position of each relevant images
  - $K_1, K_2, \dots, K_R$
- Compute Precision@K for each  $K_1, K_2, \dots, K_R$
- Average precision = average of P@K

- Example:



has average precision:



$$AP@K = \frac{1}{3} * \left( \frac{1}{1} * 1 + \frac{1}{2} * 0 + \frac{2}{3} * 1 + \frac{2}{4} * 0 + \frac{3}{5} * 1 \right) \approx 0.76$$

relevance: {0,1}

$$AP@K = \frac{\sum_{k=1}^K (P@k * rel(k))}{\# \text{relevant results}}$$

# Retrieval Evaluation

P@K and AP@K for “apple” (relevant: food)



$$\text{Average Precision @ 17} = \frac{1}{3} * (1/4 + 2/6 + 3/17)$$

# Retrieval Evaluation

## Mean Average Precision (mAP)

- Take the mean over multiple queries/rankings

$$mAP@K = \frac{1}{Q} \sum_{q=1}^Q AP@k_q$$

- Example:

Ranking #1:  $(1.0 + 0.67 + 0.75 + 0.8 + 0.83 + 0.6)/6 = 0.78$

Ranking #2:  $(0.5 + 0.4 + 0.5 + 0.57 + 0.56 + 0.6)/6 = 0.52$

$$mAP@K = \frac{1}{2} * (0.78 + 0.52)$$



Ranking #1

Recall	0.17	0.17	0.33	0.5	0.67	0.83	0.83	0.83	1.0
Precision	1.0	0.5	0.67	0.75	0.8	0.83	0.71	0.63	0.56

Ranking #2

Recall	0.0	0.17	0.17	0.17	0.33	0.5	0.67	0.67	0.83	1.0
Precision	0.0	0.5	0.33	0.25	0.4	0.5	0.57	0.5	0.56	0.6

# Today's Agenda

---

- Image Retrieval
- **Object Detection**
- Image Segmentation

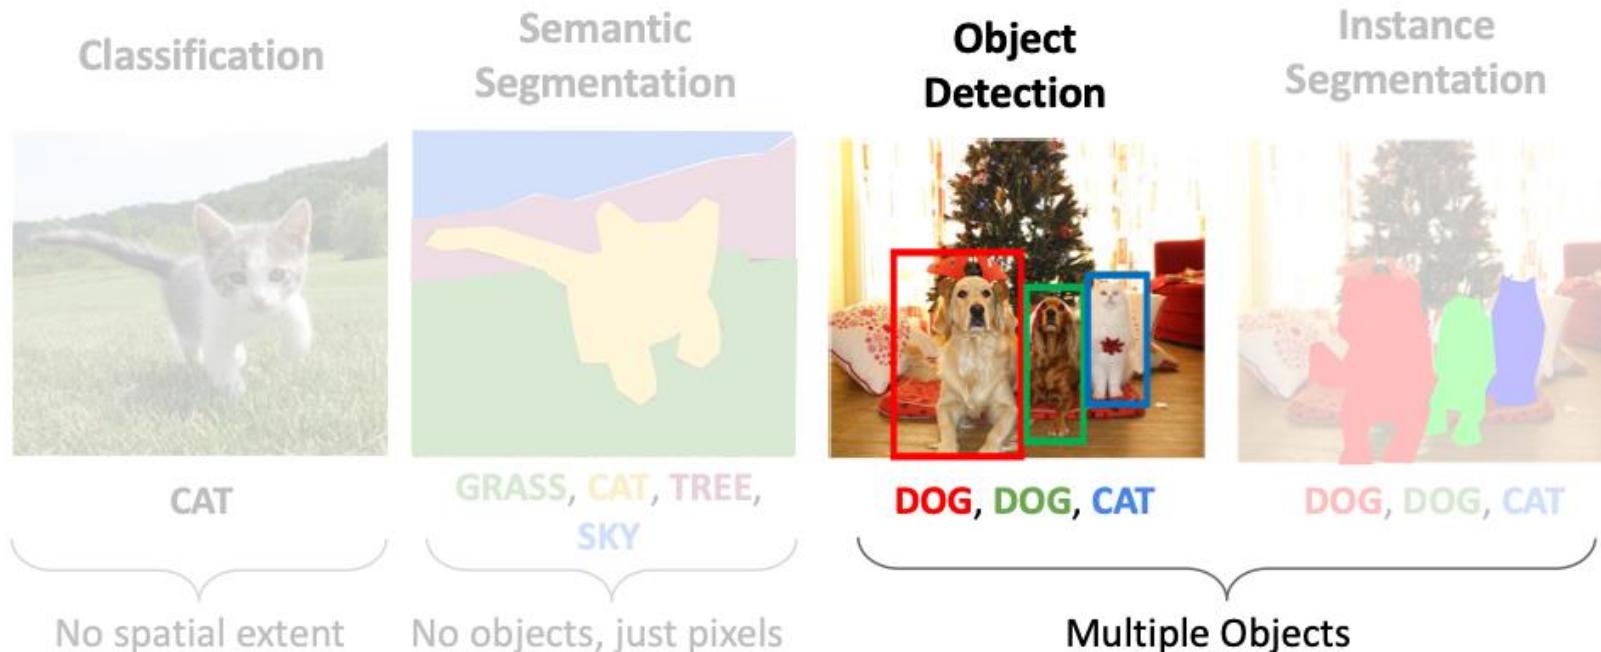
# Object Detection



# Object Detection



# Object Detection

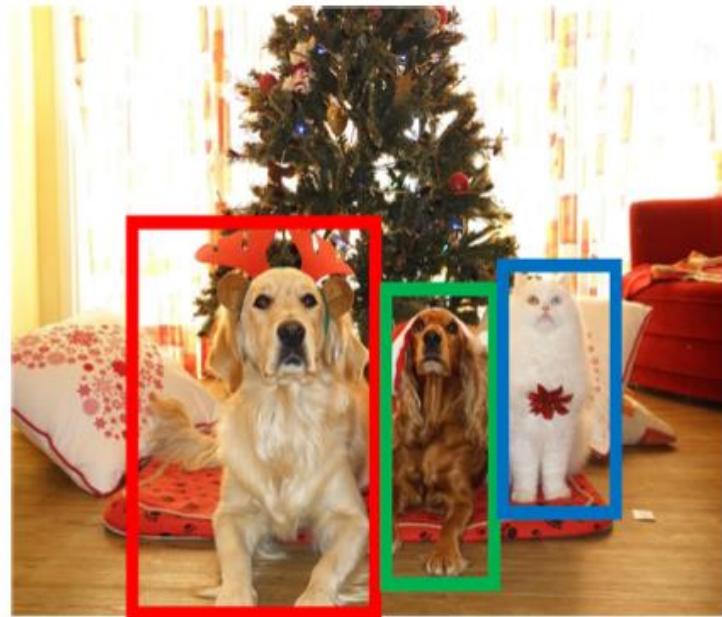


# Object Detection: Task Definition

**Input:** Single RGB Image

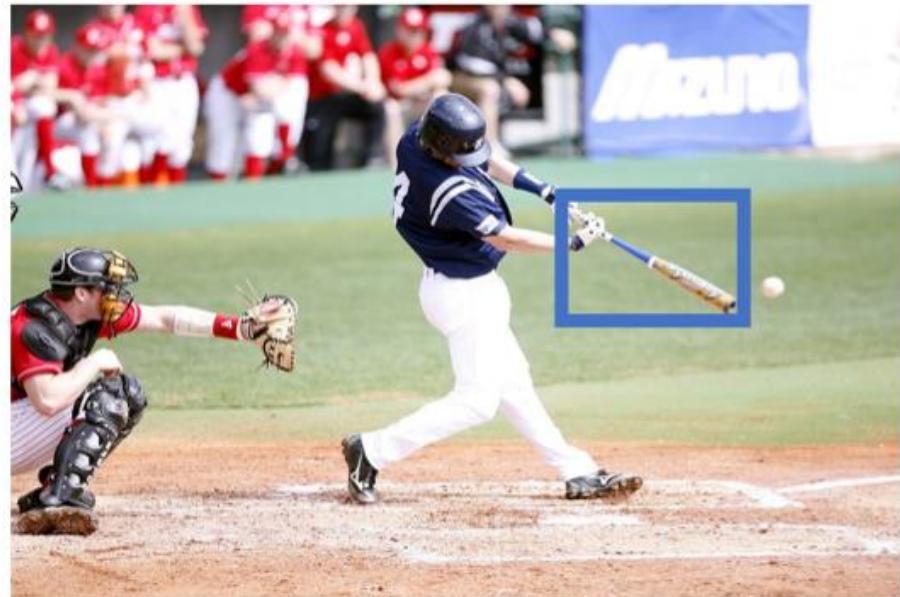
**Output:** A set of detected objects;  
For each object predict:

1. Category label (from fixed, known set of categories)
2. Bounding box (four numbers: x, y, width, height)



# Bounding Boxes

Bounding boxes are typically *axis-aligned*



# Bounding Boxes

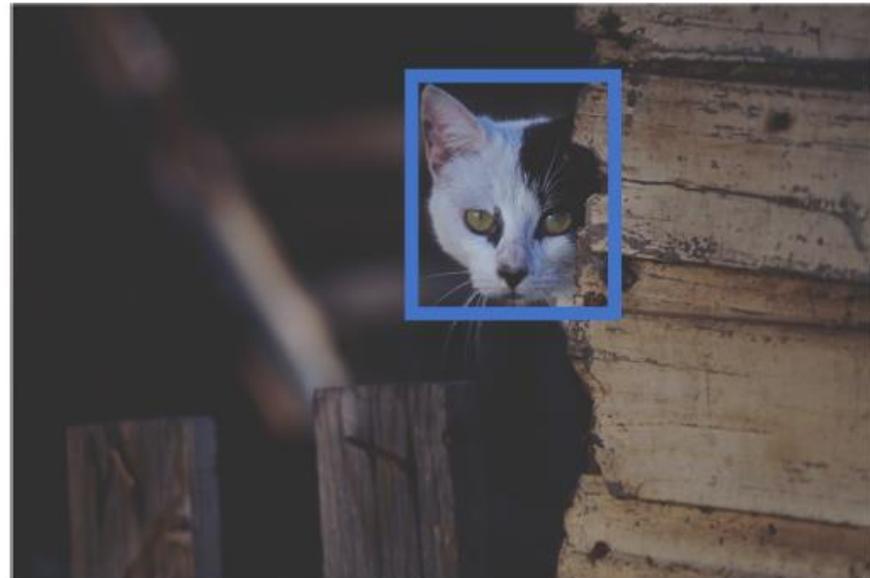
Bounding boxes are typically *axis-aligned*

*Oriented* boxes are much less common



# Bounding Boxes

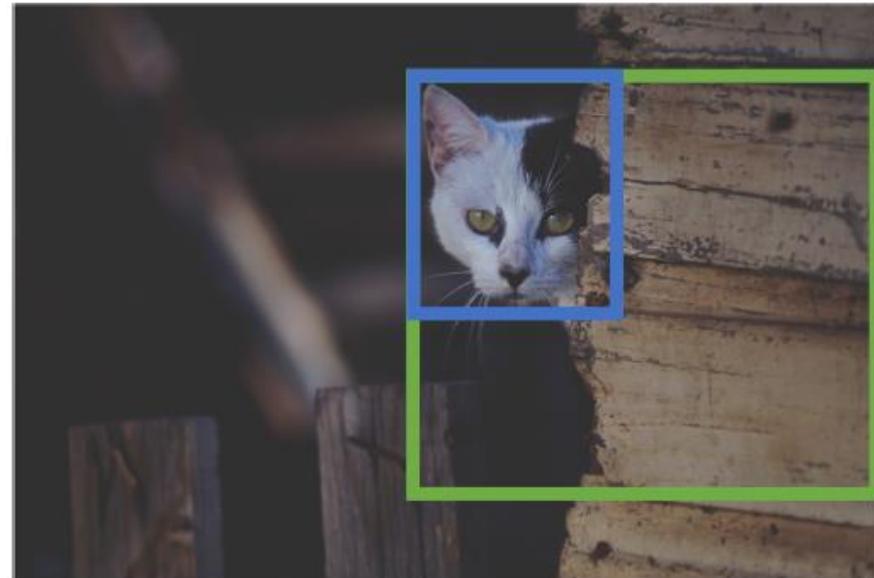
Bounding boxes (usually)  
cover only the visible  
portion of the object



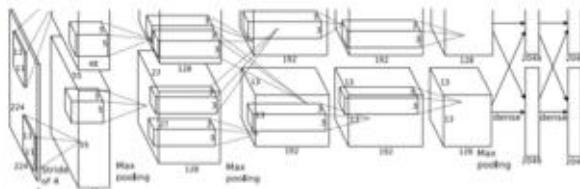
# Bounding Boxes

Bounding boxes (usually) cover only the visible portion of the object

Amodal detection:  
box covers the entire extent of the object, even occluded parts

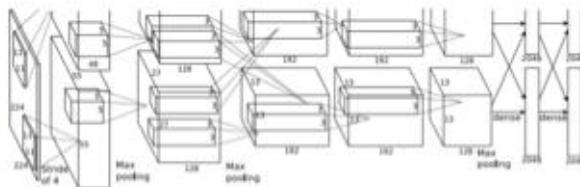


# Detecting Multiple Objects



CAT: (x, y, w, h)

4 numbers

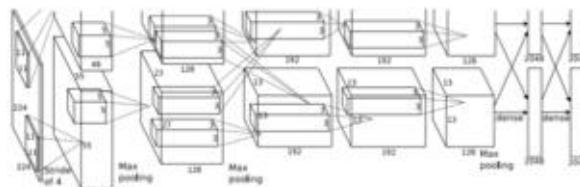


DOG: (x, y, w, h)

DOG: (x, y, w, h)

CAT: (x, y, w, h)

12 numbers



DUCK: (x, y, w, h)

Many  
numbers!

....

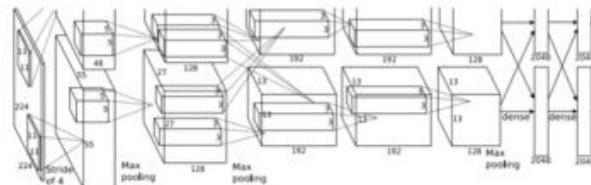
# Object Detection

---

- Sliding window
- Selective Search
- Region proposals (next lecture)

# Detecting Multiple Objects: Sliding Window

Apply a CNN to many different crops of the image, CNN classifies each crop as object or background

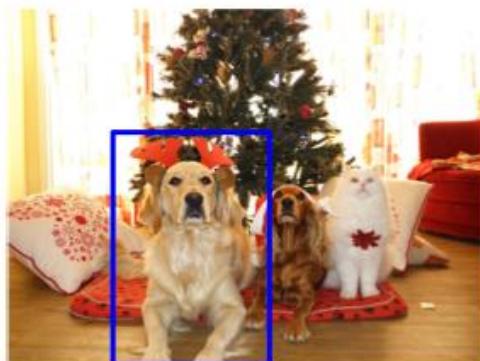


Dog? NO

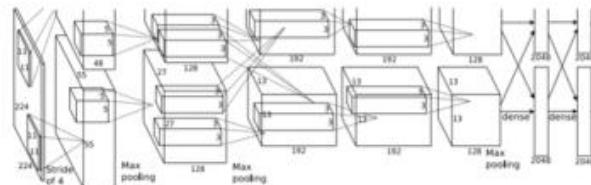
Cat? NO

Background? YES

# Detecting Multiple Objects: Sliding Window



Apply a CNN to many different crops of the image, CNN classifies each crop as object or background



Dog? YES

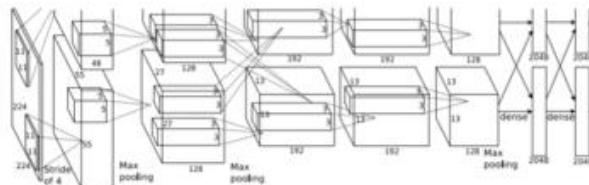
Cat? NO

Background? NO

# Detecting Multiple Objects: Sliding Window



Apply a CNN to many different crops of the image, CNN classifies each crop as object or background

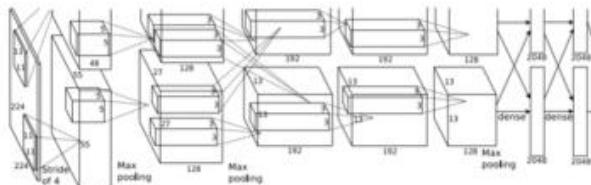


Dog? YES  
Cat? NO  
Background? NO

# Detecting Multiple Objects: Sliding Window



Apply a CNN to many different crops of the image, CNN classifies each crop as object or background



Dog? NO

Cat? YES

Background? NO

# Detecting Multiple Objects: Sliding Window



Apply a CNN to many different crops of the image, CNN classifies each crop as object or background

Question: How many possible boxes are there in an image of size  $H \times W$ ?

# Detecting Multiple Objects: Sliding Window



Apply a CNN to many different crops of the image, CNN classifies each crop as object or background

**Question:** How many possible boxes are there in an image of size  $H \times W$ ?

Consider a box of size  $h \times w$ :

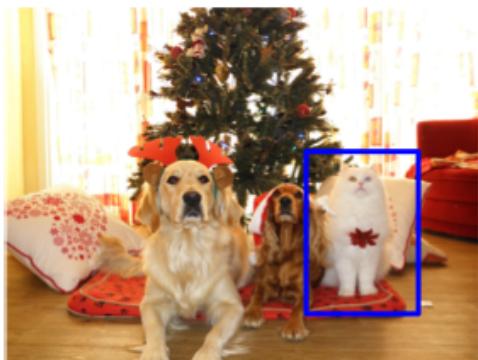
Possible x positions:  $W - w + 1$

Possible y positions:  $H - h + 1$

Possible positions:

$$(W - w + 1) * (H - h + 1)$$

# Detecting Multiple Objects: Sliding Window



Apply a CNN to many different crops of the image, CNN classifies each crop as object or background

**Question:** How many possible boxes are there in an image of size  $H \times W$ ?

Consider a box of size  $h \times w$ :

Possible x positions:  $W - w + 1$

Possible y positions:  $H - h + 1$

Possible positions:

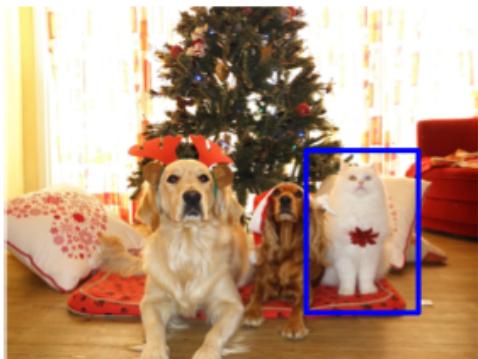
$$(W - w + 1) * (H - h + 1)$$

Total possible boxes:

$$\sum_{h=1}^H \sum_{w=1}^W (W - w + 1)(H - h + 1)$$

$$= \frac{H(H + 1)}{2} \frac{W(W + 1)}{2}$$

# Detecting Multiple Objects: Sliding Window



Apply a CNN to many different crops of the image, CNN classifies each crop as object or background

**Question:** How many possible boxes are there in an image of size  $H \times W$ ?

Consider a box of size  $h \times w$ :

Possible x positions:  $W - w + 1$

Possible y positions:  $H - h + 1$

Possible positions:

$$(W - w + 1) * (H - h + 1)$$

800 x 600 image  
has ~58M boxes!  
No way we can evaluate them all

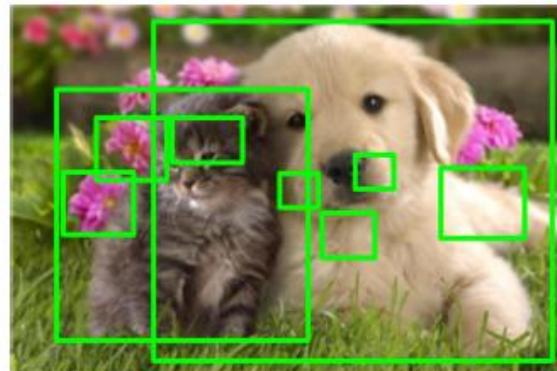
Total possible boxes:

$$\sum_{h=1}^H \sum_{w=1}^W (W - w + 1)(H - h + 1)$$

$$= \frac{H(H + 1)}{2} \frac{W(W + 1)}{2}$$

# Region Proposals

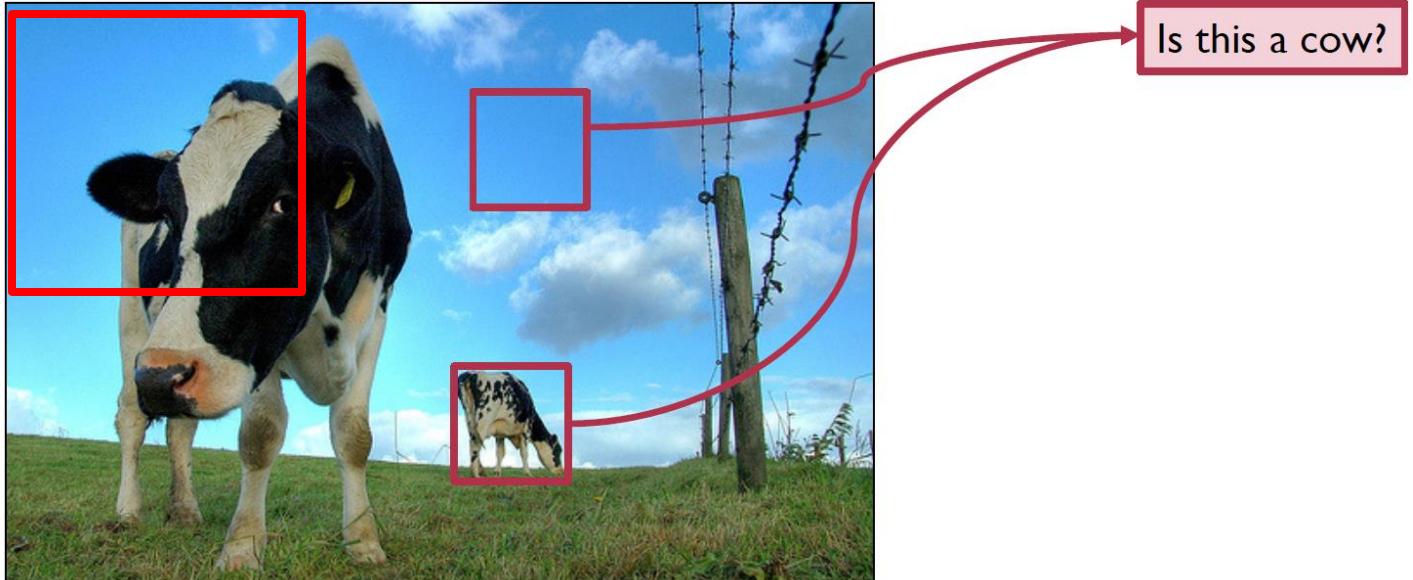
- Find a small set of boxes that are likely to cover all objects
- Often based on heuristics: e.g. look for “blob-like” image regions
- Relatively fast to run; e.g. Selective Search gives 2000 region proposals in a few seconds on CPU



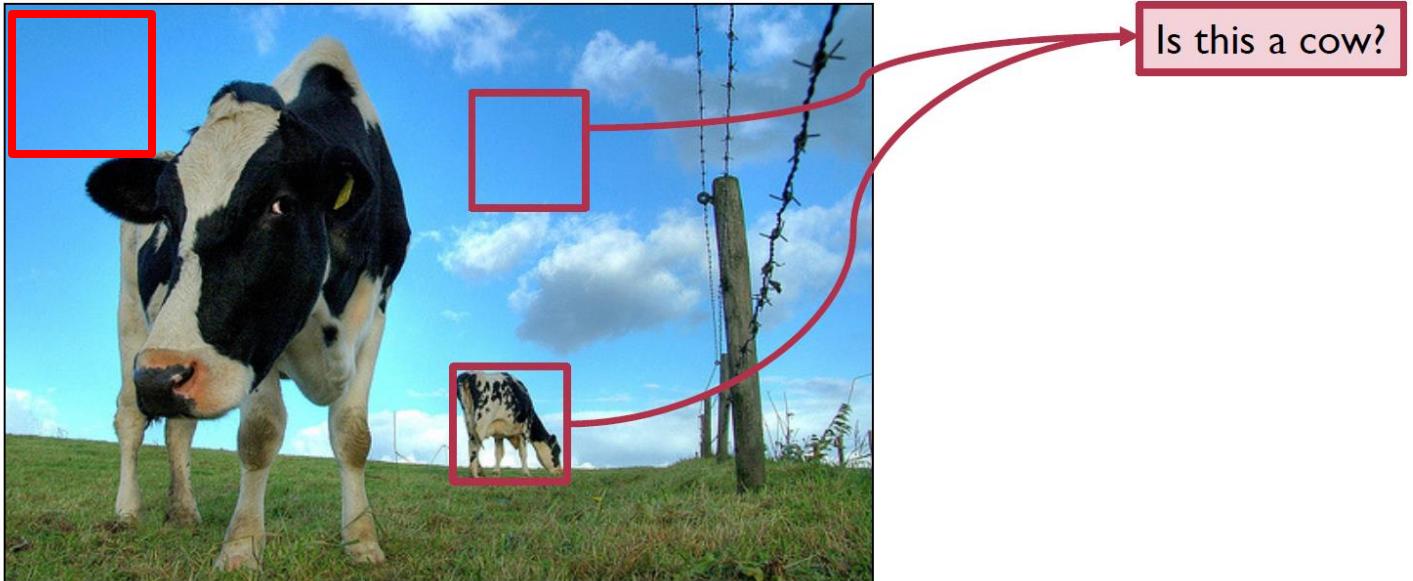
# How to Find The Bounding Box?



# Sliding Window



# Sliding Window



# Sliding Window



- + Many boxes
- + Many scales
- + Many box ratios
- + Many classes
- = Computational Expensive

# Selective Search

- Idea: Images are hierarchical

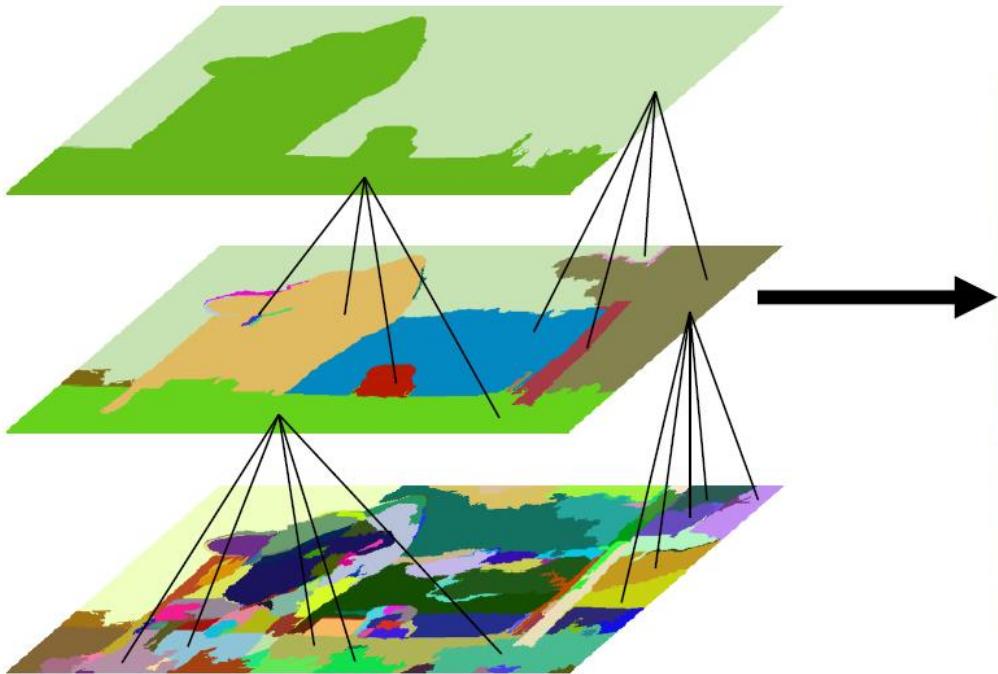


# Selective Search

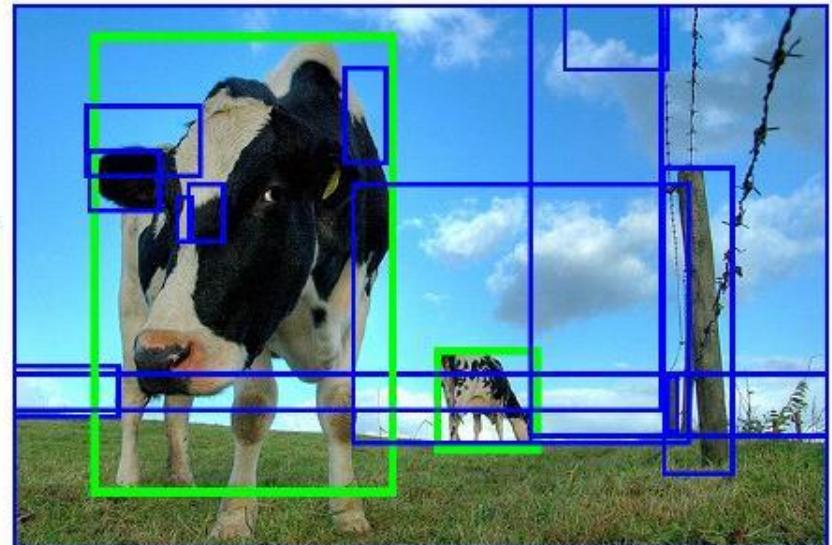
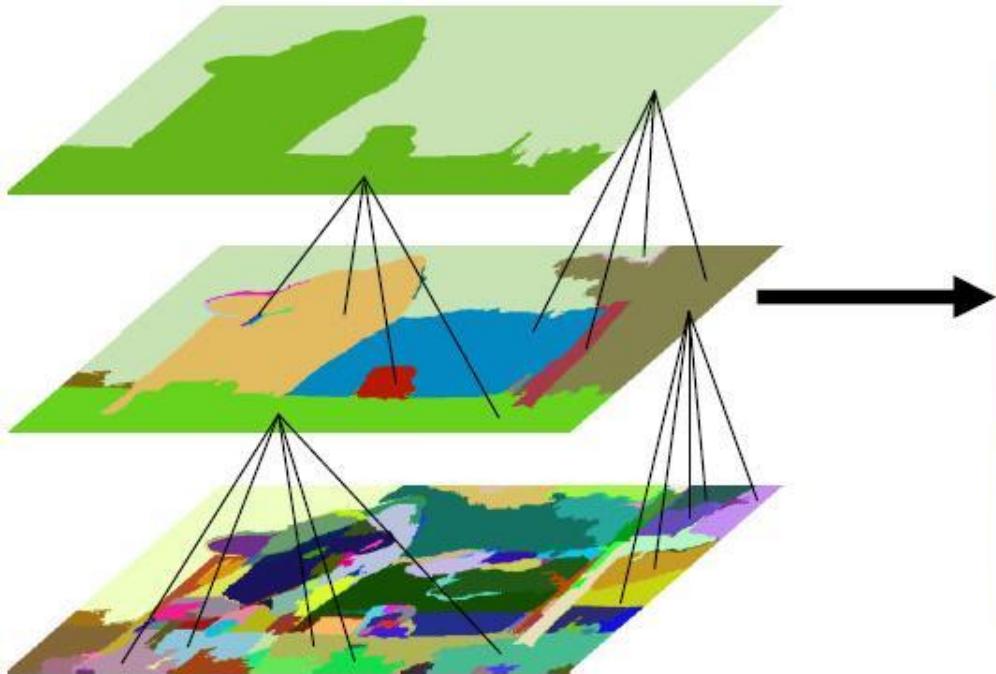
- Initial segments from over-segmentation



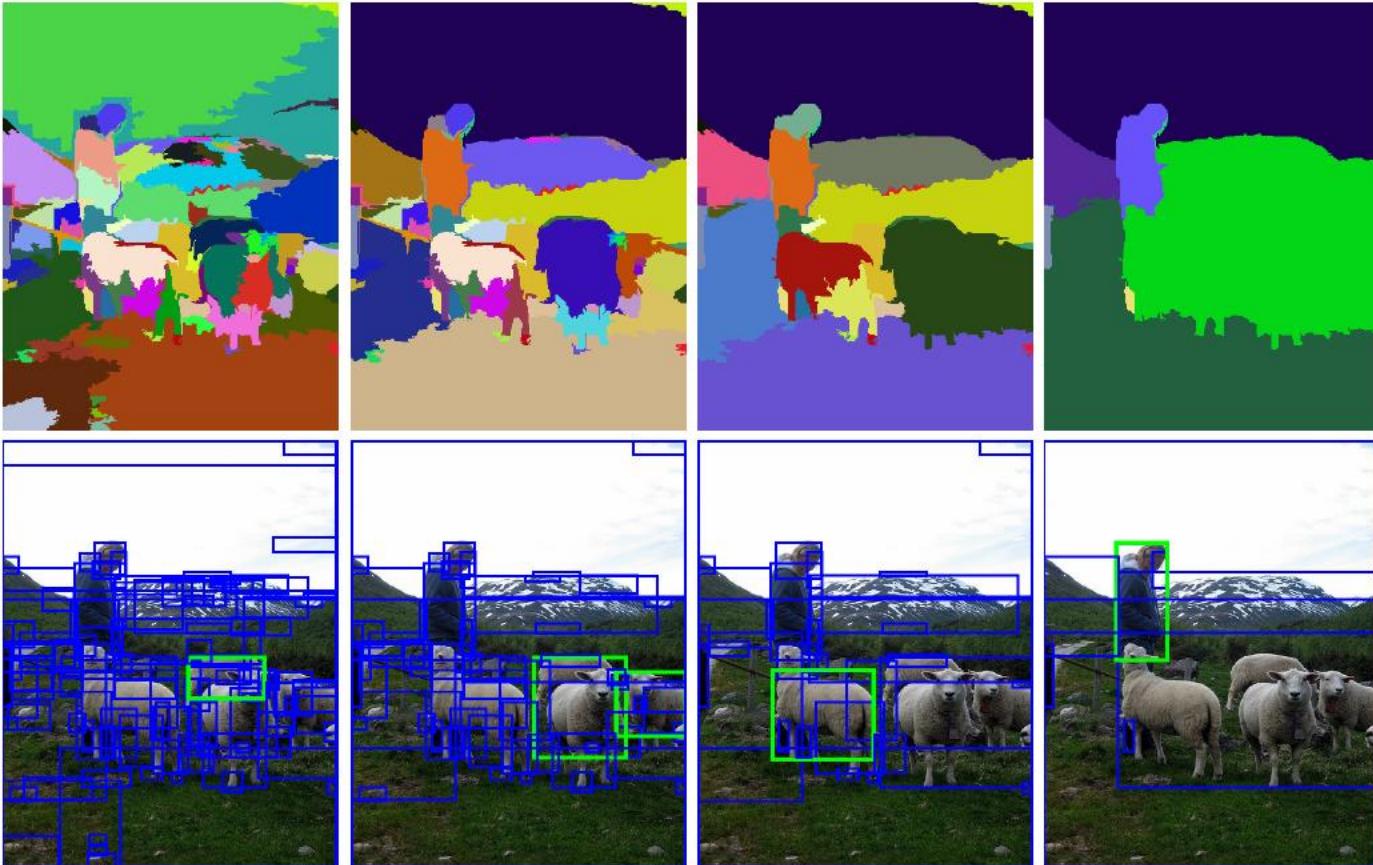
# Selective Search



# Selective Search



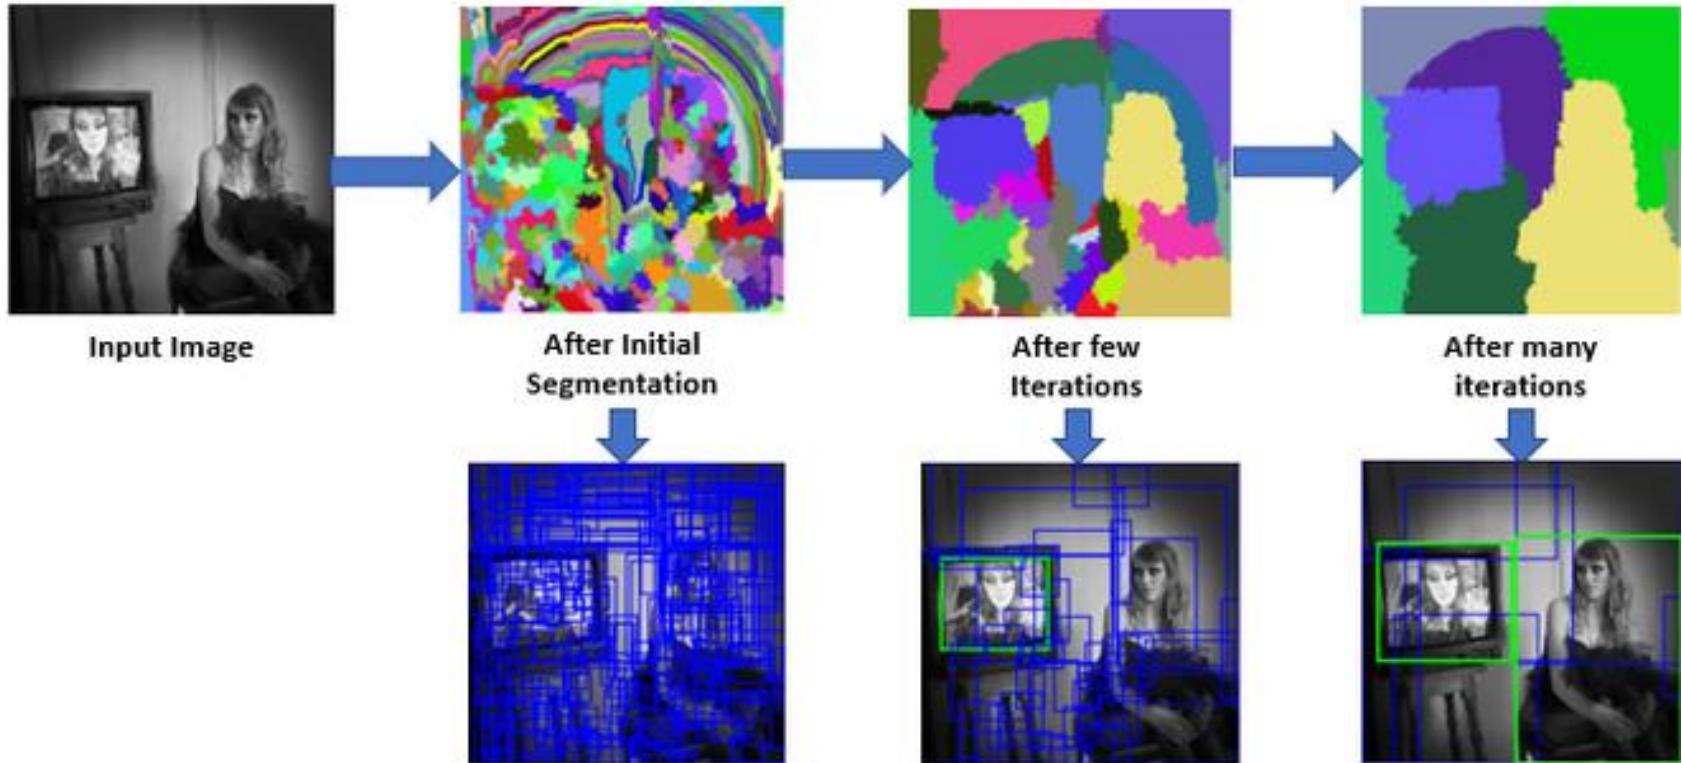
# Selective Search - Examples



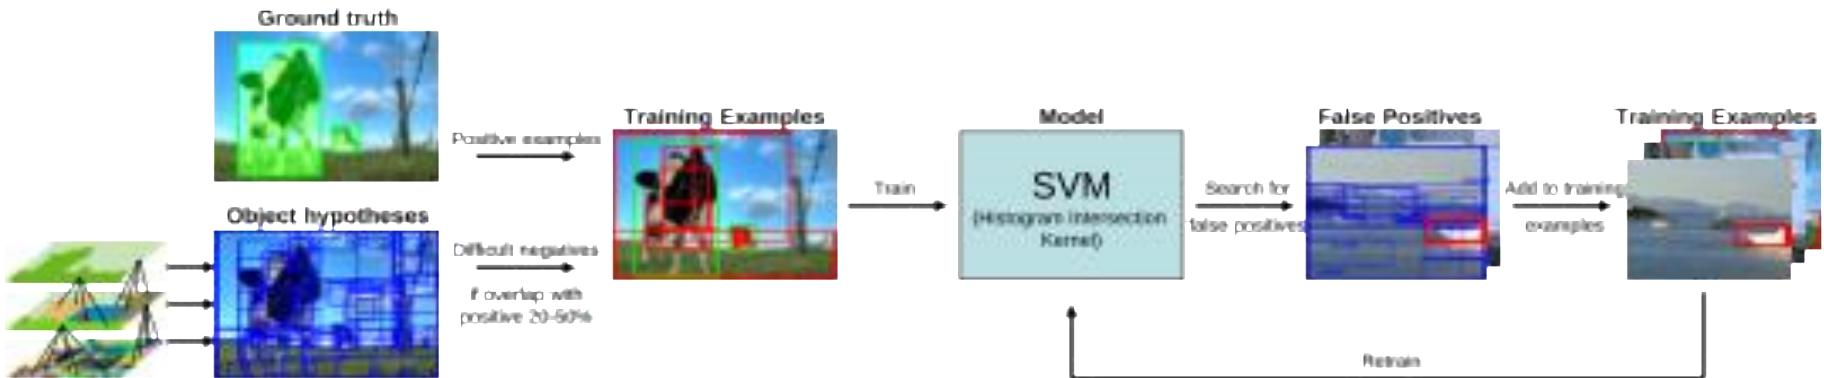
# Selective Search - Examples



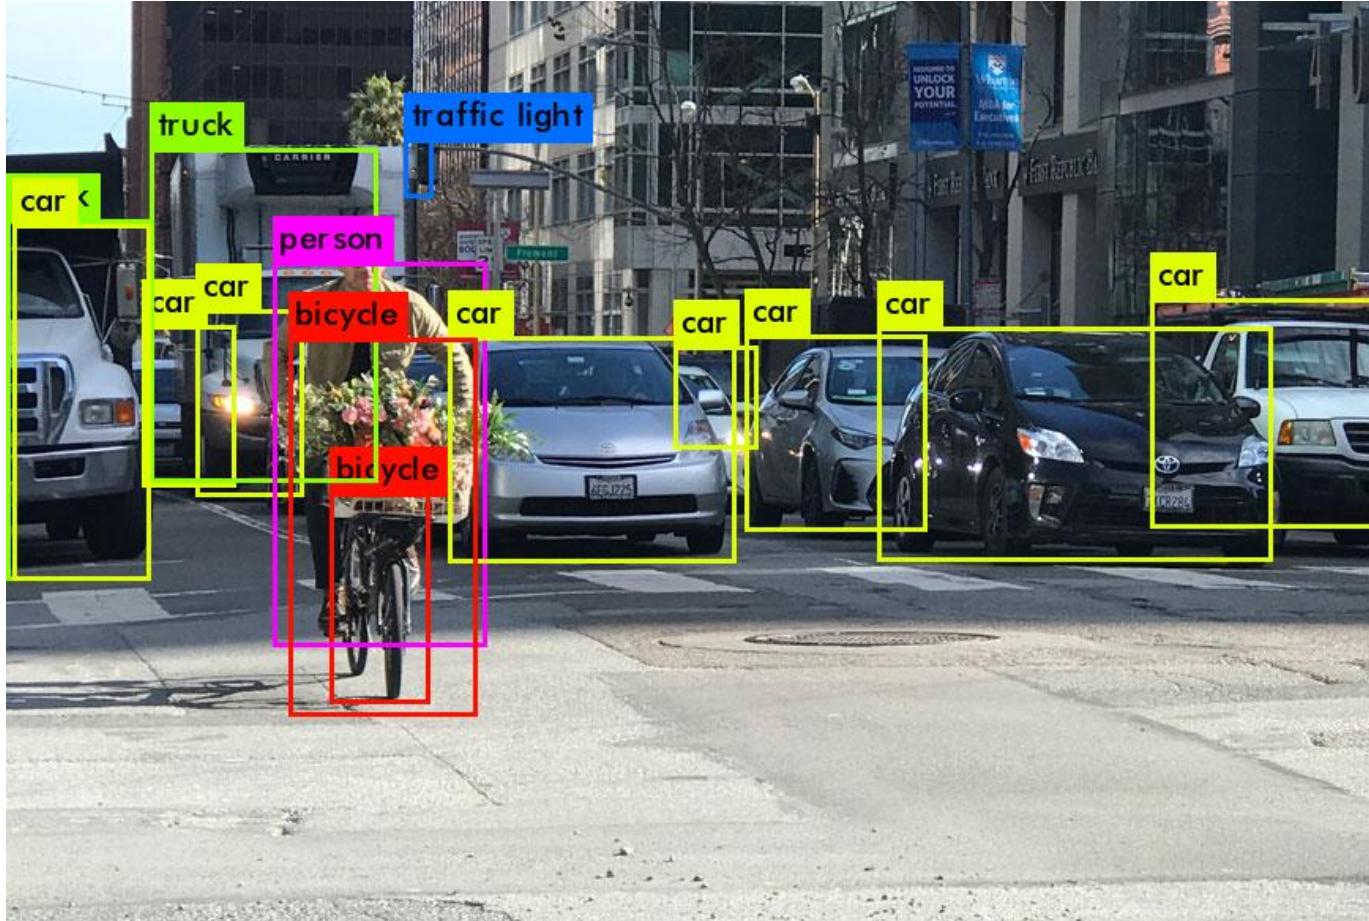
# Selective Search - Examples



# Selective Search in Object Recognition



# You only look once (YOLO)



# You only look once (YOLO)



# Autonomous Driving



# Object Detection Performance Evaluation



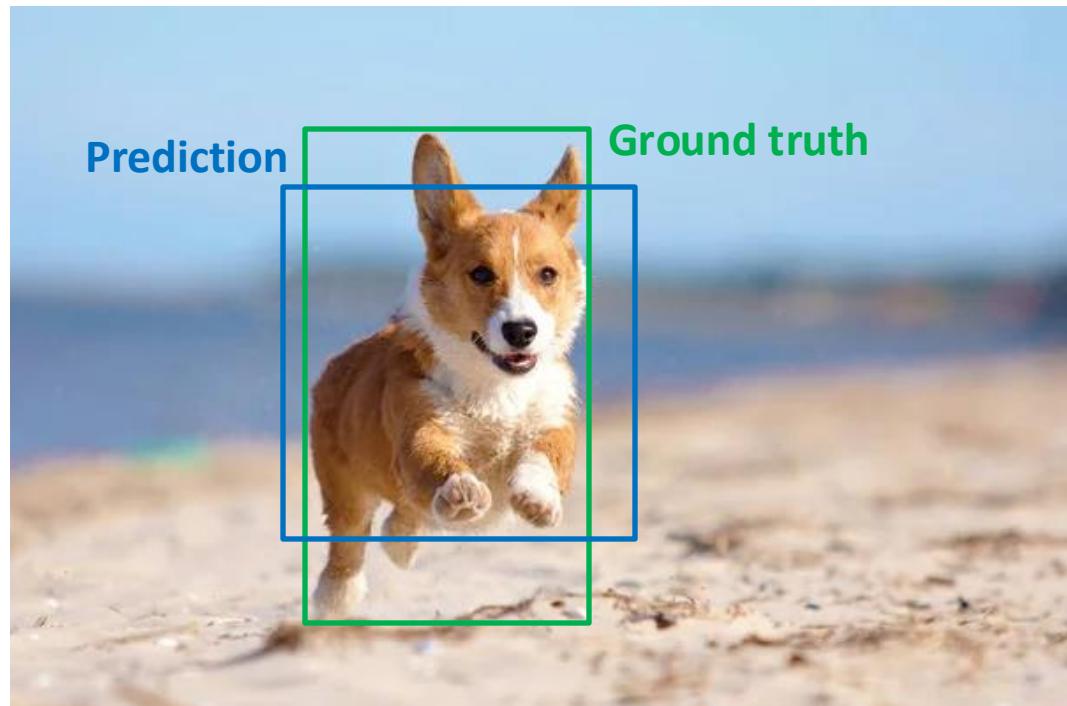
- How can we evaluate object detection?
- If we binary select inliniers and outliers of predicted bounding boxes, we can use retrieval evaluation measures:  
Precision and Recall, AP, mAP, etc.

# Comparing Bounding Boxes

- Intersection over Union (IoU)  
(a.k.a. “Jaccard similarity/index”)

$$\text{IoU} = \frac{\text{Intersection}}{\text{Union}}$$

- Typically choose a threshold  
e.g.  $\text{IoU} \geq 0.5$ , or 0.7  
between prediction and GT  
to decide about correct or  
incorrect predictions



[<https://depositphotos.com/photos/corgi-running.html>]

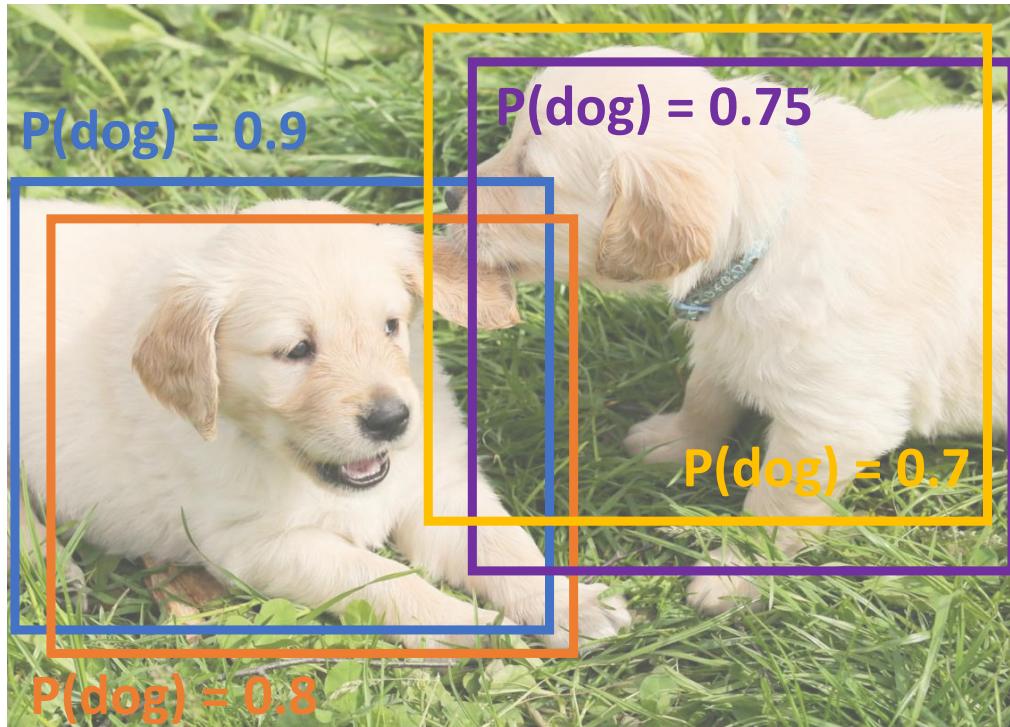
# Duplicate Filter: Non-Maximum Suppression (NMS)



**Problem:** Object detectors often output many overlapping detections

**Solution:** Post-process raw detections using **Non-Max Suppression (NMS)**:

1. Select next highest-scoring box
2. Eliminate lower-scoring boxes with  $\text{IoU} > \text{threshold}$  (e.g. 0.7)
3. If any boxes remain, GOTO 1



[Puppy image is CC0 Public Domain]

# Duplicate Filter: Non-Maximum Suppression (NMS)



**Problem:** Object detectors often output many overlapping detections

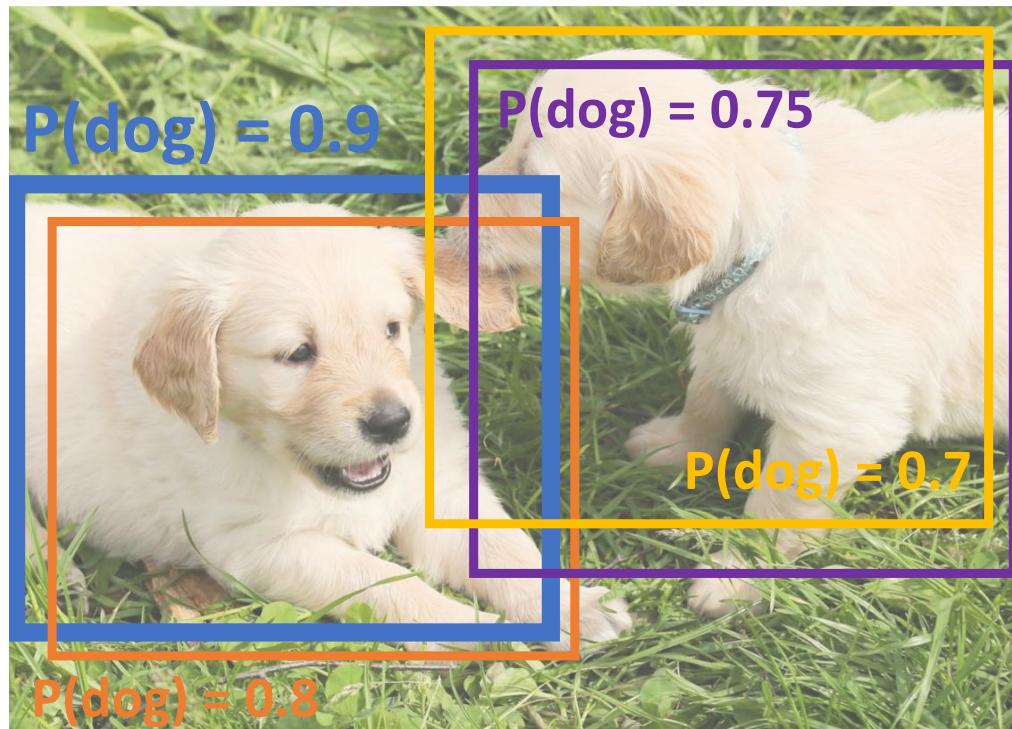
**Solution:** Post-process raw detections using **Non-Max Suppression (NMS)**:

1. Select next highest-scoring box
2. Eliminate lower-scoring boxes with  $\text{IoU} > \text{threshold}$  (e.g. 0.7)
3. If any boxes remain, GOTO 1

$$\text{IoU}(\textcolor{blue}{\square}, \textcolor{orange}{\square}) = 0.78$$

$$\text{IoU}(\textcolor{blue}{\square}, \textcolor{purple}{\square}) = 0.05$$

$$\text{IoU}(\textcolor{blue}{\square}, \textcolor{yellow}{\square}) = 0.07$$



[Puppy image is CC0 Public Domain]

# Duplicate Filter: Non-Maximum Suppression (NMS)

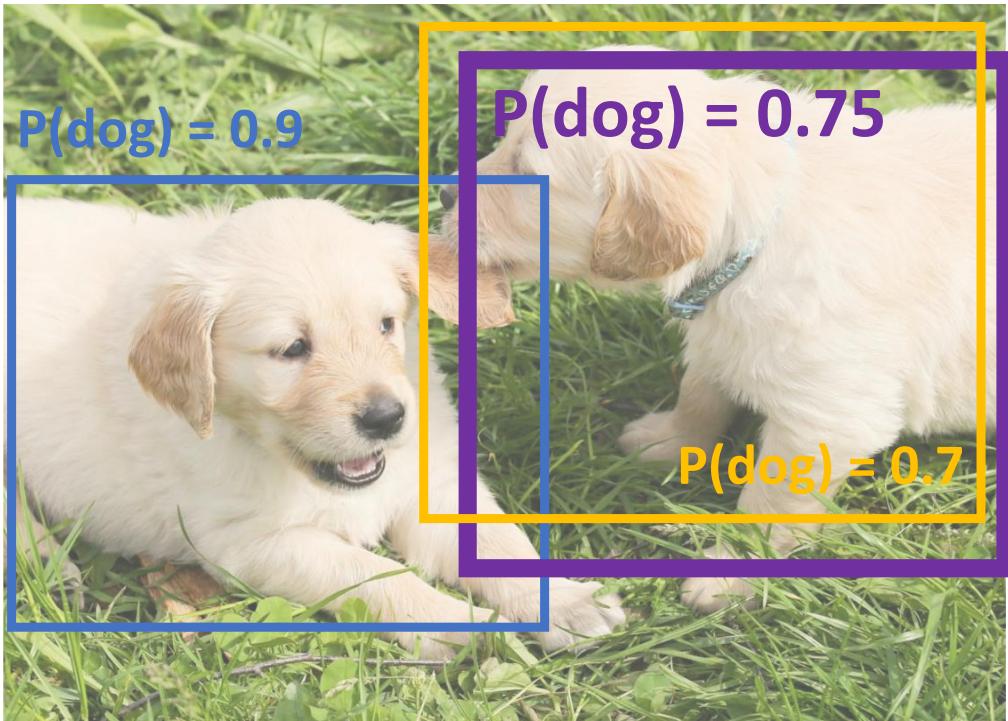


**Problem:** Object detectors often output many overlapping detections

**Solution:** Post-process raw detections using **Non-Max Suppression (NMS):**

1. Select next highest-scoring box
2. Eliminate lower-scoring boxes with  $\text{IoU} > \text{threshold}$  (e.g. 0.7)
3. If any boxes remain, GOTO 1

$$\text{IoU}(\blacksquare, \blacksquare) = 0.74$$



[Puppy image is CC0 Public Domain]

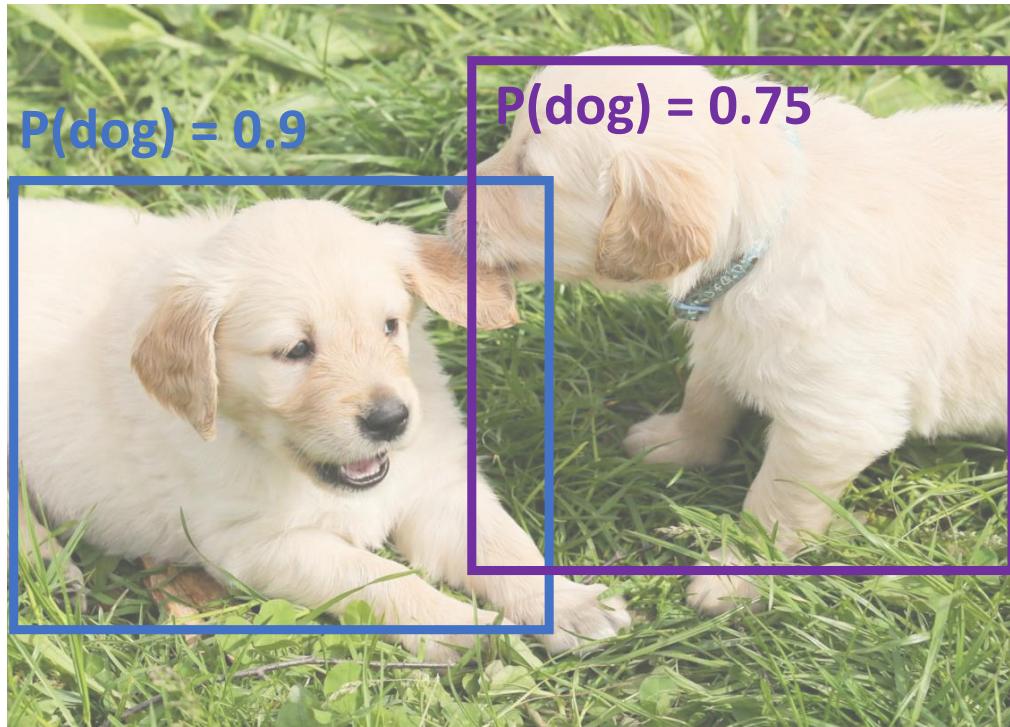
# Duplicate Filter: Non-Maximum Suppression (NMS)



**Problem:** Object detectors often output many overlapping detections

**Solution:** Post-process raw detections using **Non-Max Suppression (NMS):**

1. Select next highest-scoring box
2. Eliminate lower-scoring boxes with  $\text{IoU} > \text{threshold}$  (e.g. 0.7)
3. If any boxes remain, GOTO 1



[Puppy image is CC0 Public Domain]

# Duplicate Filter: Non-Maximum Suppression (NMS)



**Problem:** Object detectors often output many overlapping detections

**Solution:** Post-process raw detections using **Non-Max Suppression (NMS):**

1. Select next highest-scoring box
2. Eliminate lower-scoring boxes with  $\text{IoU} > \text{threshold}$  (e.g. 0.7)
3. If any boxes remain, GOTO 1



[Crowd image] is for commercial use under the [Pixabay license](#)

**Problem:** NMS may eliminate "good" boxes when objects are highly overlapping. There is not yet a good solution.

# Evaluating Object Detectors

## Mean Average Precision (mAP):

1. Run object detector on all test images (with NMS)
2. For each category, compute Average Precision (AP) =  
area under Precision vs Recall Curve
  1. For each detection (highest score to lowest score)

All dog detections sorted by score

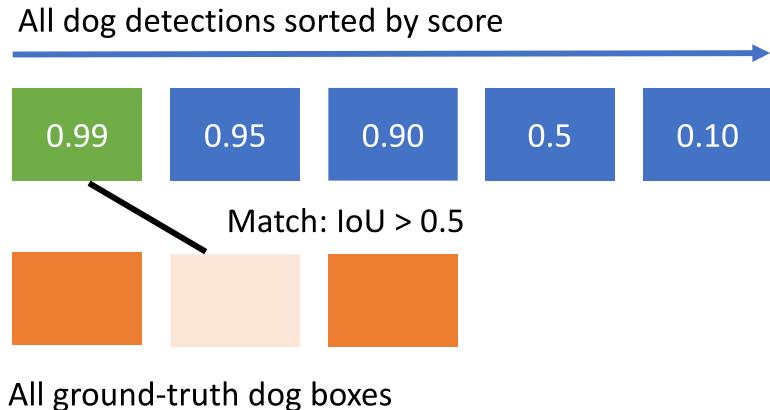


All ground-truth dog boxes

# Evaluating Object Detectors

## Mean Average Precision (mAP):

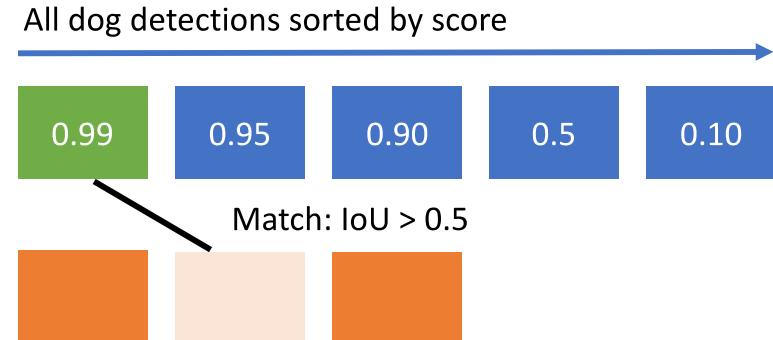
1. Run object detector on all test images (with NMS)
2. For each category, compute Average Precision (AP) = area under Precision vs Recall Curve
  1. For each detection (highest score to lowest score)
    1. If it matches some GT box with  $\text{IoU} > 0.5$ , mark it as positive and eliminate the GT
    2. Otherwise mark it as negative



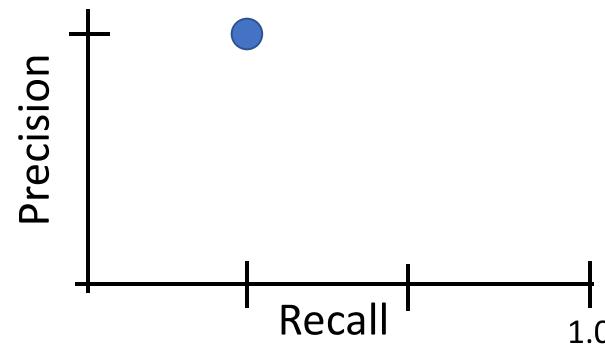
# Evaluating Object Detectors

## Mean Average Precision (mAP):

1. Run object detector on all test images (with NMS)
2. For each category, compute Average Precision (AP) = area under Precision vs Recall Curve
  1. For each detection (highest score to lowest score)
    1. If it matches some GT box with  $\text{IoU} > 0.5$ , mark it as positive and eliminate the GT
    2. Otherwise mark it as negative
    3. Plot a point on PR Curve



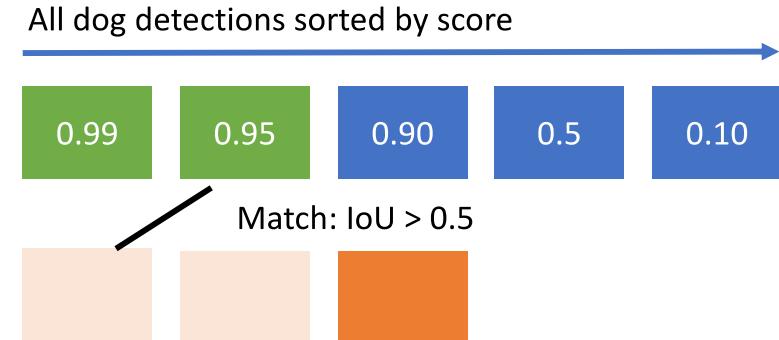
All ground-truth dog boxes



# Evaluating Object Detectors

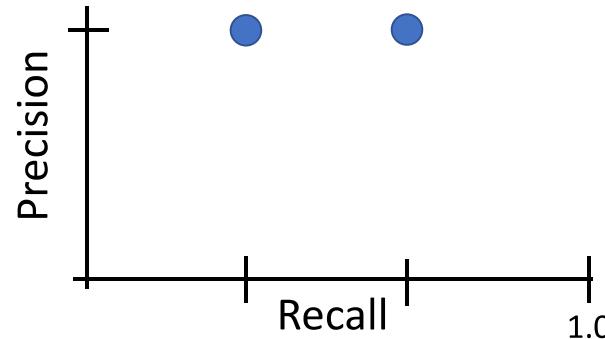
## Mean Average Precision (mAP):

1. Run object detector on all test images (with NMS)
2. For each category, compute Average Precision (AP) = area under Precision vs Recall Curve
  1. For each detection (highest score to lowest score)
    1. If it matches some GT box with  $\text{IoU} > 0.5$ , mark it as positive and eliminate the GT
    2. Otherwise mark it as negative
    3. Plot a point on PR Curve



All ground-truth dog boxes

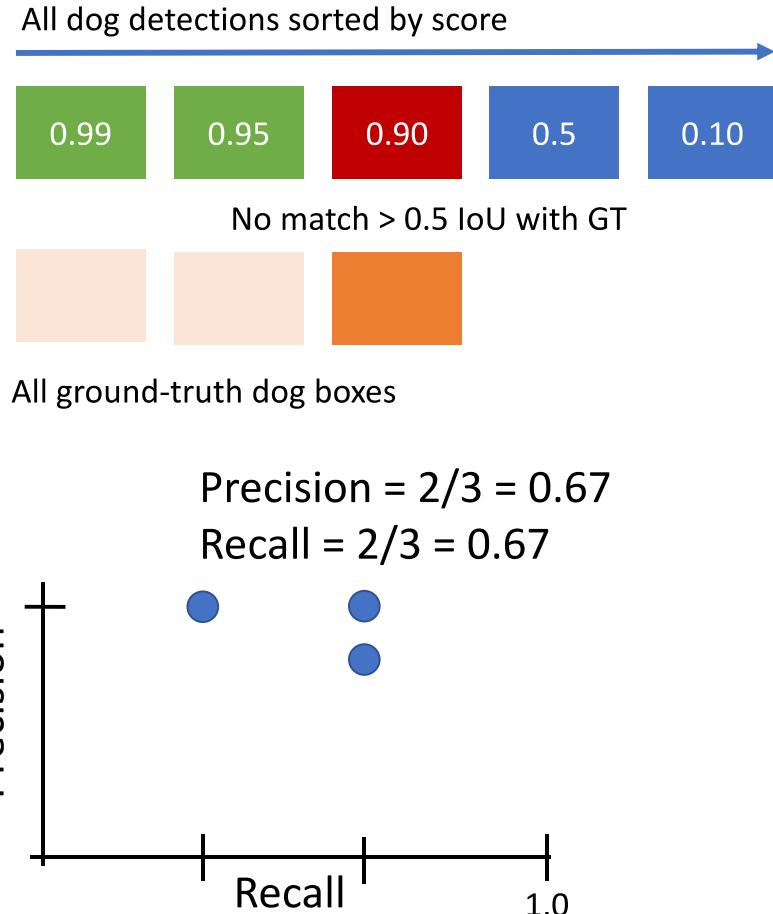
$$\begin{aligned}\text{Precision} &= 2/2 = 1.0 \\ \text{Recall} &= 2/3 = 0.67\end{aligned}$$



# Evaluating Object Detectors

## Mean Average Precision (mAP):

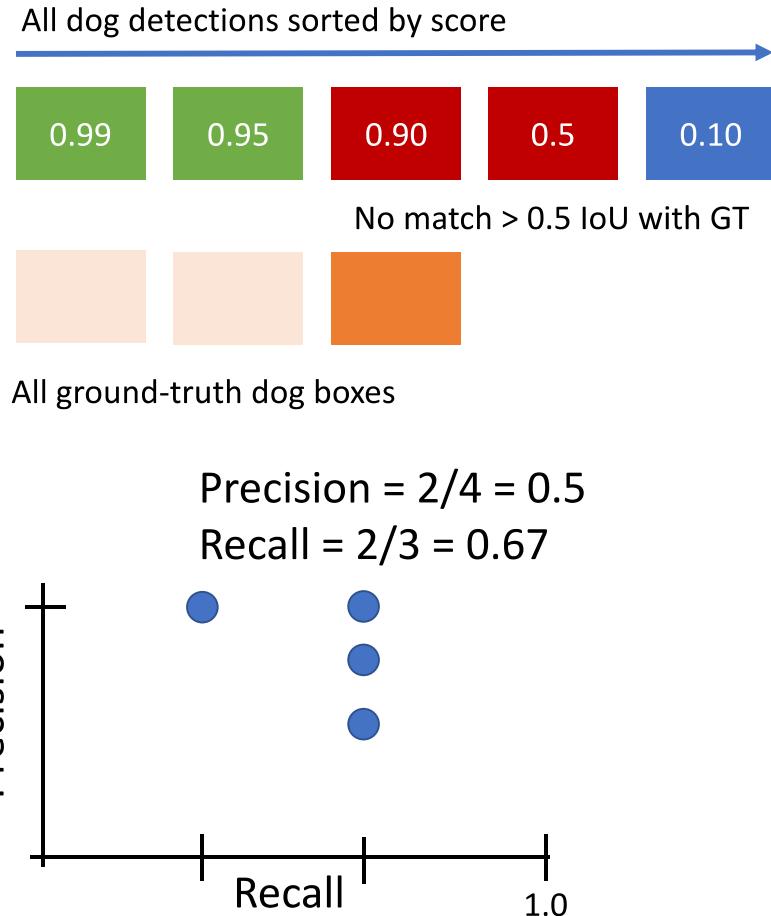
1. Run object detector on all test images (with NMS)
2. For each category, compute Average Precision (AP) = area under Precision vs Recall Curve
  1. For each detection (highest score to lowest score)
    1. If it matches some GT box with IoU > 0.5, mark it as positive and eliminate the GT
    2. Otherwise mark it as negative
    3. Plot a point on PR Curve



# Evaluating Object Detectors

## Mean Average Precision (mAP):

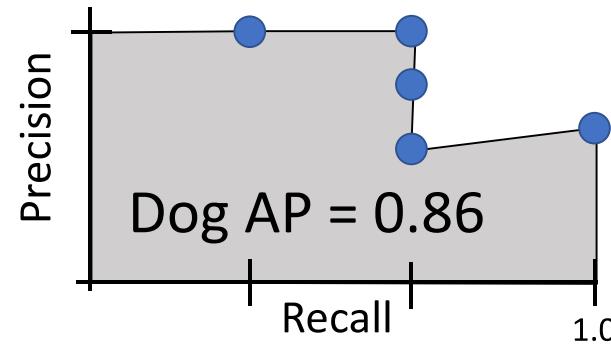
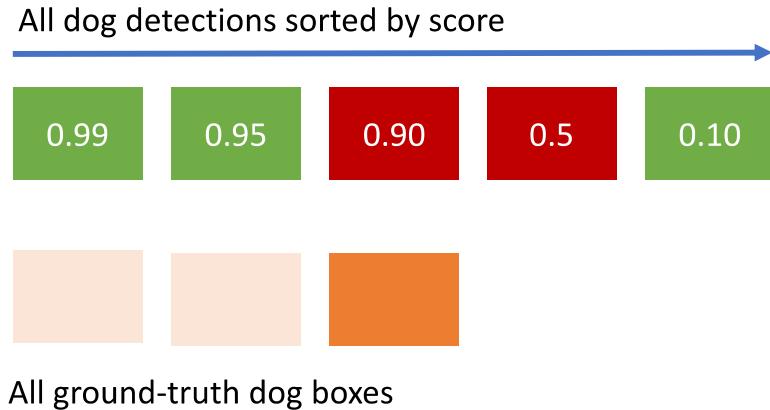
1. Run object detector on all test images (with NMS)
2. For each category, compute Average Precision (AP) = area under Precision vs Recall Curve
  1. For each detection (highest score to lowest score)
    1. If it matches some GT box with IoU > 0.5, mark it as positive and eliminate the GT
    2. Otherwise mark it as negative
    3. Plot a point on PR Curve



# Evaluating Object Detectors

## Mean Average Precision (mAP):

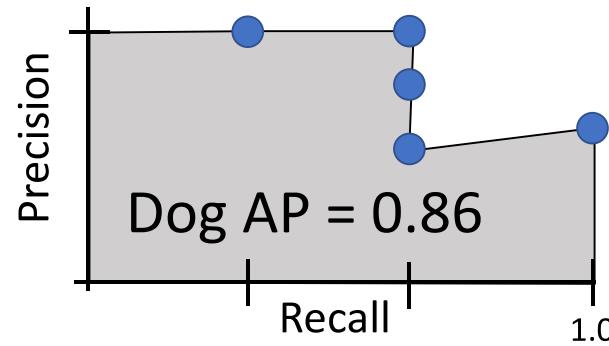
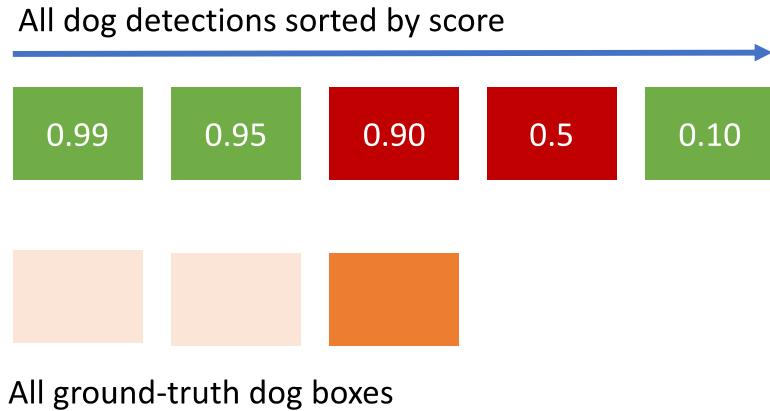
1. Run object detector on all test images (with NMS)
2. For each category, compute Average Precision (AP) = area under Precision vs Recall Curve
  1. For each detection (highest score to lowest score)
    1. If it matches some GT box with  $\text{IoU} > 0.5$ , mark it as positive and eliminate the GT
    2. Otherwise mark it as negative
    3. Plot a point on PR Curve



# Evaluating Object Detectors

## Mean Average Precision (mAP):

1. Run object detector on all test images (with NMS)
2. For each category, compute Average Precision (AP) = area under Precision vs Recall Curve
  1. For each detection (highest score to lowest score)
    1. If it matches some GT box with  $\text{IoU} > 0.5$ , mark it as positive and eliminate the GT
    2. Otherwise mark it as negative
    3. Plot a point on PR Curve



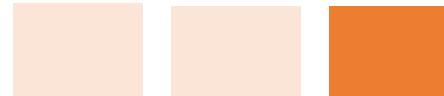
# Evaluating Object Detectors

## Mean Average Precision (mAP):

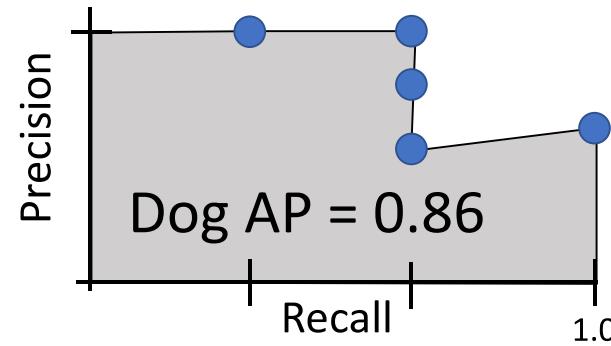
1. Run object detector on all test images (with NMS)
2. For each category, compute Average Precision (AP) = area under Precision vs Recall Curve
  1. For each detection (highest score to lowest score)
    1. If it matches some GT box with  $\text{IoU} > 0.5$ , mark it as positive and eliminate the GT
    2. Otherwise mark it as negative
    3. Plot a point on PR Curve

**How to get AP = 1.0: Hit all GT boxes with  $\text{IoU} > 0.5$ , and have no “false positive” detections ranked above any “true positives”**

All dog detections sorted by score



All ground-truth dog boxes



# Evaluating Object Detectors

## Mean Average Precision (mAP):

1. Run object detector on all test images (with NMS)
2. For each category, compute Average Precision (AP) = area under Precision vs Recall Curve
  1. For each detection (highest score to lowest score)
    1. If it matches some GT box with  $\text{IoU} > 0.5$ , mark it as positive and eliminate the GT
    2. Otherwise mark it as negative
    3. Plot a point on PR Curve
3. Mean Average Precision (mAP) = average of AP for each category

Car AP = 0.65

Cat AP = 0.80

Dog AP = 0.86

mAP@0.5 = 0.77

# Evaluating Object Detectors

## Mean Average Precision (mAP):

1. Run object detector on all test images (with NMS)
2. For each category, compute Average Precision (AP) = area under Precision vs Recall Curve
  1. For each detection (highest score to lowest score)
    1. If it matches some GT box with  $\text{IoU} > 0.5$ , mark it as positive and eliminate the GT
    2. Otherwise mark it as negative
    3. Plot a point on PR Curve
3. Mean Average Precision (mAP) = average of AP for each category
4. For “COCO mAP”: Compute mAP@thresh for each IoU threshold (0.5, 0.55, 0.6, ..., 0.95) and take average

$\text{mAP}@0.5 = 0.77$

$\text{mAP}@0.55 = 0.71$

$\text{mAP}@0.60 = 0.65$

...

$\text{mAP}@0.95 = 0.2$

COCO mAP = 0.4

# Today's Agenda

---

- Image Retrieval
- Object Detection
- Image Segmentation

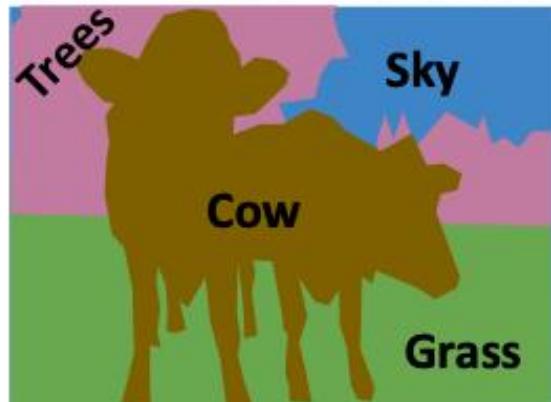
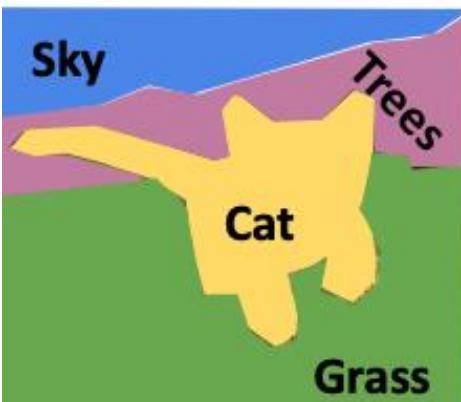
# Image Segmentation

**Goal:** Partition the image into semantically-meaningful or perceptually-similar regions.  
(In contrast to object detection: all pixels are labeled!)

**That is:** Label each pixel in the image with a category label



[This image is CC0 public domain](#)

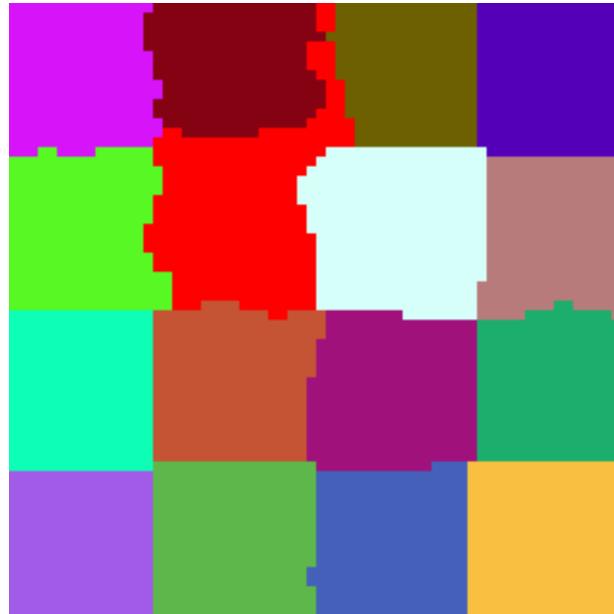
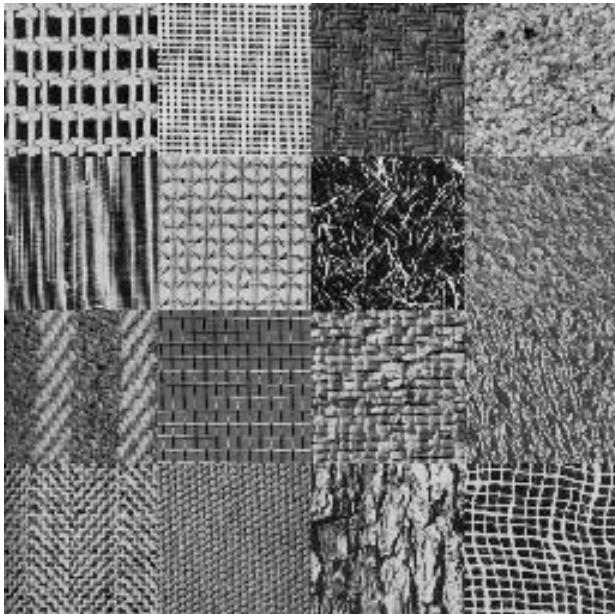


# Image Segmentation

Example: Color space clustering to reduce color palette.



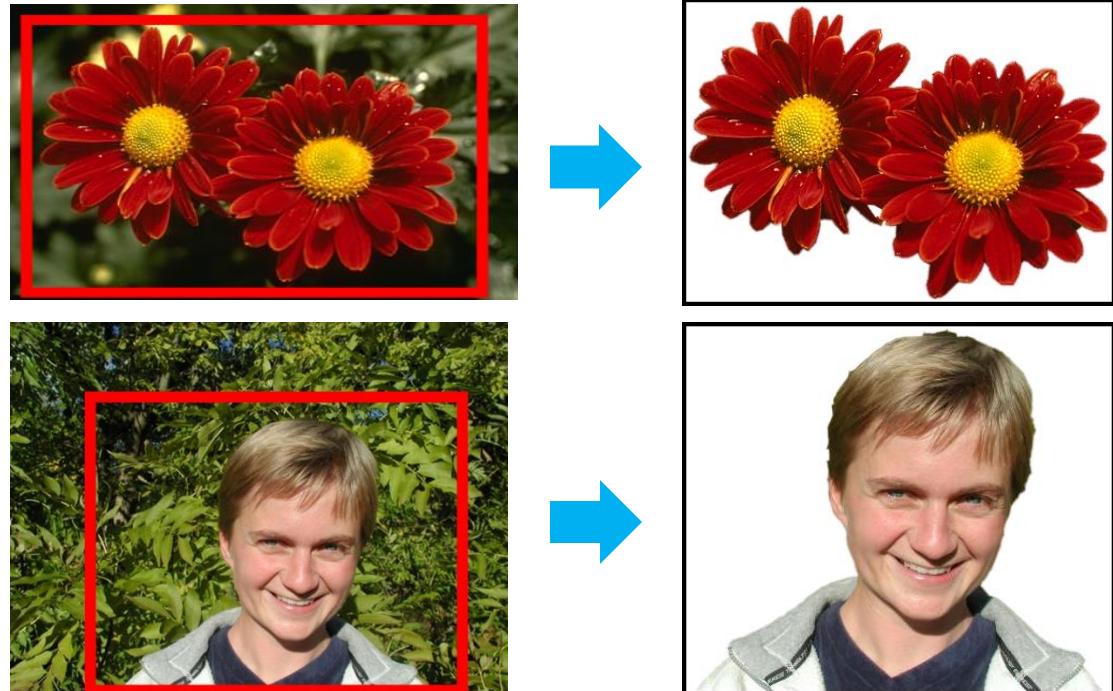
# Texture Segmentation



Courtesy Uni Bonn

# Image Segmentation

- “Grabcut”  
Graphcut-based  
Image segmentation
- Integrated into MS Office  
e.g. Powerpoint



# Image Segmentation



(a)



(b)

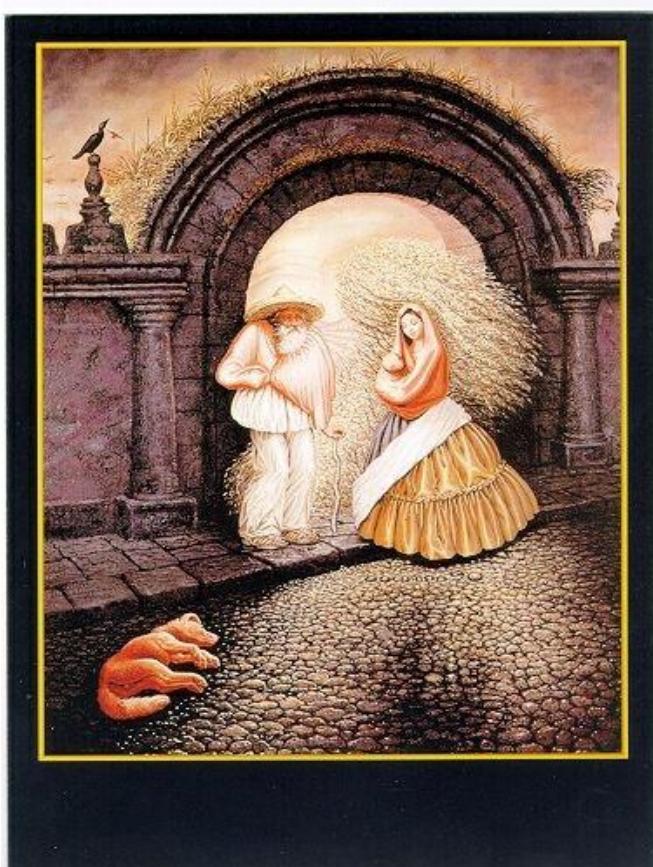


(c)



# Segment first, or recognise first? (what is the context?)

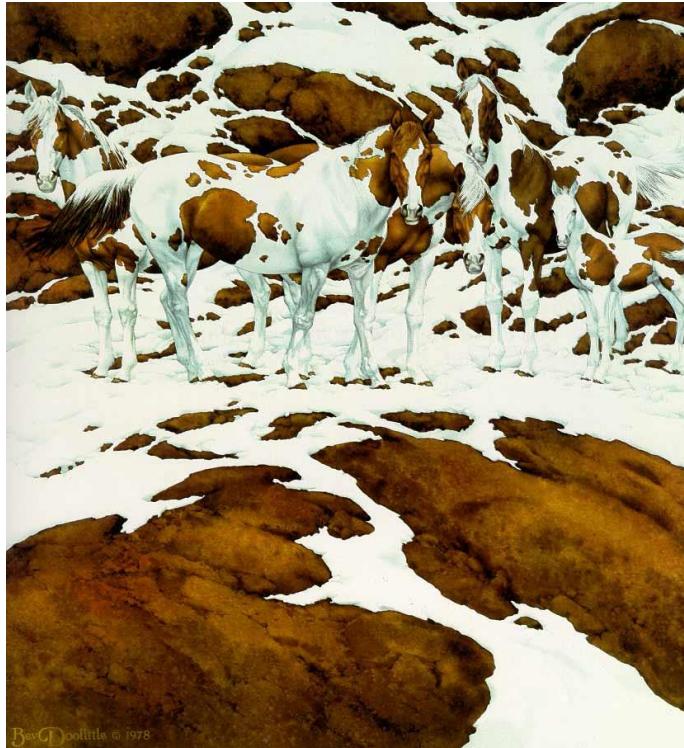
A portrait of  
an old man, or



A couple with  
a dog on the street

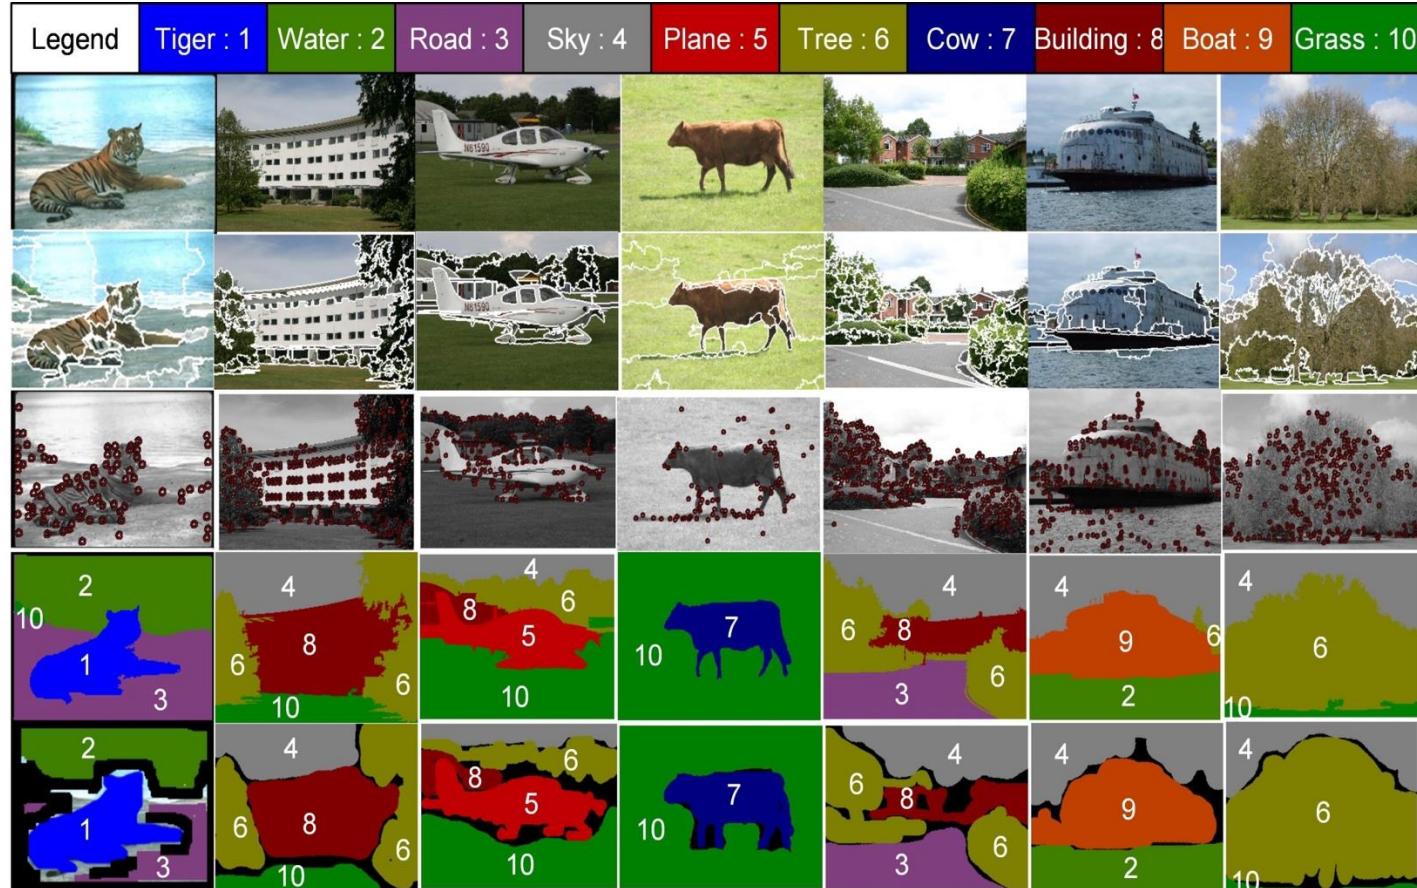
# Segment first, or recognise first? (what is the context?)

Snow and rocks,  
or



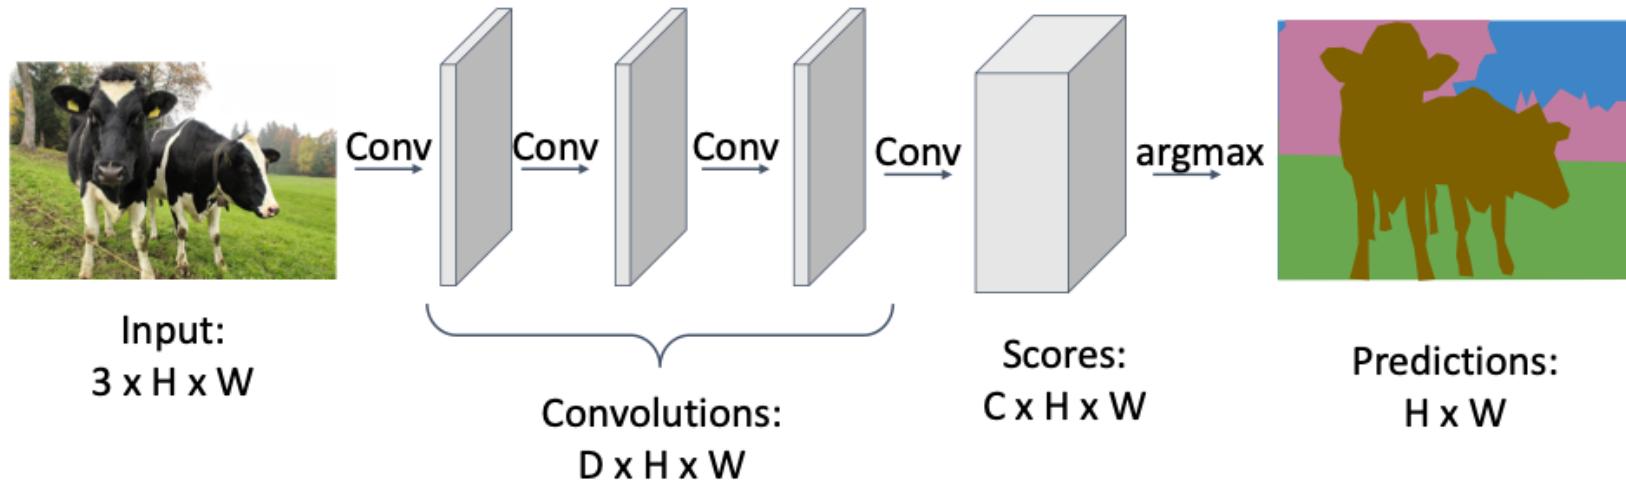
A group of horses  
In snow and rocks

# Segmentation for Recognition



# Image Segmentation

Design a network with multiple convolutional layers to make predictions for pixels all at once!



Loss function: Per-Pixel cross-entropy

# Today's Agenda

---

- Image Retrieval
- Object Detection
- **Image Segmentation**
  - K-Means

# Segmentation as Clustering

---

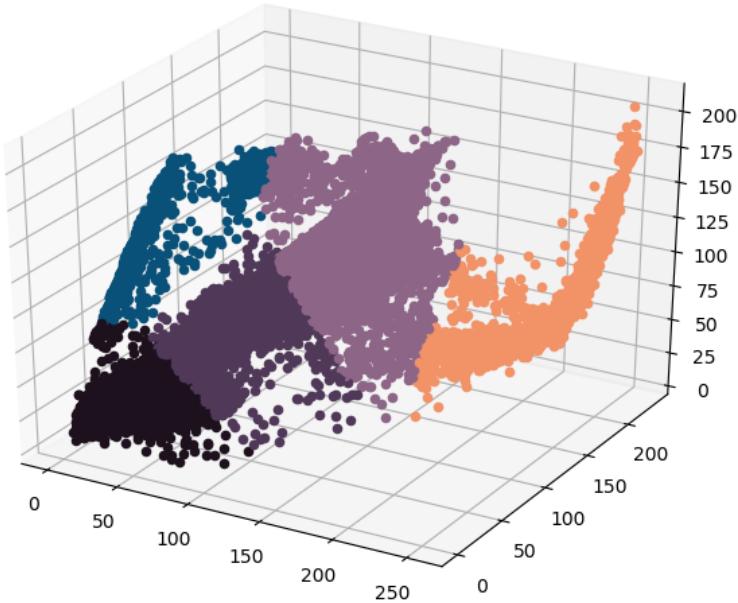
- Low- and mid-level visual cues for segmentation.
- Partition the image pixels into groups with perceptually homogeneous or similar pixel properties (e.g., intensity, colour, texture, or similar spatial location).
- Grouping pixels into perceptually homogeneous regions.

# Segmentation as Clustering

- Cluster similar pixels using color features



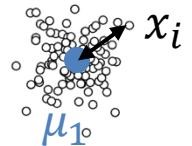
Cluster in color space (3D) or in color + position space (5D)



# K-Means

**Goal:** Minimize the distance between  $N$  given data points and  $K$  assigned cluster centers:

$$\text{minimize} \quad \sum_{k=1}^K \sum_{x_i \in C_k} \|x_i - \mu_k\|^2$$



## K-Means Algorithm

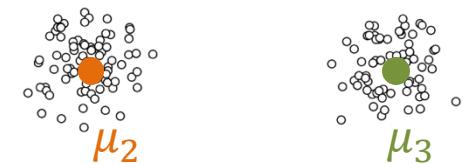
1. **Initialize**  $K$  cluster means  $\mu_1, \dots, \mu_K$
2. **Reassign Points:** For  $i = 1, \dots, N$  assign each point  $x_n$  to the closest cluster

$$z_n = \operatorname{argmin}_{k \in \{1, \dots, K\}} \|x_i - \mu_k\|^2$$

3. **Update Centroids:** Suppose  $C_k = \{x_i : z_i = k\}$ . Recompute the means as

$$\mu_k = \operatorname{mean}(C_k), \quad k = 1, \dots, K$$

4. **Repeat:** Go to step 2 if not yet converged

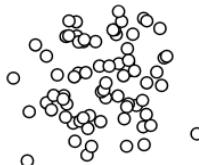
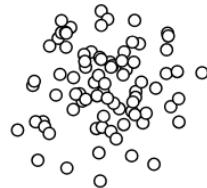
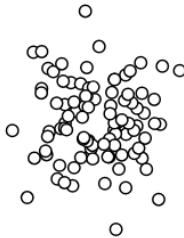


# K-Means

Next step:

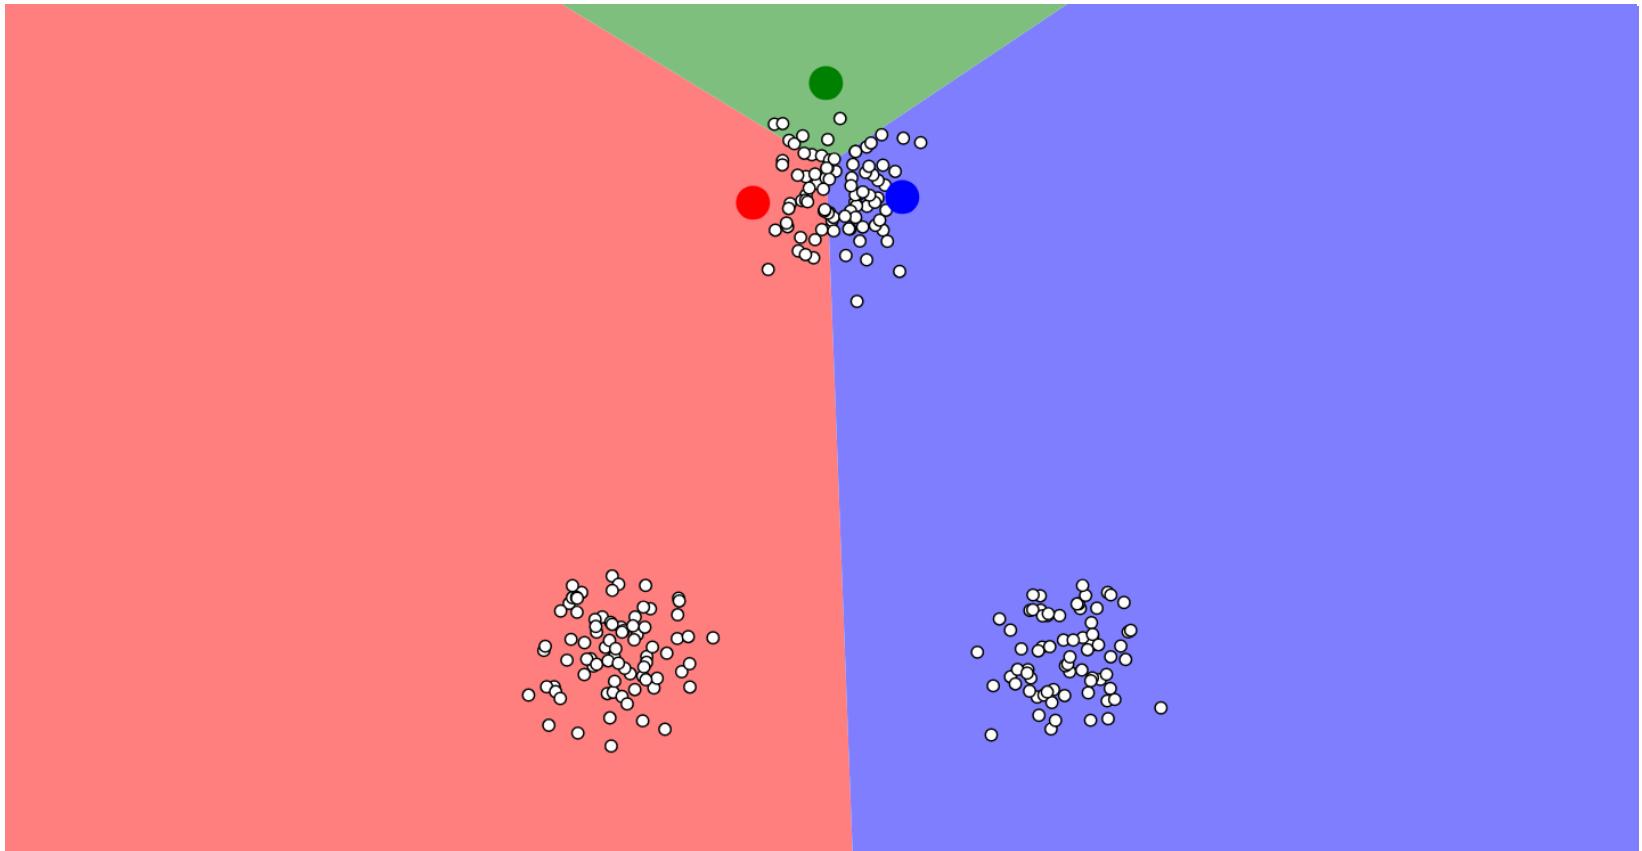
**Initialize**

**Means**



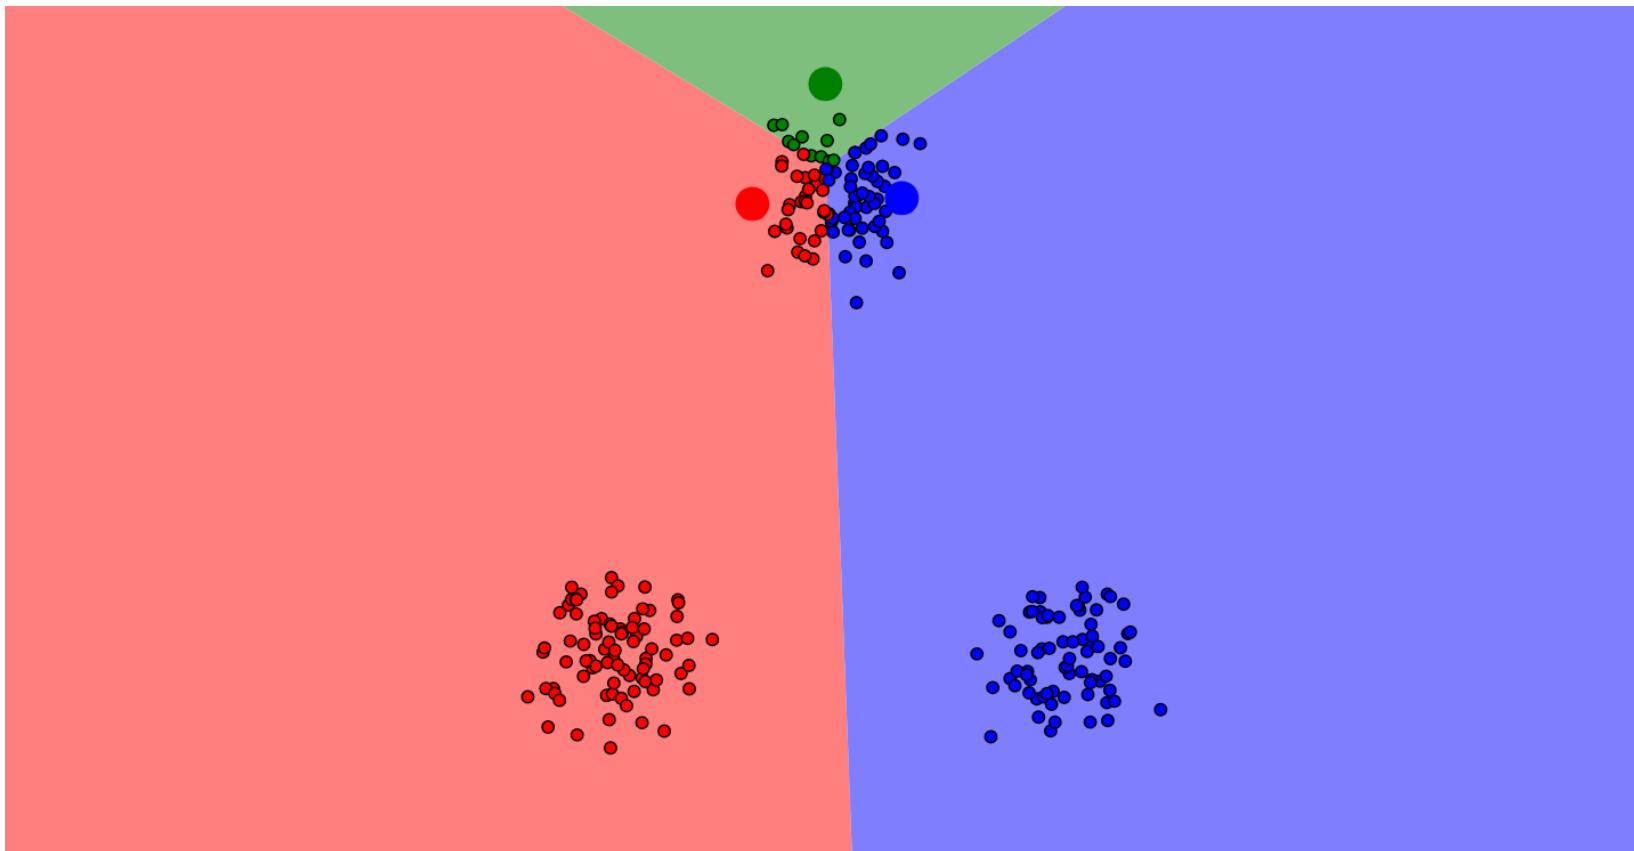
# K-Means

Next step:  
**Reassign  
Points**



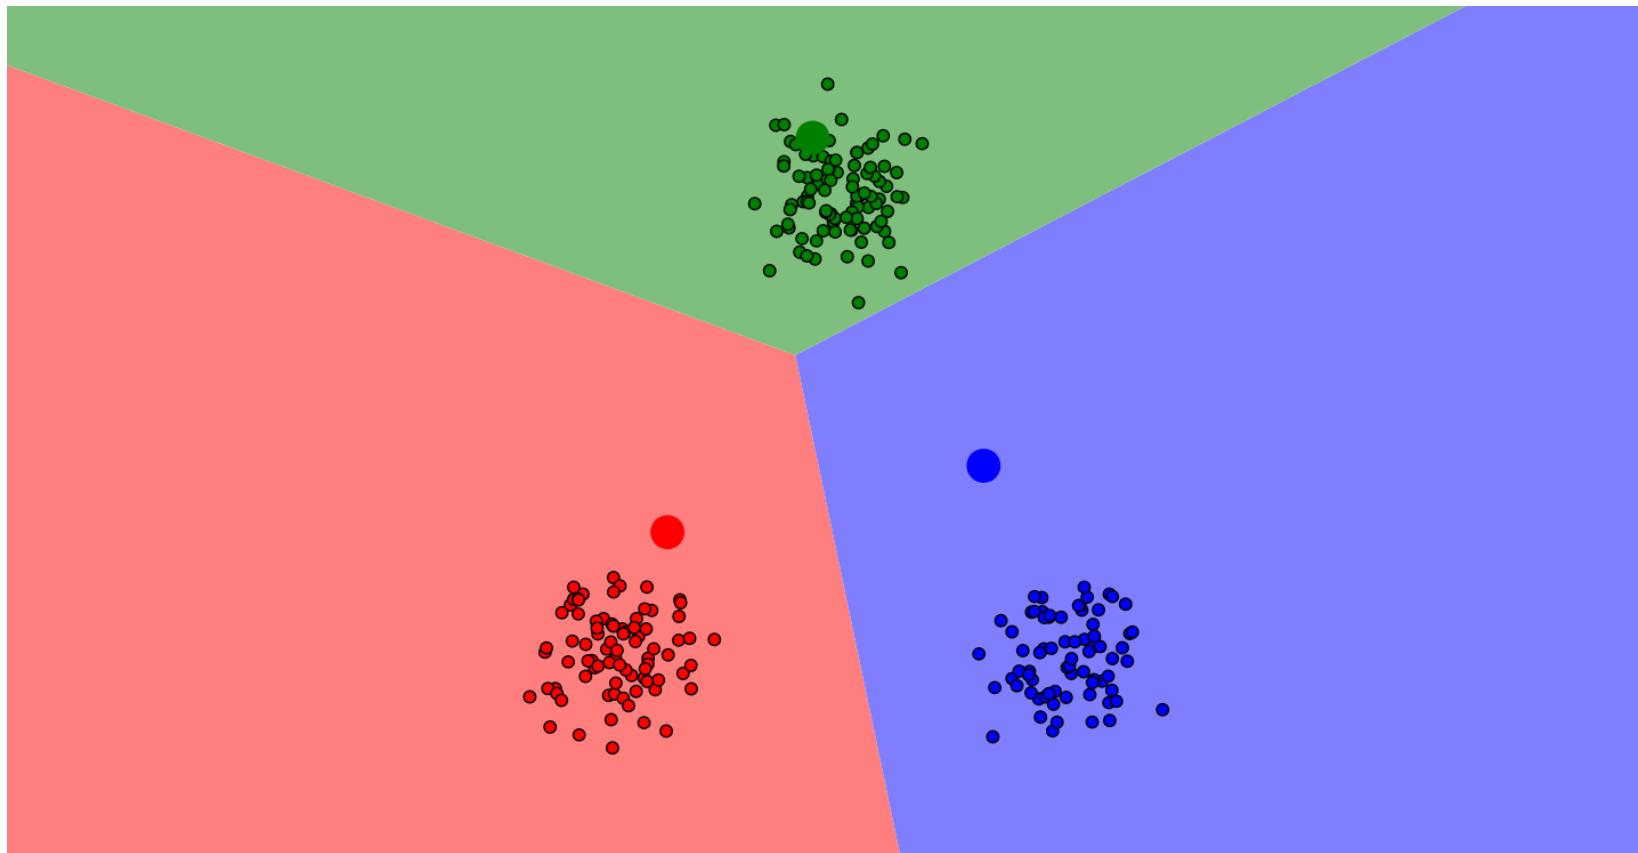
# K-Means

Next step:  
Update  
Centroids



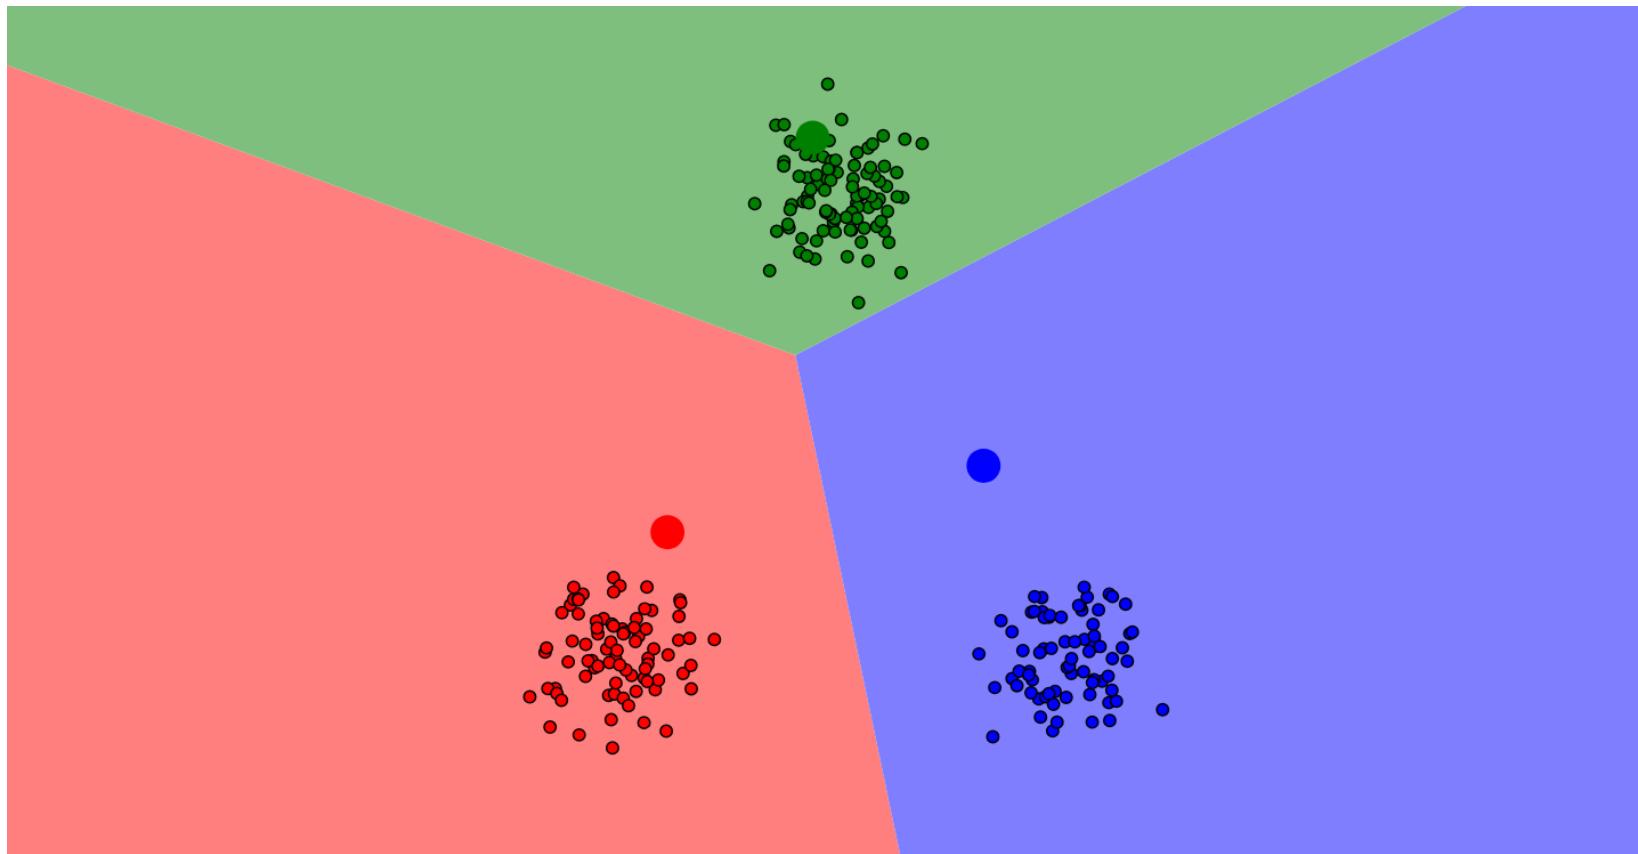
# K-Means

Next step:  
Reassign  
Points



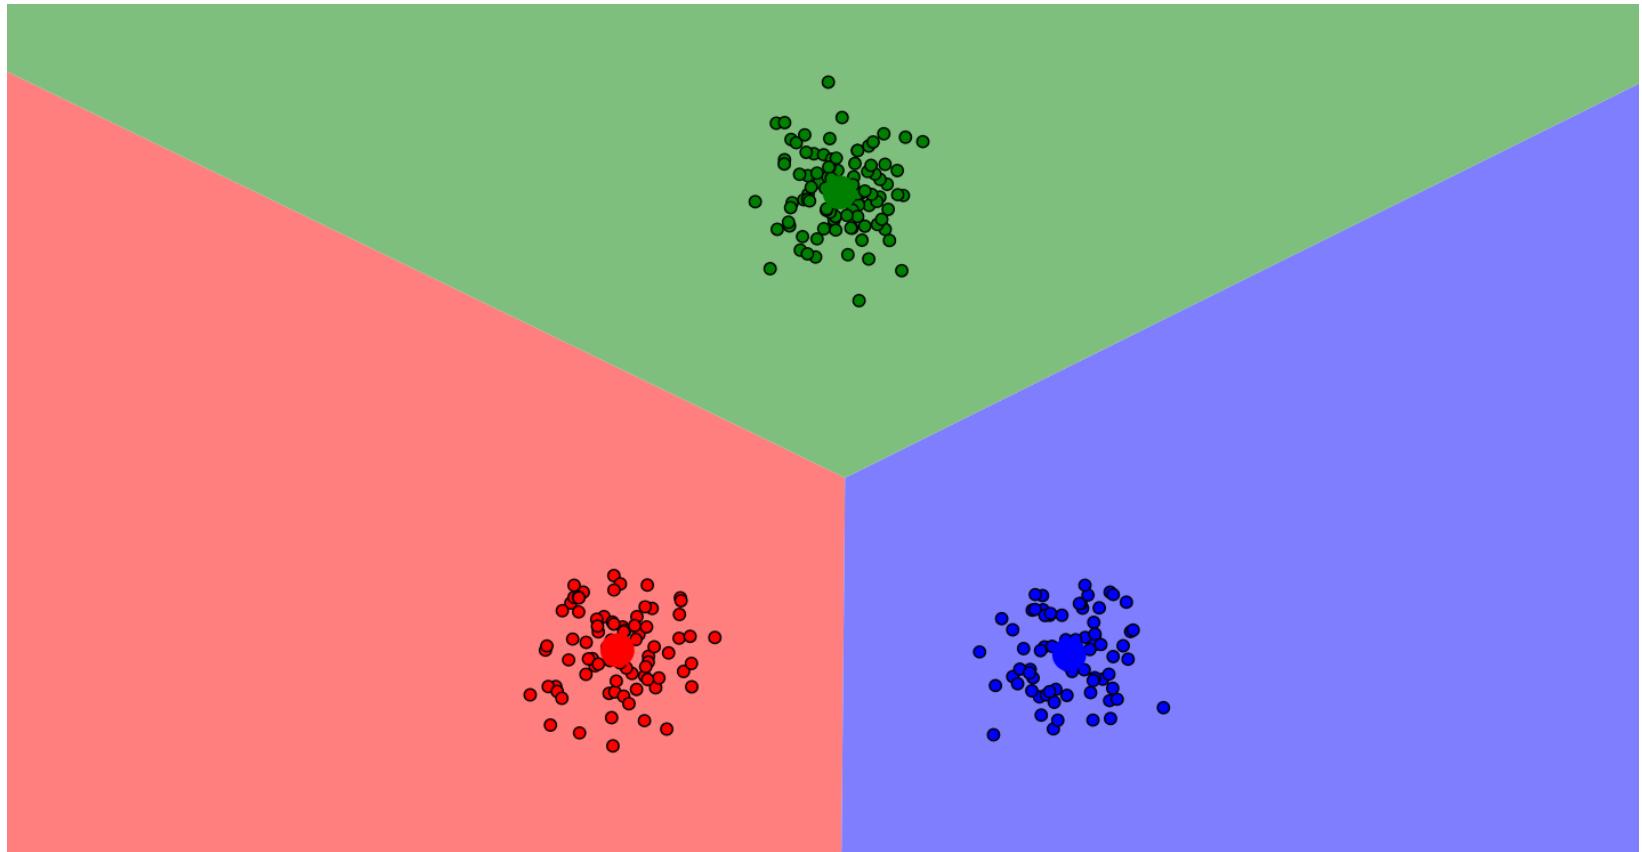
# K-Means

Next step:  
Update  
Centroids



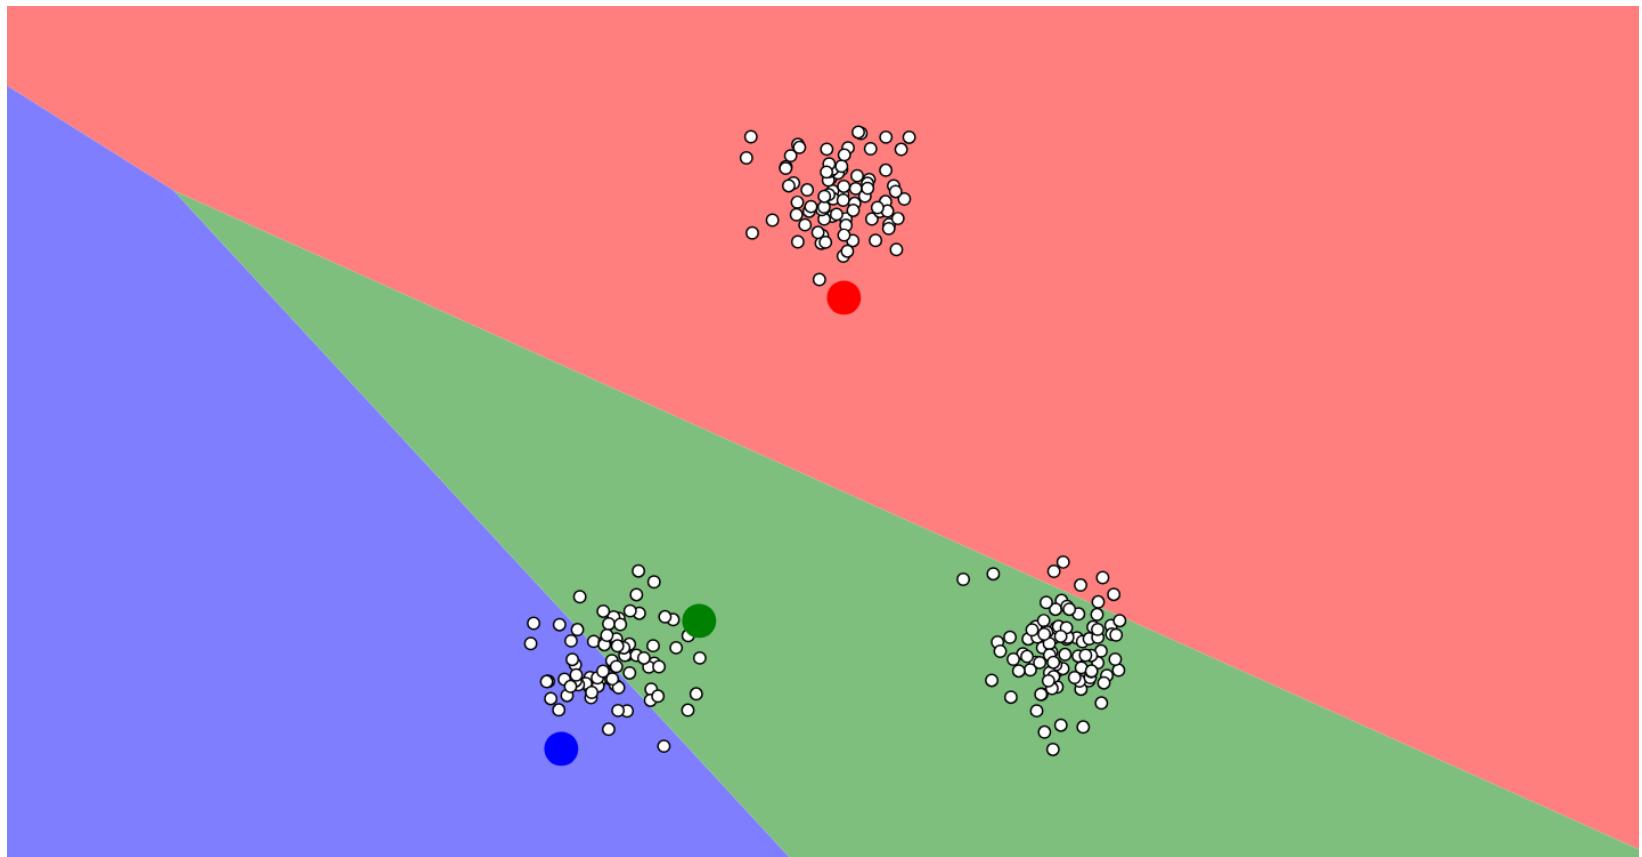
# K-Means

Next step:  
End



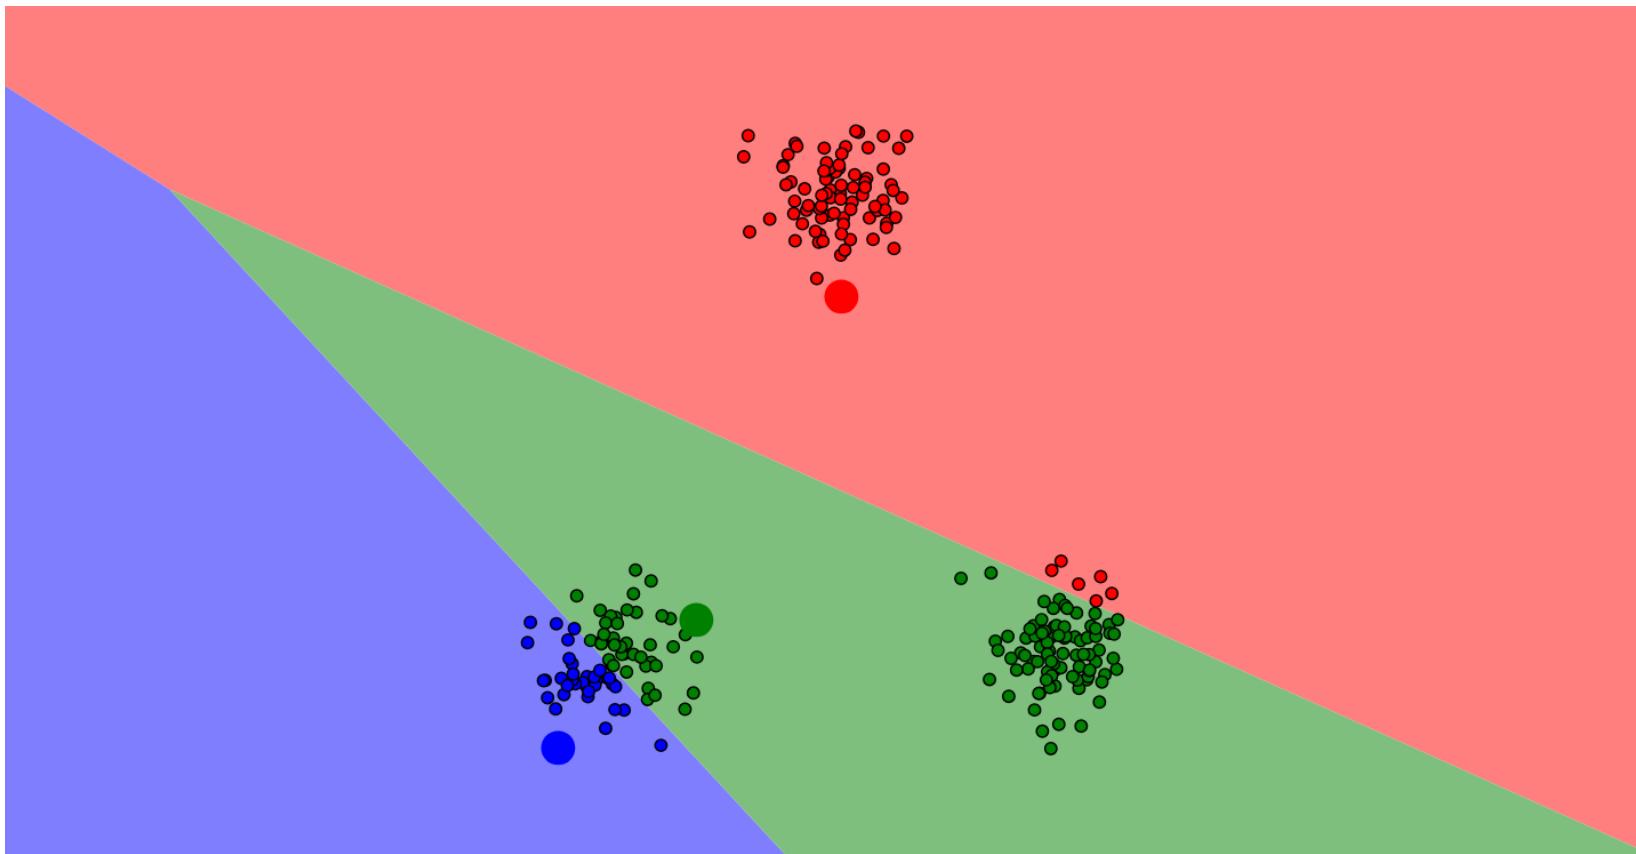
# K-Means

Next step:  
Reassign  
Points



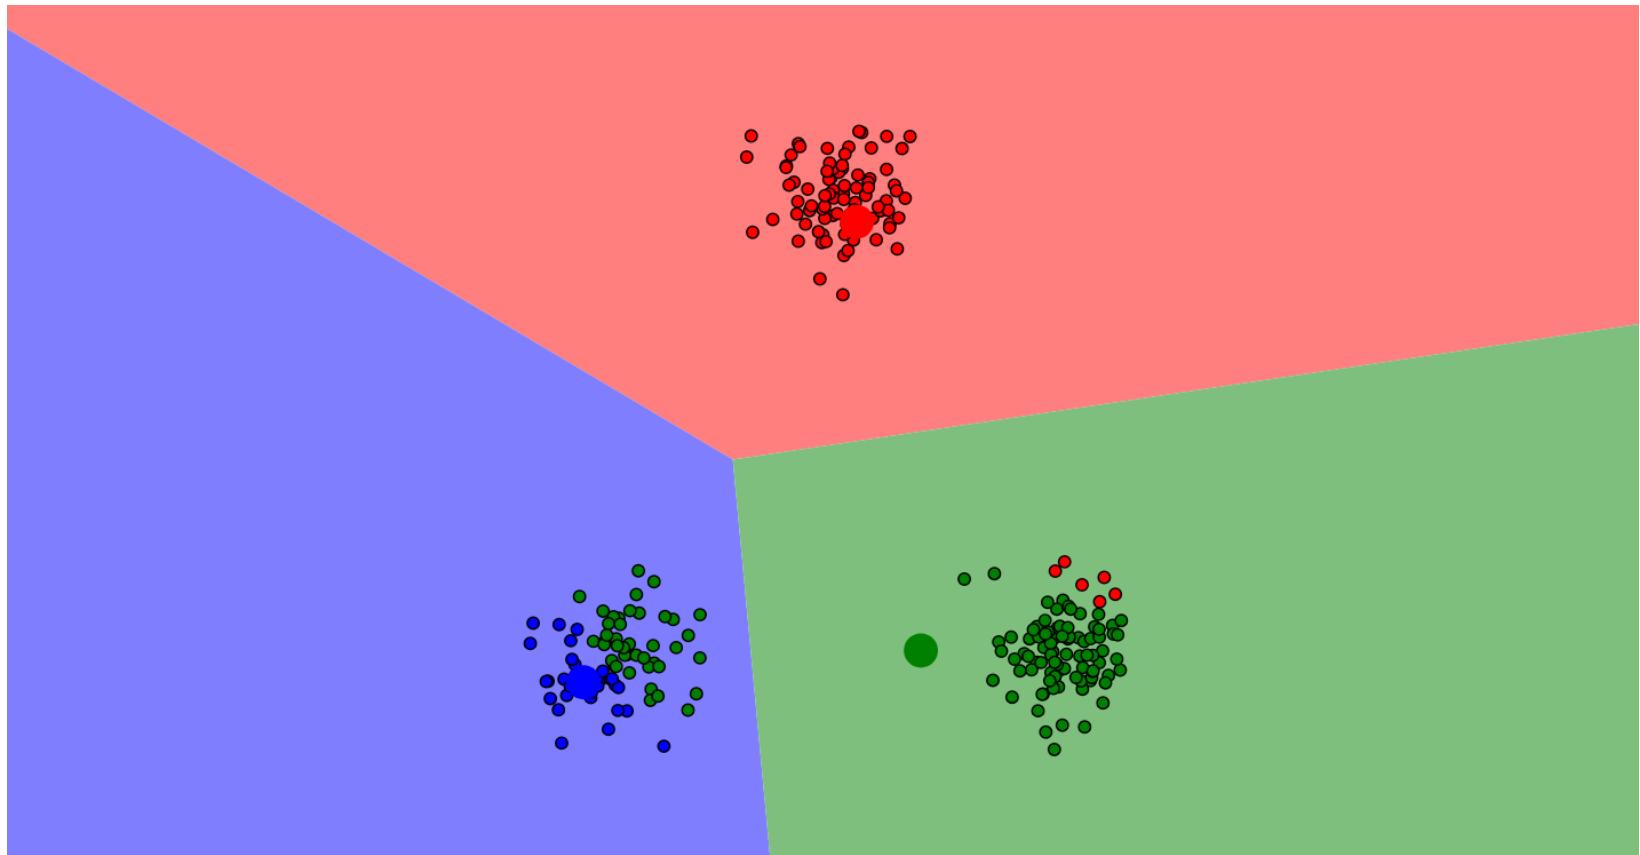
# K-Means

Next step:  
Update  
Centroids



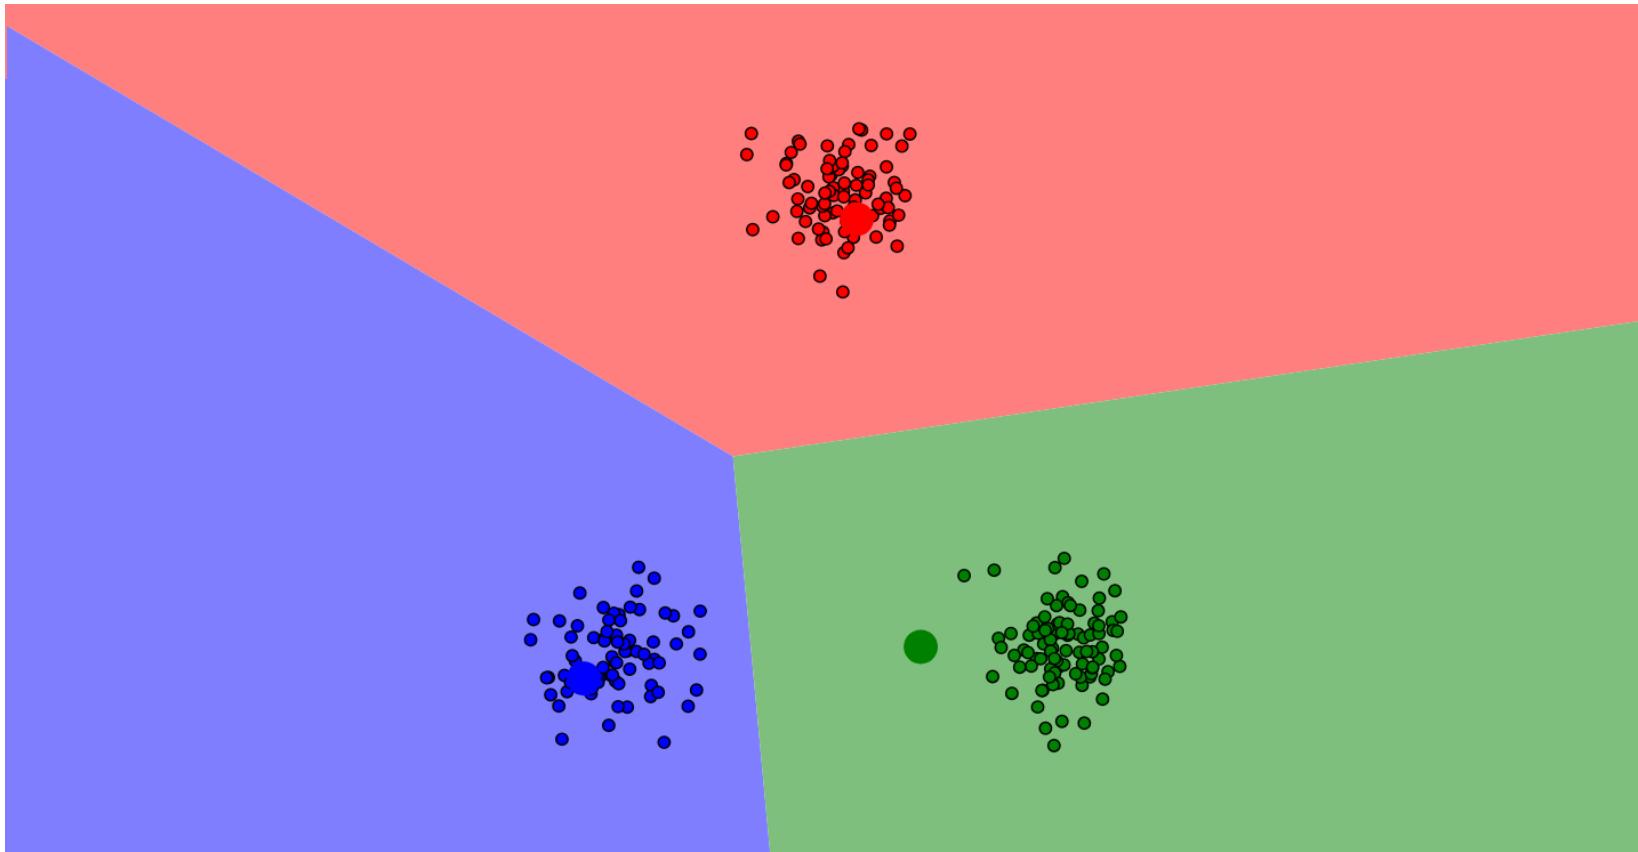
# K-Means

Next step:  
Reassign  
Points



# K-Means

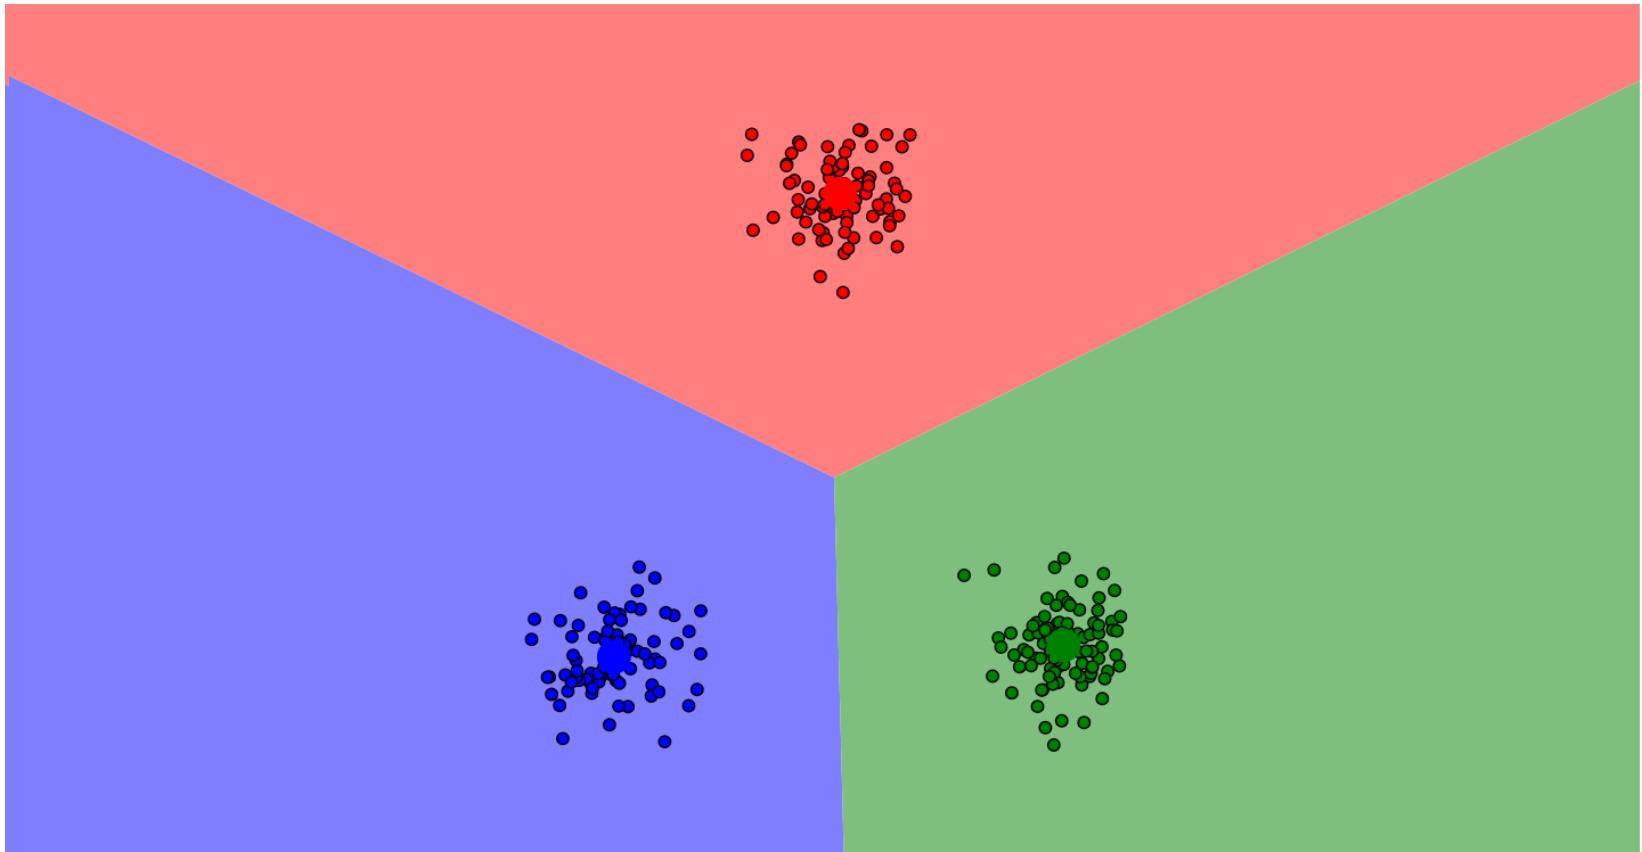
Next step:  
Update  
Centroids



# K-Means

Next step:

End

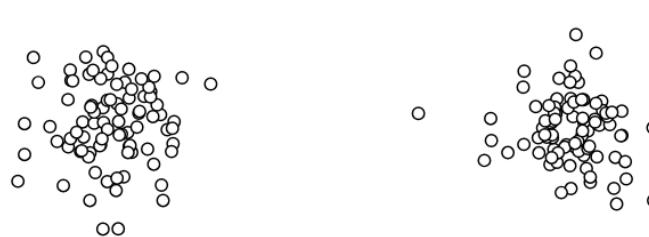
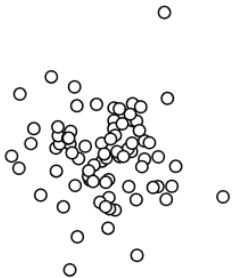


# K-Means

Next step:

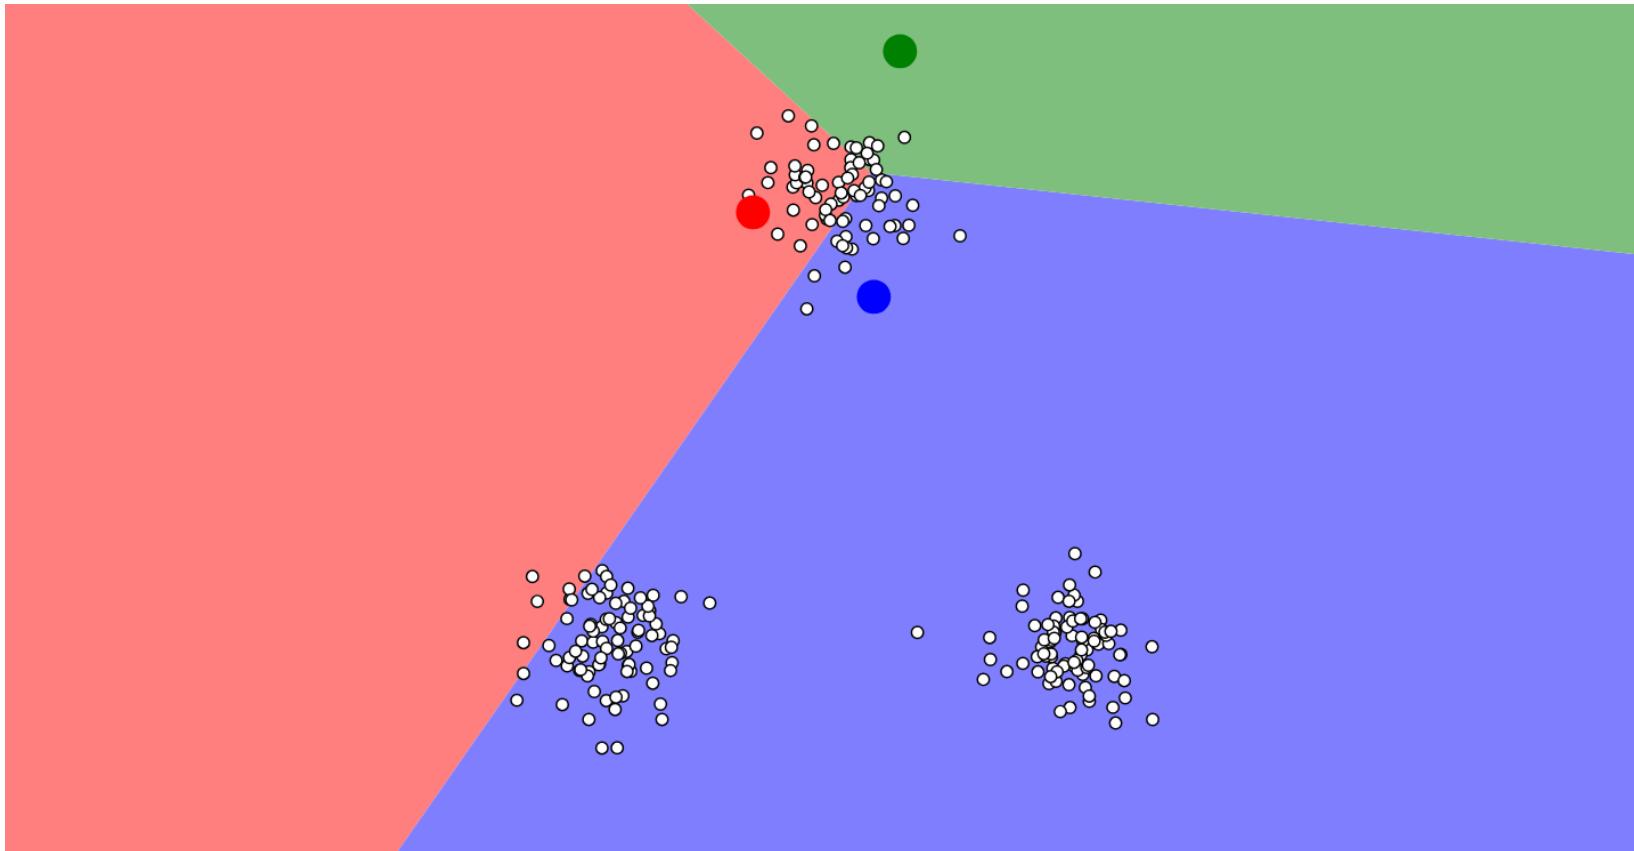
**Initialize**

**Means**



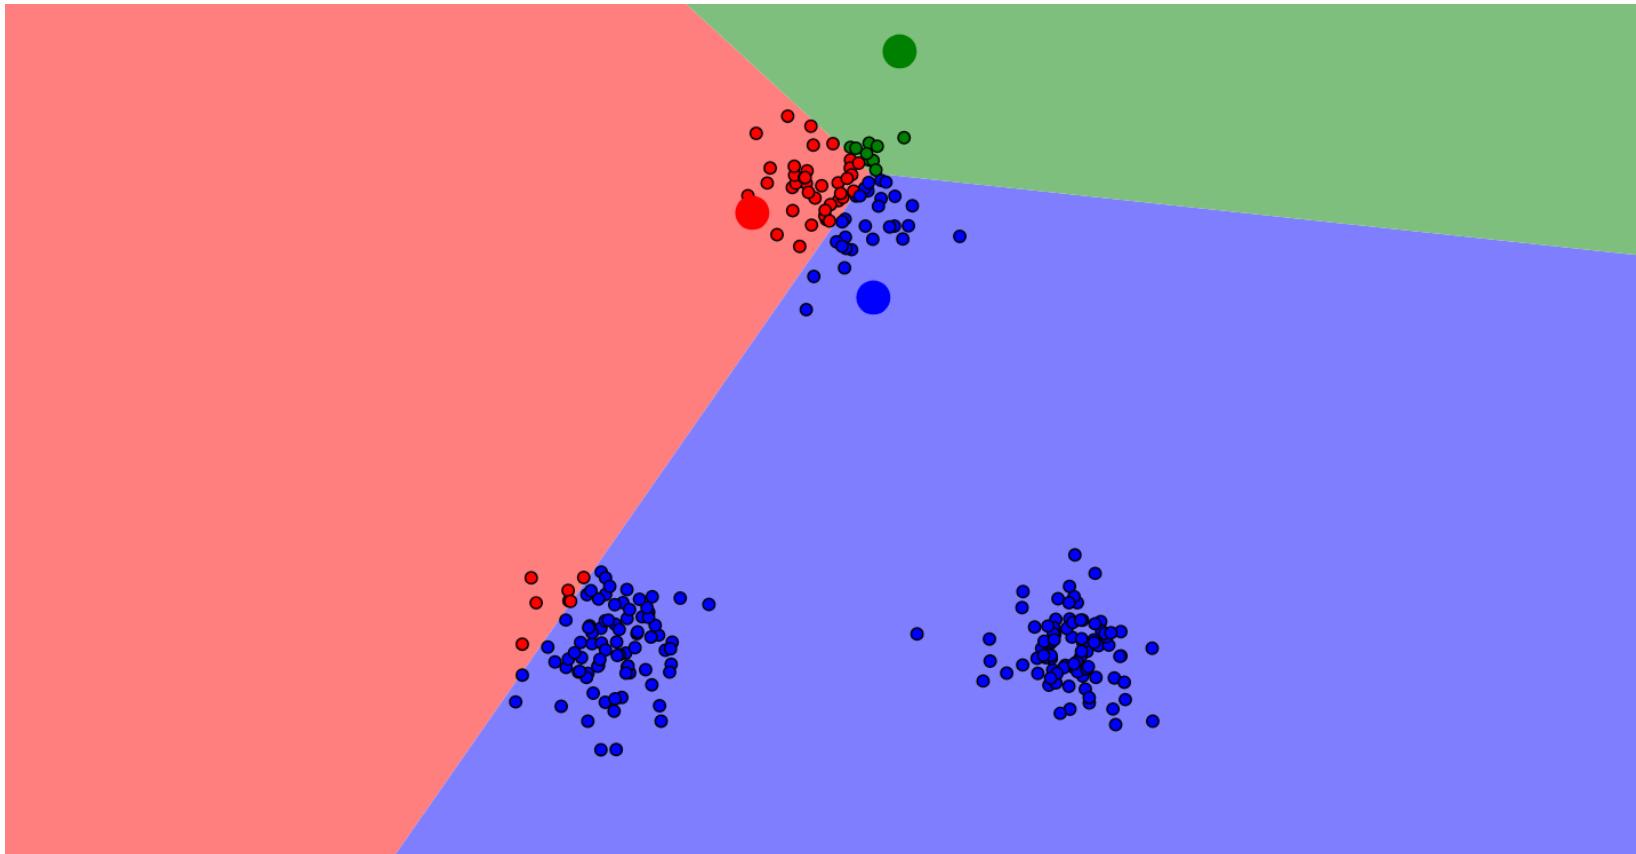
# K-Means

Next step:  
Reassign  
Points



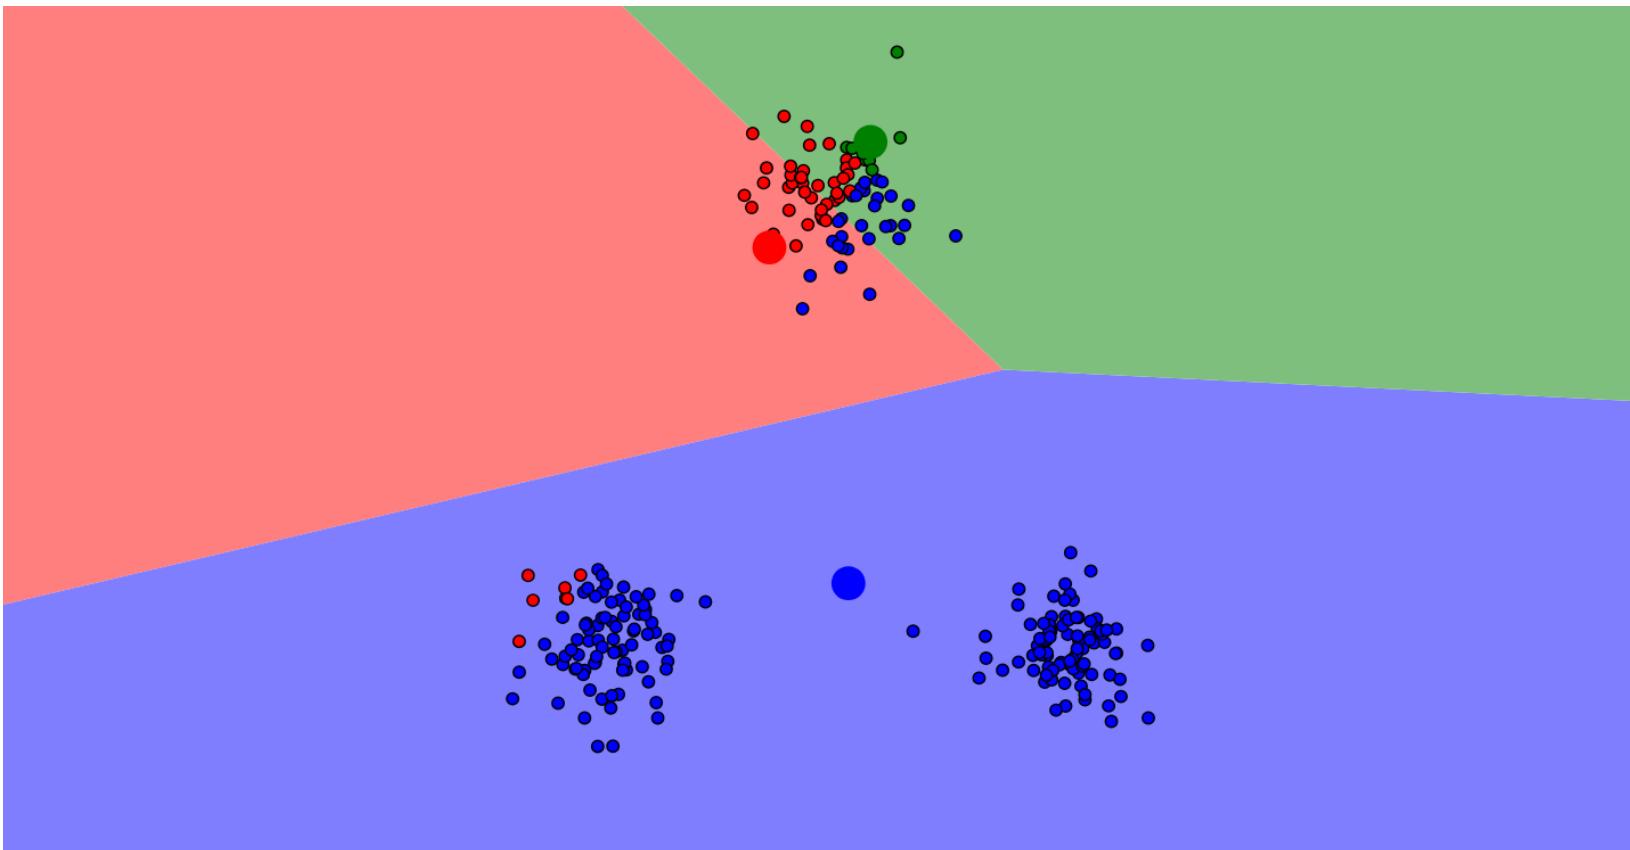
# K-Means

Next step:  
Update  
Centroids



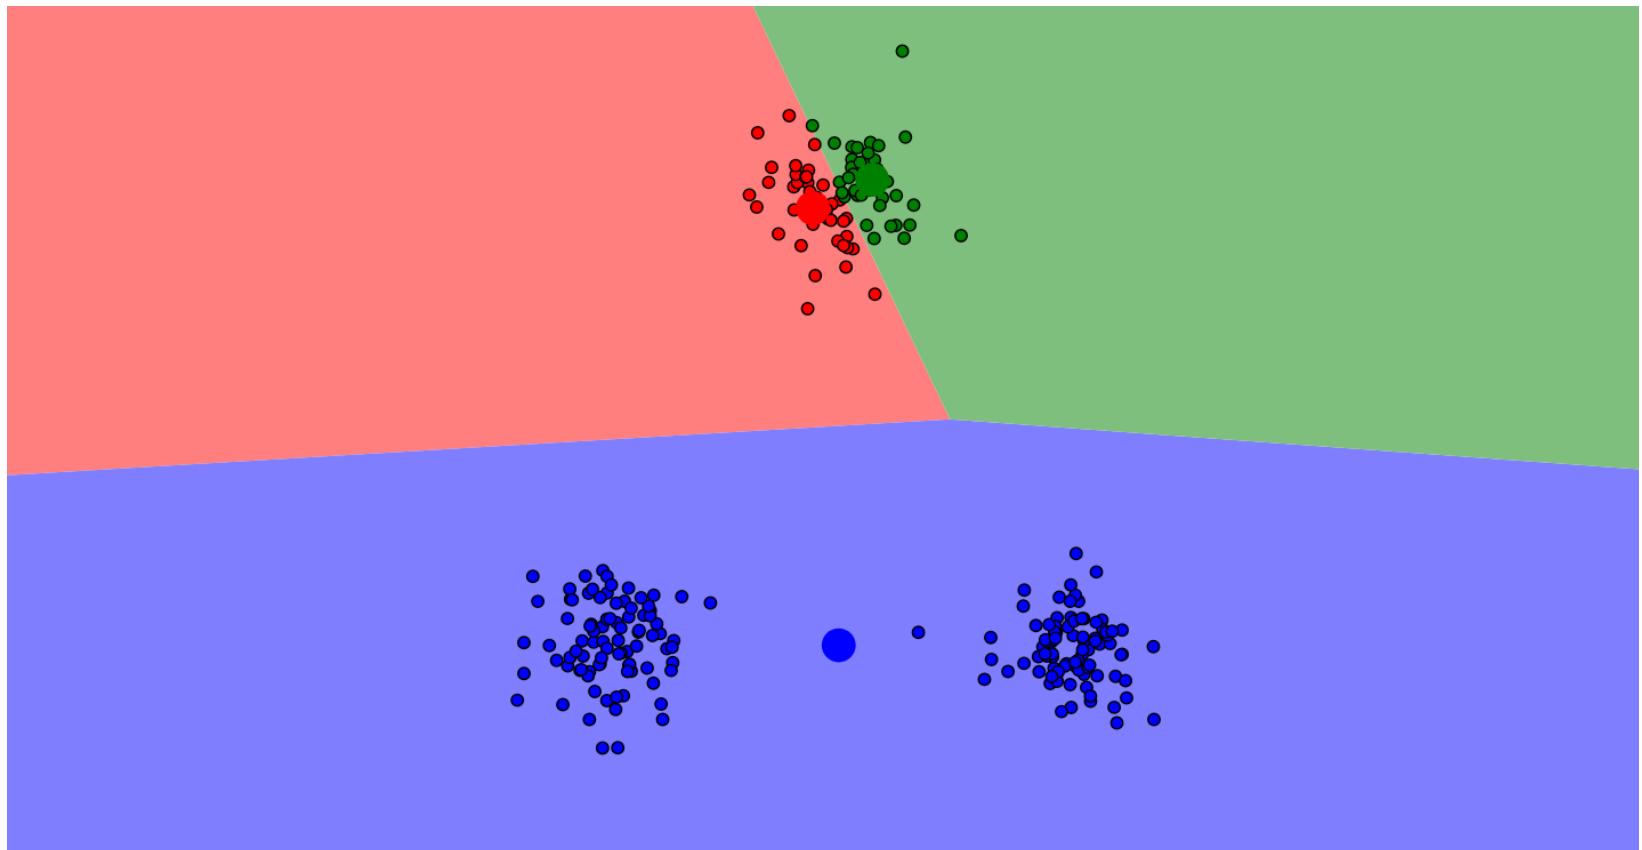
# K-Means

Next step:  
Reassign  
Points



# K-Means

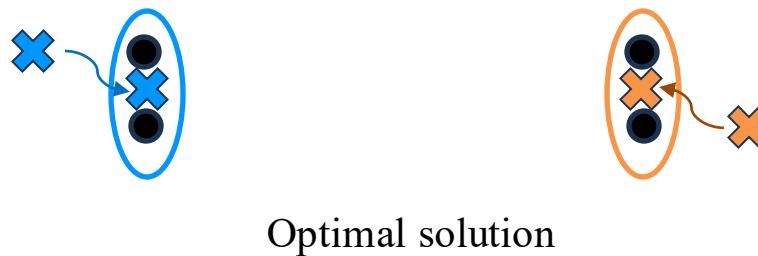
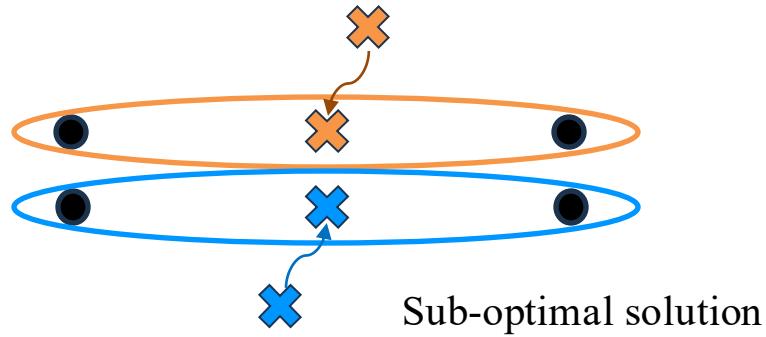
Next step:  
End



# K-Means

## Local Minima

- Sensitive to initialization  
→ sub-optimal solution



# K-Means

## Pros:

- Finds cluster centers that minimize variance (good representation of data)
- Simple to implement, widespread applications.

## Cons:

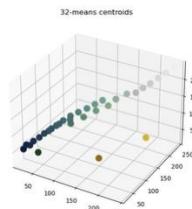
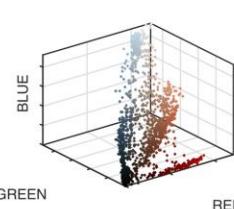
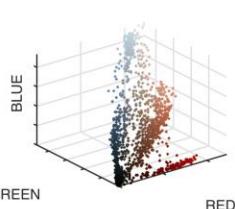
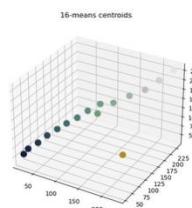
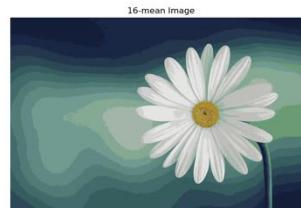
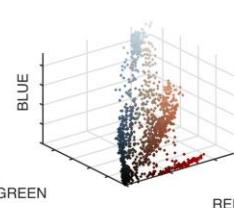
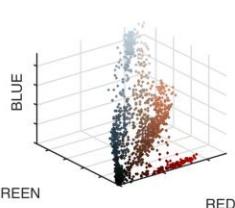
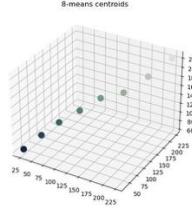
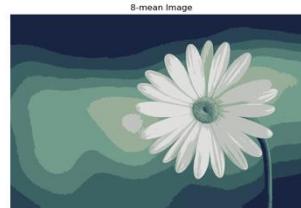
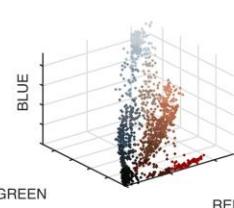
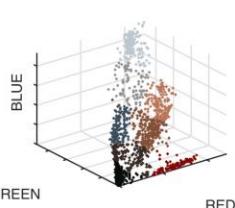
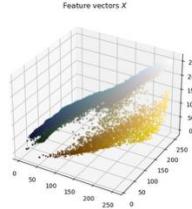
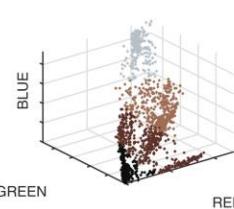
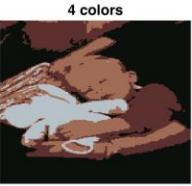
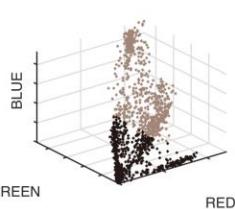
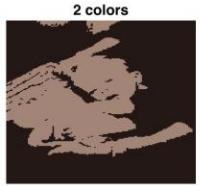
- All clusters have spherical distribution (same to all directions or isotropic)
- Prone to local minima
- Need to choose K
- Can be very slow: each iteration is  $O(KN)$  for N-dimensional points

# K-Means Quantization

---

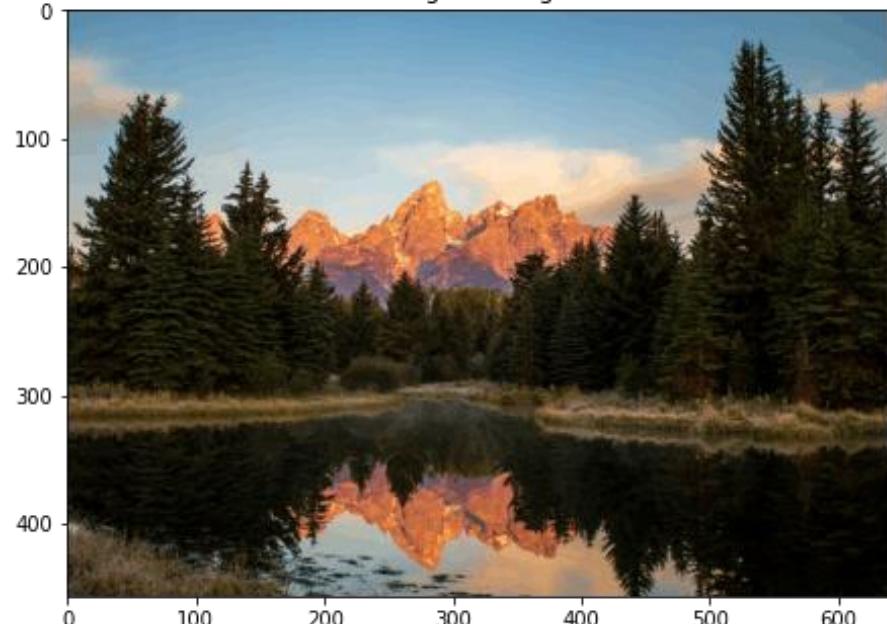
1. Select a value of K
2. Select a feature vector for every pixel (color, texture, position, or combination of these etc.)
3. Define a similarity measure between feature vectors (Usually Euclidean distance)
4. Apply the K-means algorithm to all the feature vectors

# K-Means Quantization Examples

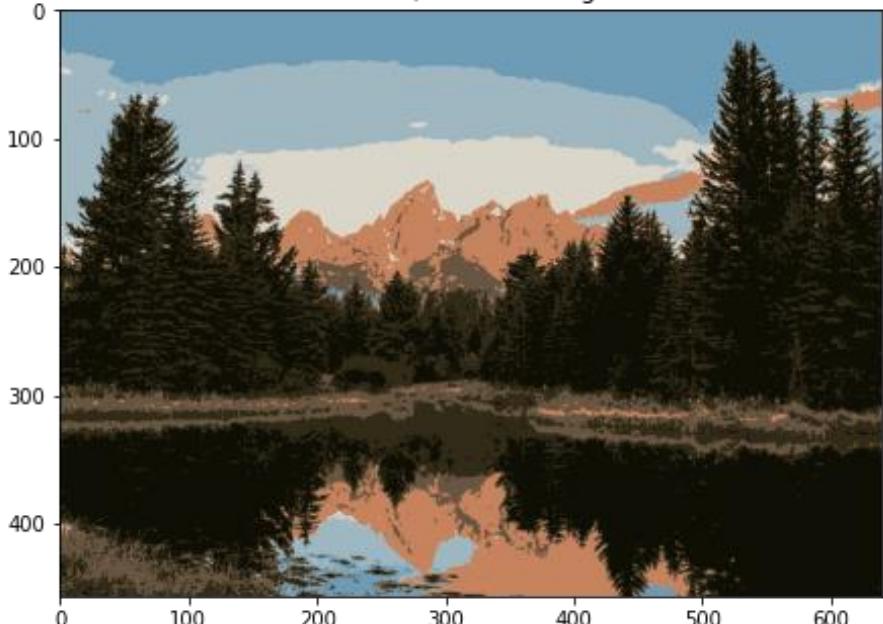


# K-Means Quantization Examples

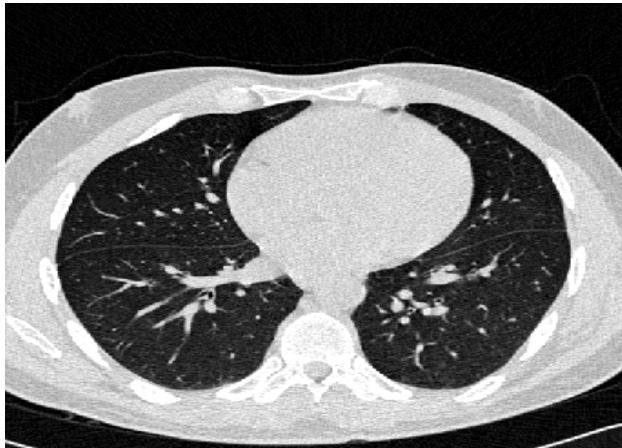
Original Image



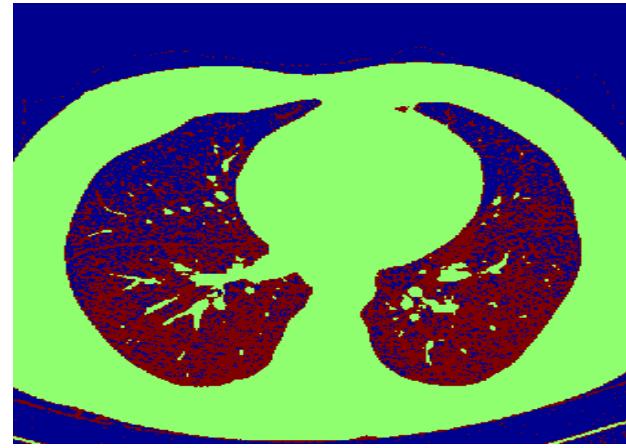
Color Quantized Image



# K-Means Quantization Examples



Input image ( $I$ )



Three-cluster image using the  
gray levels of *input image*

Matlab code:

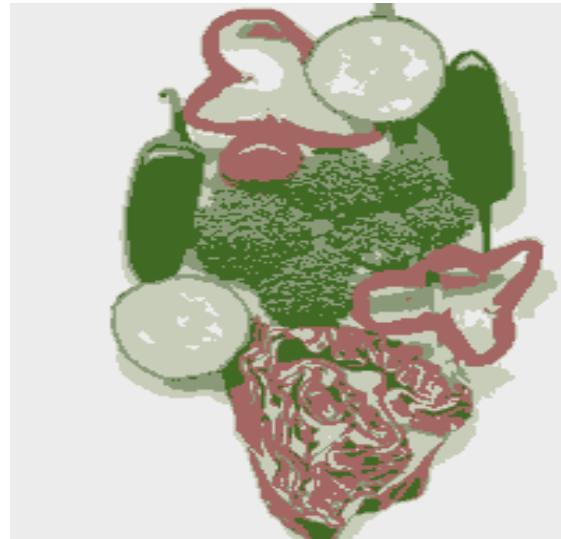
```
I = double(imread('...'));
```

```
J = reshape(kmeans(I(:,3),size(I)); % need statistics toolbox
```

# K-Means Quantization



Color Image

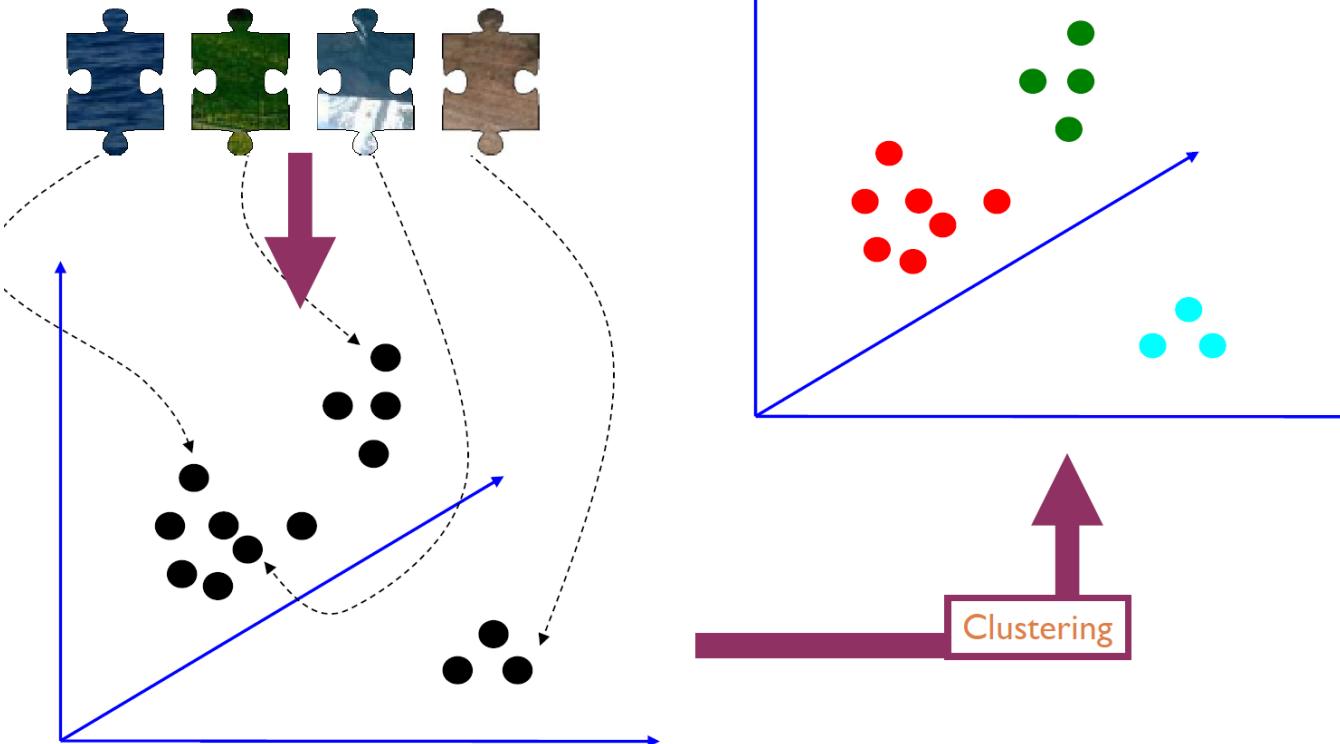


Segmented on color

K-means clustering used for reducing the number of colors (e.g. k=4)

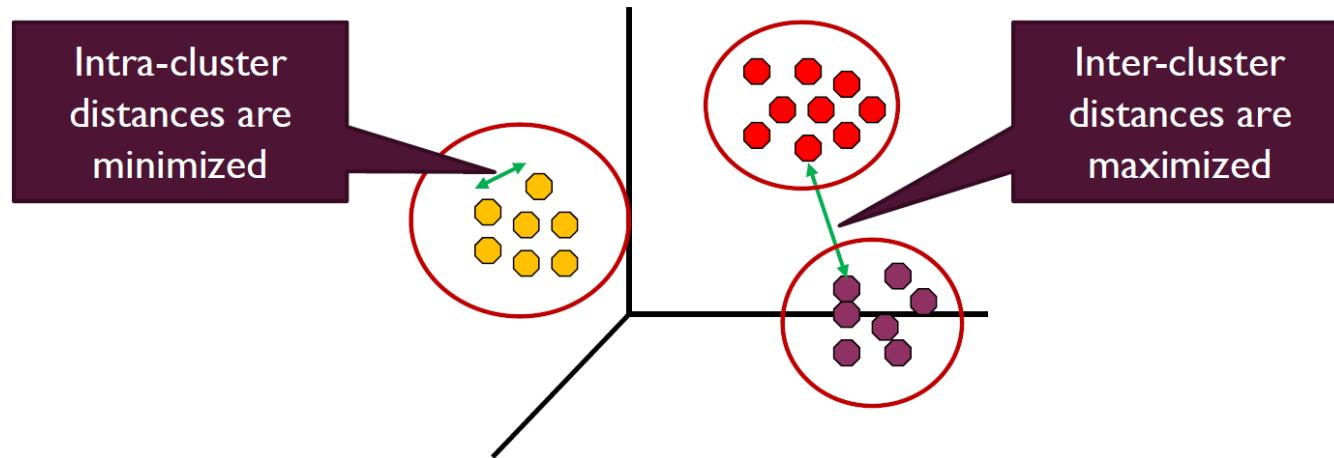
# Reminder: Feature Clustering for BoW

## CLUSTERING FOR BOW



# Reminder: Feature Clustering for BoW

Find clusters of similar descriptors

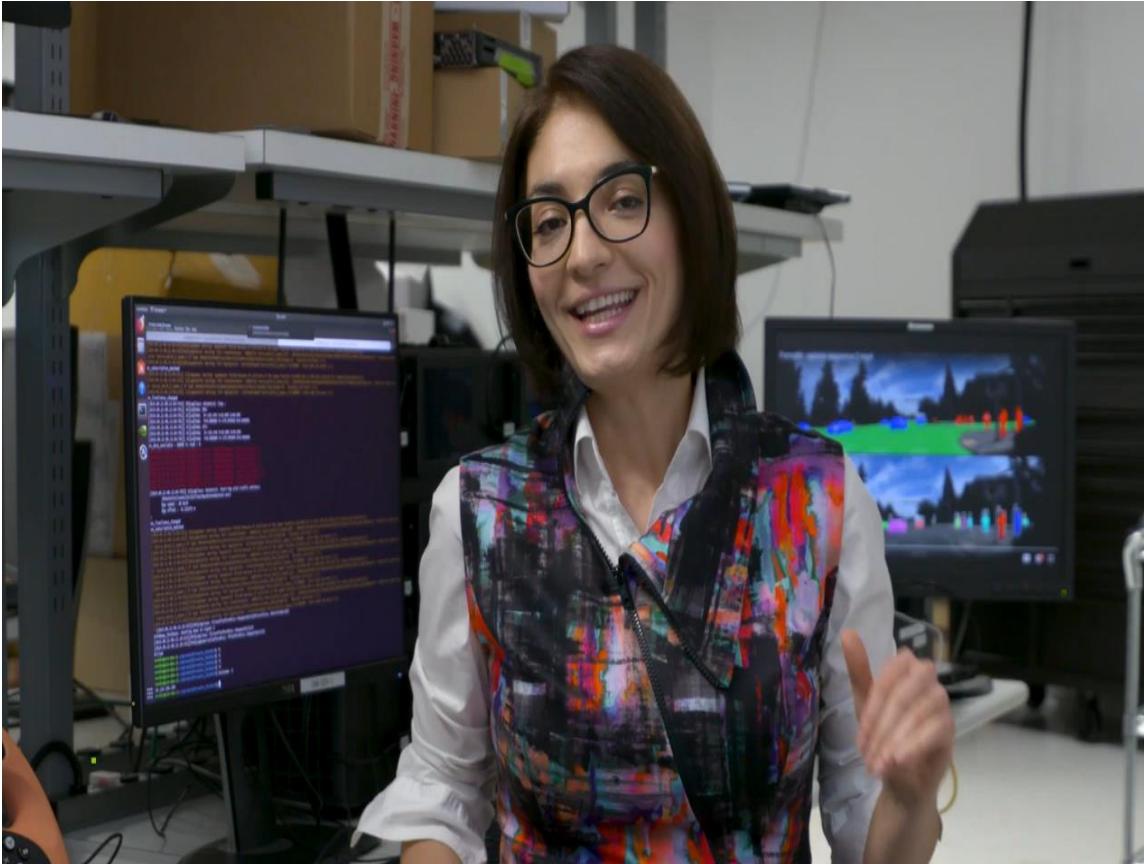


# DeepLab v3+ Semantic Segmentation on Cityscape Dataset

CV



# Panoptic Segmentation (Semantic+Instance)



---

# Questions?

# Disclaimer

---

Many of the slides used here are obtained from online resources (including many open lecture materials) without appropriate acknowledgement. They are used here for the sole purpose of classroom teaching. All the credit and all the copyrights belong to the original authors. You should not copy it, redistribute it, put it online, or use it for any other purposes than for this course.