

Computer Vision 1

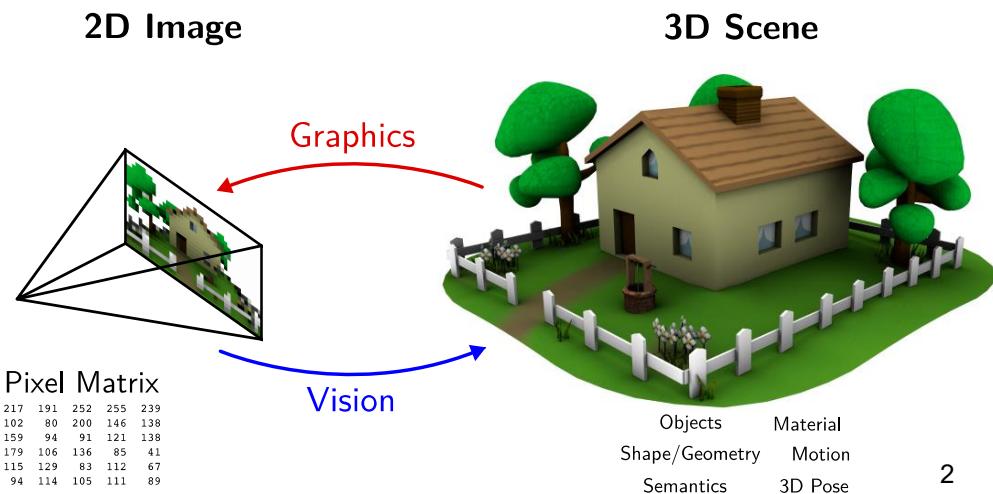
HC1b

Camera Image Formation Color

Dr. Martin Oswald, Dr. Dimitris Tzionas, Dr. Arun Mukundan,
[m.r.oswald, d.tzionas, a.mukundan]@uva.nl

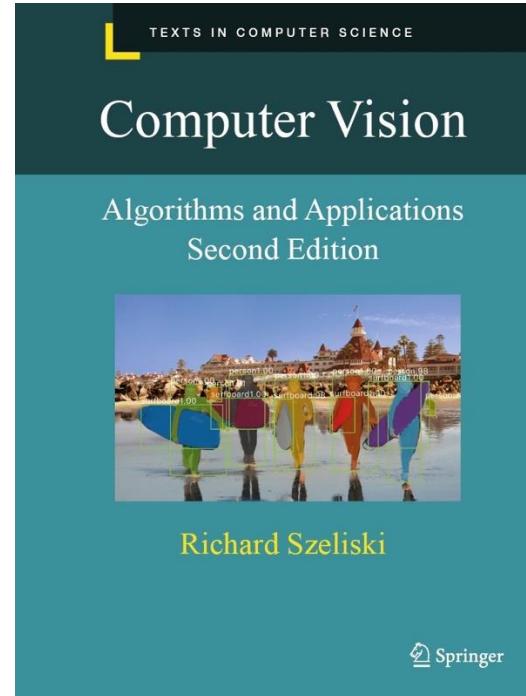
Outline

- Pinhole Camera
- Geometric Image Formation
- Photometric Image Formation
- Image Representation



Textbook – Sections

- Sec. 2.1.[1-2]
- Sec. 2.1.[4-5]
- Sec. 2.2
- Sec. 2.3.[1-2]



Human Vision – Retina

Optics of eye → Create a **focused 2D image** of 3D world **on retina**

Retina → The eye's innermost & **light-sensitive layer** of tissue

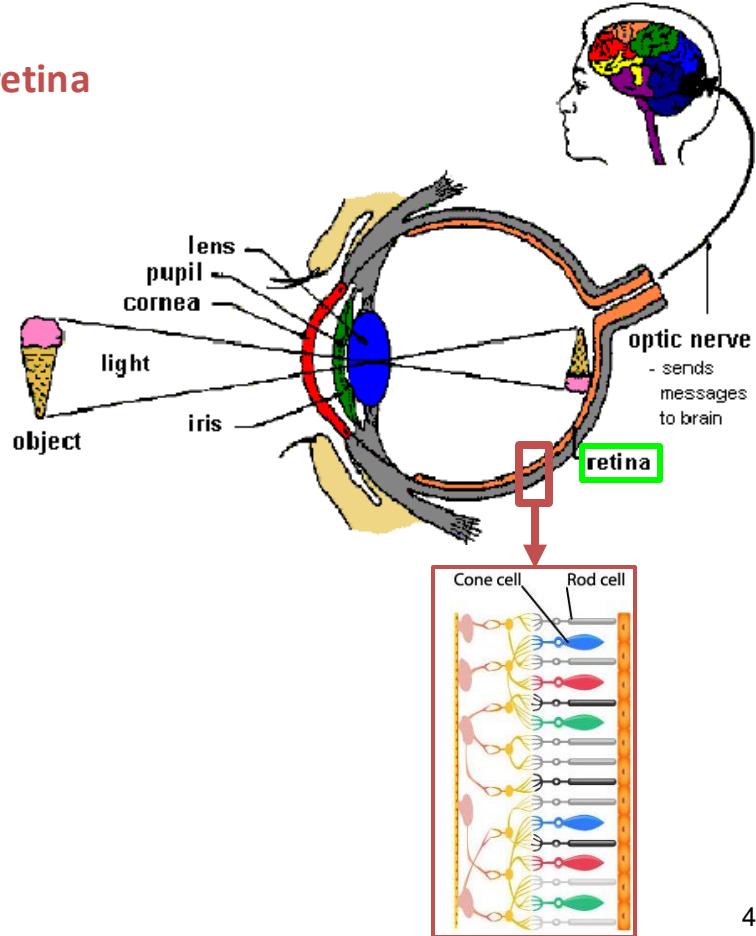
Luminance ~ Sensed in **light-sensitive cells** in retina
(visual energy)

(rods & cones)

Collection of measurements → Form 'image'

Collection of pixels on screen → Form digital image

Retina: **Processes** that image &
Send **impulses** along optic nerve to **visual cortex**
to create **visual perception**



Human Vision – Retina



Peyman Milanfar

@docmilanfar

Distinguished Scientist at Google.

...

The retina is arguably the most impressive part of the brain – it's also the only part of the brain that faces the world directly – it's a sensor and processor in one

Consumes 50% more energy per gram than the rest of the brain.

1000:1 compression from retina to optic nerve

25000:1 from optic nerve to brain

For every 1Gb collected by the retina, 1Mb is sent to the brain thru the optic nerve and < 100bits used, at a rate of about 875Kbps. Similar to broadband internet.

NIH Public Access
Author Manuscript

Published in final edited form as:
Curr Biol. 2006 July 25; 16(14): 1428-1434.

How Much the Eye Tells the Brain

Kristin Koch¹, Judith McLean¹, Ronen Segev², Michael A. Freed¹, Michael J. Berry², Vijay Balasubramanian¹, and Peter Sterling³.

¹ Department of Neuroscience University of Pennsylvania Philadelphia, Pennsylvania 19104

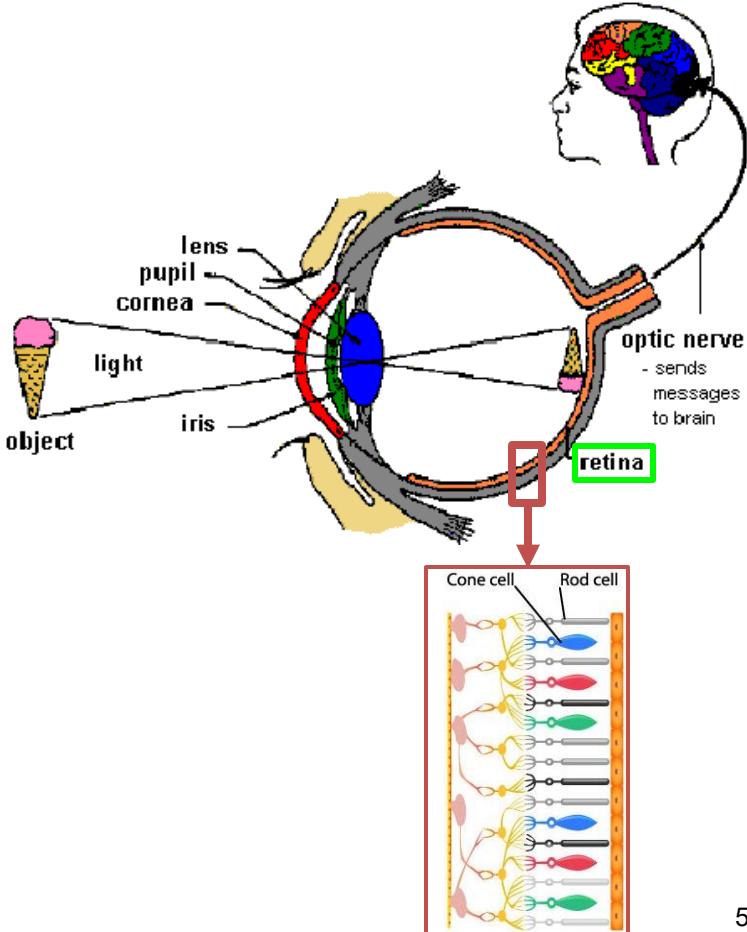
² Department of Molecular Biology Princeton University Princeton, New Jersey 08544

³ Department of Physics University of Pennsylvania Philadelphia, Pennsylvania 19104

Summary

In the classic “What the frog’s eye tells the frog’s brain,” Lettvin and colleagues [1] showed that different types of retinal ganglion cell send specific kinds of information. For example, one type responds best to a dark, coneless form moving centripetally (inify). Here we consider a complementary question: how much information does the visual field and how it appears among different cell types? Because motion information pins down a moving object in space and preserves the sense of motion in natural scenes, we measured information rates for seven types of ganglion cell. Mean rates varied across cell types ($\sim 13 \text{ bits s}^{-1}$) more than across stimuli. Sluggish cells transmitted information at lower rates than fast ones, but they did so with higher information density (per pixel) and thus had higher coding efficiency. Calculating the proportion of each cell type from receptive field size and coverage factor, we conclude (assuming independence) that the approximately 10^5 ganglion cells transmit on the order of $875,000 \text{ bits s}^{-1}$. Because sluggish cells are equally efficient but more numerous, they account for most of the information. With approximately 10^5 ganglion cells, the human retina would transmit data at roughly the rate of an Ethernet connection.

<https://x.com/docmilanfar/status/1914529429445075157>



Pinhole Camera – Toy Example

Retina @ eye ~~ Image sensor @ camera

Add light-sensitive ‘screen’ in front of an object

Lights reflects onto the object & hits the screen

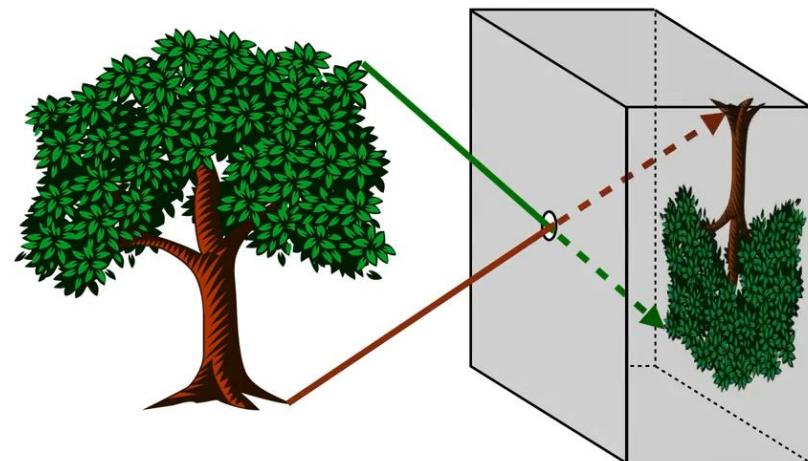
Problem: \forall object point: Emits light in **all** directions

Image **washed-out** 😞

Sun **washes out++** 😞

Solution: Box with a **hole** → **Filters** most light rays

Result: **Crisp** 😊 (but **flipped** 😞) image



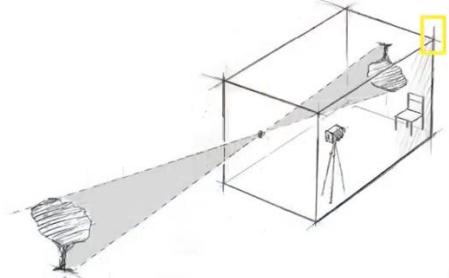
Pinhole Camera - Real Example

'Programmatically'
Flip Vertical



Pinhole Camera - Real Example

Making Your Own Room With a View | National Geographic
National Geographic



<https://youtu.be/gvzpu0Q9RTU>



Pinhole Camera – Real Example



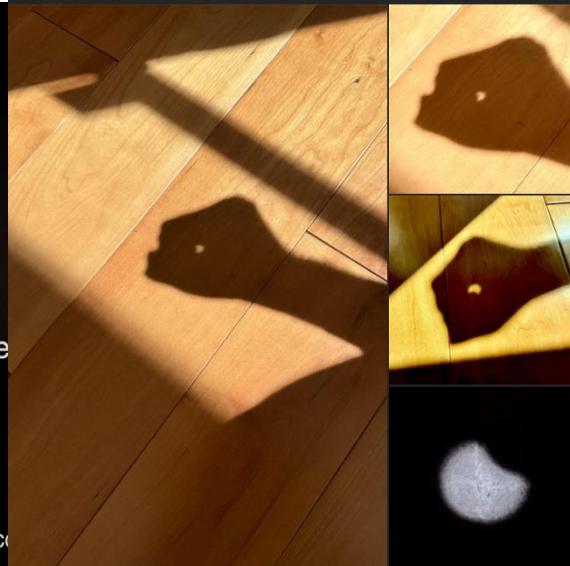
© O'Dea at Wikimedia Commons, CC BY-SA 3.0



Eric Horvitz

@erichorvitz

Chief Scientific Officer, Microsoft



Pinhole Camera – Real Example

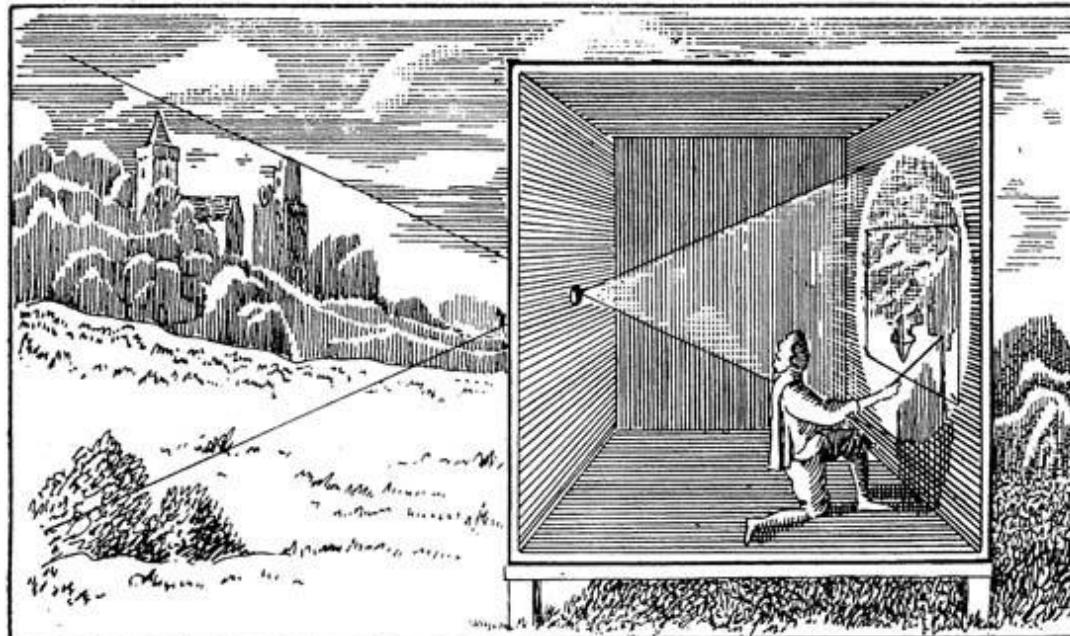
Normal
Sun



Solar
Eclipse

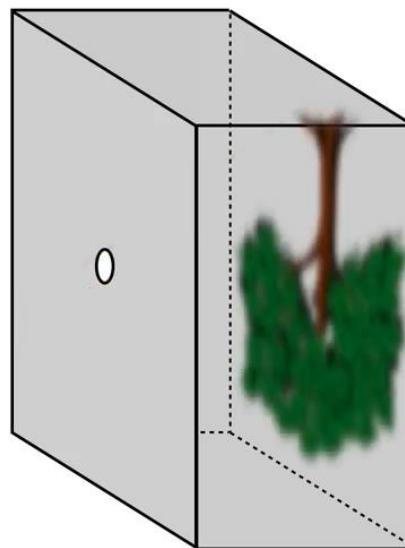


Pinhole Camera - Camera Obscura



3D → 2D projection
with a [human in the loop](#)

Pinhole Camera - Aperture



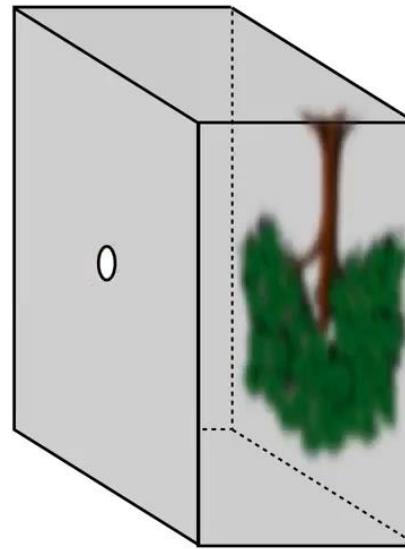
The hole (*aperture*) lets a narrow *beam of light* through

This creates a little *spot of color* on the screen
(that might also *overlap*)

Adds together many *overlapping* color *spots*

Result: ***Blurry*** image ☹

Pinhole Camera – Aperture Size – Blur



↑ Apperture size
↑ Size of color spots

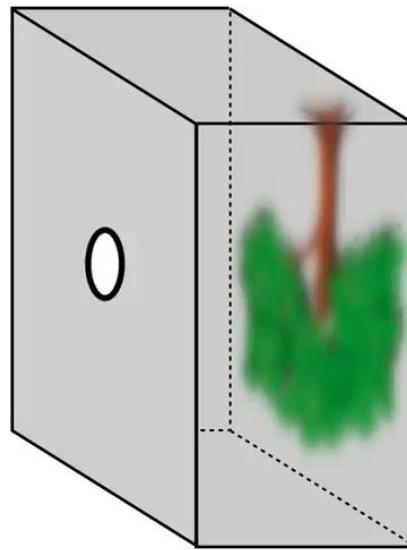
↑ Blurriness
↓ Sharpness

↓ Apperture size
↓ Size of color spots

↓ Blurriness
↑ Sharpness

Pinhole Camera – Aperture Size – Brightness

CV



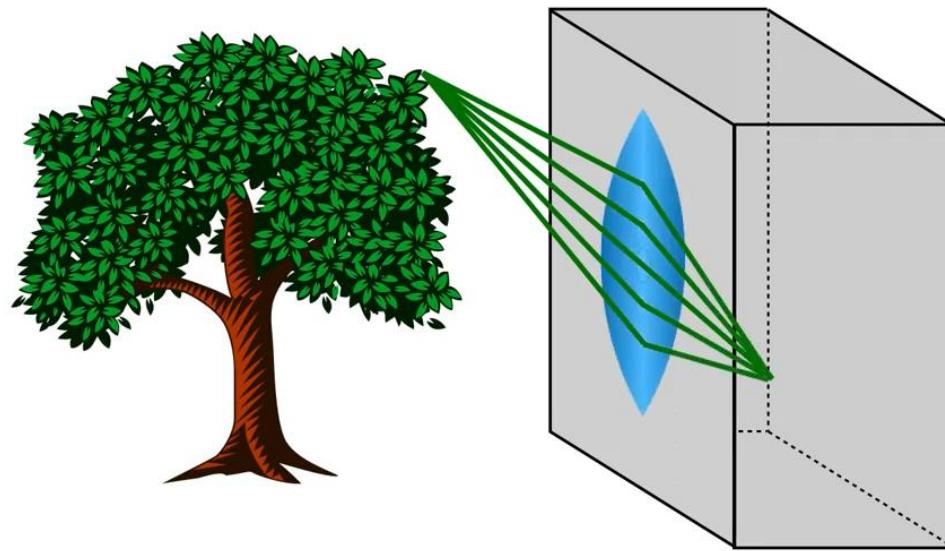
↑ Apperture size
↓ Sharpness
↑ Brightness

Trade off!

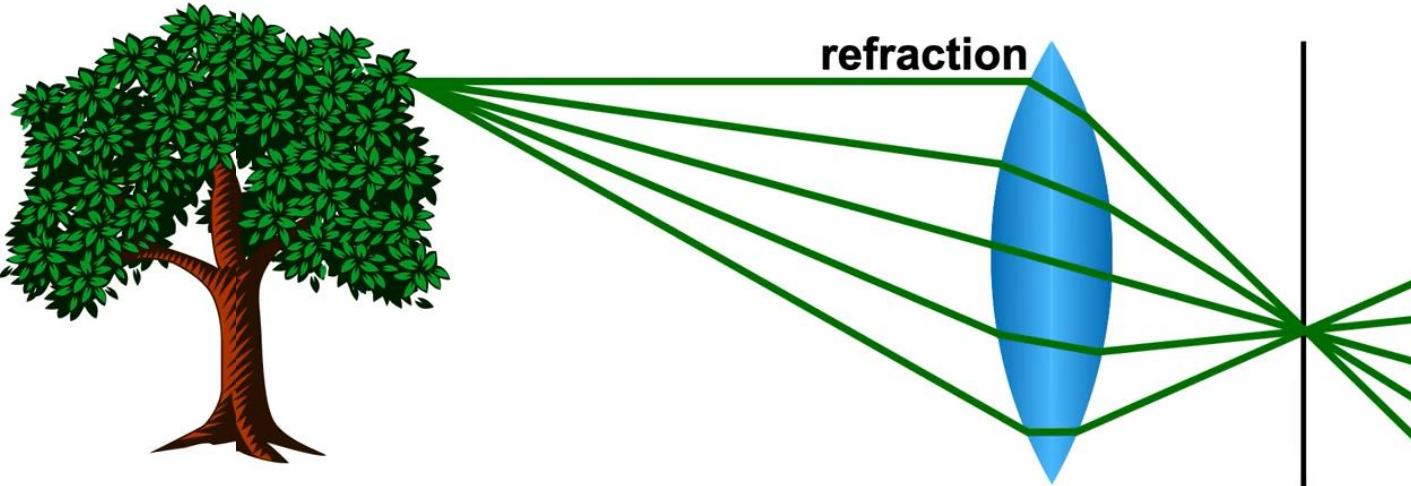
↓ Apperture size
↑ Sharpness
↓ Brightness

Pinhole Camera - Capture more Light

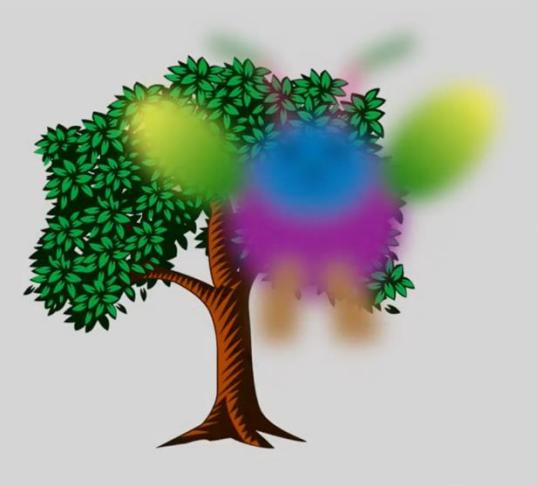
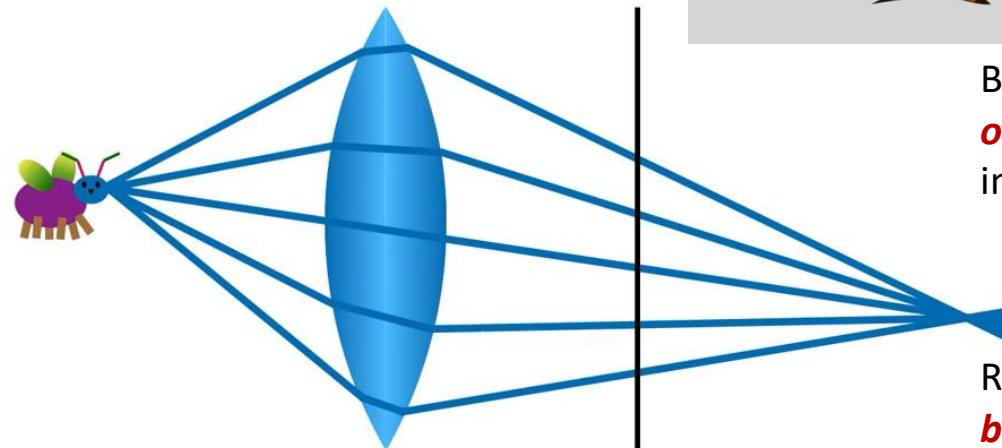
CV



Pinhole Camera - Capture more Light



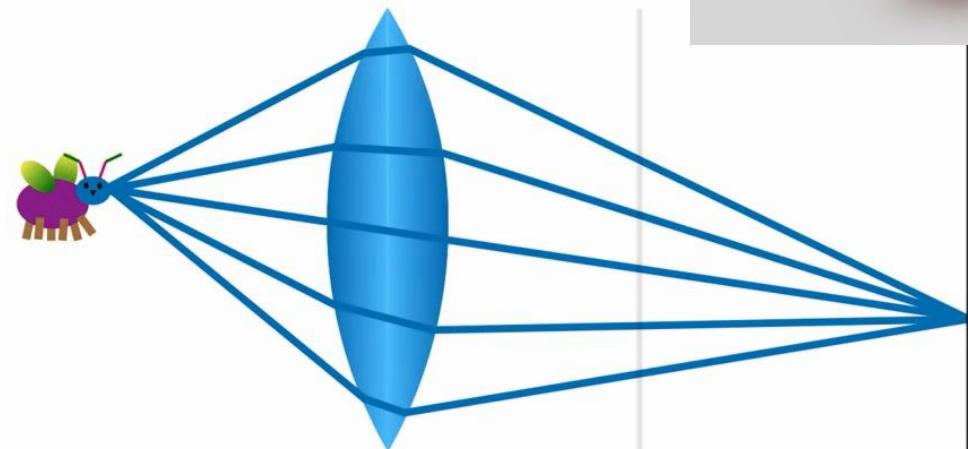
Pinhole Camera – Focus



Bug appears
out of focus
in the image

Rays converge
behind the
image plane

Pinhole Camera – Focus

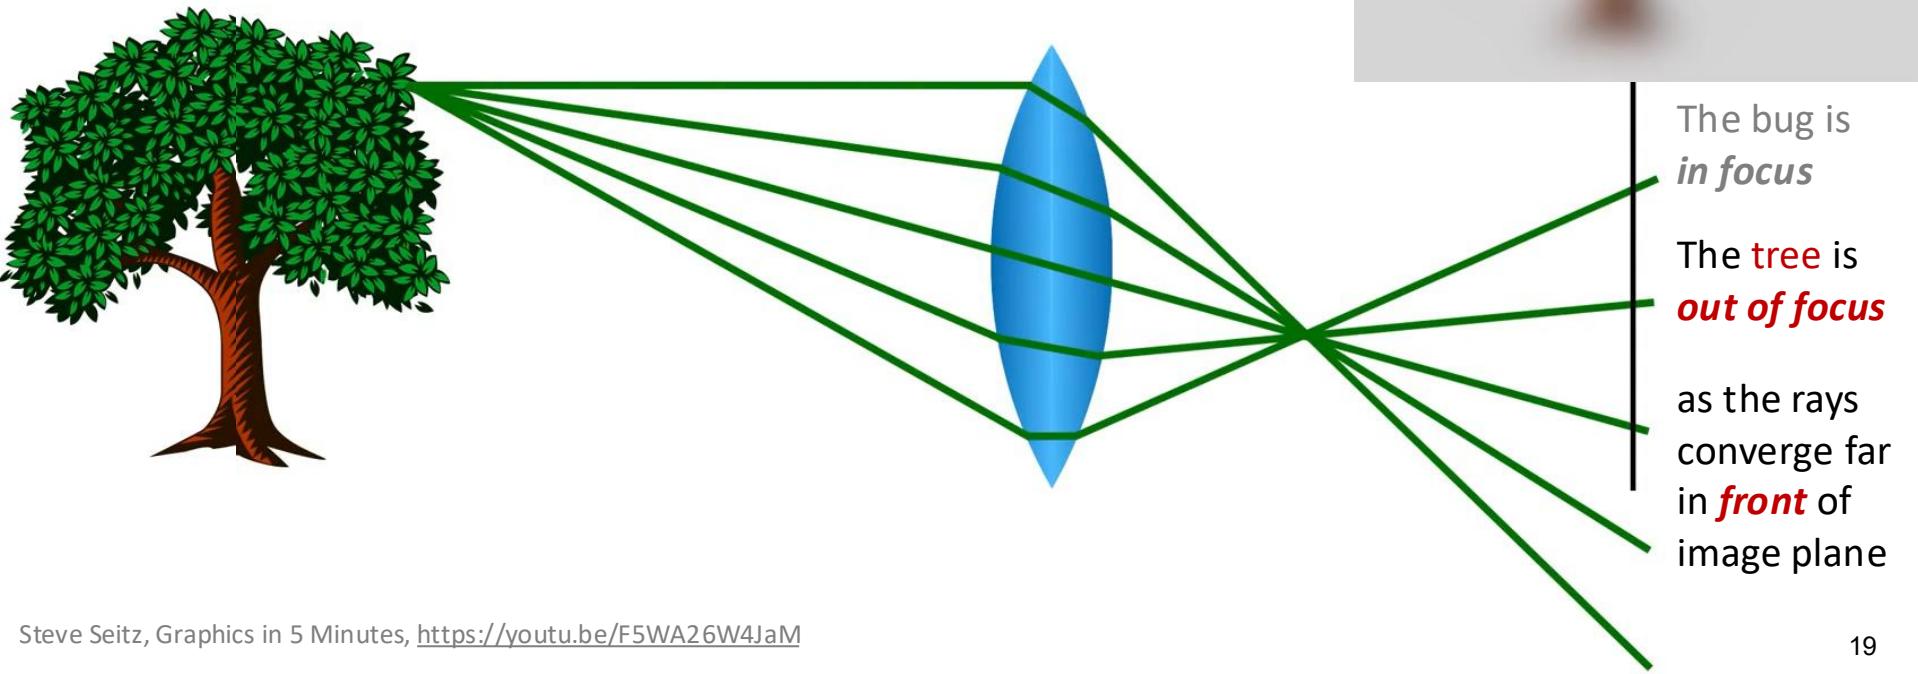


The **bug** is
in focus

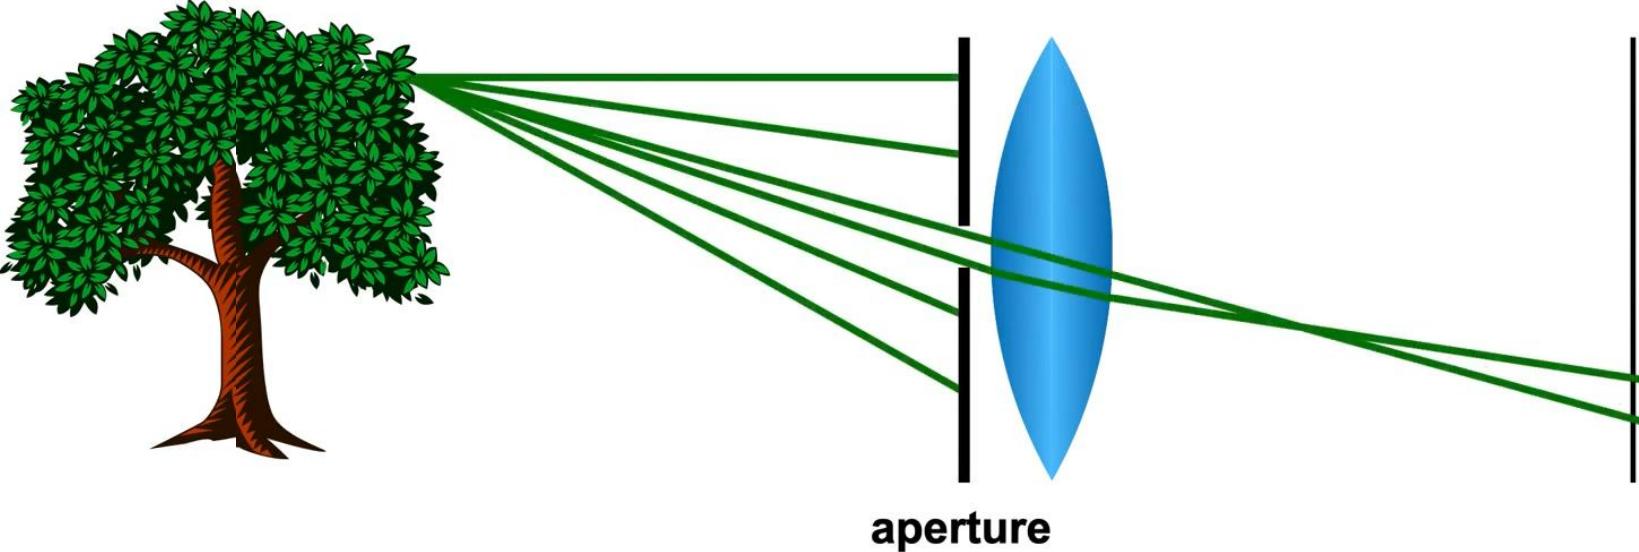
The **tree** is
out of focus

Increase
lens-2-image-
plane
distance

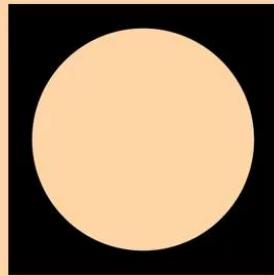
Pinhole Camera – Focus



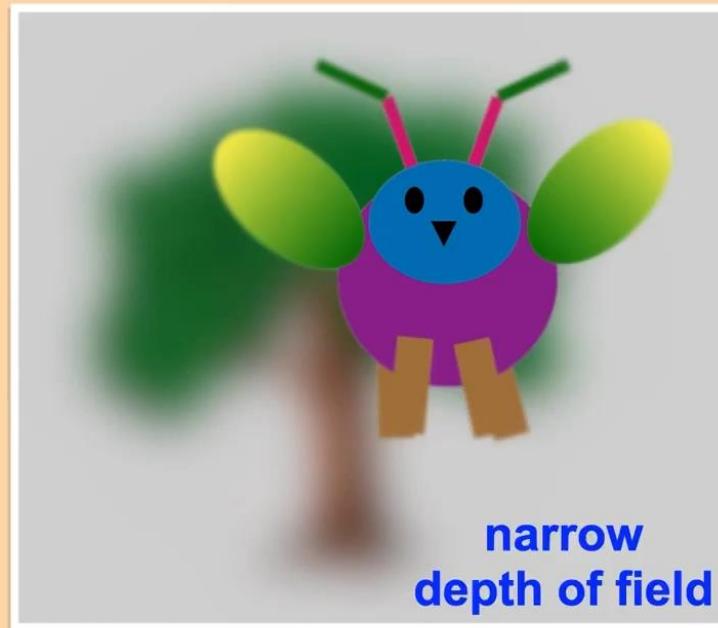
Pinhole Camera – Focus



Pinhole Camera – Focus



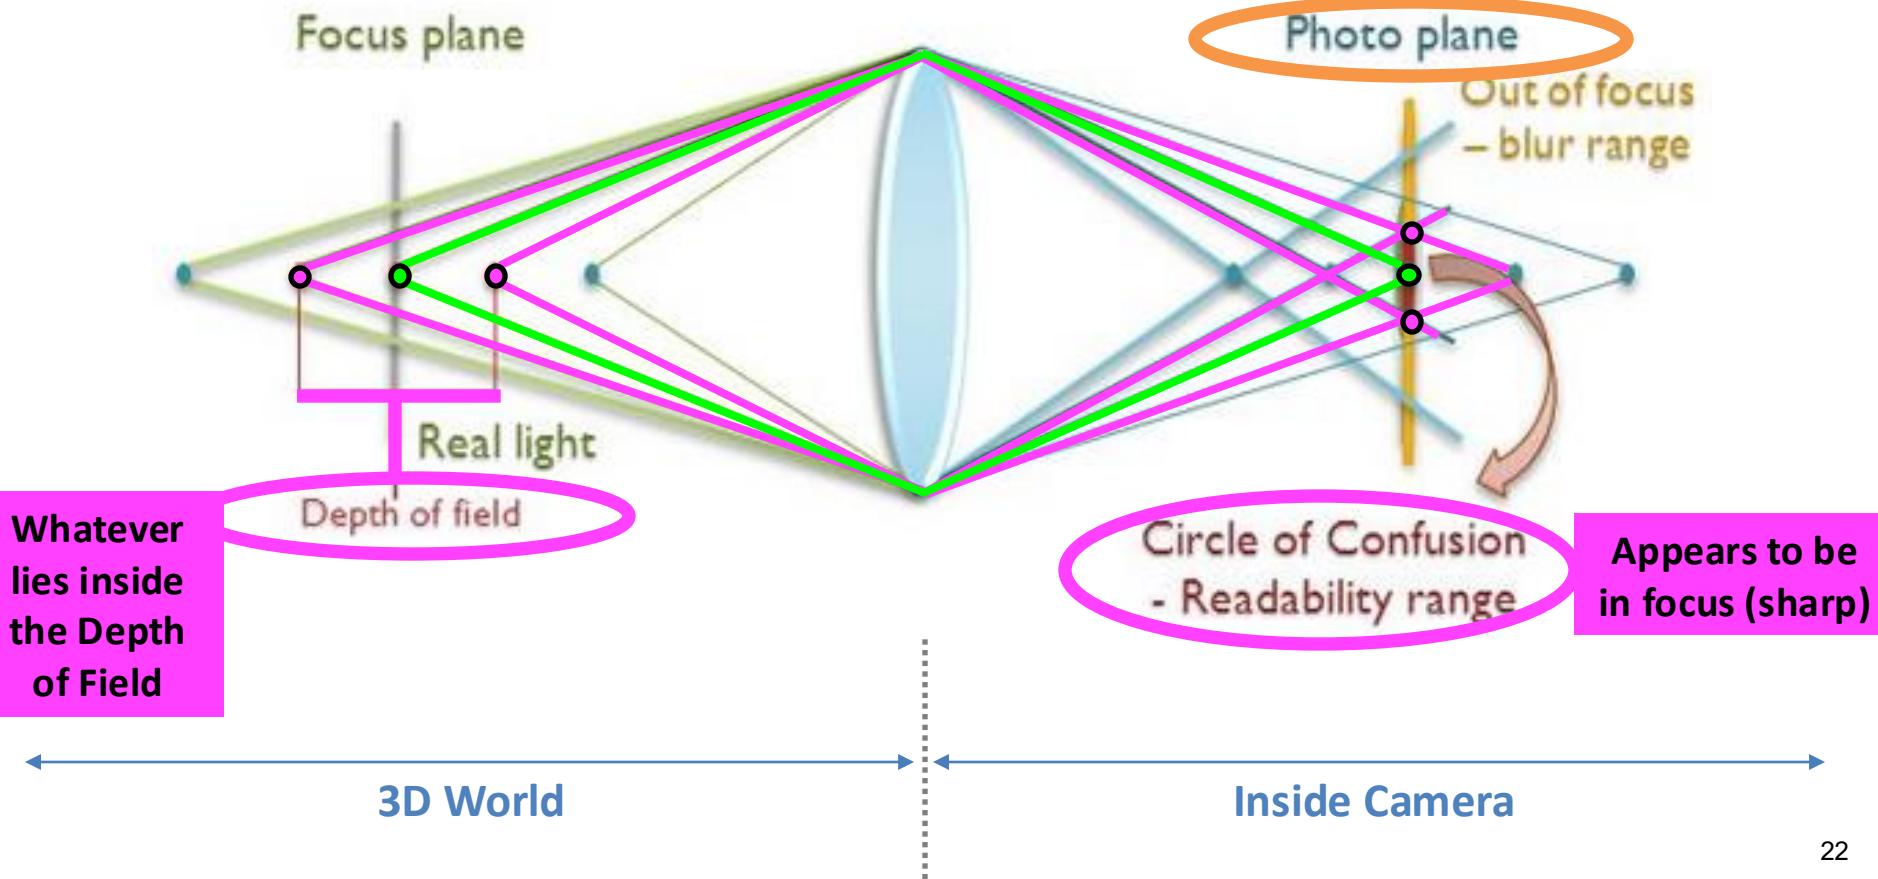
big
aperture



↑ Apperture
↓ Depth of Field

↓ Apperture
↑ Depth of Field

Aperture Adjustment – Depth of Field



Aperture Adjustment – Depth of Field

Disclaimer: The f here is just a dummy symbol. It should not be mistaken with the 'focal length f' defined in later slides



Aperture wide open → f/1.8



f/2.8



f/4.0



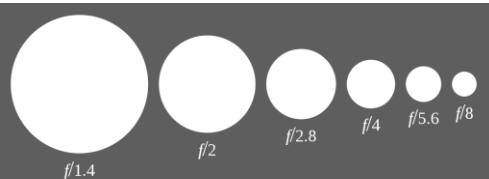
f/5.6



f/16

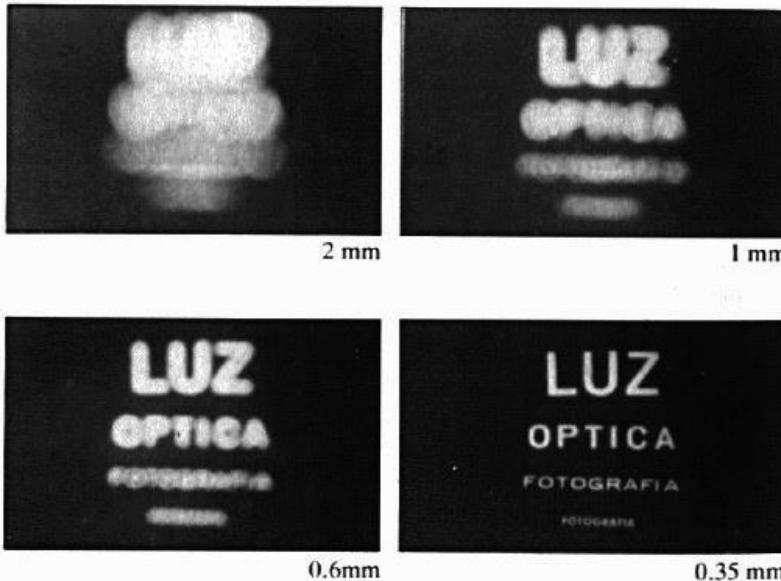


f/22 ← 'Narrow' Aperture



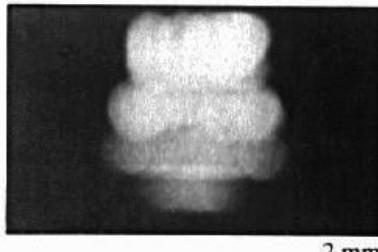
<https://en.wikipedia.org/wiki/F-number>

Aperture Shrinking



Why not make the aperture
as small as possible?

Aperture Shrinking



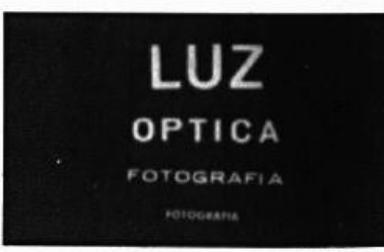
2 mm



1 mm



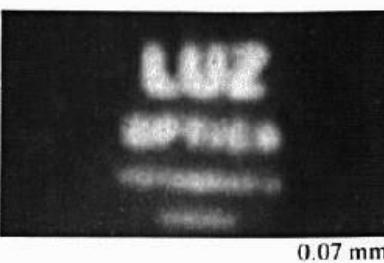
0.6mm



0.35 mm

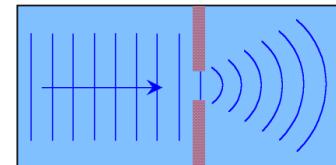


0.15 mm



0.07 mm

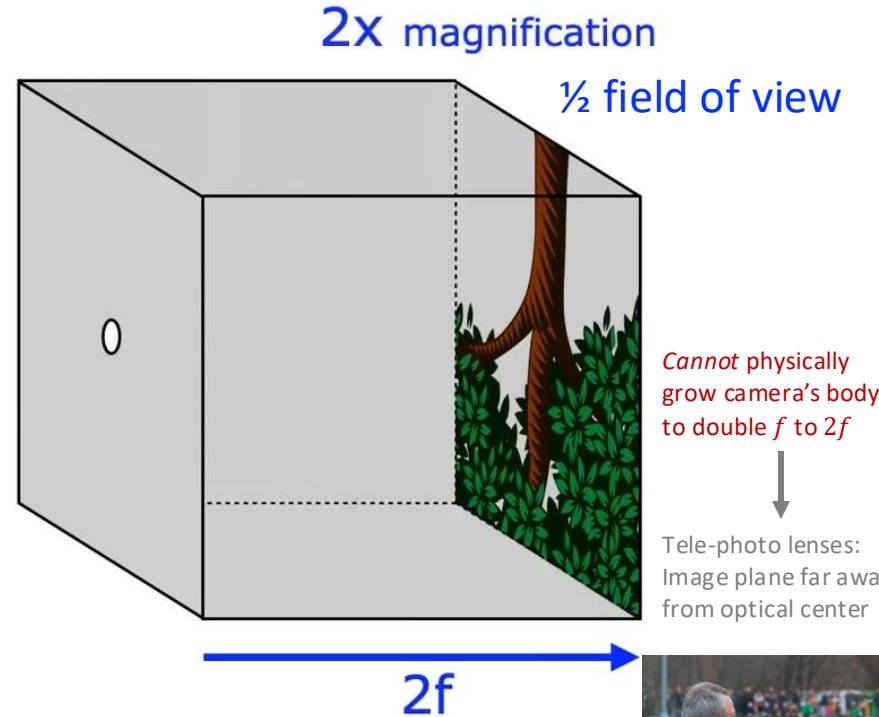
Less light gets through →
Diffraction effects →



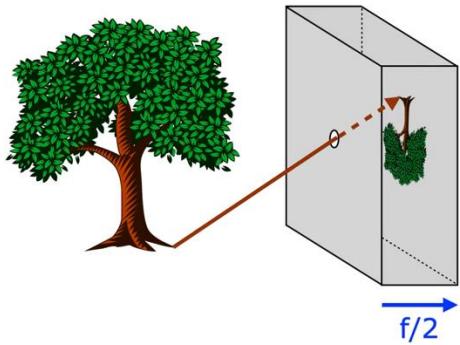
Pinhole Camera – Focal Length



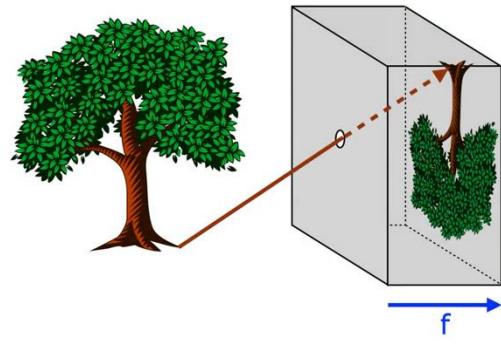
2x magnification



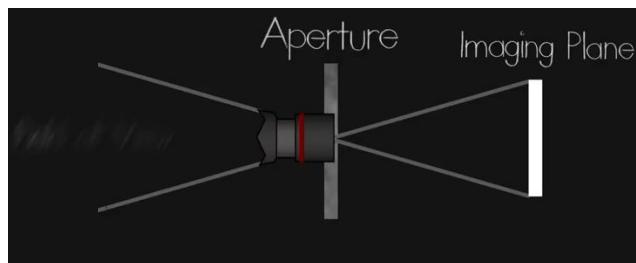
Pinhole Camera - Focal Length



- ↓ Focal length
- ↓ Object size (shrinking)
- ↑ Field of view (more things fit the image)



Impossible to move the image plain!
 'Moveable' lens & 'bend' light rays & change focus



- ↑ Focal length
- ↑ Object size (magnification)
- ↓ Field of view (less things fit the image)

Outline

- Pinhole Camera
- Geometric Image Formation
- Photometric Image Formation
- Image Representation

Pinhole Camera – Projection Matrix

Camera Coordinate Frame:

- Origin O at camera center (hole)
- Z axis perpendicular to Image Plane I (towards outside the camera)
- X and Y axes parallel to Image Plane
- Projection plane is the plane $Z = -f$ (*)
 $Z = +f$ (**)

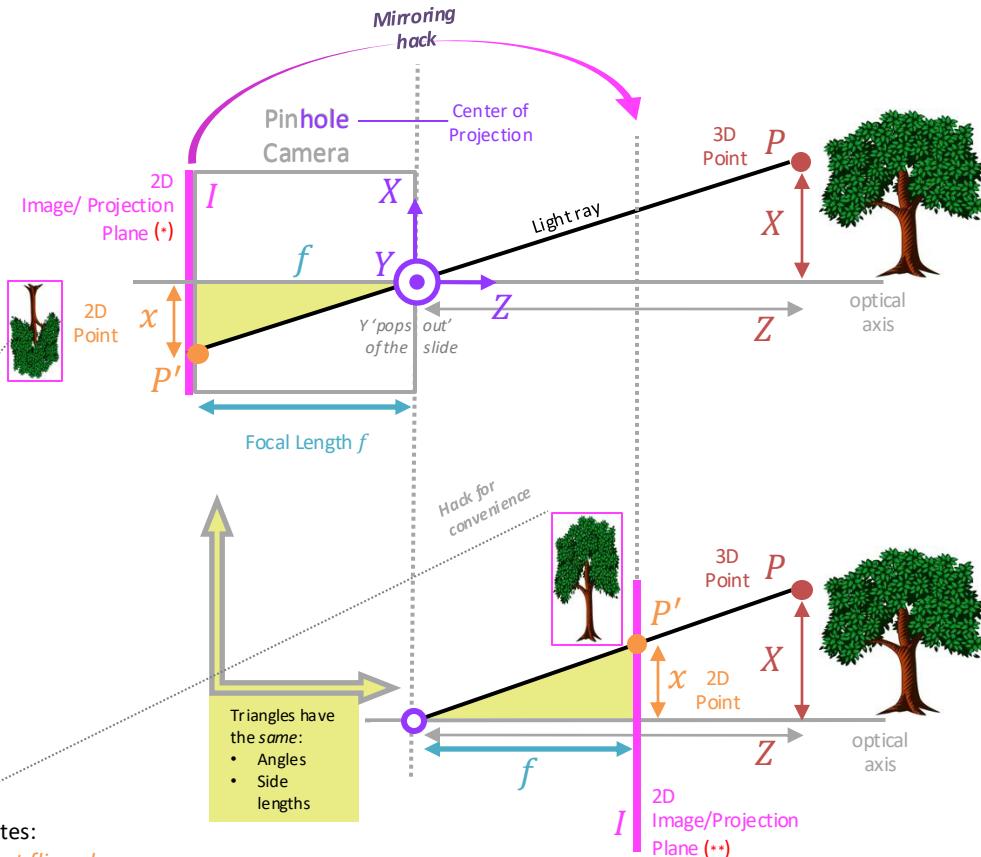
3D point $\xrightarrow{\text{Projection}}$ 2D point

$$P = \begin{bmatrix} X \\ Y \\ Z \end{bmatrix}$$

$$P' = \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} -f \frac{X}{Z} \\ f \frac{Y}{Z} \\ -f \frac{Z}{Z} \end{bmatrix} \xrightarrow{\text{mirroring hack}} \begin{bmatrix} f \frac{X}{Z} \\ f \frac{Y}{Z} \\ f \frac{Z}{Z} \end{bmatrix}$$

Minus sign denotes:
projected image is flipped
(horiz. & vertically)

Plus sign denotes:
projected image is not flipped
(see bottom-right plot in this slide)



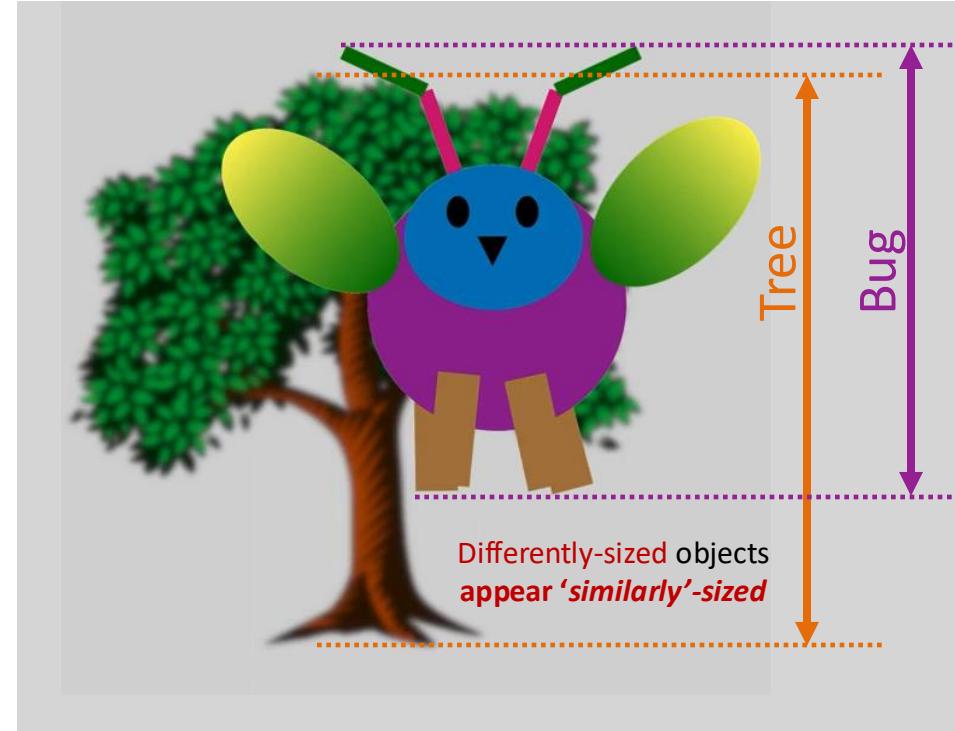
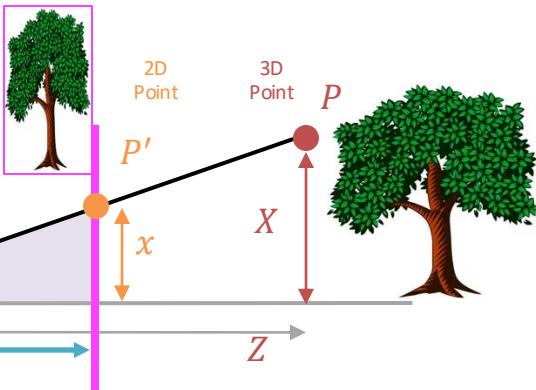
Pinhole Camera - Object Size

Scaled by $\frac{1}{Z}$

$$\text{3D point } P = \begin{bmatrix} X \\ Y \\ Z \end{bmatrix}$$

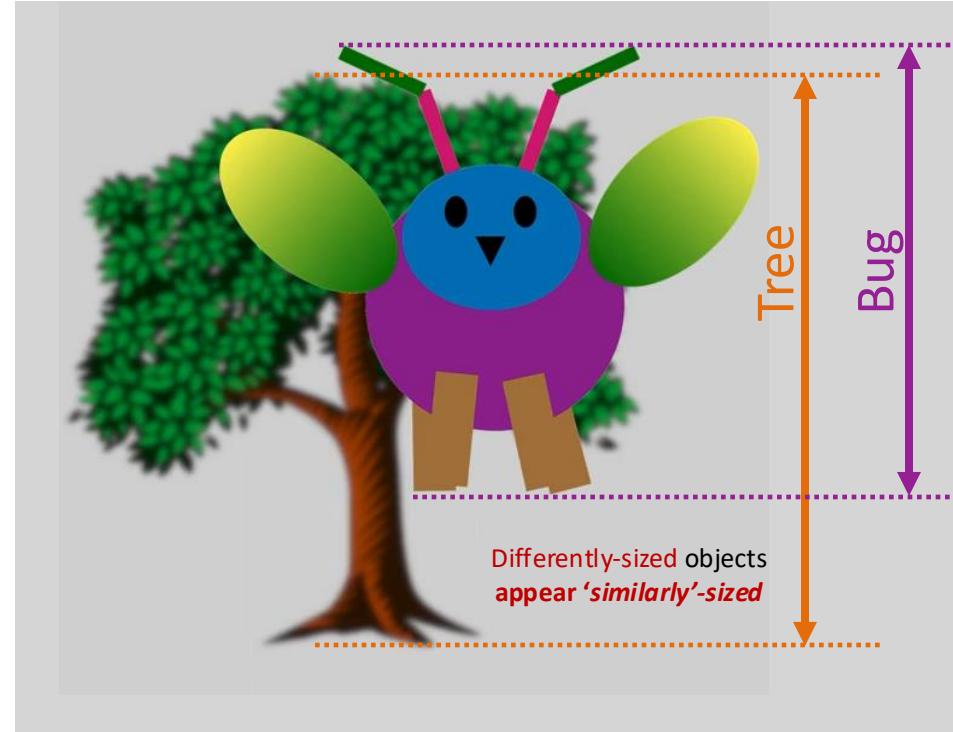
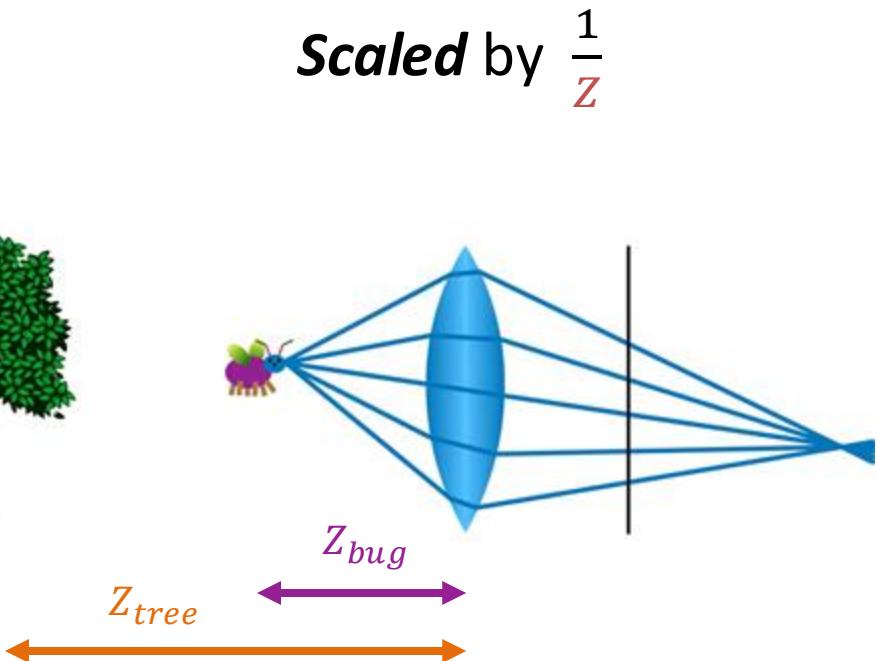
Projection

$$\text{2D point } P' = \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} f \frac{X}{Z} \\ f \frac{Y}{Z} \end{bmatrix}$$



Steve Seitz, Graphics in 5 Minutes, <https://youtu.be/F5WA26W4JaM>

Pinhole Camera - Object Size

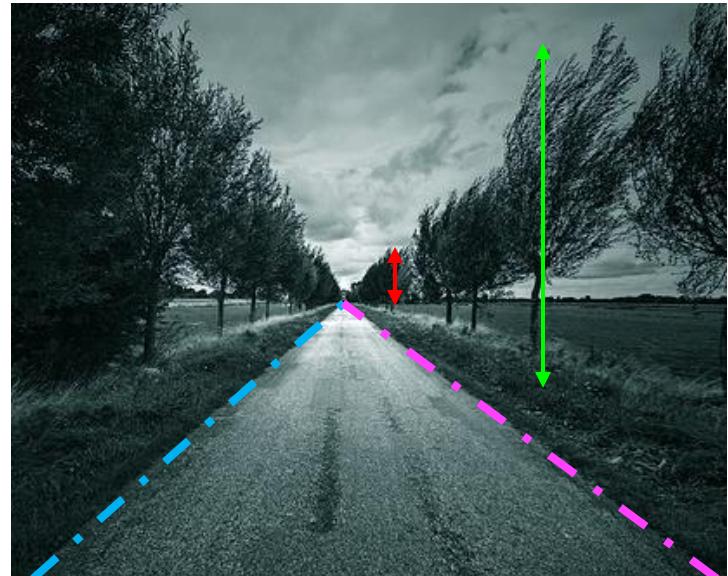


Steve Seitz, Graphics in 5 Minutes, <https://youtu.be/F5WA26W4JaM>

Projective Geometry

Projective transformation preserves:

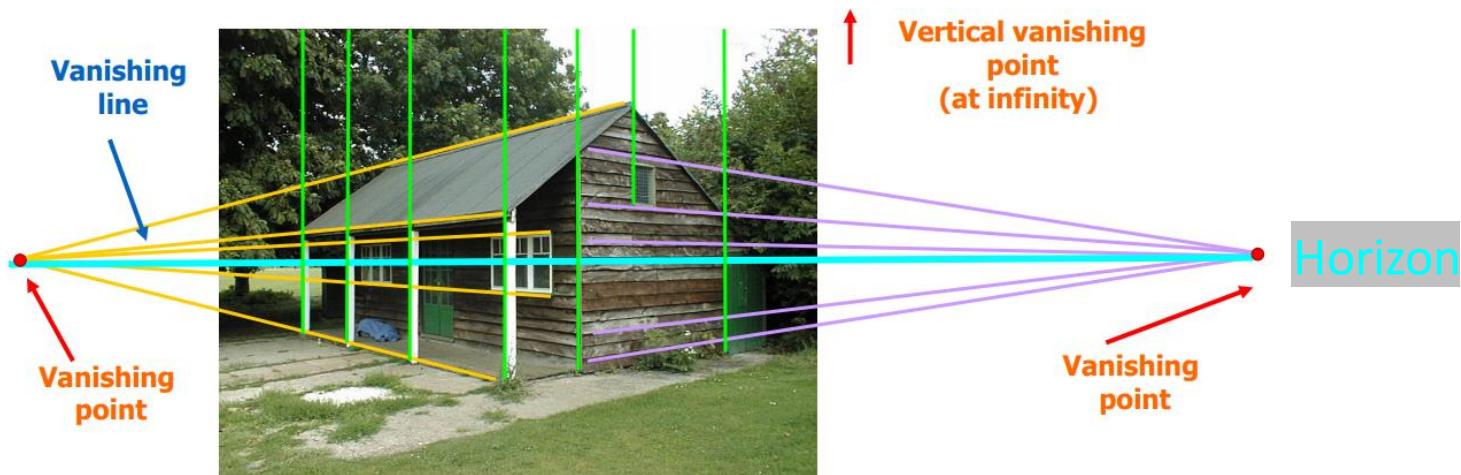
- **Collinearity** of points (straight lines remain straight)
- **Nothing else!** Changes:
 - **Parallelism:** Parallel lines in 3D, after projection meet @ a vanishing point in 2D
 - **Length:** Same-height Trees: far-away VS nearby appear: (small) (big)
 - **Angles:** Depend on viewpoint



Vanishing Points & Lines

Parallel lines 'converge' @ vanishing point

Can lie both inside & outside the image



Vanishing Points & Lines



Parallel lines
intersect @
vanishing point



[Rob Hoeijmakers / robhoeij]
[Amsterdam-Rijn kanaal]
[https://x.com/Rainmaker1973/
status/1927160448169783414](https://x.com/Rainmaker1973/status/1927160448169783414)

Perspective in Art

Oblique Projection



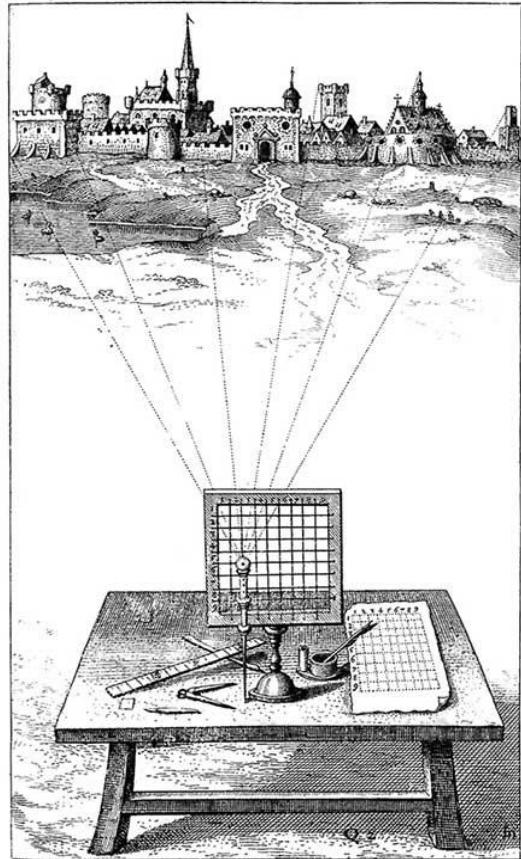
The Birth of Saint John the Baptist: Predella Panel
Giovanni di Paolo, 1454

One-eye Perspective



The Healing of the Cripple and Raising of Tabitha
Masolino, 1427

Perspective in Art



Robert Fludd's [sighting grid](#) (1617)

'Dimensionality Reduction'
Machine ($3D \rightarrow 2D$)

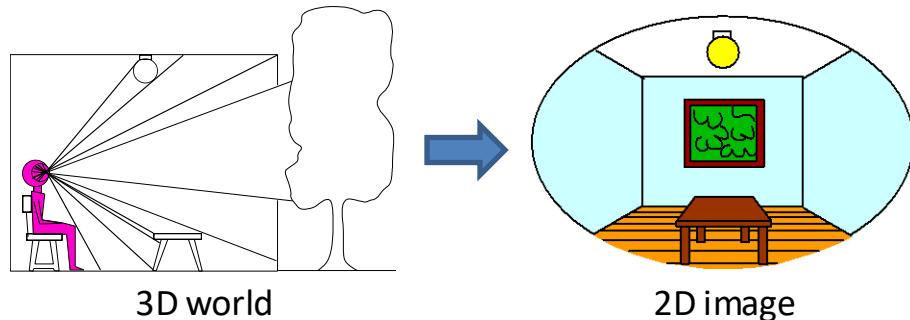


Figure: Stephen E.
Palmer, 2002

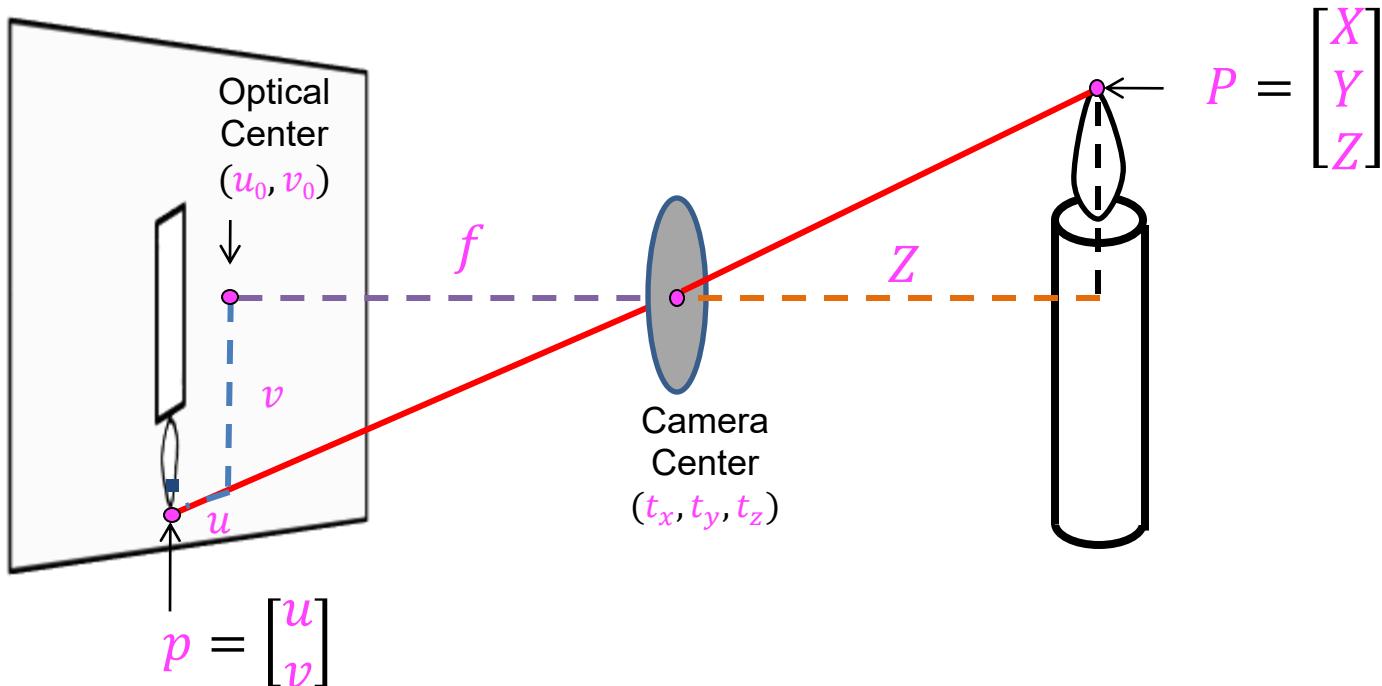
Perspective in Art



Perspective in Art



Projection: 3D \rightarrow 2D coordinates



Projection: 3D \rightarrow 2D coordinates

Length & area size
are not preserved!

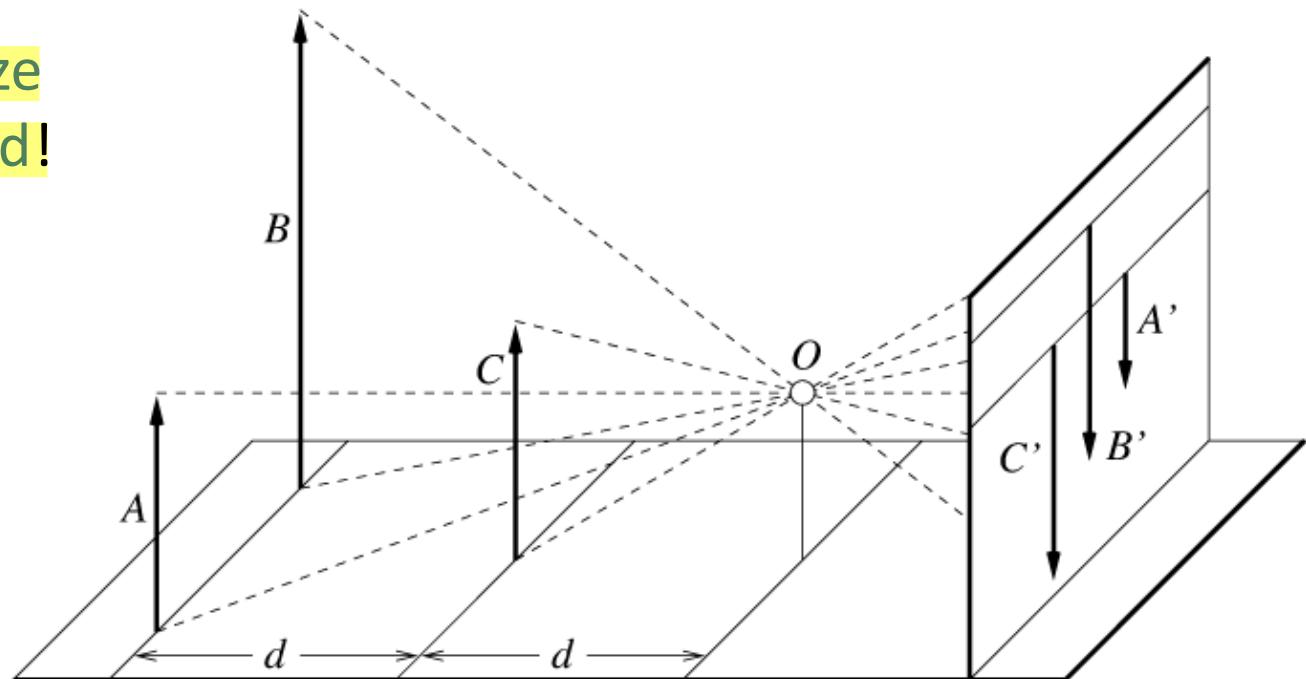


Figure by
David Forsyth

Outline

- Pinhole Camera
- Geometric Image Formation
 - Homogeneous Coordinates
- Photometric Image Formation
- Image Representation

Homogeneous Coordinates

Cartesian Coordinates

- **Rotation** → Applied as matrix mult.
- **Translation** → Cannot



We need a new tool
to simplify math

Homogeneous Coordinates

- **Rotation** → Applied as matrix mult.
- **Translation** → Applied as matrix mult.

Cartesian → Homog. coordinates

$$(x, y) \Rightarrow \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} \quad (x, y, z) \Rightarrow \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

2D (image) coordinates

3D (scene) coordinates

Homog. → Cartesian coordinates

$$\begin{bmatrix} x \\ y \\ z \\ w \end{bmatrix} \Rightarrow (x/w, y/w, z/w)$$

2D (image) coordinates

3D (scene) coordinates

Homogeneous Coordinates

Invariant
to scaling

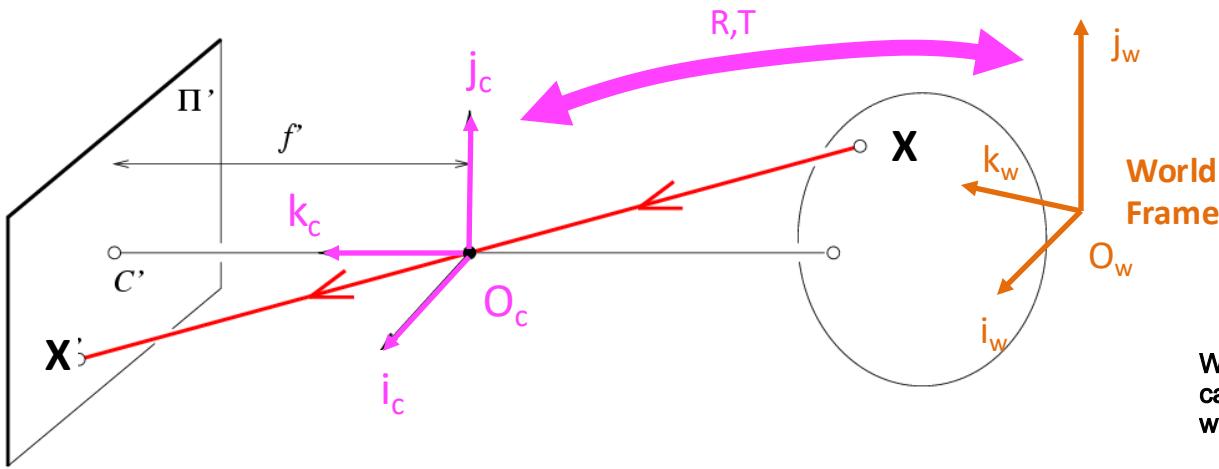
$$k \begin{bmatrix} x \\ y \\ w \end{bmatrix} = \begin{bmatrix} kx \\ ky \\ kw \end{bmatrix} \Rightarrow \begin{bmatrix} \frac{kx}{kw} \\ \frac{ky}{kw} \\ \frac{kw}{kw} \end{bmatrix} = \begin{bmatrix} \frac{x}{w} \\ \frac{y}{w} \\ 1 \end{bmatrix} \quad k \neq 0$$

Homogeneous
Coordinates

Cartesian
Coordinates

A Point in Cartesian Coordinates is ...
a Ray in Homogeneous Coordinates

3D→2D Projection Matrix



**Camera Frame
posed (R, T) w.r.t.
World Frame**

We have two coordinate systems so our calculation might change a bit depending on which one we are considering

$$\mathbf{x} = \mathbf{K} \begin{bmatrix} \mathbf{I} & \mathbf{0} \end{bmatrix} \mathbf{X}$$

Intrinsic Extrinsic
 Matrix Matrix

\mathbf{O}_w : Center of World Frame
 \mathbf{O}_c : Center of Camera Frame

\mathbf{X} : 3D point in World Coordinates: $(X, Y, Z, 1)$
 \mathbf{x} : 2D Image Coordinates: $(u, v, 1)$, up to scale w

\mathbf{K} : Intrinsic Matrix (3x3)
 \mathbf{R} : Rotation (3x3)
 \mathbf{t} : Translation (3x1)

*Note: different books
use different notation!*

3D→2D Projection Matrix

Extrinsic Assumptions

- Camera == World Frame
- No rotation → $R = I$
- Camera at → $T = [0,0,0]$

Intrinsic Assumptions

- Unit aspect ratio → $f_x = f_y = f$
- Optical center at (0,0) → $u_0 = v_0 = 0$
- No skew → $s = 0$

The optical center is
the hole

K tells us how the internals of the camera
works

$$\mathbf{x} = \mathbf{K} \begin{bmatrix} \text{Rot} & \text{Transl} \\ \mathbf{I} & \mathbf{0} \end{bmatrix} \mathbf{X} \quad \Rightarrow \quad \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} f & 0 & 0 & 0 \\ 0 & f & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

Projection Matrix

Cartesian Coord.
(after dividing with last element)

$$u = x \frac{f}{z}$$

$$v = y \frac{f}{z}$$



3D→2D Projection Matrix

Extrinsic Assumptions

- Camera == World Frame
- No rotation → $R = I$
- Camera at → $T = [0,0,0]$

Intrinsic Assumptions

- Unit aspect ratio → $f_x = f_y = f$
- Optical center at (0,0) → $u_0 = v_0 = 0$
- No skew → $s = 0$

$$\mathbf{x} = \mathbf{K} \underbrace{\begin{bmatrix} \mathbf{I} & \mathbf{0} \end{bmatrix}}_{\text{Projection Matrix}} \mathbf{X} \rightarrow \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \mathbf{K} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

$\mathbf{K} = \begin{bmatrix} f & 0 & u_0 & 0 \\ 0 & f & v_0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix}$

3D→2D Projection Matrix

Extrinsic Assumptions

- Camera == World Frame
- No rotation → $R = I$
- Camera at → $T = [0,0,0]$

Intrinsic Assumptions

- Unit aspect ratio → $f_x = f_y = f$
- Optical center at $(0,0)$ → $u_0 = v_0 = 0$
- No skew → $s = 0$

$$\mathbf{x} = \mathbf{K} \underbrace{\begin{bmatrix} \mathbf{I} & \mathbf{0} \end{bmatrix}}_{\text{Projection Matrix}} \mathbf{X} \rightarrow \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \mathbf{K} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

\mathbf{K}

3D→2D Projection Matrix

Extrinsic Assumptions

- Camera == World Frame
- No rotation → $R = I$
- Camera at → $T = [0,0,0]$

Intrinsic Assumptions

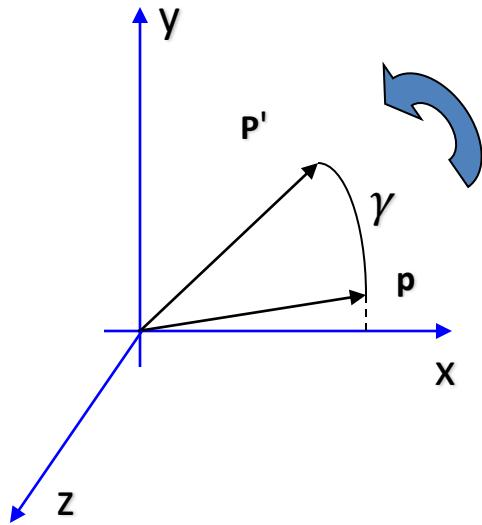
- Unit aspect ratio → $f_x = f_y = f$
- Optical center at (0,0) → $u_0 = v_0 = 0$
- No skew → $s = 0$

The most general
form of Intrinsics K

$$\mathbf{x} = \mathbf{K} \begin{bmatrix} \mathbf{I} & \mathbf{0} \end{bmatrix} \mathbf{X} \quad \Rightarrow \quad \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} f_x & s & u_0 & 0 \\ 0 & f_y & v_0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

Projection
Matrix

3D Rotation of Points



**Rotation around
coordinate axes**
(counter-clockwise)

Each of R_x , R_y , R_z has 1 DoF
(number of free parameters)

Complex Rotations → Multiply R_x & R_y & R_z
(the order matters)

$\alpha \rightarrow$ Angle around axis-x
 $\beta \rightarrow$ Angle around axis-y
 $\gamma \rightarrow$ Angle around axis-z

$$R_x(\alpha) = \begin{bmatrix} 1 & 0 & 0 \\ 0 & \cos \alpha & -\sin \alpha \\ 0 & \sin \alpha & \cos \alpha \end{bmatrix}$$

$$R_y(\beta) = \begin{bmatrix} \cos \beta & 0 & \sin \beta \\ 0 & 1 & 0 \\ -\sin \beta & 0 & \cos \beta \end{bmatrix}$$

$$R_z(\gamma) = \begin{bmatrix} \cos \gamma & -\sin \gamma & 0 \\ \sin \gamma & \cos \gamma & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

Intrinsics & Extrinsics

Camera Frame

Rotated & Translated
w.r.t. World Frame

$$\mathbf{x} = \mathbf{K}[\mathbf{R} \quad \mathbf{t}] \mathbf{X}$$



$$w \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} f_x & s & u_0 \\ 0 & f_y & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} r_{11} & r_{12} & r_{13} & t_x \\ r_{21} & r_{22} & r_{23} & t_y \\ r_{31} & r_{32} & r_{33} & t_z \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

Degrees of Freedom (DoF) →

5

+

6

= 11 DoFs in total

Intrinsics & Extrinsics

What if
no rotation?

$$\mathbf{x} = \mathbf{K}[\mathbf{R} \quad \mathbf{t}] \mathbf{X}$$



$$w \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} f_x & s & u_0 \\ 0 & f_y & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} r_{11} & r_{12} & r_{13} & t_x \\ r_{21} & r_{22} & r_{23} & t_y \\ r_{31} & r_{32} & r_{33} & t_z \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

Intrinsics & Extrinsics

Identity Rotation →
 $R = I$

$$\mathbf{x} = \mathbf{K}[\mathbf{I}] \mathbf{t} \mathbf{X}$$



$$w \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} f_x & 0 & u_0 \\ 0 & f_y & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} t_x \\ t_y \\ t_z \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

What if
 no rotation &
 no translation?



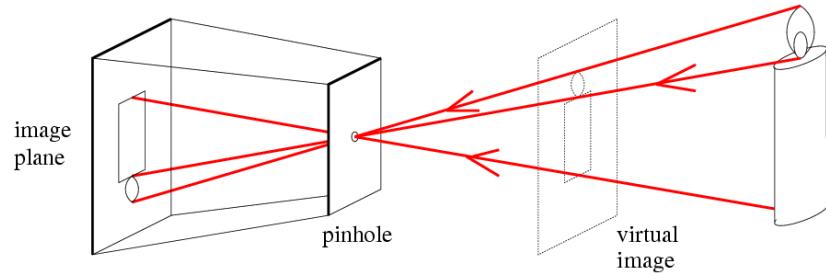
The World & Camera
 Frames coincide



Only the **Intrinsics**
suffice for projection

Quick Summary

- Pinhole Camera model
- Camera Matrices for Projection
- Homogeneous Coordinates



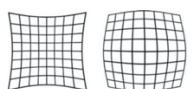
$$\mathbf{x} = \mathbf{K}[\mathbf{R} \quad \mathbf{t}] \mathbf{X}$$

Intrinsics Extrinsics

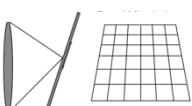
$$(x, y) \Rightarrow \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} \qquad (x, y, z) \Rightarrow \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

Intrinsics – Distorted Image

Front
'fisheye' camera
of FrodoBot
(IRLab, UvA)



Radial distortion

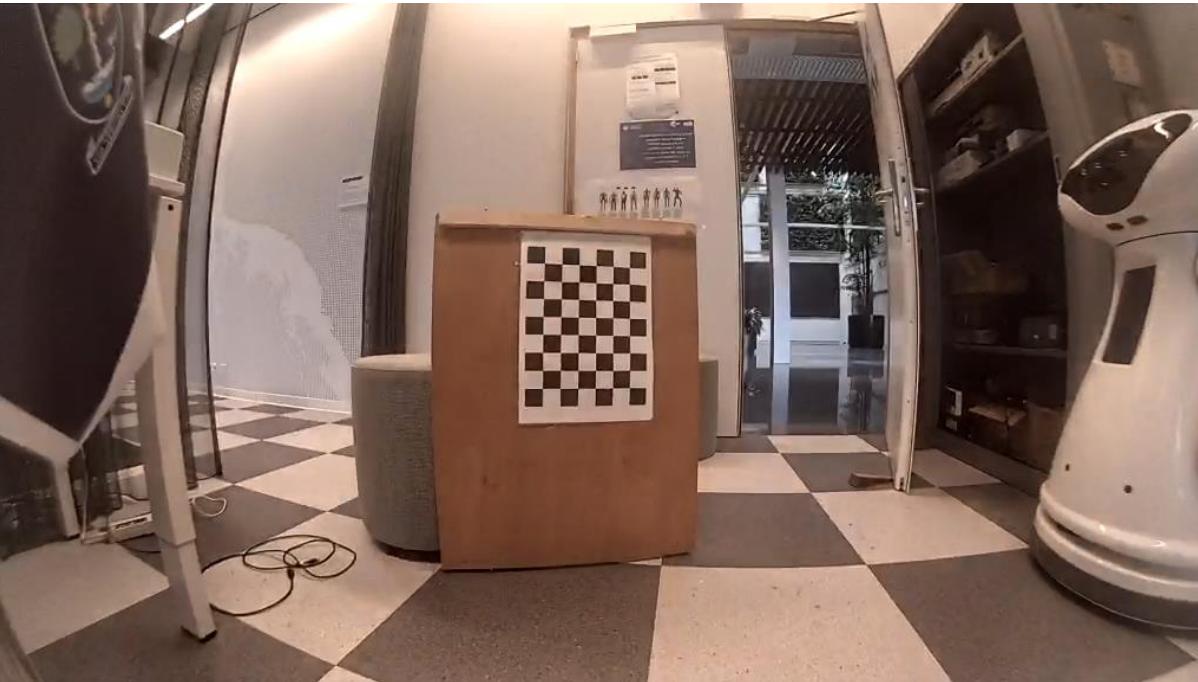


Tangential distortion

Camera is **uncalibrated** 😞

Need to: **Calibrate** camera → estimate **intrinsics**
Undistort image via estimated **intrinsics**

Intrinsics – Distorted Image

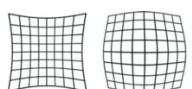


OpenCV tutorial on
Camera Calibration
& image Undistortion

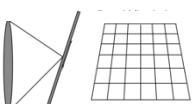
https://docs.opencv.org/4.x/dc/dbb/tutorial_py_calibration.html



Distortion coeffs =
 $(k_1 \quad k_2 \quad p_1 \quad p_2 \quad k_3)$



Radial distortion → $x_{distorted} = x(1 + k_1r^2 + k_2r^4 + k_3r^6)$ $y_{distorted} = y(1 + k_1r^2 + k_2r^4 + k_3r^6)$



Tangential distortion → $x_{distorted} = x + [2p_1xy + p_2(r^2 + 2x^2)]$ $y_{distorted} = y + [p_1(r^2 + 2y^2) + 2p_2xy]$

Intrinsics – Undistorted Image – 0 Elements

CV



```
dist_coeff = [-0.000, 0.0000, 0.00000, -0.00000, -0.00000]
```

Intrinsics – Undistorted Image – 1 Element



```
dist_coeff = [-0.217, 0.0000, 0.00000, -0.00000, -0.00000]
```

Intrinsics – Undistorted Image – 2 Elements

CV



```
dist_coeff = [-0.217, 0.0537, 0.00000, -0.00000, -0.00000]
```

Intrinsics – Undistorted Image – 3 Elements

CV



```
dist_coeff = [-0.217, 0.0537, 0.00185, -0.00000, -0.00000]
```

Intrinsics – Undistorted Image – 4 Elements



```
dist_coeff = [-0.217, 0.0537, 0.00185, -0.00210, -0.00000]
```

Intrinsics – Undistorted Image – 5 Elements



```
dist_coeff = [-0.217, 0.0537, 0.00185, -0.00210, -0.00600]
```

Intrinsics – Undistorted Image



Intrinsics – Undistorted Image



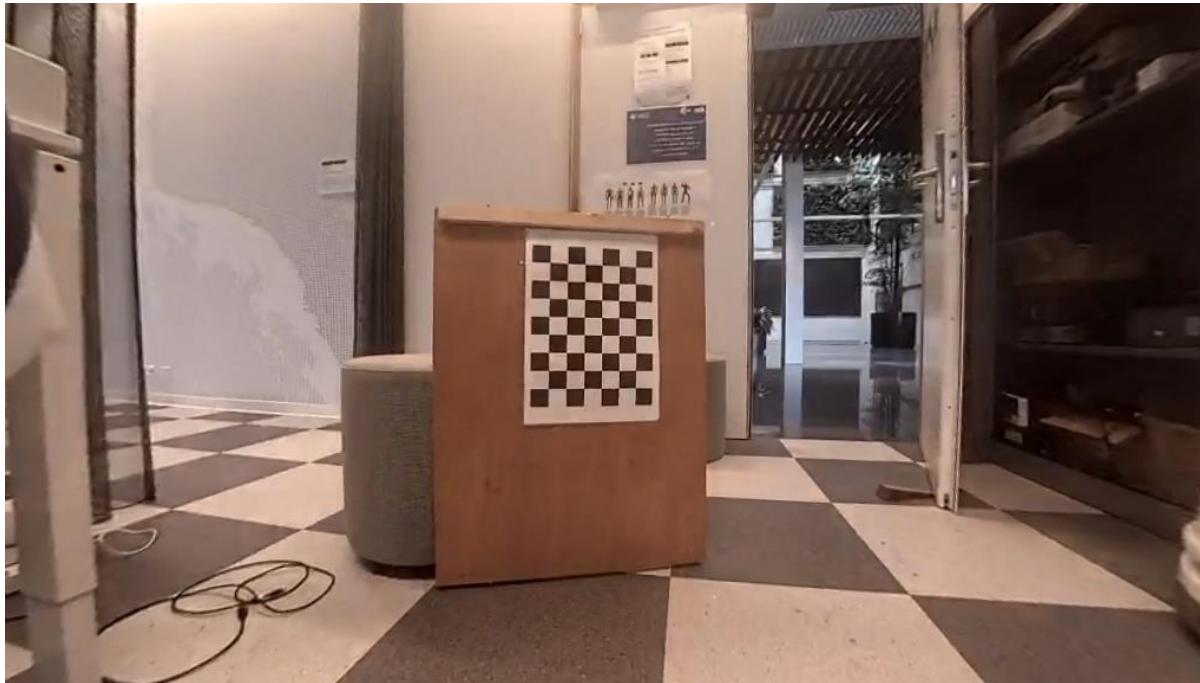
camera matrix =

$$\begin{bmatrix} f_x & 0 & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{bmatrix}$$

h,w = [1024, 576]

f_x, f_y, c_x, c_y = [407.860, 407.866, 533.301, 278.699]

Intrinsics - Undistorted Image



camera matrix =

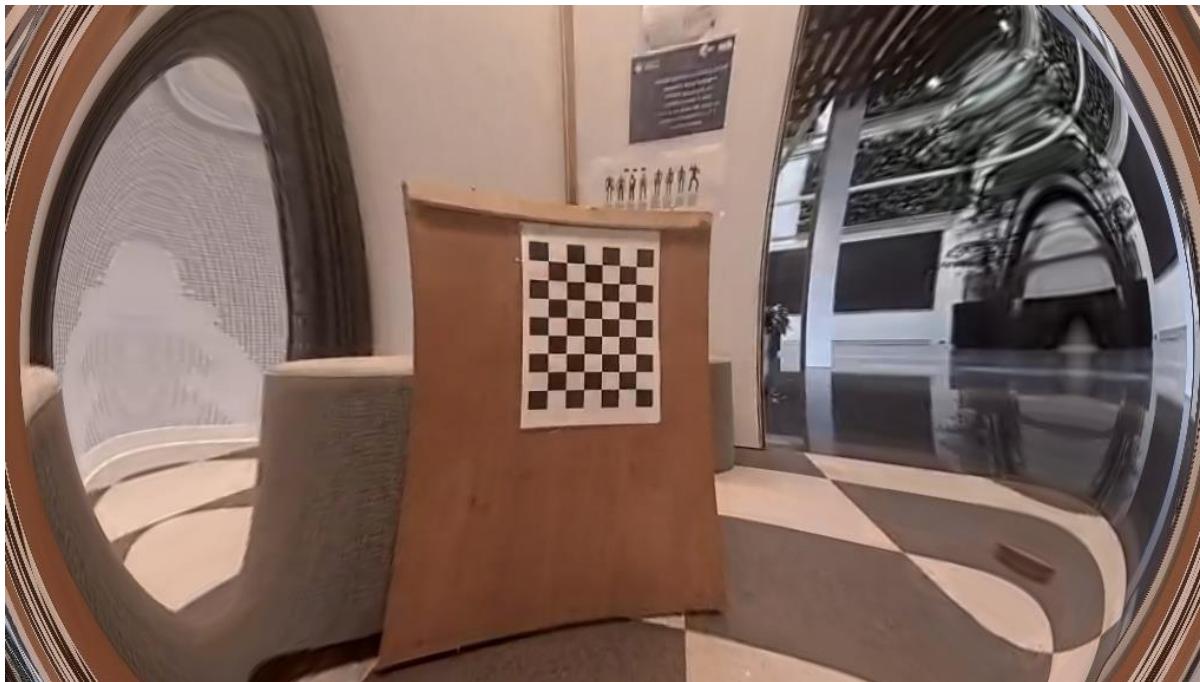
$$\begin{bmatrix} f_x & 0 & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{bmatrix}$$

h,w = [1024, 576]

Mess up: $f_x, f_y, c_x, c_y = [407.860, 407.866, 512.000, 288.000]$

Original: $f_x, f_y, c_x, c_y = [407.860, 407.866, 533.301, 278.699]$

Intrinsics - Undistorted Image



camera matrix =

$$\begin{bmatrix} f_x & 0 & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{bmatrix}$$

h,w = [1024, 576]

Mess up: $f_x, f_y, c_x, c_y = [200.000, 200.000, 512.000, 288.000]$

Original: $f_x, f_y, c_x, c_y = [407.860, 407.866, 533.301, 278.699]$

Intrinsics - Undistorted Image



camera matrix =

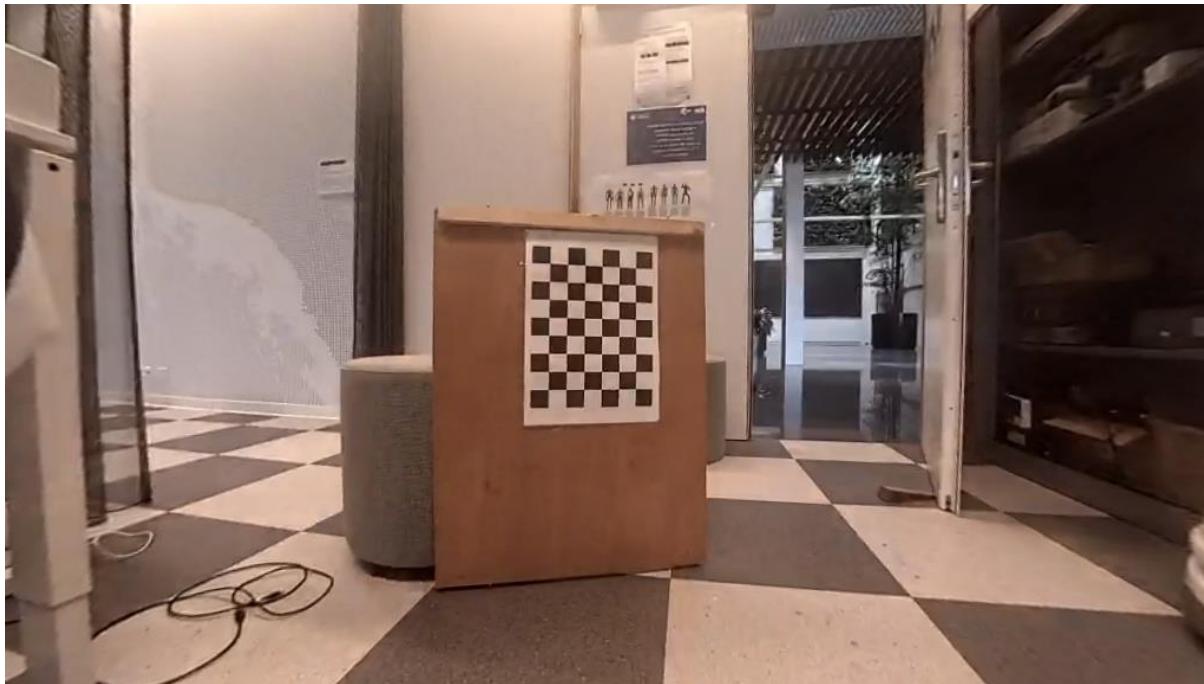
$$\begin{bmatrix} f_x & 0 & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{bmatrix}$$

h,w = [1024, 576]

Mess up: $f_x, f_y, c_x, c_y = [300.000, 300.000, 512.000, 288.000]$

Original: $f_x, f_y, c_x, c_y = [407.860, 407.866, 533.301, 278.699]$

Intrinsics - Undistorted Image



camera matrix =

$$\begin{bmatrix} f_x & 0 & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{bmatrix}$$

h,w = [1024, 576]

Mess up: $f_x, f_y, c_x, c_y = [400.000, 400.000, 512.000, 288.000]$

Original: $f_x, f_y, c_x, c_y = [407.860, 407.866, 533.301, 278.699]$

Intrinsics - Undistorted Image



camera matrix =

$$\begin{bmatrix} f_x & 0 & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{bmatrix}$$

h,w = [1024, 576]

Mess up: $f_x, f_y, c_x, c_y = [600.000, 600.000, 512.000, 288.000]$

Original: $f_x, f_y, c_x, c_y = [407.860, 407.866, 533.301, 278.699]$

Intrinsics - Undistorted Image



Mess up: $f_x, f_y, c_x, c_y = [800.000, 800.000, 512.000, 288.000]$

Original: $f_x, f_y, c_x, c_y = [407.860, 407.866, 533.301, 278.699]$

Intrinsics – Undistorted Image - Skew



Skew

$s = [0\ 0\ 0\ .0]$

camera matrix =

$$\begin{bmatrix} f_x & \boxed{0} & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{bmatrix}$$

$h, w = [1024, 576]$

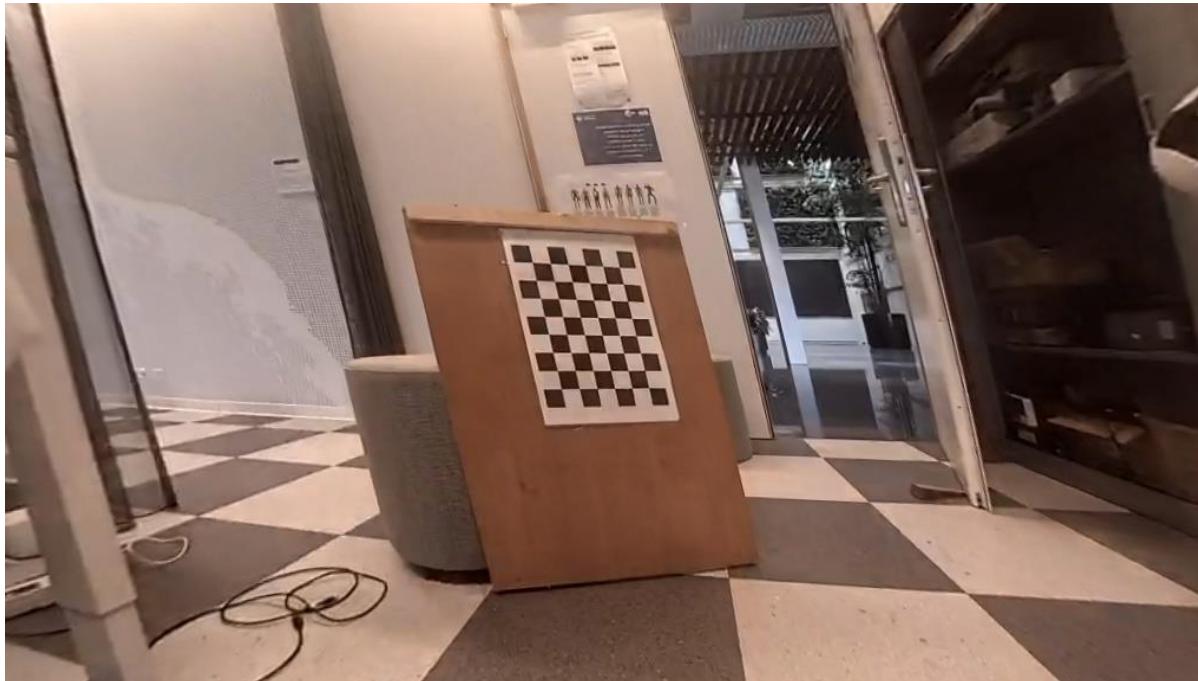
$f_x, f_y, c_x, c_y = [407.860, 407.866, 533.301, 278.699]$

Intrinsics - Undistorted Image - Skew

Mess up:

Skew

$s = [100.0]$



camera matrix =

$$\begin{bmatrix} f_x & \boxed{S} & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{bmatrix}$$

$h, w = [1024, 576]$

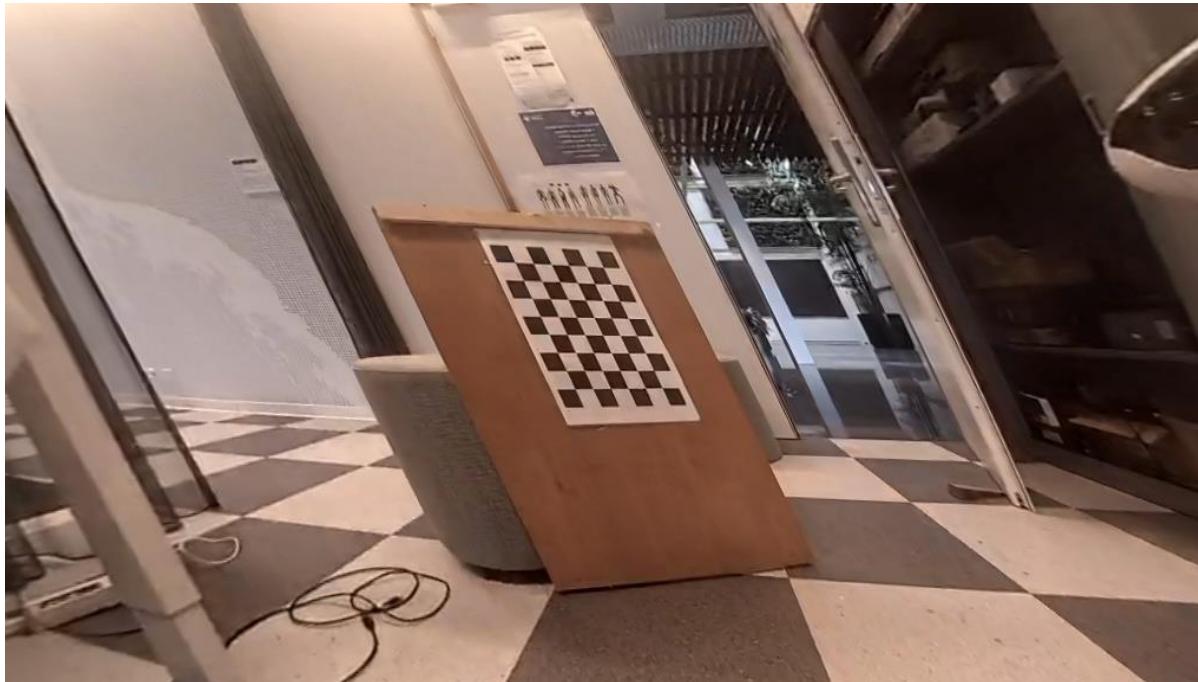
$f_x, f_y, c_x, c_y = [407.860, 407.866, 533.301, 278.699]$

Intrinsics - Undistorted Image - Skew

Mess up:

Skew

$s = [200.0]$



camera matrix =

$$\begin{bmatrix} f_x & \boxed{S} & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{bmatrix}$$

$h, w = [1024, 576]$

$f_x, f_y, c_x, c_y = [407.860, 407.866, 533.301, 278.699]$

Intrinsics - Undistorted Image - Skew

Mess up:

Skew

$s = [400.0]$



camera matrix =

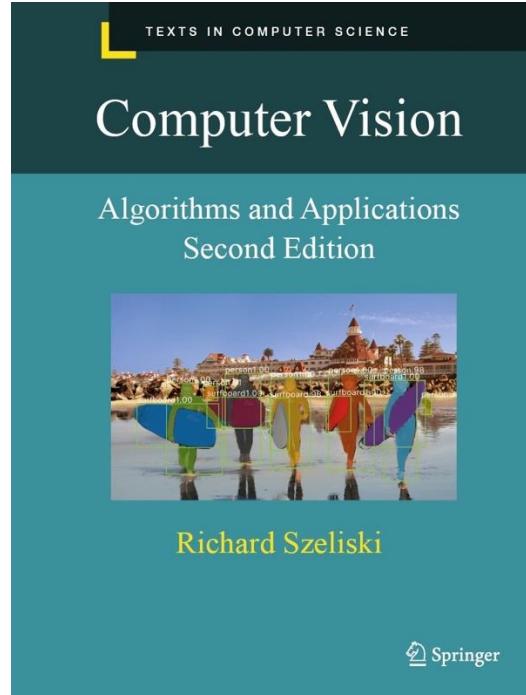
$$\begin{bmatrix} f_x & \boxed{S} & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{bmatrix}$$

$h, w = [1024, 576]$

$f_x, f_y, c_x, c_y = [407.860, 407.866, 533.301, 278.699]$

TextBook – Sections

- 2.1.1
- 2.1.2
- 2.1.4
- 2.2
- 2.3.1
- 2.3.2



Outline

- Camera Model
 - Pinhole Camera
 - Geometric Image Formation
 - Photometric Image Formation
 - Image Representation
- Color
 - Physical & Biological Model
 - Light Source
 - Object
 - Observer
 - Tristimulus Theory
 - Colour Systems

Image Formation

Scenes are
complex



Image Formation

Four main factors influence image intensity values

- Illumination of scene
- Geometry of scene
- Reflectance of visible surfaces
- Camera view & optics

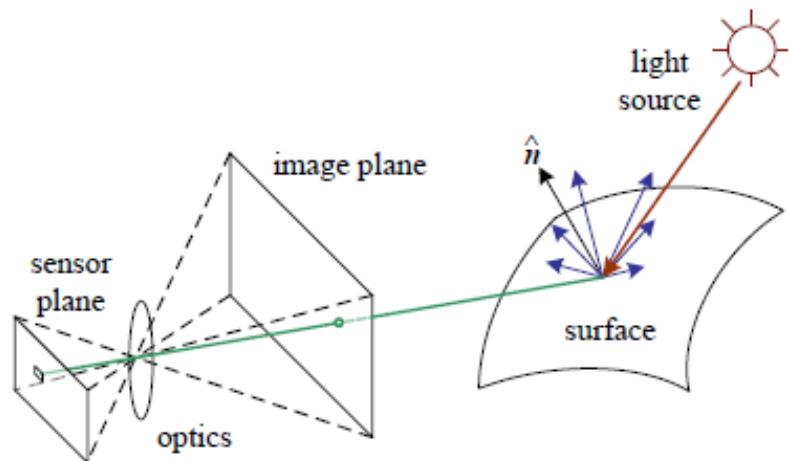


Image Formation – Photometry

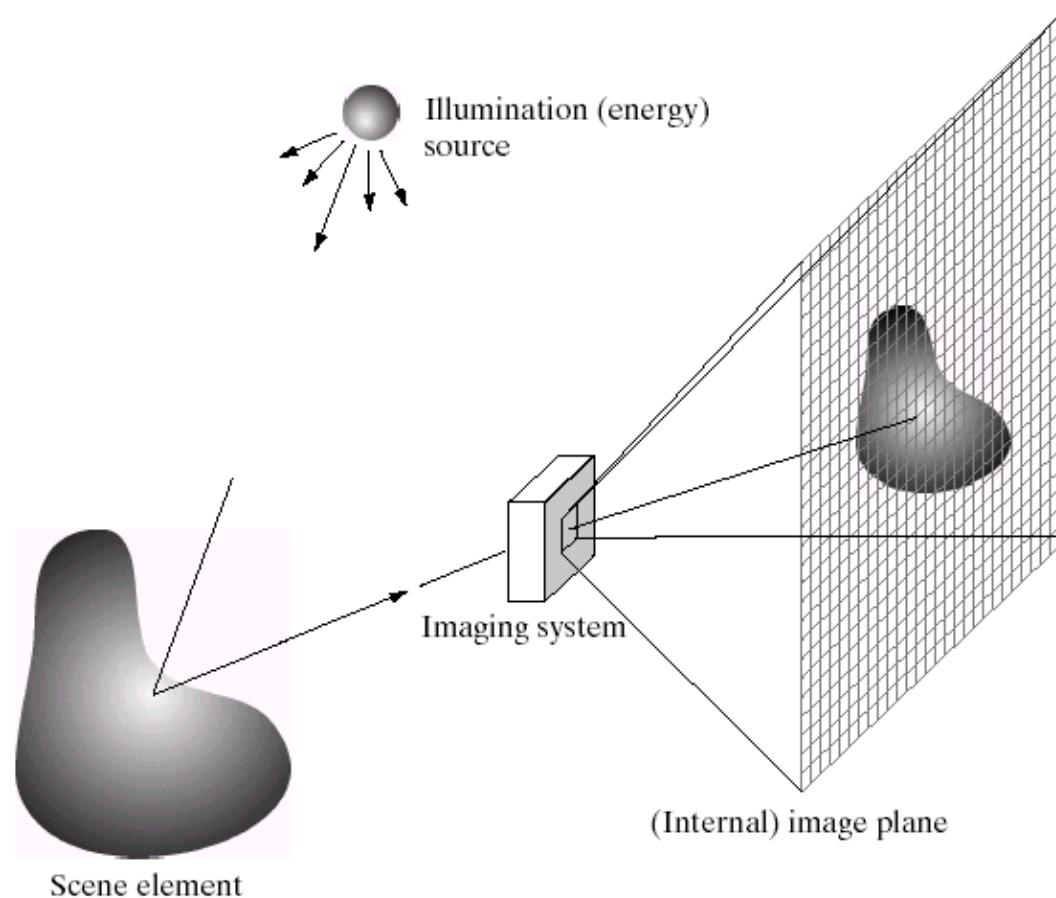
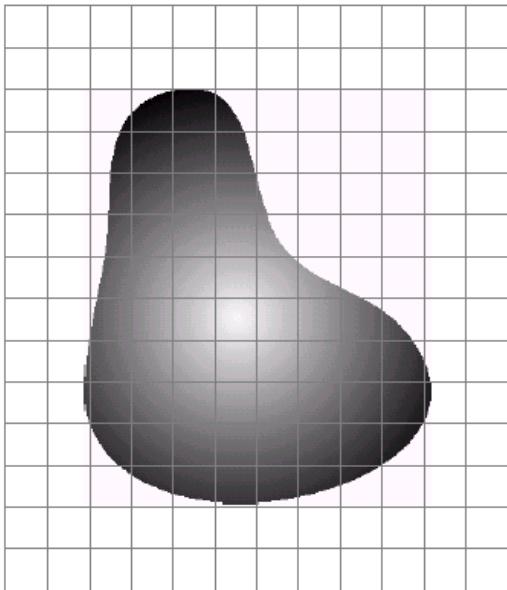
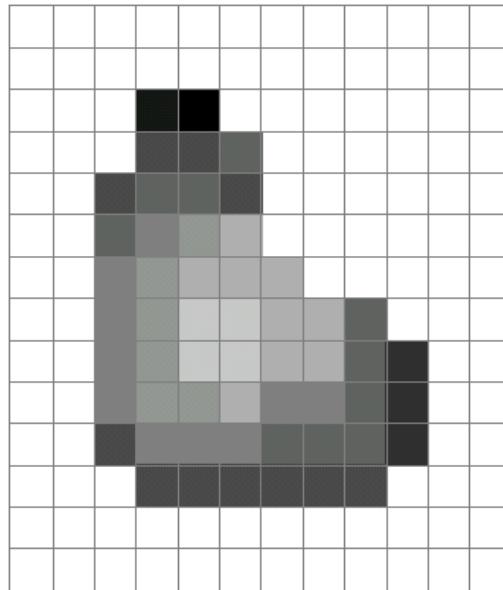


Image Formation – Sensor Array

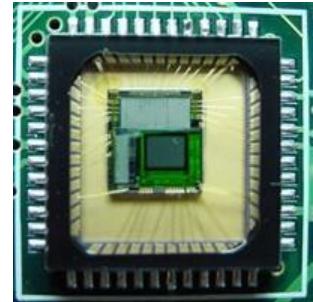


Continuous image
projected onto a
sensor array



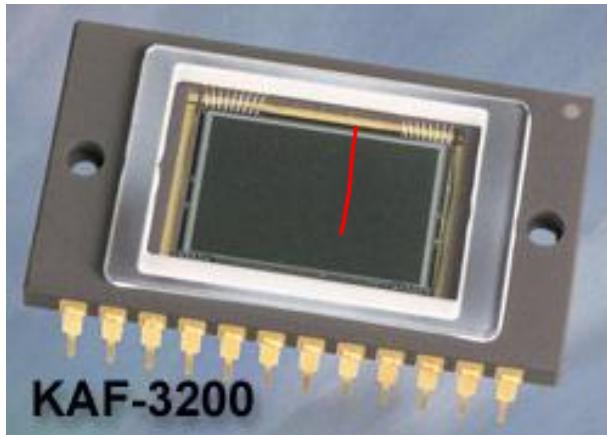
Result of image
sampling &
quantization

→ Spatial position
→ Numerical value



CCD/CMOS sensor

Image Formation – Sensor Array



CCD KAF-3200E from Kodak
(2184 x 1472 pixels)
(Pixel size 6.8 microns²)

Charge-Coupled Device (CCD)

- Converts **continuous** image → **digital** image
- Contains an **array** of **light-sensor** units
- Converts **photons** → **electric charges** accumulated in each sensor unit

Image Formation – Sensor – Imperfections



Low light



Dynamic Range



CCD overflow

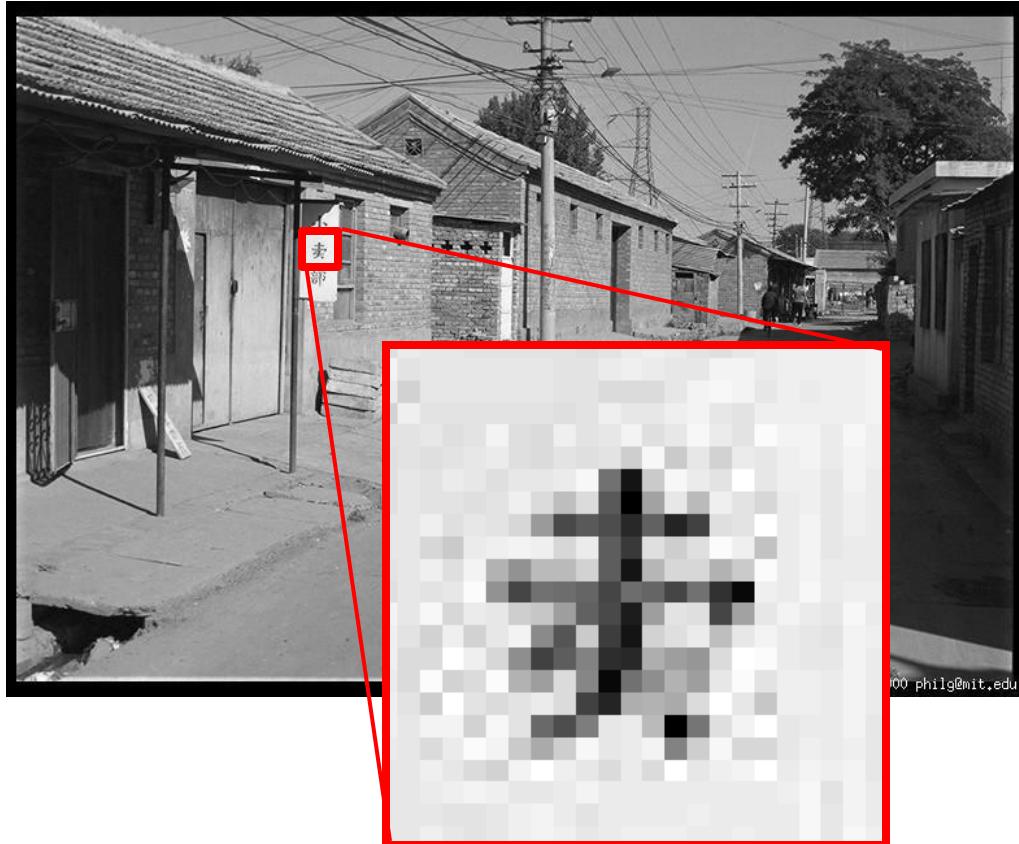


Rolling Shutter

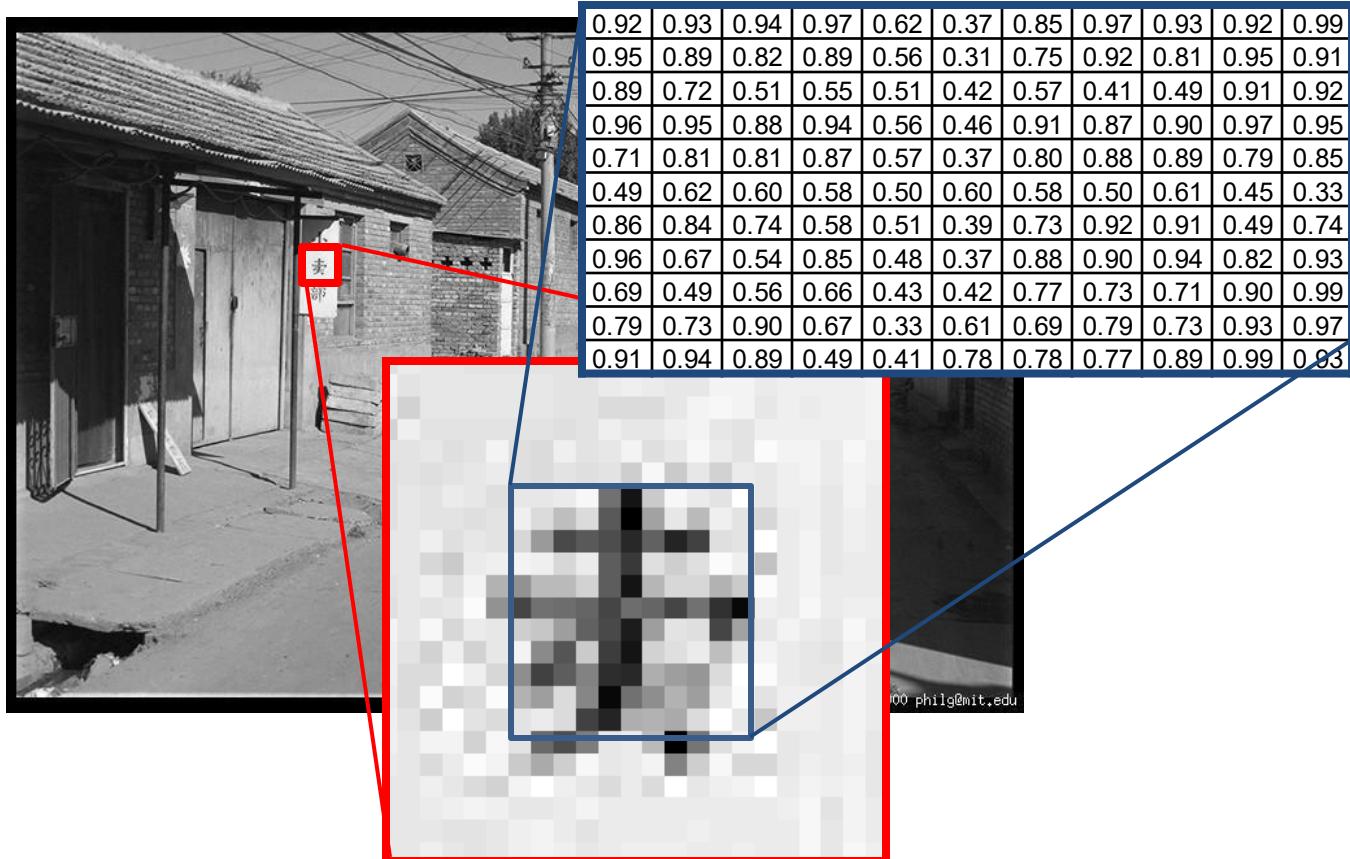
Outline

- Camera Model
 - Pinhole Camera
 - Geometric Image Formation
 - Photometric Image Formation
 - Image Representation
- Color
 - Physical & Biological Model
 - Light Source
 - Object
 - Observer
 - Tristimulus Theory
 - Colour Systems

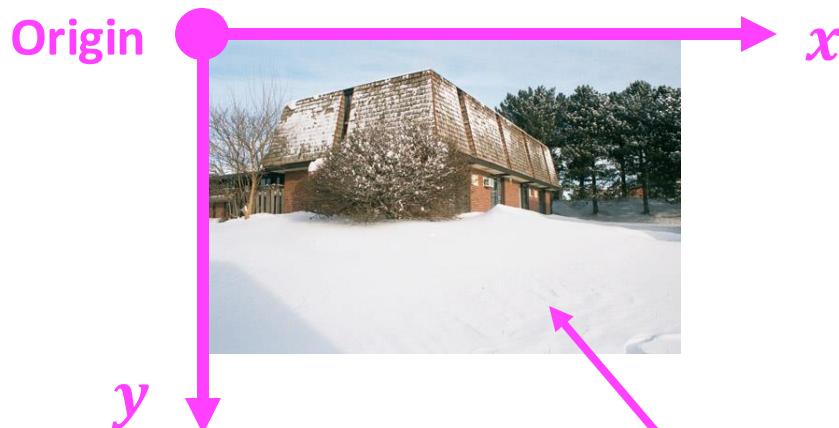
Digital Image – Pixel Matrix



Digital Image – Pixel Matrix

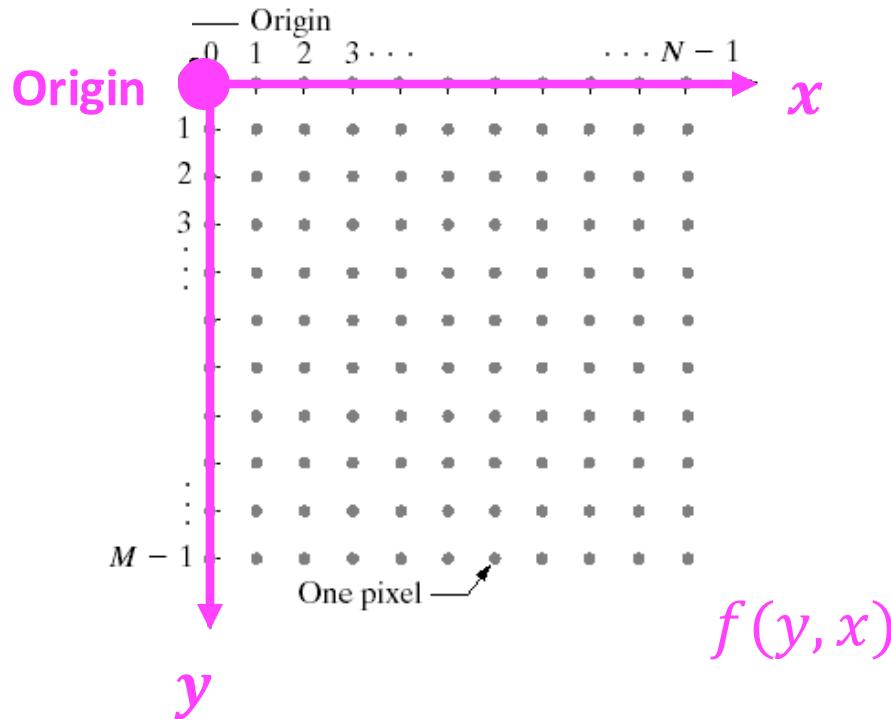


Digital image – Fundamentals



- **Image == a function** of spatial coordinates $f(y, x)$
- **Spatial coordinates:** (x, y) for 2D case – e.g. photograph
 (x, y, t) for video
 (x, y, z) for 3D case – e.g. CT scans
- The function f may represent intensity (for greyscale images) or color (for color images) or other associated values

Digital image – Coordinate System

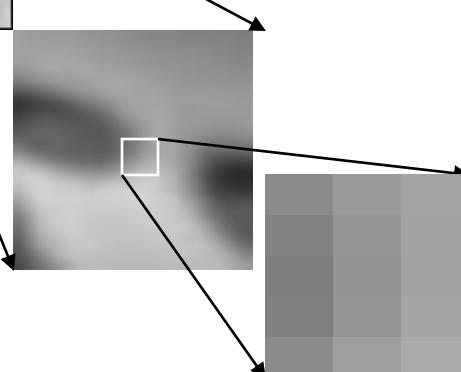
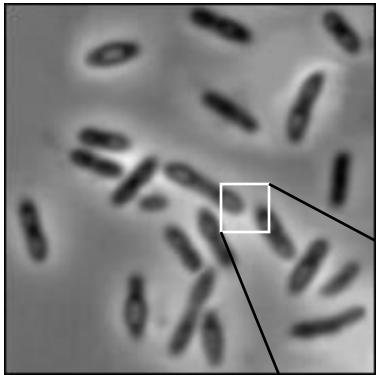


Regular grid of sampling points

Sampling point called a **Pixel**
(PICTure EElement)

Digital Image – Grayscale

Intensity- or
Grayscale-
image



Each pixel → Light intensity in the range of

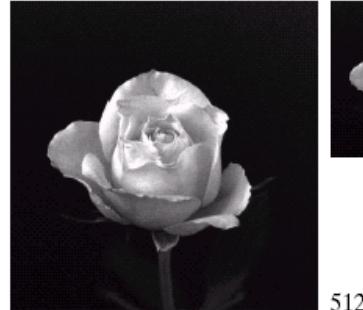
- $[0.0, \dots, 1.0]$ (`float`)
- $[0, \dots, 255]$ (`uint8`)

as [black, ..., white]

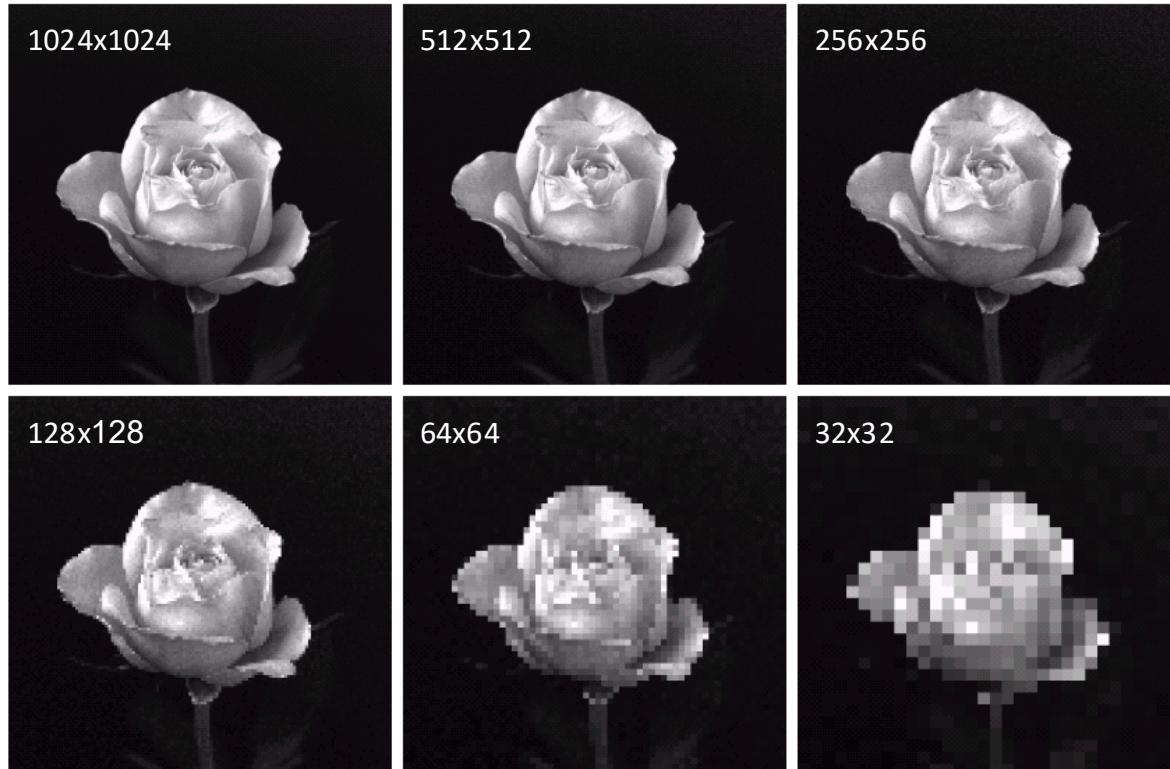
Gray scale values

| | | | |
|----|----|----|----|
| 10 | 10 | 16 | 28 |
| 9 | 6 | 26 | 37 |
| 15 | 25 | 13 | 22 |
| 32 | 15 | 87 | 39 |

Digital Image – Spatial Resolution



Digital Image – Spatial Resolution



Images from
Rafael C. Gonzalez and Richard E. Wood,
Digital Image Processing, 2nd Edition

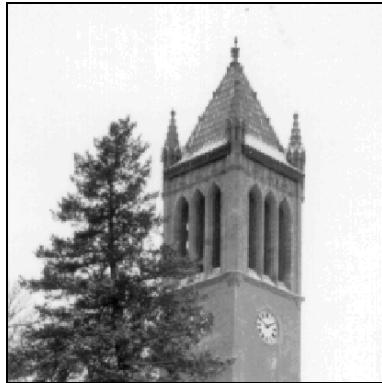
Digital Image – Quantization Levels



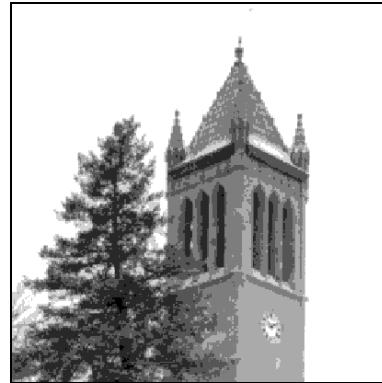
256 levels (8-bit)



128 levels (7-bit)

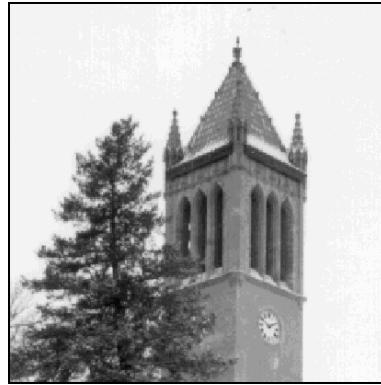


64 levels (6-bit)

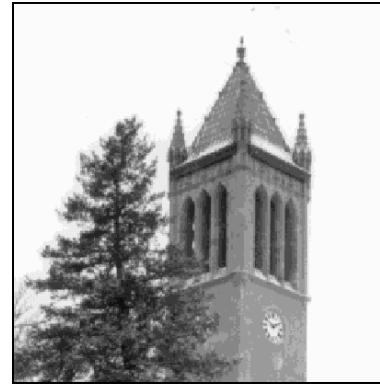


32 levels (5-bit)

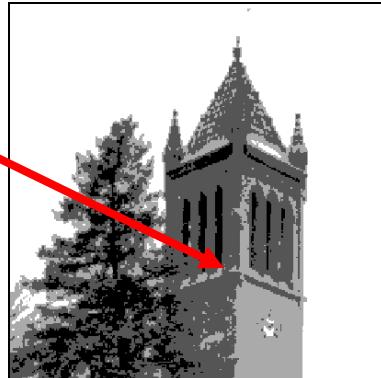
Digital Image – Quantization Levels



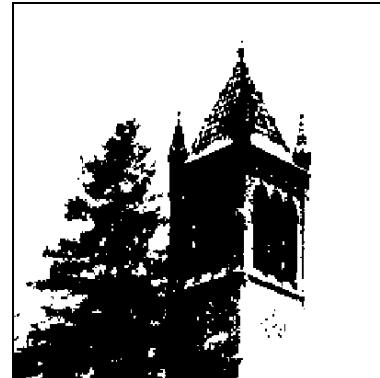
16 levels (4-bit)



8 levels (3-bit)



4 levels (2-bit)

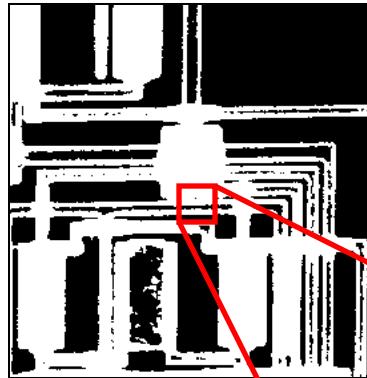


2 levels (1-bit)

In this image
it is easy to see
a 'false' contour

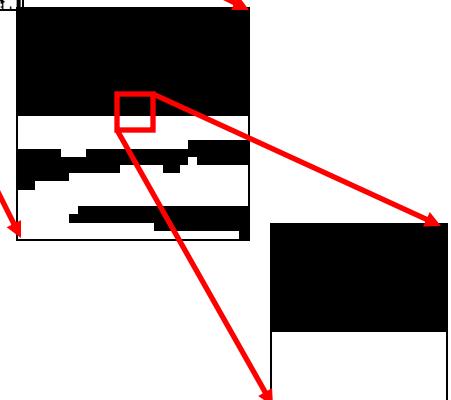
Digital Image – Binary Image

Binary- or
Black-n-White-
image



Each pixel → Contains 1 bit:

- 0 → black
- 1 → white

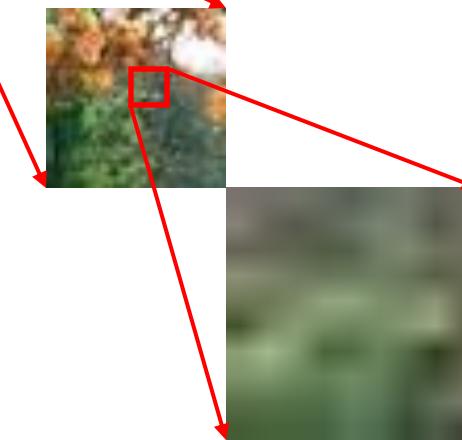


Binary data

| | | | |
|---|---|---|---|
| 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 |
| 1 | 1 | 1 | 1 |
| 1 | 1 | 1 | 1 |

Digital Image – Color

Color (RGB)
image



Each
pixel → Contains a 3-element vector:

- $[0, \dots, 255] \rightarrow$ Red
- $[0, \dots, 255] \rightarrow$ Green
- $[0, \dots, 255] \rightarrow$ Blue

RGB components

| | | | |
|----|----|----|----|
| 10 | 10 | 16 | 28 |
| 9 | 65 | 70 | 56 |
| 15 | 32 | 99 | 70 |
| 32 | 21 | 56 | 43 |
| | 54 | 78 | |
| | 85 | 60 | 90 |
| | 32 | 96 | 67 |
| | 54 | 85 | 43 |
| | 85 | 85 | 92 |
| | 32 | 32 | 99 |
| | 65 | 65 | 87 |
| | 54 | 87 | 99 |

Digital Image – Matlab

row ↓

column →

| | | | | | | | | | | |
|------|------|------|------|------|------|------|------|------|------|------|
| 0.92 | 0.93 | 0.94 | 0.97 | 0.62 | 0.37 | 0.85 | 0.97 | 0.93 | 0.92 | 0.99 |
| 0.95 | 0.89 | 0.82 | 0.89 | 0.56 | 0.31 | 0.75 | 0.92 | 0.81 | 0.95 | 0.91 |
| 0.89 | 0.72 | 0.51 | 0.55 | 0.51 | 0.42 | 0.57 | 0.41 | 0.49 | 0.91 | 0.92 |
| 0.96 | 0.95 | 0.88 | 0.94 | 0.56 | 0.46 | 0.91 | 0.87 | 0.90 | 0.97 | 0.95 |
| 0.71 | 0.81 | 0.81 | 0.87 | 0.57 | 0.37 | 0.80 | 0.88 | 0.89 | 0.79 | 0.85 |
| 0.49 | 0.62 | 0.60 | 0.58 | 0.50 | 0.60 | 0.58 | 0.50 | 0.61 | 0.45 | 0.33 |
| 0.86 | 0.84 | 0.74 | 0.58 | 0.51 | 0.39 | 0.73 | 0.92 | 0.91 | 0.49 | 0.74 |
| 0.96 | 0.67 | 0.54 | 0.85 | 0.48 | 0.37 | 0.88 | 0.90 | 0.94 | 0.82 | 0.93 |
| 0.69 | 0.49 | 0.56 | 0.66 | 0.43 | 0.42 | 0.77 | 0.73 | 0.71 | 0.90 | 0.99 |
| 0.79 | 0.73 | 0.90 | 0.67 | 0.33 | 0.61 | 0.69 | 0.79 | 0.73 | 0.93 | 0.97 |
| 0.91 | 0.94 | 0.89 | 0.49 | 0.41 | 0.78 | 0.78 | 0.77 | 0.89 | 0.99 | 0.93 |
| | | | | | 0.79 | 0.73 | 0.90 | 0.67 | 0.33 | 0.61 |
| | | | | | 0.91 | 0.94 | 0.89 | 0.49 | 0.41 | 0.78 |
| | | | | | | 0.79 | 0.73 | 0.90 | 0.67 | 0.33 |
| | | | | | | 0.91 | 0.94 | 0.89 | 0.49 | 0.41 |
| | | | | | | | 0.79 | 0.73 | 0.90 | 0.67 |
| | | | | | | | 0.91 | 0.94 | 0.89 | 0.49 |
| | | | | | | | | 0.79 | 0.73 | 0.90 |
| | | | | | | | | 0.91 | 0.94 | 0.89 |
| | | | | | | | | | 0.79 | 0.73 |
| | | | | | | | | | 0.91 | 0.94 |

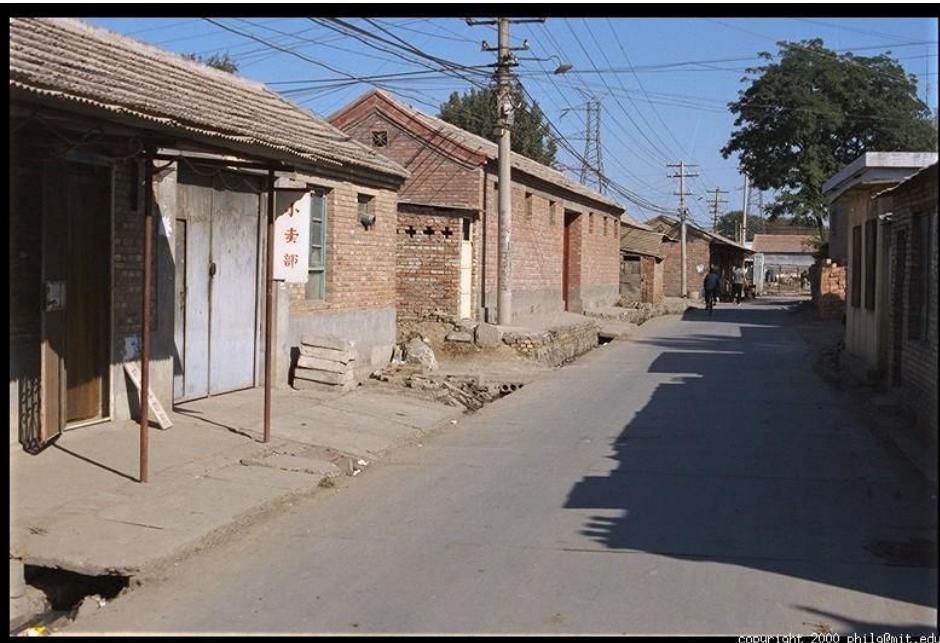


In OpenCV
the channel
order is **BGR**!

- Suppose we have a NxM RGB image called “im”
 - im(1,1,1) = top-left pixel value in R-channel
 - im(y, x, b) = y pixels down, x pixels to right in the bth channel
- imread(filename) returns a uint8 image (values 0 to 255)
 - For processing: Convert to double format (values 0 to 1) with im2double() → important!

Digital Image

Shape information mostly contained
in **intensity** (grayscale) image



Color image



Grayscale image

Outline

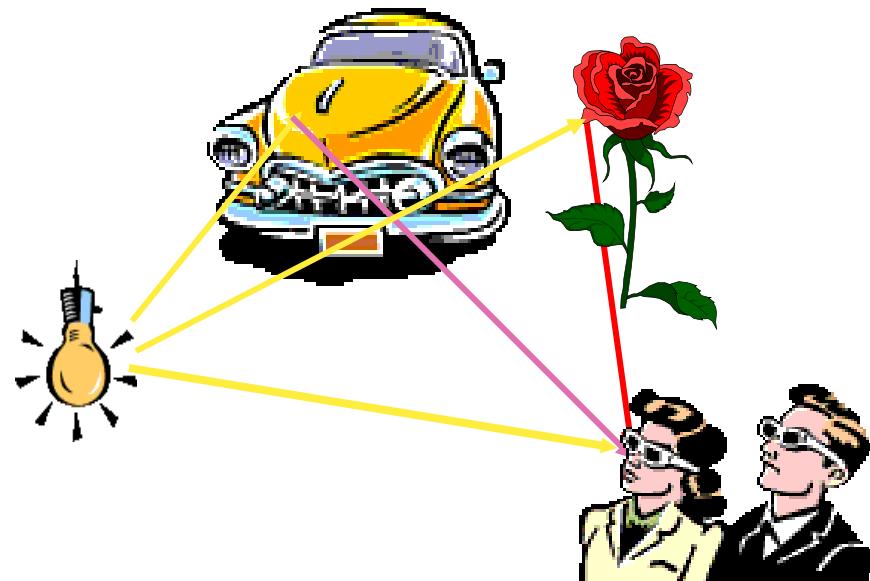
- Camera Model
 - Pinhole Camera
 - Geometric Image Formation
 - Photometric Image Formation
 - Image Representation
- Color
 - Physical & Biological Model
 - Light Source
 - Object
 - Observer
 - Tristimulus Theory
 - Colour Systems

What makes for an image?

Light source

Object (s)

Observer / Sensor

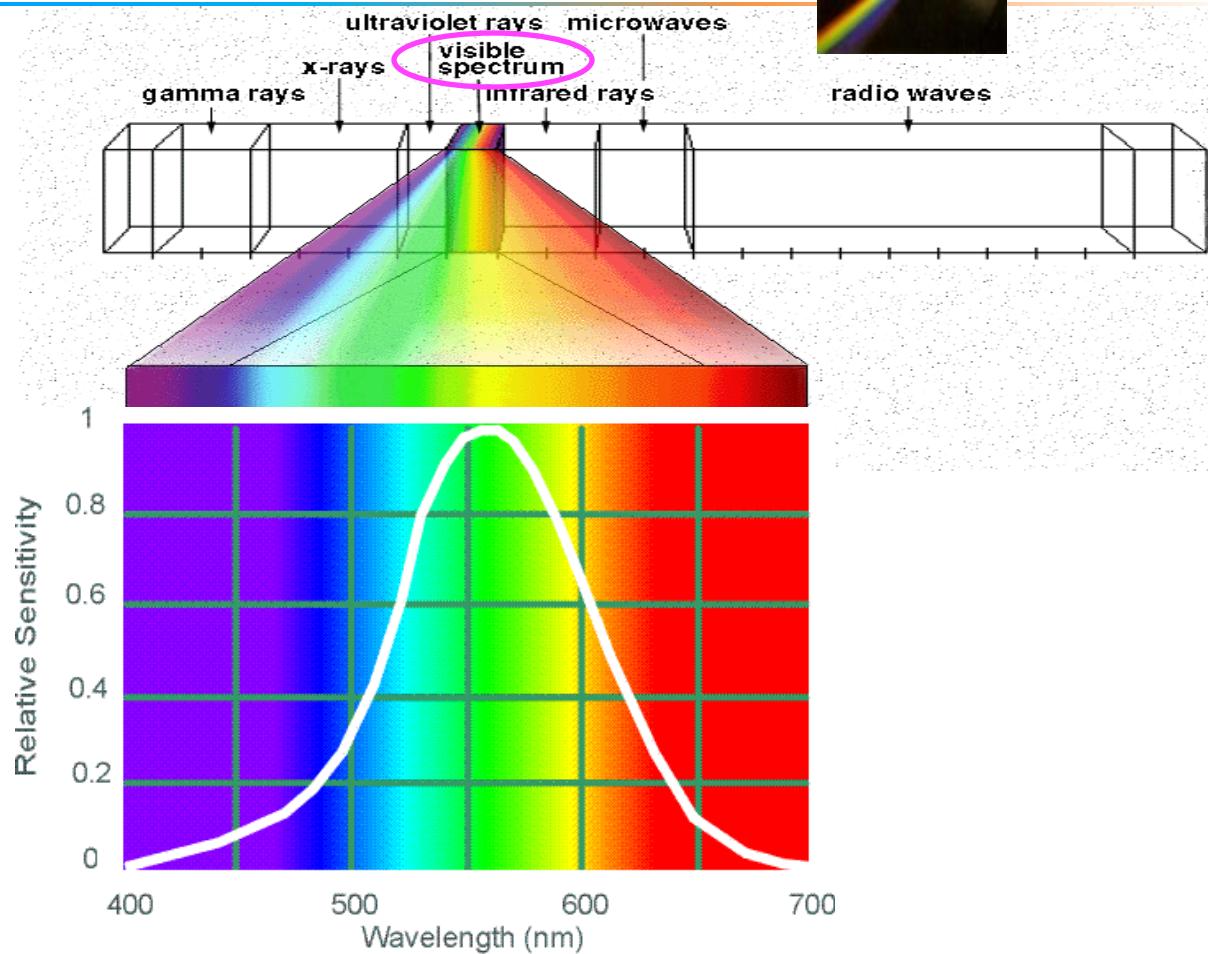


Humans – Sensitivity Function

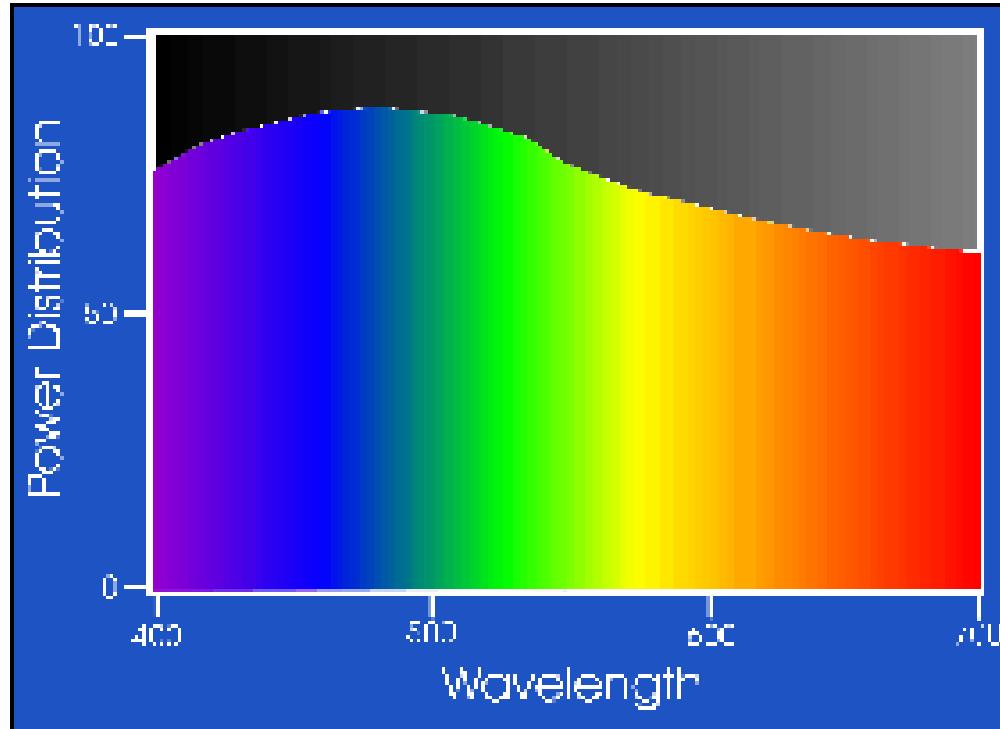


Most sensitive at
middle wavelengths

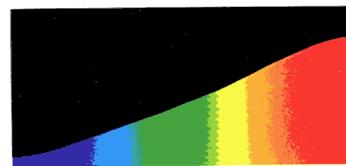
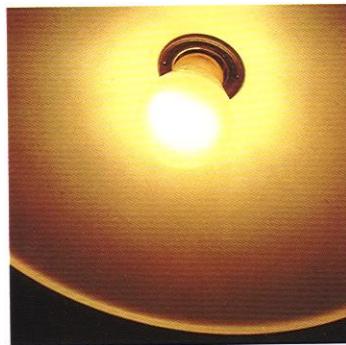
Sensitivity falls off
for long & short
wavelengths



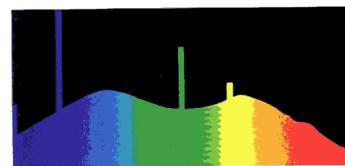
Light Sources – Spectral Power Distribution



Light Sources – Types



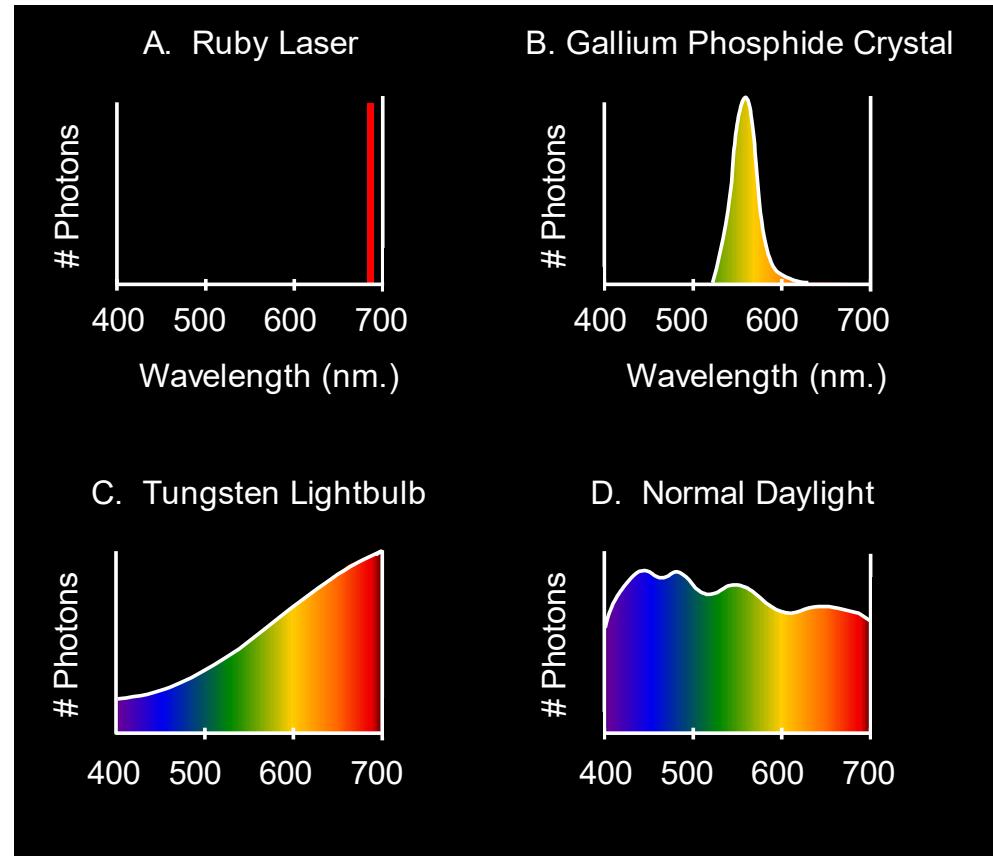
Incandescent
lamp



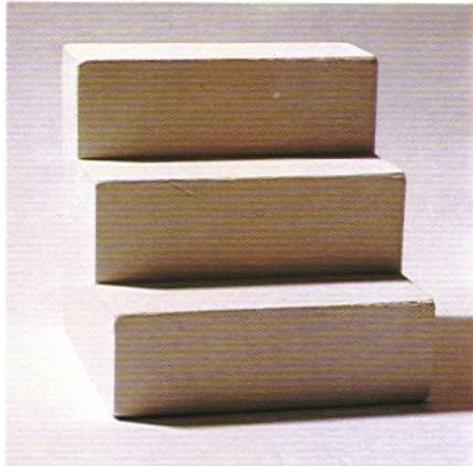
Fluorescent
lamp

Light Sources – Types

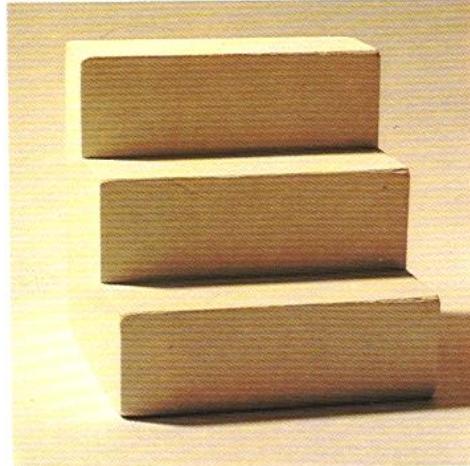
Spectra examples of various light sources



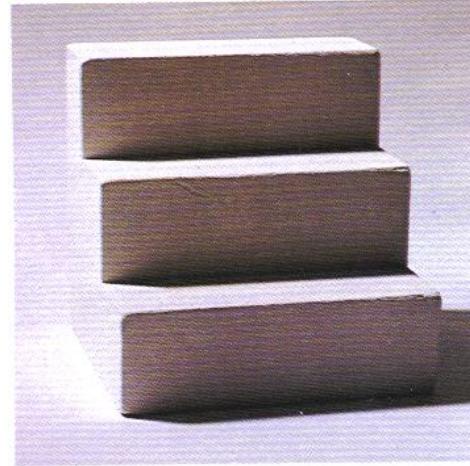
Light Sources – Influence



Average
daylight

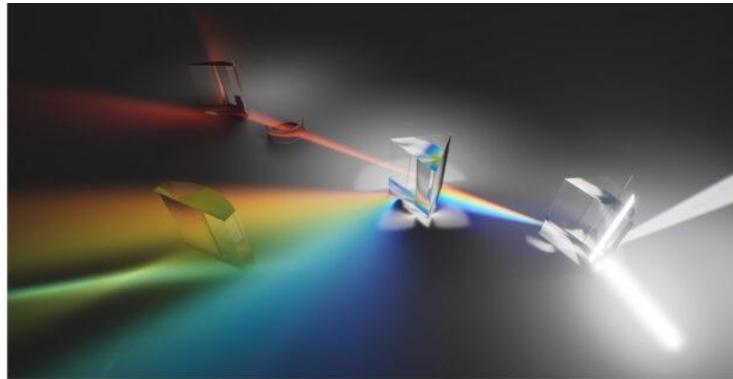


Incandescent
lamp



Fluorescent
lamp

Light Sources – Computer Graphics



Outline

- Camera Model
 - Pinhole Camera
 - Geometric Image Formation
 - Photometric Image Formation
 - Image Representation
- Color
 - Physical & Biological Model
 - Light Source
 - Object
 - Observer
 - Tristimulus Theory
 - Colour Systems

Object Colors

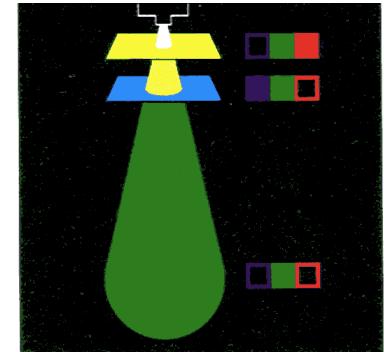
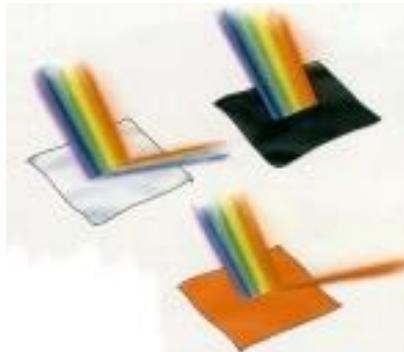
Materials:

Transparent

Opaque

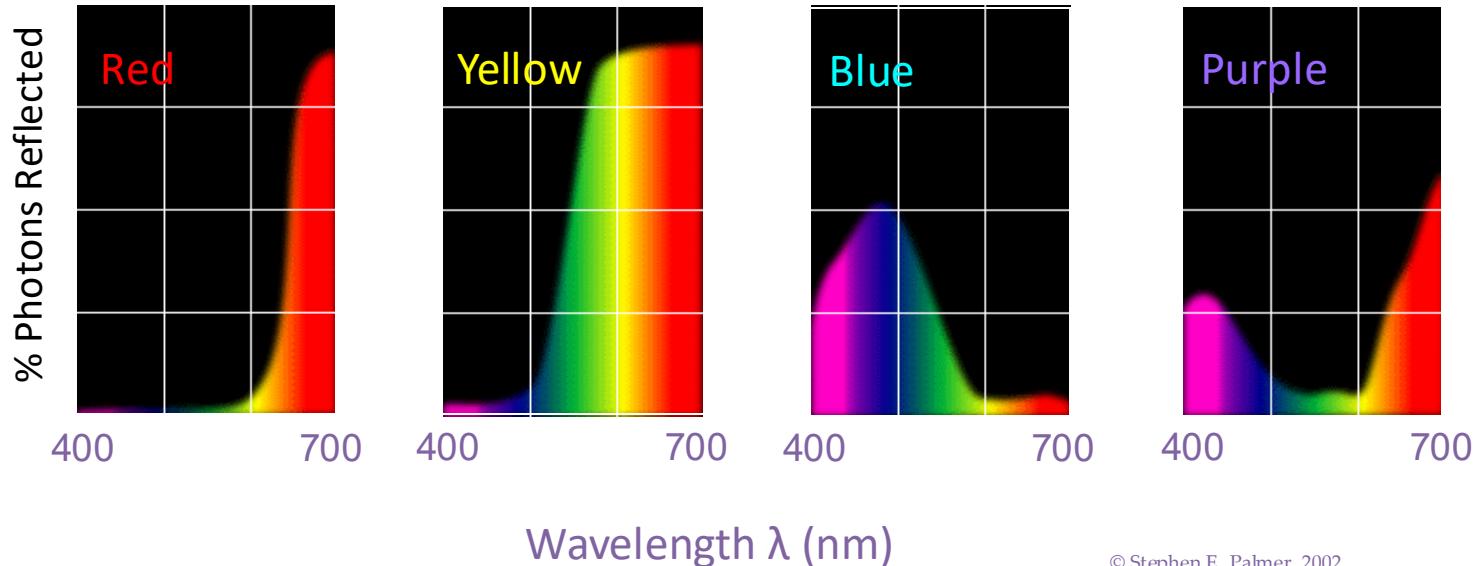
Spectral Reflectance $\rho(\lambda)$

Wavelength
of light



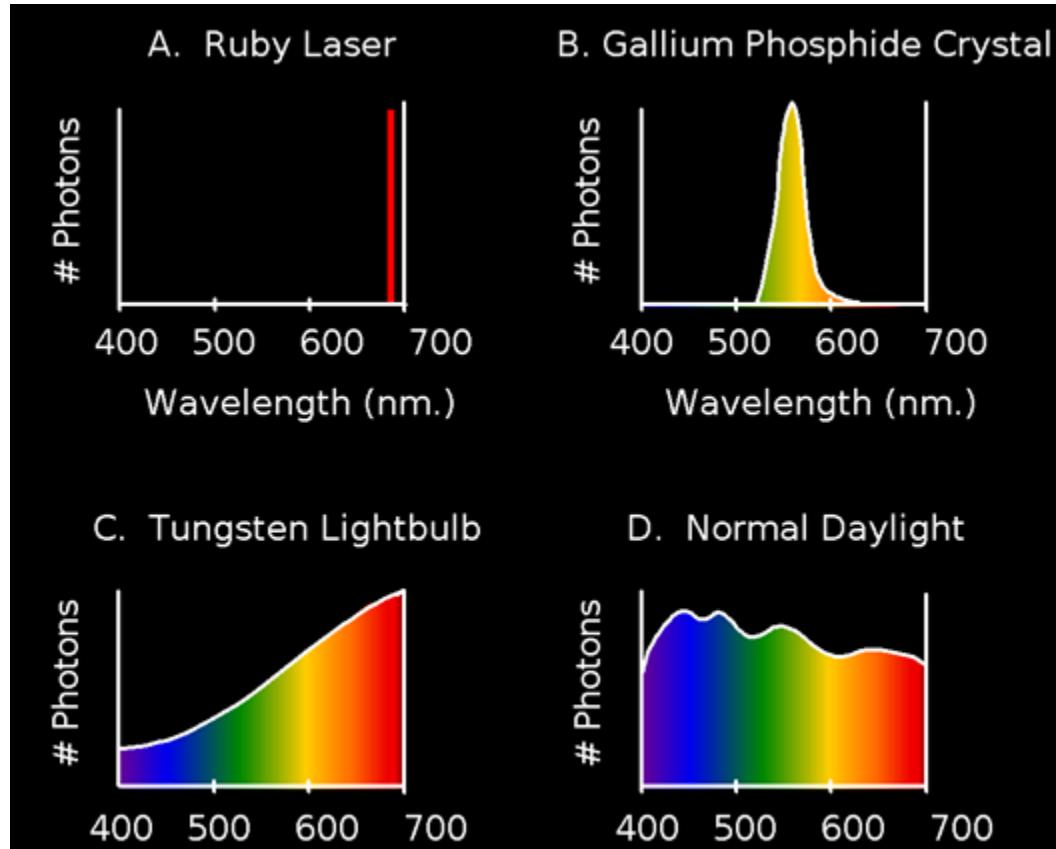
Object Colors

Examples of reflectance spectra of surfaces



Object Colors

Light sources
are diverse!



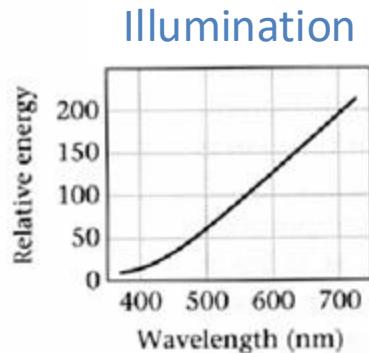
Object Colors

What color is the object?

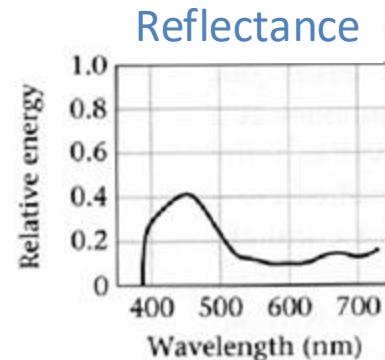


Depends on **both**:

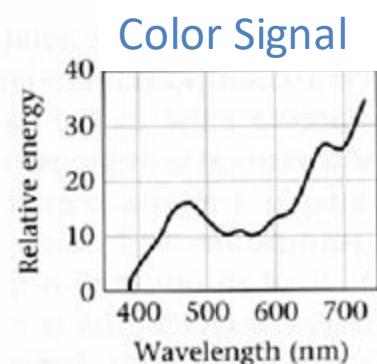
- Incident light – **Illumination**
- Object's **Reflectance**



• *



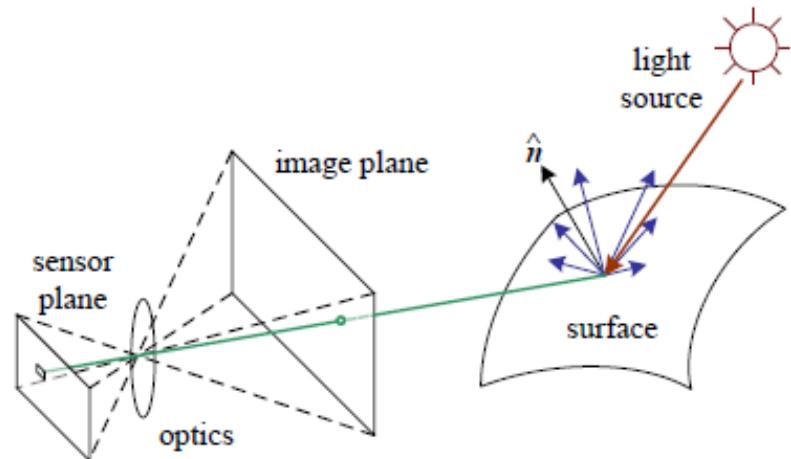
=



Object Colors

Four main factors influence image intensity values

- Illumination of scene
- Geometry of scene
- Reflectance of visible surfaces
- Camera view & optics



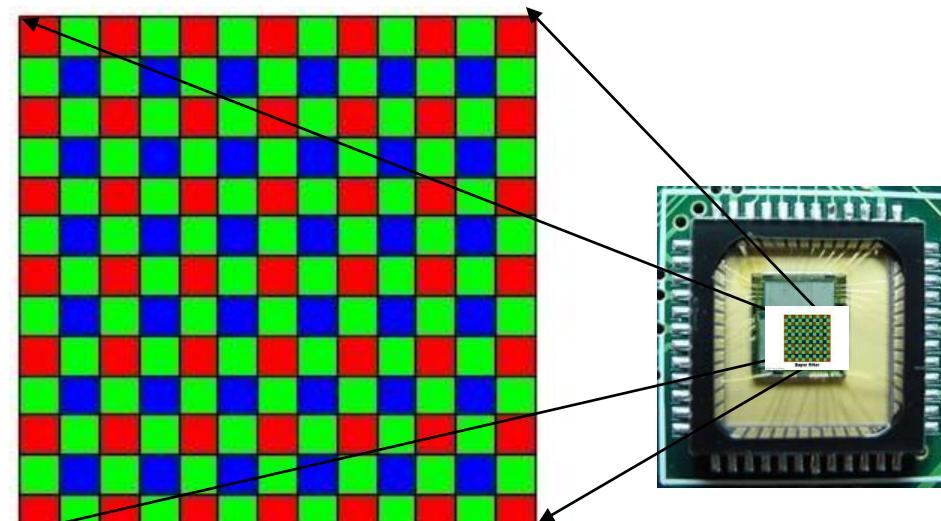
Outline

- Camera Model
 - Pinhole Camera
 - Geometric Image Formation
 - Photometric Image Formation
 - Image Representation
- Color
 - Physical & Biological Model
 - Light Source
 - Object
 - Observer
 - Tristimulus Theory
 - Colour Systems

Color images – How to form?

Bayer filter

- Green fills in half of checkerboard
- Red and Blue fill the rest.



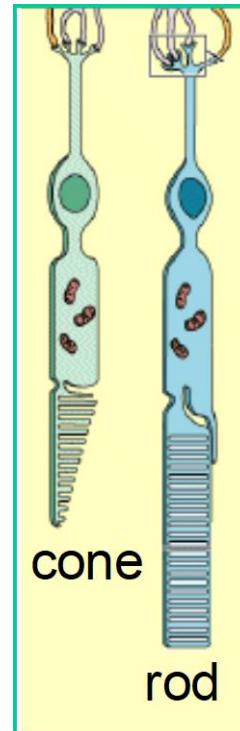
Bayer filter



Rods & Cones – Light-Sensitive Cells

Cones

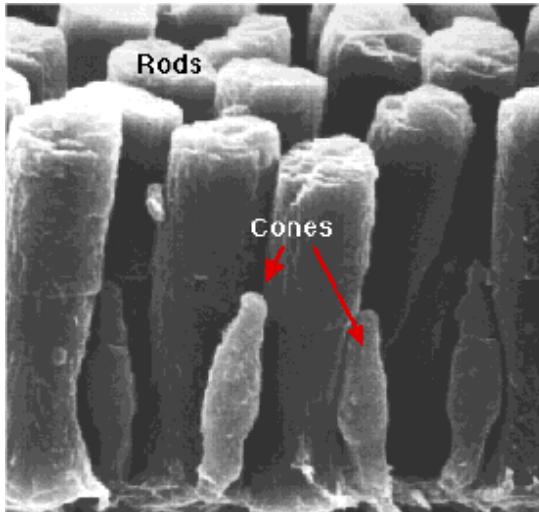
- cone-shaped
- less sensitive
- operate in strong light
- color vision



Rods

- rod-shaped
- highly sensitive
- operate at night
- gray-scale vision

Rods & Cones – Attributes



Rods:

- **120 million** rods in retina
- **Low-resolution peripheral** vision
- Sense **B/W brightness** in low illumination
- 1000X more light-sensitive than Cones
- Short wave-length sensitive

Cones:

- **6-7 million** cones in the retina
- High-resolution vision @ **Fovea**
- Sense **Colors**
- **3 types: 64% red, 32% green, 2% blue)**
- Sensitive to **any combination** of 3 colors

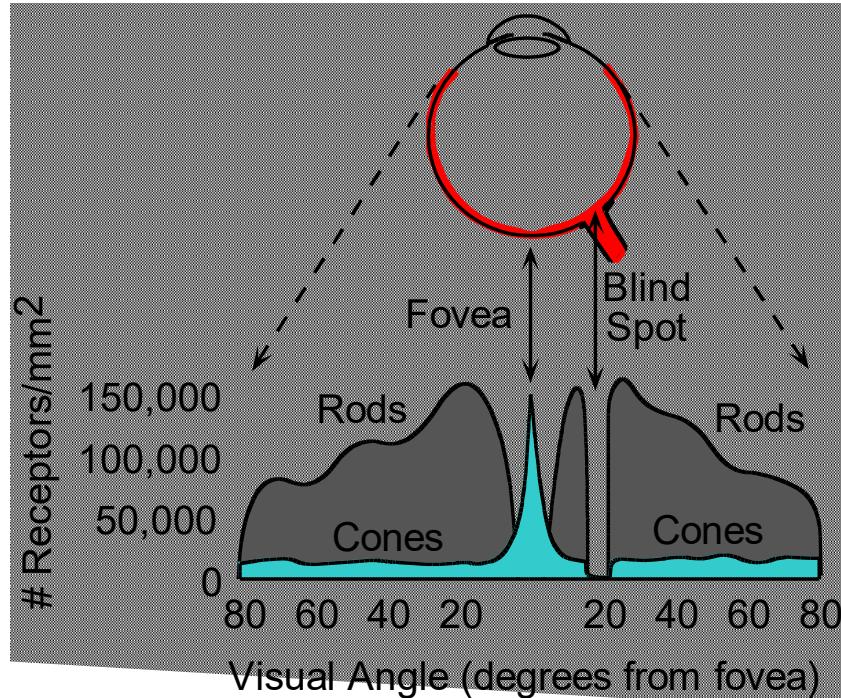
Rods & Cones – Distribution

Night Sky:

Why do we 'see' more stars appearing 'off-center' in our field of view?

Averted vision

Peripheral vision-areas on retina have more rod cells that are sensitive in low-light conditions

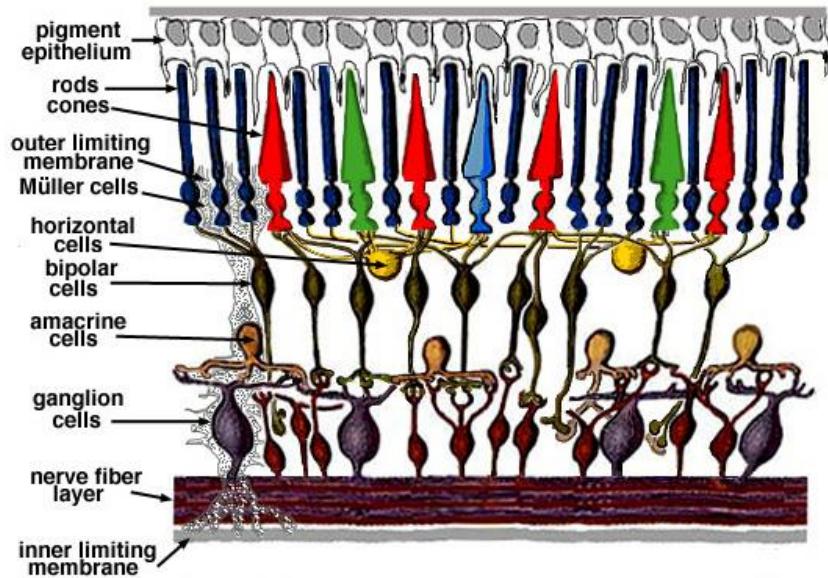


http://en.wikipedia.org/wiki/Averted_vision

Slide source: James Hays

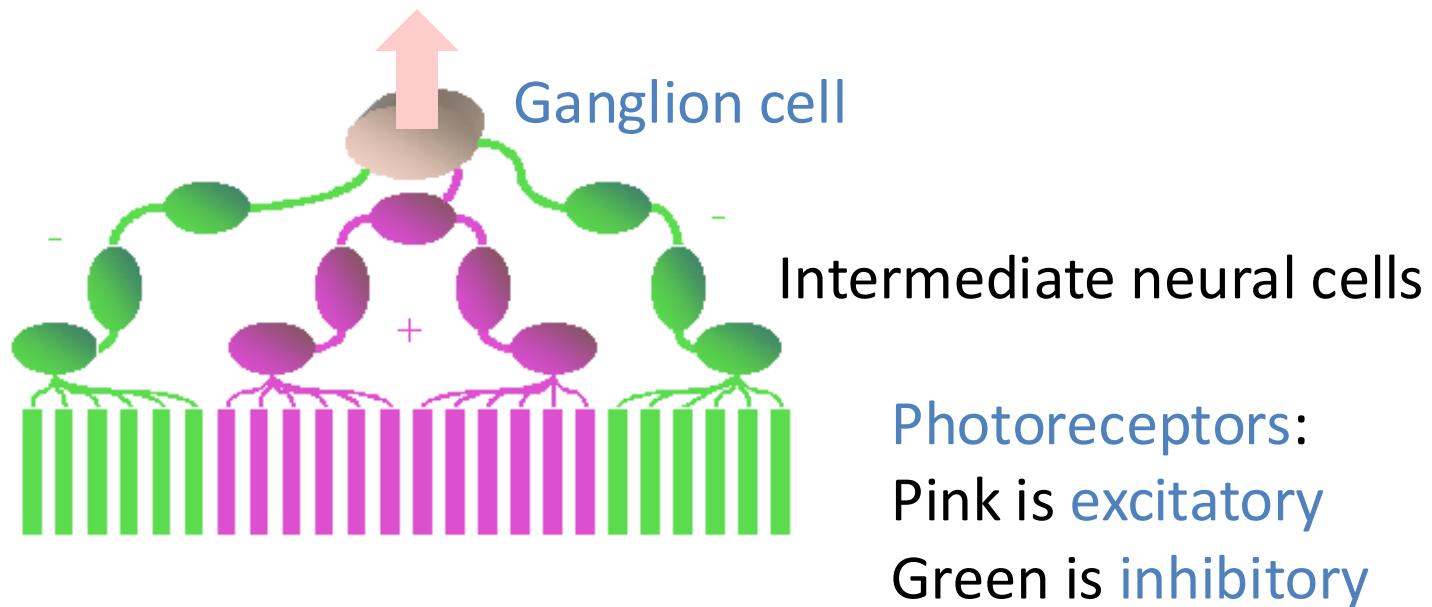
Retina

- 0.5 mm thick
- Photosensors (rods & cones) lie outermost in the retina
- Interneurons
- Ganglion cells (retina's output neurons) lie innermost in the retina closest to the lens & front of the eye

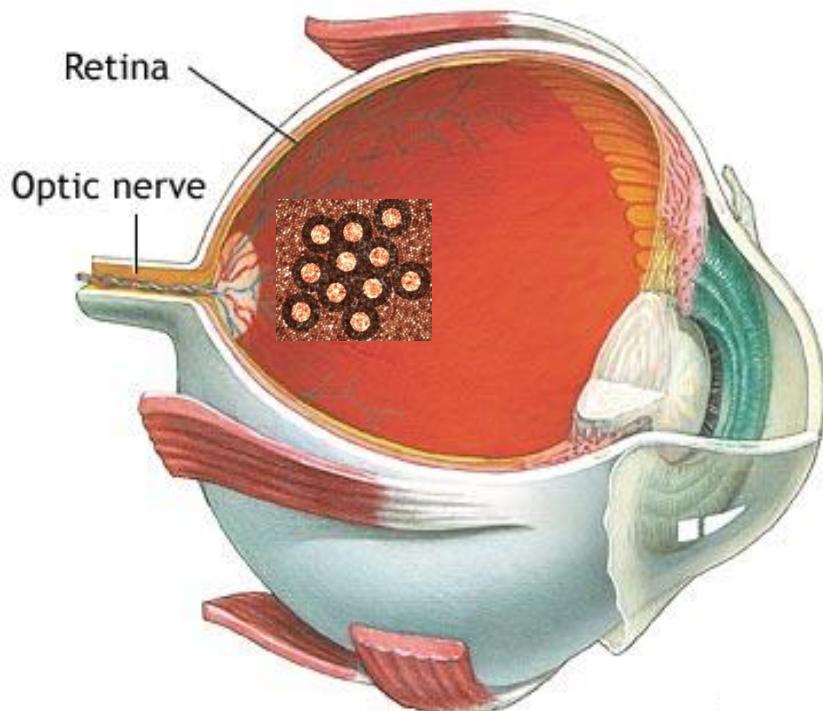


Ganglion Cell

The **ganglion cell** produces some background response even when there is no light on its receptive field

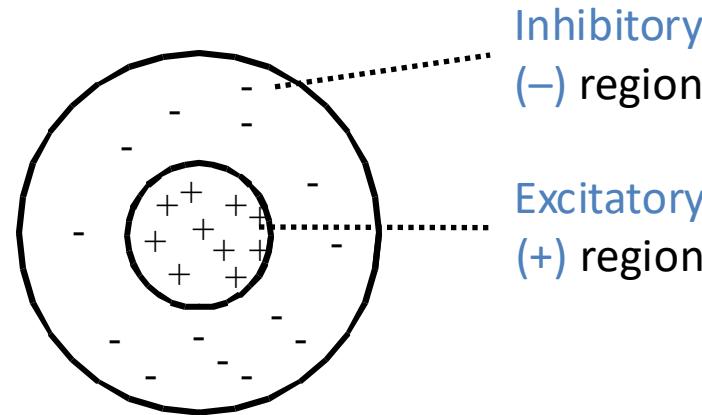


Ganglion Cell – Receptive Fields



The size of receptors & receptive field shown here much larger than actual size!

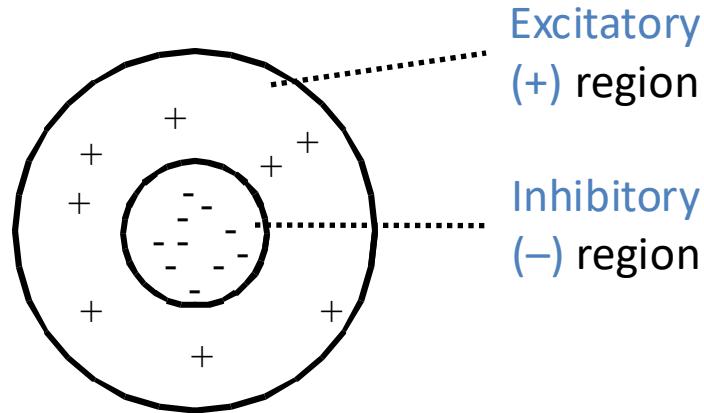
Ganglion Cell – Receptive Fields – On-Center



Responds maximally to light increments @ center
and light decrements in the surround

Ganglion Cell – Receptive Fields – Off-Center

CV

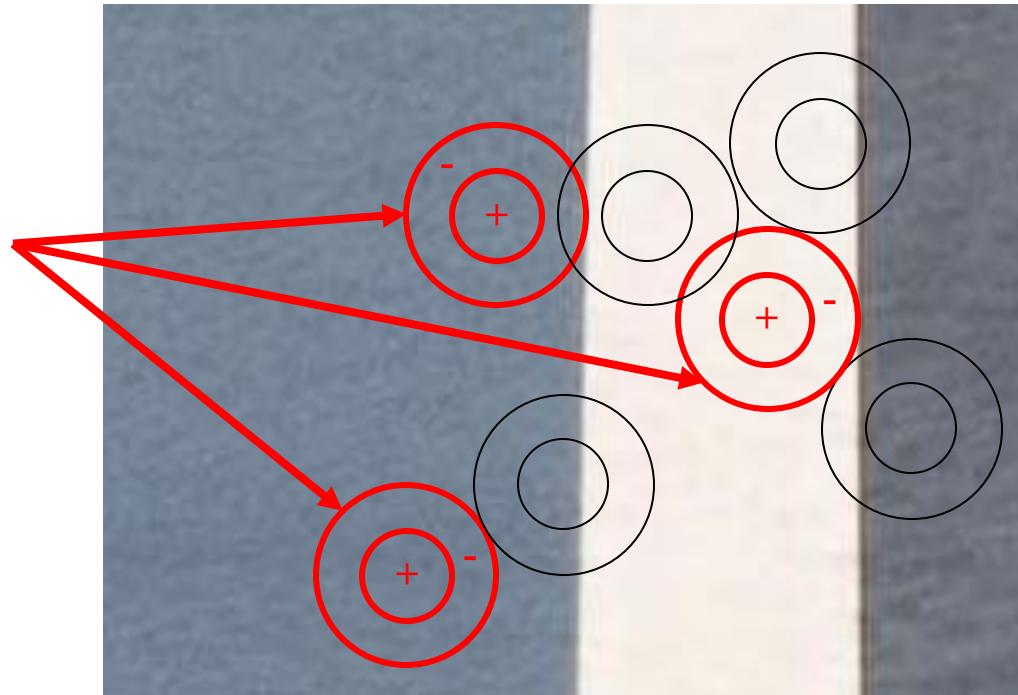


Responds maximally to light decrements in the center
and light increments in the surround

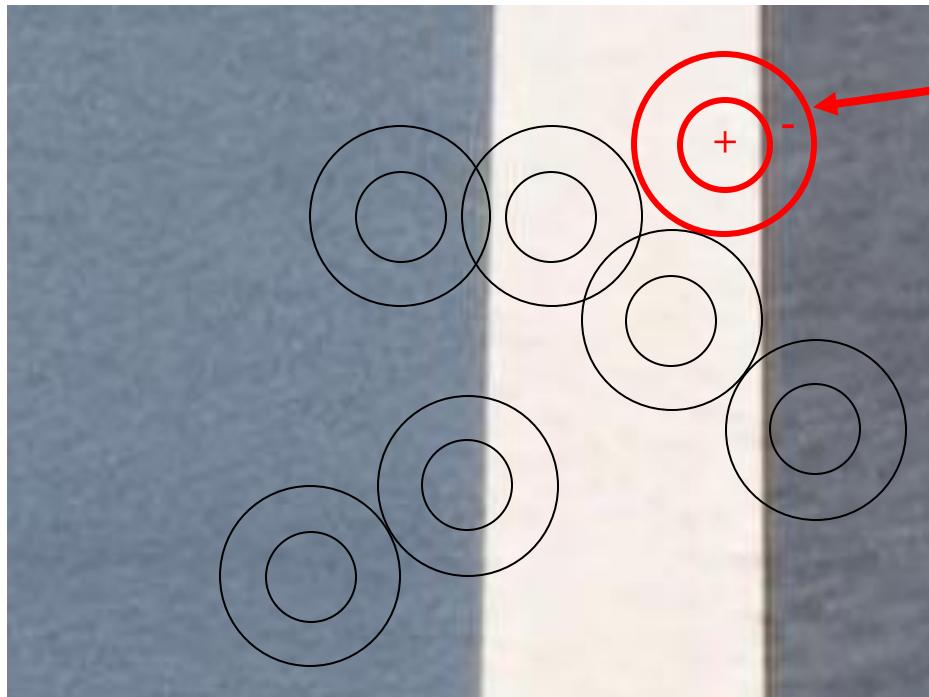
Ganglion Cell – Responses – On-Center

Receptive field
uniformly
illuminated

Response is
unchanged



Ganglion Cell – Responses – On-Center

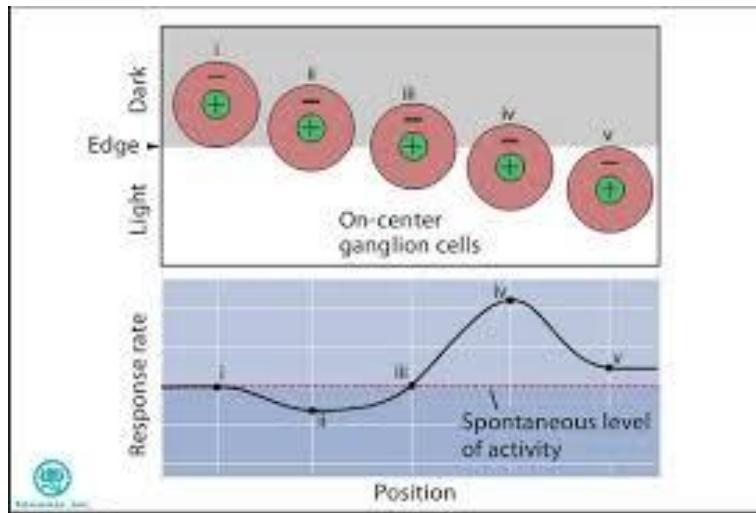


Entire excitatory reg.
illuminated

Part of inhibitory reg.
not illuminated

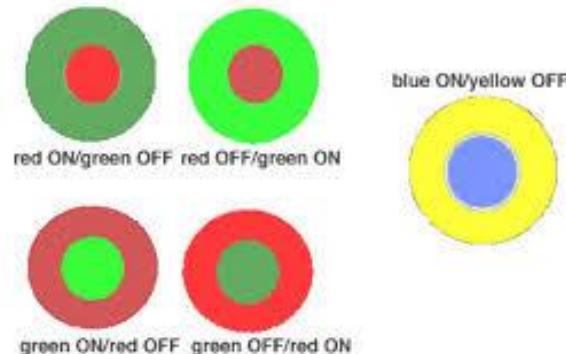
Response is
increased

Ganglion Cell – Responses – On-Center



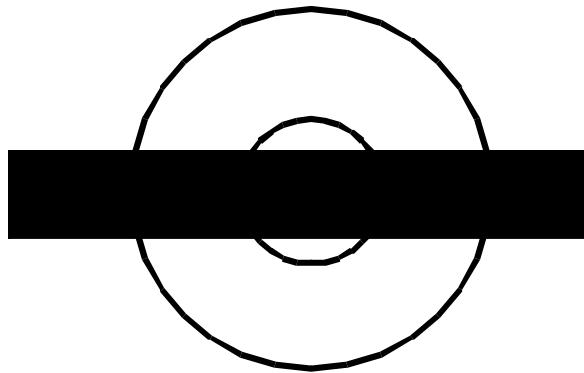
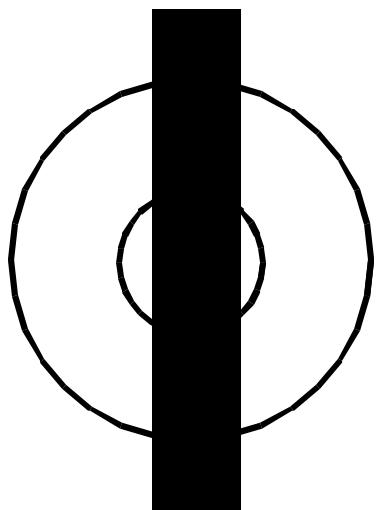
Ganglion Cell – Opponent Colors

Color opponent ganglion cells



Ganglion Cell – Rotational Equivariance

Ganglion cells have no orientation preference



Edge Responses – Ganglion Cell

Input image
(cornea)



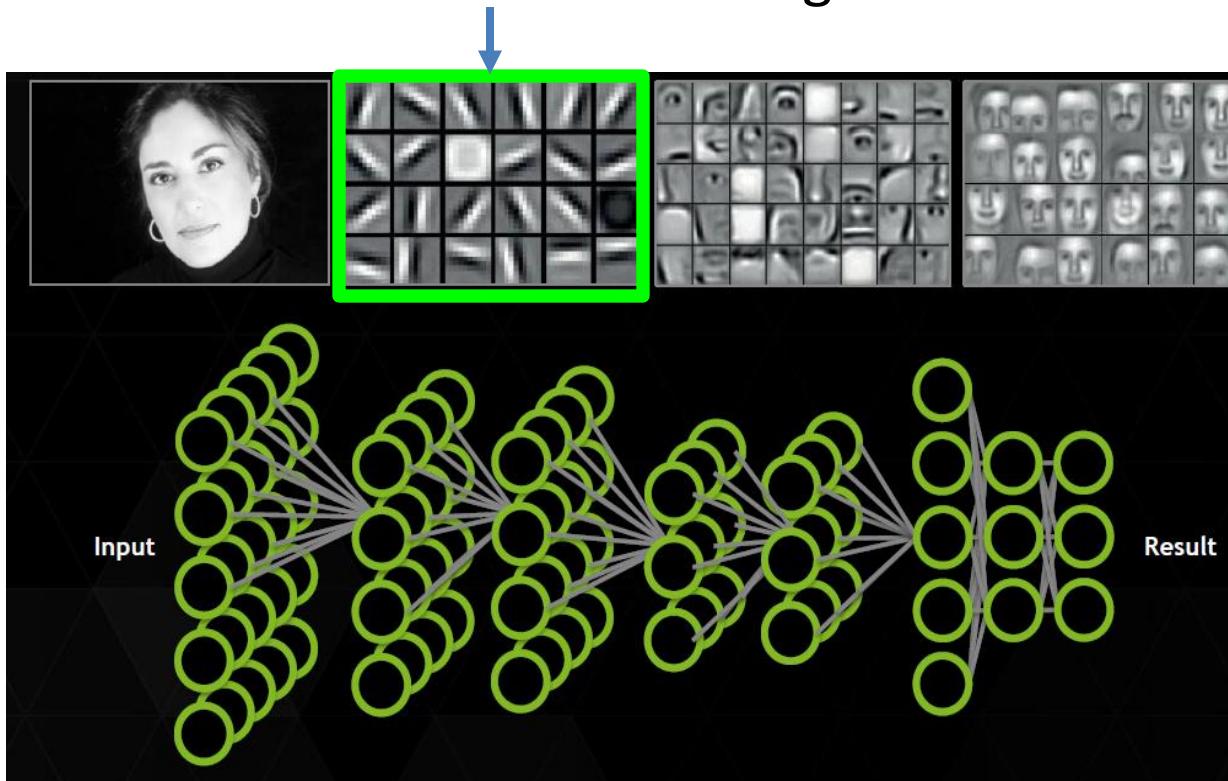
“Neural image”
(retinal ganglion cells)



Ganglion cells respond to edges

Edge Responses – Neural Nets

Neural Nets: **Similar activations** to Ganglion Cells at shallow layers

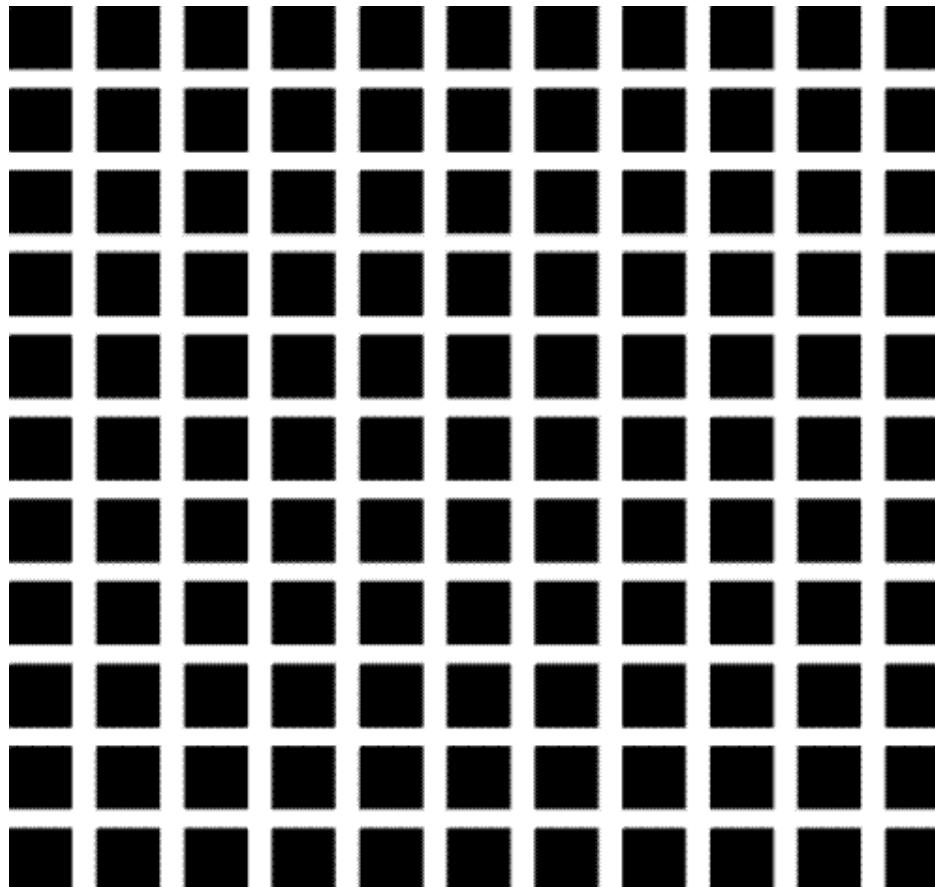


Ganglion Cells – Illusion – Hermann Grid

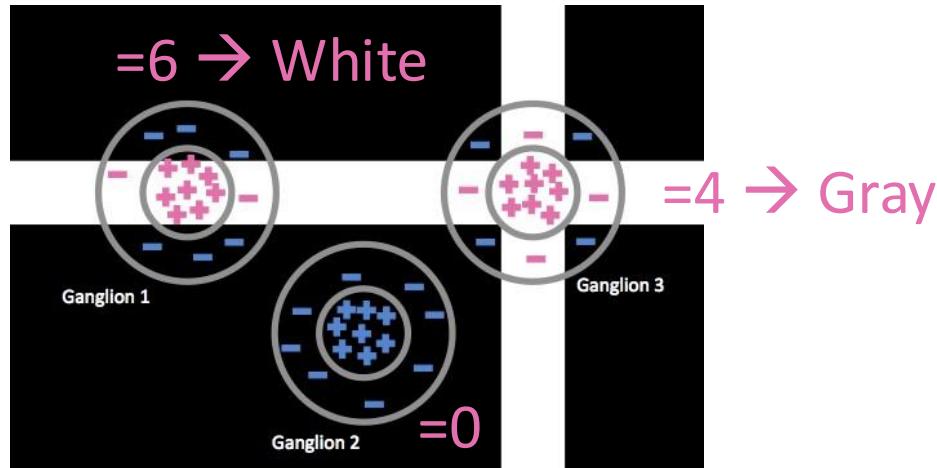
White grid on
black background

Illusion

Faint **grey** spots @
intersections
in **peripheral** areas!



Ganglion Cells – Illusion – Hermann Grid

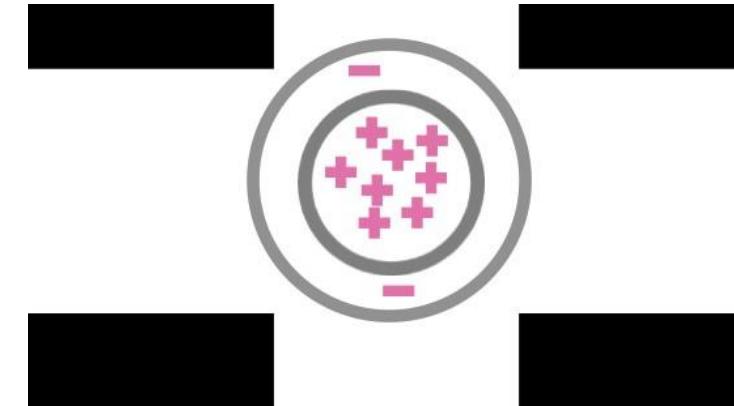


Illusion – Grey spots @ peripheral intersections

Pink → Inputs **stimulated** by light

Blue → Inputs that **not stimulated**

Compute **net signal** between activated + and -



Illusion **disappears** for intersection we directly look at!

Fovea → Receptive fields smaller
→ Cells @ intersection no longer surrounded by inhibitory components of grid

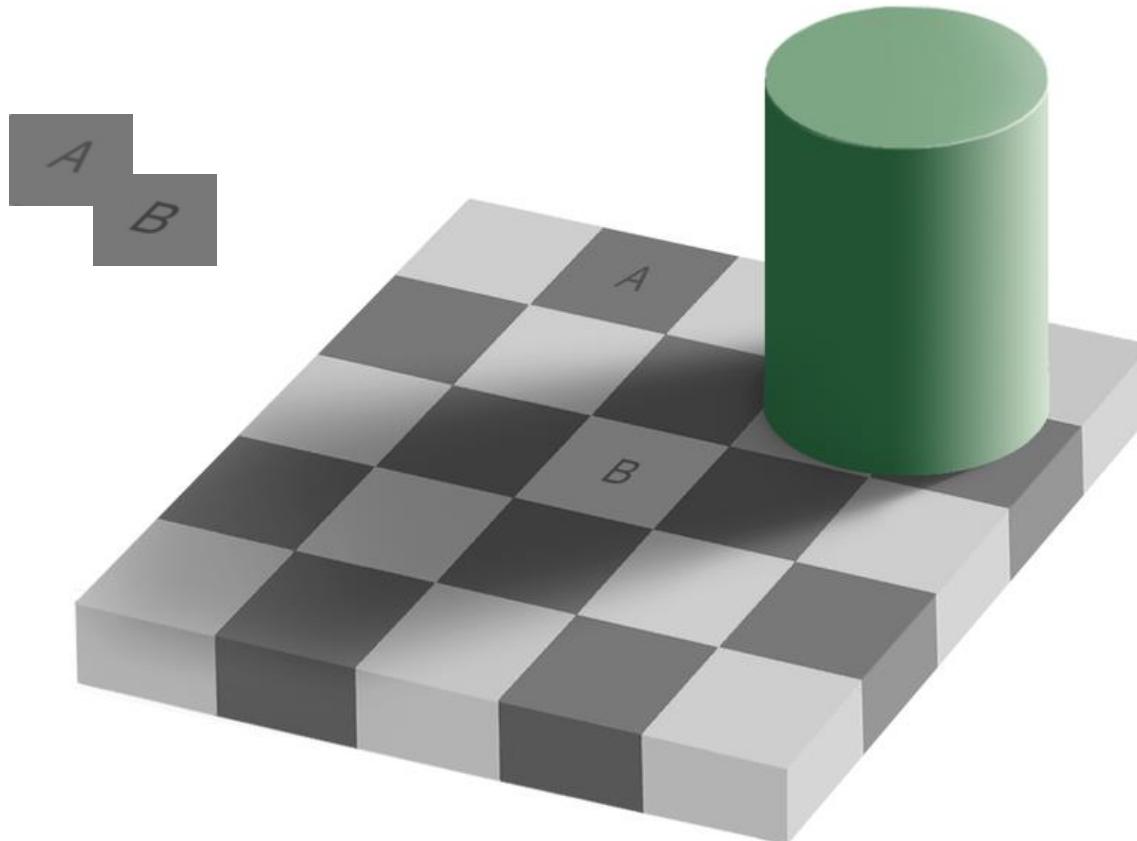


Simultaneous Lightness Contrast



- Occurs when the **lightness of an area is influenced by neighboring regions**
- Our perception of lightness is **not objective**, but depends on the **surrounding area**
- The center square on the right looks lighter because the surrounding area is a darker gray

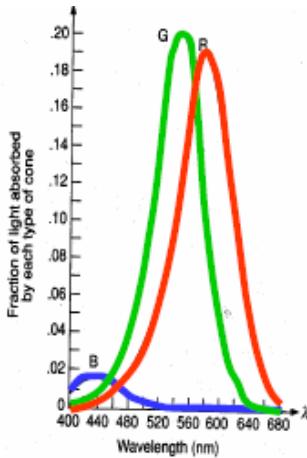
Simultaneous Lightness Contrast



Outline

- Camera Model
 - Pinhole Camera
 - Geometric Image Formation
 - Photometric Image Formation
 - Image Representation
- Color
 - Physical & Biological Model
 - Light Source
 - Object
 - Observer
 - Tristimulus Theory
 - Colour Systems

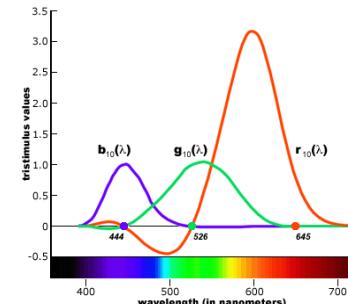
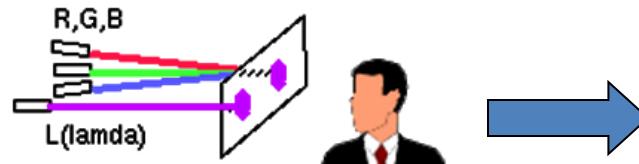
Color – Tristimulus Theory



Spectral-response functions
of each of 3 cone types

Color matching function based on RGB

- Any spectral color – Linear combination of these primary colors
- Perceptual experiment – Human observers try to match a color of a given wavelength λ by mixing 3 pure wavelengths



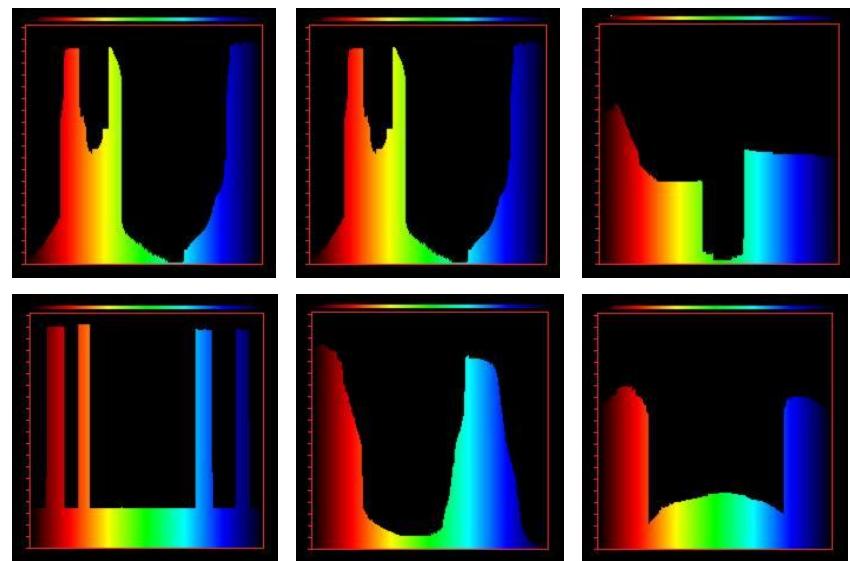
Problem: Sometimes red to be added to the target before a match can be achieved. Shown on the graph with a negative R value

Spectral Energy Distribution

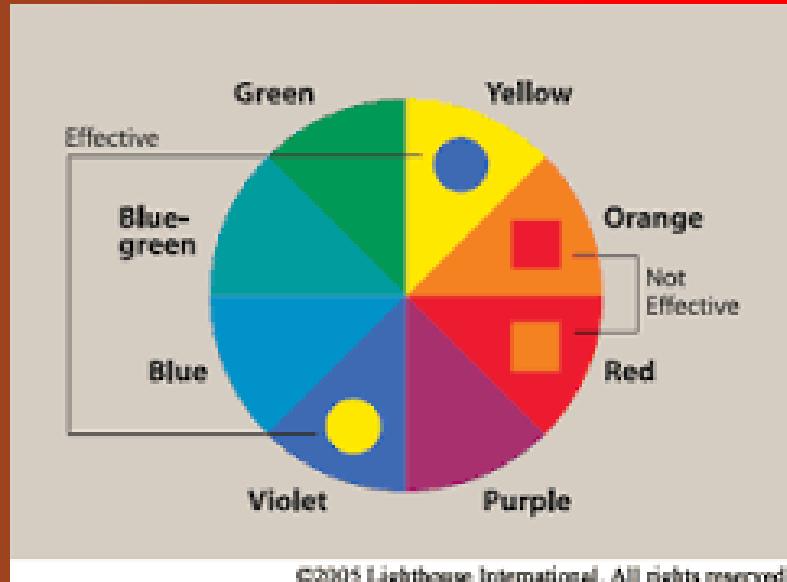
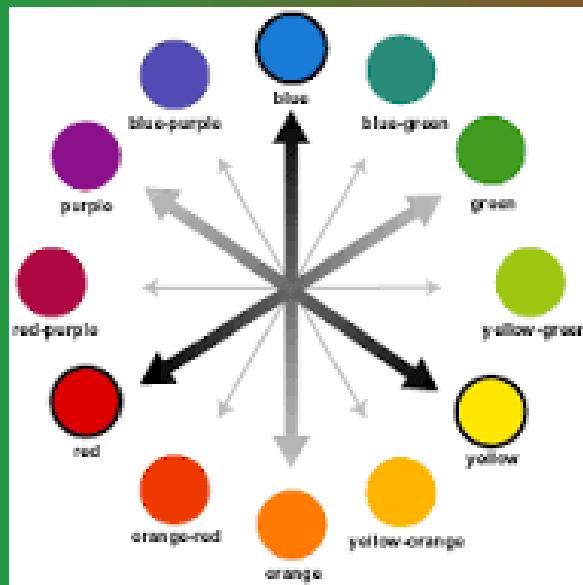


Subjective Perception:

These 6 spectra **look** like the
'same purple' to *normal*
color-vision people

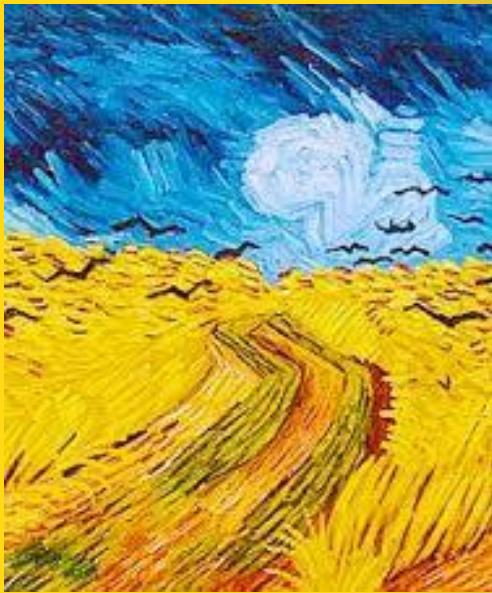


Opponent Colors



©2005 Lighthouse International. All rights reserved.

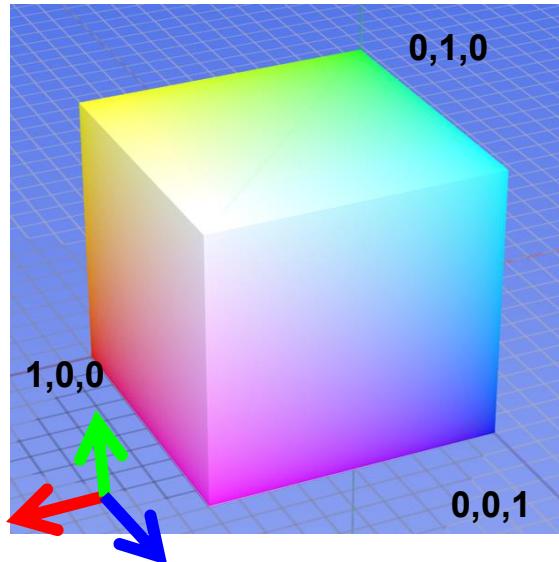
Opponent Colors



Outline

- Camera Model
 - Pinhole Camera
 - Geometric Image Formation
 - Photometric Image Formation
 - Image Representation
- Color
 - Physical & Biological Model
 - Light Source
 - Object
 - Observer
 - Tristimulus Theory
 - Colour Systems

RGB space



Default color space

Drawbacks:

- Strongly **correlated** channels
- Perceptually **non-meaningful**



R
(G=0,B=0)

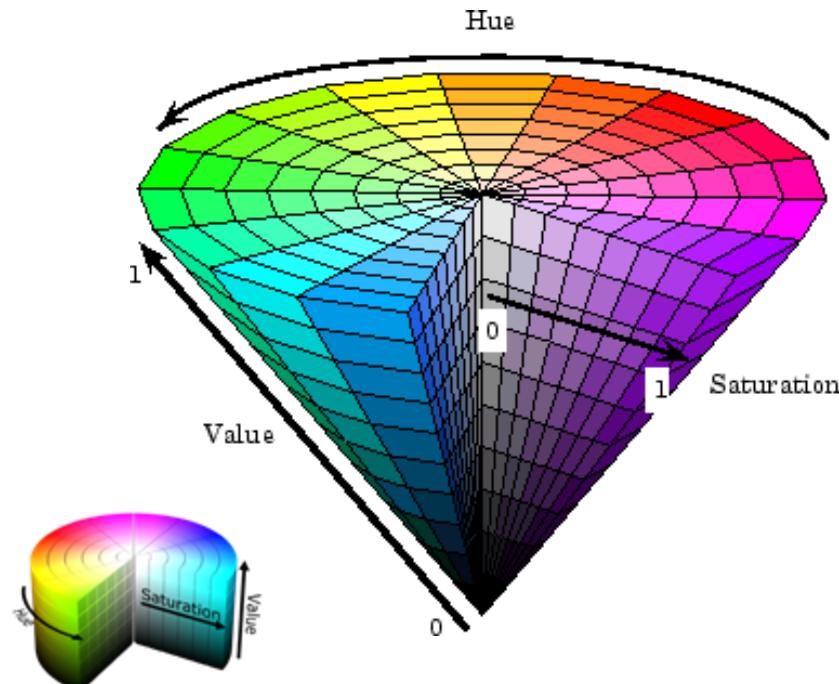
G
(R=0,B=0)

B
(R=0,G=0)

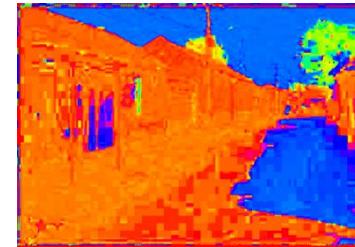


OpenCV order is
BGR

HSV Color Space



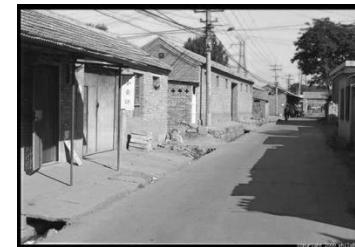
Intuitive color space



H
($S=1, V=1$)



S
($H=1, V=1$)



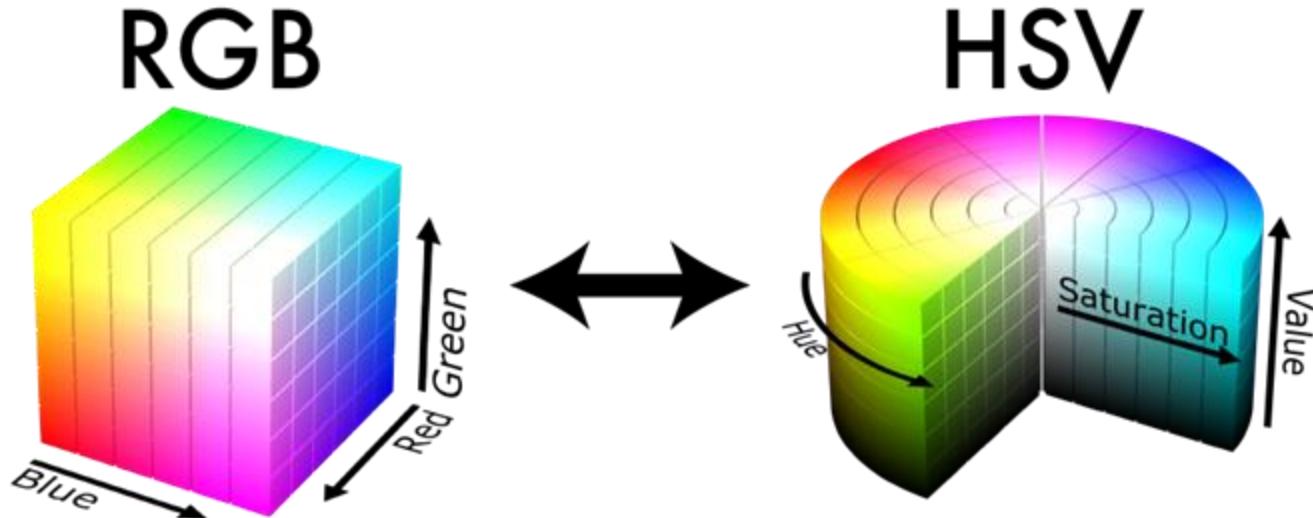
V
($H=1, S=0$)

Hue

Saturation

Value

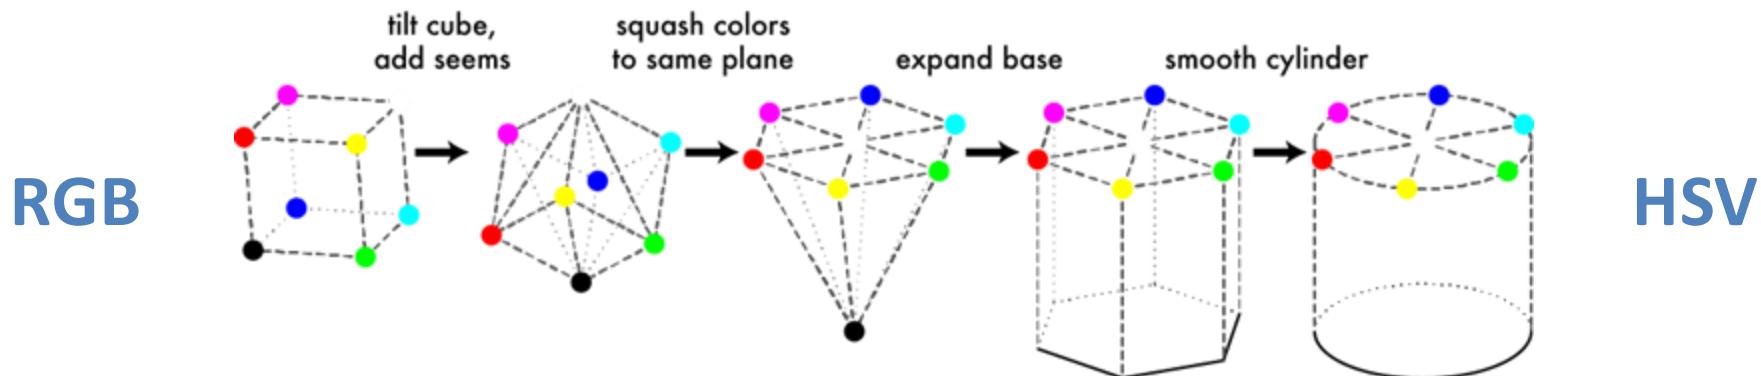
Color Space Conversion – RGB → HSV



https://en.wikipedia.org/wiki/HSL_and_HSV

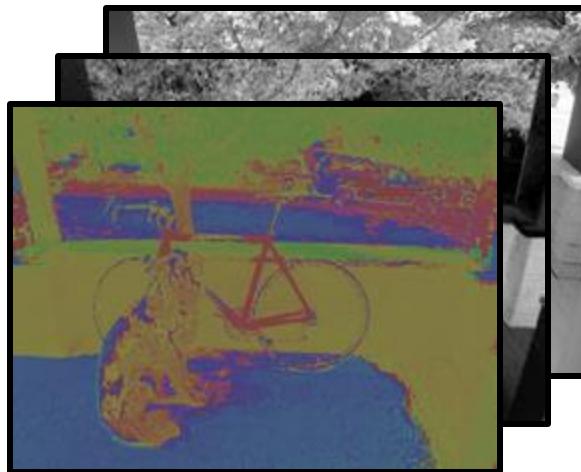
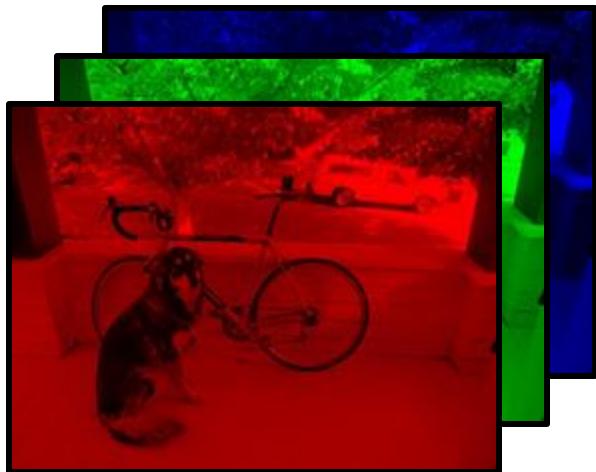
Color Space Conversion – RGB → HSV

Geometric Interpretation



https://en.wikipedia.org/wiki/HSL_and_HSV

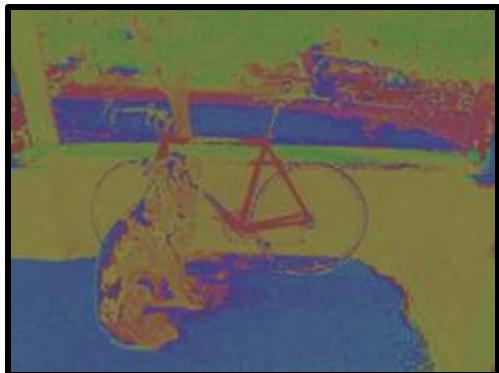
Color Space Conversion – RGB → HSV



Still 3D tensors

but encode **different info**

HSV Color Space



Hue

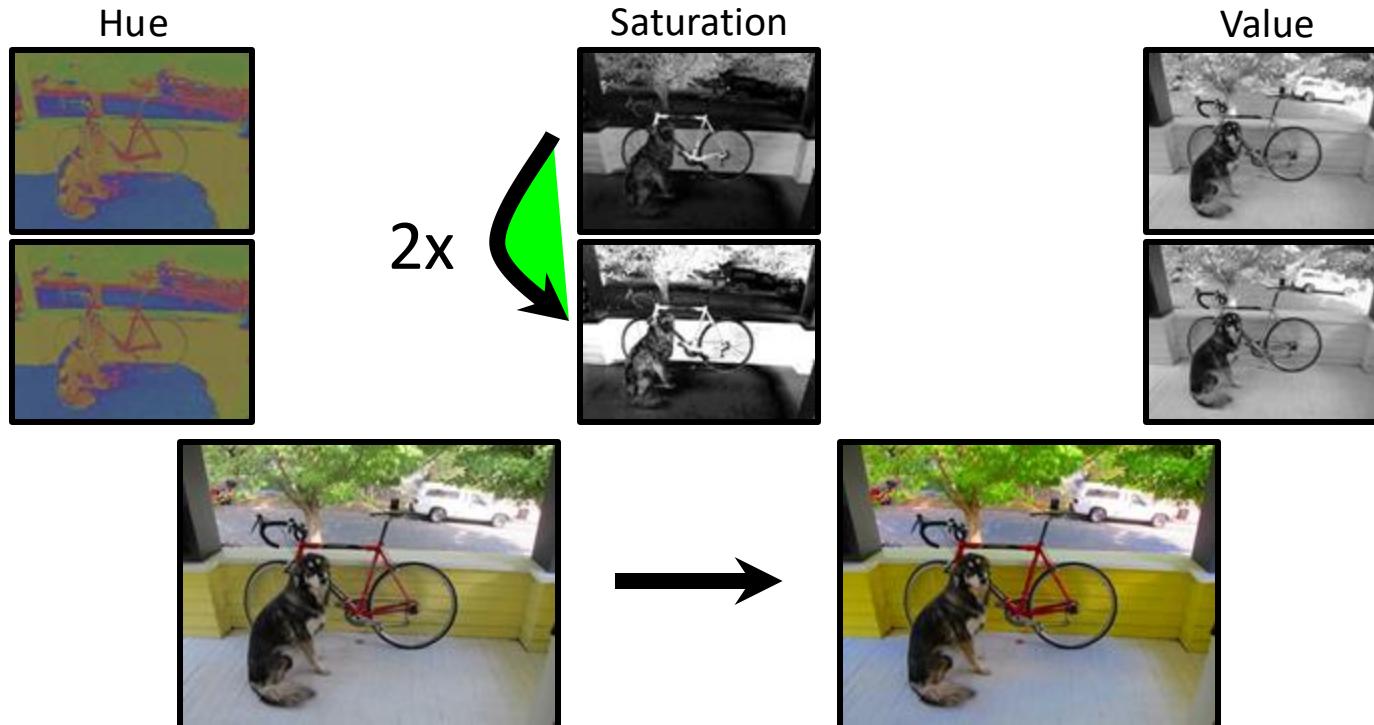


Saturation



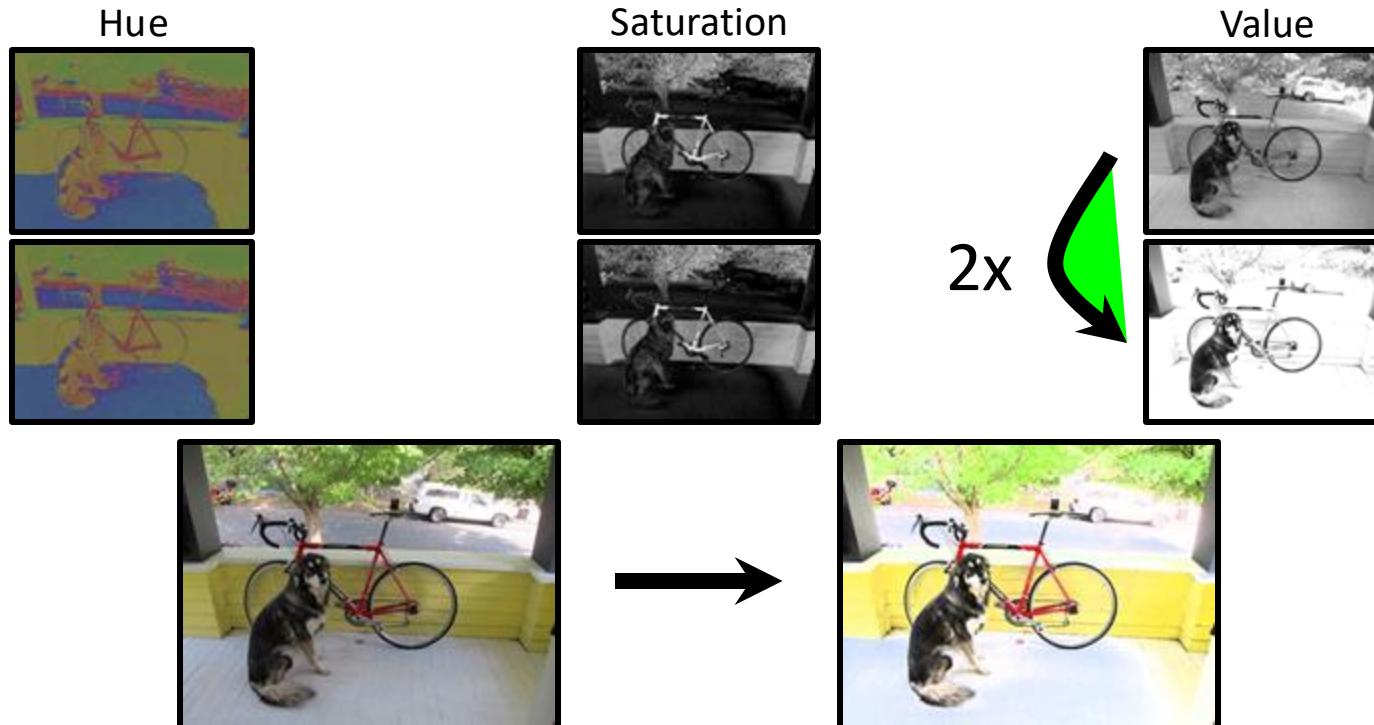
Value

HSV Color Space



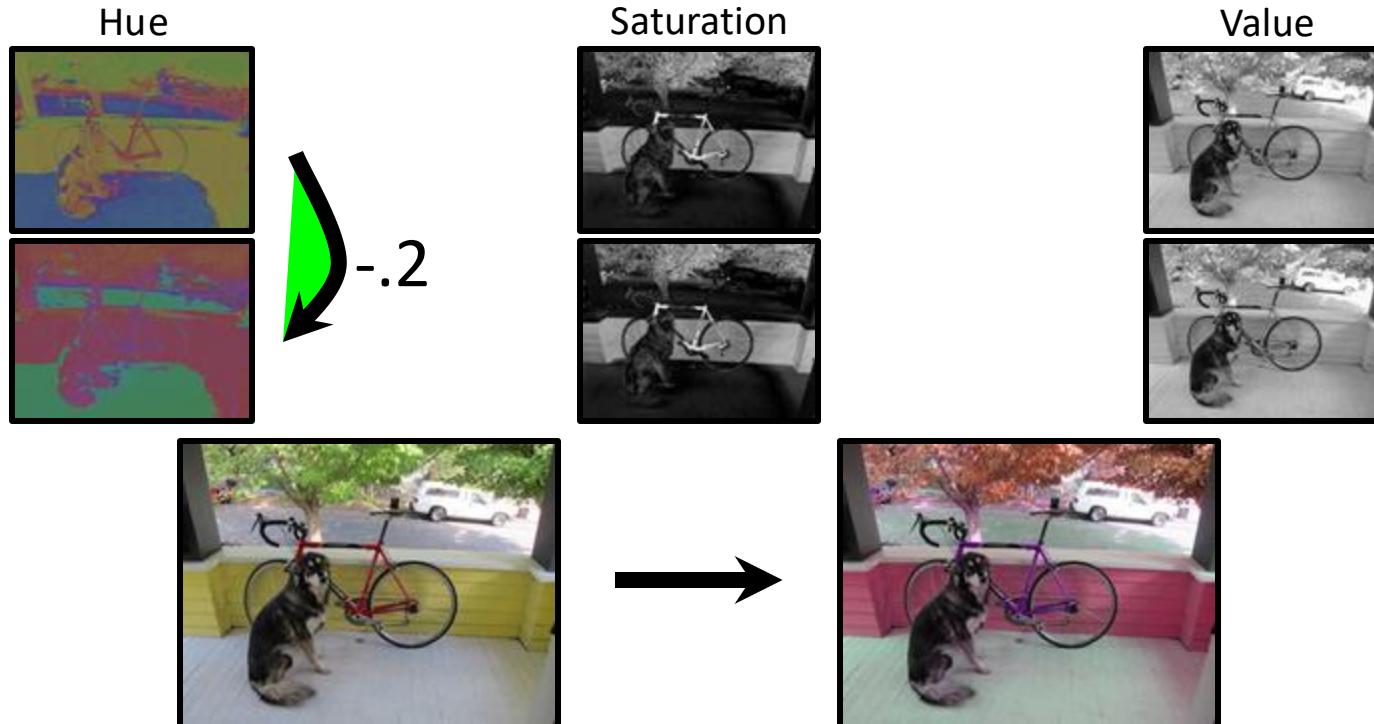
Higher Saturation → More intense colors

HSV Color Space



Higher Value → Brighter colors

HSV Color Space

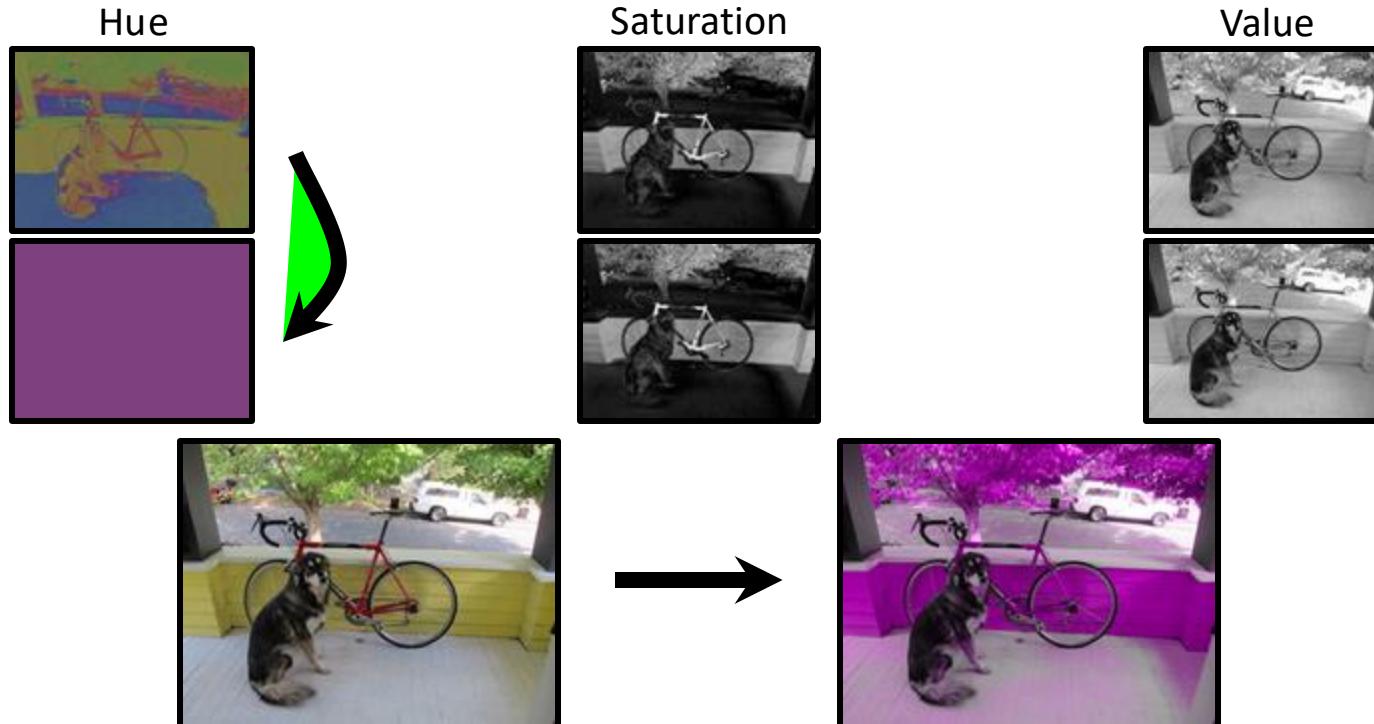


Value



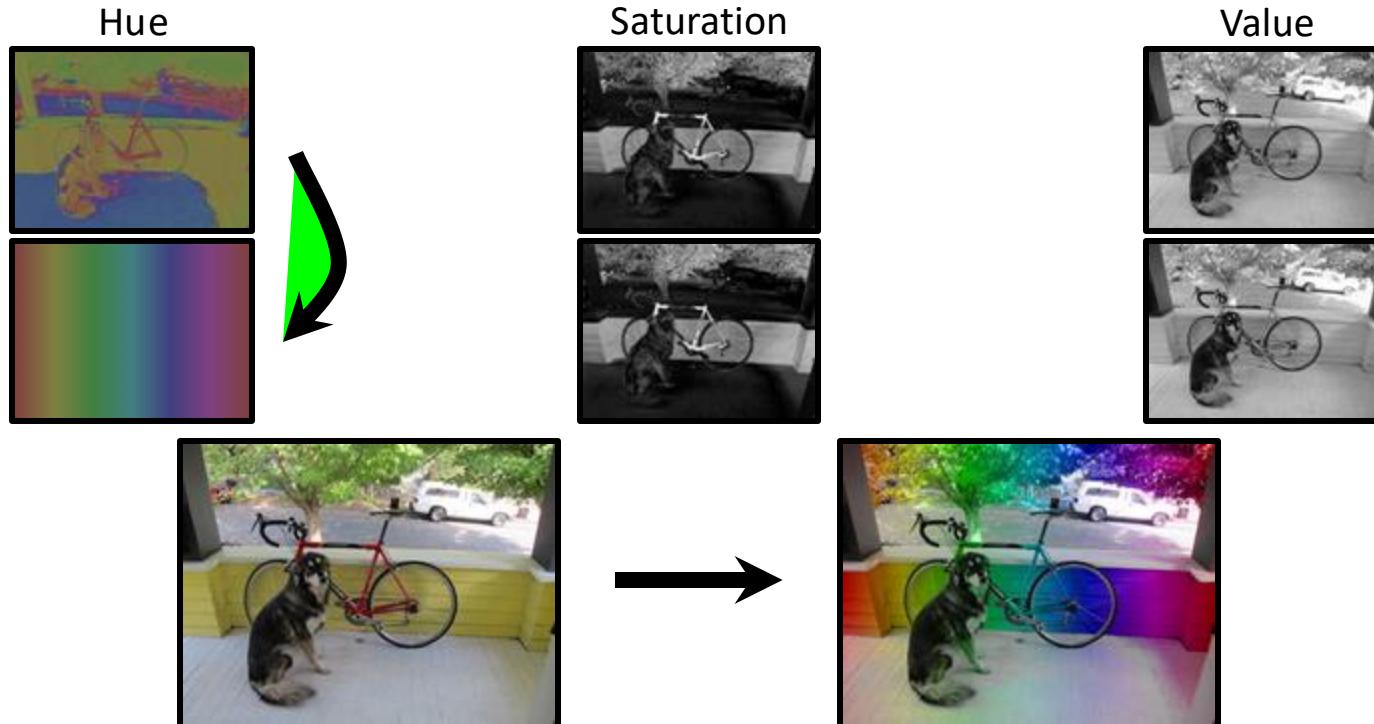
Change Hue → Change Colors

HSV Color Space



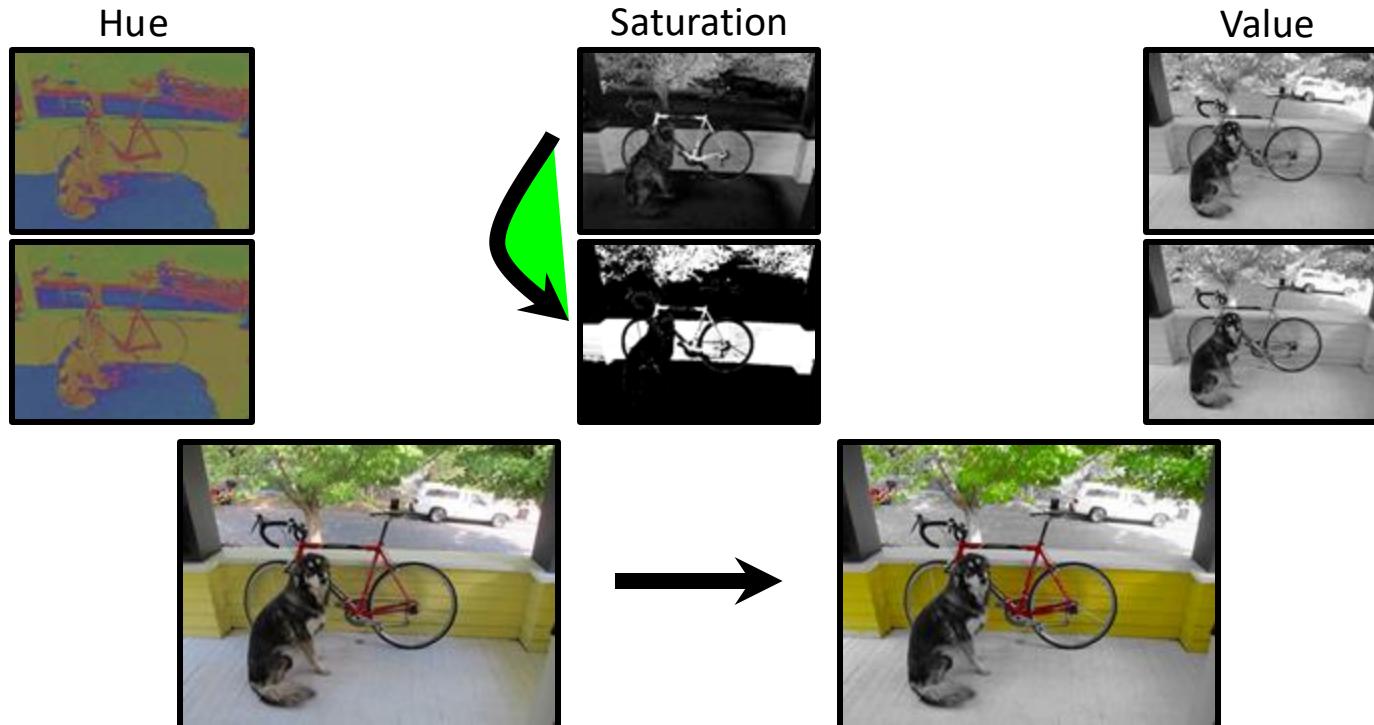
Set Hue to your favorite color!

HSV Color Space



Set **Hue** to your favorite pattern!

HSV Color Space



Increase & Threshold Saturation

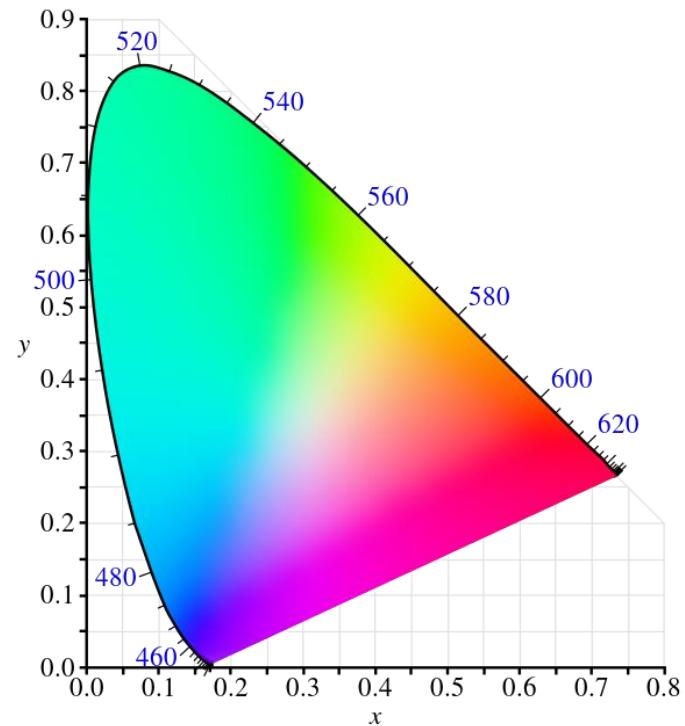
HSV Color Space



CIE Color Space – Chromaticity Diagram

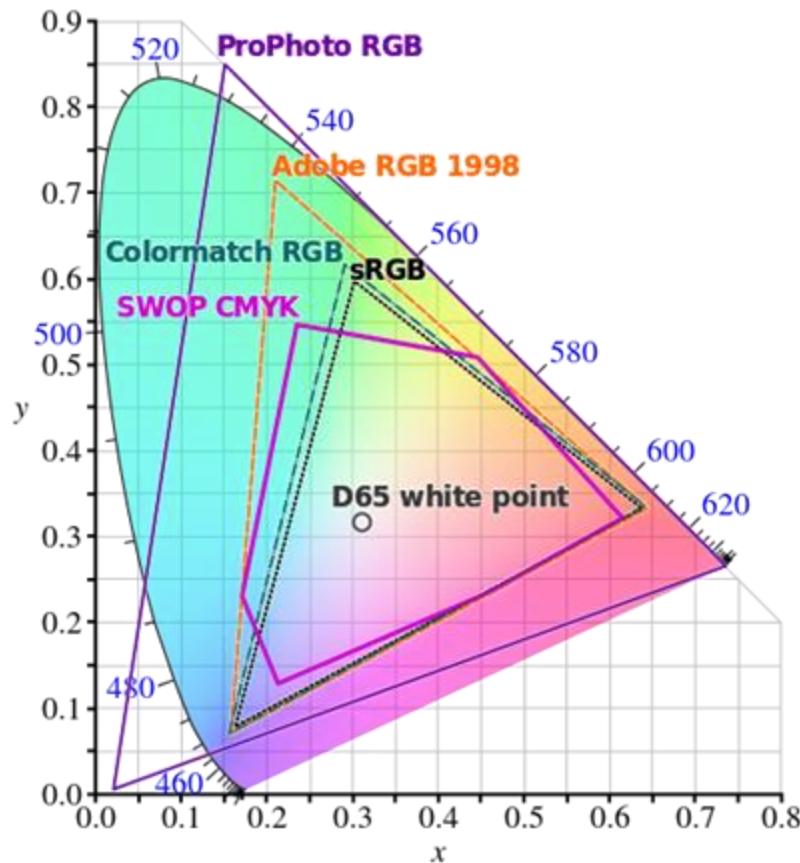
Wikipedia:

- CIE 1931 color space [chromaticity diagram](#) with [wavelengths in nanometers](#).
- The [colors depicted](#) depend on the [color space of the device](#) on which the image is viewed.



[International Commission on Illumination](#)
Commission Internationale de l'Éclairage (CIE)

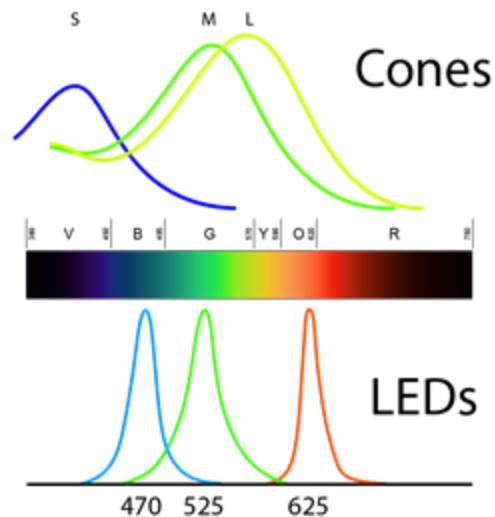
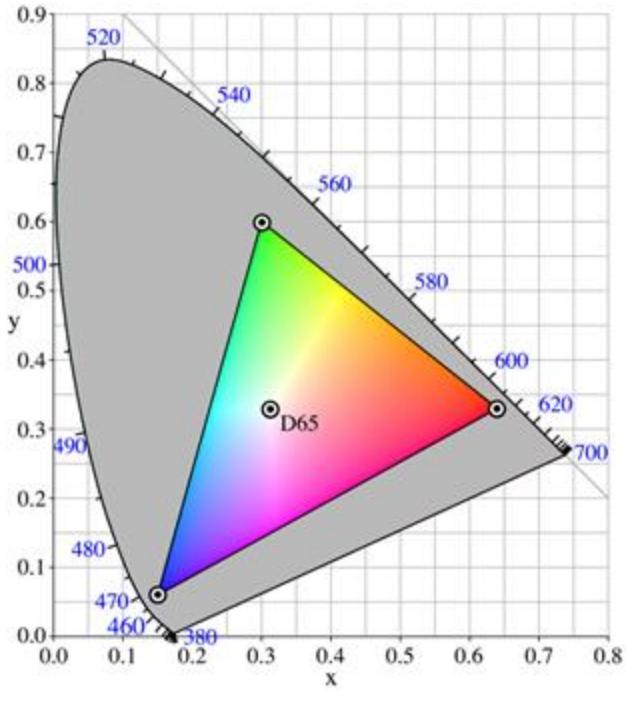
MANY Color Spaces



sRGB Primaries

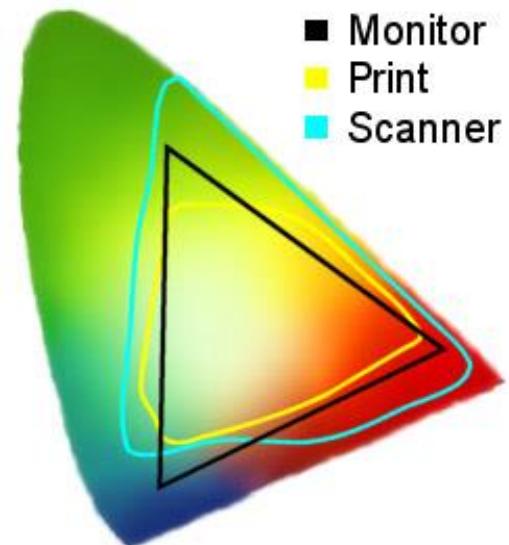
Most widely used color space

Not all colors can be represented with RGB (or any Color Space)



Monitor / Printer / Scanner Gamut

- The **gamut** is a triangle – All colors a device generates are *positive linear combinations* from *3 light sources* (e.g. a monitor with 3 LED types for R,G,B)
- The 3 light sources are the **triangle's vertices**
- The 3 light sources are usually not pure color (single wavelength), so that the vertices are usually not on the boundary of the CIE-xy chart
- In a physical color context, they are not pure color.
However, on the device, a single LED color is the purest color it can generate, so they are usually referred as 'pure color' in such context



What does this mean for computers?

- We represent images as grids of pixels
 - Each pixel has a 3-component color: RGB
- Can represent color with 3 numbers
 - #ff00ff; (1.0, 0.0, 1.0); (255, 0, 255); etc...
- Not every color can be represented in RGB
- RGB is made to ‘trick’ humans – not to be accurate

Closing Remarks

Action points

- Check out & start working on Lab0 & Lab1 assignments
- Feel free to discuss on Ed
- Lecture 3 (HC2a) covers the relevant theory

Disclaimer

Many of the slides used here are obtained from online resources (including many open lecture materials) without appropriate acknowledgement. They are used here for the sole purpose of classroom teaching. All the credit and all the copyrights belong to the original authors. You should not copy it, redistribute it, put it online, or use it for any other purposes than for this course.

Pinhole Camera

Pinhole camera:

Simple model, approx.
perspective projection
for image formation

Pinhole is a point:

Only 1 ray from any given
point enters the camera

