



# Deep Learning 1

2025-2026 – Pascal Mettes

**Lecture 7**

*Graph Deep Learning*

# Previous lecture

Lecture	Title	Lecture	Title
1	Intro and history of deep learning	2	AutoDiff
3	Deep learning optimization I	4	Deep learning optimization II
5	Convolutional deep learning	6	Attention-based deep learning
7	Graph deep learning	8	From supervised to unsupervised deep learning
9	Multi-modal deep learning	10	Generative deep learning
11	What doesn't work in deep learning	12	Non-Euclidean deep learning
13	Q&A	14	Deep learning for videos

I have to leave a 10:10 on the dot to be able to make it to a PhD committee at 11:00 in the city center.

# This lecture

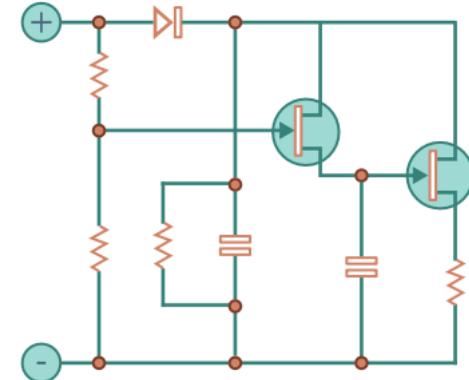
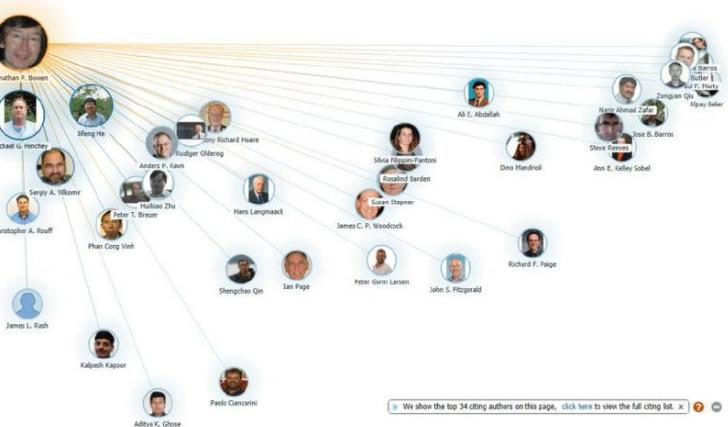
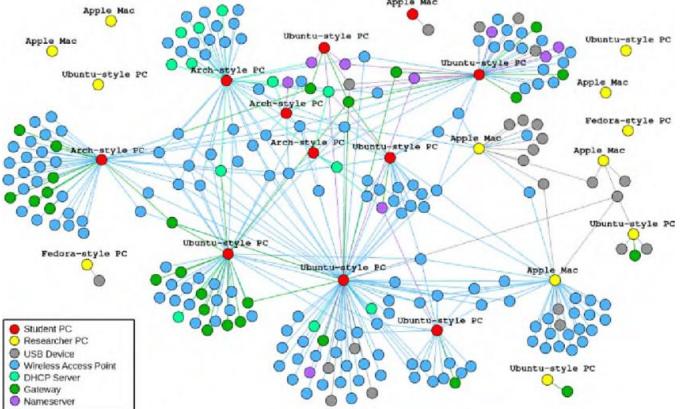
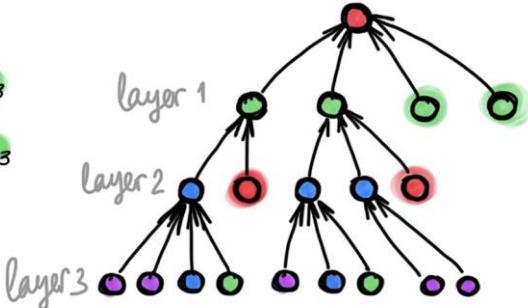
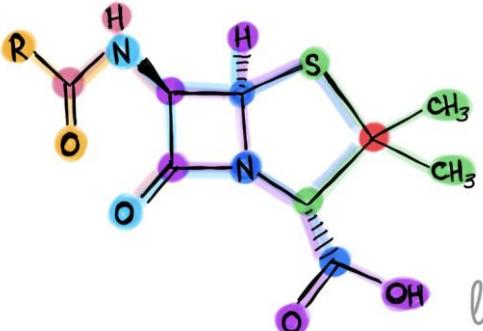
Graphs

Graph convolutions

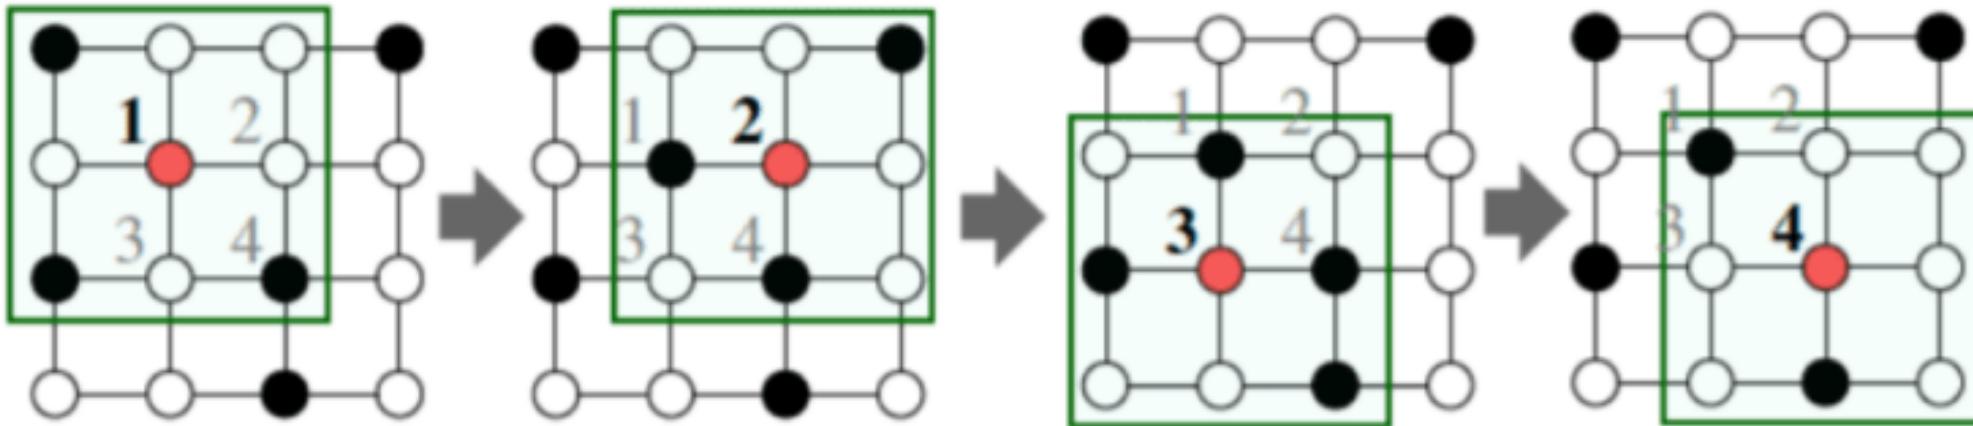
Graph attention

Graph applications

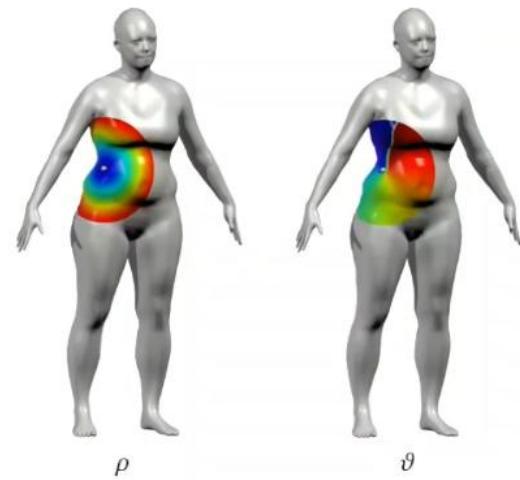
# Graphs, more common than you think



# Many structures are special cases of graphs

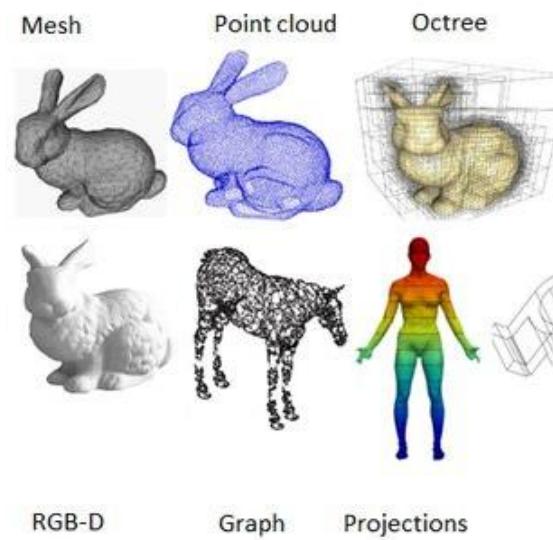
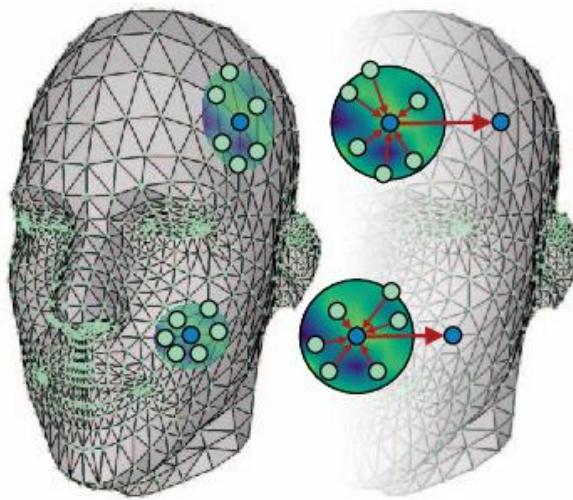
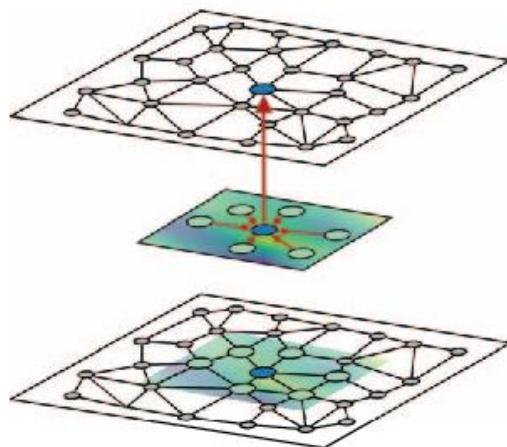
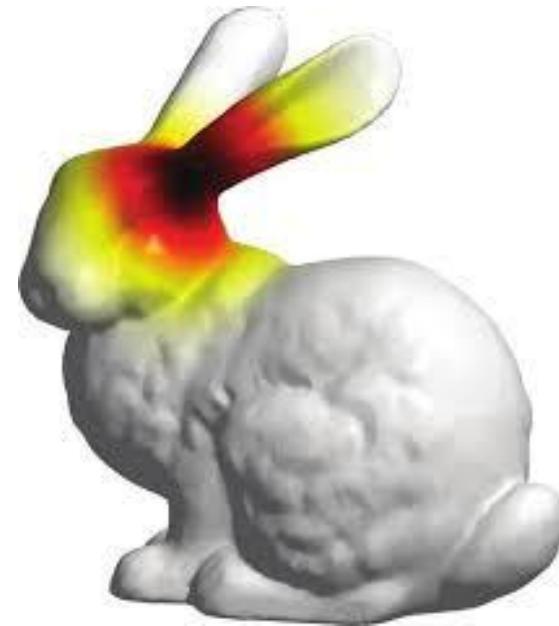


# Graphs as geometry



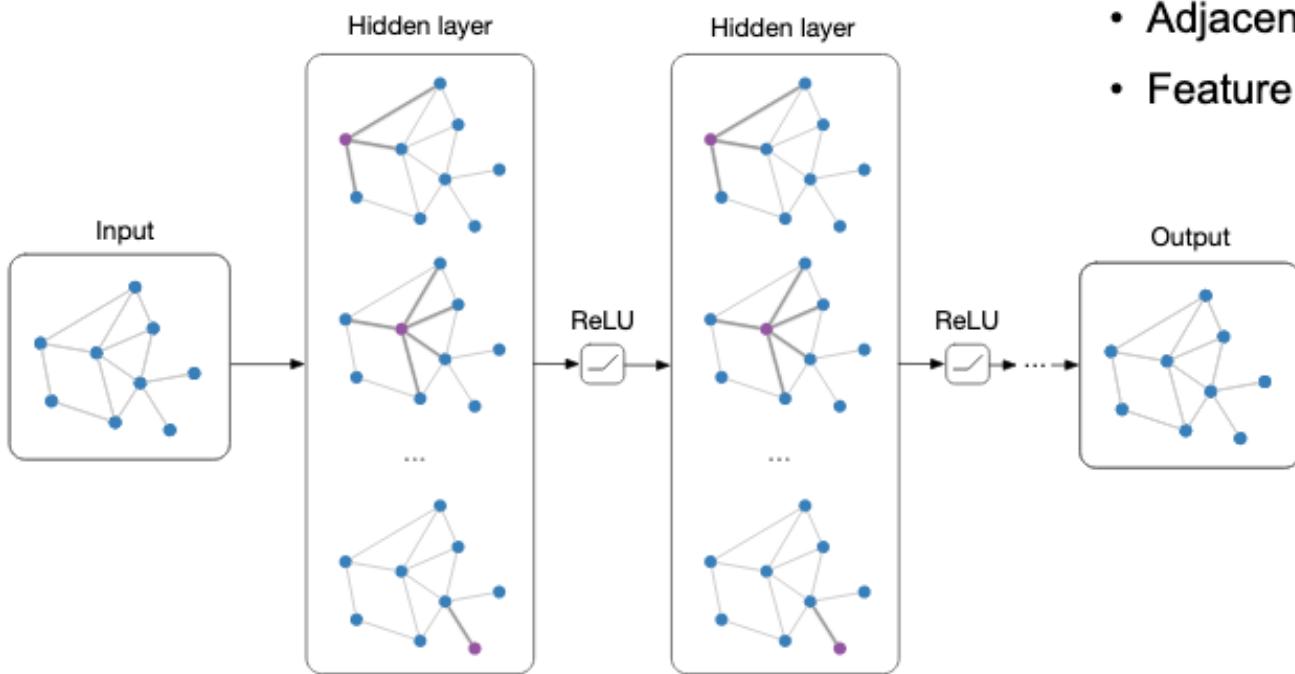
$\rho$

$\vartheta$



# What are graph networks?

**The bigger picture:**



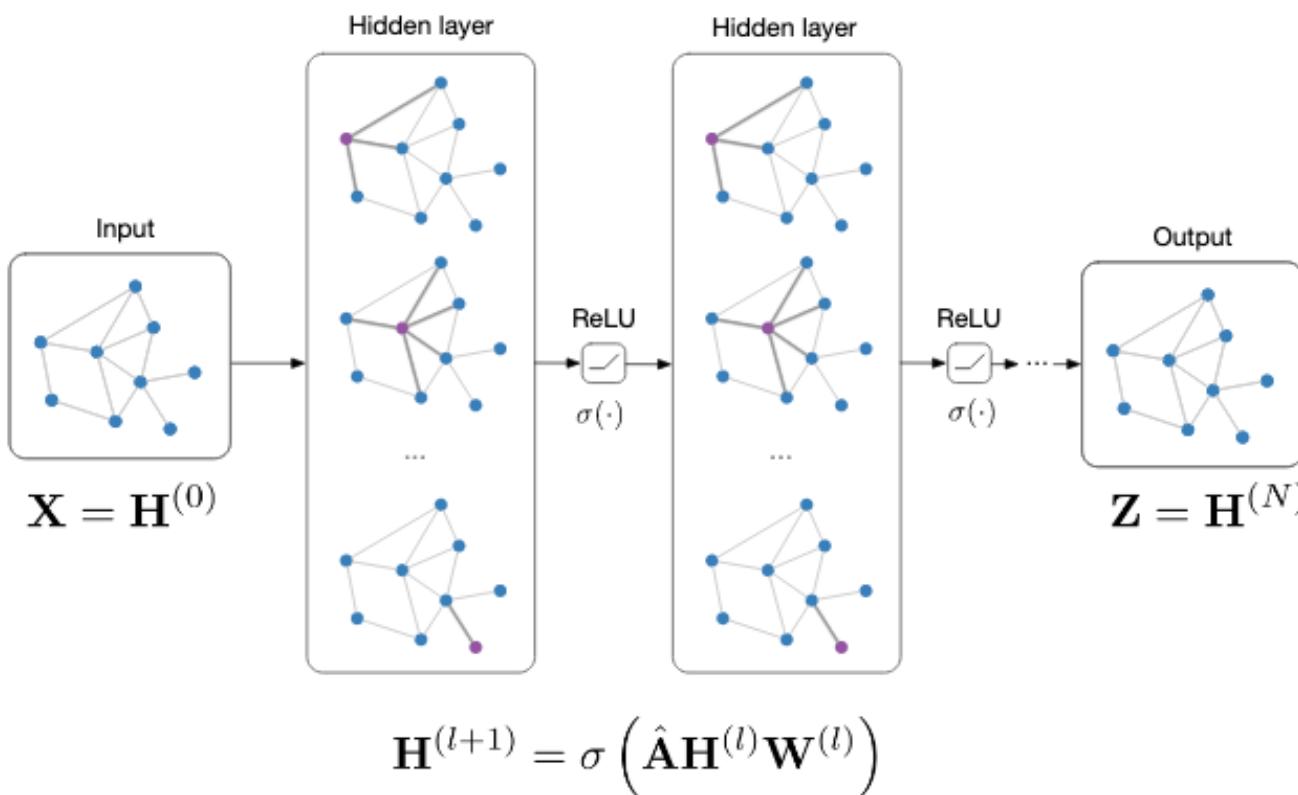
**Notation:**  $\mathcal{G} = (\mathbf{A}, \mathbf{X})$

- Adjacency matrix  $\mathbf{A} \in \mathbb{R}^{N \times N}$
- Feature matrix  $\mathbf{X} \in \mathbb{R}^{N \times F}$

**Main idea:** Pass messages between pairs of nodes & agglomerate

# Graph networks

**Input:** Feature matrix  $\mathbf{X} \in \mathbb{R}^{N \times E}$ , preprocessed adjacency matrix  $\hat{\mathbf{A}}$



**Node classification:**

$$\text{softmax}(\mathbf{z}_n)$$

e.g. Kipf & Welling (ICLR 2017)

**Graph classification:**

$$\text{softmax}(\sum_n \mathbf{z}_n)$$

e.g. Duvenaud et al. (NIPS 2015)

**Link prediction:**

$$p(A_{ij}) = \sigma(\mathbf{z}_i^T \mathbf{z}_j)$$

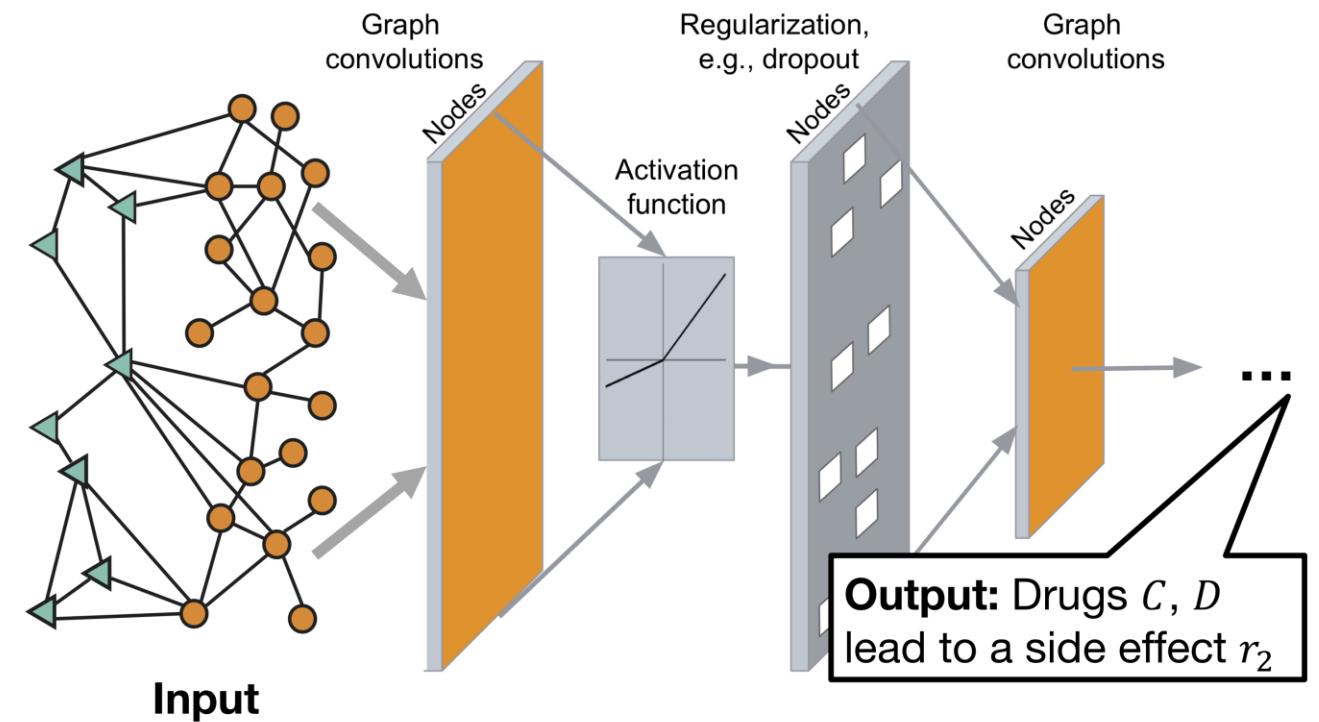
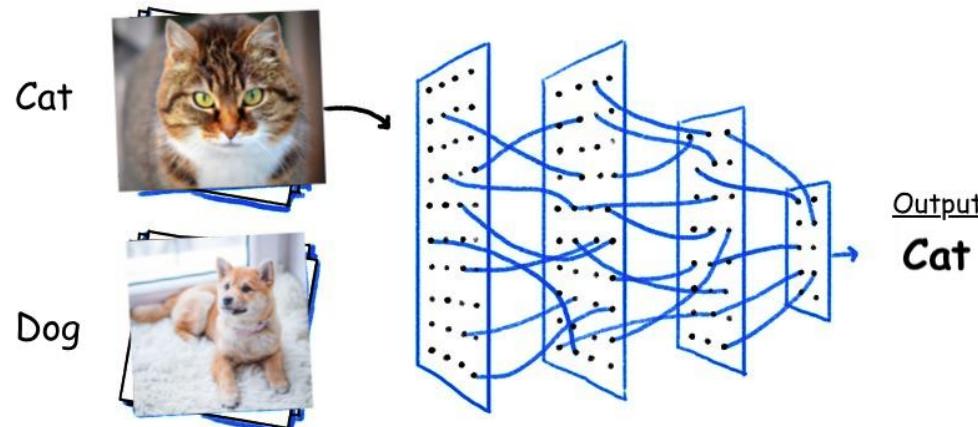
Kipf & Welling (NIPS BDL 2016)

“Graph Auto-Encoders”

# 1) Graph classification

Make a prediction over the entire graph.

Akin to assigning a label to an entire image.



## 2) Node classification

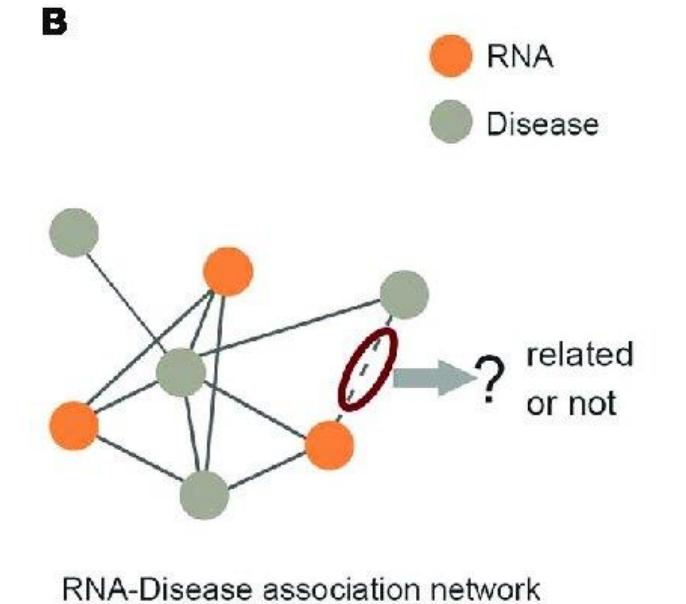
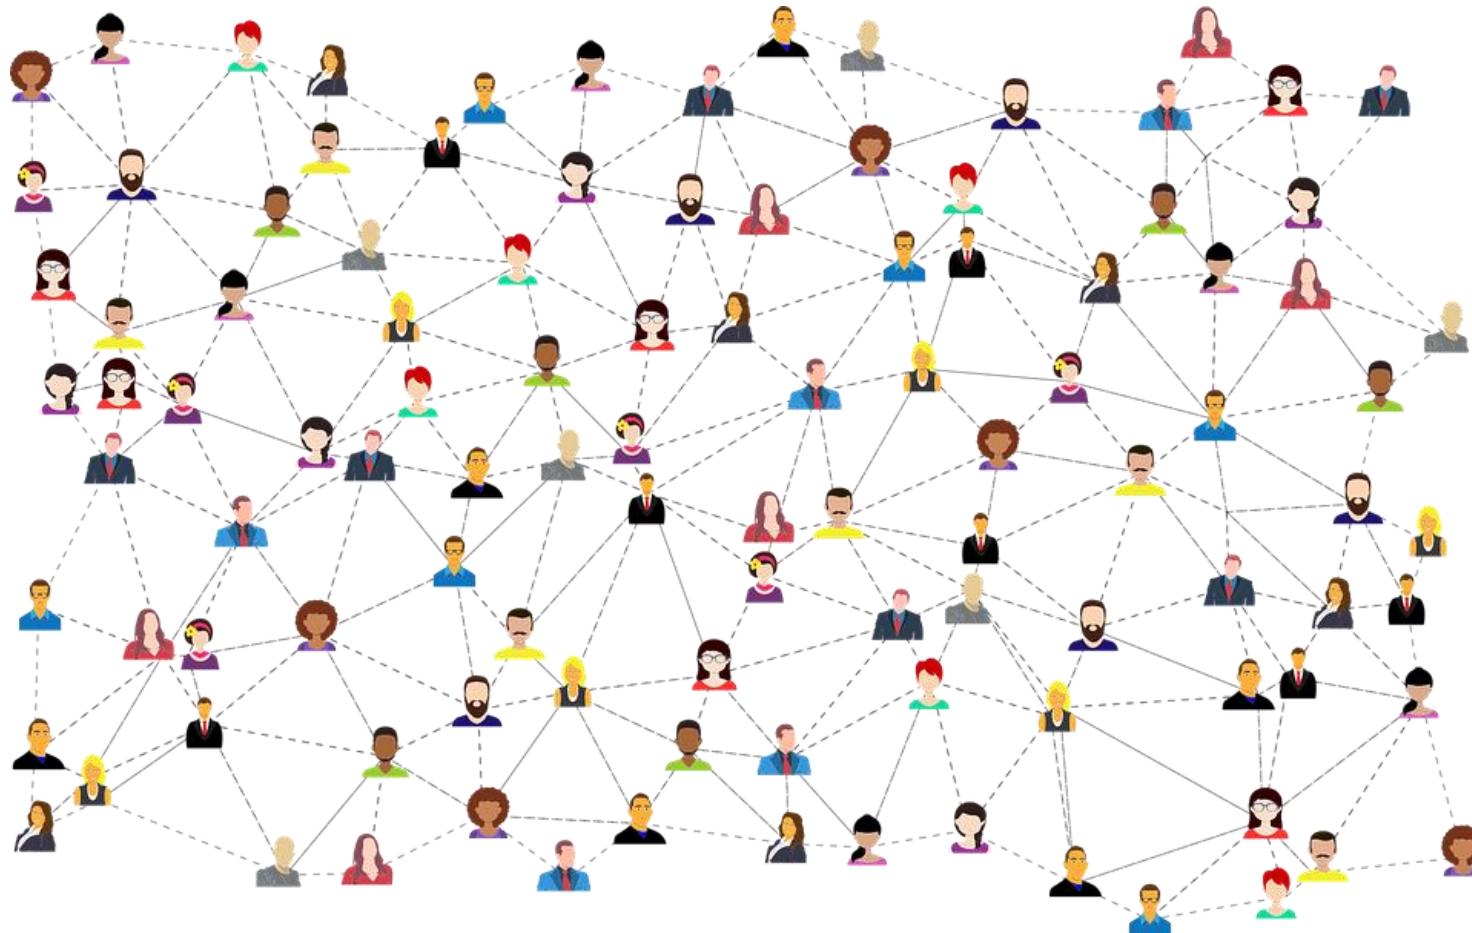
Make a prediction for each individual node.

Akin to segmentation for images.



# 3) Link prediction

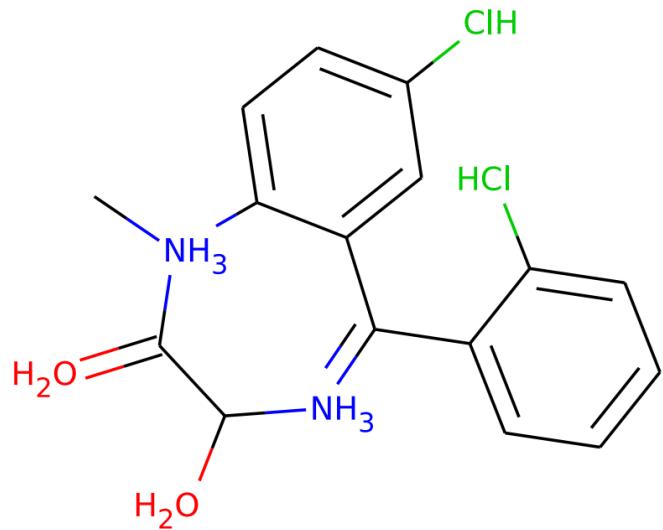
Make a prediction for each edge between two nodes.



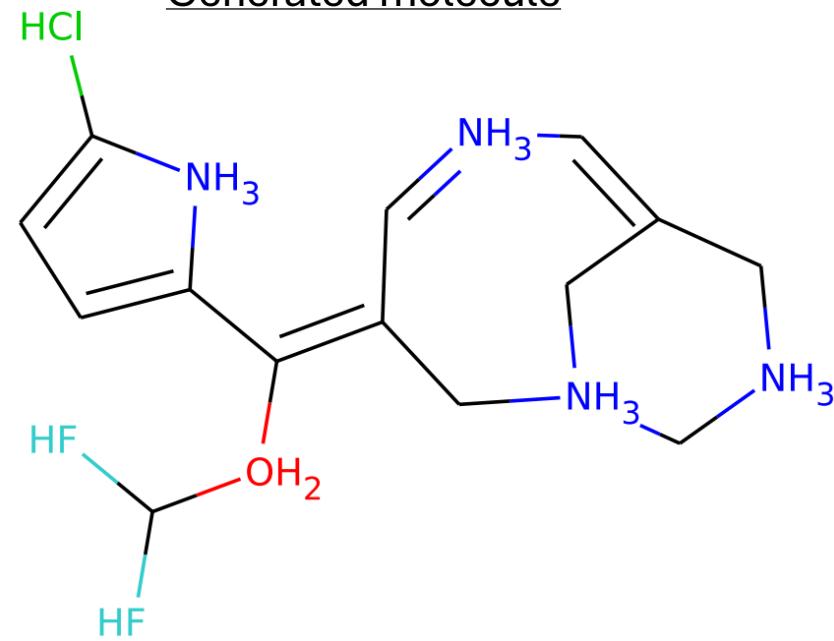
# 4) Graph generation

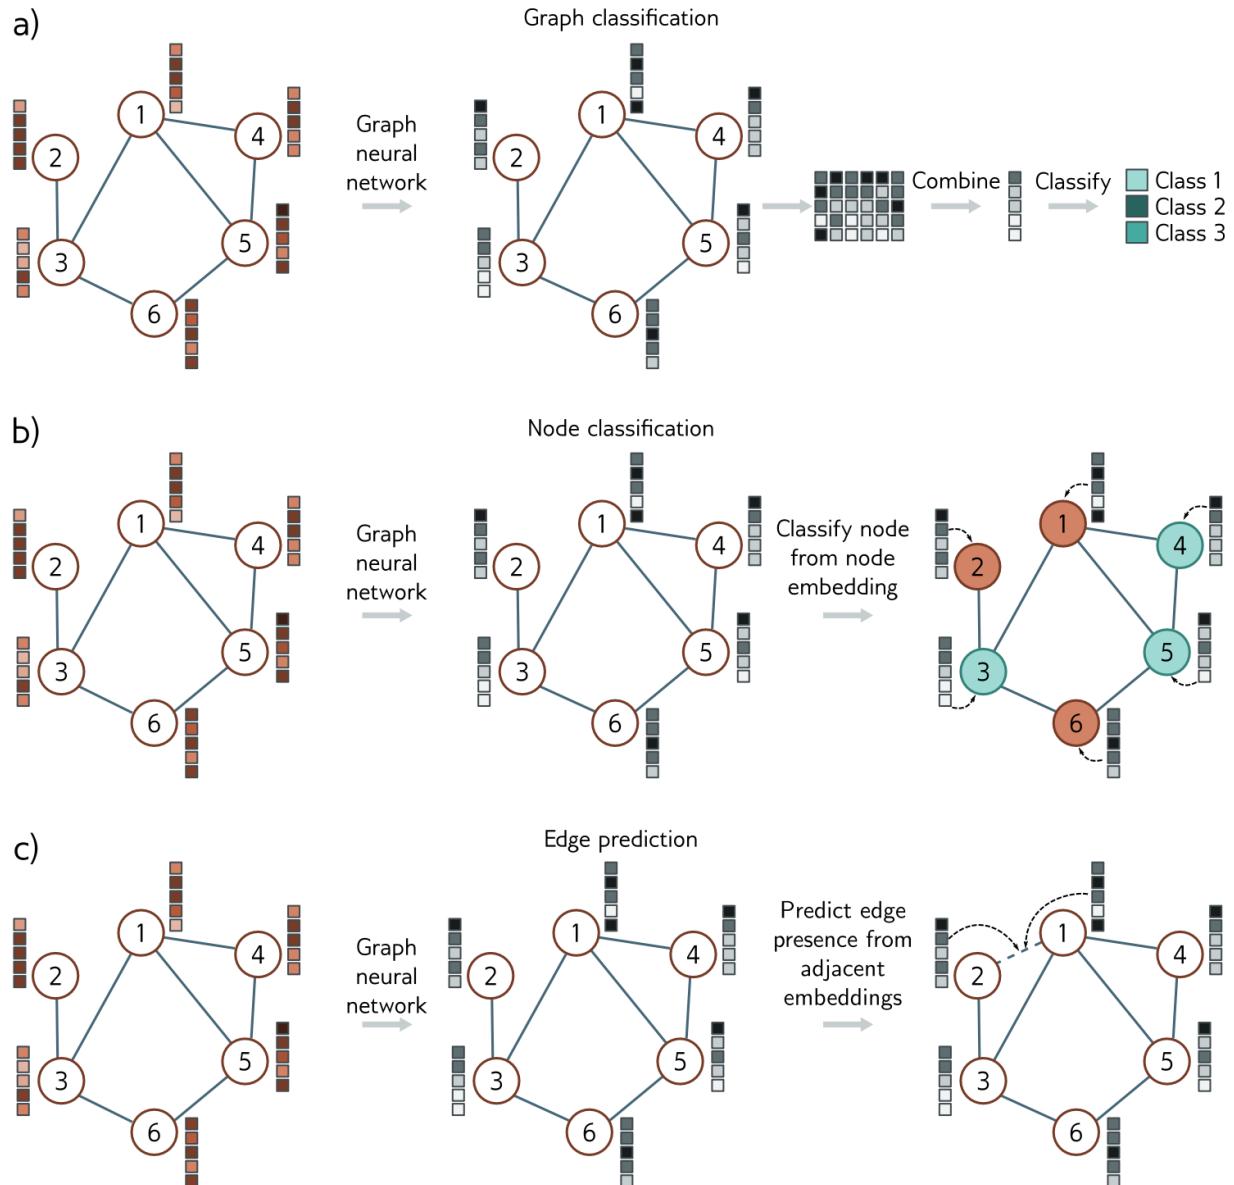
Similar in spirit to image/text generation, topic of next week.

Example molecule



Generated molecule

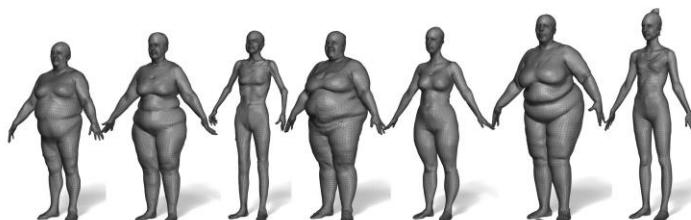
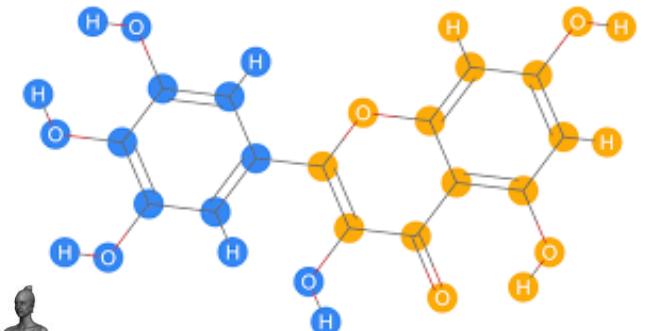




# Graphs can be dynamic

Graphs have fixed structures.

but the features are very variable, and also this is an assumption, the graph structure could change, nodes and edges could appear and disappear



But many are subject to change.

In practice, this change can even be gradual and continuous.

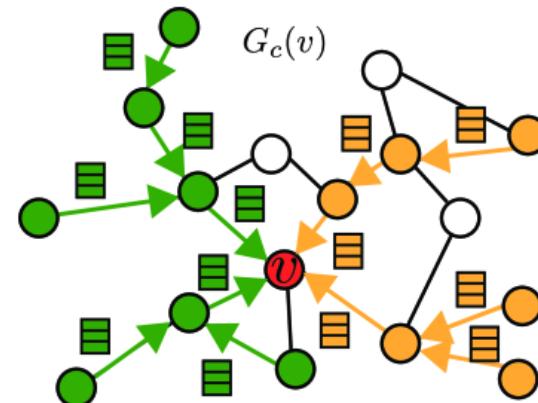


# Regular structures and graphs

Regular structures are a subset of graphs. I.e., images are grid graphs.



- Convolution + pooling
- Local neighborhood: fixed window
- Constant number of neighbors
- With fixed ordering
- Translation equivariance



- Message passing + coarsening
- Local neighborhood: 1-hop
- Different number of neighbors
- No ordering of neighbors
- Local permutation equivariance

Labelled graph	Degree matrix	Adjacency matrix	Laplacian matrix
	$\begin{pmatrix} 2 & 0 & 0 & 0 & 0 & 0 \\ 0 & 3 & 0 & 0 & 0 & 0 \\ 0 & 0 & 2 & 0 & 0 & 0 \\ 0 & 0 & 0 & 3 & 0 & 0 \\ 0 & 0 & 0 & 0 & 3 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 \end{pmatrix}$	$\begin{pmatrix} 0 & 1 & 0 & 0 & 1 & 0 \\ 1 & 0 & 1 & 0 & 1 & 0 \\ 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 & 1 & 1 \\ 1 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 \end{pmatrix}$	$\begin{pmatrix} 2 & -1 & 0 & 0 & -1 & 0 \\ -1 & 3 & -1 & 0 & -1 & 0 \\ 0 & -1 & 2 & -1 & 0 & 0 \\ 0 & 0 & -1 & 3 & -1 & -1 \\ -1 & -1 & 0 & -1 & 3 & 0 \\ 0 & 0 & 0 & -1 & 0 & 1 \end{pmatrix}$

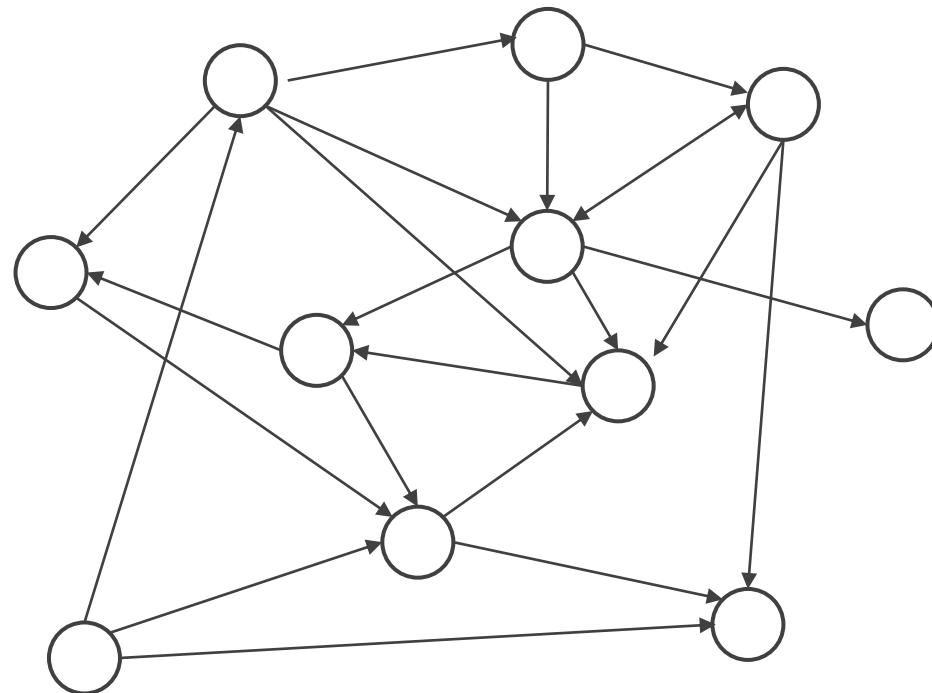
# Definition of a graph

(in deep learning)

# Directed graphs

Vertices  $\mathcal{V} = \{1, \dots, n\}$ , also called “nodes”

Edges  $\mathcal{E} = \{(i, j) : i, j \in \mathcal{V}\} \subseteq \mathcal{V} \times \mathcal{V}$  (directed)

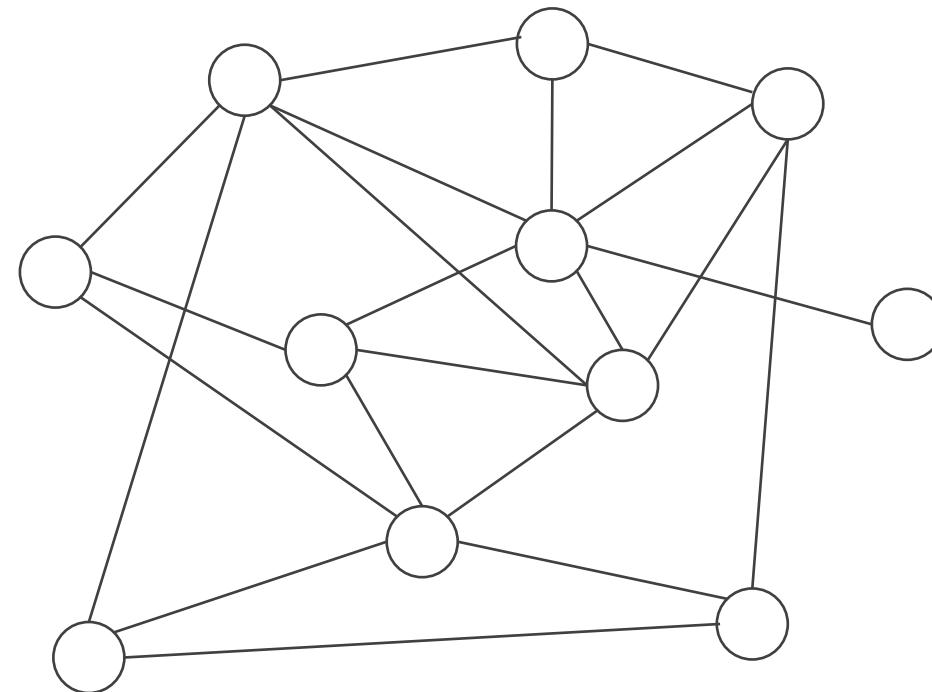


# Undirected graphs

Vertices  $\mathcal{V} = \{1, \dots, n\}$

Edges  $\mathcal{E} = \{(i, j) : i, j \in \mathcal{V}\} \subseteq \mathcal{V} \times \mathcal{V}$  (directed)

Edges  $\mathcal{E} = \{\{i, j\} : i, j \in \mathcal{V}\} \subseteq \mathcal{V} \times \mathcal{V}$  (undirected)



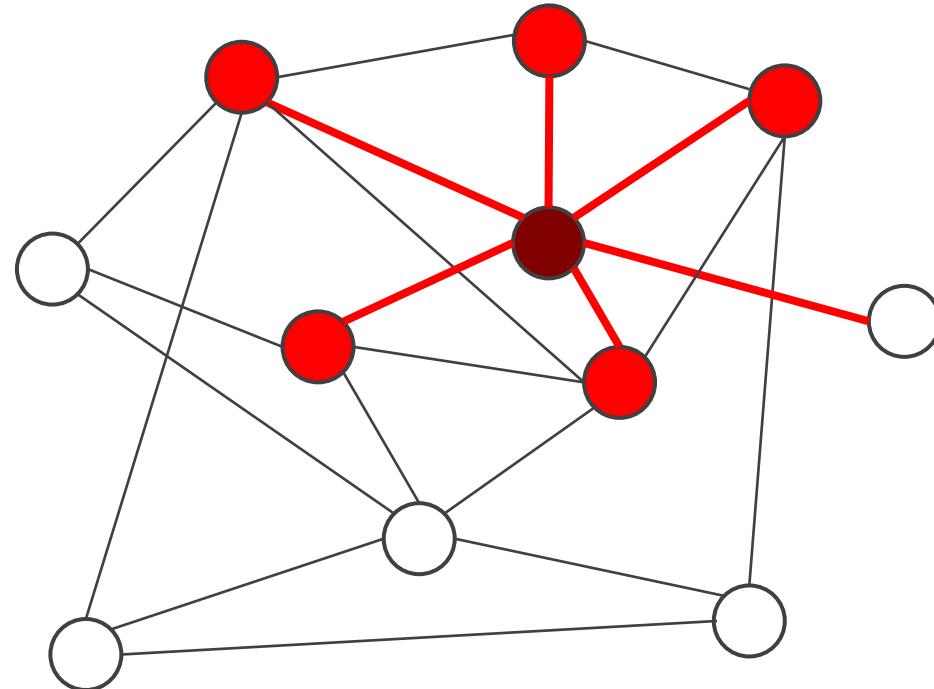
# Graph neighborhood

The neighborhood of a node consists of all nodes directly connected to it

$$\mathcal{N}(i) = \{j: (i, j) \in \mathcal{E}\}$$

The **degree** of a node is the number of neighbors:  $d_i = |\mathcal{N}(i)|$

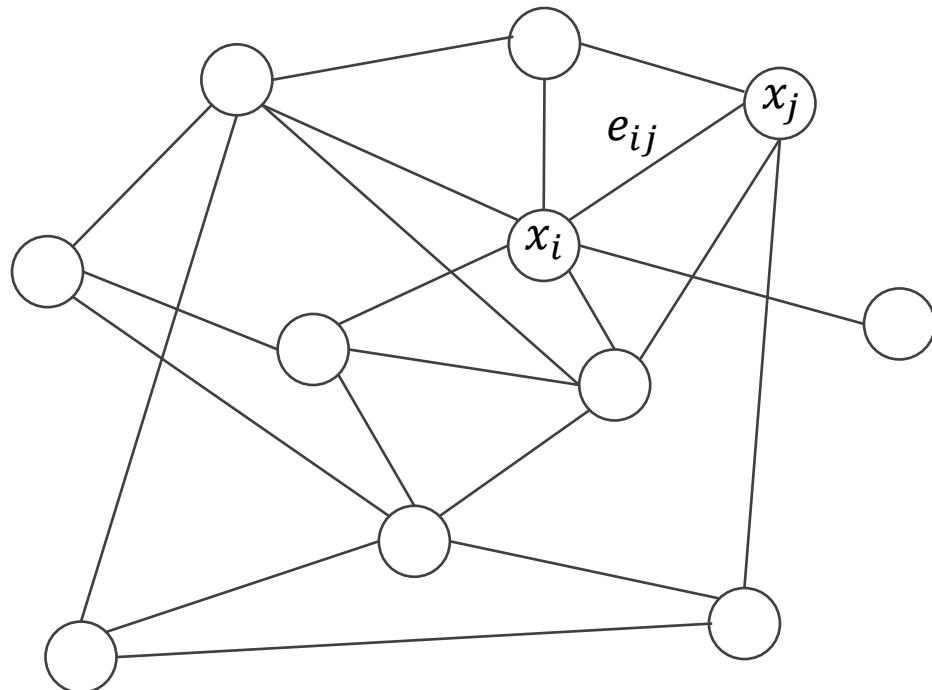
*The diagonal matrix  $D$  contains all degrees per node*



# Attributes

Node features  $\mathbf{x}: \mathcal{V} \rightarrow \mathbb{R}^d, X = (\mathbf{x}_1, \dots, \mathbf{x}_n)$

Edge features  $e_{ij}: \mathcal{E} \rightarrow \mathbb{R}^{d'}$



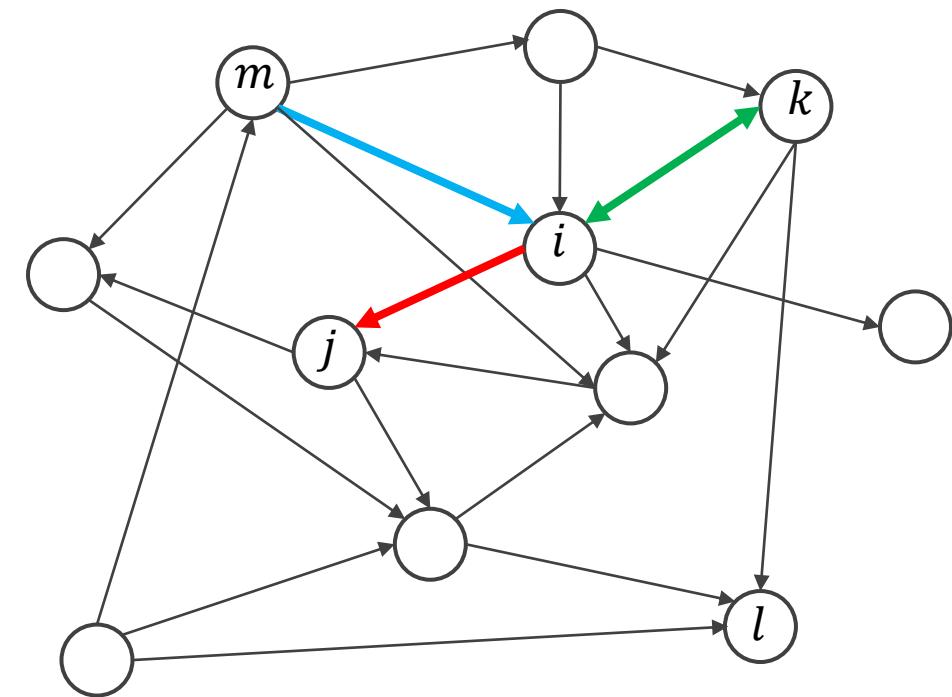
# Adjacency matrix

An  $n \times n$  matrix  $A$ , for  $n$  nodes

$$A_{ij} = \begin{cases} 1 & \text{if } (i,j) \in \varepsilon \\ 0 & \text{if } (i,j) \notin \varepsilon \end{cases}$$

$(A^z)_{ij}$ : number of paths that go from  $i$  to  $j$  in  $z$  steps

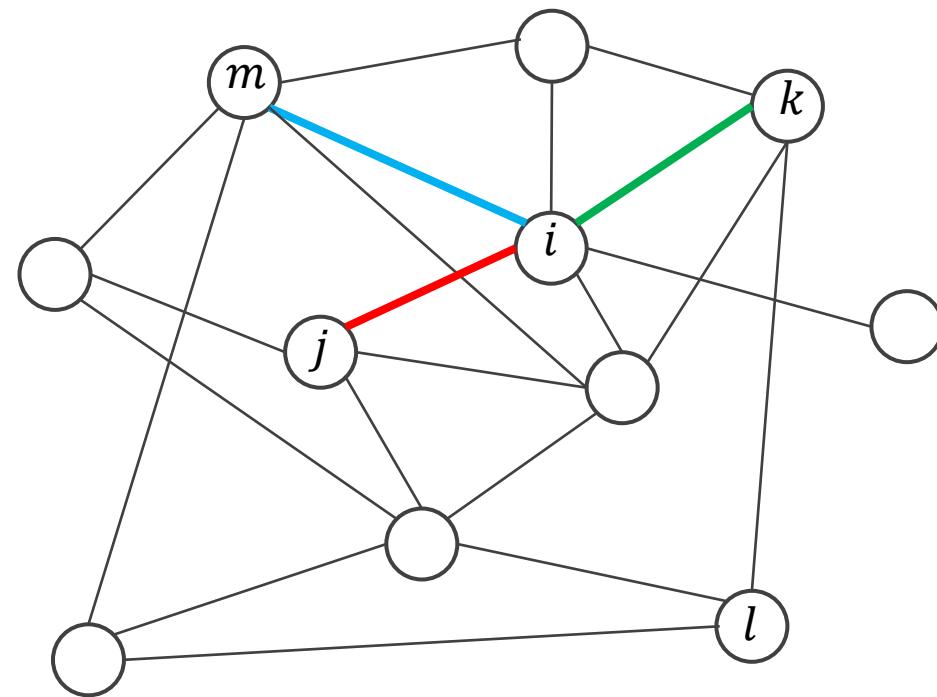
	$i$	$j$	$k$	$l$	$m$
$i$		1	1	0	
$j$					
$k$	1				
$l$	0			1	
$m$					



# Adjacency matrix for undirected graphs

The adjacency matrix is symmetric for undirected graphs.

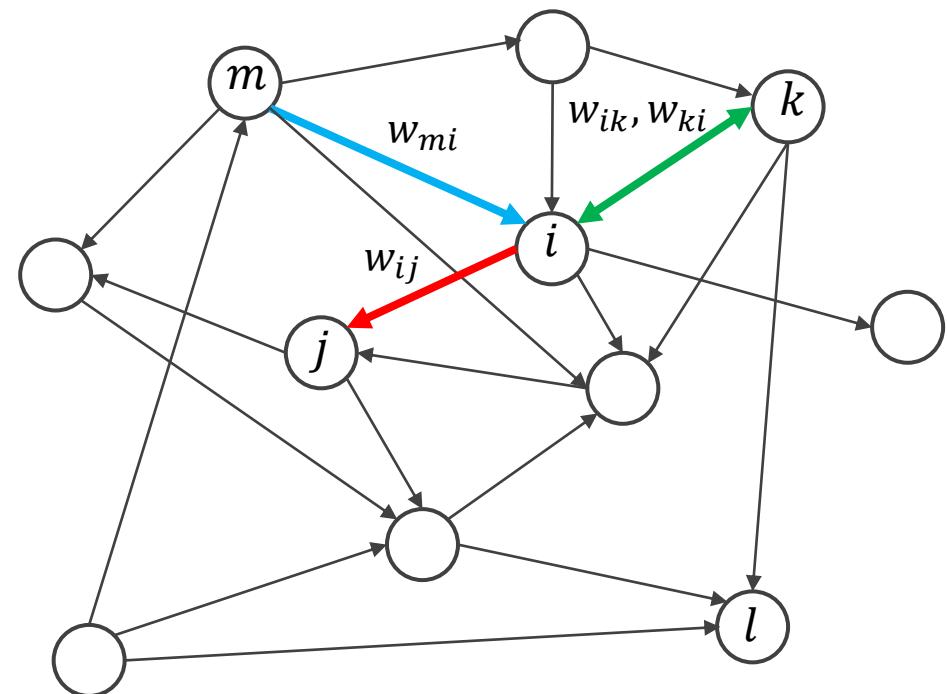
	$i$	$j$	$k$	$l$	$m$
$i$		1	1	0	1
$j$	1				
$k$		1			
$l$		0			
$m$		1			



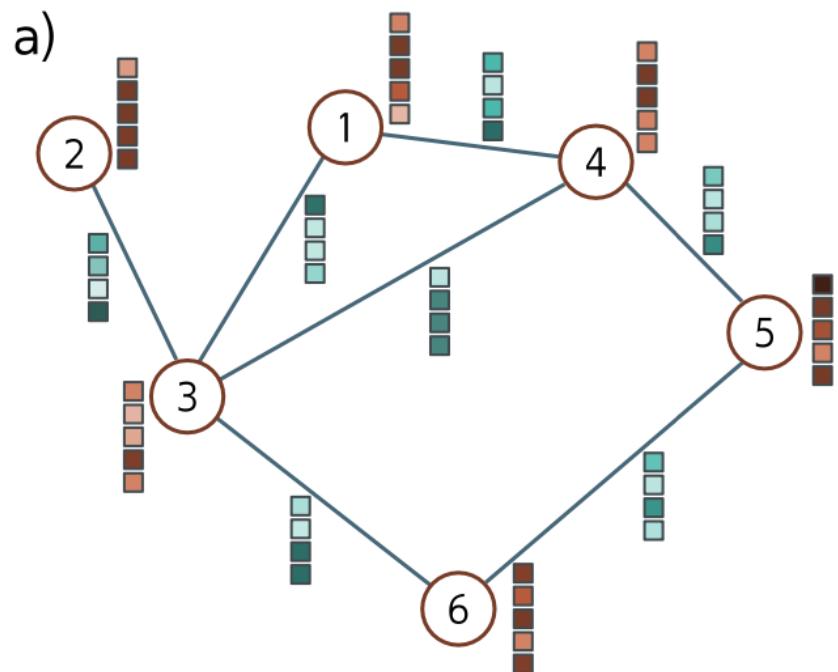
# Weighted adjacency matrix

When the edges have weights, so does the adjacency matrix.

	$i$	$j$	$k$	$l$	$m$
$i$		$w_{ij}$	$w_{ik}$	0	
$j$					
$k$		$w_{ki}$			
$l$	0				
$m$			$w_{mi}$		



# Final graph input representation



b)

Adjacency matrix,  $\mathbf{A}$   
 $N \times N$

	1	2	3	4	5	6
1	■	□	□	■	□	□
2	□	■	□	■	□	□
3	■	■	■	□	■	■
4	■	■	■	■	■	□
5	■	■	■	■	■	■
6	■	■	■	■	■	■

c)

Node data,  $\mathbf{X}$   
 $D \times N$

1	2	3	4	5	6
■	■	■	■	■	■
■	■	■	■	■	■
■	■	■	■	■	■
■	■	■	■	■	■
■	■	■	■	■	■
■	■	■	■	■	■

d)

Edge data,  $\mathbf{E}$   
 $D_E \times E$

1	1	2	3	3	4	5
3	■	■	■	■	■	■
4	■	■	■	■	■	■
3	■	■	■	■	■	■
6	■	■	■	■	■	■
5	■	■	■	■	■	■
6	■	■	■	■	■	■

# Back to graph networks

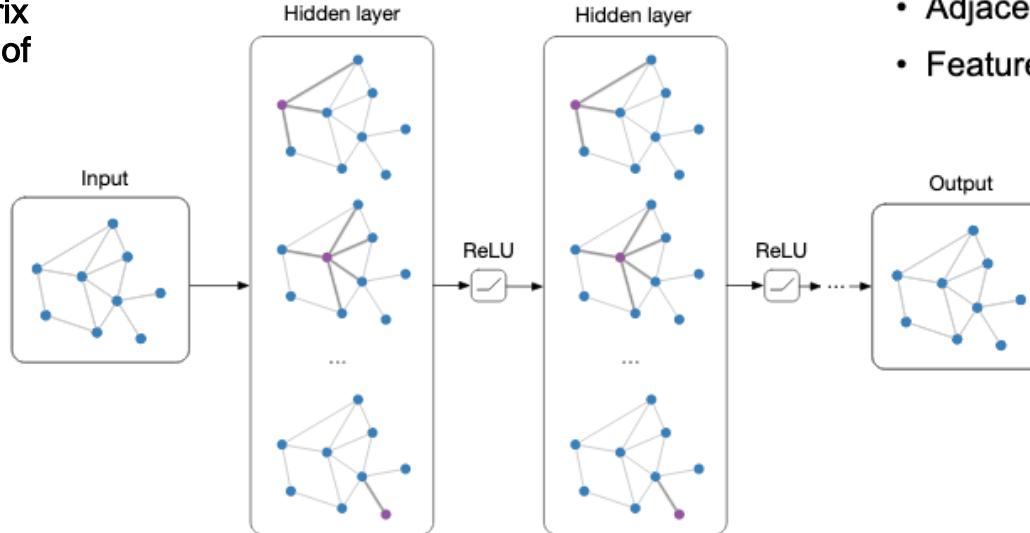
We can do a lot of processing on this data structure.

But the pre-defined features are raw inputs.

Graph networks do 1 thing: transform the feature vector per node over layers.

## The bigger picture:

every layer we use the adjacency matrix  
to help changing the representation of  
the nodes (feature matrix)

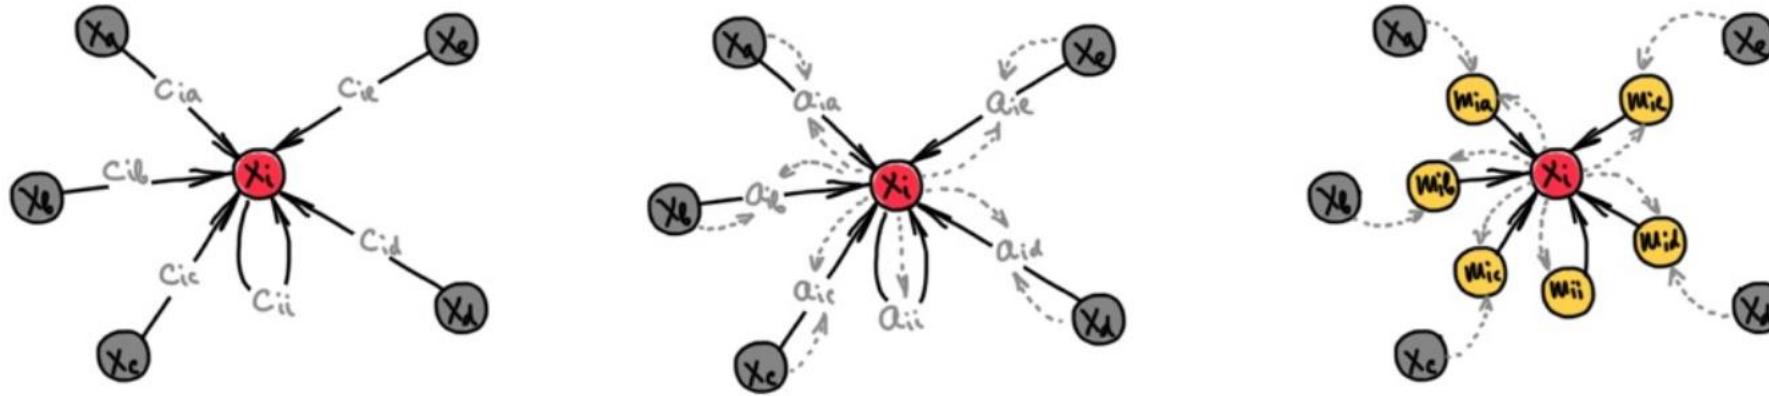


## Notation: $\mathcal{G} = (\mathbf{A}, \mathbf{X})$

- Adjacency matrix  $\mathbf{A} \in \mathbb{R}^{N \times N}$
- Feature matrix  $\mathbf{X} \in \mathbb{R}^{N \times F}$

In the end we obtain an NxN  
output feature matrix

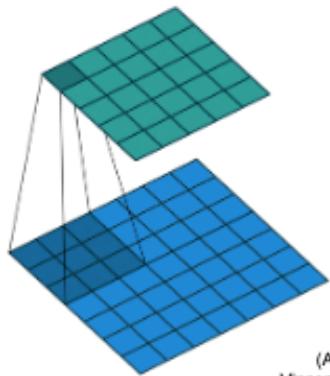
# Three perspectives to graph networks



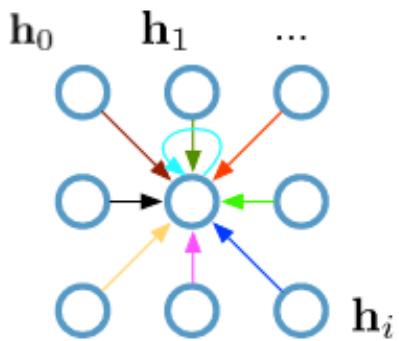
Three “flavours” of GNNs, left-to-right: convolutional, attentional, and general nonlinear message passing flavours. All are forms of message passing. Figure adapted from P. Veličković.

# Graph layer as a convolution layer

**Single CNN layer  
with 3x3 filter:**



(Animation by  
Vincent Dumoulin)



**Update for a single pixel:**

- Transform messages individually  $\mathbf{W}_i h_i$
- Add everything up  $\sum_i \mathbf{W}_i h_i$

$h_i \in \mathbb{R}^F$  are (hidden layer) activations of a pixel/node

**Full update:**

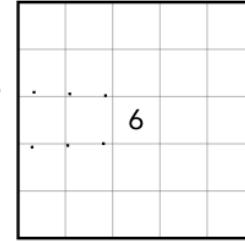
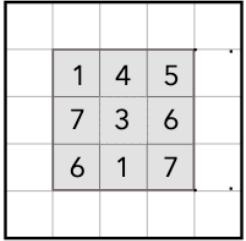
$$\mathbf{h}_4^{(l+1)} = \sigma \left( \mathbf{W}_0^{(l)} \mathbf{h}_0^{(l)} + \mathbf{W}_1^{(l)} \mathbf{h}_1^{(l)} + \dots + \mathbf{W}_8^{(l)} \mathbf{h}_8^{(l)} \right)$$

# Which assumptions from images are no longer valid?

Number of neighbors per node no longer fixed.

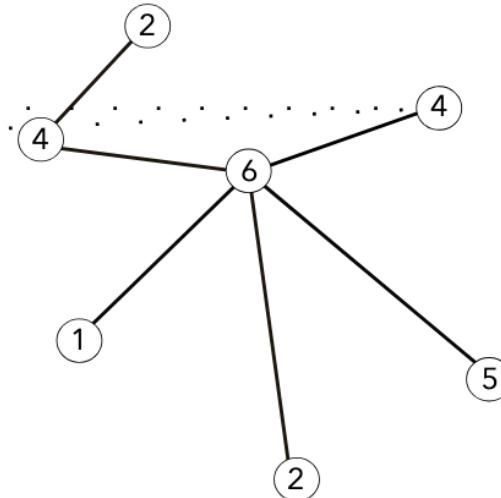
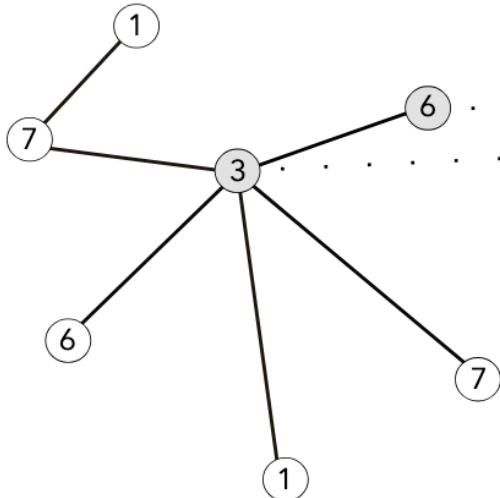
No more ordering between neighbours.

# Extending convolutions to graphs



Convolution in CNNs

CNNs perform localized convolutions. Neighbours participating in the convolution at the center pixel are highlighted in gray.



Localized Convolution in GNNs

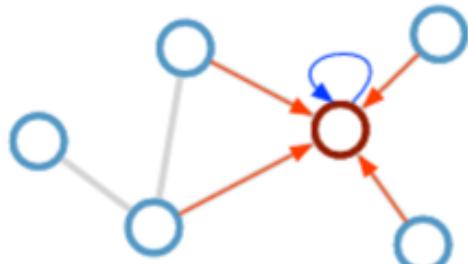
GNNs can perform localized convolutions mimicking CNNs. Hover over a node to see its immediate neighbourhood highlighted on the left. The structure of this neighbourhood changes from node to node.

# Graph convolution layer

Consider this undirected graph:



Calculate update for node in red:



Update rule:  $\mathbf{h}_i^{(l+1)} = \sigma \left( \mathbf{h}_i^{(l)} \mathbf{W}_0^{(l)} + \sum_{j \in \mathcal{N}_i} \frac{1}{c_{ij}} \mathbf{h}_j^{(l)} \mathbf{W}_1^{(l)} \right)$

Scalability: subsample messages [Hamilton et al., NIPS 2017]

**Desirable properties:**

- Weight sharing over all locations
- Invariance to permutations
- Linear complexity  $O(E)$
- Applicable both in transductive and inductive settings

**Limitations:**

- Requires gating mechanism / residual connections for depth
- Only indirect support for edge features

for one layer we define one degree of separation.

adding layer we increase the degree of separation i.e how much information one node capture from distant nodes

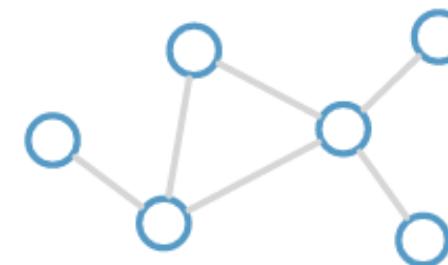
# Stacking graph convolution layers

Each layer aggregates information from their direct neighbors.

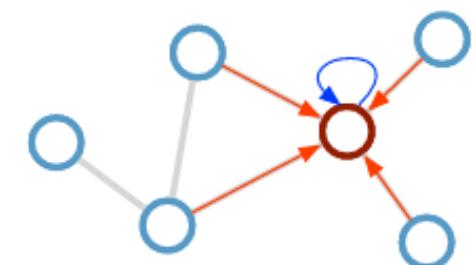
At the end of each layer, we add a non-linearity such as a ReLU.

We can increase complexity and receptive field simply by stacking multiple layers.

Consider this undirected graph:



Calculate update for node in red:



**Update rule:** 
$$\mathbf{h}_i^{(l+1)} = \sigma \left( \mathbf{h}_i^{(l)} \mathbf{W}_0^{(l)} + \sum_{j \in \mathcal{N}_i} \frac{1}{c_{ij}} \mathbf{h}_j^{(l)} \mathbf{W}_1^{(l)} \right)$$

# Graph convolution layer in matrix form

$$f(X, A) := \sigma \left( D^{-1/2} (A + I) D^{-1/2} X W \right)$$

$A \in \mathbb{R}^{n \times n}$  := The adjacency matrix

$I \in \mathbb{R}^{n \times n}$  := The identity matrix

$D \in \mathbb{R}^{n \times n}$  := The degree matrix of  $A + I$

$X \in \mathbb{R}^{n \times d}$  := The input data (i.e., the per-node feature vectors)

$W \in \mathbb{R}^{d \times w}$  := The layer's weights

$\sigma(\cdot)$  := The activation function (e.g., ReLU)

# Let's break it down

$$f(\mathbf{X}, \mathbf{A}) := \sigma\left(\mathbf{D}^{-1/2}(\mathbf{A} + \mathbf{I})\mathbf{D}^{-1/2}\mathbf{XW}\right)$$

The equation is enclosed in a large dashed rectangular bracket. Inside this bracket, there are three smaller nested brackets: a top bracket labeled "Add self-loops", a middle bracket labeled "Normalize adjacency matrix", and a bottom bracket labeled "Aggregate". To the right of the "Aggregate" bracket, there is another small bracket labeled "Update".

# Let's break it down

Add ones to diagonal, needed because each node should pass its own vector through.

the diagonal of adjacency matrix has all zero, but we want to use the informations about self nodes so we add a diagonal of 1 with ones

$$f(\mathbf{X}, \mathbf{A}) := \sigma\left(\mathbf{D}^{-1/2} (\mathbf{A} + \mathbf{I}) \mathbf{D}^{-1/2} \mathbf{XW}\right)$$

The diagram illustrates the components of the equation  $f(\mathbf{X}, \mathbf{A})$ . It consists of three nested dashed boxes:

- The innermost box contains the term  $(\mathbf{A} + \mathbf{I})$ , with the identity matrix  $\mathbf{I}$  highlighted by a red rectangular box. This is labeled "Add self-loops".
- The middle box contains the term  $\mathbf{D}^{-1/2} (\mathbf{A} + \mathbf{I}) \mathbf{D}^{-1/2}$ , with the entire expression enclosed in parentheses. This is labeled "Normalize adjacency matrix".
- The outermost box contains the full expression  $\sigma(\mathbf{D}^{-1/2} (\mathbf{A} + \mathbf{I}) \mathbf{D}^{-1/2} \mathbf{XW})$ . This is labeled "Aggregate".

Below the middle box, there is an additional label "Update".

# Let's break it down

This step essentially normalizes the adjacency matrix. I will show how in a few slides.

The matrix  $D$  is a diagonal matrix with the degree of the nodes on the diagonal. nodes with huge grades obtain a higher sum from neighbours, this creates an imbalance in activations between nodes that needs to be normalized

$$f(\mathbf{X}, \mathbf{A}) := \sigma\left(\boxed{\mathbf{D}^{-1/2}} (\mathbf{A} + \mathbf{I}) \boxed{\mathbf{D}^{-1/2}} \mathbf{XW}\right)$$

The diagram illustrates the components of the equation:

- $\boxed{\mathbf{D}^{-1/2}}$ : Add self-loops
- $(\mathbf{A} + \mathbf{I})$ : Normalize adjacency matrix
- $\boxed{\mathbf{D}^{-1/2}}$ : Aggregate
- $\mathbf{XW}$ : Update

# Let's break it down

Just a standard linear layer and a non-linearity.

Cant see the convolution here

$$f(\mathbf{X}, \mathbf{A}) := \sigma(\mathbf{D}^{-1/2}(\mathbf{A} + \mathbf{I})\mathbf{D}^{-1/2}\mathbf{XW})$$

The diagram illustrates the components of the function  $f(\mathbf{X}, \mathbf{A})$  from left to right:

- Add self-loops:** A bracket under the term  $\mathbf{A} + \mathbf{I}$  indicates the addition of identity matrices to all nodes.
- Normalize adjacency matrix:** A bracket under the term  $\mathbf{D}^{-1/2}(\mathbf{A} + \mathbf{I})\mathbf{D}^{-1/2}$  indicates the normalization of the adjacency matrix by its degree matrix.
- Aggregate:** A bracket under the term  $\mathbf{XW}$  indicates the aggregation of node features.
- Update:** A bracket under the entire expression  $\mathbf{XW}$  indicates the final update step.

Red boxes highlight the non-linearity  $\sigma$  and the linear layer  $\mathbf{XW}$ .

# Let's break it down

Just a standard linear layer and a non-linearity.

$$f(\mathbf{X}, \mathbf{A}) := \sigma\left(\boxed{\mathbf{D}^{-1/2}(\mathbf{A} + \mathbf{I})\mathbf{D}^{-1/2}\mathbf{XW}}\right)$$

The diagram illustrates the components of the equation  $f(\mathbf{X}, \mathbf{A})$ . It shows a red box around the term  $\mathbf{D}^{-1/2}(\mathbf{A} + \mathbf{I})\mathbf{D}^{-1/2}$ . Below this box, a bracket labeled "Add self-loops" points to the identity matrix  $\mathbf{I}$ . To the left of the red box, another bracket labeled "Normalize adjacency matrix" points to the product of  $\mathbf{D}^{-1/2}$  and  $\mathbf{A}$ . Below the red box, a bracket labeled "Aggregate" points to the product of  $\mathbf{D}^{-1/2}$  and  $\mathbf{XW}$ . At the bottom, a bracket labeled "Update" points to the entire expression  $\sigma(\text{red box})$ .

# Rewriting into 2 steps

$$\tilde{\mathbf{A}} := \mathbf{D}^{-1/2}(\mathbf{A} + \mathbf{I})\mathbf{D}^{-1/2}$$

$$\tilde{A}_{i,j} := \begin{cases} \frac{1}{\sqrt{d_{i,i}d_{j,j}}}, & \text{if there is an edge between node } i \text{ and } j \\ 0, & \text{otherwise} \end{cases}$$

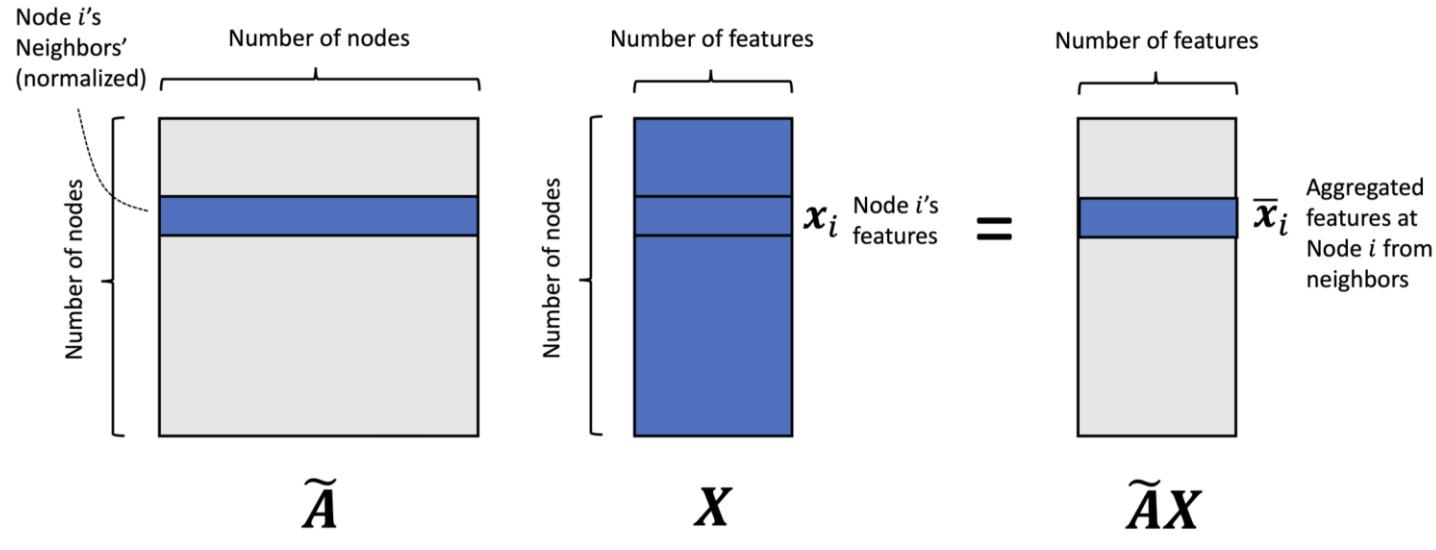
$$f(X, A) := \sigma(\tilde{\mathbf{A}}\mathbf{X}\mathbf{W})$$

$$\mathbf{D} := \begin{bmatrix} d_{1,1} & 0 & 0 & \dots & 0 \\ 0 & d_{2,2} & 0 & \dots & 0 \\ 0 & 0 & d_{3,3} & \dots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \dots & d_{n,n} \end{bmatrix}$$

$$\mathbf{D}^{-1/2} := \begin{bmatrix} \frac{1}{\sqrt{d_{1,1}}} & 0 & 0 & \dots & 0 \\ 0 & \frac{1}{\sqrt{d_{2,2}}} & 0 & \dots & 0 \\ 0 & 0 & \frac{1}{\sqrt{d_{3,3}}} & \dots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \dots & \frac{1}{\sqrt{d_{n,n}}} \end{bmatrix}$$

$$f(X, A) := \sigma(\tilde{A}XW)$$

# Left side of the equation



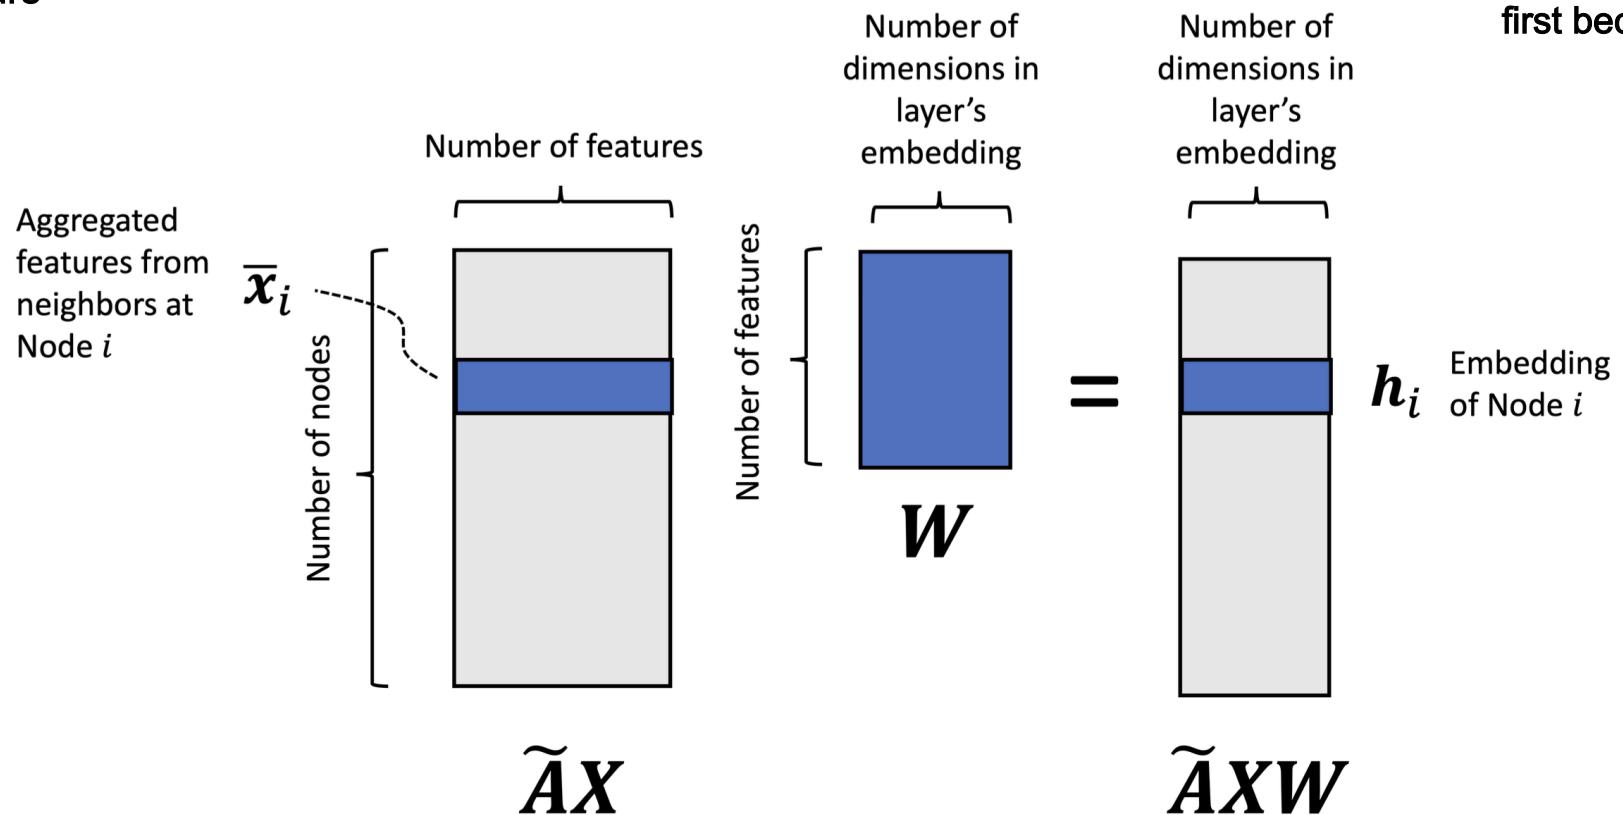
$$\begin{aligned}\bar{x}_i &= \sum_{j=1}^n \tilde{a}_{i,j} x_j \\ &= \sum_{j \in \text{Neigh}(i)} \tilde{a}_{i,j} x_j \\ &= \sum_{j \in \text{Neigh}(i)} \frac{1}{\sqrt{d_{i,i} d_{j,j}}} x_j\end{aligned}$$

For each feature  $i$  I'm summing all the values of that specific feature from the nodes in the neighborhood

$$f(X, A) := \sigma(\tilde{A}XW)$$

# Right side of the equation

the core concept is doing an MLP and summing the neighbours



It can be easier to think as a linear transformations of the whole feature and then apply to it the graph structure captured by  $A$

but it is more efficient to do  $AX$  first because  $A$  may be sparse

# Why add a normalization step?

Do we even need it? Let's see what happens without it:

$$\hat{A} := A + I$$

Normalization dropped, only  
self-loop retrained

$$f_{\text{unnormalized}}(X, A) := \sigma(\hat{A}XW)$$

Layer update remains the same

$$\bar{x}_i = \sum_{j=1}^n \hat{a}_{i,j}x_j$$

$$= \sum_{j=1}^n \mathbb{I}(j \in \text{Neigh}(i))x_j$$

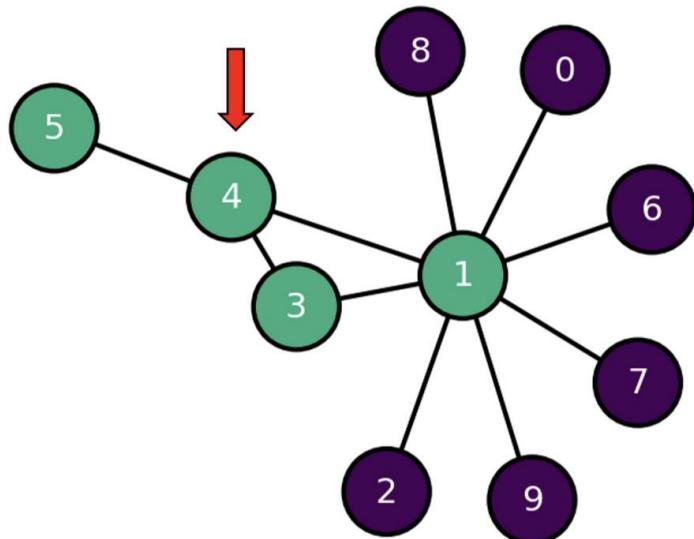
$$= \sum_{j \in \text{Neigh}(i)} x_j$$

Problem! More neighbors = bigger sum.  
Huge bias when training graph networks.

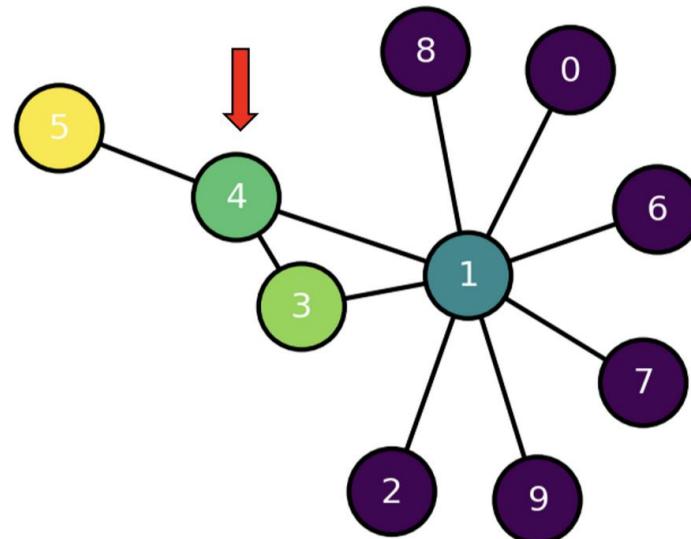
# Why not simply divide by the node degree?

this is a totally viable way  
to implement a graph  
neural network. here I have  
a uniform weighting

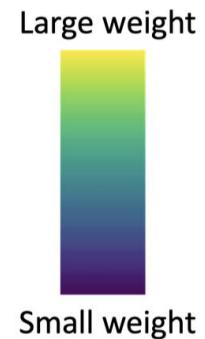
$$D^{-1}\hat{A}$$



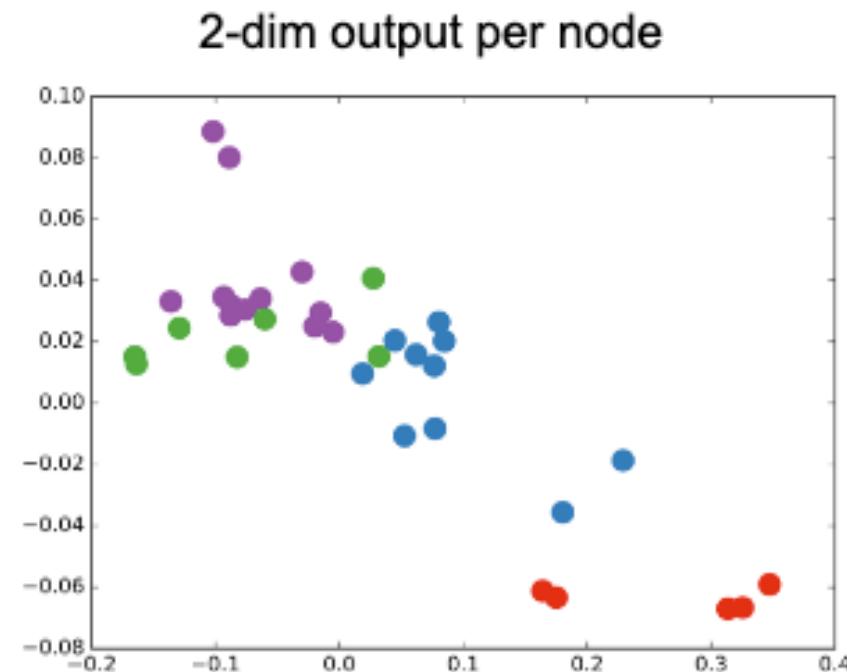
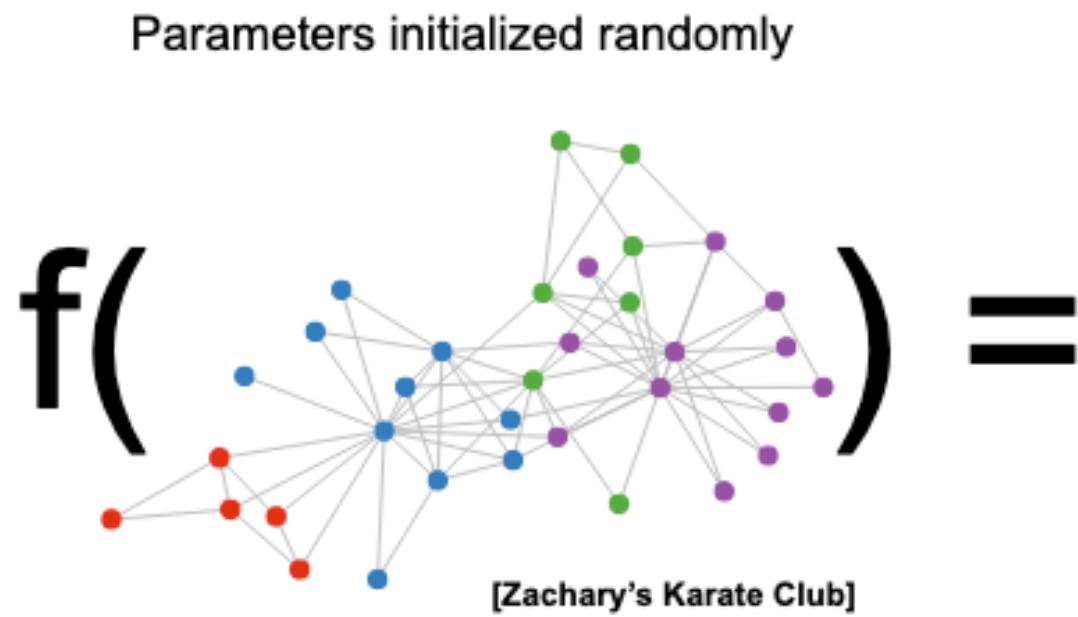
$$D^{-1/2}\hat{A}D^{-1/2}$$



but this one, use the joint  
degree between edges. It  
makes something more unique  
more important.  
This implementation is  
incorporating some second  
order information



# Visualizing node representations



## Alternative: graph layer as attention

Similar but including attention as *aggregation*:  $y_i = h\left(\sum_{j \in \mathcal{N}(i)} a_{ij} \mathbf{z}_j\right)$

Using self-attention:

$$a_{ij} = \frac{\exp(e_{ij})}{\sum_{k \in \mathcal{N}(i)} \exp(e_{ik})},$$

where  $e_{ij}$  are the self-attention weights (like query == key)

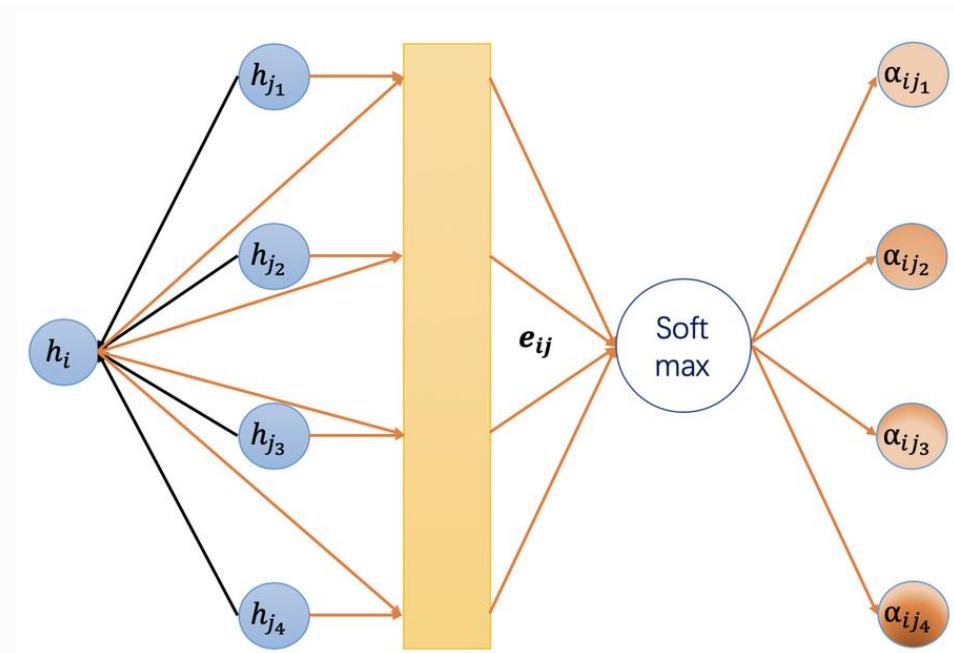
$$e_{ij} = \text{LeakyReLU}\left([\mathbf{x}_i \mathbf{W}, \mathbf{x}_j \mathbf{W}] \cdot \mathbf{u}\right)$$

$\mathbf{u}$  is a weight vector.

This bracket notation stands for concatenation

Flexibility in Learning Relationships: Concatenation allows the attention mechanism to learn asymmetric relationships between nodes. When you concatenate  $[\mathbf{x}_j \mathbf{W}, \mathbf{x}_i \mathbf{W}]$  the network can learn that the importance of node  $j$  to node  $i$  might be different from the importance of  $i$  to  $j$  even in undirected graphs.

# The four steps of a graph attention layer



$$z_i^{(l)} = W^{(l)} h_i^{(l)}, \quad (1)$$

$$e_{ij}^{(l)} = \text{LeakyReLU}(\vec{a}^{(l)T} (z_i^{(l)} || z_j^{(l)})), \quad (2)$$

$$\alpha_{ij}^{(l)} = \frac{\exp(e_{ij}^{(l)})}{\sum_{k \in \mathcal{N}(i)} \exp(e_{ik}^{(l)})}, \quad (3)$$

$$h_i^{(l+1)} = \sigma \left( \sum_{j \in \mathcal{N}(i)} \alpha_{ij}^{(l)} z_j^{(l)} \right), \quad (4)$$

# Connecting graphs, convolutions, and transformers

Transformers operate on a complete graph (adjacency matrix with all 1's).

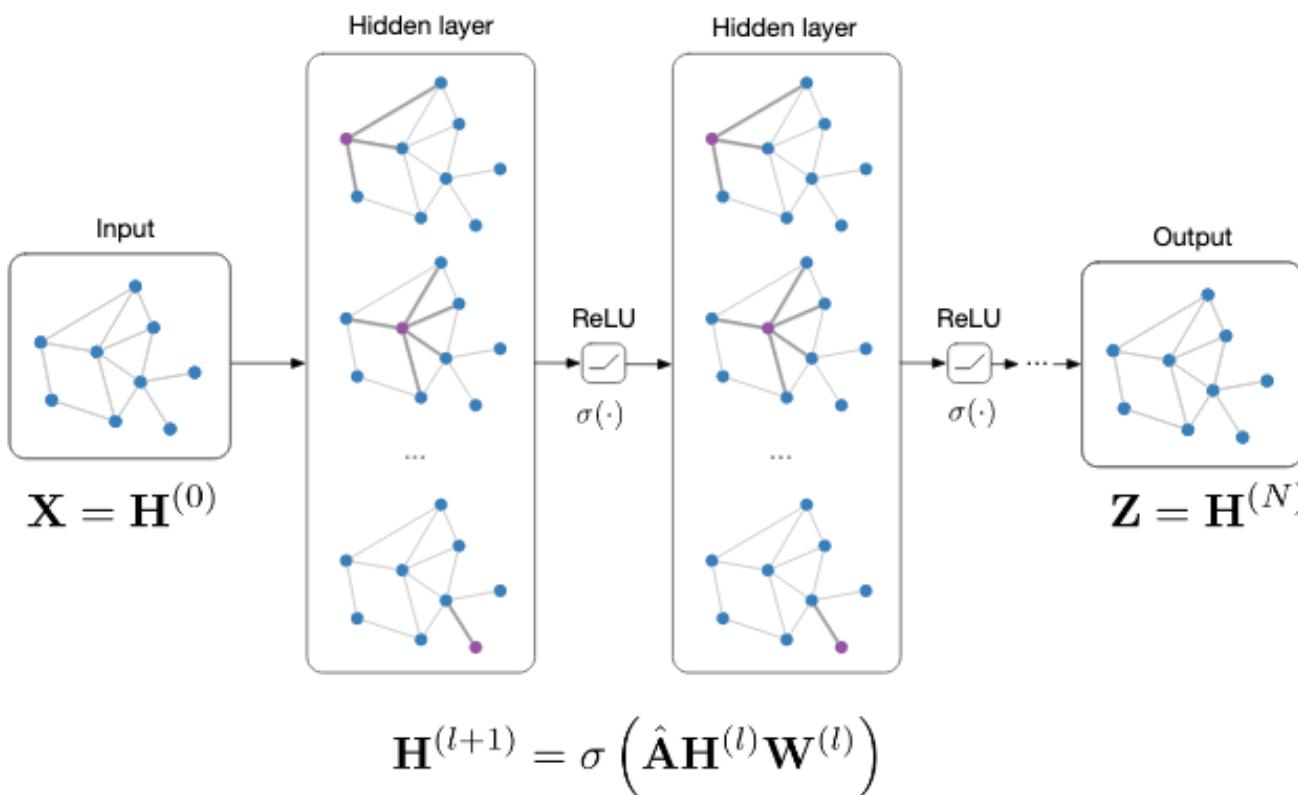
With attention-based GCN, we recover the Transformer.

equivariance:		
Architecture	Domain $\Omega$	Symmetry group $\mathfrak{G}$
CNN	Grid	Translation
<i>Spherical CNN</i>	Sphere / SO(3)	Rotation SO(3)
<i>Intrinsic / Mesh CNN</i>	Manifold	Isometry $\text{Iso}(\Omega)$ / Gauge symmetry SO(2)
GNN	Graph	Permutation $\Sigma_n$
<i>Deep Sets</i>	Set	Permutation $\Sigma_n$
<i>Transformer</i>	Complete Graph	Permutation $\Sigma_n$
LSTM	1D Grid	Time warping

# Optimizing graph networks

**Input:** Feature matrix  $\mathbf{X} \in \mathbb{R}^{N \times E}$ , preprocessed adjacency matrix  $\hat{\mathbf{A}}$

after K layers I have only new node feature, the structure of the graphs remain the same



**Node classification:**

$$\text{softmax}(\mathbf{z}_n)$$

e.g. Kipf & Welling (ICLR 2017)

**Graph classification:**

$$\text{softmax}(\sum_n \mathbf{z}_n)$$

e.g. Duvenaud et al. (NIPS 2015)

**Link prediction:**

$$p(A_{ij}) = \sigma(\mathbf{z}_i^T \mathbf{z}_j)$$

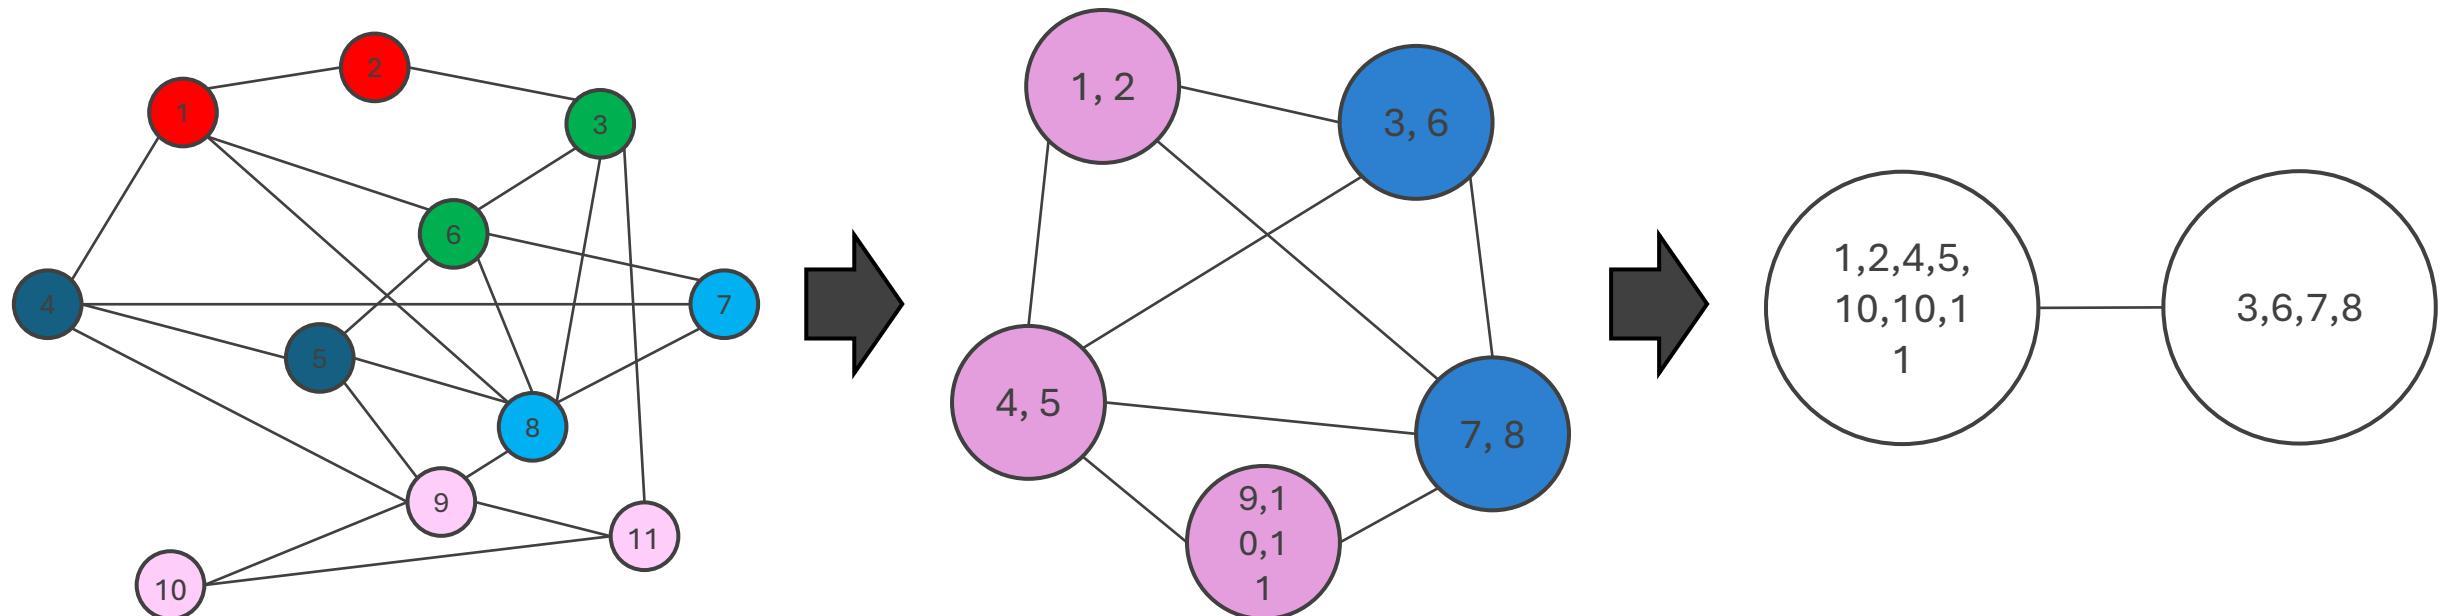
Kipf & Welling (NIPS BDL 2016)

“Graph Auto-Encoders”

# Pooling in graph networks

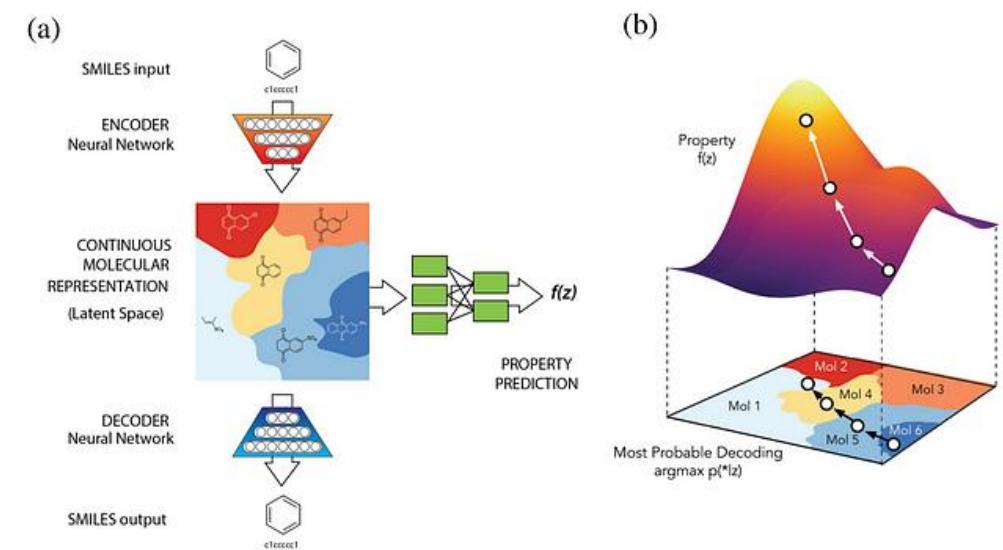
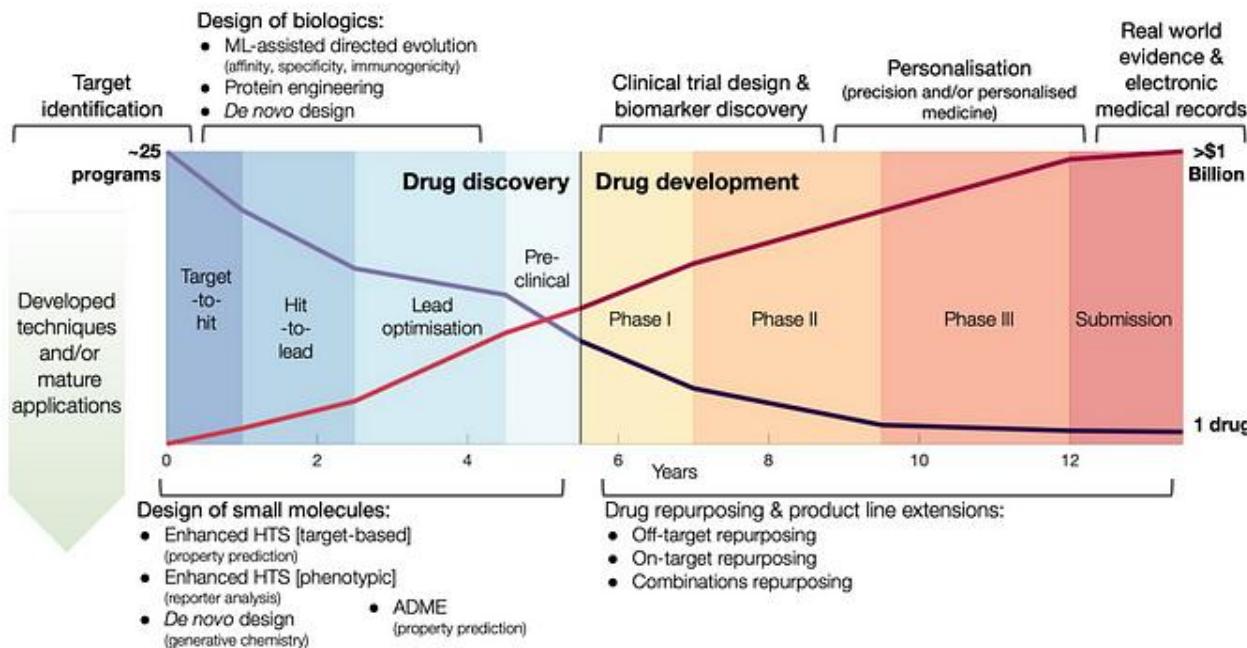
Specifically for graph classification, pooling is an optional operators.

Pool nodes together to save compute, requires updating the adjacency matrix.

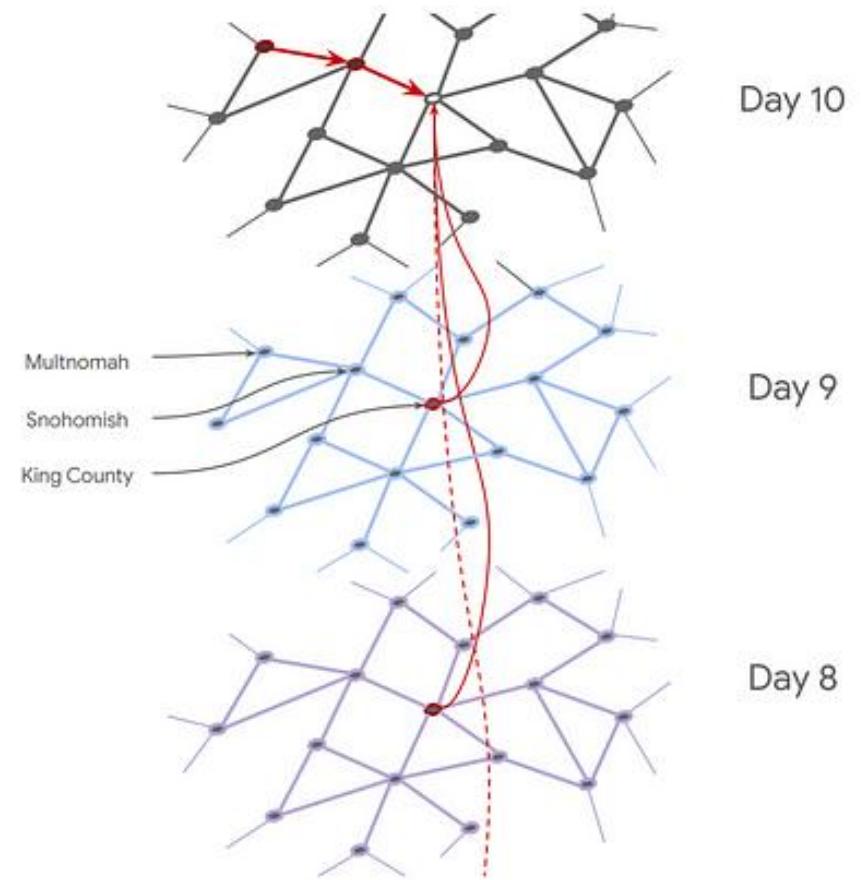
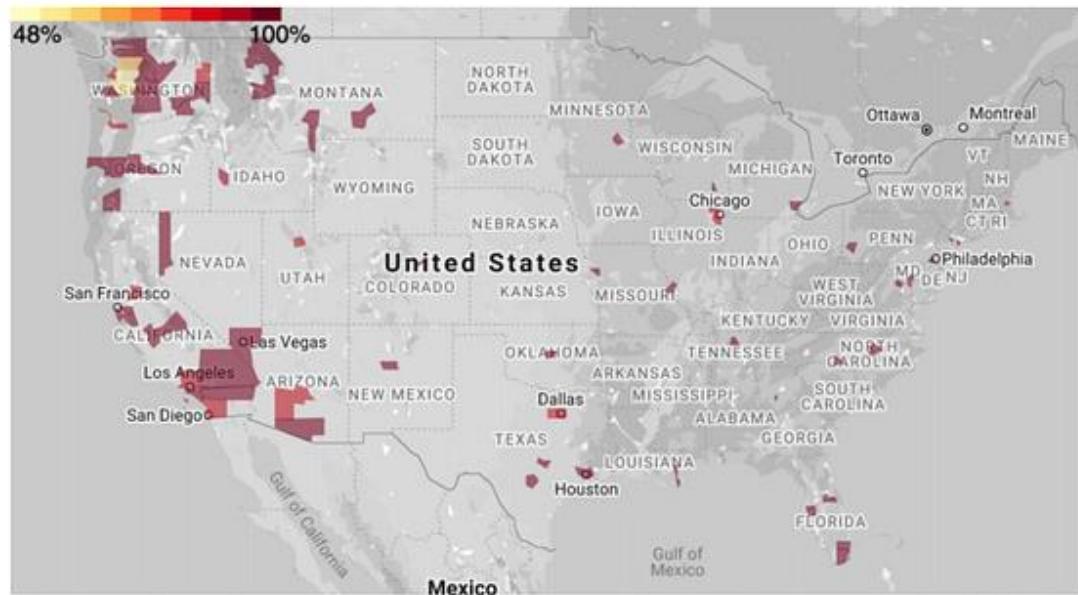


# Applications of graph networks

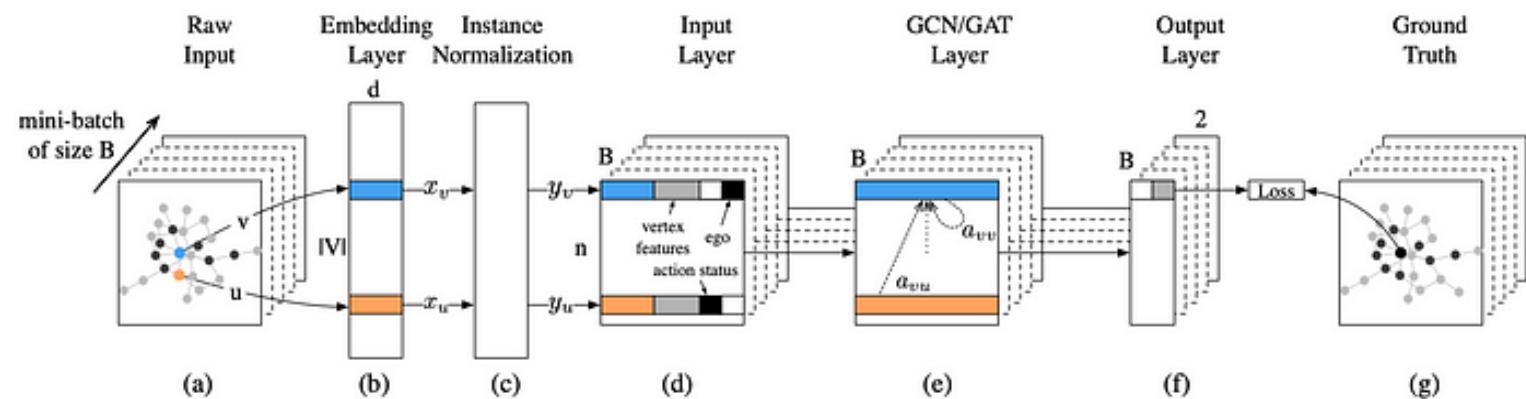
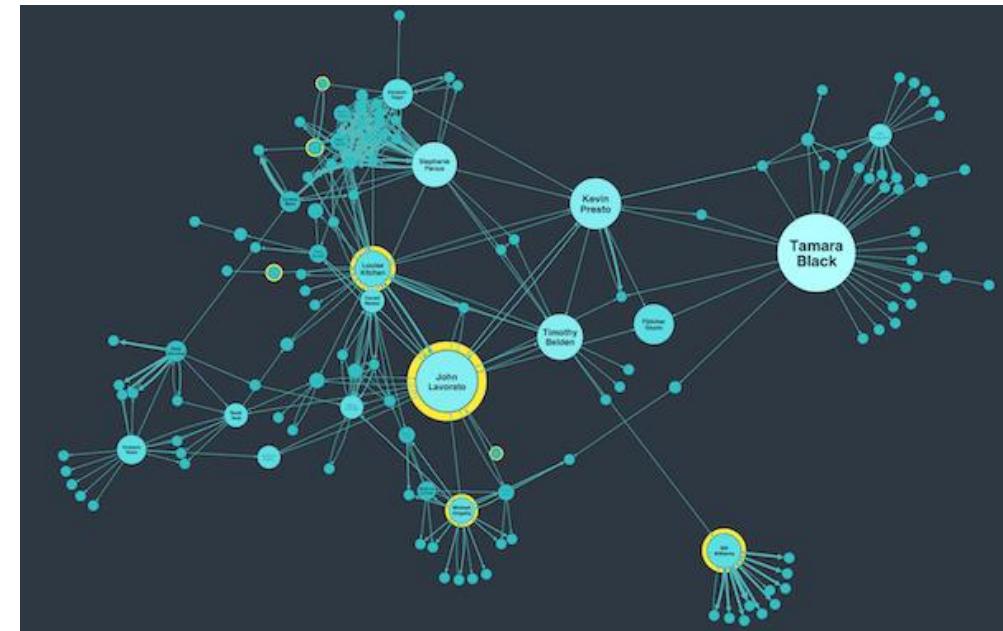
# Drug discovery



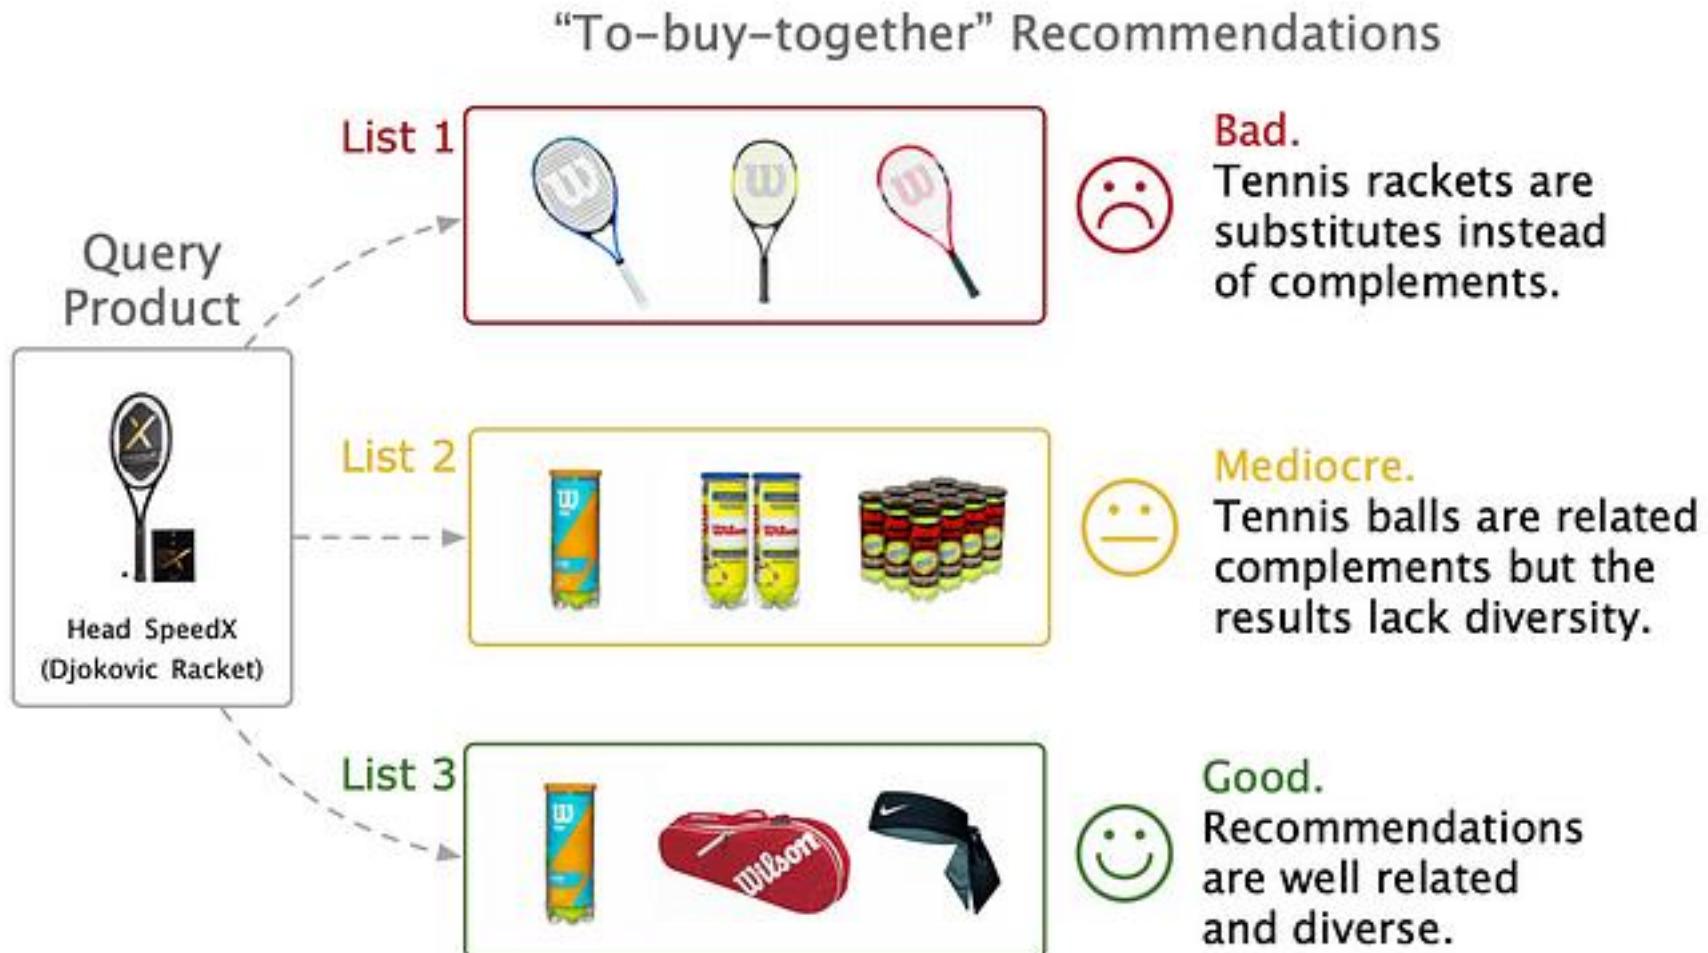
# Modeling the spread of deceases



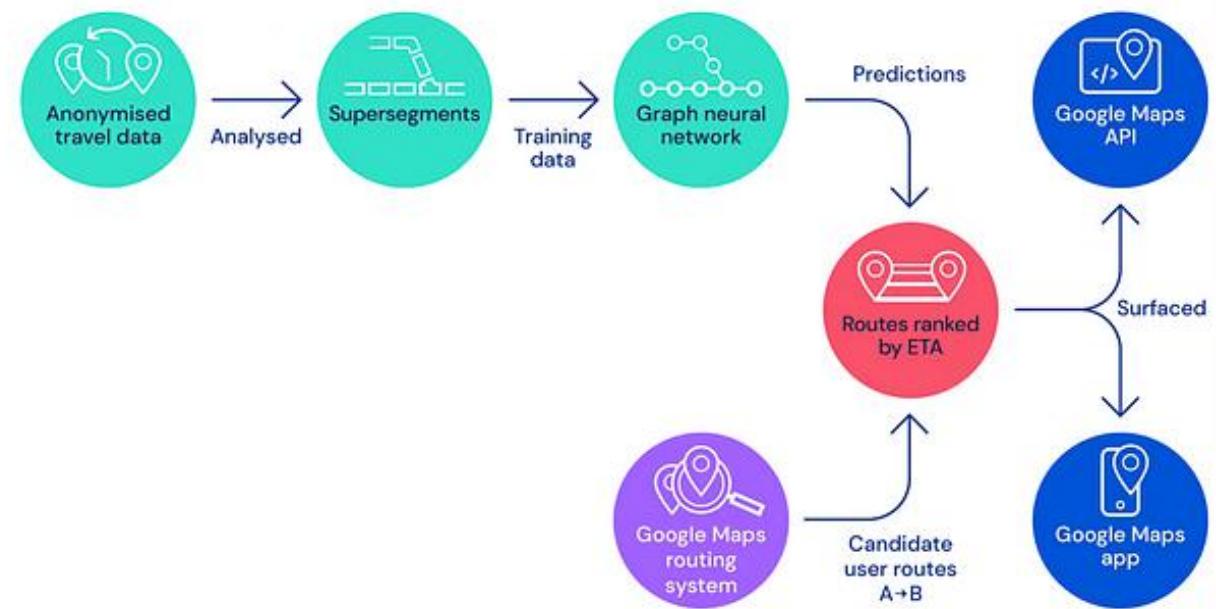
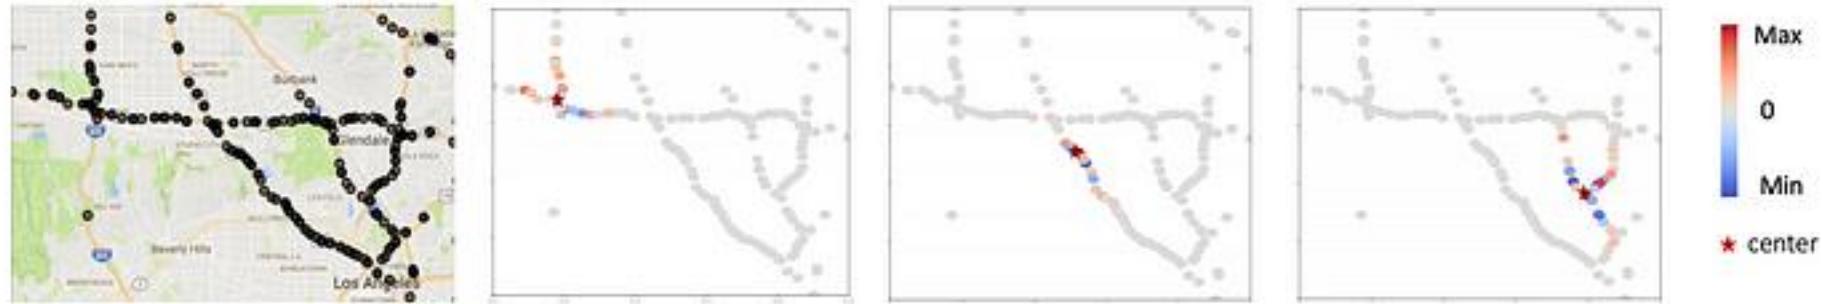
# Social networks



# Recommendation

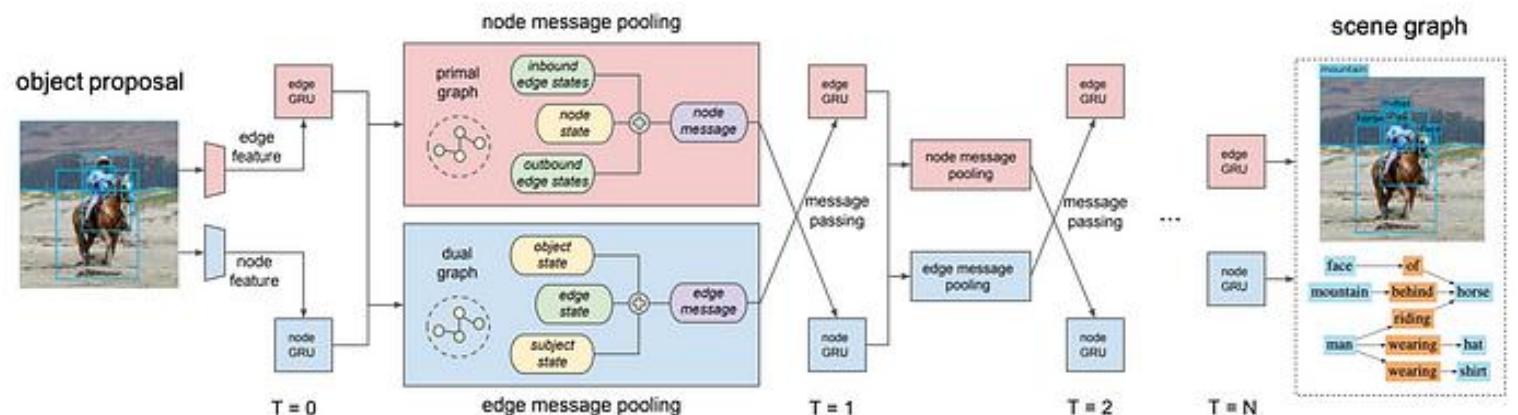
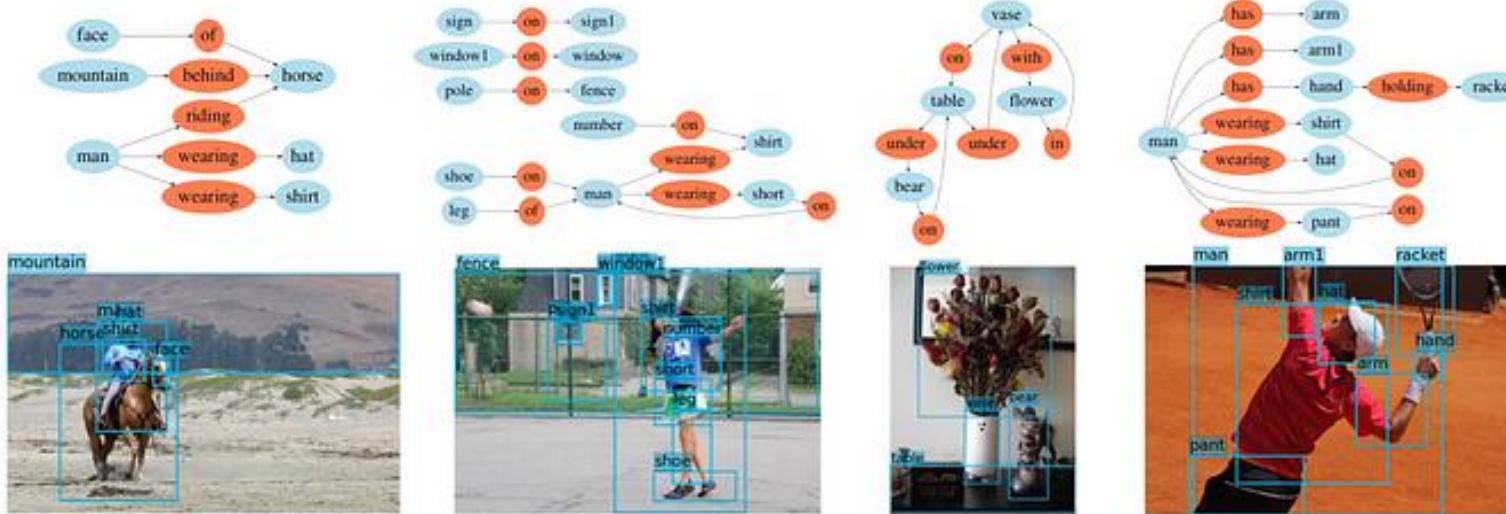


# Traffic forecasting

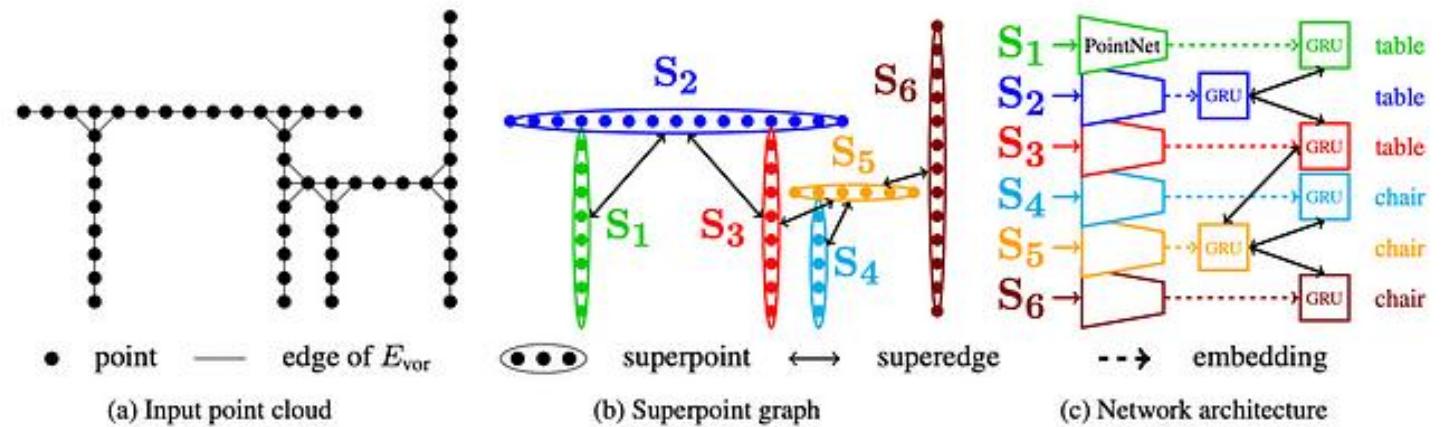
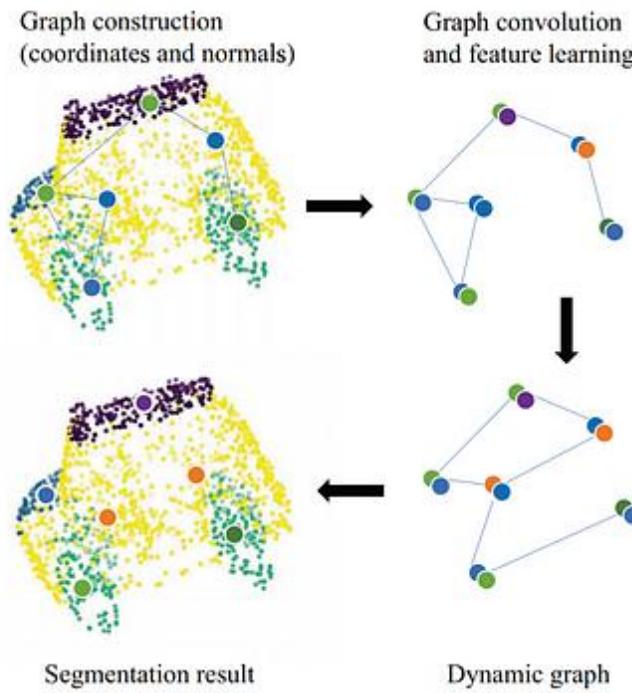


The model architecture for determining optimal routes and their travel time.

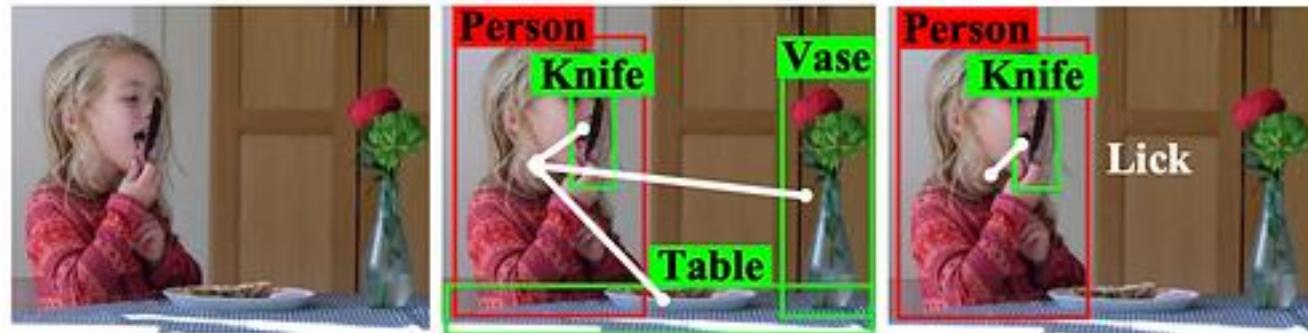
# Scene graph generation of visual data



# Point cloud classification



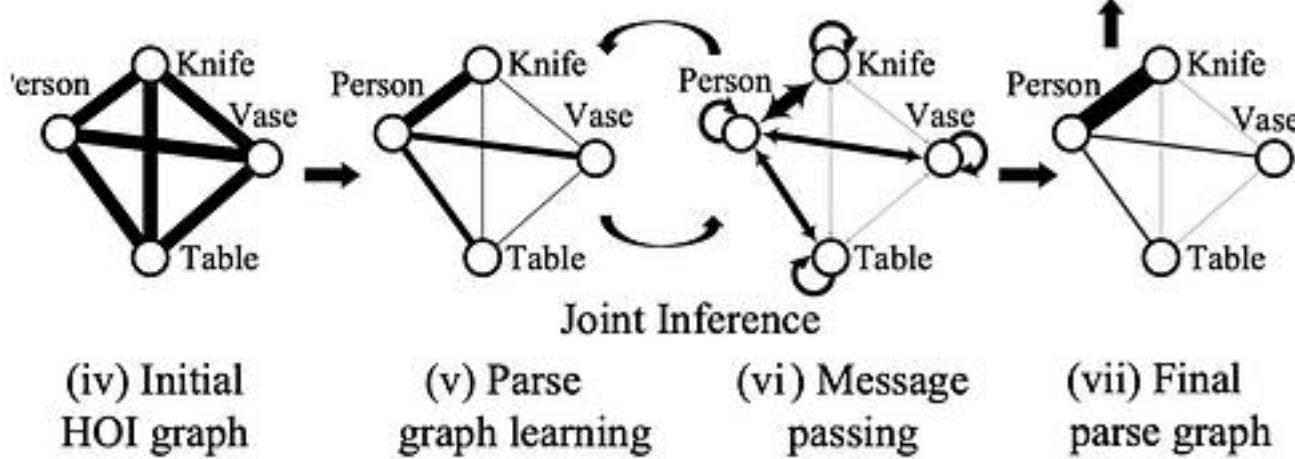
# Object interactions



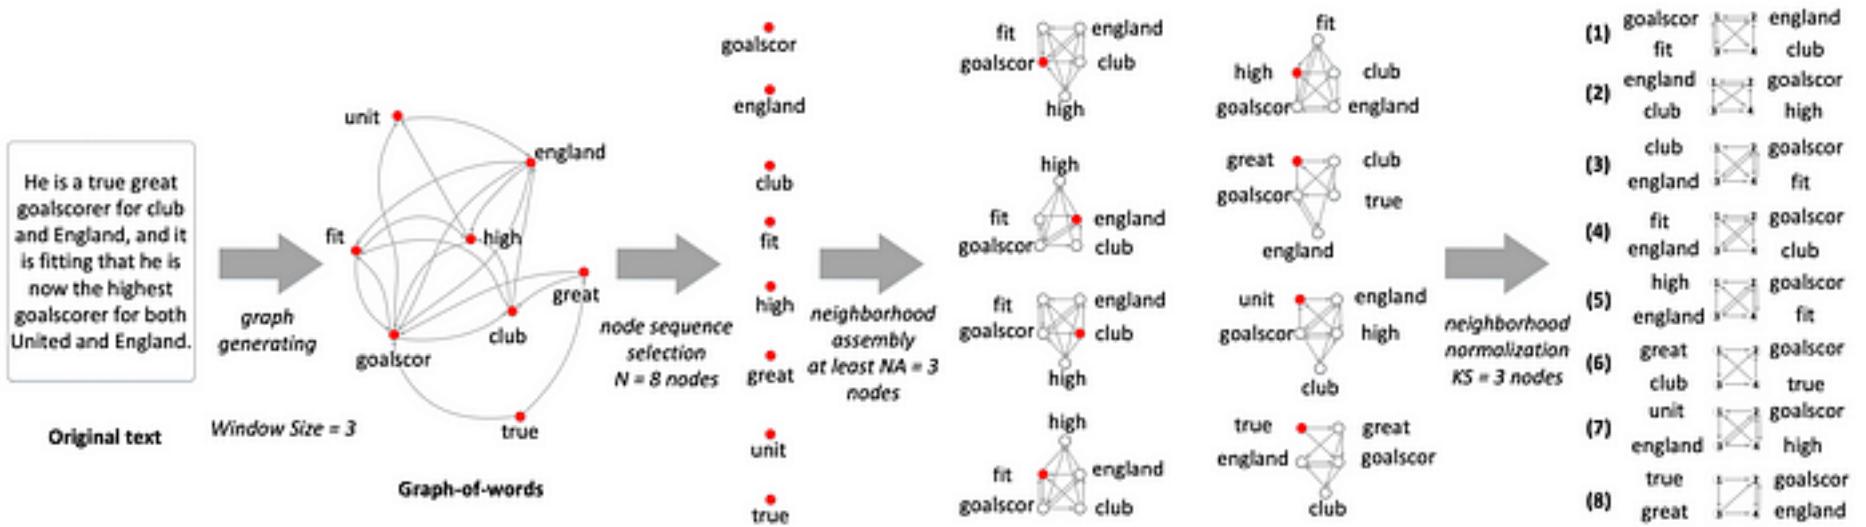
(i) Image

(ii) HOI candidates

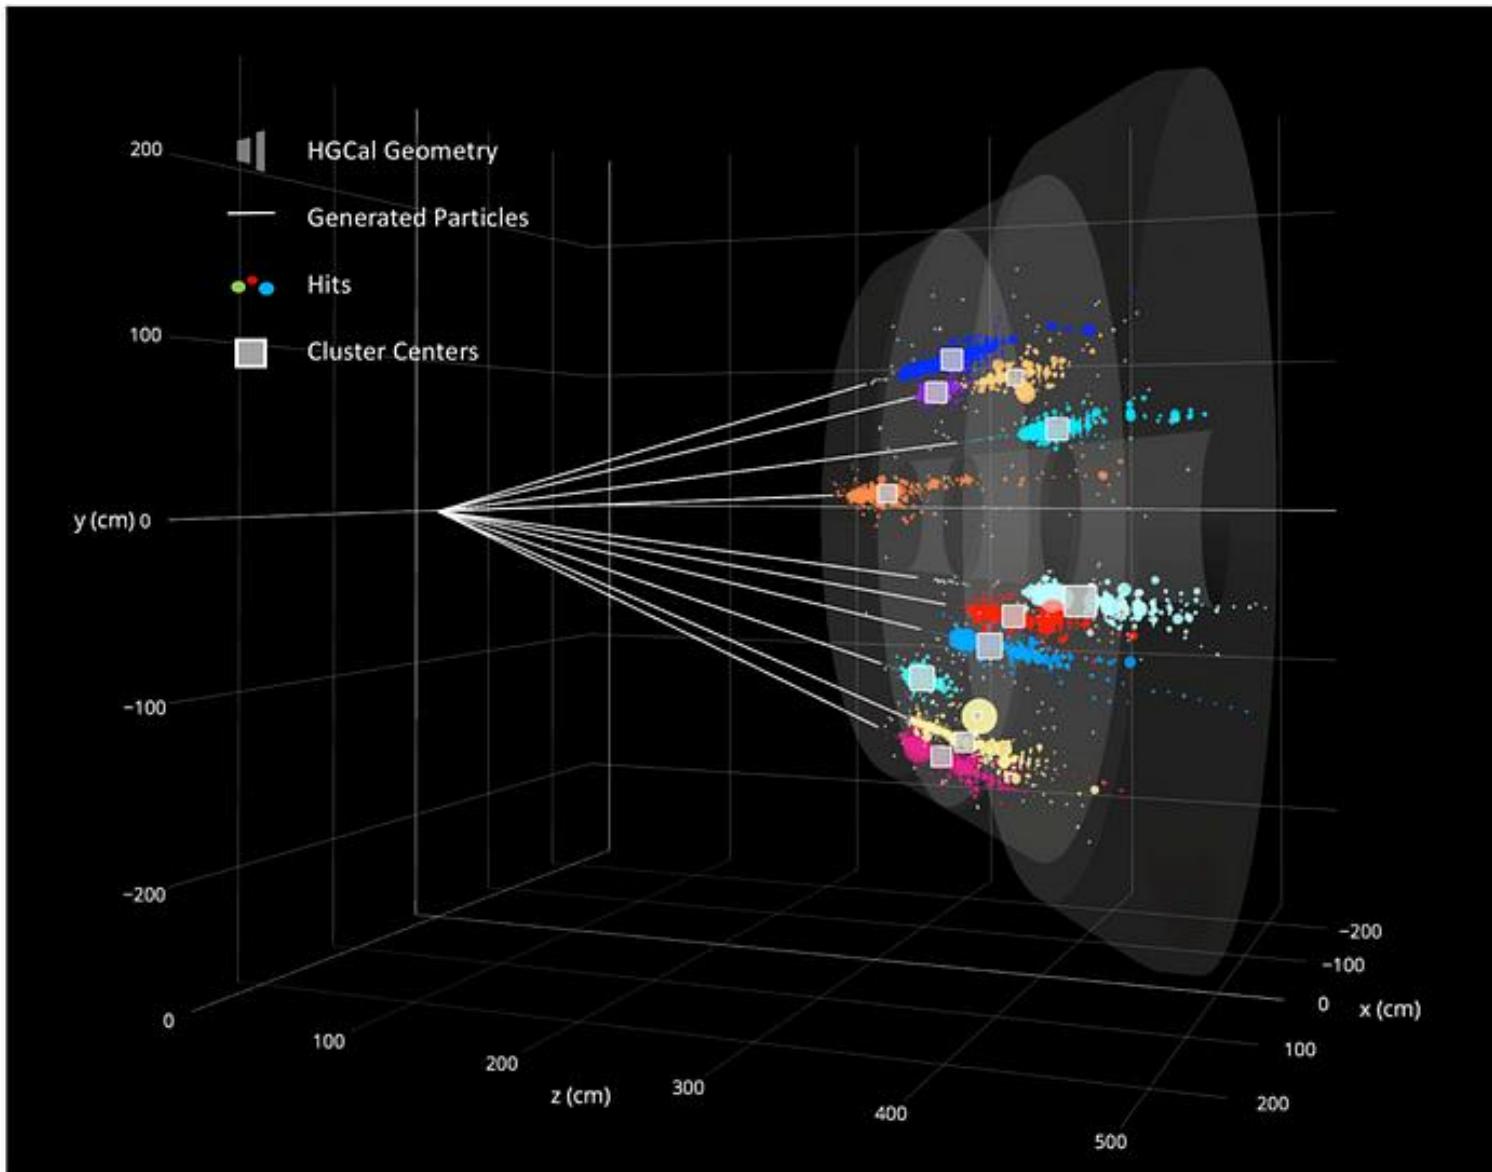
(iii) HOI result



# Text classification



# Particle physics



# Next lecture

Lecture	Title	Lecture	Title
1	Intro and history of deep learning	2	AutoDiff
3	Deep learning optimization I	4	Deep learning optimization II
5	Convolutional deep learning	6	Attention-based deep learning
7	Graph deep learning	8	From supervised to unsupervised deep learning
9	Multi-modal deep learning	10	Generative deep learning
11	What doesn't work in deep learning	12	Non-Euclidean deep learning
13	Q&A	14	Deep learning for videos

# Learning and reflection

Understanding Deep Learning: Chapter 13

Thank you!