

NLP1

Neural sequence modelling

Wilker Aziz

waziz@uva.nl

probabll.github.io

ILLC

Fall 2025

Table of contents

1. Sequence modelling
2. Parameterisation
3. Parameter estimation
4. Predictions
5. Design Choices
6. Puzzle

Sequence modelling

Neural models of sequence prediction

Many NLP tasks involve conditioning on text and predicting sequences

- part-of-speech tagging [Ling et al., 2015]
- named-entity recognition [Lample et al., 2016]
- machine translation [Sutskever et al., 2014]
- text summarisation [Rush et al., 2015]
- entity retrieval [Cao et al., 2021]
- information extraction [Josifoski et al., 2022]

when the model has
a lot of
independence
assumption it
becomes really task
specific

Neural models of sequence prediction

Many NLP tasks involve conditioning on text and predicting sequences

- part-of-speech tagging [Ling et al., 2015]
- named-entity recognition [Lample et al., 2016]
- machine translation [Sutskever et al., 2014]
- text summarisation [Rush et al., 2015]
- entity retrieval [Cao et al., 2021]
- information extraction [Josifoski et al., 2022]

Deploying a system for any of these tasks requires a lot of expert knowledge (about task, datasets, design decisions, etc.), but most solutions employ a similar backbone: a neural model of sequence prediction.

Sequence-to-sequence

we focus on the task of sequence to sequence (ex translation)

We are interested in modelling a specific relationship between pairs of sequences:

- an input sequence x from an input space \mathcal{X}
- an output sequence y from an output space \mathcal{Y}

We will assume this relationship can be modelled directionally
 $(x \rightarrow y)$ in a non-deterministic way.¹

¹ **Notation** capital letters for random variables (e.g., Y), lowercase letters for their assignments (e.g., y), calligraphic letters for sample spaces (e.g., \mathcal{Y}). We use Y_j to denote a step in a random sequence and $Y_{<j}$ to denote a prefix sequence (up until but not including the Y_j). P_Y is the distribution of Y , $P_{Y|X=x}$ is the distribution of Y given $X = x$. $P(Y = y|X = x)$ is the probability of observing $Y = y$ given $X = x$.

Probabilistic modelling

We will treat y as an observation for a random variable (rv) Y , which we draw conditionally given an observation x for the rv X .

The probability $P(Y = y|X = x)$ with which we observe $Y = y$ conditioned on $X = x$ is given by a parametric function with parameters θ :

$$P(Y = y|X = x) = f(y|x; \theta) \quad (1)$$

Our first job, as modellers, is to design this probability mass function (pmf). Once it is in place, we will discuss how to estimate parameters for it, and, finally, how to use it to make predictions.

it is completely different from NB and HMM, we do not model the joint, I assume to have well formed X and want to predict Y

Challenges

Designing a pmf involves

1. specifying the parametric family
2. picking a value for its parameters

Let's concentrate on (1), assuming that we will be employing a form of gradient-based optimisation for (2).

we need the function to be differentiable in order to use gradient-based optimisation

Parameterisation

Conditional probability distributions (cpds) for structures

Given any $x \in \mathcal{X}$, we want to be able to parameterise a distribution over outcomes of Y . There are 2 key challenges here:

- the input space \mathcal{X} is very large (typically infinite)
- the output space \mathcal{Y} is very large (either infinite or it grows combinatorially with the size of input x)

When we were modelling classifiers the output space was very easy

So to give some intuitions let's assume for a moment that the output space is simple (follow on the next slide)

Structured \mathcal{X} , unstructured \mathcal{Y}

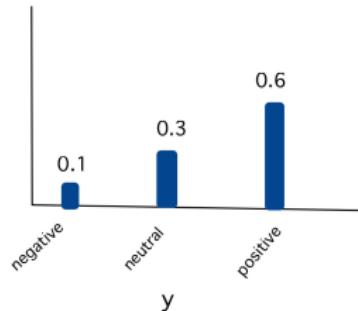
Pretend for a moment that $\mathcal{Y} = \{1, \dots, K\}$. To prescribe a cpd for $Y|X = x$, we need K probabilities for any given $x \in \mathcal{X}$.

Structured \mathcal{X} , unstructured \mathcal{Y}

Pretend for a moment that $\mathcal{Y} = \{1, \dots, K\}$. To prescribe a cpd for $Y|X = x$, we need K probabilities for any given $x \in \mathcal{X}$.

For a single $x \in \mathcal{X}$, this is not so difficult (we could store K probability values):

Categorical($y|0.1, 0.3, 0.6$)

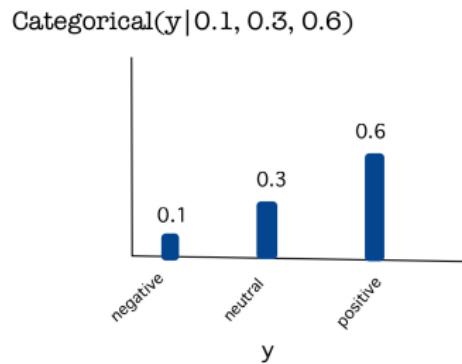


x = “Great breakfast and service. A bit far from the centre, but you get a quiet area.”

Structured \mathcal{X} , unstructured \mathcal{Y}

Pretend for a moment that $\mathcal{Y} = \{1, \dots, K\}$. To prescribe a cpd for $\mathcal{Y}|\mathcal{X} = x$, we need K probabilities for any given $x \in \mathcal{X}$.

For a single $x \in \mathcal{X}$, this is not so difficult (we could store K probability values):



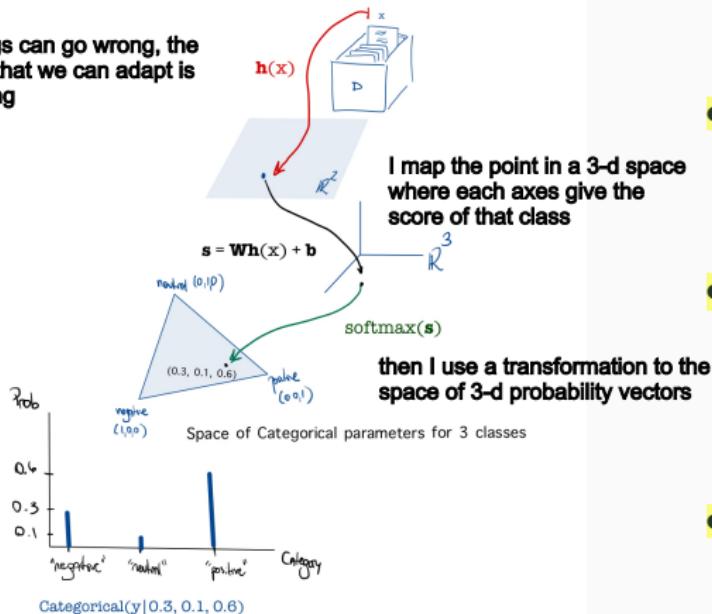
$x = \text{"Great breakfast and service. A bit far from the centre, but you get a quiet area."}$

But doing so for each and every possible $x \in \mathcal{X}$, including those we've never seen, requires a bit more ingenuity.

we can't just store all the probabilities for each x , we need to parametrize it

Log-linear (or logistic) cpds

many things can go wrong, the only thing that we can adapt is the mapping



- map x to a fixed number of features $\mathbf{h}(x) \in \mathbb{R}^D$
- map $\mathbf{h}(x)$ to K scores (a.k.a. logits), for example, linearly: $\mathbf{W}\mathbf{h}(x) + \mathbf{b}$
- constrain the outputs to the probability simplex

This will map *any* x that we can 'featurise' to a Categorical pmf. Crucially, no matter how large \mathcal{X} is, it only takes $K \times D + K$ parameters.

Encoding functions

The ability to ‘encode’ an arbitrary x into a D -dimensional space is essential for our parameterisation.

Pre-2010 these functions were handmade feature functions.

Nowadays they are part of the parameterisation. That is, we use NNs to represent the input and map it to output probability values.

they started to be less linguistic specific and more agnostic. By tuning the parameters of the network the features become more or less interesting to the problem at hand.

Rather than storing features is better to parametrize a little mechanism which scales as a function fo the sequence length. The parameters of NN is fixed, is a design destino from the beginning.

A neural text classifier

K -way text classifier

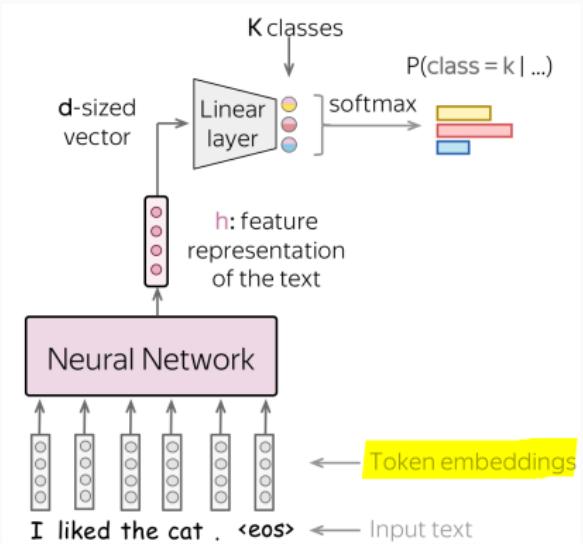


Image from Lena Voita's NLP course For You

$$Y|X = x \sim \text{Cat}(\mathbf{g}(x; \theta))$$

Encoder-decoder

suppose $I = |x|$

$$\mathbf{e}_i = \text{embed}_D(x_i; \theta_{\text{inp}}) \quad i = 1, \dots, I$$

$$\mathbf{r}_{1:I} = \text{LSTM}_H(\mathbf{e}_{1:I}; \theta_{\text{enc}})$$

$$\mathbf{h} = \text{avgpool}(\mathbf{r}_{1:I})$$

$$\mathbf{s} = \text{linear}_K(\mathbf{h}; \theta_{\text{out}})$$

$$\mathbf{g}(x; \theta) = \text{softmax}(\mathbf{s})$$

this NN cannot be a MLP because it does not scale with input length. But we can combine MLPs to combine indefinite length sequence by computing a state which encode the progress and the current input (so a RNN/LSTM)

Classification rule

$$\arg \max_{k \in [K]} g_k(x; \theta)$$

The parameters θ include the embedding matrix, the LSTM parameters, as well as the final linear transformation.

the dimension of the embedding grows with the size of the vocabulary

Conditional probability distributions (cpds) for structures

Given any $x \in X$, we want to be able to parameterise a distribution over outcomes of Y . There are 2 key challenges here:

- the input space \mathcal{X} is very large (typically infinite)
- the output space \mathcal{Y} is very large (either infinite or it grows with size of input x)

Structured \mathcal{Y}

We exploit a decomposition into parts, where each part is drawn from a ‘small’ sample space.

Structured \mathcal{Y}

We exploit a decomposition into parts, where each part is drawn from a ‘small’ sample space.

For example, a **POS tag sequence** can be decomposed into a sequence of *word categories*:

$$x = \langle \text{I, am, going, home} \rangle$$

$$y = \langle \text{PRP, VBP, VBG, NN} \rangle$$

x

Structured \mathcal{Y}

We exploit a decomposition into parts, where each part is drawn from a ‘small’ sample space.

For example, a **POS tag sequence** can be decomposed into a sequence of *word categories*:

$$x = \langle \text{I, am, going, home} \rangle$$

$$y = \langle \text{PRP, VBP, VBG, NN} \rangle$$



Structured \mathcal{Y}

We exploit a decomposition into parts, where each part is drawn from a ‘small’ sample space.

For example, a **POS tag sequence** can be decomposed into a sequence of *word categories*:

$$x = \langle \text{I, am, going, home} \rangle$$

$$y = \langle \text{PRP, VBP, VBG, NN} \rangle$$



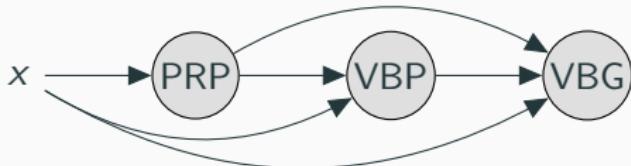
Structured \mathcal{Y}

We exploit a decomposition into parts, where each part is drawn from a ‘small’ sample space.

For example, a **POS tag sequence** can be decomposed into a sequence of *word categories*:

$$x = \langle \text{I, am, going, home} \rangle$$

$$y = \langle \text{PRP, VBP, VBG, NN} \rangle$$



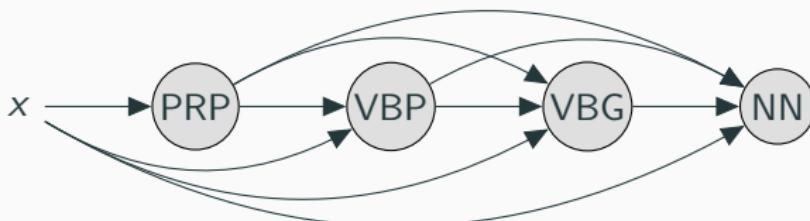
Structured \mathcal{Y}

We exploit a decomposition into parts, where each part is drawn from a ‘small’ sample space.

For example, a **POS tag sequence** can be decomposed into a sequence of *word categories*:

$$x = \langle \text{I, am, going, home} \rangle$$

$$y = \langle \text{PRP, VBP, VBG, NN} \rangle$$



Another example

A **translation** can be decomposed into a sequence of *target words*:

$$x = \langle \text{How, are, you, doing, ?} \rangle$$

$$y = \langle \text{Hoe, gaat, het, ?} \rangle$$

x

Another example

A **translation** can be decomposed into a sequence of *target words*:

$$x = \langle \text{How, are, you, doing, ?} \rangle$$

$$y = \langle \text{Hoe, gaat, het, ?} \rangle$$

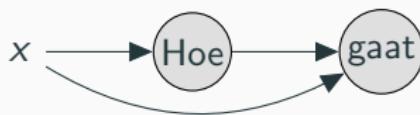


Another example

A **translation** can be decomposed into a sequence of *target words*:

$$x = \langle \text{How, are, you, doing, ?} \rangle$$

$$y = \langle \text{Hoe, gaat, het, ?} \rangle$$

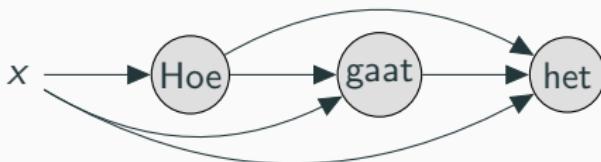


Another example

A **translation** can be decomposed into a sequence of *target words*:

$$x = \langle \text{How, are, you, doing, ?} \rangle$$

$$y = \langle \text{Hoe, gaat, het, ?} \rangle$$

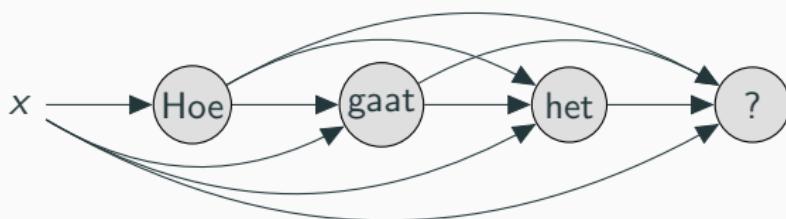


Another example

A **translation** can be decomposed into a sequence of *target words*:

$$x = \langle \text{How, are, you, doing, ?} \rangle$$

$$y = \langle \text{Hoe, gaat, het, ?} \rangle$$

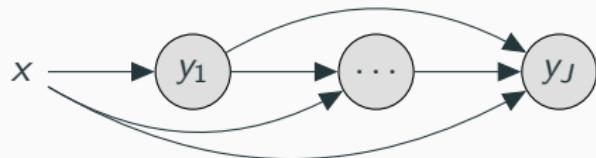


In general

For an input-output pair:

$$x = \langle x_1, \dots, x_I \rangle$$

$$y = \langle y_1, \dots, y_J \rangle$$



$$P(Y = y|X = x) = \prod_{j=1}^J P(Y_j = y_j | \underbrace{X = x, Y_{<j} = y_{<j}}_{\text{parents of } j\text{th rv}}) \quad (2)$$

In the decomposition, conditioned on increasingly complex context, each part is drawn from a small sample space. Models factorised this way are often referred to as *autoregressive*.

autoregressive conditional language model (autoregressive means that we are making assumptions about the past)

One C -way classifier, J steps

A general model of sequence prediction is in fact obtained by repeated application of something very similar to a shared text classifier:

$$Y_j | \underbrace{X = x, Y_{<j} = y_{<j}}_{\text{conditioning context}} \sim \text{Cat}(\underbrace{\mathbf{g}(x, y_{<j}; \theta)}_{\text{outcome probs}}) \quad (3)$$

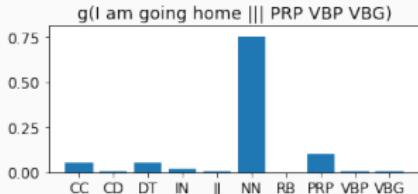
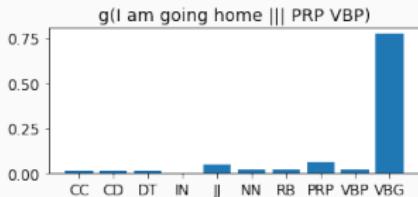
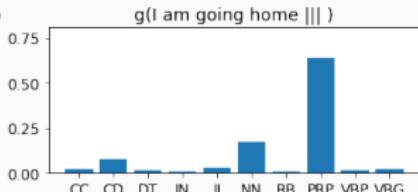
Here, \mathbf{g} maps from an input x and a partial output $y_{<j}$ to the probabilities of the possible outcomes for the j th step. Typically, the sample space across steps (e.g., all tags, all words).

can I map a sequence and a partial sequence to a vector of scores of the vocabulary size?

POS tagging example

here I'm inferring the first tag, the ||| is the separator

Same NN $g(\cdot; \theta)$ reused over and over (as many times as there are steps in the input sequence), each time mapping from a growing context (x and $y_{<j}$) to a probability distribution over the same categorical space (i.e., space of tags).



A neural tagger

Statistical model let the function \mathbf{g} map from an input x and prefix $y_{<j}$ to a distribution over K tags:

$$Y_j | X = x, Y_{<j} = y_{<j} \sim \text{Cat}(\mathbf{g}(x, y_{<j}; \theta)) \quad (4)$$

A neural tagger

Statistical model let the function \mathbf{g} map from an input x and prefix $y_{<j}$ to a distribution over K tags:

$$Y_j | X = x, Y_{<j} = y_{<j} \sim \text{Cat}(\mathbf{g}(x, y_{<j}; \theta)) \quad (4)$$

Encoder-decoder $I = |x|$

I embed each word in a space of arbitrary D dimension $\mathbf{e}_i = \text{embed}_D(x_i; \theta_{\text{words}}) \quad i = 1, \dots, I$

A neural tagger

Statistical model let the function \mathbf{g} map from an input x and prefix $y_{<j}$ to a distribution over K tags:

$$Y_j | X = x, Y_{<j} = y_{<j} \sim \text{Cat}(\mathbf{g}(x, y_{<j}; \theta)) \quad (4)$$

Encoder-decoder $I = |x|$

$$\mathbf{e}_i = \text{embed}_D(x_i; \theta_{\text{words}}) \quad i = 1, \dots, I$$

I concat the embedding in a sequence and feed them to a bidirection LSTM

$$\mathbf{c}_{1:I} = \text{BiLSTM}_H(\mathbf{e}_{1:I}; \theta_{\text{enc}})$$

A neural tagger

Statistical model let the function \mathbf{g} map from an input x and prefix $y_{<j}$ to a distribution over K tags:

$$Y_j | X = x, Y_{<j} = y_{<j} \sim \text{Cat}(\mathbf{g}(x, y_{<j}; \theta)) \quad (4)$$

Encoder-decoder $I = |x|$

$$\mathbf{e}_i = \text{embed}_D(x_i; \theta_{\text{words}}) \quad i = 1, \dots, I$$

$$\mathbf{c}_{1:I} = \text{BiLSTM}_H(\mathbf{e}_{1:I}; \theta_{\text{enc}})$$

$$\mathbf{t}_{j-1} = \text{embed}_D(y_{j-1}; \theta_{\text{tags}})$$

A neural tagger

Statistical model let the function \mathbf{g} map from an input x and prefix $y_{<j}$ to a distribution over K tags:

$$Y_j | X = x, Y_{<j} = y_{<j} \sim \text{Cat}(\mathbf{g}(x, y_{<j}; \theta)) \quad (4)$$

Encoder-decoder $I = |x|$

$$\mathbf{e}_i = \text{embed}_D(x_i; \theta_{\text{words}}) \quad i = 1, \dots, I$$

$$\mathbf{c}_{1:I} = \text{BiLSTM}_H(\mathbf{e}_{1:I}; \theta_{\text{enc}})$$

$$\mathbf{t}_{j-1} = \text{embed}_D(y_{j-1}; \theta_{\text{tags}})$$

$$\mathbf{h}_j = \text{rnnstep}_H(\mathbf{h}_{j-1}, [\mathbf{c}_j, \mathbf{t}_{j-1}]; \theta_{\text{dec}})$$

A neural tagger

Statistical model let the function \mathbf{g} map from an input x and prefix $y_{<j}$ to a distribution over K tags:

$$Y_j | X = x, Y_{<j} = y_{<j} \sim \text{Cat}(\mathbf{g}(x, y_{<j}; \theta)) \quad (4)$$

Encoder-decoder $I = |x|$

$$\mathbf{e}_i = \text{embed}_D(x_i; \theta_{\text{words}}) \quad i = 1, \dots, I$$

$$\mathbf{c}_{1:I} = \text{BiLSTM}_H(\mathbf{e}_{1:I}; \theta_{\text{enc}})$$

$$\mathbf{t}_{j-1} = \text{embed}_D(y_{j-1}; \theta_{\text{tags}})$$

$$\mathbf{h}_j = \text{rnnstep}_H(\mathbf{h}_{j-1}, [\mathbf{c}_j, \mathbf{t}_{j-1}]; \theta_{\text{dec}})$$

$$\mathbf{s}_j = \text{linear}_K(\mathbf{h}_j; \theta_{\text{out}})$$

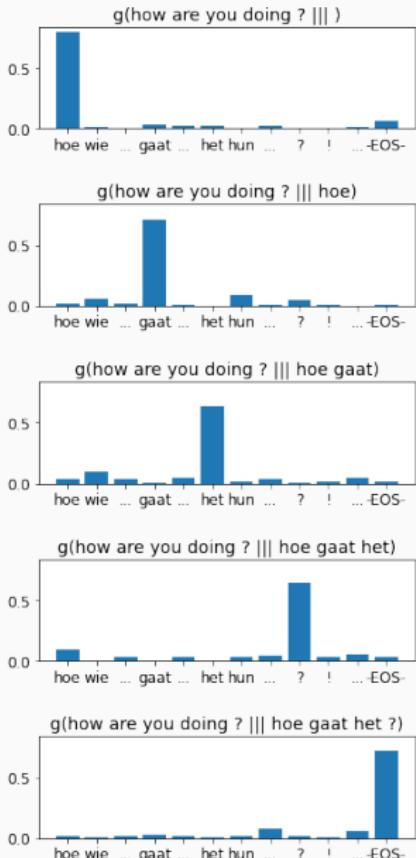
$$\mathbf{g}(x, y_{<j}; \theta) = \text{softmax}(\mathbf{s}_j)$$

the parameters θ include the embedding matrices (words and tags), the parameters of the BiLSTM encoder and the LSTM decoder, as well as the final linear transformation.

Translation example

Same NN $g(\cdot; \theta)$ reused over and over, each time mapping from a growing context (x and $y_{<j}$) to a probability distribution over the same categorical space (i.e., space of words).

In translation the output length is not determined by the length of the input, instead we repeat this process until a special terminating symbol is observed or generated.



A neural translation model

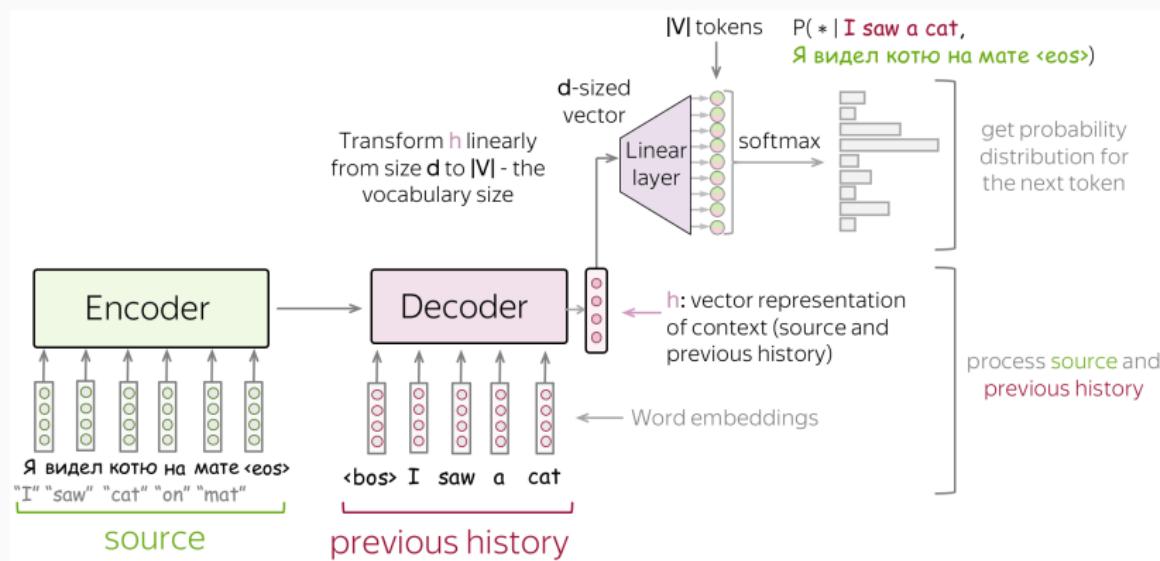
Statistical model let the function \mathbf{g} map from an input x and prefix $y_{<j}$ to a distribution over V words:

$$Y_j | X = x, Y_{<j} = y_{<j} \sim \text{Cat}(\mathbf{g}(x, y_{<j}; \theta)) \quad (5)$$

A neural translation model

Statistical model let the function \mathbf{g} map from an input x and prefix $y_{<j}$ to a distribution over V words:

$$Y_j | X = x, Y_{<j} = y_{<j} \sim \text{Cat}(\mathbf{g}(x, y_{<j}; \theta)) \quad (5)$$



Text Summarisation

NLP task: generate a short version of a (collection of) document(s) that contains the most important information.

- Extractive: identify and extract important/relevant sentences from the input document(s).
- Abstractive: interpret the content of the document and generate an ‘original’ summary.

Natural language processing

Natural language processing is an interdisciplinary subfield of computer science and linguistics. It is primarily concerned with giving computers the ability to support and manipulate human language. [Wikipedia](#)

People also search for

[View 10+ more](#)



Artificial
intelligence



Machine
learning



Data
science



Deep
learning

Natural language processing

[Article](#) [Talk](#)

From Wikipedia, the free encyclopedia

For other uses, see [NLP](#).

This article is about natural language processing done by computers. For the natural language processing done by the human brain, see [Language processing in the brain](#).

Natural language processing (NLP) is an [interdisciplinary](#) subfield of [computer science](#) and [linguistics](#). It is primarily concerned with giving computers the ability to support and manipulate human language. It involves processing [natural language](#) datasets, such as [text corpora](#) or speech corpora, using either rule-based or probabilistic (i.e., statistical and, most recently, neural network-based) [machine learning](#) approaches. The goal is a computer capable of “understanding” the contents of documents, including the [contextual](#) nuances of the language within them. The technology can then accurately extract information and insights contained in the documents as well as categorize and organize the documents themselves.

Challenges in natural language processing frequently involve [speech recognition](#), [natural-language understanding](#), and [natural-language generation](#).

History [\[edit\]](#)

Further information: [History of natural language processing](#)

Natural language processing has its roots in the 1950s. Already in 1950, Alan Turing published an article titled “[Computing Machinery and Intelligence](#)” which proposed what is now called the [Turing test](#) as a criterion of intelligence, though at the time that was not articulated as a problem separate from artificial intelligence. The proposed test includes a task that involves the automated interpretation and generation of natural language.



Illustration of the field of NLP and AI, showing a brain, a computer monitor displaying code, and various icons representing different sub-fields like speech recognition, text processing, and neural networks.

A neural abstractive summarisation model

Statistical model: just like in translation, we condition on the input documents(s) x and the prefix of already generated tokens in the summary $y_{<j}$ and parameterise a probability distribution over the space of possible tokens for the j th position.

the attention predict a score for each position based on the encoded tokens. It compute what it is more salient to the decoder at this point

memory networks, the encoder encodes the memory while the decoder manage the query, their interpolation produce the result

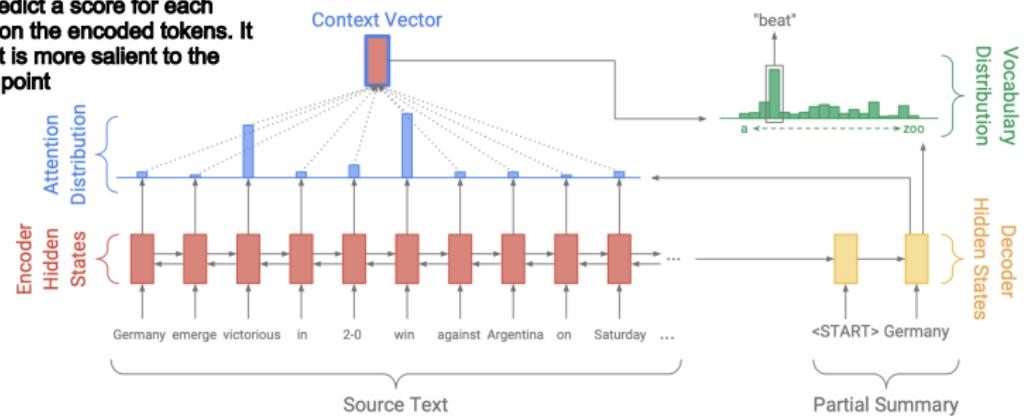


Figure 2: Baseline sequence-to-sequence model with attention. The model may attend to relevant words in the source text to generate novel words, e.g., to produce the novel word *beat* in the abstractive summary *Germany beat Argentina 2-0* the model may attend to the words *victorious* and *win* in the source text.

An improved neural abstractive summarisation model

Through the architecture, we try to bias the model towards ‘preferred’ solutions. For example, we can give this model a push towards sometimes performing extractive summarisation.

Figure from [See et al., 2017]

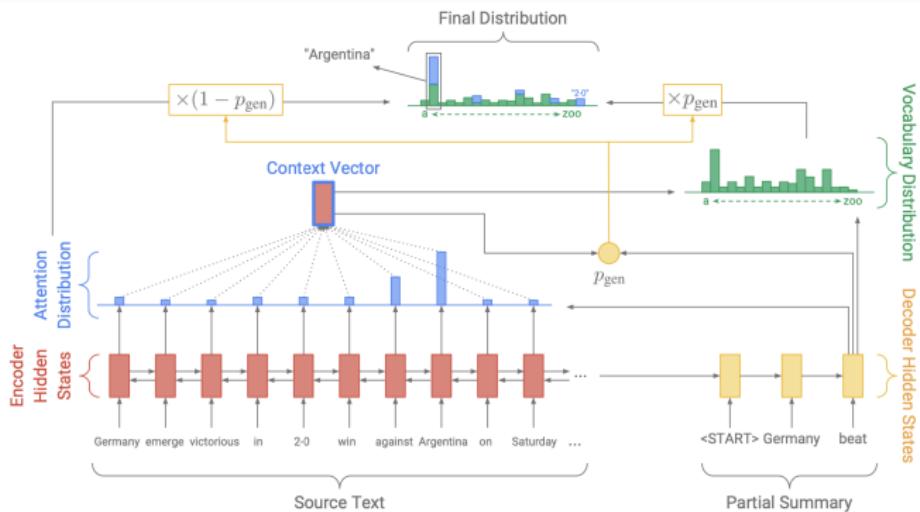


Figure 3: Pointer-generator model. For each decoder timestep a generation probability $p_{\text{gen}} \in [0, 1]$ is calculated, which weights the probability of *generating* words from the vocabulary, versus *copying* words from the source text. The vocabulary distribution and the attention distribution are weighted and summed to obtain the final distribution, from which we make our prediction. Note that out-of-vocabulary article words such as 2-0 are included in the final distribution. Best viewed in color.

Parameter estimation

the transformer is less expressive than a lstm, non recursive model are not infinitely expressive.

exam question, given an architecture satte his prons and cons, or talk about independence assumptions

Data and Task

Data a collection of pairs (x, y) where both x and y can be treated as a sequence of outcomes from small discrete sets.

Statistical task observe x and predict a conditional distribution over all possible sequences.

NLP task map a sequence x to a sequence y : for example via $\arg \max_{y \in \mathcal{Y}} P(Y = y | X = x)$.

Formalisation of statistical task

Statistical model let the function \mathbf{g} map from x to a chain rule factorisation of the conditional distribution $Y|X = x$:

$$Y_j|X = x, Y_{<j} = y_{<j} \sim \text{Cat}(\mathbf{g}(x, y_{<j}; \theta)) \quad (6)$$

θ collectively refers to all trainable parameters in the model.

Formalisation of statistical task

Statistical model let the function \mathbf{g} map from x to a chain rule factorisation of the conditional distribution $Y|X = x$:

$$Y_j|X = x, Y_{<j} = y_{<j} \sim \text{Cat}(\mathbf{g}(x, y_{<j}; \theta)) \quad (6)$$

θ collectively refers to all trainable parameters in the model.

Statistical objective maximum likelihood of model given a dataset of observations \mathcal{D} :

$$\mathcal{L}(\theta|\mathcal{D}) = \sum_{(x,y) \in \mathcal{D}} \log \underbrace{P(Y = y|X = x)}_{f(y|x; \theta)}$$

Formalisation of statistical task

Statistical model let the function \mathbf{g} map from x to a chain rule factorisation of the conditional distribution $Y|X = x$:

$$Y_j|X = x, Y_{<j} = y_{<j} \sim \text{Cat}(\mathbf{g}(x, y_{<j}; \theta)) \quad (6)$$

θ collectively refers to all trainable parameters in the model.

Statistical objective maximum likelihood of model given a dataset of observations \mathcal{D} :

$$\begin{aligned} \mathcal{L}(\theta|\mathcal{D}) &= \sum_{(x,y) \in \mathcal{D}} \log \underbrace{P(Y = y|X = x)}_{f(y|x;\theta)} \\ &= \sum_{(x,y) \in \mathcal{D}} \underbrace{\sum_{j=1}^{|y|} \log g_{y_j}(x, y_{<j}; \theta)}_{\log f(y|x;\theta)} \end{aligned} \quad (7)$$

Formalisation of statistical task

Statistical model let the function \mathbf{g} map from x to a chain rule factorisation of the conditional distribution $Y|X = x$:

$$Y_j|X = x, Y_{<j} = y_{<j} \sim \text{Cat}(\mathbf{g}(x, y_{<j}; \theta)) \quad (6)$$

θ collectively refers to all trainable parameters in the model.

Statistical objective maximum likelihood of model given a dataset of observations \mathcal{D} :

$$\begin{aligned} \mathcal{L}(\theta|\mathcal{D}) &= \sum_{(x,y) \in \mathcal{D}} \log \underbrace{P(Y = y|X = x)}_{f(y|x;\theta)} \\ &= \sum_{(x,y) \in \mathcal{D}} \underbrace{\sum_{j=1}^{|y|} \log g_{y_j}(x, y_{<j}; \theta)}_{\log f(y|x;\theta)} \end{aligned} \quad (7)$$

Algorithm For concave \mathcal{L} (or convex negative log-likelihood), find θ such that $\nabla_\theta \mathcal{L}(\theta|\mathcal{D}) = \mathbf{0}$.

Parameter estimation

Algorithm Solve the equation $\nabla_{\theta}\mathcal{L}(\theta|\mathcal{D}) = \mathbf{0}$ for θ . There is no closed form solution. But an optimum can be found via a fixed-point iteration: $\theta \leftarrow \theta + \gamma \nabla_{\theta}\mathcal{L}(\theta|\mathcal{D})$ for $\gamma > 0$.

What if my dataset is massive?

The time and memory necessary to compute $\nabla_{\theta}\mathcal{L}(\theta|\mathcal{D})$ grows linearly with the size of the data.

What if my dataset is massive?

The time and memory necessary to compute $\nabla_{\theta}\mathcal{L}(\theta|\mathcal{D})$ grows linearly with the size of the data.

Luckily, stochastic optimisation will converge in finite time even with gradient estimates, as long as they are unbiased, and as long as we use careful learning rate schedules [Robbins and Monro, 1951, Bottou and Cun, 2004].

What if my dataset is massive?

The time and memory necessary to compute $\nabla_{\theta}\mathcal{L}(\theta|\mathcal{D})$ grows linearly with the size of the data.

Luckily, stochastic optimisation will converge in finite time even with gradient estimates, as long as they are unbiased, and as long as we use careful learning rate schedules [Robbins and Monro, 1951, Bottou and Cun, 2004].

If \mathcal{B} is a random subset ('mini batch') of \mathcal{D} , it holds that:

$$\nabla_{\theta}\mathcal{L}(\theta|\mathcal{D}) = \mathbb{E}_{\mathcal{B} \sim \mathcal{D}}[\nabla_{\theta}\mathcal{L}(\theta|\mathcal{B})] \quad (8)$$

Thus we can take iterative steps, each based on a small random subset of the data.

A soup of names

You probably heard of the *cross entropy loss*, which is identical to the *negative* of the quantity in the previous slide.

Some people will also call it the categorical cross entropy loss, or the softmax loss.

'Softmax loss' is a bit odd, softmax is a vector-valued function, it's hard to imagine it as a loss.

Categorical cross entropy is clear, cross entropy can be clear enough in context.

Predictions

Making decisions

Our final job, as modellers, is to find a reasonable way to form predictions.

That is, given an input x , our model outputs a representation of an entire probability distribution $P_{Y|X=x}$ (i.e., over all of \mathcal{Y}).

We are now confronted with the task to map from $P_{Y|X=x}$ to a single output y . This is often formulated as a search, or discrete optimisation, problem.

Most probable output

A common algorithm for making decisions is to search for the candidate output c which is assigned highest probability:

$$y^* = \arg \max_{c \in \mathcal{Y}} f(c|x; \theta) \quad (9)$$

This can also be done in log space: $\arg \max_{c \in \mathcal{Y}} \log f(c|x; \theta)$.

Most probable output

A common algorithm for making decisions is to search for the candidate output c which is assigned highest probability:

the space c is unbounded, be infinite is not a
good argument because some models
factorise so simply that even with infinite space
they can compute it easily, so the right answer
is unbounding!

$$y^* = \arg \max_{c \in \mathcal{Y}} f(c|x; \theta) \quad (9)$$

This can also be done in log space: $\arg \max_{c \in \mathcal{Y}} \log f(c|x; \theta)$.

This is intractable for most models (incl. autoregressive models)!

Common approximations include *greedy decoding* (iteratively argmax each step) or *beam search decoding* (maintain a fixed number of high-scoring candidates).

greedy decoding select the max at each step without looking back, beam search does looks back but in a certain fixed range

Let $u(y, c; x)$ quantify the utility of c when y is known to be a valid output for x .
 the utility compare a candidate translation to a reference translation via a similarity function

In decision theory, a rational decision maker acts by maximising expected utility under the model:

$$y^* = \arg \max_{c \in \mathcal{Y}} \mathbb{E}[u(Y, c; x)] \quad (10)$$

I rank the candidates over utility

this is as intractable as before but can be approximated if we can draw samples from the model using the below method

Expected utility can be approximated via Monte Carlo (MC):

$$\mathbb{E}[u(Y, c; x)] \stackrel{\text{MC}}{\approx} \frac{1}{K} \sum_{k=1}^K u(y^{(k)}, c; x) \quad (11)$$

with $y^{(k)} \sim P_{Y|X=x}$.

See [Eikema and Aziz, 2020, 2022].

We can sort candidates on expected utility, which is given by some similarity function.

Enumerating candidates

- Greedy decoding
- Beam search
- Ancestral sampling
- Top-p and top-k sampling [Holtzman et al., 2019]
- Sampling without replacement [Kool et al., 2019]

Evaluation

Statistical: does the model fit the data well?

- perplexity
- statistics of model samples

see [Giulianelli et al., 2023]

Task-driven: does the model support good decisions (in a benchmark)?

- exact-match/precision/recall/F1 for short generations (QA, entity linking, information extraction)
- string similarity (e.g., BLEU, ROUGE, METEOR, BEER)
- semantic similarity (e.g., COMET, BLEURT)

Design Choices

Architectures

Many choices of encoders and decoders:

- CNNs [Gehring et al., 2017]
- GCNs [Bastings et al., 2017]
- Transformers [Vaswani et al., 2017]

The Transformer, in particular, is today's architecture of choice.

Transformer

Transformers are based on a stack of (parallel) self-attention heads³ and feed-forward networks. For any one Transformer layer, the outputs $\mathbf{h}_1^{(k+1)}, \dots, \mathbf{h}_{\ell}^{(k+1)}$ are such that each output $\mathbf{h}_j^{(k+1)}$ depends on no other output $\mathbf{h}_i^{(k+1)}$ and on *at most* all of the layer's inputs $\mathbf{h}_1^{(k)}, \dots, \mathbf{h}_{\ell}^{(k)}$.

- **Encoder:** $\mathbf{h}_j^{(k+1)}$ depends on all inputs $\mathbf{h}_{1:\ell}^{(k)}$
- **Decoder:** $\mathbf{h}_j^{(k+1)}$ depends on inputs $\mathbf{h}_{<j}^{(k)}$ prior to position j , as to ensure autoregressiveness

³Illustration for multiheaded attention.

Transformer

Transformers are based on a stack of (parallel) self-attention heads³ and feed-forward networks. For any one Transformer layer, the outputs $\mathbf{h}_1^{(k+1)}, \dots, \mathbf{h}_{\ell}^{(k+1)}$ are such that each output $\mathbf{h}_j^{(k+1)}$ depends on no other output $\mathbf{h}_i^{(k+1)}$ and on *at most* all of the layer's inputs $\mathbf{h}_1^{(k)}, \dots, \mathbf{h}_{\ell}^{(k)}$.

- **Encoder:** $\mathbf{h}_j^{(k+1)}$ depends on all inputs $\mathbf{h}_{1:\ell}^{(k)}$
- **Decoder:** $\mathbf{h}_j^{(k+1)}$ depends on inputs $\mathbf{h}_{<j}^{(k)}$ prior to position j , as to ensure autoregressiveness

Contrast this with RNNs

- Encoder (e.g., BiLSTM): $\mathbf{h}_j^{(k+1)}$ depends directly on $\mathbf{h}_{j-1}^{(k+1)}$ and $\mathbf{h}_{j+1}^{(k+1)}$ and hence, recursively, on all other outputs.
- Decoder: $\mathbf{h}_j^{(k+1)}$ depends directly on $\mathbf{h}_{j-1}^{(k+1)}$ and hence, recursively, on all $\mathbf{h}_{<j}^{(k+1)}$, ensuring autoregressiveness.

³Illustration for multiheaded attention.

Transformers vs. RNNs

RNNs are 'stateful' and Transformers are 'stateless'. This has certain implications:

At training time all inputs (past and future) are already known.

- For the Transformer, this is an opportunity for parallelism: compute all outputs in parallel.
- The recursive nature of RNNs impose sequential processing.

At test time, regardless of which architecture we use, no inputs other than past (already generated) inputs are available for observation, hence sequential computation is unavoidable.

There are also implications for learning dynamics (related to chain rule of derivatives), but that's more of a DL1 topic.

Alternative Factorisation

Many alternatives to chain rule:

- Latent variables (LVMs) [Zhang et al., 2016, Eikema and Aziz, 2019]

We use **marginalisation** to overcome factorisation assumptions:

$$P_{Y|X}(y|x) = \int \mathcal{N}(z|0, I) \underbrace{f(y|x, z; \theta)}_{\text{autoregressive}} dz$$

- Non-autoregressive models

when something is
not autoregressive
it enables a lot of
parallelism.

- CRFs (for sequence labelling tasks) [Ma and Hovy, 2016]
We factorise with **strong (bigram-like) assumptions**, but in an undirected manner: $P_{Y|X}(y|x) \propto \prod_{j=1}^{\ell} \exp(g(x, y_{j-1}, y_j; \theta))$
- Combine LVMs and strong conditional independences [Gu et al., 2018, Ghazvininejad et al., 2019]
- Energy-based models [Song and Kingma, 2021]: regress to a score without factorisation and **normalise**

$$P_{Y|X}(y|x) = \frac{\exp(g(x, y; \theta))}{\sum_{y'} \exp(g(x, y'; \theta))}$$

Moving away from chain rule?

Strong conditional independences are often unrealistic.

Marginals and normalising constants are typically intractable to compute. This complicates learning, as the probability of observed data (necessary for training via MLE) isn't tractable to compute, and can complicate prediction (e.g., search and sampling are very hard in EBMs [Eikema et al., 2022]).

These topics are covered in detail in DL2.

Some background material

- Probabilistic graphical models [Koller and Friedman, 2009] (esp, part I on representation of probability distributions).
- Decision theory [Berger, 2013].
- On the origin of softmax: see Chapter 3 of Vlad Niculae's PhD thesis.⁶

Related courses

- DL4NLP (Christof Monz): state-of-the art architectures for most major sequence-to-sequence tasks.
- DL2 (Efstratios Gavves and Wilker Aziz): check it online.

⁶Learning Deep Models with Linguistically-Inspired Structure

Puzzle

Puzzle

Consider our typical sequence-to-sequence model, that is, a neural parameterisation of a chain rule factorisation of the joint distribution of our output random sequence Y given an outcome of an input random sequence $X = x$.

Suppose the output random sequence has length J . The probability of any outcome $\langle y_1, \dots, y_J \rangle$ is given by

$$P(Y = \langle y_1, \dots, y_J \rangle | X = x) = \prod_{j=1}^J P(Y_j = y_j | X = x, Y_{<j} = y_{<j}) \quad (12)$$

Suppose we obtain an output sequence but the k th step is missing. How can we generate outcomes for Y_k given the assignments of all other variables?

It is sufficient to solve this for an example: $X = x$, and $\langle Y_1 = \text{the}, Y_3 = \text{dog}, Y_4 = \text{EOS} \rangle$, with missing Y_2 .

Puzzle - Solution

Let's denote by O the set of output random variables we observe (that is, $\{Y_1, Y_3, Y_4\}$ in the example) and o their observed values (that is, $\{\text{the}, \text{dog}, \text{EOS}\}$), and by U the set of variables that are unobserved (that is $\{Y_2\}$ in the example). We want to express the probability that U takes on some value u (e.g., $Y_2 = w$ for any word w in the vocabulary, in one case w might be the word 'cute' for example) given the assignments of the observed variables O and $X = x$.

We start by application of the definition of conditional probability:

$$P(U = u|O = o, X = x) = \frac{P(O = o, U = u|X = x)}{P(O = o|X = x)} \quad (13)$$

Puzzle - Solution

Note that the numerator is exactly the joint distribution we have access to. In our example: $P(O = o, U = u | X = x) = \underbrace{P(Y_1 = \text{the}, Y_3 = \text{dog}, Y_4 = \text{EOS}, Y_2 = w | X = x)}_{O=o} \underbrace{|}_{U=u}$.

Note the the denominator is the marginal of the numerator, where we marginalise out all possible assignments of U . In our example, we would marginalise out all possibilities for the second token. If \mathcal{V} is the entire vocabulary, we would compute

$$\sum_{t \in \mathcal{V}} P(Y_1 = \text{the}, Y_2 = t, Y_3 = \text{dog}, Y_4 = \text{EOS} | X = x).$$

Suppose in general an output sequence has length J and the vocabulary has size V . What's the computational complexity of evaluating the conditional probability of an assignment of some Y_k given everything else?

Puzzle - Solution

Probabilities that condition on past context are simple because those are directly predicted by our NN (provided we have access to observations for all variables in the past). Probabilities that condition on future are much more difficult because we need to assess the joint probability for every possible assignment of the unobserved variable (in the denominator of conditional probability).

Each joint probability takes J calls to our NN $\mathbf{g}(\cdot; \theta)$ (one per token in the sequence). We have to perform this computation V times, once per possible value of Y_k . So the total computation takes time proportional to $\mathcal{O}(JV)$.

References

- Jasmijn Bastings, Ivan Titov, Wilker Aziz, Diego Marcheggiani, and Khalil Sima'an. Graph convolutional encoders for syntax-aware neural machine translation. In *Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing*, pages 1957–1967, Copenhagen, Denmark, September 2017. Association for Computational Linguistics. doi: 10.18653/v1/D17-1209. URL <https://aclanthology.org/D17-1209>.
- James O Berger. *Statistical decision theory and Bayesian analysis*. Springer Science & Business Media, 2013.

Léon Bottou and Yann L. Cun. Large scale online learning. In S. Thrun, L. K. Saul, and B. Schölkopf, editors, *Advances in Neural Information Processing Systems 16*, pages 217–224. MIT Press, 2004.

Nicola De Cao, Gautier Izacard, Sebastian Riedel, and Fabio Petroni. Autoregressive entity retrieval. In *International Conference on Learning Representations*, 2021. URL <https://openreview.net/forum?id=5k8F6UU39V>.

Bryan Eikema and Wilker Aziz. Auto-encoding variational neural machine translation. In *Proceedings of the 4th Workshop on Representation Learning for NLP (RepL4NLP-2019)*, pages 124–141, Florence, Italy, August 2019. Association for Computational Linguistics. doi: 10.18653/v1/W19-4315. URL <https://aclanthology.org/W19-4315>.

Bryan Eikema and Wilker Aziz. Is MAP decoding all you need? the inadequacy of the mode in neural machine translation. In *Proceedings of the 28th International Conference on Computational Linguistics*, pages 4506–4520, Barcelona, Spain (Online), December 2020. International Committee on Computational Linguistics. doi:

10.18653/v1/2020.coling-main.398. URL

<https://aclanthology.org/2020.coling-main.398>.

Bryan Eikema and Wilker Aziz. Sampling-based minimum Bayes risk decoding for neural machine translation. In *EMNLP*, 2022.

Bryan Eikema, Germán Kruszewski, Christopher R Dance, Hady Elsahar, and Marc Dymetman. An approximate sampler for energy-based models with divergence diagnostics. *Transactions on Machine Learning Research*, 2022. ISSN 2835-8856. URL <https://openreview.net/forum?id=VW4IrC0n0M>.

Jonas Gehring, Michael Auli, David Grangier, and Yann Dauphin.
A convolutional encoder model for neural machine translation.
In *Proceedings of the 55th Annual Meeting of the Association
for Computational Linguistics (Volume 1: Long Papers)*, pages
123–135, Vancouver, Canada, July 2017. Association for
Computational Linguistics. doi: 10.18653/v1/P17-1012. URL
<https://aclanthology.org/P17-1012>.

Marjan Ghazvininejad, Omer Levy, Yinhan Liu, and Luke
Zettlemoyer. Mask-predict: Parallel decoding of conditional
masked language models. In *Proceedings of the 2019
Conference on Empirical Methods in Natural Language
Processing and the 9th International Joint Conference on
Natural Language Processing (EMNLP-IJCNLP)*, pages
6112–6121, Hong Kong, China, November 2019. Association for

Computational Linguistics. doi: 10.18653/v1/D19-1633. URL
<https://aclanthology.org/D19-1633>.

Mario Julianelli, Joris Baan, Wilker Aziz, Raquel Fernández, and Barbara Plank. What comes next? evaluating uncertainty in neural text generators against human production variability. In *EMNLP*, 2023.

Jiatao Gu, James Bradbury, Caiming Xiong, Victor O.K. Li, and Richard Socher. Non-autoregressive neural machine translation. In *International Conference on Learning Representations*, 2018.
URL <https://openreview.net/forum?id=B1l8BtlCb>.

Ari Holtzman, Jan Buys, Li Du, Maxwell Forbes, and Yejin Choi. The curious case of neural text degeneration. *arXiv preprint arXiv:1904.09751*, 2019.

Martin Josifoski, Nicola De Cao, Maxime Peyrard, Fabio Petroni, and Robert West. GenIE: Generative information extraction. In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 4626–4643, Seattle, United States, July 2022. Association for Computational Linguistics. doi: 10.18653/v1/2022.naacl-main.342. URL <https://aclanthology.org/2022.naacl-main.342>.

Daphne Koller and Nir Friedman. *Probabilistic graphical models: principles and techniques*. MIT press, 2009.

Wouter Kool, Herke Van Hoof, and Max Welling. Stochastic beams and where to find them: The gumbel-top-k trick for sampling sequences without replacement. In *International Conference on Machine Learning*, pages 3499–3508. PMLR, 2019.

Guillaume Lample, Miguel Ballesteros, Sandeep Subramanian, Kazuya Kawakami, and Chris Dyer. Neural architectures for named entity recognition. In *Proceedings of the 2016 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 260–270, San Diego, California, June 2016. Association for Computational Linguistics. doi: 10.18653/v1/N16-1030. URL <https://aclanthology.org/N16-1030>.

Wang Ling, Chris Dyer, Alan W Black, Isabel Trancoso, Ramón Fermandez, Silvio Amir, Luís Marujo, and Tiago Luís. Finding function in form: Compositional character models for open vocabulary word representation. In *Proceedings of the 2015 Conference on Empirical Methods in Natural Language Processing*, pages 1520–1530, Lisbon, Portugal, September 2015. Association for Computational Linguistics. doi:

10.18653/v1/D15-1176. URL

<https://aclanthology.org/D15-1176>.

Xuezhe Ma and Eduard Hovy. End-to-end sequence labeling via bi-directional LSTM-CNNs-CRF. In *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1064–1074, Berlin, Germany, August 2016. Association for Computational Linguistics. doi: 10.18653/v1/P16-1101. URL

<https://aclanthology.org/P16-1101>.

Herbert Robbins and Sutton Monro. A stochastic approximation method. *Ann. Math. Statist.*, 22(3):400–407, 1951. doi: 10.1214/aoms/1177729586. URL

<http://dx.doi.org/10.1214/aoms/1177729586>.

Alexander M. Rush, Sumit Chopra, and Jason Weston. A neural attention model for abstractive sentence summarization. In *Proceedings of the 2015 Conference on Empirical Methods in Natural Language Processing*, pages 379–389, Lisbon, Portugal, September 2015. Association for Computational Linguistics. doi: 10.18653/v1/D15-1044. URL
<https://aclanthology.org/D15-1044>.

Abigail See, Peter J. Liu, and Christopher D. Manning. Get to the point: Summarization with pointer-generator networks. In Regina Barzilay and Min-Yen Kan, editors, *Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1073–1083, Vancouver, Canada, July 2017. Association for Computational Linguistics. doi: 10.18653/v1/P17-1099. URL
<https://aclanthology.org/P17-1099>.

Yang Song and Diederik P Kingma. How to train your energy-based models. *arXiv preprint arXiv:2101.03288*, 2021.

Ilya Sutskever, Oriol Vinyals, and Quoc V. V Le. Sequence to sequence learning with neural networks. In Z. Ghahramani, M. Welling, C. Cortes, N.D. Lawrence, and K.Q. Weinberger, editors, *NIPS, 2014*, pages 3104–3112. Montreal, Canada, 2014.

Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *NeurIPS*, pages 6000–6010, 2017. URL <http://papers.nips.cc/paper/7181-attention-is-all-you-need>.

Biao Zhang, Deyi Xiong, Jinsong Su, Hong Duan, and Min Zhang.
Variational neural machine translation. In *Proceedings of the
2016 Conference on Empirical Methods in Natural Language
Processing*, pages 521–530, Austin, Texas, November 2016.
Association for Computational Linguistics. doi:
10.18653/v1/D16-1050. URL
<https://aclanthology.org/D16-1050>.